

# Entering the Fray

## Chapel's Computer Language Benchmarks Game Entry

Brad Chamberlain, Ben Albrecht, Lydia Duncan, Ben Harshbarger,  
Elliot Ronaghan, Preston Sahabu, Mike Noakes, and Laura Delaney

CHIUW 2017, Orlando, FL

June 2, 2017

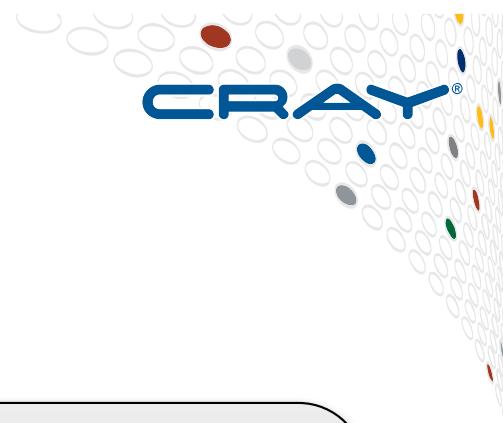


---

COMPUTE

| STORE

| ANALYZE



# Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.



## **CLBG: What it is**



- A suite of 13 “toy” benchmarks
    - exercise key features like...
      - ...memory management
      - ...tasking and synchronization
      - ...arbitrary-precision math
      - ...vectorization
      - ...strings and regular expressions
    - single-node
    - serial, vectorizable, or multicore parallel

# The Computer Language Benchmarks Game

## 64-bit quad core data set

Will your toy benchmark program be faster if you write it in a different programming language? It depends how you write it!

## Which programs are fast?

Which are succinct? Which are efficient?

<u>Ada</u>	<u>C</u>	<u>Chapel</u>	<u>C#</u>	<u>C++</u>	<u>Dart</u>
<u>Erlang</u>	<u>F#</u>	<u>Fortran</u>	<u>Go</u>	<u>Hack</u>	
<u>Haskell</u>	<u>Java</u>	<u>JavaScript</u>	<u>Lisp</u>	<u>Lua</u>	
<u>OCaml</u>	<u>Pascal</u>	<u>Perl</u>	<u>PHP</u>	<u>Python</u>	
<u>Racket</u>	<u>Ruby</u>	<u>JRuby</u>	<u>Rust</u>	<u>Smalltalk</u>	
	<u>Swift</u>	<u>TypeScript</u>			





## But wait...

- **This is IPDPS / HPC / Chapel...**

...do we really care about a single-node benchmark suite?

- **Yes:**

- success at the largest scales depends on good scalar performance
- despite its focus on large-scale systems, Chapel is also intended for productive programming on workstations
- several CLBG features match early user wishes
  - memory management
  - tasking and lightweight synchronization
  - arbitrary precision arithmetic
  - strings and regular expressions
  - vectorization
  - ...
- who doesn't enjoy a good game?



# CLBG: What it is

- A suite of 13 “toy” benchmarks

- exercise key features like...
  - ...memory management
  - ...tasking and synchronization
  - ...arbitrary-precision math
  - ...vectorization
  - ...strings and regular expressions
- single-node
- serial, vectorizable, or multicore parallel

- Imagine a 3D ragged matrix:

- with 13 benchmarks
  - x ~28 languages
  - x as many impls as are interesting
- each entry contains:
  - source code
  - performance statistics
  - “code size”

## The Computer Language Benchmarks Game

### 64-bit quad core data set

Will your toy benchmark program be faster if you write it in a different programming language? It depends how you write it!

### Which programs are fast?

Which are succinct? Which are efficient?

<u>Ada</u>	<u>C</u>	<u>Chapel</u>	<u>C#</u>	<u>C++</u>	<u>Dart</u>
<u>Erlang</u>	<u>F#</u>	<u>Fortran</u>	<u>Go</u>	<u>Hack</u>	
<u>Haskell</u>	<u>Java</u>	<u>JavaScript</u>	<u>Lisp</u>	<u>Lua</u>	
<u>OCaml</u>	<u>Pascal</u>	<u>Perl</u>	<u>PHP</u>	<u>Python</u>	
<u>Racket</u>	<u>Ruby</u>	<u>JRuby</u>	<u>Rust</u>	<u>Smalltalk</u>	
		<u>Swift</u>	<u>TypeScript</u>		
{ for <u>researchers</u> }			<u>fast-faster-fastest</u>		
<u>stories</u>					





# Timeline

**Feb 2016:** Inquired about adding Chapel to the contest

**Apr 2016:** Chapel entries began to be accepted

**Our approach:**

- Submit codes that strive for performance without sacrificing elegance
- Submit codes that would serve as good models for learning the benchmark

**May 2016:** First program accepted

**Sept 2016:** Last program accepted, Chapel added to the site



COMPUTE

|

STORE

|

ANALYZE

# CLBG: What it is

- A suite of 13 “toy” benchmarks
  - exercise key features like...
    - ...memory management
    - ...tasking and synchronization
    - ...arbitrary-precision math
    - ...vectorization
    - ...strings and regular expressions
  - single-node
  - serial, vectorizable, or multicore parallel

- Imagine a 3D ragged matrix:

- with 13 benchmarks
  - x ~200 lines of C code
  - x a few lines of Chapel
- each entry contains:
  - source code
  - performance information
  - “code size”

**Chapel added to site in  
September 2016**

## The Computer Language Benchmarks Game

### 64-bit quad core data set

Will your toy benchmark program be faster if you write it in a different programming language? It depends how you write it!

### Which programs are fast?

Which are succinct? Which are efficient?

<u>Ada</u>	<u>C</u>	<u>Chapel</u>	<u>C#</u>	<u>C++</u>	<u>Dart</u>
<u>Erlang</u>	<u>F#</u>	<u>Fortran</u>	<u>Go</u>	<u>Hack</u>	
<u>Haskell</u>	<u>Java</u>	<u>JavaScript</u>	<u>Lisp</u>	<u>Lua</u>	
<u>OCaml</u>	<u>Pascal</u>	<u>Perl</u>	<u>PHP</u>	<u>Python</u>	
<u>Racket</u>	<u>Ruby</u>	<u>JRuby</u>	<u>Rust</u>	<u>Smalltalk</u>	
	<u>Swift</u>	<u>TypeScript</u>			
{ for <u>researchers</u> }			<u>fast-faster-fastest stories</u>		



# CLBG: What it is

- A suite of 13 “toy” benchmarks

- exercise key features like...
  - ...memory management
  - ...tasking and synchronization
  - ...arbitrary-precision math
  - ...vectorization
  - ...strings and regular expressions
- single-node
- serial, vectorizable, or multicore parallel

- Imagine a 3D ragged matrix:

- with 13 benchmarks
  - x ~28 languages
  - x as many impls as are interesting
- each entry contains:
  - source code
  - performance information
  - “code size”

## The Computer Language Benchmarks Game

### 64-bit quad core data set

Will your toy benchmark program be faster if you write it in a different programming language? It depends how you write it!


### Which programs are fast?

Which are succinct? Which are efficient?


<u>Ada</u>	<u>C</u>	<u>Chapel</u>	<u>C#</u>	<u>C++</u>	<u>Dart</u>
<u>Erlang</u>	<u>F#</u>	<u>Fortran</u>	<u>Go</u>	<u>Hack</u>	
<u>Haskell</u>	<u>Java</u>	<u>JavaScript</u>	<u>Lisp</u>	<u>Lua</u>	
<u>OCaml</u>	<u>Pascal</u>	<u>Perl</u>	<u>PHP</u>	<u>Python</u>	
<u>Racket</u>	<u>Ruby</u>	<u>JRuby</u>	<u>Rust</u>	<u>Smalltalk</u>	
		<u>Swift</u>	<u>TypeScript</u>		
{ for <u>researchers</u> }			<u>fast-faster-fastest</u>		
			<u>stories</u>		




# CLBG: Fast-faster-fastest graph (Sep 2016)




Site summary: relative performance (sorted by geometric mean)



# CLBG: Fast-faster-fastest graph (May 2017)



## Site summary: relative performance (sorted by geometric mean)



# CLBG: Viewing per-benchmark results

Can sort results by execution time, code size, memory or CPU use:

The Computer Language Benchmarks Game						
chameneos-redux						
<u>description</u>						
program source code, command-line and measurements						
x	source	secs	mem	gz	cpu	cpu load
1.0	<a href="#">C gcc #5</a>	<b>0.60</b>	820	2863	2.37	100% 100% 98% 100%
1.2	<a href="#">C++ g++ #5</a>	<b>0.70</b>	3,356	1994	2.65	100% 100% 91% 92%
1.7	<a href="#">Lisp SBCL #3</a>	<b>1.01</b>	55,604	2907	3.93	97% 96% 99% 99%
2.3	<a href="#">Chapel #2</a>	<b>1.39</b>	76,564	1210	5.43	99% 99% 98% 99%
3.3	<a href="#">Rust #2</a>	<b>2.01</b>	56,936	2882	7.81	97% 98% 98% 98%
5.6	<a href="#">C++ g++ #2</a>	3.40	1,880	2016	11.88	100% 51% 100% 100%
6.8	<a href="#">Chapel</a>	4.09	66,584	1199	16.25	100% 100% 100% 100%
8.0	<a href="#">Java #4</a>	<b>4.82</b>	37,132	1607	16.73	98% 98% 54% 99%
8.5	<a href="#">Haskell GHC</a>	<b>5.15</b>	8,596	989	9.26	79% 100% 2% 2%
10	<a href="#">Java</a>	6.13	53,760	1770	8.78	42% 45% 41% 16%
10	<a href="#">Haskell GHC #4</a>	6.34	6,908	989	12.67	99% 100% 2% 1%
11	<a href="#">C# .NET Core</a>	<b>6.59</b>	86,076	1400	22.96	99% 82% 78% 91%
11	<a href="#">Go</a>	<b>6.90</b>	832	1167	24.19	100% 96% 56% 100%
13	<a href="#">Go #2</a>	7.59	1,384	1408	27.65	91% 99% 99% 78%
13	<a href="#">Java #3</a>	7.94	53,232	1267	26.86	54% 96% 98% 94%

The Computer Language Benchmarks Game						
chameneos-redux						
<u>description</u>						
program source code, command-line and measurements						
x	source	secs	mem	gz	cpu	cpu load
1.0	<a href="#">Erlang</a>	58.90	28,668	<b>734</b>	131.19	62% 60% 51% 53%
1.0	<a href="#">Erlang HiPE</a>	59.39	25,784	<b>734</b>	131.58	60% 56% 56% 54%
1.1	<a href="#">Perl #4</a>	5 min	14,084	<b>785</b>	7 min	40% 40% 29% 28%
1.1	<a href="#">Racket</a>	5 min	132,120	<b>791</b>	5 min	1% 0% 0% 100%
1.1	<a href="#">Racket #2</a>	175.88	116,488	842	175.78	100% 1% 1% 0%
1.2	<a href="#">Python 3 #2</a>	236.84	7,908	<b>866</b>	5 min	24% 48% 27% 45%
1.3	<a href="#">Ruby</a>	90.52	9,396	<b>920</b>	137.53	35% 35% 35% 34%
1.3	<a href="#">Ruby JRuby</a>	48.78	628,968	<b>928</b>	112.15	65% 60% 49% 58%
1.3	<a href="#">Go #5</a>	11.05	832	<b>957</b>	32.48	75% 74% 75% 73%
1.3	<a href="#">Haskell GHC #4</a>	6.34	6,908	<b>989</b>	12.67	99% 100% 2% 1%
1.3	<a href="#">Haskell GHC</a>	5.15	8,596	989	9.26	79% 100% 2% 2%
1.6	<a href="#">OCaml #3</a>					32% 38% 37% 39%
1.6	<a href="#">Go</a>					100% 96% 56% 100%
1.6	<a href="#">Chapel</a>					0% 100% 100% 100%
1.6	<a href="#">Chapel #2</a>					99% 99% 98% 99%

gz == code size metric  
strip comments and extra whitespace, then gzip



COMPUTE

STORE

ANALYZE



# Timeline

**Feb 2016: Inquired about adding Chapel to the contest**

**Apr 2016: Chapel entries began to be accepted**

**Our approach:**

- Submit codes that strive for performance without sacrificing elegance
- Submit codes that would serve as good models for learning the benchmark

**May 2016: First program accepted**

**Sept 2016: Last program accepted, Chapel added to the site**

**Oct 2016: Upgraded to 1.14**




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# CLBG: Improvements due to 1.14

## 1.14 improved many benchmarks with no code changes:


- **thread-ring:** benefitted from qthread sync variable improvements
  - climbed ~16 slots  $\Rightarrow$  5<sup>th</sup> fastest after Haskell, Go, F#, Scala
  - 1<sup>st</sup> most compact code followed by Ruby, Racket, Erlang, Ocaml, Python
- specifically, Chapel 1.14...
  - ...extended Qthreads sync vars to handle all Chapel operations
  - ...mapped Chapel sync vars directly to Qthreads sync vars (for simple types)



# CLBG: Improvements due to 1.14

## 1.14 improved many benchmarks with no code changes:

- **fannkuch-redux**: benefitted from optimized array accesses
  - climbed from ~#22 to #6 in performance
  - ~1.5–2x more compact than most other top entries
- specifically, Chapel 1.14...
  - ...optimized an unnecessary multiply out of typical array accesses



- this helped several other performance benchmarks as well
- Chapel 1.15 made this optimization more precise and robust





# CLBG: Improvements due to 1.14

## 1.14 improved many benchmarks with no code changes:

- **chameneos-redux**: benefitted from tasking improvements
  - climbed from ~#11 to #8 in terms of performance
- **binary-trees**: benefitted from jemalloc improvements
  - climbed ~2 performance slots as a result
  - still ~5x off from top entries which use explicit memory pools
- **n-body**: saw marginal improvements, but climbed ~17 slots
- **regex-dna, revcomp**: saw marginal improvements, climbed ~3 slots
- **meteor**: saw marginal improvements, climbed ~1 slot



COMPUTE

|

STORE

|

ANALYZE



# Chapel CLBG Standings (Oct 17<sup>th</sup>)

- 8 / 13 programs in top-20 fastest:
  - one #1 fastest:  
**pidigits**
  - 2 others in the top-5 fastest:  
**meteor-contest**  
**thread-ring**
  - 2 others in the top-10 fastest:  
**chameneos-redux**  
**fannkuch-redux**
  - 3 others in the top-20 fastest:  
**binary-trees**  
**n-body**  
**spectral-norm**
- 8 / 13 programs in top-20 smallest:
  - two #1 smallest:  
**n-body**  
**thread-ring**
  - 2 others in the top-5 smallest:  
**pidigits**  
**spectral-norm**
  - 4 others in the top-20 smallest:  
**chameneos-redux**  
**mandelbrot**  
**meteor-contest**  
**regex-dna**





# Chapel CLBG Standings (Apr 20<sup>th</sup>)

- 12 /13 programs in top-20 fastest:
  - one #1 fastest:  
**pidigits**
  - 3 others in the top-5 fastest:  
**chameneos-redux**  
**meteor-contest**  
**thread-ring**
  - 3 others in the top-10 fastest:  
**fannkuch-redux**  
**fasta**  
**mandelbrot**
  - 5 others in the top-20 fastest:  
**binary-trees**  
**k-nucleotide**  
**n-body**  
**regex-redux**  
**spectral-norm**
- 8 / 13 programs in top-20 smallest:
  - two #1 smallest:  
**n-body**  
**thread-ring**
  - 2 others in the top-5 smallest:  
**pidigits**  
**spectral-norm**
  - 1 other in the top-10 smallest:  
**regex-redux**
  - 3 others in the top-20 smallest:  
**chameneos-redux**  
**mandelbrot**  
**meteor-contest**





# Timeline

**Feb 2016: Inquired about adding Chapel to the contest**

**Apr 2016: Chapel entries began to be accepted**

**Our approach:**

- Submit codes that strive for performance without sacrificing elegance
- Submit codes that would serve as good models for learning the benchmark

**May 2016: First program accepted**

**Sept 2016: Last program accepted, Chapel added to the site**

**Oct 2016: Upgraded to 1.14**

**ongoing: Improved programs themselves in spare time**

**Apr 2017: Upgraded to 1.15**





# What's new with the CLBG since then?

- Two programs changed their official definitions:

## binary-trees:

- improved checksum to avoid false positives at 1/2, 1/4, 1/8 the memory
- eliminated per-node data field
- changed what trees are allocated and freed, slightly
- increased the problem size

## regex:

- changed the regular expression used
- renamed the test to regex-redux
- several versions are not currently passing due to these changes
  - our current standings may be due in part to this






# What's new with the Chapel CLBG entries?

- We've submitted some new versions:


binary-trees: used an initializer rather than a factory type method



# What's new with the Chapel CLBG entries?

- We've submitted some new versions:

binary-trees: used an initializer rather than a factory type method  
chameneos-redux: increased parallelism and tuned a spin-wait



# What's new with the Chapel CLBG entries?


- We've submitted some new versions:

**binary-trees**: used an initializer rather than a factory type method

**chameneos-redux**: increased parallelism and tuned a spin-wait

**fasta**: implemented a parallel version and tuned for clarity and speed

- also, changed some 'var' declarations due to const-checking improvements



# What's new with the Chapel CLBG entries?

- We've submitted some new versions:


**binary-trees**: used an initializer rather than a factory type method

**chameneos-redux**: increased parallelism and tuned a spin-wait

**fasta**: implemented a parallel version and tuned for clarity and speed

- also, changed some 'var' declarations due to const-checking improvements

**mandelbrot**: accelerated by hoisting values and using tuples of values





# What's new with the Chapel CLBG entries?

- **We've submitted some new versions:**

**binary-trees**: used an initializer rather than a factory type method

**chameneos-redux**: increased parallelism and tuned a spin-wait

**fasta**: implemented a parallel version and tuned for clarity and speed

- also, changed some 'var' declarations due to const-checking improvements

**mandelbrot**: accelerated by hoisting values and using tuples of values

**meteor-fast**: fixed a race condition caused by array memory changes

- textbook example of an array being used by a 'begin' task

**pidigits**: submitted a version that uses 'bigint's

- currently the #1 fastest version, and also quite elegant

- **Note that some of these changes followed the 1.15 release**

- As such, not all are found in examples/benchmarks/shootout/ for 1.15





# CLBG: Comparing Pairs of Languages

Can also compare languages pair-wise (for performance only):

## The Computer Language Benchmarks Game

Chapel programs versus Go  
all other Chapel programs & measurements

by benchmark task performance

### regex-redux

source	secs	mem	gz	cpu	cpu load
Chapel	<b>10.02</b>	1,022,052	477	19.68	99% 72% 14% 12%
Go	29.51	352,804	798	61.51	77% 49% 43% 40%

### binary-trees

source	secs	mem	gz	cpu	cpu load
Chapel	<b>14.32</b>	324,660	484	44.15	100% 58% 78% 75%
Go	34.77	269,068	654	132.04	95% 97% 95% 95%

### fannkuch-redux

source	secs	mem	gz	cpu	cpu load
Chapel	<b>11.38</b>	46.056	728	45.18	100% 99% 99% 100%

*Happily, all the data is open!*



COMPUTE

STORE

ANALYZE

Copyright 2017 Cray Inc.



# CLBG Scatter Plots

- **The following graphs use the CLBG's normalized ratios**
  - Graphs were created using April 20<sup>th</sup> data (current at time of creation)
    - things have continued to be in flux again since that date...



COMPUTE


|

STORE

|

ANALYZE

# Chapel entries (Apr 2017)



COMPUTE


|

STORE

|

ANALYZE

# Chapel entries (Apr 2017, noting outliers)



# Chapel vs. 9 other languages




COMPUTE | STORE | ANALYZE

# Chapel vs. 9 other languages (zoomed out)



COMPUTE | STORE | ANALYZE

# Cross-Language Summary





COMPUTE

STORE

ANALYZE

# Cross-Language Summary (no Python)



COMPUTE

STORE

ANALYZE

Copyright 2017 Cray Inc.



# Chapel CLBG Standings as of Apr 20<sup>th</sup>

- 12 /13 programs in top-20 fastest:
  - one #1 fastest:  
**pidigits**
  - 3 others in the top-5 fastest:  
**chameneos-redux**  
**meteor-contest**  
**thread-ring**
  - 3 others in the top-10 fastest:  
**fannkuch-redux**  
**fasta**  
**mandelbrot**
  - 5 others in the top-20 fastest:  
**binary-trees**  
**k-nucleotide**  
**n-body**  
**regex-redux**  
**spectral-norm**
- 8 / 13 programs in top-20 smallest:
  - two #1 smallest:  
**n-body**  
**thread-ring**
  - 2 others in the top-5 smallest:  
**pidigits**  
**spectral-norm**
  - 1 other in the top-10 smallest:  
**regex-redux**
  - 3 others in the top-20 smallest:  
**chameneos-redux**  
**mandelbrot**  
**meteor-contest**



# Comparing Chapel vs. C Chameneos

Can also browse program source code (but this requires actual thought):

```

proc main() {
    printColorEquations();

    const group1 = [i in 1..popSize1] new Chameneos(i, ((i-1)%3):Color);
    const group2 = [i in 1..popSize2] new Chameneos(i, colors10[i]);

    cobegin {
        holdMeetings(group1, n);
        holdMeetings(group2, n);
    }

    print(group1);
    print(group2);

    for c in group1 do delete c;
    for c in group2 do delete c;
}

// Print the results of getNewColor() for all color pairs.
// proc printColorEquations() {
//     for c1 in Color do
//         for c2 in Color do
//             writeln(c1, " + ", c2, " -> ", getNewColor(c1, c2));
//     writeln();
// }

// Hold meetings among the population by creating a shared meeting
// place, and then creating per-chameneos tasks to have meetings.
// proc holdMeetings(population, numMeetings) {
const place = new MeetingPlace(numMeetings);

coforall c in population do          // create a task per chameneos
    c.haveMeetings(place, population);

    delete place;
}

```

excerpt from 1210 gz 4<sup>th</sup>-place Chapel entry

```

void get_affinity(int* is_smp, cpu_set_t* affinity1, cpu_set_t* affinity2)
{
    cpu_set_t
    FILE*
    char
    char const*
    int
    int
    int
    int
    int
    size_t
    size_t

    active_cpus;
    f;
    buf [2048];
    pos;
    cpu_idx;
    physical_id;
    core_id;
    cpu_cores;
    apic_id;
    cpu_count;
    i;

    char const*
    size_t
    char const*
    size_t
    char const*
    size_t
    char const*
    size_t

    processor_str      = "processor";
    processor_str_len = strlen(processor_str);
    physical_id_str   = "physical id";
    physical_id_str_len = strlen(physical_id_str);
    core_id_str        = "core id";
    core_id_str_len   = strlen(core_id_str);
    cpu_cores_str      = "cpu cores";
    cpu_cores_str_len = strlen(cpu_cores_str);

    CPU_ZERO(&active_cpus);
    sched_getaffinity(0, sizeof(active_cpus), &active_cpus);
    cpu_count = 0;
    for (i = 0; i != CPU_SETSIZE; i += 1)
    {
        if (CPU_ISSET(i, &active_cpus))
        {
            cpu_count += 1;
        }
    }

    if (cpu_count == 1)
    {
        is_smp[0] = 0;
        return;
    }

    is_smp[0] = 1;
    CPU_ZERO(affinity1);
}

```

excerpt from 2863 gz 1<sup>st</sup>-place C gcc entry



# Comparing Chapel vs. C Chameneos

Can also browse program source code (but this requires actual thought):

```

proc main() {
    printColorEquations();

    const group1 = [i in 1..popSize1] new Ch
    const group2 = [i in 1..popSize2] new Ch

    cobegin {
        holdMeetings(group1, n);
        holdMeetings(group2, n);
    }

    print(group1);
    print(group2);

    for c in group1 do delete c;
    for c in group2 do delete c;
}

// Print the results of getNewColor() for all color pairs.
// 
proc printColorEquations() {
    for c1 in Color do
        for c2 in Color do
            writeln(c1, " + ", c2, " -> ", getNewColor(c1, c2));
            writeln();
}

// Hold meetings among the population by creating a
// place, and then creating per-chameneos tasks to
// handle each meeting
// 
proc holdMeetings(population, numMeetings) {
    const place = new MeetingPlace(numMeetings);

    coforall c in population do          // create a
        c.haveMeetings(place, population);

    delete place;
}

```

```

cobegin {
    holdMeetings(group1, n);
    holdMeetings(group2, n);
}

```

```

void get_affinity(int* is_smp, cpu_set_t* affinity1, cpu_set_t* affinity2)
{
    int i, j, k, l, m, n, p, q, r, s, t, u, v, w, x, y, z;
    int live_cpus;
    int max_cpus = 2048;
    int cpus[2048];
    int cpus_idx;
    int sical_id;
    int re_id;
    int cores;
    int c_id;
    int count;

    char const* processor_str = "processor";
    size_t processor_str_len = strlen(processor_str);
    char const* physical_id_str = "physical id";
    size_t physical_id_str_len = strlen(physical_id_str);
    char const* core_id_str = "core id";
    size_t core_id_str_len = strlen(core_id_str);
    char const* cpu_cores_str = "cpu cores";
    size_t cpu_cores_str_len = strlen(cpu_cores_str);

proc holdMeetings(population, numMeetings) {
    const place = new MeetingPlace(numMeetings);

    coforall c in population do          // create a
        c.haveMeetings(place, population);

    delete place;
}

```

excerpt from 1210 gz 4<sup>th</sup>-place Chapel entry

excerpt from 2863 gz 1<sup>st</sup>-place C gcc entry



# Comparing Chapel vs. C Chameneos

Can also browse program source code (but this requires actual thought):

```

proc main() {
    printColorEquations();

    const group1 = [i in 1..popSize1] new Chameneos(i, ((i-1)%3):Color);

    char const* core_id_str =;
    size_t core_id_str_len =;
    char const* cpu_cores_str =;
    size_t cpu_cores_str_len =;

    CPU_ZERO(&active_cpus);
    sched_getaffinity(0, sizeof(active_cpus), &active_cpus);
    cpu_count = 0;
    for (i = 0; i != CPU_SETSIZE; i += 1)
    {
        if (CPU_ISSET(i, &active_cpus))
        {
            cpu_count += 1;
        }
    }

    if (cpu_count == 1)
    {
        is_smp[0] = 0;
        return;
    }

    const place = new MeetingPlace(numMeetings);

    coforall c in population do // create e_task_per_chameneos
        c.haveMeetings(place, population);

    delete place;
}

```

excerpt from 1210 gz 4<sup>th</sup>-place Chapel entry

```

void get_affinity(int* is_smp, cpu_set_t* affinity1, cpu_set_t* affinity2)
{
    cpu_set_t
    FILE*
    char
    char const*
    int
    int
    int
    int
    size_t
    size_t

    active_cpus;
    f;
    buf [2048];
    pos;
    cpu_idx;
    physical_id;
    core_id;
    cpu_cores;
    apic_id;
    cpu_count;
    i;

    char const* processor_str = "processor";
    size_t processor_str_len = strlen(processor_str);
    physical_id_str = "physical id";
    physical_id_str_len = strlen(physical_id_str);
    core_id_str = "core id";
    core_id_str_len = strlen(core_id_str);
    cpu_cores_str = "cpu cores";
    cpu_cores_str_len = strlen(cpu_cores_str);

    CPU_ZERO(&active_cpus);
    sched_getaffinity(0, sizeof(active_cpus), &active_cpus);
    cpu_count = 0;
    for (i = 0; i != CPU_SETSIZE; i += 1)
    {
        if (CPU_ISSET(i, &active_cpus))
        {
            cpu_count += 1;
        }
    }

    if (cpu_count == 1)
    {
        is_smp[0] = 0;
        return;
    }

    is_smp[0] = 1;
    CPU_ZERO(affinity1);
}

```

excerpt from 2863 gz 1<sup>st</sup>-place C gcc entry





# Chapel CLBG Standings as of Apr 20<sup>th</sup>

- 12 /13 programs in top-20 fastest:
  - one #1 fastest:  
**pidigits**
  - 3 others in the top-5 fastest:  
**chameneos-redux**  
**meteor-contest**  
**thread-ring**
  - 3 others in the top-10 fastest:  
**fannkuch-redux**  
**fasta**  
**mandelbrot**
  - 5 others in the top-20 fastest:  
**binary-trees**  
**k-nucleotide**  
**n-body**  
**regex-redux**  
**spectral-norm**
- 8 / 13 programs in top-20 smallest:
  - two #1 smallest:  
**n-body**  
**thread-ring**
  - 2 others in the top-5 smallest:  
**pidigits**  
**spectral-norm**
  - 1 other in the top-10 smallest:  
**regex-redux**
  - 3 others in the top-20 smallest:  
**chameneos-redux**  
**mandelbrot**  
**meteor-contest**



# Comparing Chapel vs. C pidigits

```

use BigInteger;

config const n = 50;           // Compute n digits of pi, 50 by default

proc main() {
    param digitsPerLine = 10;

    // Generate n digits, printing them in groups of digitsPerLine
    for (d, i) in genDigits(n) {
        write(d);
        if i % digitsPerLine == 0 then
            writeln("t:", i);
    }

    // Pad out any trailing digits for the final line
    if n % digitsPerLine then
        writeln(" " * (digitsPerLine - n % digitsPerLine), "\t:", n);
    }

iter genDigits(numDigits) {
    var numer, denom: bigint = 1,
        accum, tmp1, tmp2: bigint;

    var i, k = 1;
    while i <= numDigits {
        nextTerm(k);
        k += 1;
        if numer <= accum {
            const d = extractDigit(3);
            if d == extractDigit(4) {
                yield(d, i);
                eliminateDigit(d);
                i += 1;
            }
        }
    }

    proc nextTerm(k) {
        const k2 = 2 * k + 1;

        accum.addmul(numer, 2);
        accum *= k2;
        denom *= k2;
        numer *= k;
    }

    proc extractDigit(nth) {
        tmp1.mul(numer, nth);
        tmp2.add(tmp1, accum);
        tmp1.div_q(tmp2, denom);

        return tmp1: int;
    }

    proc eliminateDigit(d) {
        accum.submul(denom, d);
        accum *= 10;
        numer *= 10;
    }
}

```

```

#include <stdio.h>
#include <stdlib.h>
#include <gmp.h>

mpz_t tmp1, tmp2, acc, den, num;
typedef unsigned int ui;

ui extract_digit(ui nth) {
    // juggling between tmp1 and tmp2, so GMP won't have to use temp buffers
    mpz_mul_ui(tmp1, num, nth);
    mpz_add(tmp2, tmp1, acc);
    mpz_tdiv_q(tmp1, tmp2, den);

    return mpz_get_ui(tmp1);
}

void eliminate_digit(ui d) {
    mpz_submul_ui(acc, den, d);
    mpz_mul_ui(acc, acc, 10);
    mpz_mul_ui(num, num, 10);
}

void next_term(ui k) {
    ui k2 = k * 2U + 1U;

    mpz_addmul_ui(acc, num, 2U);
    mpz_mul_ui(acc, acc, k2);
    mpz_mul_ui(den, den, k2);
    mpz_mul_ui(num, num, k);
}

int main(int argc, char **argv) {
    ui d, k, i;
    int n = atoi(argv[1]);

    mpz_init(tmp1);
    mpz_init(tmp2);

    mpz_init_set_ui(acc, 0);
    mpz_init_set_ui(den, 1);
    mpz_init_set_ui(num, 1);

    for (i = k = 0; i < n;) {
        next_term(++k);
        if (mpz_cmp(num, acc) > 0)
            continue;

        d = extract_digit(3);
        if (d != extract_digit(4))
            continue;

        putchar('0' + d);
        if (++i % 10 == 0)
            printf("\t:%u\n", i);
        eliminate_digit(d);
    }

    return 0;
}

```

excerpt from 423 gz 1<sup>st</sup>-place Chapel entry

excerpt from 448 gz 4<sup>th</sup>-place C gcc entry



# Comparing Chapel vs. C pidigits

```

use BigInteger;

config const n = 50;           // Compute n digits of pi, 50 by default

proc main() {
    param digitsPerLine = 10;

    // Generate n digits, printing them in groups of digitsPerLine
    for(d, i) in genDigits(n) {
        write(d);
        if i % digitsPerLine == 0 then
            writeln("t:", i);
    }

    // Pad out any trailing digits for the final line
    if n % digitsPerLine then
        writeln(" " * (digitsPerLine - n % digitsPerLine), "\t:""\n");
    }

iter genDigits(numDigits) {
    var numer, denom: bigint = 1,
        accum, tmp1, tmp2: bigint;

    var i, k = 1;
    while i <= numDigits {
        nextTerm(k);
        k += 1;
        if numer <= accum {
            const d = extractDigit(3);
            if d == extractDigit(4)
                yield(d, i);
            eliminateDigit(d);
            i += 1;
        }
    }
}

proc nextTerm(k) {
    const k2 = 2 * k + 1;

    accum.addmul(numer, 2);
    accum *= k2;
    denom *= k2;
    numer *= k;
}

proc extractDigit(nth) {
    tmp1.mul(numer, nth);
    tmp2.add(tmp1, accum);
    tmp1.div_q(tmp2, denom);

    return tmp1: int;
}

proc eliminateDigit(d) {
    accum.submul(denom, d);
    accum *= 10;
    numer *= 10;
}

```

```

#include <stdio.h>
#include <stdlib.h>
#include <gmp.h>

mpz_t tmp1, tmp2, acc, den, num;
typedef unsigned int ui;

proc nextTerm(k) {
    const k2 = 2 * k + 1;

    accum.addmul(numer, 2);
    accum *= k2;
    denom *= k2;
    numer *= k;
}

int main(int argc, char **argv) {
    ui d, k, i;
    int n = atoi(argv[1]);

    mpz_init(tmp1);
    mpz_init(tmp2);

    mpz_init_set_ui(acc, 0);
    mpz_init_set_ui(den, 1);
    mpz_init_set_ui(num, 1);

    for (i = k = 0; i < n;) {
        nextTerm(++k);
        if (mpz_cmp(num, acc) > 0)
            continue;

        d = extract_digit(3);
        if (d != extract_digit(4))
            continue;

        putchar('0' + d);
        if (++i % 10 == 0)
            printf("\t:%u\n", i);
        eliminate_digit(d);
    }

    return 0;
}

```

excerpt from 423.gz 1<sup>st</sup>-place Chapel entry

excerpt from 448.gz 4<sup>th</sup>-place C gcc entry



# Comparing Chapel vs. C pidigits

```
use BigInteger;

config const n = 50;           // Compute n digits of pi, 50 by default

proc main() {
    param digitsPerLine = 10;

    // Generate n digits, printing them in groups of digitsPerLine
    for (d, i) in genDigits(n) {
        write(d);
        if i % digitsPerLine == 0 then
            writeln("\t:", i);
    }

    // Pad out any trailing digits for the final line
    if n % digitsPerLine then
        writeln(" " * (digitsPerLine - n % digitsPerLine), "\t:", n);
}

iter genDigits(numDigits) {
    var numer, denom: bigint = 1,
        accum, tmp1, tmp2: bigint;
    ... Ver. i, j, k ...
}
```

```
void next_term(ui k) {
    ui k2 = k * 2U + 1U;

    mpz_addmul_ui(acc, num, 2U);
    mpz_mul_ui(acc, acc, k2);
    mpz_mul_ui(den, den, k2);
    mpz_mul_ui(num, num, k);
}
```

```
tmp1.div_q(tmp2, denom);

return tmp1: int;
}

proc eliminateDigit(d) {
    accum.submul(denom, d);
    accum *= 10;
    numer *= 10;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <gmp.h>

mpz_t tmp1, tmp2, acc, den, num;
typedef unsigned int ui;

ui extract_digit(ui nth) {
    // juggling between tmp1 and tmp2, so GMP won't have to use temp buffers
    mpz_mul_ui(tmp1, num, nth);
    mpz_add(tmp2, tmp1, acc);
    mpz_tdiv_q(tmp1, tmp2, den);

    return mpz_get_ui(tmp1);
}

void eliminate_digit(ui d) {
    mpz_submul_ui(acc, den, d);
    mpz_mul_ui(acc, acc, 10);
    mpz_mul_ui(num, num, 10);
}

void next_term(ui k) {
    ui k2 = k * 2U + 1U;

    mpz_addmul_ui(acc, num, 2U);
    mpz_mul_ui(acc, acc, k2);
    mpz_mul_ui(den, den, k2);
    mpz_mul_ui(num, num, k);
}

int main(int argc, char **argv) {
    ui d, k, i;
    int n = atoi(argv[1]);

    mpz_init(tmp1);
    mpz_init(tmp2);

    mpz_init_set_ui(acc, 0);
    mpz_init_set_ui(den, 1);
    mpz_init_set_ui(num, 1);

    for (i = k = 0; i < n;) {
        next_term(++k);
        if (mpz_cmp(num, acc) > 0)
            continue;

        d = extract_digit(3);
        if (d != extract_digit(4))
            continue;

        putchar('0' + d);
        if (++i % 10 == 0)
            printf("\t:%u\n", i);
        eliminate_digit(d);
    }

    return 0;
}
```

excerpt from 423 gz 1<sup>st</sup>-place Chapel entry

excerpt from 448 gz 4<sup>th</sup>-place C gcc entry





# CLBG: Next Steps

- Additional Performance Improvements

- Improve vectorization support
- Optimize idioms used by string-related benchmarks
  - strings, associative domains/arrays, byte arrays
- Support memory pools?

- How to shine a light on these qualitative comparisons?
  - Chapel blog articles?






# CLBG: Next Major Steps

- How can we create a similar competition within HPC?  
(where “we” == “the HPC community”, not Chapel)
  - multi-language
  - ongoing
  - open
  - addictive
- Intel Parallel Research Kernels (PRK) as a possible basis
  - My EMBRACE talk this morning has related thoughts



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.




# Questions?



---

COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.




# CLBG Scatter Plots



COMPUTE | STORE | ANALYZE


Copyright 2016 Cray Inc.

# Chapel vs. C




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# Chapel vs. C (zoomed out)




COMPUTE

|


STORE

|

ANALYZE




# Chapel vs. C++



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. C++ (zoomed out)




COMPUTE

STORE

ANALYZE

# Chapel vs. Fortran




COMPUTE

STORE

ANALYZE


# Chapel vs. Fortran (zoomed out)



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. Go



COMPUTE


|

STORE

|

ANALYZE


# Chapel vs. Go (zoomed out)



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.


# Chapel vs. Rust




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. Rust (zoomed out)




# Chapel vs. Swift




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. Swift (zoomed out)




# Chapel vs. Java



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.


# Chapel vs. Java (zoomed out)



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. Scala




COMPUTE


STORE

ANALYZE


# Chapel vs. Scala (zoomed out)



# Chapel vs. Python



# Chapel vs. Python (zoomed out)



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# Legal Disclaimer

*Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.*

*Cray Inc. may make changes to specifications and product descriptions at any time, without notice.*

*All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.*

*Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.*

*Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.*

*Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.*

*The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, and URIKA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.*





**CRAY**  
THE SUPERCOMPUTER COMPANY