

# Implementing Stencil Problems in Chapel: An Experience Report

---

Per Fuchs   Pieter Hijma   Clemens Grelck

Saturday 22 June 2019  
CHI UW 2019

- Teach Chapel in course
  - Programming Concurrent Systems (UvA)
  - Parallel Programming Practical (VU)
- Paper based on student report by Per Fuchs
- 4 versions
  - sequential
  - global view single locale
  - global view multi locale
  - local view multi locale

## System

- DAS-5
  - Dual socket, 8 cores, 2 hyperthreads
  - 48 Gbps InfiniBand
  - CentOS 7.4, 3.10.0 GNU/Linux
- Software
  - Chapel 1.19 (latest)
  - GCC 6.4.0 (latest version on DAS-5)

- Over the years disappointing multi-locale performance
- Our attempt to get to the bottom of this
  - from users' perspective
- CHI UW as target in mind
- Can serve as explanation to (former) students

# Conversation Users/Implementers

File Edit View Go Message Events and Tasks Tools Help

Inbox - p.hijma@uvs Q Fwd: Re: CHIUIW page limit

Get Messages Write Chat Address Book Tag Quick Filter Search <Ctrl+K> CardBook

Unread Starred CardBook Contact Tags Attachment Filter these messages <Ctrl+Shift+K>

Subject	Correspondents	Date	Location
CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	27/04/2019, 04:07	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	29/04/2019, 13:34	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	30/04/2019, 00:40	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	30/04/2019, 09:21	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	02/05/2019, 17:17	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	02/05/2019, 18:41	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	02/05/2019, 19:49	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	02/05/2019, 21:11	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 02:40	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 15:36	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 16:18	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 16:34	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 17:07	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 17:45	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 17:47	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 21:02	Sent
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	03/05/2019, 21:26	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	04/05/2019, 09:39	Inbox
Re: CHIUIW: possible performance gotchas for your stencil computation	Brad Chamberlain	04/05/2019, 13:27	Sent

From Brad Chamberlain

Subject: CHIUIW: possible performance gotchas for your stencil computation

To: Clemens Greick, Pieter Hijma, Per Fuchs

27/04/2019, 04:07

Unread: 0 Today Pane

# Textbook implementation

```
1  config const N = 8192;
2  config const M = 8192;
3  config const I = 500;
4  config const E = 0.01;
5
6  const HaloDomain: domain(2) = {0..N+1, 0..M+1};
7  const CylinderDomain: subdomain(HaloDomain) = {1..N, 1..M};
8
9  const LeftHalo: subdomain(HaloDomain) = {1..N, 0..0};
10 const RightHalo: subdomain(HaloDomain) = {1..N, M+1..M+1};
11 const UpperHalo: subdomain(HaloDomain) = {0..0, 0..M+1};
12 const LowerHalo: subdomain(HaloDomain) = {N+1..N+1, 0..M+1};
13
14 const LeftColumn: subdomain(HaloDomain) = {1..N, 1..1};
15 const RightColumn: subdomain(HaloDomain) = {1..N, M..M};
16 const UpperRow: subdomain(HaloDomain) = {1..1, 0..M+1};
17 const LowerRow: subdomain(HaloDomain) = {N..N, 0..M+1};
18
19 class Cylinder {
20   var temperature : [HaloDomain] real;
21 }
```

# Textbook implementation

```
1 config const N = 8192;
2 config const M = 8192;
3 config const I = 500;
4 config const E = 0.01;
5
6 const HaloDomain: domain(2) = {0..N+1, 0..M+1};
7 const CylinderDomain: subdomain(HaloDomain) = {1..N, 1..M};
8
9 const LeftHalo: subdomain(HaloDomain) = {1..N, 0..0};
10 const RightHalo: subdomain(HaloDomain) = {1..N, M+1..M+1};
11 const UpperHalo: subdomain(HaloDomain) = {0..0, 0..M+1};
12 const LowerHalo: subdomain(HaloDomain) = {N+1..N+1, 0..M+1};
13
14 const LeftColumn: subdomain(HaloDomain) = {1..N, 1..1};
15 const RightColumn: subdomain(HaloDomain) = {1..N, M..M};
16 const UpperRow: subdomain(HaloDomain) = {1..1, 0..M+1};
17 const LowerRow: subdomain(HaloDomain) = {N..N, 0..M+1};
18
19 class Cylinder {
20   var temperature : [HaloDomain] real;
21 }
```

# Textbook implementation

```
1 config const N = 8192;
2 config const M = 8192;
3 config const I = 500;
4 config const E = 0.01;
5
6 const HaloDomain: domain(2) = {0..N+1, 0..M+1};
7 const CylinderDomain: subdomain(HaloDomain) = {1..N, 1..M};
8
9 const LeftHalo: subdomain(HaloDomain) = {1..N, 0..0};
10 const RightHalo: subdomain(HaloDomain) = {1..N, M+1..M+1};
11 const UpperHalo: subdomain(HaloDomain) = {0..0, 0..M+1};
12 const LowerHalo: subdomain(HaloDomain) = {N+1..N+1, 0..M+1};
13
14 const LeftColumn: subdomain(HaloDomain) = {1..N, 1..1};
15 const RightColumn: subdomain(HaloDomain) = {1..N, M..M};
16 const UpperRow: subdomain(HaloDomain) = {1..1, 0..M+1};
17 const LowerRow: subdomain(HaloDomain) = {N..N, 0..M+1};
18
19 class Cylinder {
20     var temperature : [HaloDomain] real;
21 }
```



# Textbook implementation

```
1  for iteration in 1..I {
2
3      forall (i,j) in zip(LeftHalo, RightColumn) {
4          src.temperature[i] = src.temperature[j];
5      }
6      forall (i,j) in zip(RightHalo, LeftColumn) {
7          src.temperature[i] = src.temperature[j];
8      }
9
10     forall (i, j) in CylinderDomain {
11         var weight = conductivity[i, j];
12         var remaining_weight = 1 - weight;
13
14         dst.temperature[i, j] =
15             weight * src.temperature[i, j] +
16             // four direct neighbors
17             remaining_weight * factor_direct_neighbors *
18             (src.temperature[i-1, j] +
```

```
19             src.temperature[i, j+1] +
20             src.temperature[i+1, j] +
21             src.temperature[i, j-1]) +
22             // four diagonal neighbors
23             remaining_weight * factor_diagonal_neighbors *
24             (src.temperature[i-1, j-1] +
25             src.temperature[i-1, j+1] +
26             src.temperature[i+1, j+1] +
27             src.temperature[i+1, j-1]);
28     }
29
30     max_difference = max reduce [ij in CylinderDomain]
31         abs(dst.temperature[ij] - src.temperature[ij]);
32
33     if max_difference < E break;
34
35     src <=> dst;
36 }
```

# Textbook implementation

```
1  for iteration in 1..I {
2
3      forall (i,j) in zip(LeftHalo, RightColumn) {
4          src.temperature[i] = src.temperature[j];
5      }
6      forall (i,j) in zip(RightHalo, LeftColumn) {
7          src.temperature[i] = src.temperature[j];
8      }
9
10     forall (i, j) in CylinderDomain {
11         var weight = conductivity[i, j];
12         var remaining_weight = 1 - weight;
13
14         dst.temperature[i, j] =
15             weight * src.temperature[i, j] +
16             // four direct neighbors
17             remaining_weight * factor_direct_neighbors *
18             (src.temperature[i-1, j] +
```

```
19         src.temperature[i, j+1] +
20         src.temperature[i+1, j] +
21         src.temperature[i, j-1]) +
22         // four diagonal neighbors
23         remaining_weight * factor_diagonal_neighbors *
24         (src.temperature[i-1, j-1] +
25         src.temperature[i-1, j+1] +
26         src.temperature[i+1, j+1] +
27         src.temperature[i+1, j-1]);
28     }
29
30     max_difference = max reduce [ij in CylinderDomain]
31         abs(dst.temperature[ij] - src.temperature[ij]);
32
33     if max_difference < E break;
34
35     src <=> dst;
36 }
```

# Textbook implementation

```
1  for iteration in 1..I {
2
3  forall (i,j) in zip(LeftHalo, RightColumn) {
4      src.temperature[i] = src.temperature[j];
5  }
6  forall (i,j) in zip(RightHalo, LeftColumn) {
7      src.temperature[i] = src.temperature[j];
8  }
9
10 forall (i, j) in CylinderDomain {
11     var weight = conductivity[i, j];
12     var remaining_weight = 1 - weight;
13
14     dst.temperature[i, j] =
15         weight * src.temperature[i, j] +
16         // four direct neighbors
17         remaining_weight * factor_direct_neighbors *
18         (src.temperature[i-1, j] +
```

```
19         src.temperature[i, j+1] +
20         src.temperature[i+1, j] +
21         src.temperature[i, j-1]) +
22         // four diagonal neighbors
23         remaining_weight * factor_diagonal_neighbors *
24         (src.temperature[i-1, j-1] +
25         src.temperature[i-1, j+1] +
26         src.temperature[i+1, j+1] +
27         src.temperature[i+1, j-1]);
28     }
29
30     max_difference = max reduce [ij in CylinderDomain]
31         abs(dst.temperature[ij] - src.temperature[ij]);
32
33     if max_difference < E break;
34
35     src <=> dst;
36 }
```

- vectorization has a large impact

## Feedback from Chapel

- LLVM backend will give more control over vectorization

# Global View Single Locale

```
1  for iteration in 1..I {
2
3    forall (i,j) in zip(LeftHalo, RightColumn) {
4      src.temperature[i] = src.temperature[j];
5    }
6    forall (i,j) in zip(RightHalo, LeftColumn) {
7      src.temperature[i] = src.temperature[j];
8    }
9
10   forall (i, j) in CylinderDomain {
11     var weight = conductivity[i, j];
12     var remaining_weight = 1 - weight;
13
14     dst.temperature[i, j] =
15       weight * src.temperature[i, j] +
16       // four direct neighbors
17       remaining_weight * factor_direct_neighbors *
18       (src.temperature[i-1, j] +
```

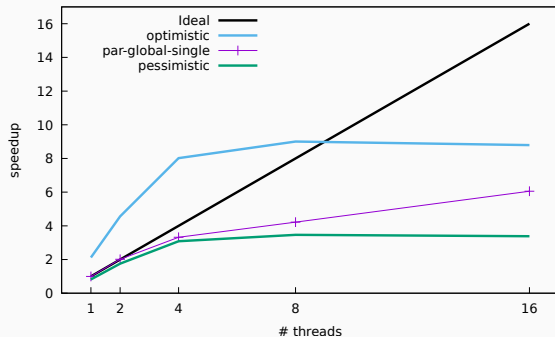
```
19       src.temperature[i, j+1] +
20       src.temperature[i+1, j] +
21       src.temperature[i, j-1]) +
22       // four diagonal neighbors
23       remaining_weight * factor_diagonal_neighbors *
24       (src.temperature[i-1, j-1] +
25       src.temperature[i-1, j+1] +
26       src.temperature[i+1, j+1] +
27       src.temperature[i+1, j-1]);
28   }
29
30   max_difference = max reduce [ij in CylinderDomain]
31     abs(dst.temperature[ij] - src.temperature[ij]);
32
33   if max_difference < E break;
34
35   src <=> dst;
36 }
```

# Global View Single Locale

```
1  for iteration in 1..I {
2
3    forall (i,j) in zip(LeftHalo, RightColumn) {
4      src.temperature[i] = src.temperature[j];
5    }
6    forall (i,j) in zip(RightHalo, LeftColumn) {
7      src.temperature[i] = src.temperature[j];
8    }
9
10   forall (i, j) in CylinderDomain {
11     var weight = conductivity[i, j];
12     var remaining_weight = 1 - weight;
13
14     dst.temperature[i, j] =
15       weight * src.temperature[i, j] +
16       // four direct neighbors
17       remaining_weight * factor_direct_neighbors *
18       (src.temperature[i-1, j] +
```

```
19       src.temperature[i, j+1] +
20       src.temperature[i+1, j] +
21       src.temperature[i, j-1]) +
22       // four diagonal neighbors
23       remaining_weight * factor_diagonal_neighbors *
24       (src.temperature[i-1, j-1] +
25       src.temperature[i-1, j+1] +
26       src.temperature[i+1, j+1] +
27       src.temperature[i+1, j-1]);
28   }
29
30   max_difference = max reduce [ij in CylinderDomain]
31     abs(dst.temperature[ij] - src.temperature[ij])
32
33   if max_difference < E break;
34
35   src <=> dst;
36 }
```

# Global View Single Locale



Threads	1	2	4	8	16
time (s)	152.20	75.23	45.58	35.85	24.99
Speedup	0.99	2.01	3.32	4.22	6.06
BW STREAM (GB/s)	18.8	40.4	71.1	79.9	78.0
Required BW (GB/s)	22.9	46.4	76.6	97.3	139.6

# Global View Multi Locale

```
1 use StencilDist;
2
3 var MyLocaleView = {0..numLocales-1, 0..0};
4 var MyLocales =
5   reshape(Locales[0..numLocales-1], MyLocaleView);
6
7 const HaloDomain: domain(2) dmapped Stencil(
8   boundingBox = {0..N+1, 0..M+1},
9   targetLocales = MyLocales,
10  fluff=(1,1)) = {0..N+1, 0..M+1};
```



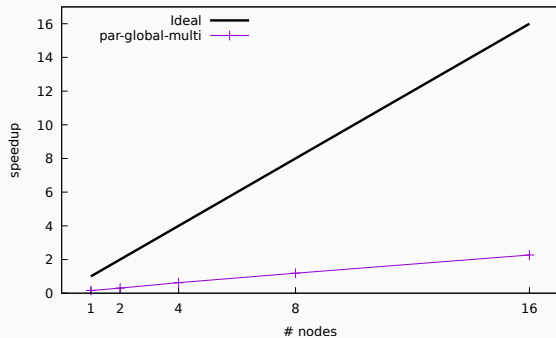
# Global View Multi Locale

```
1 for iteration in 1..I {
2
3   ref s = src.temperature;
4   ref d = dst.temperature;
5
6   // exchange columns
7
8   forall (i, j) in CylinderDomain {
9     local {
10       var weight = conductivity[i, j];
11       var remaining_weight = 1 - weight;
12
13       d[i, j] =
14         weight * s[i, j] +
15         // neighbor computations
16       }
17     }
18     d.updateFluff();
19
20     max_difference = max reduce [ij in CylinderDomain]
21       abs d[ij] - s[ij];
22
23     if max_difference < E break;
24
25     src <=> dst;
26 }
```

# Global View Multi Locale

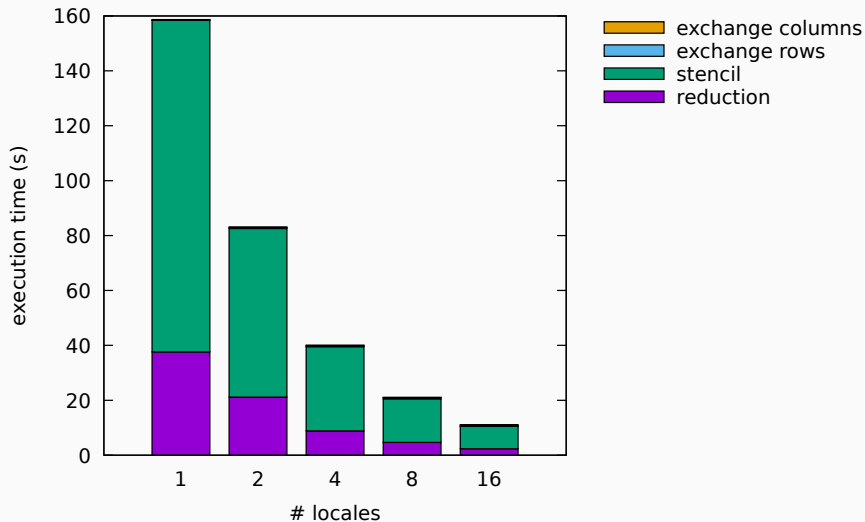
```
1 for iteration in 1..I {
2
3   ref s = src.temperature;
4   ref d = dst.temperature;
5
6   // exchange columns
7
8   forall (i, j) in CylinderDomain {
9     local {
10       var weight = conductivity[i, j];
11       var remaining_weight = 1 - weight;
12
13       d[i, j] =
14         weight * s[i, j] +
15         // neighbor computations
16       }
17   }
18   d.updateFluff();
19
20   max_difference = max reduce [ij in CylinderDomain]
21     abs(d[ij] - s[ij]);
22
23   if max_difference < E break;
24
25   src <=> dst;
26 }
```

# Global View Multi Locale



Nodes	1	2	4	8	16
time (s)	158.68	83.13	40.08	21.10	11.03
Speedup	0.16	0.30	0.62	1.18	2.27
Efficiency (%)	0.16	0.15	0.16	0.15	0.14

# Global View Multi Locale



## Locality

- `local` blocks are deprecated
- `localAccess` expressions

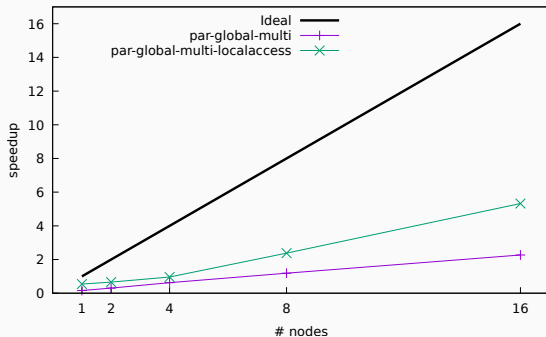
## NUMA awareness

- InfiniBand uses GASNet as communication layer
- GASNet is not NUMA aware
- Registering memory for the network is required
- Happens only on one NUMA-node
- Other NUMA-node suffers from this

```
1 for iteration in 1..I {
2
3   ref s = src.temperature;
4   ref d = dst.temperature;
5
6   // exchange columns
7
8   forall (i, j) in CylinderDomain {
9     local {
10       var weight = conductivity[i, j];
11       var remaining_weight = 1 - weight;
12
13       d[i, j] =
14         weight * s[i, j] +
15         // neighbor computations
16     }
17   }
18   d.updateFluff();
19
20   max_difference = max reduce [ij in CylinderDomain]
21     abs(d[ij] - s[ij]);
22
23   if max_difference < E break;
24
25   src <=> dst;
26 }
```

```
1 for iteration in 1..I {
2
3   ref s = src.temperature;
4   ref d = dst.temperature;
5
6   // exchange columns
7
8   forall (i, j) in CylinderDomain {
9
10    var weight = conductivity.localAccess[i, j];
11    var remaining_weight = 1 - weight,
12
13    d.localAccess[i, j] =
14      weight * s.localAccess[i, j] +
15      // neighbor computations
16
17    }
18    d.updateFluff();
19
20    max_difference = max reduce [ij in CylinderDomain]
21      abs(d.localAccess[ij] - s.localAccess[ij]);
22
23    if max_difference < E break;
24
25    src <=> dst;
26 }
```

# Performance difference



- factor 3.5 faster for 1 locale
- factor 2.4 faster for 16 locales



# Local View Multi Locale

```
1 const Row : domain(1) = {0..M+1};
2
3 class Communicator {
4   var firstRow: [Row] real;
5   var lastRow: [Row] real;
6 }
7
8 var communicators: [LocaleSpace] Communicator;
```

# Local View Multi Locale

```
1 var globalIterations = 0;
2 var globalMaxDiff: real = min(real);
3
4 forall l in Locales with
5   (ref globalIterations, ref globalMaxDiff) {
6   on l {
7     const myCommunicator = new unmanaged Communicator();
8     communicators[here.id] = myCommunicator;
9
10    allLocalesBarrier.barrier();
11    const isLast = here.id == LocaleSpace.last;
12    const isFirst = here.id == LocaleSpace.first;
13
14    const beforeCommunicator = if !isFirst
15      then communicators[here.id - 1]
16      else new unmanaged Communicator();
17    const nextCommunicator = if !isLast
18      then communicators[here.id + 1]
19      else new unmanaged Communicator();
```


```
1 const LocalCylinderDomain = CylinderDomain.localSubdomain();  
2 const LocalHaloDomain = LocalCylinderDomain.expand(1, 1);  
3  
4 const LocalUpperRow = LocalHaloDomain.dim(1).first;  
5 const LocalLowerRow = LocalHaloDomain.dim(1).last;
```

# Local View Multi Locale

```
1 do {
2   ref s = src.temperature;
3   ref d = dst.temperature;
4
5   local {
6     local_max_difference = min(real);
7
8     forall (i, j) in LocalCylinderDomain with
9       (max reduce local_max_difference) {
10      const weight = localConductivity[i, j];
11      const remaining_weight = 1 - weight;
12      const oldTemp = s[i, j];
13
14      const newTemp = weight * oldTemp +
15        // neighbor computations
16
17      d[i, j] = newTemp;
18      local_max_difference = max(local_max_difference,
19        abs(new_temp - old_temp));
20    }
21    // exchange columns
22  }
23
24  nextCommunicator.firstRow = d[LocalLowerRow - 1, ..];
```

```
25 beforeCommunicator.lastRow = d[LocalUpperRow + 1, ..];
26
27 src <=> dst;
28
29 // located on locale 0
30 max_diffs[here.id] = local_max_difference;
31
32 allLocalesBarrier.barrier();
33
34 if (!isLast) {
35   forall col in Row {
36     d[LocalLowerRow, col] = myCommunicator.lastRow[col];
37   }
38 }
39 if (!isFirst) {
40   forall col in Row {
41     d[LocalUpperRow, col] =
42   }
43 }
44
45 local_max_difference = max reduce max_diffs;
46
47 local_iteration += 1;
48 } while (local_max_difference > E && local_iteration < I);
```

# Local View Multi Locale




Nodes	1	2	4	8	16
time (s)	47.15	24.74	14.55	7.62	4.33
Speedup	0.53	1.01	1.72	3.28	5.78
Efficiency	0.53	0.51	0.43	0.41	0.36

## Same issues

- locality control
- NUMA awareness of the communication layer

# Performance difference




- factor 1.46 faster for 1 locale
- factor 1.36 faster for 16 locales

- We can explain the disappointing performance
- We cannot achieve high performance on our system
  - NUMA awareness of communication layer remains a problem
  - new OpenFabrics layer is promising
- Chapel is a very powerful and expressive language



# Performance




## localAccess more fine-grained

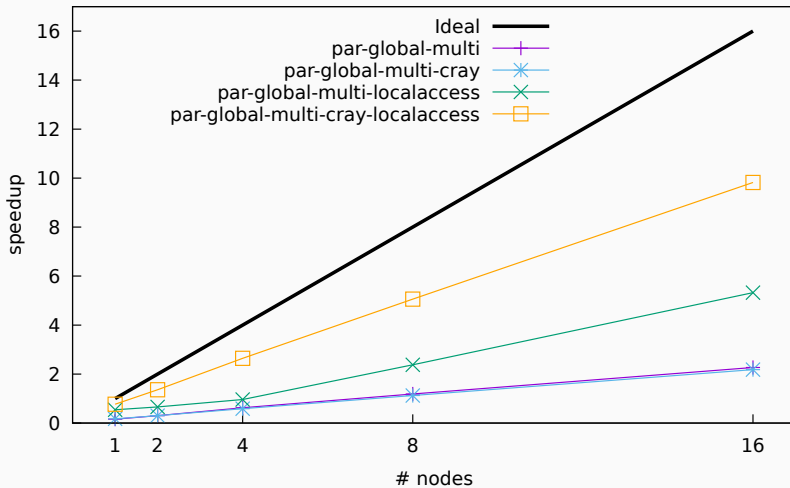
```
1 for iteration in 1..I {
2
3   ref s = src.temperature;
4   ref d = dst.temperature;
5
6   // exchange columns
7
8   forall (i, j) in CylinderDomain {
9     local {
10      var weight = conductivity[i, j];
11      var remaining_weight = 1 - weight;
12
13      d[i, j] =
14        weight * s[i, j] +
15        // neighbor computations
16      }
17    }
18    d.updateFluff();
19
20    max_difference = max reduce [ij in CylinderDomain]
21      abs(d[ij] - s[ij]);
22
23    if max_difference < E break;
24
25    src <=> dst;
26 }
```

```
1 var tmp : [CylinderDomain] real;
2
3 local {
4   tmp = [ij in CylinderDomain] abs(d[ij] - s[ij]);
5 }
6
7 max_difference = max reduce tmp;
```

# Local View Multi Locale



# All Attempts Multi Locale Global View (speculative)



## All Attempts Multi Locale Local View (speculative)

