

# Arachne: An Open-Source Framework for Interactive Massive-Scale Graph Analytics



**David A. Bader**

 [@Prof\\_DavidBader](https://twitter.com/Prof_DavidBader)

<http://www.cs.njit.edu/~bader>



# David A. Bader

Distinguished Professor and  
Director, Institute for Data Science

- IEEE Fellow, ACM Fellow, SIAM Fellow, AAAS Fellow
- IEEE Sidney Fernbach Award
- 2022 inductee into University of Maryland's Innovation Hall of Fame, A. James Clark School of Engineering
- Recent Service:
  - White House's National Strategic Computing Initiative (NSCI) panel
  - Computing Research Association Board
  - Chair, NSF Committee of Visitors for Office of Advanced Cyberinfrastructure
  - NSF Advisory Committee on Cyberinfrastructure
  - Council on Competitiveness HPC Advisory Committee
  - IEEE Computer Society Board of Governors
  - IEEE IPDPS Steering Committee
  - Editor-in-Chief, ACM Transactions on Parallel Computing
  - Editor-in-Chief, IEEE Transactions on Parallel and Distributed Systems
- Over \$188M of research awards
- 300+ publications,  $\geq 13,900$  citations, h-index  $\geq 66$
- National Science Foundation CAREER Award recipient
- Directed: Facebook AI Systems
- Directed: NVIDIA GPU Center of Excellence, NVIDIA AI Lab (NVAIL)
- Directed: Sony-Toshiba-IBM Center for the Cell/B.E. Processor
- Founder: Graph500 List benchmarking “Big Data” platforms
- Recognized as a “RockStar” of High Performance Computing by InsideHPC in 2012 and as HPCwire’s People to Watch in 2012 and 2014.



# Today's talk is dedicated to my uncle

## Elliot Norman Ashrey

- B: 23 Sep 1931 in New York City, New York, USA
- D: 26 Apr 2024 in New York City, New York, USA



# 2021 IEEE Sidney Fernbach Award

David Bader cited for the development of Linux-based massively parallel production computers and for pioneering contributions to scalable discrete parallel algorithms for real-world applications.



2022 IEEE Computer Society President Bill Gropp presents David Bader with the Sidney Fernbach Award at SC21

# 1998: Bader Invents the Linux Supercomputer

DEPARTMENT EDITOR: Alex Magoun, annals-anecdotes@computer.org

## ANECDOTES

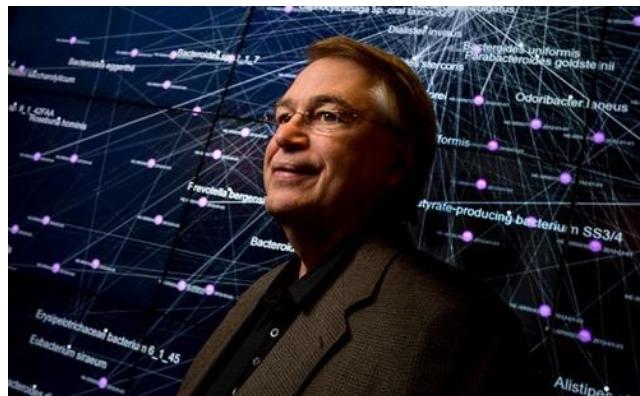
### Linux and Supercomputing: How My Passion for Building COTS Systems Led to an HPC Revolution

David A. Bader , Ying Wu College of Computing, New Jersey Institute of Technology, Newark, NJ, USA

**B**ack in the early 1990s, when I was a graduate student in electrical and computer engineering at the University of Maryland, the term "supercomputer" meant Single Instruction, Multiple Data (SIMD) vector processor machines (the Cray-1 was the most popular), or massively parallel multiprocessor systems, such as the Thinking Machine CM-5. These systems were bulky—a Cray-1 occupied 2.7m × 2m of floor area and contained 60 miles of wires; expensive, selling for several million dollars; and required significant expertise to program and operate.

required a simultaneous development of scalable, high performance algorithms and services. Otherwise, application developers would be forced to develop algorithms from scratch every time vendors introduced a newer, faster, hardware platform.

By the late 1990s, the term "cluster computing" was common among computer science researchers and several of these systems had received significant publicity. One of the first cluster approaches to attract interest was Beowulf, which cost from a tenth to a third of the price of a traditional supercomputer. A typical setup



Source: UC San Diego  
<https://ucsdnews.ucsd.edu/pressrelease/pioneering-scientist-and-innovator-larry-smarr-retires>



"This effort of yours has enormous historic resonance,"  
— Larry Smarr, Distinguished Professor Emeritus, UC San Diego  
Founding Director of NCSA, Founding Director of Calit2

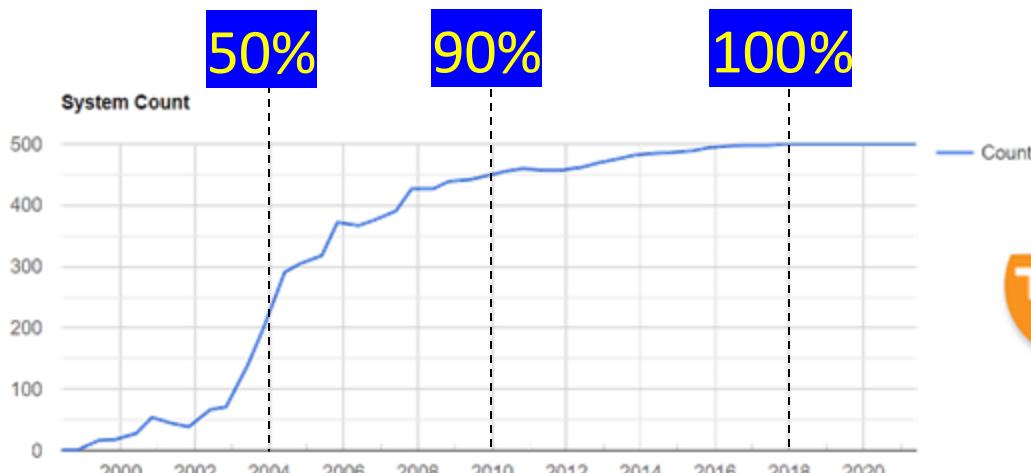
# Impact: Top500 Supercomputers Running Linux



**“Today, 100% of the Top 500 supercomputers in the world are Linux HPC systems, based on Bader’s technical contributions and leadership. This is one of the most significant technical foundations of HPC.”**

— Steve Wallach is a guest scientist for Los Alamos National Laboratory and 2008 IEEE CS Seymour Cray Computer Engineering Award recipient.

Photo credit: Information Week, 2008



N J I T

New Jersey Institute  
of Technology

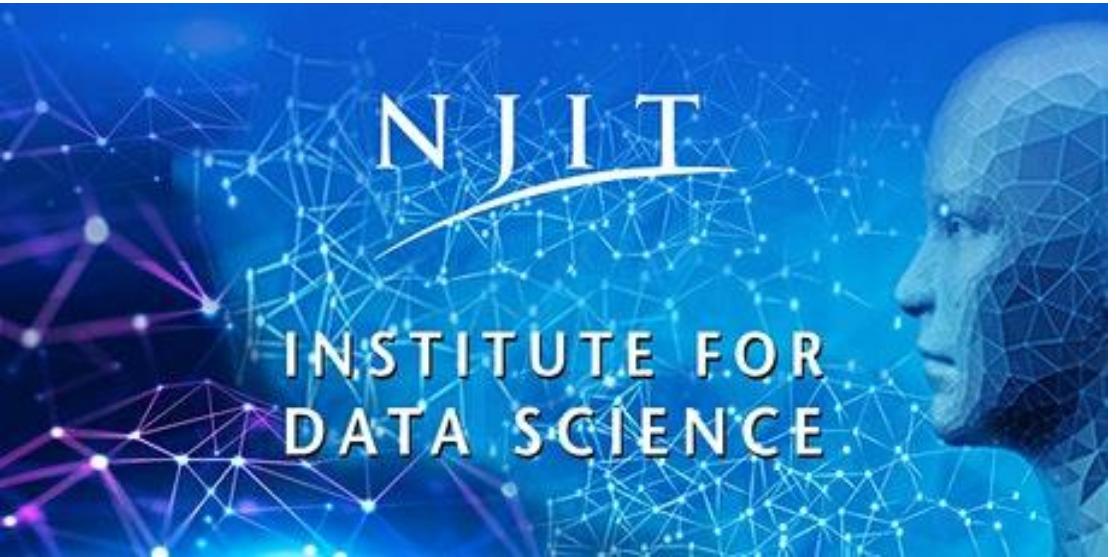
# New Jersey Institute of Technology



“NJIT Climbs the Rankings of U.S. News & World Report, A Top 50 Public University”  
– 13 Sep 2021

“NJIT Named As One of Nation's 'Best Colleges' for 2022, The Princeton Review Says”  
– 6 Sep 2021

“Wall Street Journal/College Pulse Ranks NJIT No. 2 Public University in the US”  
– 6 Sep 2023



Launched in **July 2019**, with inaugural director  
**David A. Bader**  
(~40 faculty in current centers)

Solving  
real-world  
challenges

- Urban sustainability
- Healthcare analytics
- Trustworthy, Free and Fair Elections
- Insider threat detection
- Utility infrastructure protection
- Cyberattack defense
- Disease outbreak and epidemic monitoring

#### Center for Big Data

- Big Data Analytics, Systems, and Tools
- Cyberinfrastructure

#### Cybersecurity Research Center

- Practical encryption
- Privacy technologies
- Information Assurance

#### The Structural Analysis of Biomedical Ontologies Center

- Medical Informatics
- NIH / National Cancer Institute

#### FinTech Group

- Financial Services
- Insurance Industry

#### Machine Learning & AI

- Real-world technologies
- Industrial partnerships

# Data-Quad



Known

Objects

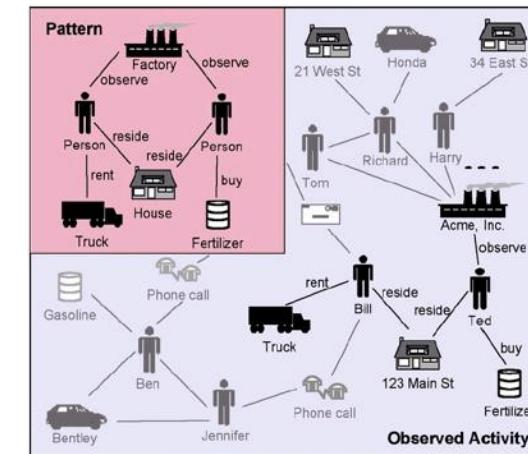


Unknown



Unknown

Patterns  
↓  
Known



# Graph Data Science: Real-world challenges

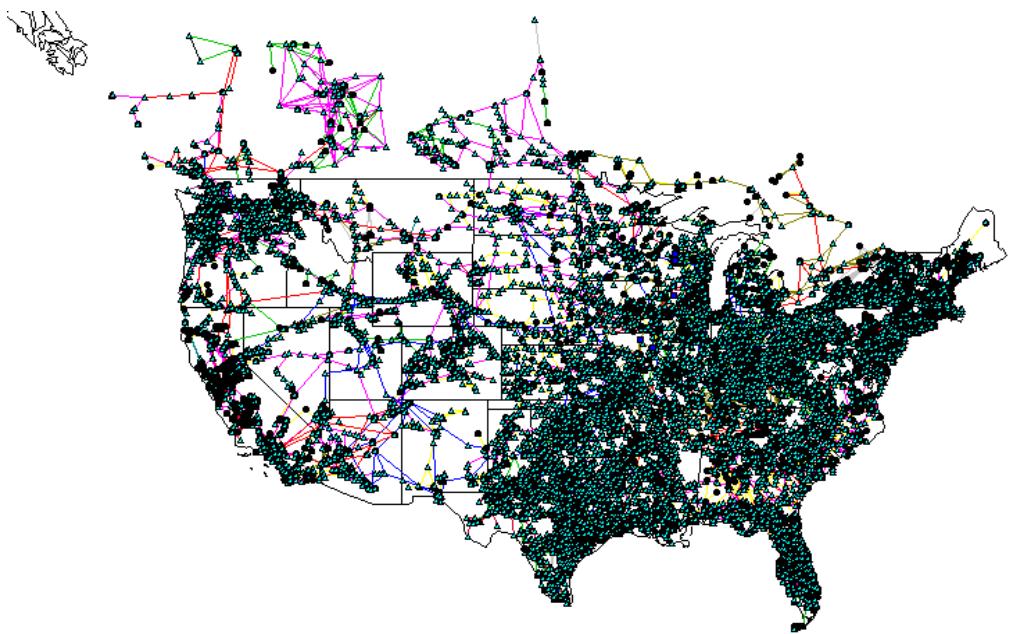
## All involve exascale streaming graphs:

- **Health care** → disease spread, detection and prevention of epidemics/pandemics (e.g. SARS, Avian flu, H1N1 “swine” flu)
- **Massive social networks** → understanding communities, intentions, population dynamics, pandemic spread, transportation and evacuation
- **Intelligence** → business analytics, anomaly detection, security, knowledge discovery from massive data sets
- **Systems Biology** → understanding complex life systems, drug design, microbial research, unravel the mysteries of the HIV virus; understand life, disease,
- **Electric Power Grid** → communication, transportation, energy, water, food supply
- **Modeling and Simulation** → Perform full-scale economic-social-political simulations

**REQUIRES PREDICTING / INFLUENCE CHANGE IN REAL-TIME AT SCALE**

# Massive Data Analytics: Infrastructure

- The U.S. high-voltage transmission grid has >150,000 miles of line.
- Real-time detection of changes and anomalies in the grid is a large-scale problem.
- May mitigate impact of widespread blackouts due to equipment failure or intentional damage.



The New York Times  
Thursday, September 4, 2008

## Report on Blackout Is Said To Describe Failure to React

By MATTHEW L. WALD  
Published: November 12, 2003

A report on the Aug. 14 blackout identifies specific lapses by various parties, including FirstEnergy's failure to react properly to the loss of a transmission line, people who have seen drafts of it say.

A working group of experts from eight states and Canada will meet in private on Wednesday to evaluate the report, people involved in the investigation said Thursday. The report, which the Energy Department

- [!\[\]\(e7b6800370c498f39807640753e19040\_img.jpg\) E-MAIL](#)
- [!\[\]\(fad5ce07162b51183e4dc4cc396a4edb\_img.jpg\) PRINT](#)
- [!\[\]\(49046ffdeb37439fbcb7169bec7c5390\_img.jpg\) SINGLE-PAGE](#)
- [!\[\]\(cefd22a4076605c4f689f2ddaad2c0f3\_img.jpg\) REPRINTS](#)
- [!\[\]\(1a952a9ed8a03f706b992745a2b8b624\_img.jpg\) SAVE](#)
- [!\[\]\(594ed574e0bfabdc993114c34c54f018\_img.jpg\) SHARE](#)

# Network Analysis for Intelligence and Surveillance

- [Krebs '04] Post 9/11 Terrorist Network Analysis from public domain information
- Plot masterminds correctly identified from interaction patterns: **centrality**

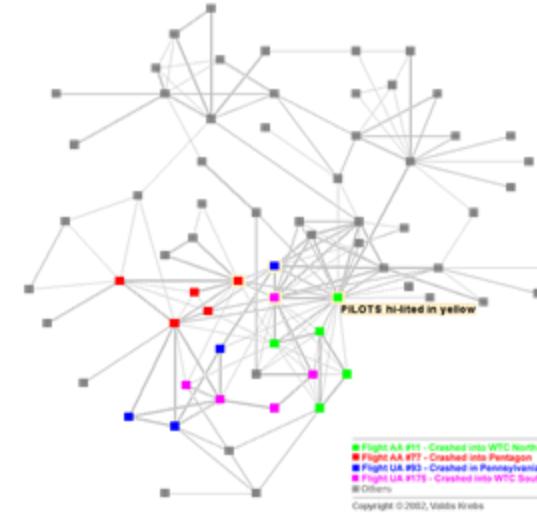


Image Source: <http://www.orgnet.com/hijackers.html>

- A global view of entities is often more insightful
- Detect anomalous activities by exact/approximate **graph matching**

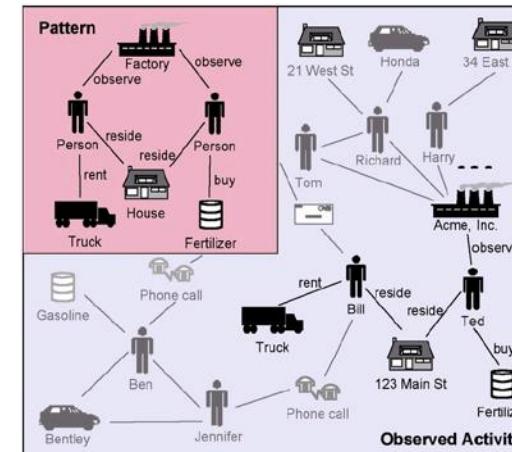
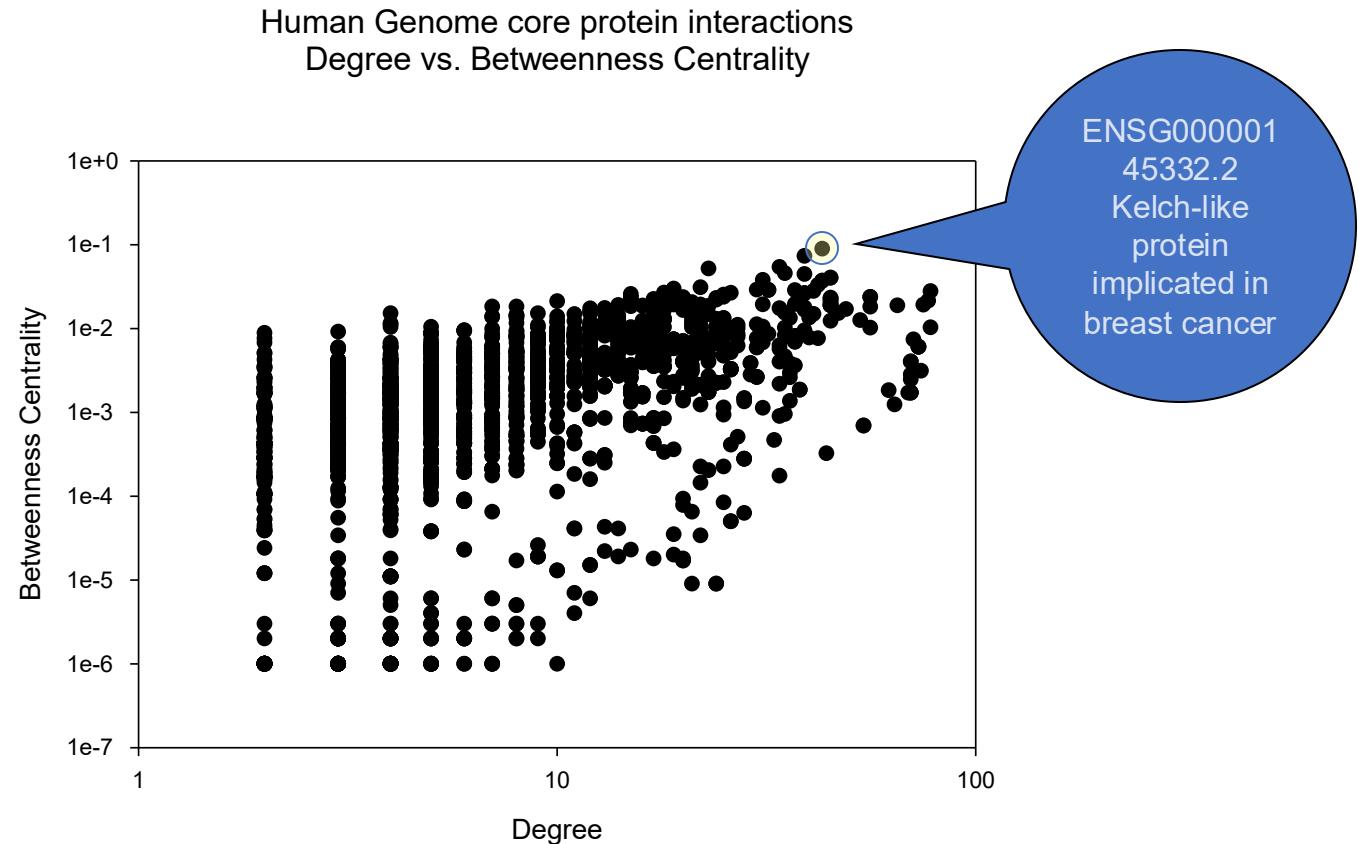


Image Source: T. Coffman, S. Greenblatt, S. Marcus, Graph-based technologies for intelligence analysis, CACM, 47 (3, March 2004): pp 45-47

# Massive Data Analytics: Public Health

- CDC/national-scale surveillance of public health
- Cancer genomics and drug design
  - Computed Betweenness Centrality of Human Proteome



# Mining Twitter for Social Good

ICPP 2010

## Massive Social Network Analysis: Mining Twitter for Social Good

David Ediger Karl Jiang  
Jason Riedy David A. Bader  
Georgia Institute of Technology  
Atlanta, GA, USA

Courtney Corley Rob Farber  
Pacific Northwest National Lab.  
Richland, WA, USA

William N. Reynolds  
Least Squares Software, Inc.  
Albuquerque, NM, USA

**Abstract**—Social networks produce an enormous quantity of data. Facebook consists of over 400 million active users sharing over 5 billion pieces of information each month. Analyzing this vast quantity of unstructured data presents challenges for software and hardware. We present GraphCT, a *Graph Characterization Toolkit* for massive graphs representing social network data. On a 128-processor Cray XMT, GraphCT estimates the betweenness centrality of an artificially generated (R-MAT) 537 million vertex, 8.6 billion edge graph in 35 minutes and a real-world graph (Kwak, et al.) with 61.6 million vertices and 1.47 billion edges in 105 minutes. We use GraphCT to analyze public data from Twitter, a microblogging network. Twitter's message connections appear primarily tree-structured as a news dissemination system. Within the

involves over 400 million active users with an average of 120 ‘friendship’ connections each and sharing 5 billion references to items each month [11].

One analysis approach treats the interactions as graphs and applies tools from graph theory, social network analysis, and scale-free networks [29]. However, the volume of data that must be processed to apply these techniques overwhelms current computational capabilities. Even well-understood analytic methodologies require advances in both hardware and software to process the growing corpus of social media.

Social media provides staggering amounts of data.

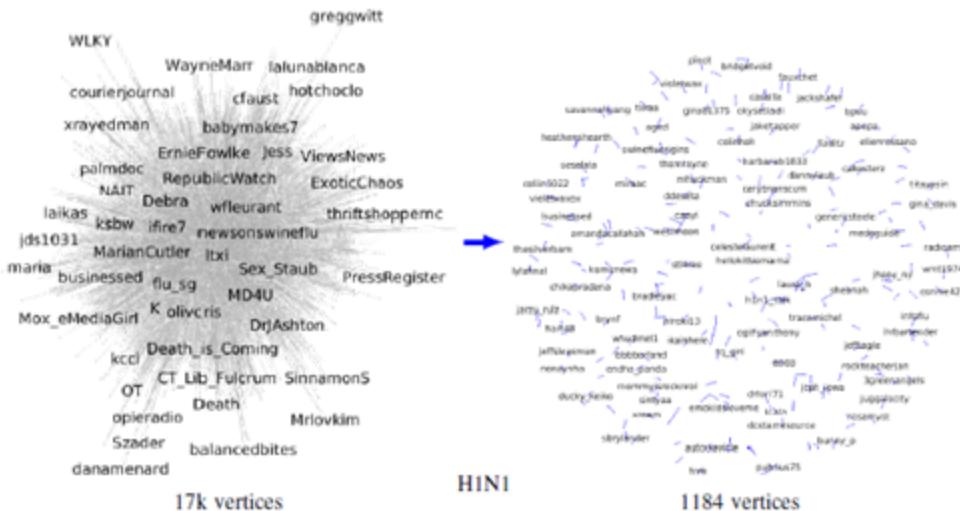


Fig. 3 Subcommunity filtering on Twitter data sets

## TOP 15 USERS BY BETWEENNESS CENTRALITY

Rank	HIN1	Data Set
1	@CDCFlu	@ajc
2	@addthis	@driveafastercar
3	@Official_PAX	@ATLCheap
4	@FluGov	@TWCi
5	@nytimes	@HelloNorthGA
6	@tweetmeme	@11AliveNews
7	@mercola	@WSB_TV
8	@CNN	@shaunking
9	@backstreetboys	@Carl
10	@EllieSmith_x	@SpaceyG
11	@TIME	@ATLINTownPaper
12	@CDCEmergency	@TJsDJs
13	@CDC_eHealth	@ATLien
14	@perezhilton	@MarshallRamsey
15	@billmaher	@Kanye

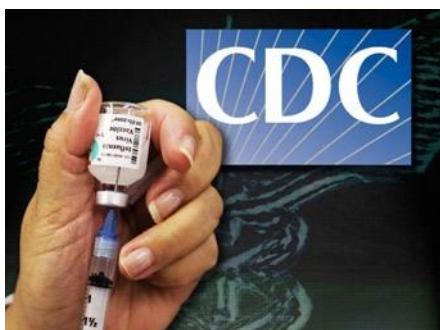


Image credit: bioethicsinstitute.org

David A. Bader



15

N J I T  
New Jersey Institute  
of Technology

# Arachne: Interactive Property Graph Analytics at Scale



Image Credit: Matias Del Carmen

# Institute for Data Science Aims to Democratize Supercomputing With NSF Grant

Written by: Evan Koblenz

Published: Wednesday, March 17, 2021



New algorithms from at NJIT can make supercomputer power available to almost anyone

Ordinary people could soon have greater ability to analyze massive amounts of information, based on new algorithms and software tools being designed at NJIT, intended to simplify

<https://news.njit.edu/institute-data-science-aims-democratize-supercomputing-nsf-grant>

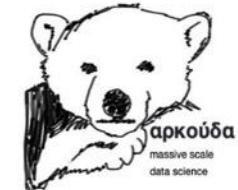
## High Performance Algorithms for Interactive Data Science at Scale

(PI: Bader)

\$2.2M

3/2021 – 6/2024

NSF CCF-2109988



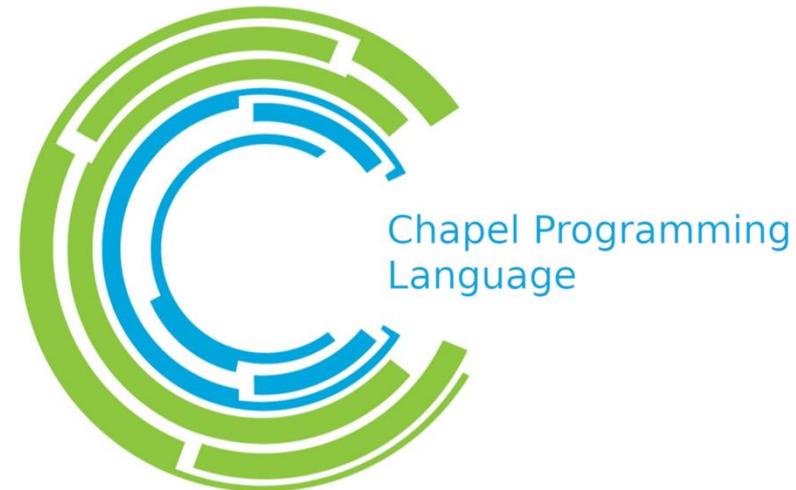
# Arkouda: Dedication to Michael H. Merrill

(June 2, 1964 ~ November 8, 2022)



**“Mike was a dedicated civil servant. He was a Computer Scientist at the Department of Defense for 34 years and was recognized in 2022 with a Distinguished Civilian Service Medal. He loved computers and technology, especially high performance computing. Mike was a problem solver and innovative thinker; he was recognized for inspiring and leading numerous large projects over the course of his career. He loved to share his knowledge and mentored many colleagues over the years—sometimes calling them his kids, sometimes his minions, but always calling them his friend.”**

# Productivity + Performance



# Connect with Chapel: the parallel programming language powering Arkouda

## ChapelCon – free virtual event

- June 5<sup>th</sup> – Tutorial Day
- June 6<sup>th</sup> – Coding Day
- June 7<sup>th</sup> – Conference Day

## Come code or chat with us!

- GitHub - <https://github.com/chapel-lang/chapel>
- Gitter - <https://gitter.im/chapel-lang/chapel>
- Discourse - <https://chapel.discourse.group>
- StackOverflow - <https://stackoverflow.com/questions/tagged/chapel>

## Follow us on social media

- LinkedIn - <https://www.linkedin.com/company/chapel-programming-language>
- YouTube - <https://www.youtube.com/@ChapelLanguage>
- Twitter/X - <https://x.com/ChapelLanguage>
- Facebook - <https://www.facebook.com/ChapelLanguage>

REGISTER FOR CHAPELCON



TAKE THE CHAPEL  
COMMUNITY SURVEY



# Arkouda + Arachne Framework

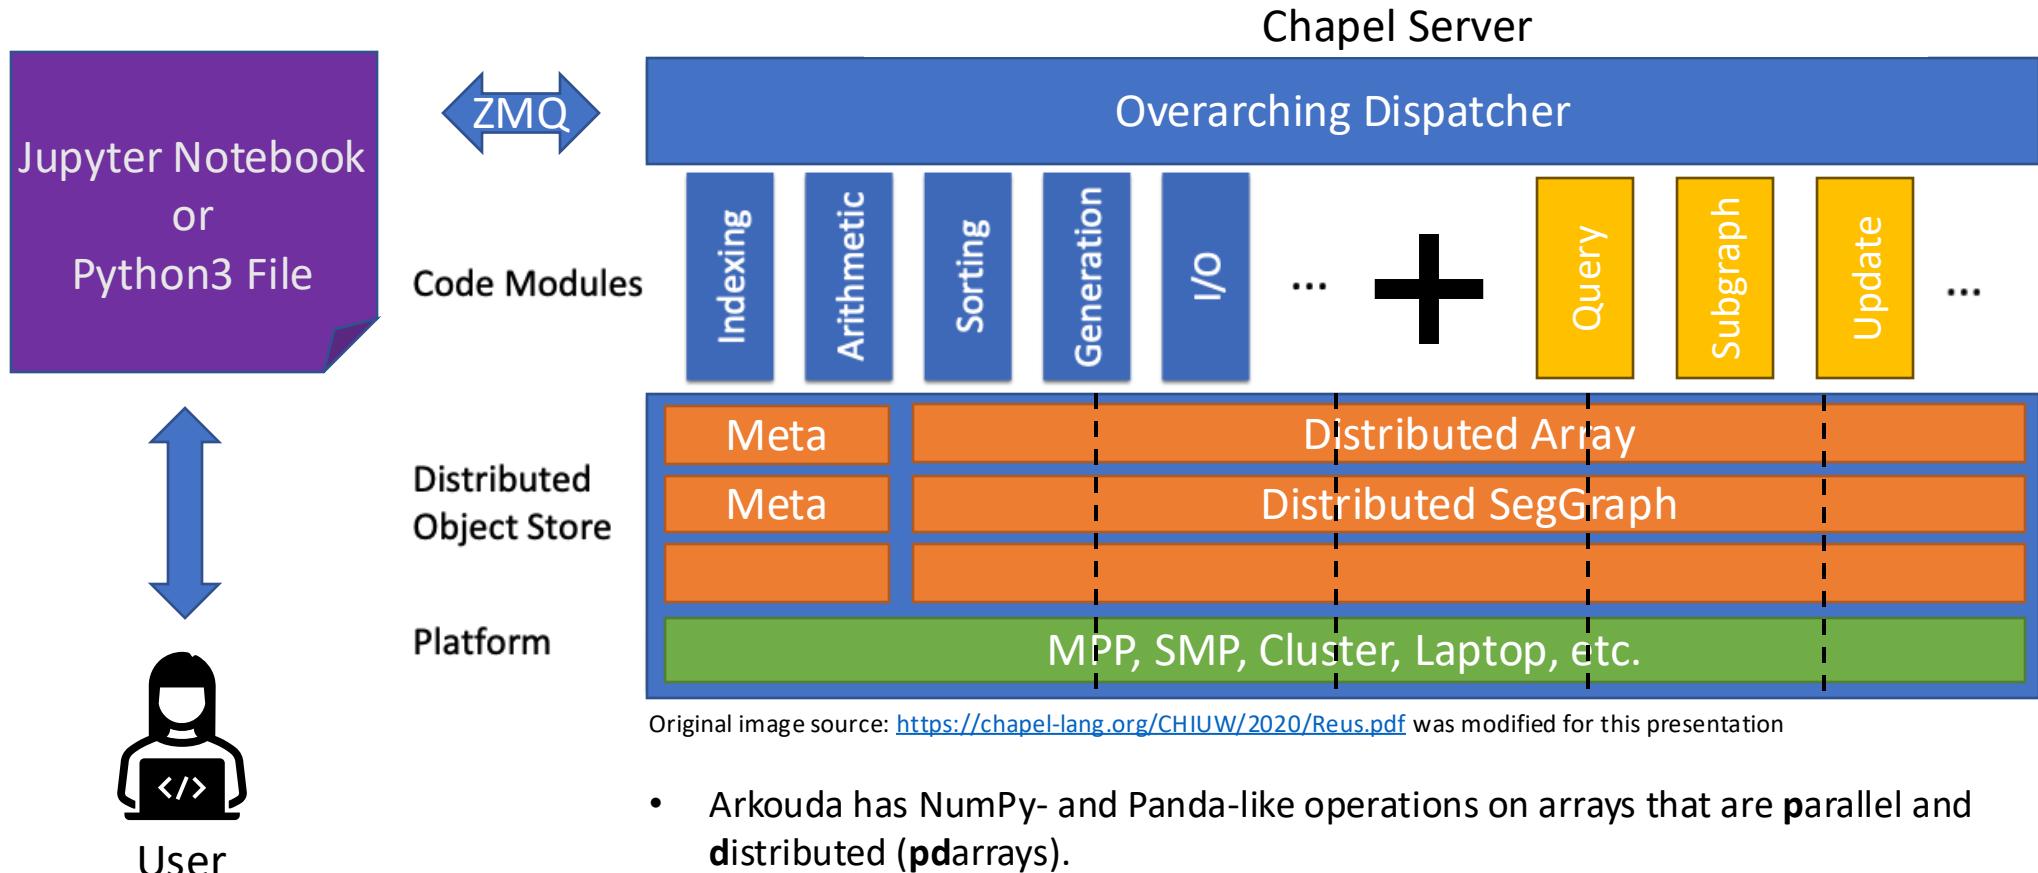
## Arkouda

- an existing open-source Python framework that allows for array and dataframe operations on data that is terabytes in size but lacks graph processing operations.

## Arachne

- an open-source extension to Arkouda to convert massive-scale dataframes to graphs with high-performance graph kernels and property graph capabilities while maintaining a NetworkX-like API for new Python users to easily transition to utilizing it.

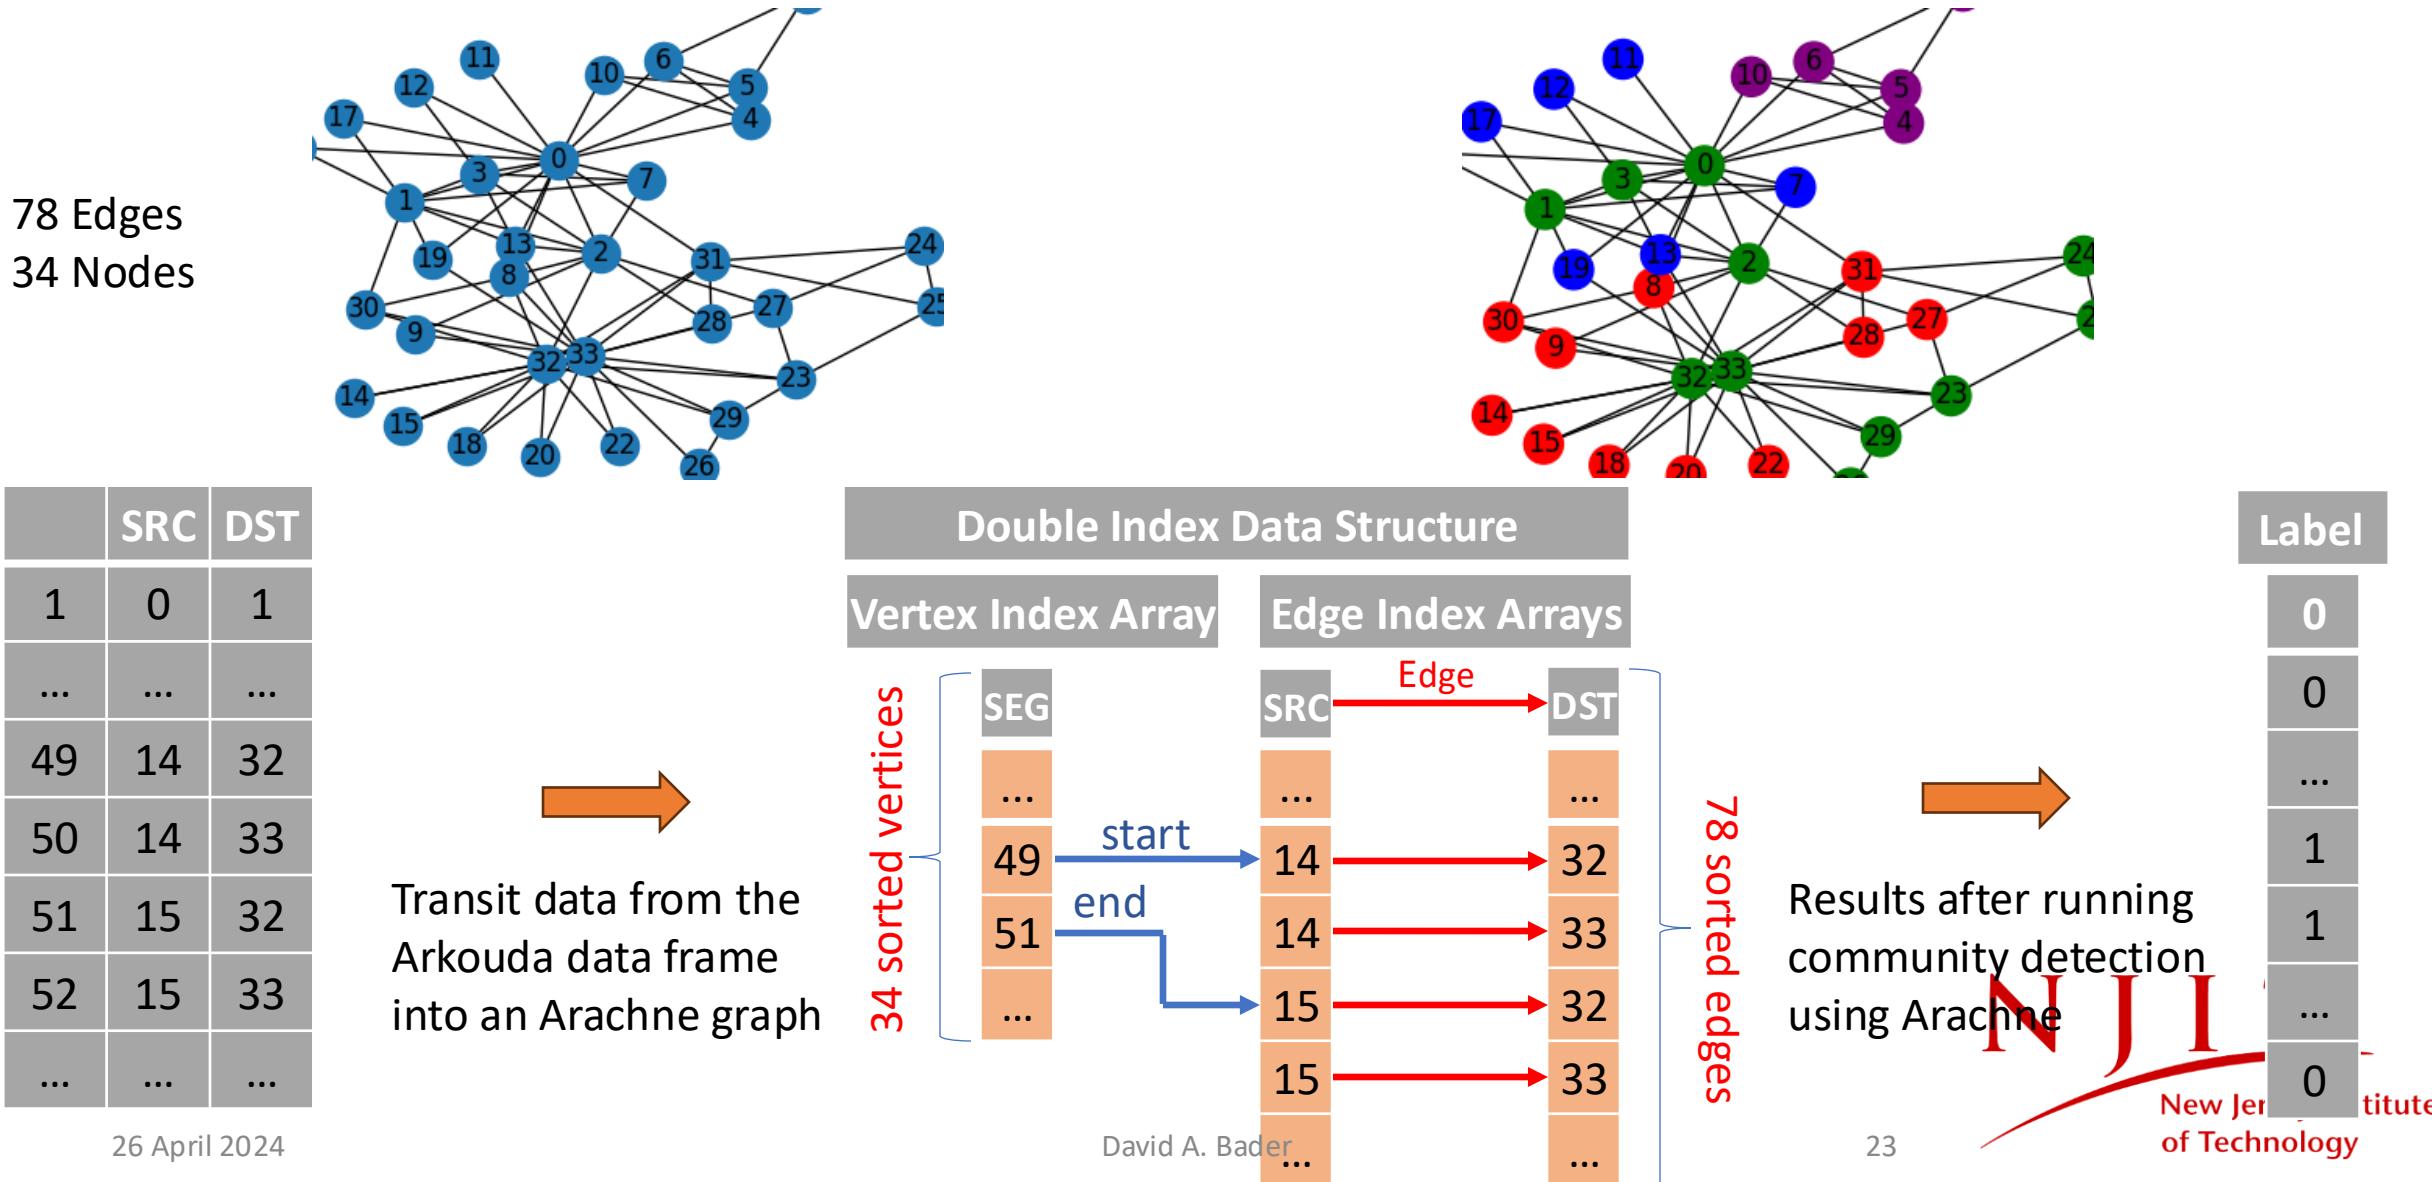
# Arkouda + Arachne Framework



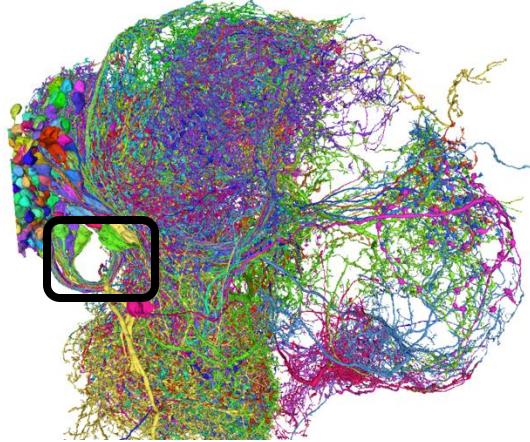
Original image source: <https://chapel-lang.org/CHIUW/2020/Reus.pdf> was modified for this presentation

- Arkouda has NumPy- and Panda-like operations on arrays that are **parallel** and **distributed** (**pdarrays**).
- Arachne extends Arkouda with graph capabilities.
- This work extends Arachne to store massive-scale graphs.
- Arachne can be thought of as a wrapper that creates a logical graph.

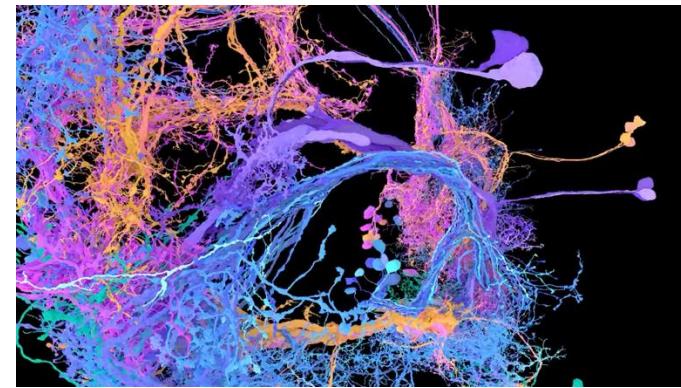
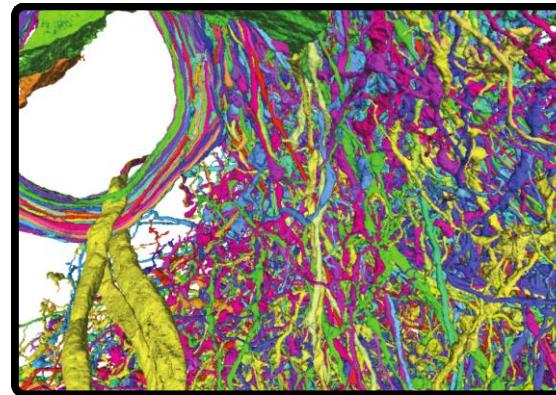
# Karate Club Graph Example



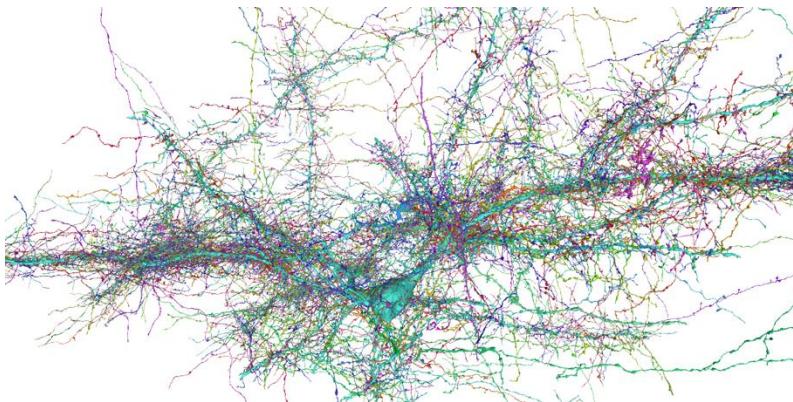
# The Connectome Project



Drosophila Hemibrain Dataset, [Scheffer et al. 2020]

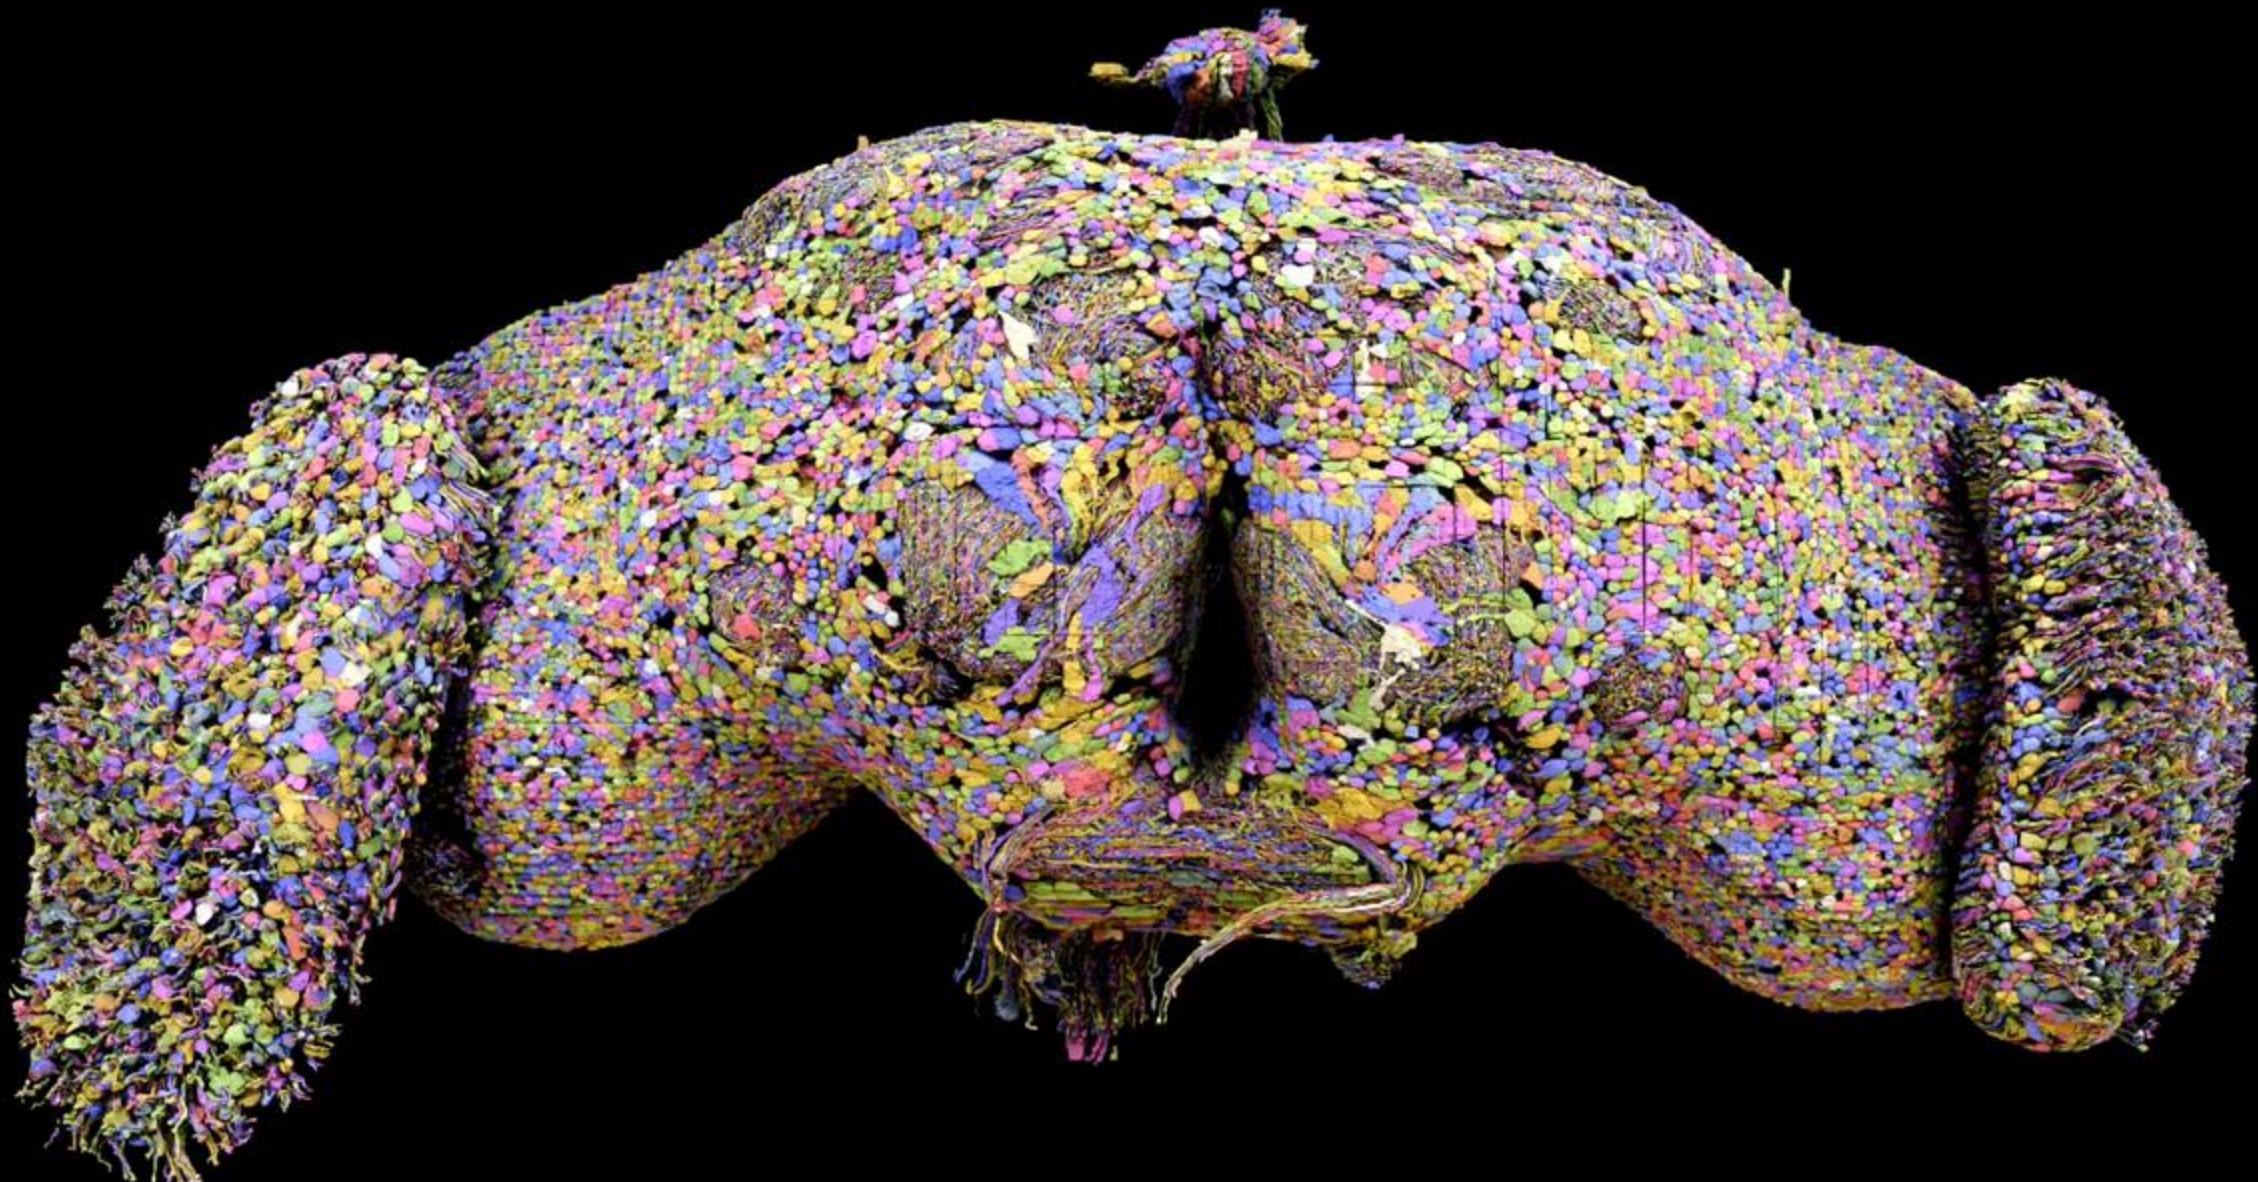


Drosophila Auditory Circuit [Baker et al. 2022]  
Video: Amy Sterling, FlyWire



- Using Arkouda, we can convert connectome datasets with one hundred million rows of JSON objects to distributable HDF5 files in under two hours.
  - Using Arachne, a graph of this size can be queried in seconds to create smaller subgraphs for deeper analysis.

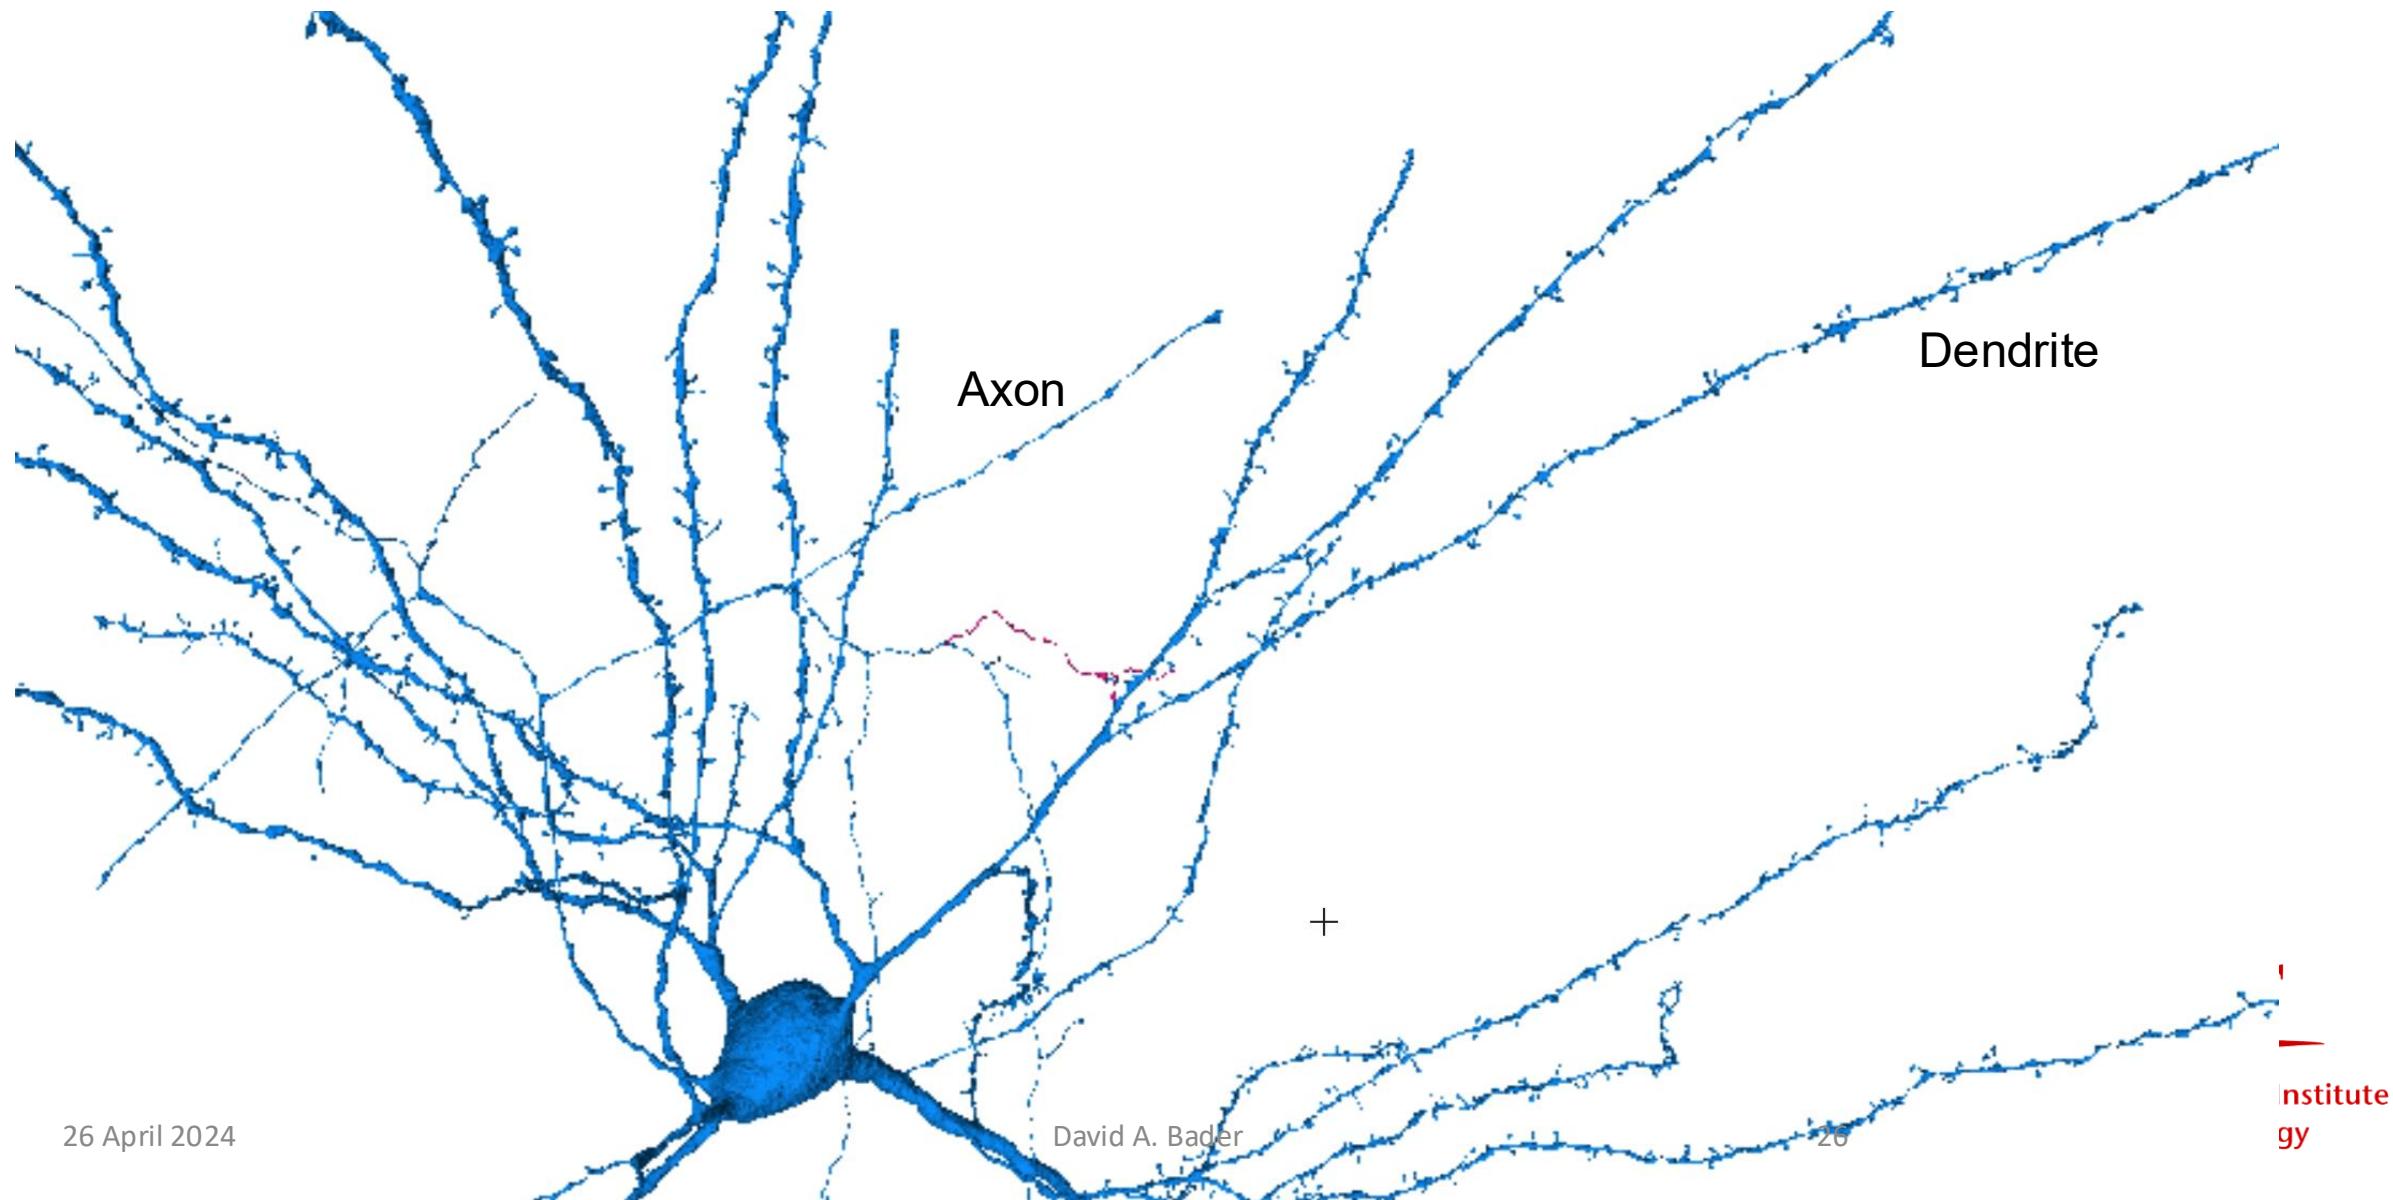
Slide credit: Jakob Troidl, Hanspeter Pfister, Jeff Lichtman (Harvard University)



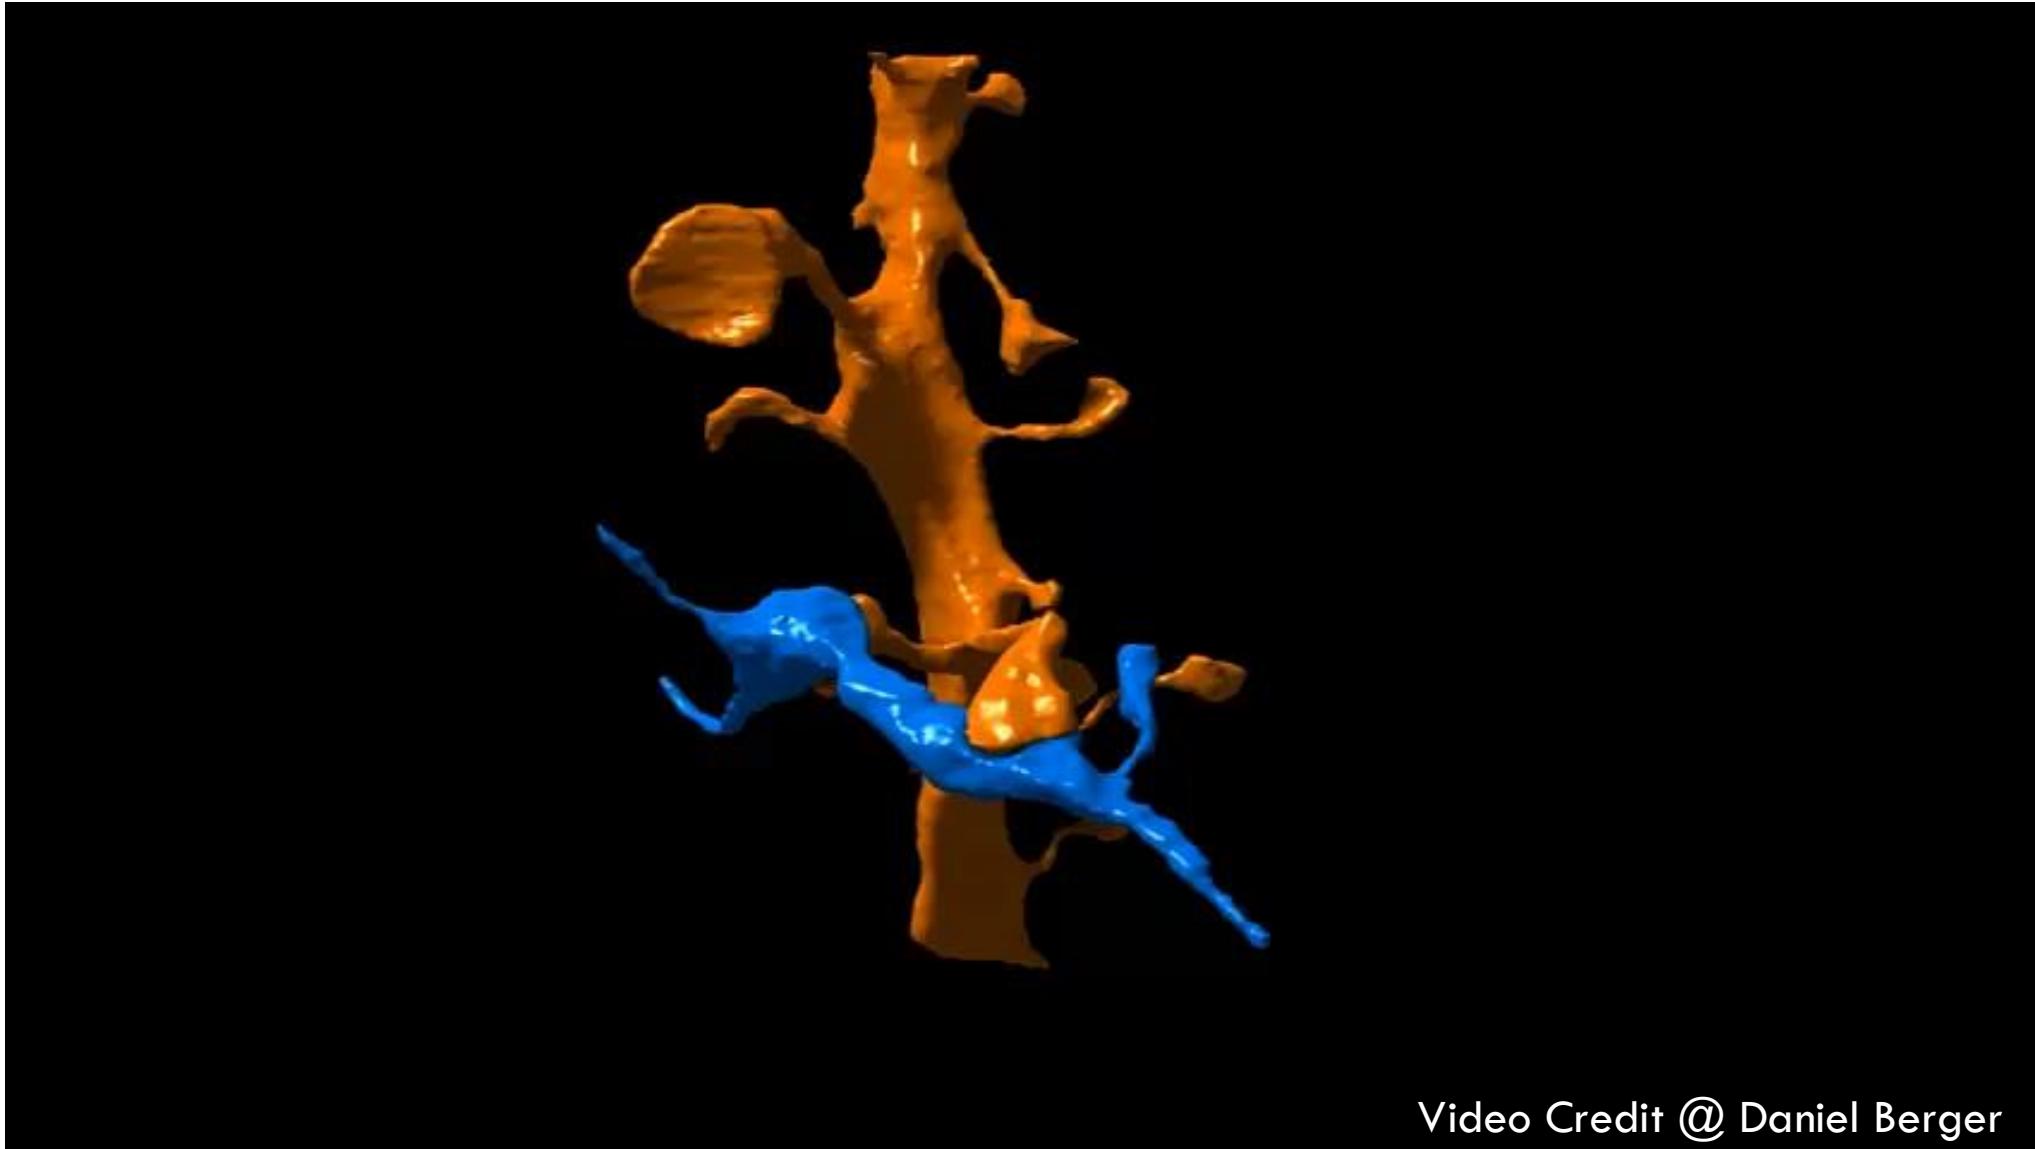
26 April 2024

Dorkenwald et. al 2023 – bioArxiv<sup>25</sup>, Animation by Tyler Sloan

# Wiring Diagram of the Brain

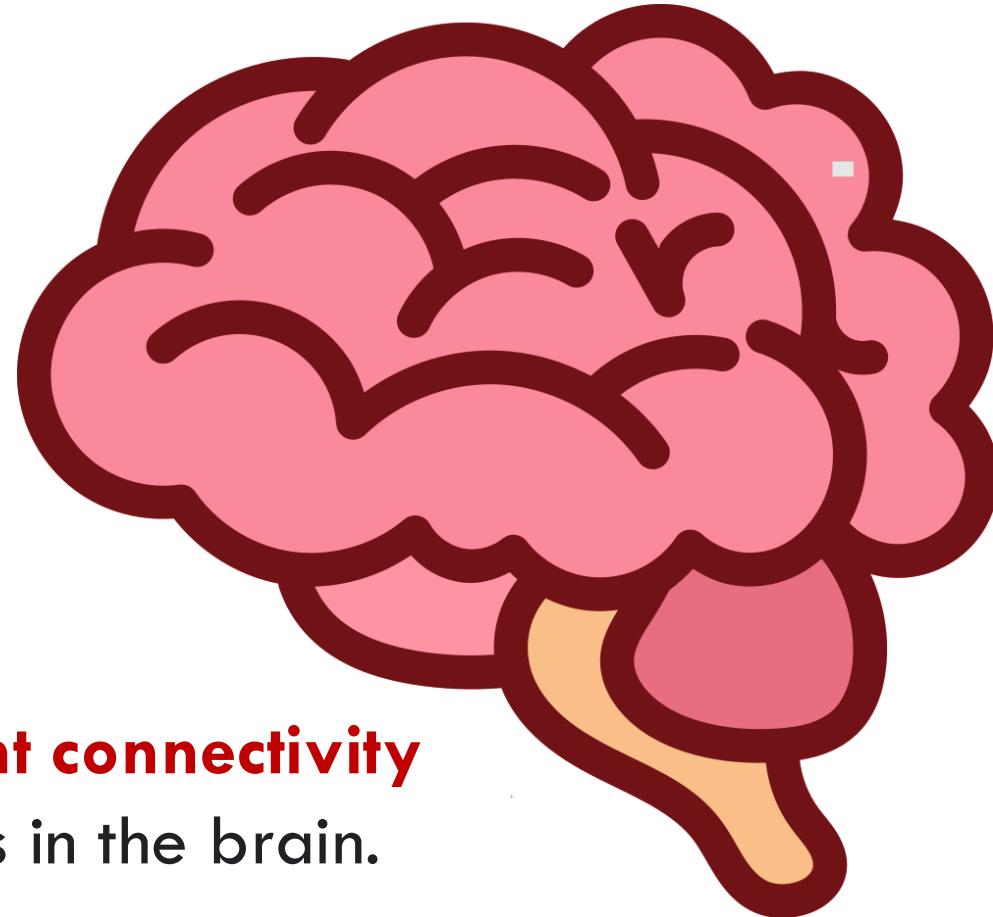


# Spatial Neighborhood Analysis



Video Credit @ Daniel Berger

# Using Graph Analytics to Understand the Brain



Motifs are **recurrent connectivity**  
**patterns** of neurons in the brain.

# Connectome: Requires Exascale Graph Analytics

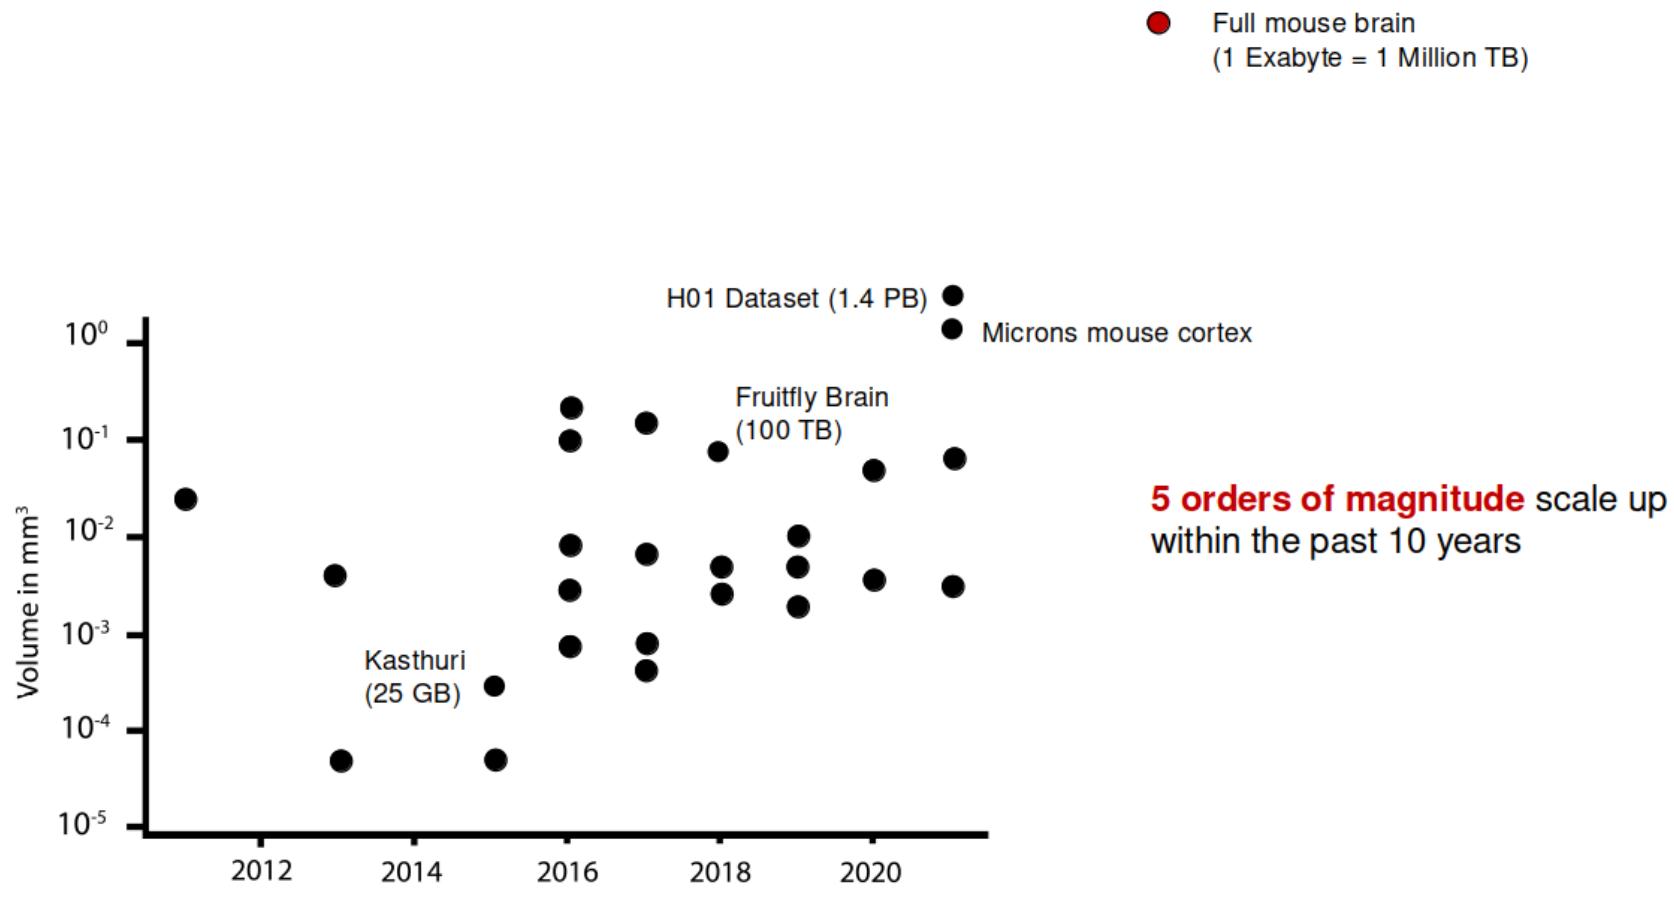
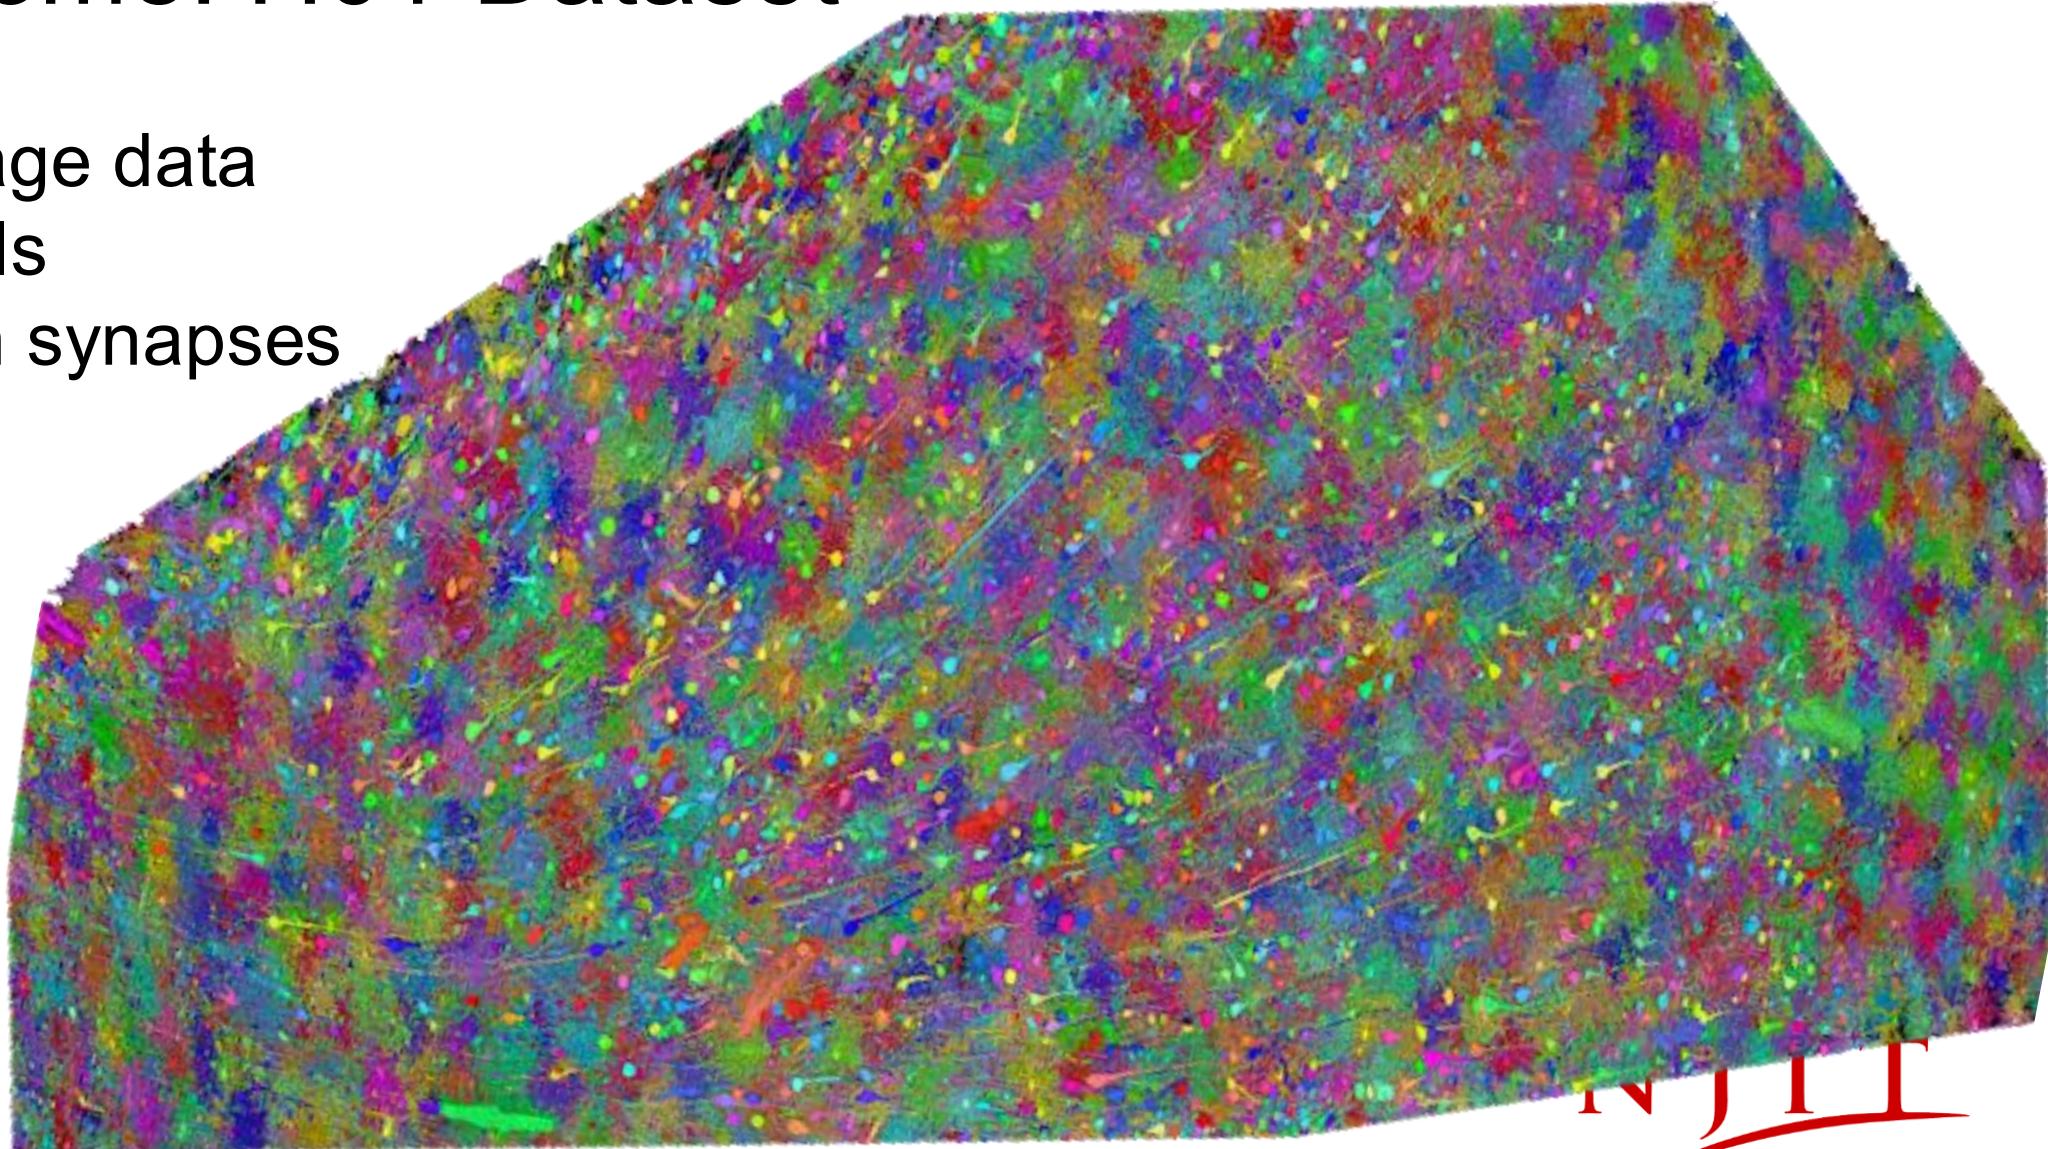


Image credit: Jakob Troidl, Hanspeter Pfister

# Connectome: H01 Dataset

~ 1.4 PB image data  
~ 57,000 cells  
~ 133 Million synapses



# Connectivity motifs

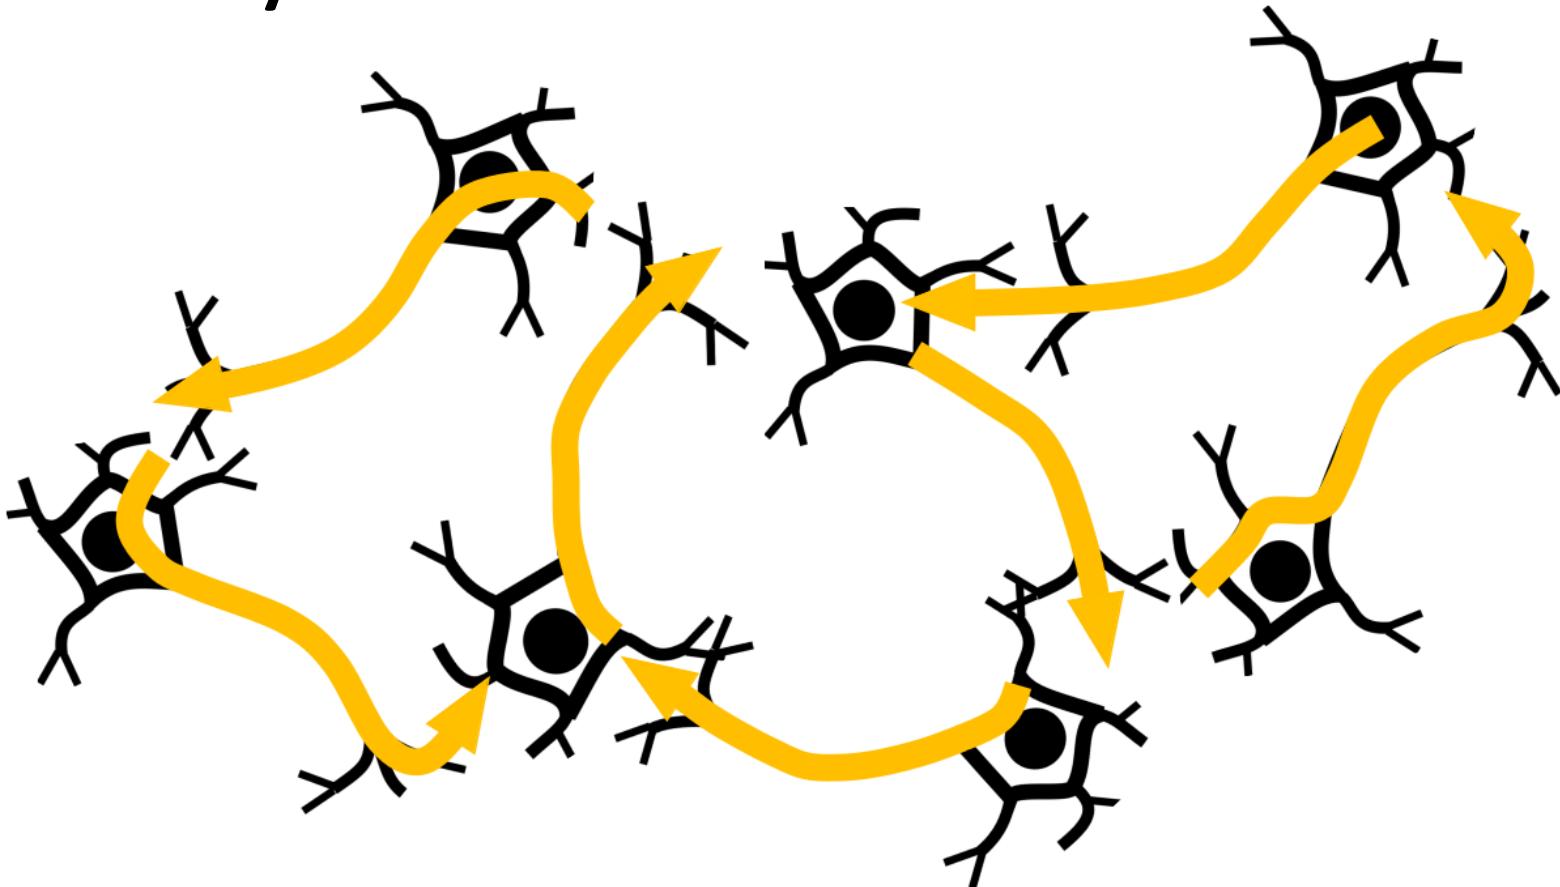


Image Credit: Wikipedia

Slide credit: Jakob Troidl, Hanspeter Pfister

26 April 2024

David A. Bader

31

# Connectivity motifs

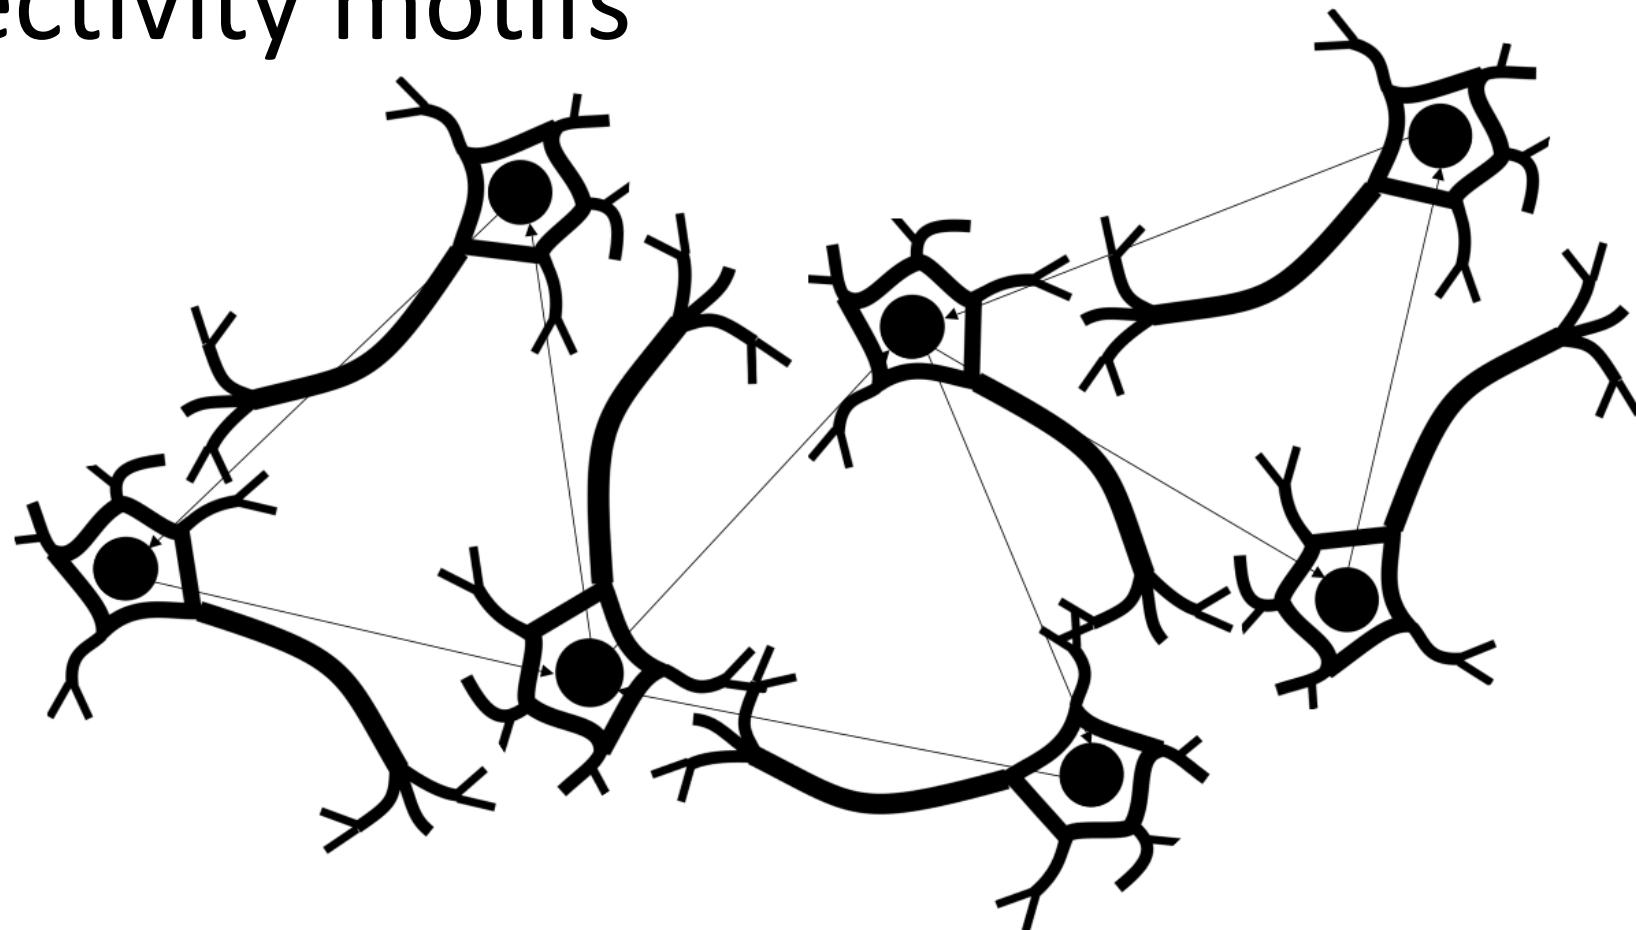


Image Credit: Wikipedia

Slide credit: Jakob Troidl, Hanspeter Pfister

26 April 2024

David A. Bader

32

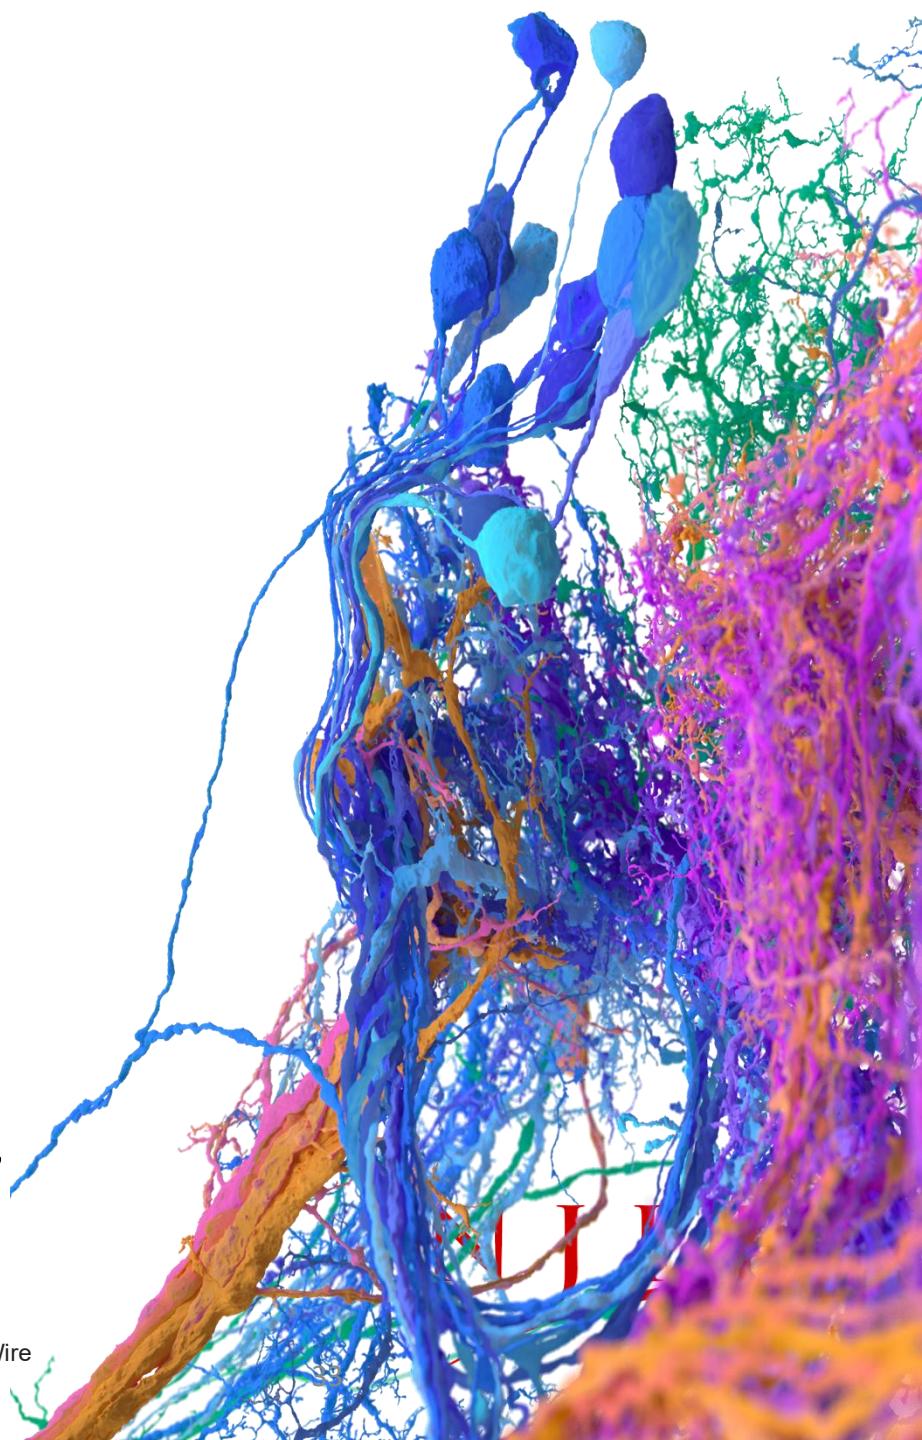
# Motif finding

- **Large Networks.** ~57,000 nodes and ~130 million edges.
- **Expensive Computation.** Verifying the existence of a motif in a larger network is NP-complete.
- **Complex 3D structure.** Neurons span long volumes and form complex branching patterns.
- Algorithms:
  - **Ullmann** (2010) which is a recursive backtracking algorithm for solving the subgraph isomorphism problem
  - **Cordella** (2004) another algorithm based on Ullmann's, VF2, which improves the refinement process using different heuristics and uses significantly less memory.

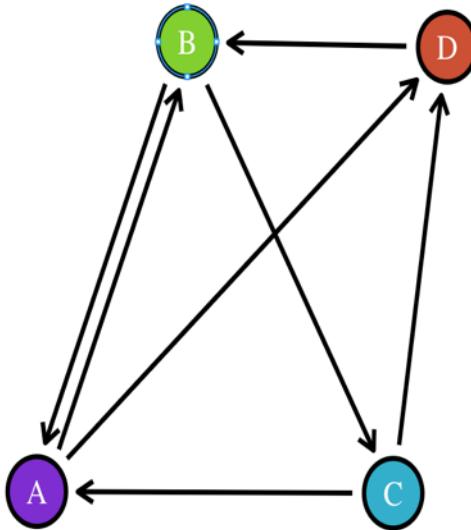
Slide modified from: Jakob Troidl, Hanspeter Pfister

26 April 2024

Image Credit: Amy Sterling @ FlyWire  
David A. Bader



# Neuroscientists with to correlate motif connectivity to neuron morpholog



?

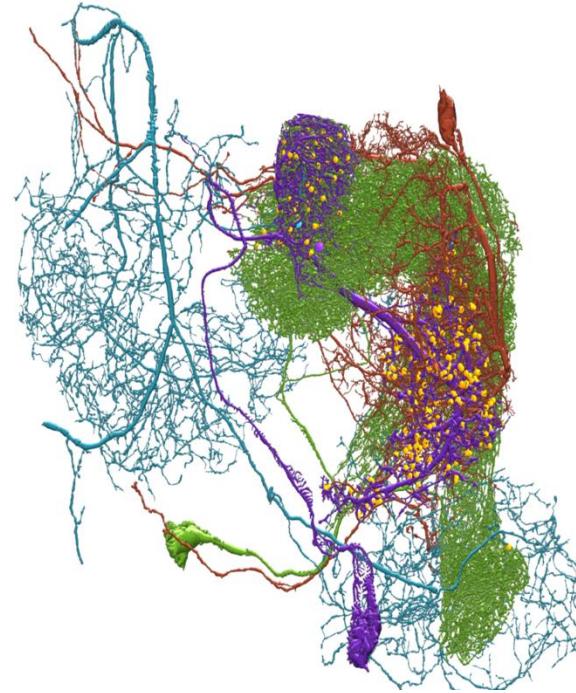
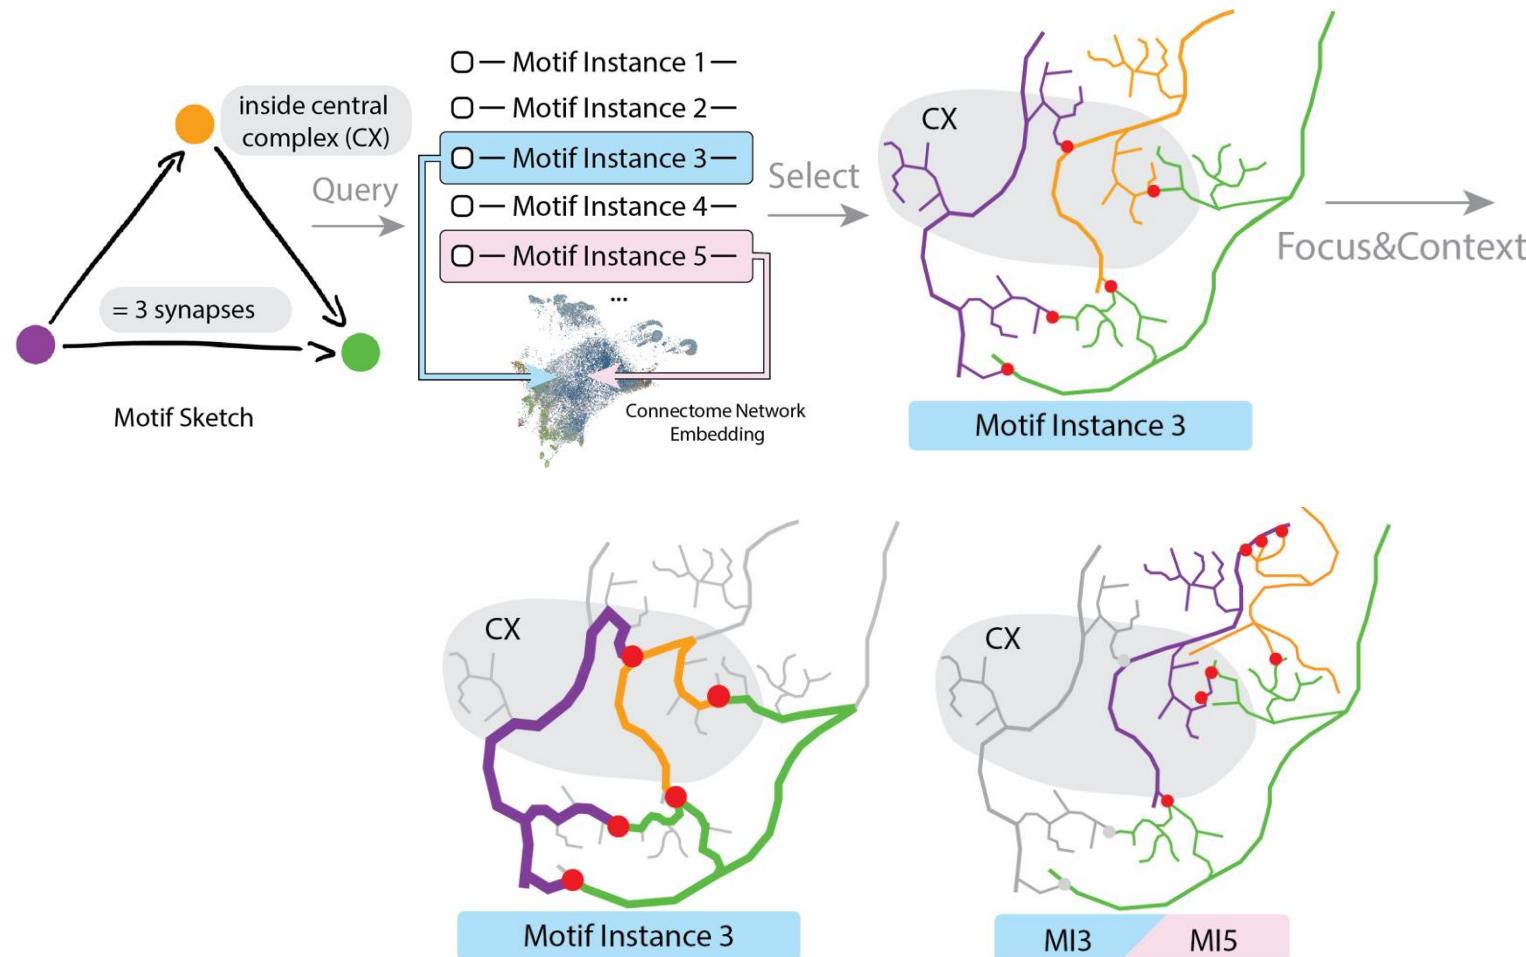


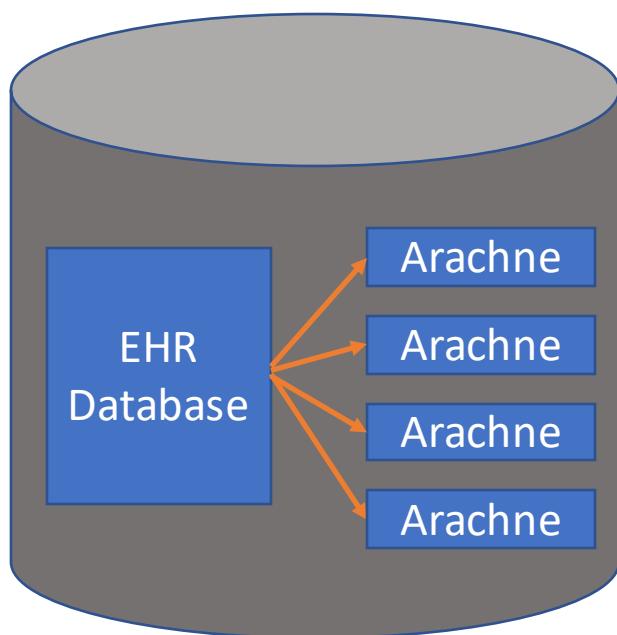
Image credit: Jakob Troidl, Hanspeter Pfister



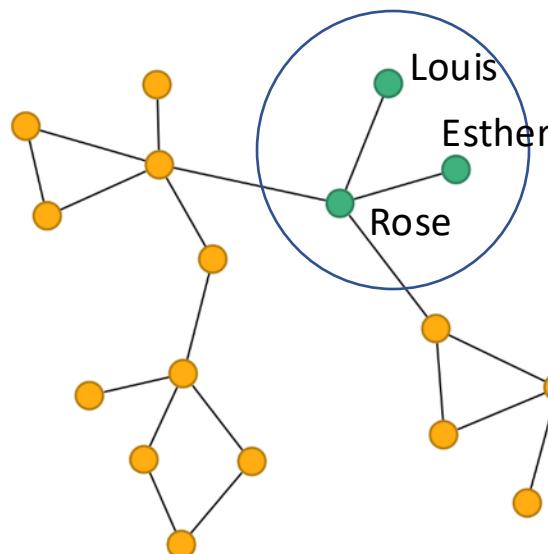
Slide credit: Jakob Troidl, Hanspeter Pfister

# Finding Patterns in Clinical Patient Records

- The adoption of electronic health record (EHR) systems has simultaneously changed clinical practice.
- In data from 2019 and 2021, 96% of general acute care hospitals had adopted EHR\*



Data pre-processed by different tasks on multicore Arachne server

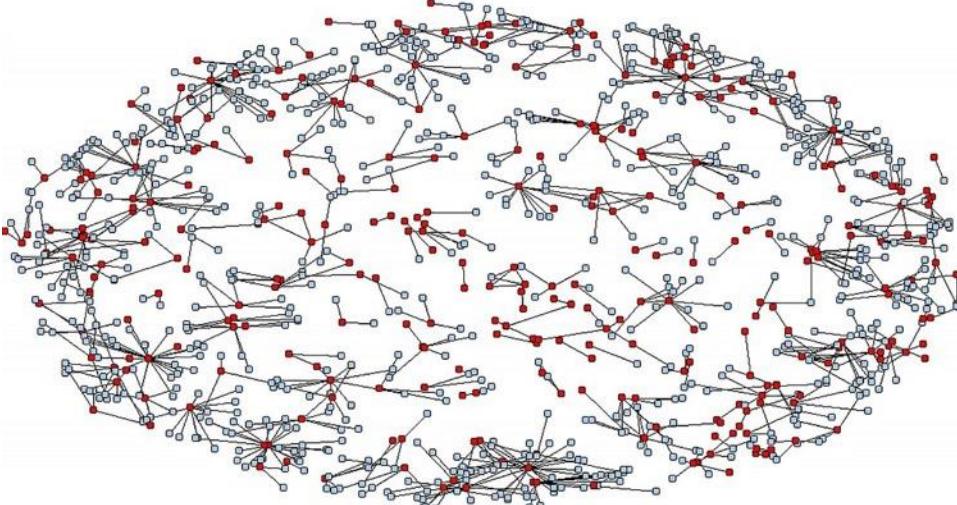


Vertex: Patient  
Edge: Shared clinical features

\* Office of the National Coordinator for Health Information Technology.  
Adoption of Electronic Health Records by Hospital Service Type 2019-2021,  
Health IT Quick Stat #60. April 2022.

- Utilize **community detection** algorithms to identify groups of vertices.
- These communities may correspond to subpopulations of patients with similar clinical characteristics or disease trajectories.

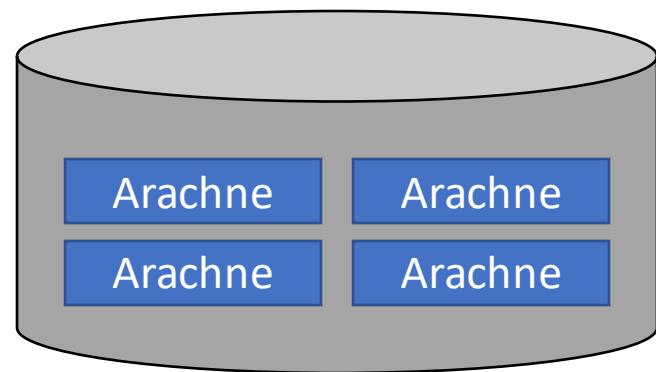
# Contact Tracing Networks (COVID, HIV, etc.)



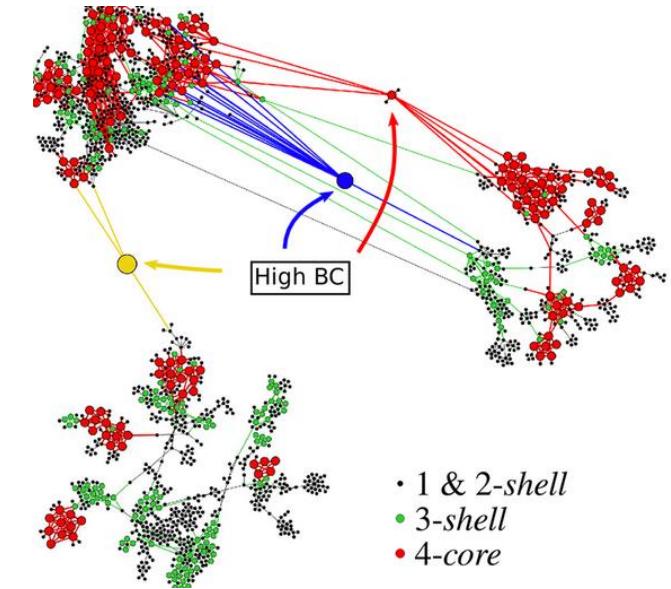
[Cushman, 2020]

Usual characteristics of graphs:

- million+ edges and vertices give rise to graphs that take more than 2GB to store in memory
- storage of in-between steps of certain algorithms like Jaccard coefficients can exceed 512GB



David A. Bader



[Serafino, Monteiro, et al. 2022]

Current:

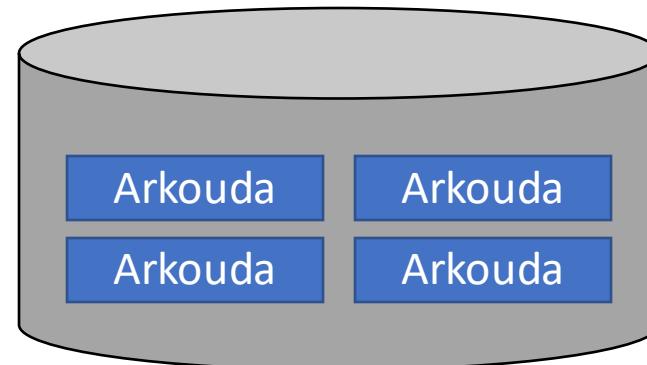
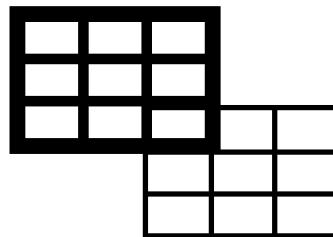
- k-truss
- triangle counting
- triangle centrality

Planned:

- re-implementation of bc
- k-core

# Population Health Data Analysis

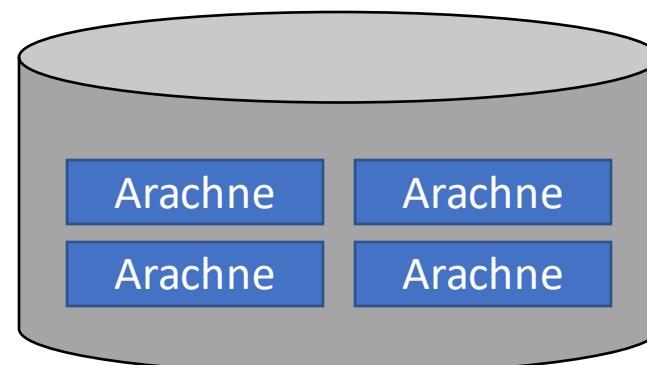
Tabular Data



Analysis on terabytes of tabular data!

- Arachne works with Arkouda as an add-on for graph analysis.
- Data can be taken from Arkouda and created into graphs by specifying columns as edge sources and destinations.

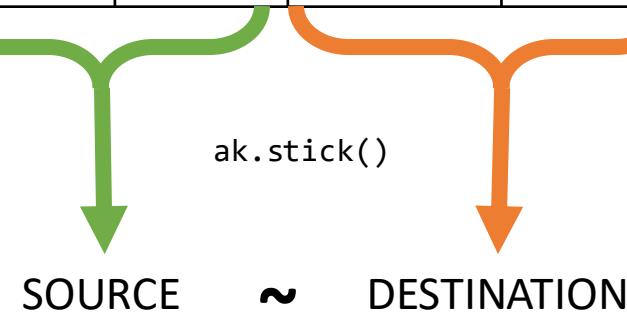
build (property) graphs



- run graph kernels
- planned: property graph algorithms

# The Arkouda-Arachne Netflow Data Pipeline

IPV4_SRC_ADDR	L4_SRC_PORT	IPV4_DST_ADDR	L4_DST_PORT	PROTOCOL	L7_PROTO	IN_BYTES	OUT_BYTES	IN_PKTS	OUT_PKTS	TCP_FLAGS	FLOW_DURATION_MS	Attack
192.168.100.6	52670	192.168.100.1	53	17	5.21	71	126	1	1	0	4294966	Benign
192.168.100.6	49160	192.168.100.149	444	6	0	217753000	199100	4521	4049	24	4176249	Theft
192.168.100.46	3456	192.168.100.5	80	17	0	8508021	8918372	9086	9086	0	4175916	Benign
192.168.100.3	80	192.168.100.55	8080	6	7	8442138	9013406	9086	9086	0	4175916	Benign



IPV4 source addresses and ports together make up the source vertex of the edge and respectively the same columns for the destination vertex of the edge.

26 April 2024

integer id gen  
ak.GroupBy()  
ak.groupby.broadcast()

IPV4_SRC	IPV4_SRC_id	IPV4_DST	IPV4_DST_id
192.168.100.6:52670	3473	192.168.100.1:53	3455
192.168.100.6:49160	4234	192.168.100.149:444	3233

David A. Bader

39



# Back-End Storage and Querying

Relationships	SRC	DST	IND	SETS
	0	1	0	{1}
	0	2	1	{0}
	1	0	2	{1}
	1	2	3	{0,2}
	2	2	4	{0}

Edge Properties

INTS	REALS
REF	REF

- Relationships are stored in sets per edge. User specifies a query, and we search the edge set in a massively parallel manner, probing the sets in amortized constant time.
- Properties are stored split by type and for each edge-type pair that exists, we store an associative domain where we extract the data by simply doing an access `edge_prop[col_id]`.

R-MAP	IND
benign	0
theft	1
DDoS	2

INT

1	2
VAL	VAL

REAL

0
VAL

COL-MAP	IND	TYPE
L7_PROTO	0	FLOAT
IN_BYTES	1	INT
OUT_BYTES	2	INT

- All searching is guaranteed to be  $O(m/p)$  since it only involves iterating over the edge set in parallel with each processor.
- 26 April 2024

- All REFS point to arrays that are of the specified column type and that store the key-value pairs for column identifier to data.
- Query hits are returned in a Boolean array specified which edges matched.

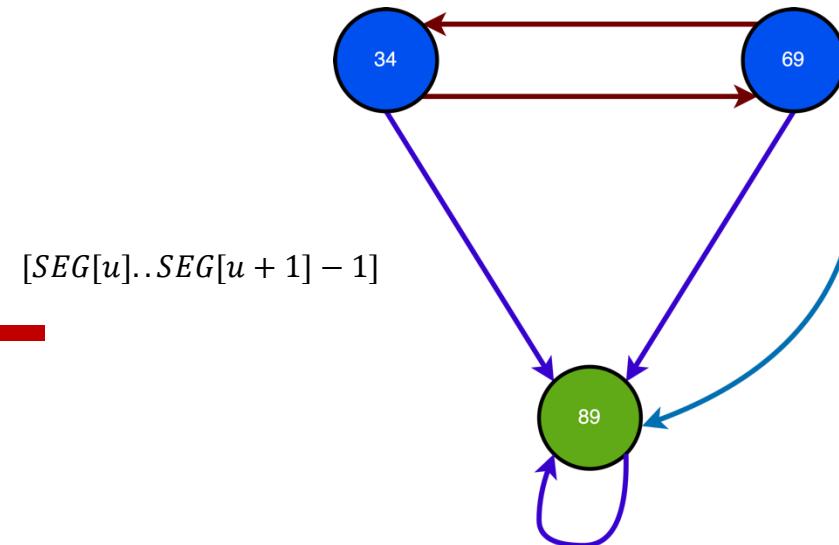
# Code Example for Python Scripts & Jupyter

```
1. import arkouda as ak
2. import arachne as ar
3.
4. ## Get src and dst from input file.
5. graph = ar.PropGraph()
6. graph.add_edges_from(src,dst)
7.
8. ## Generate relationships_df and edge_properties_df from input file.
9. graph.add_edge_relationships(relationships_df)
10. graph.add_edge_properties(edge_properties_df)
11.
12.## User generates relationships_to_find and property query.
13.returned_edges_rel = graph.query_relationships(relationships_to_find)
14.returned_edges_prop = graph.query_edge_properties("COLUMN", 67, ">")
15.
16.returned_edges = ak.intersect1d(returned_edges_rel, returned_edges_prop)
17.subgraph_src = returned_edges[0]
18.subgraph_dst = returned_edges[1]
19.
20.subgraph = ar.Graph()
21.subgraph.add_edges_from(subgraph_src, subgraph_dst)
22.bfs = ar.bfs_layers(subgraph)
23.cc = ar.connected_components(subgraph)
24.tris = ar.triangles(subgraph)
25.squares = ar.squares(subgraph)
26.## And more!!!!
```

- Line 6 input is generated from input files from types such as HDF5, CSV, Parquet, etc.
- Lines 9 and 10 input is generated from input files as well.
- Lines 13 and 14 relationships and properties to find are generated by the user.
- Lines 16-19 use Arkouda operations and slicing to get the edges that are returned by both queries.
- Lines 20 and 21 create a new Arachne simple graph with the returned edges of the queries.
- Lines 22-25 run some of the other algorithms available in Arachne!
- Arachne can also return arrays composed of the edges which can be converted to Python lists or NumPy arrays so they can be loaded into NetworkX for further analysis!

# Arachne DI Data Structure [Du et al. 2021]

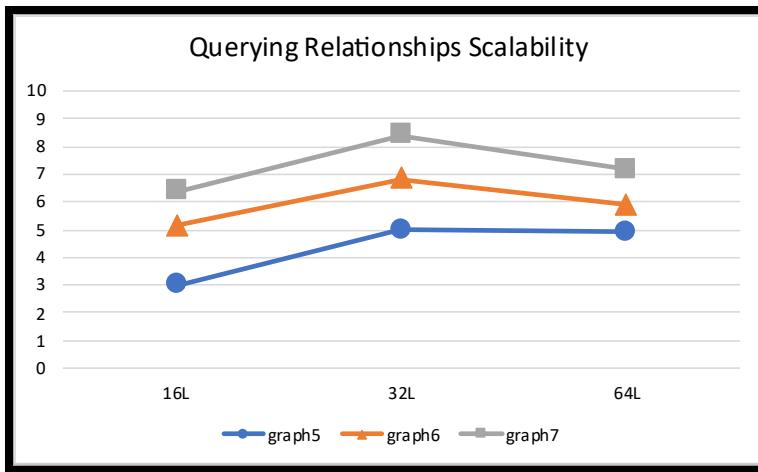
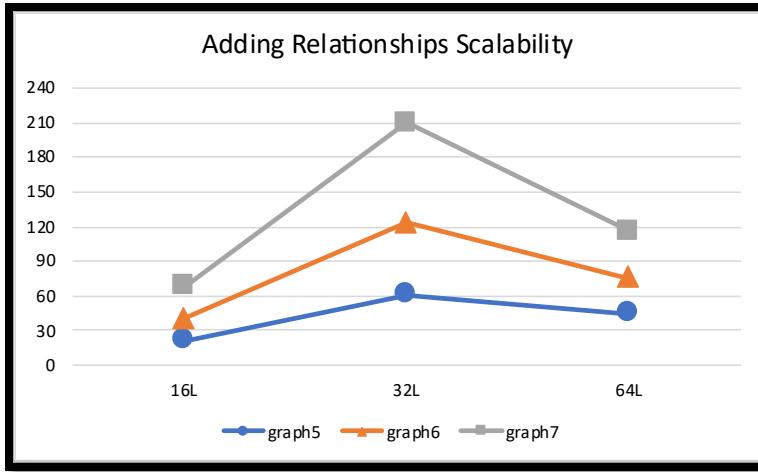
Vertices			Edges		
IND	MAP	SEG	IND	SRC	DST
0	34	0	0	0	1
1	69	2		0	2
2	89	4	2	1	0
3		5		1	2
			4	2	2



- Allows for simple, compact, **distributable** storage of vertex and edge sets.
- Given an edge index, all vertices that make up that edge are found in **constant time**, avoiding a binary search into SRC (CSR offsets index equivalent).
- MAP allows explicitly storing original vertex labels, returning original graph involves creating arrays and place values of SRC[MAP] and DST[MAP] into new arrays.

# Property Graph Results

	$n$	$m$	
graph5	864,648,454	1,000,000,000	●
graph6	2,161,664,289	2,500,000,000	▲
graph7	1,408,892,291	5,000,000,000	■



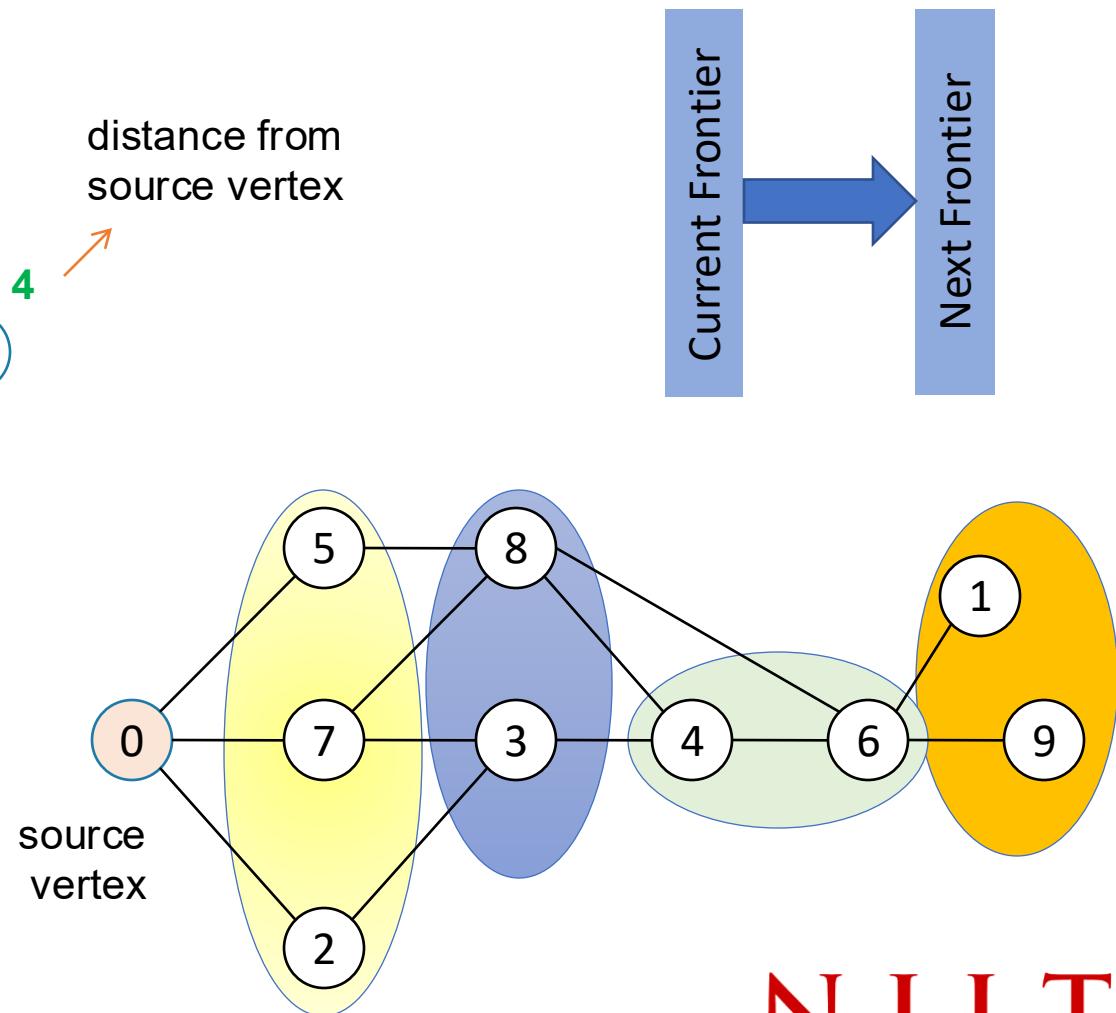
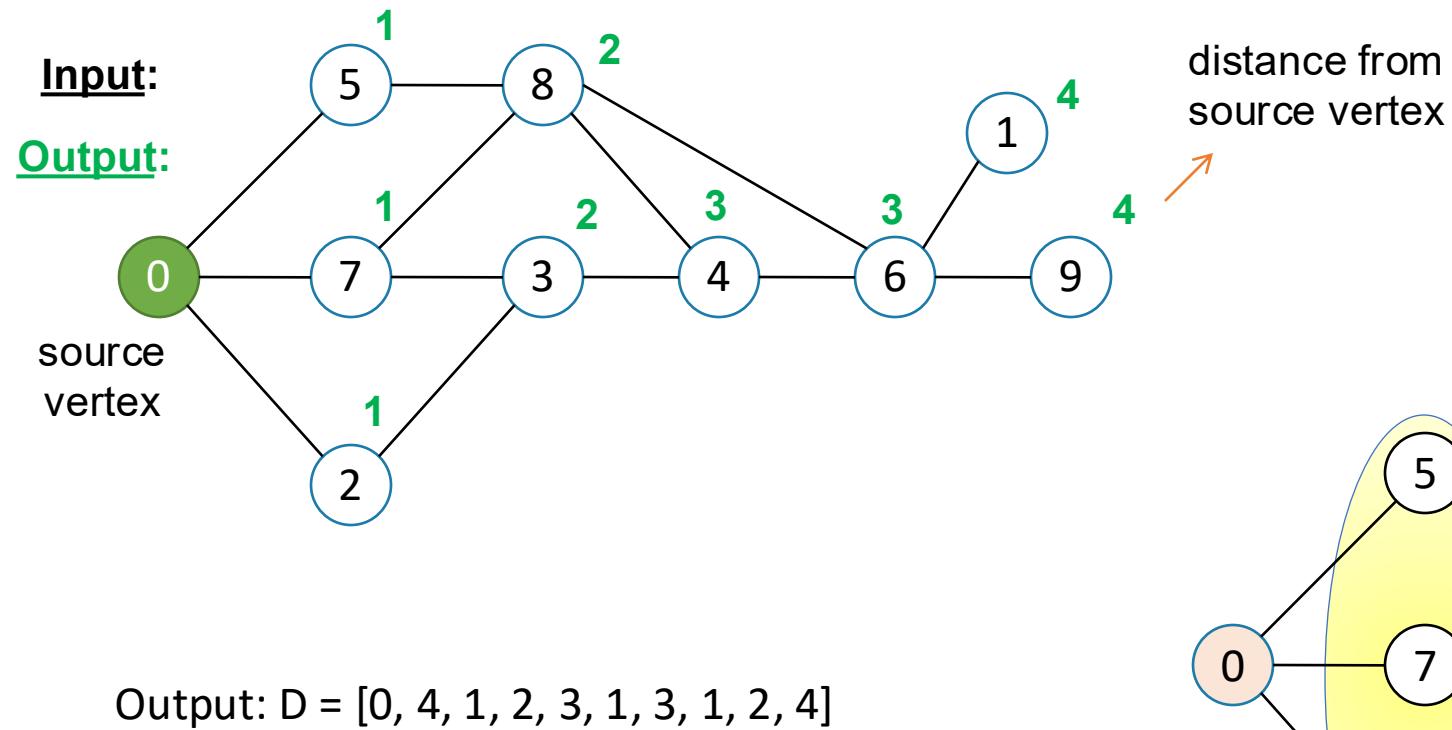
- Experiments conducted on a cluster where each compute node was composed of 128 cores (64 per AMD EPYC 7713 CPUs), 1TB DDR4 RAM, and an Infiniband HDR 200 GB/s node interconnect.
  - At time of results, some nodes had performance issues, hence the weird elbows.
  - Fifty random relationships were made and randomly assigned to edge indices meaning some edges could be picked more than once and some none at all.
- Querying involved searching for the edges that included three of the fifty relationships, each list performed a set and operation with the search space.

**Takeaway: Building a graph of five billion edges takes under 60 seconds, running ETL to insert relationships takes under 4 minutes, and querying it under 10 seconds.**

# Graph Algorithms in Arachne

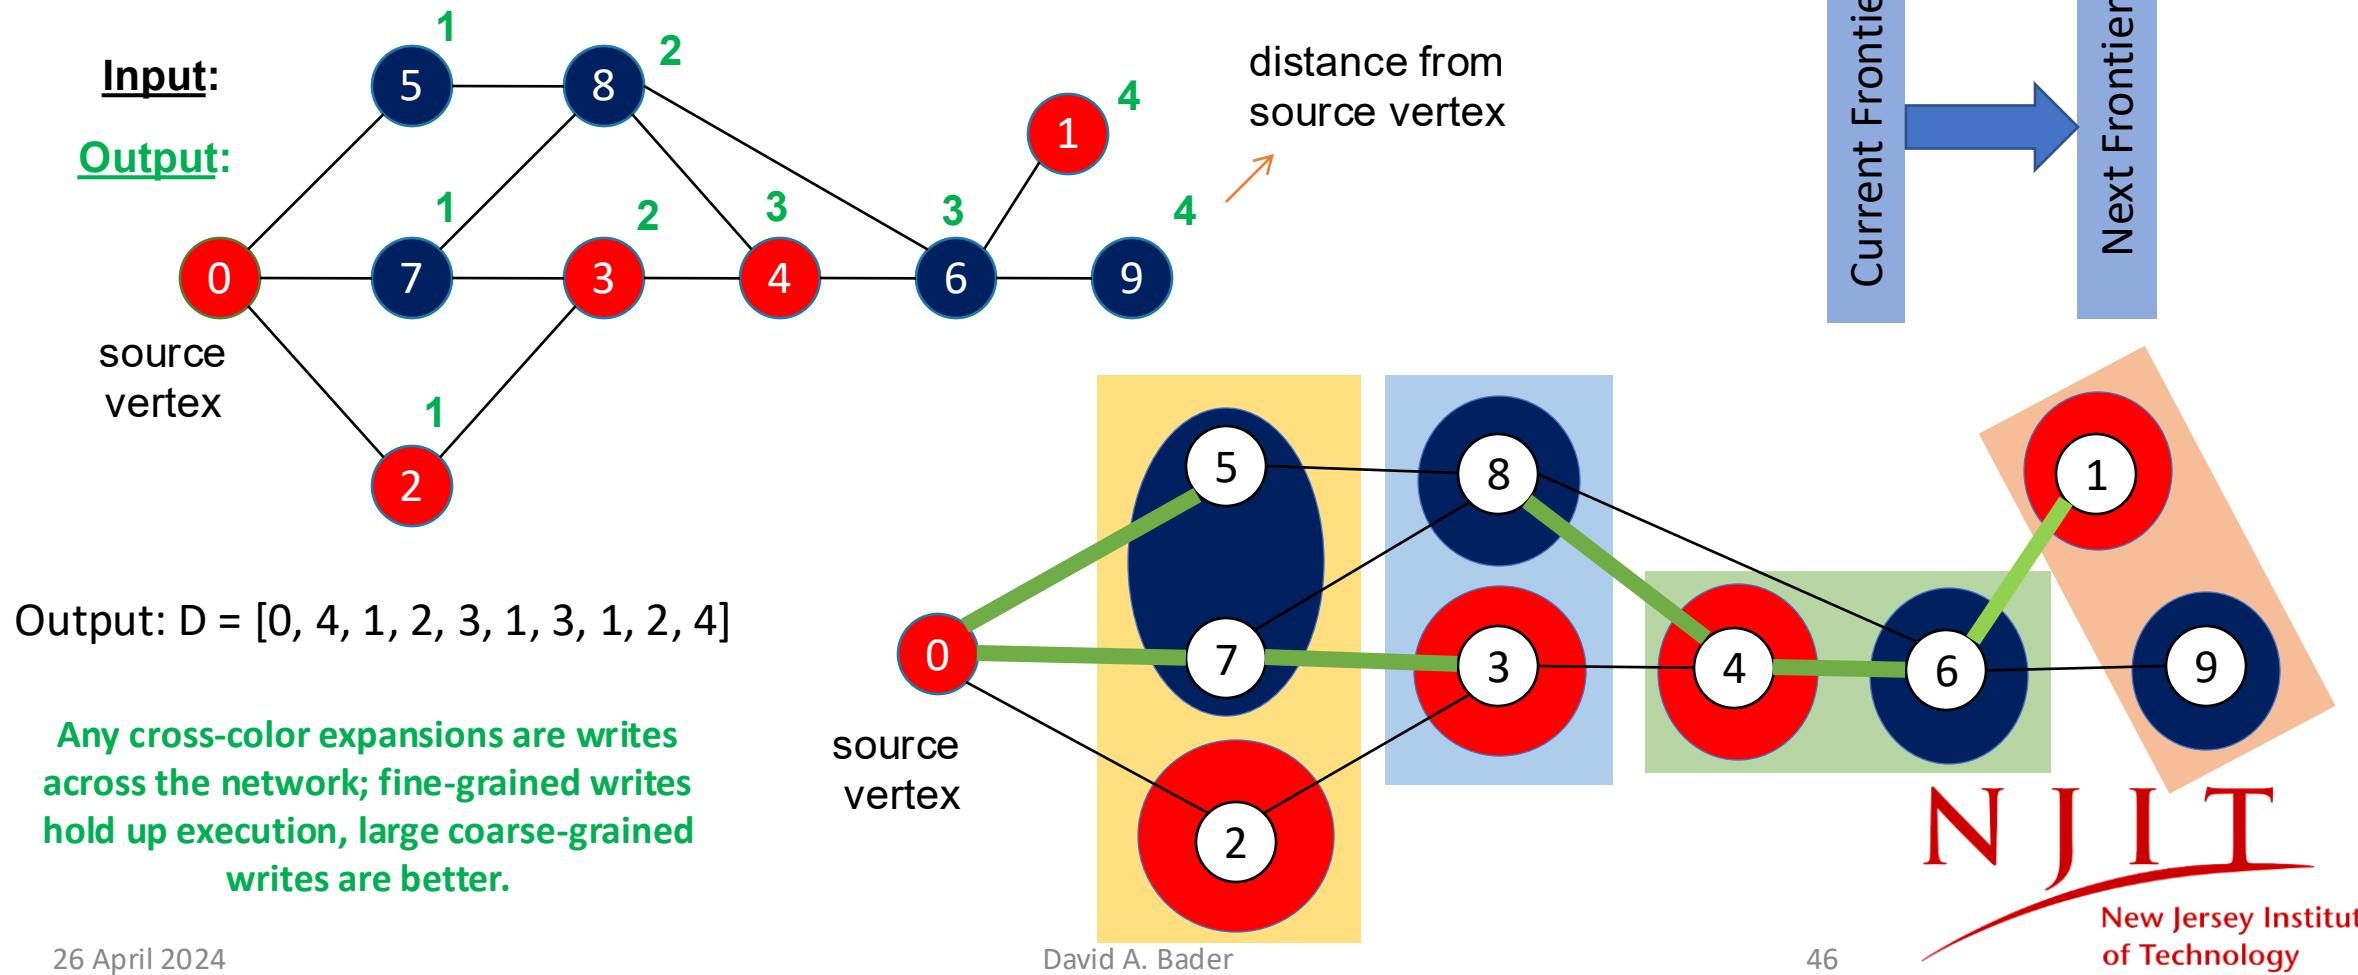
- **Breadth-first search (BFS)** [Du, Alvarado Rodriguez, Bader 2021]  
Returns an array of size  $n$  with how many hops away some vertex  $v$  is from an initial vertex  $u$ .
- **Connected components**  
Returns an array of size  $n$  where all vertices who belong to the same component have the same value  $x$ . The value of  $x$  is the label of the largest vertex in the component.
- **Triangle counting** [Du, Alvarado Rodriguez, Patchett, Bader 2021]  
Returns the number of triangles in a graph.
- **Truss Analytics** [Du, Patchett, Bader 2021][Du, Patchett, Alvarado Rodriguez, Li, Bader 2022]  
K-truss returns every edge in the truss where each edge must be a part of  $k - 2$  triangles that are made up of nodes in that truss. Max truss returns the maximum  $k$ . Truss decomposition returns the maximum  $k$  for each edge.
- **Square counting** [Burkhardt, Harris 2023]  
Returns the number of four-cycles in the graph.
- **Triangle centrality** [Patchett, Du, Bader 2022][Patchett, 2022]  
Returns an array of size  $n$  with the proportion of triangles centered at a vertex  $v$ .
- **Subgraph isomorphism** [Dindoost, Bader, 2023, in progress]  
Finds instances of a pattern in a larger graph.

# Shared-Memory Parallel Breadth-First Search



# Distributed-Memory Parallel Breadth-First Search

Assume our edge list is split down the middle, then the neighborhood of some vertices will live on one compute node while the rest live on another compute node.



# Breadth-First Search Communication Volume Results

locale	get			put		
	di	di-norev	di-agg-ls	di	di-norev	di-agg-ls
0	15672640	7873842	639827	5629422	2749193	138070
1	15834332	7939017	687156	1952226	1016946	127936
2	15715554	7722659	226754	1942839	962031	45217
3	15817879	7723971	226880	1951313	962201	45060
4	15964559	7724880	226691	1961552	962199	51217
5	15739226	7726504	230024	1940688	962439	52714
6	15569450	7727678	229096	1925536	962680	51977
7	15341933	7736094	225083	1904757	963418	48413

di: 84 seconds

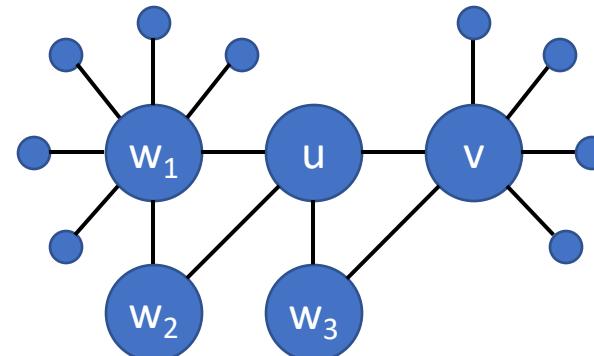
di-agg-ls: 3.36 seconmds

delaunayn20 is a graph with 3 million edges and a large diameter

**Takeaway:** Aggregating writes drastically reduces communication volumes, improving performance, all done easily through Chapel by adapting aggregators for different uses.

# Minimum Search Triangle Counting

- Given an edge  $(u, v)$  we assume that  $|Adj(u)| \leq |Adj(v)|$ .
- Then, for  $\forall w \in Adj(u)$  we spawn  $|Adj(u)| - 1$  parallel threads to check if we can form a complete triangle with  $(u, v, w)$ .
- If  $|Adj(w)| < |Adj(v)|$  we will check if  $v \in Adj(w)$ , else, we check if  $w \in Adj(v)$ .



Adj(x)	Value
Adj(u)	4
Adj(v)	6
Adj(w1)	7
Adj(w2)	2
Adj(w3)	2

Thread  $w_1$ : search for  $w_1$  in  $Adj(v)$ , no match, kill.

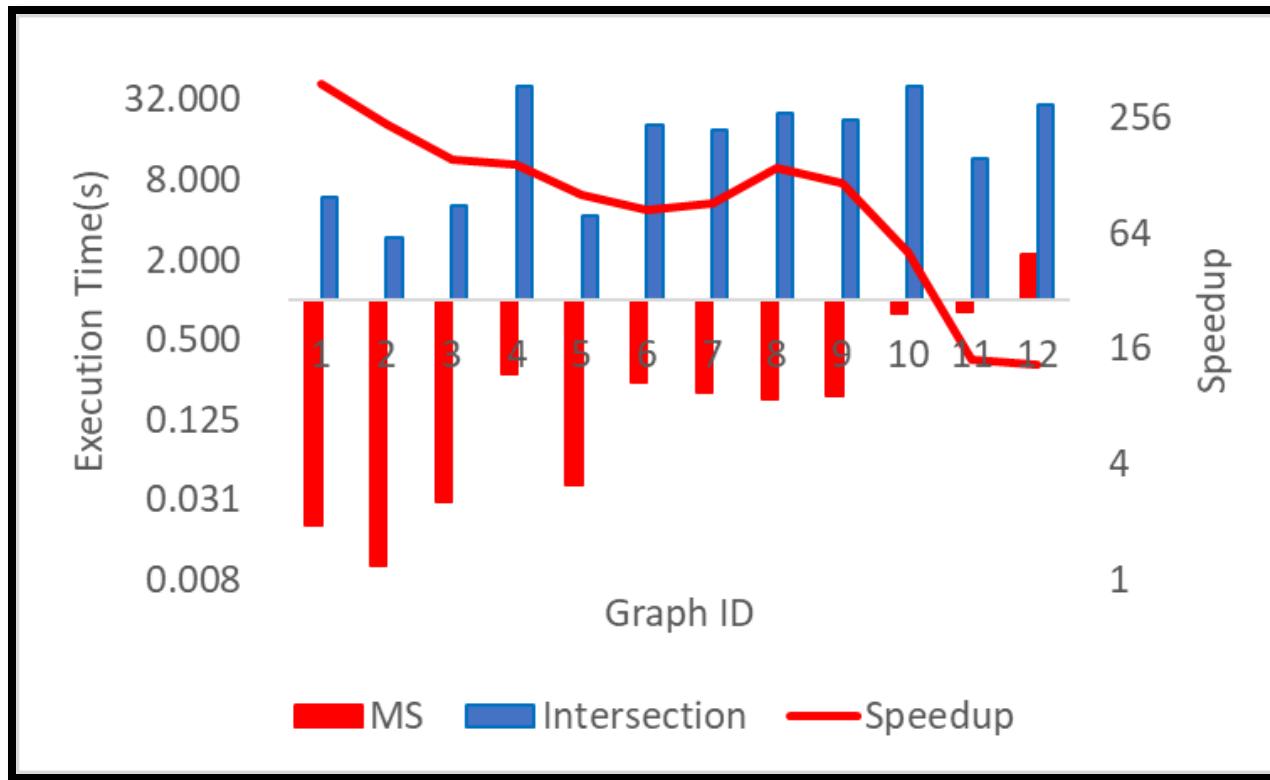
Thread  $w_2$ : search for  $v$  in  $Adj(w_2)$ , no match, kill.

Thread  $w_3$ : search for  $v$  in  $Adj(w_3)$ , match! Increment count.

# Minimum Search Triangle Counting Operation Count Comparison

- Assume  $|Adj_u| < |Adj_v|$  and we spawn threads for every  $w \in |Adj_u|$ 
  - Minimum search:  $\max_{w \in Adj_u} \log_2(\min(|Adj_w|, |Adj_v|))$
  - List Intersection:  $\log_2(|Adj_v|)$
- Say we have the following information for our vertices:
  - $|Adj_u| = 4$  and  $|Adj_v| = 1024$
  - For every  $w$  in  $Adj_u$ ,  $|Adj_w| \leq 8$
- **List intersection:** 4 threads amounting to  $\lceil \log_2 1024 \rceil = 10$  operations each.
- **Minimum search:** 4 threads amounting to  $\lceil \log_2 8 \rceil = 3$  operations each.

# Triangle Counting Results



- Our method outperforms with the Conte method with a highest speedup of 385.8 and an average speedup of 128.

**Takeaway:** Truss decomposition with minimum search triangle counting outperforms a C++ method coded with OpenMP, with SSE-Acceleration, binary searching on adjacency list, and no atomic operations.

Graphs used were a variety of real-world graphs available for view in paper.

# Major Contributions

- **Arachne**, a large-scale graph analysis framework, extends Arkouda for productive graph analysis. **Arachne** is built on a novel sparse graph data structure.
- **Arachne** leverages **productivity** through Python with **high performance** backed by Chapel.
- **Arachne**, Arkouda, Chapel are all open-source.
  - <https://github.com/Bears-R-Us/arkouda-njit>
  - <https://github.com/Bears-R-Us/arkouda>
  - <https://github.com/chapel-lang/chapel>
- Experimental results on real-world and synthetic graphs demonstrate that **Arachne** works for graphs with billions of edges.

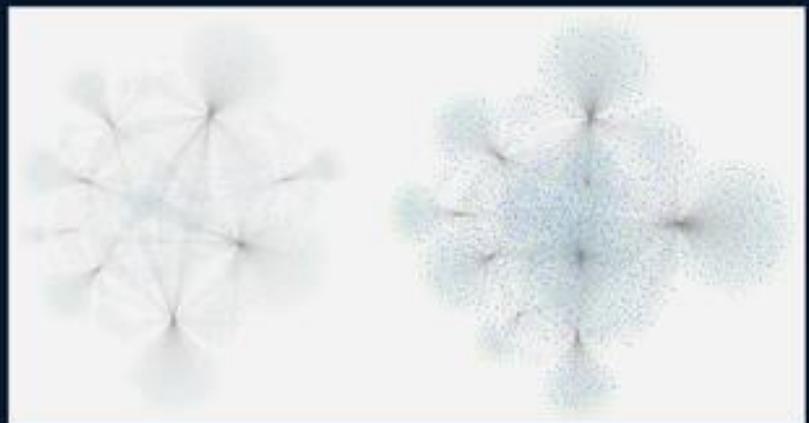
# Publications

- **Oliver Alvarado Rodriguez, Zhihui Du, Joseph Patchett, Fuhuan Li, David Bader (2022). Arachne: An Arkouda Package for Large-Scale Graph Analytics. IEEE HPEC.**
- Oliver Alvarado Rodriguez, Fernando Vera Buschmann, Zhihui Du, David Bader (2023). Property Graphs in Arachne. IEEE HPEC.
- Soroush Vahidi, Baruch Schieber, Zhihui Du, David Bader (2023). Parallel Longest Common SubSequence Analysis In Chapel. IEEE HPEC.
- Joseph Patchett, Zhihui Du, Fuhuan Li, David Bader (2022). Triangle Centrality in Arkouda. IEEE HPEC.
- Zhihui Du, Oliver Alvarado Rodriguez, David Bader (2021). Large Scale String Analytics In Arkouda. IEEE HPEC.
- Zhihui Du, Oliver Alvarado Rodriguez, David Bader (2021). Enabling Exploratory Large Scale Graph Analytics through Arkouda. IEEE HPEC.
- Joseph Patchett, Zhihui Du, David Bader (2021). K-Truss Implementation in Arkouda (Extended Abstract). IEEE HPEC.
- Zhihui Du, Oliver Alvarado Rodriguez, Joseph Patchett, David Bader (2021). Interactive Graph Stream Analytics in Arkouda. Algorithms.
- Zhihui Du, Oliver Alvarado Rodriguez, David A. Bader, Michael Merrill, William Reus (2021). Exploratory Large Scale Graph Analytics in Arkouda. CHI UW.

# Conclusions & Further Work

- We can design and develop high performance graph analysis algorithms using Arkouda/Chapel quickly and efficiently.
- We plan to work on optimizing all current methods to work as efficiently as possible in single locale and multi locale environments.
- We plan to implement new novel algorithms such as stringology, a communication-efficient triangle counting, large-scale community detection, and machine learning.

# MASSIVE GRAPH ANALYTICS

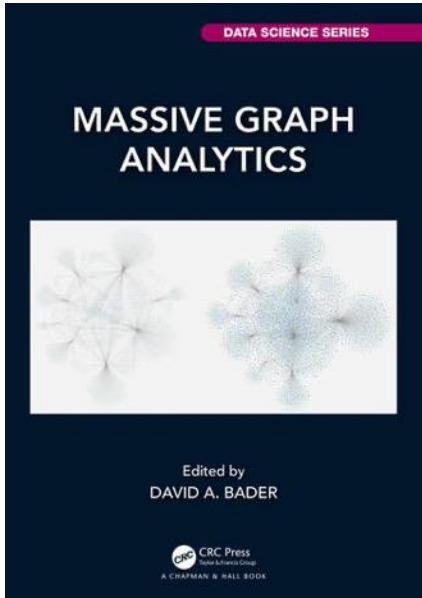


Edited by  
DAVID A. BADER



Massive Graph Analytics  
Edited By David A. Bader  
Copyright Year 2022  
ISBN 9780367464127  
Published July 20, 2022 by Chapman & Hall  
616 Pages 207 B/W Illustrations

# Chapters



## Algorithms: Search and Paths

### **A Work-Efficient Parallel Breadth-First Search Algorithm (or How to Cope With the Nondeterminism of Reducers)**

*Charles E. Leiserson and Tao B. Schardl*

### **Multi-Objective Shortest Paths**

*Stephan Erb, Moritz Kobitzsch, Lawrence Mandow, and Peter Sanders*

## Algorithms: Structure

### **Multicore Algorithms for Graph Connectivity Problems**

*George M. Slota, Sivasankaran Rajamanickam, and Kamesh Madduri*

### **Distributed Memory Parallel Algorithms for Massive Graphs**

*Maksudul Alam, Shaikh Arifuzzaman, Hasanuzzaman Bhuiyan, Maleq Khan, V.S. Anil Kumar, and Madhav Marathe*

### **Efficient Multi-core Algorithms for Computing Spanning Forests and Connected Components**

*Fredrik Manne, Md. Mostofa Ali Patwary*

### **Massive-Scale Distributed Triangle Computation and Applications**

*Geoffrey Sanders, Roger Pearce, Benjamin W. Priest, Trevor Steil*

## Algorithms and Applications

### **Computing Top-k Closeness Centrality in Fully-dynamic Graphs**

*Eugenio Angriman, Patrick Bisenius, Elisabetta Bergamini, Henning Meyerhenke*

### **Ordering Heuristics for Parallel Graph Coloring**

*William Hasenplaugh, Tim Kaler, Tao B. Schardl, and Charles E. Leiserson*

### **Partitioning Trillion Edge Graphs**

*George M. Slota, Karen Devine, Sivasankaran Rajamanickam, Kamesh Madduri*

### **New Phenomena in Large-Scale Internet Traffic**

*Jeremy Kepner, Kenjiro Cho, KC Claffy, Vijay Gadepally, Sarah McGuire, Peter Michaleas, Lauren Milechin*

### **Parallel Algorithms for Butterfly Computations**

*Jessica Shi and Julian Shun*

## Models

### **Recent Advances in Scalable Network Generation**

*Manuel Penschuck, Ulrik Brandes, Michael Hamann, Sebastian Lamm, Ulrich Meyer, Ilya Safro, Peter Sanders, and Christian Schulz*

### **Computational Models for Cascades in Massive Graphs: How to Spread a Rumor in Parallel**

*Ajitesh Srivastava, Charalampos Chelmis, Viktor K. Prasanna*

### **Executing Dynamic Data-Graph Computations Deterministically Using Chromatic Scheduling**

*Tim Kaler, William Hasenplaugh, Tao B. Schardl, and Charles E. Leiserson*

## Frameworks and Software

### **Graph Data Science Using Neo4j**

*Amy E. Hodler, Mark Needham*

### **The Parallel Boost Graph Library 2.0**

*Nicholas Edmonds and Andrew Lumsdaine*

### **RAPIDS cuGraph**

*Alex Fender, Bradley Rees, Joe Eaton*

### **A Cloud-based approach to Big Graphs**

*Paul Burkhardt and Christopher A. Waring*

### **Introduction to GraphBLAS**

*Jeremy Kepner, Peter Aaltonen, David Bader, Aydin Buluc, Franz Franchetti, John Gilbert, Dylan Hutchinson, Manoj Kumar, Andrew Lumsdaine, Henning Meyerhenke, Scott McMillian, Jose Moreira, John D. Owens, Carl Yang, Marcin Zalewski, and Timothy G. Mattson*

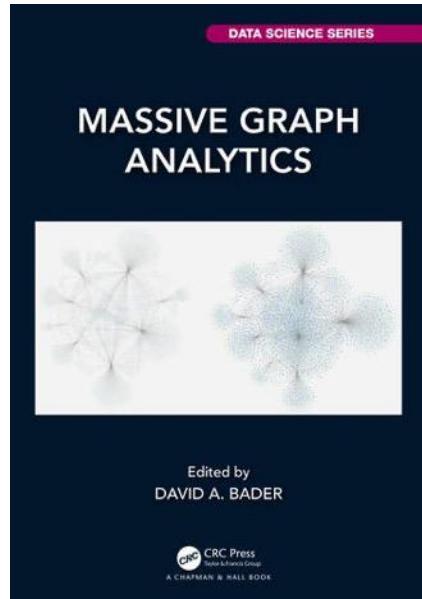
### **Graphulo: Linear Algebra Graph Kernels**

*Vijay Gadepally, Jake Bolewski, Daniel Hook, Shana Hutchison, Benjamin A Miller, Jeremy Kepner*

### **Interactive Graph Analytics at Scale in Arkouda**

*Zihui Du, Oliver Alvarado Rodriguez, Joseph Patchett, and David A. Bader*

# 85 Contributors



Peter Aaltonen	Joe Eaton	Charles E. Leiserson	Andrew Prou
Maksudul Alam	Nicholas Edmonds	Andrew Lumsdaine	Sivasankaran Rajamanickam
Md. Mostofa Ali Patwary	Stephan Erb	Kamesh Madduri	Bradley Rees
Eugenio Angriman	Alex Fender	Lawrence Mandon	Albert Reuther
V.S. Anil Kumar	Franz Franchetti	Fredrik Manne	Oliver Alvarado Rodriguez
William Arcand	Vijay Gadepally	Madhav V. Marathe	Antonio Rosa
Shaikh Arifuzzaman	John Gilbert	Timothy G. Mattson	Ilya Safro
Elisabetta Bergamini	Michael Hamann	Sarah McGuire	Siddharth Samsi
William Bergeron	William Hasenplaugh	Scott McMillan	Geoffrey Sanders
David Bestor	Amy E. Hodle	Ulrich Meyer	Peter Sanders
Hasanuzzaman Bhuiyan	Michael Houle	Henning Meyerhenke	Tao B. Schardl
Patrick Bisenius	Matthew Hubbell	Peter Michaleas	Christian Schulz
Ulrik Brandes	Shana Hutchison	Lauren Milechi	Jessica Shi
Aydin Buluc	Hayden Jananthan	Jose Moreira	Julian Shun
Paul Burkhardt	Michael Jones	Mark Needham	George M. Slota
Chansup Byun	Tim Kaler	John D. Owens	Ajitesh Srivastava
Charalampos Chelmis	Jeremy Kepner	Joseph Patchett	Trevor Steil
Kenjiro Cho	Maleq Khan	Roger Pearce	Christopher A. Waring
K.C. Claffy	Moritz Kobitzsch	Manuel Penschuck	Carl Yang
Karen Devine	Manoj Kumar	Viktor K. Prasanna	Charles Yee
Zhihui Du	Sebastian Lamm	Benjamin W. Priest	Marcin Zalewski

# Bader Research Lab

## Research Faculty:



Zhihui Du



Fatemeh  
Ramezani  
Khozestani



Fernando  
Vera  
Buschmann  
(Fulbright  
Scholar)



Fuhuan Li



Mohammad  
Dindoost



Oliver Alvarado  
Rodriguez



Asha Saxena



Soroush Vahidi  
(co-advised with  
Baruch Schieber)

## Undergraduates:



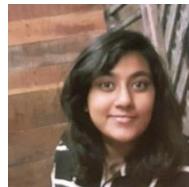
Vanshika  
Agrawal



Palina  
Paulichenka



Davor Petrovikj



Antonita  
Rachael



Naren  
Khatwani



Tania  
Majumder

## Current & Past High School Interns:



Pranhav  
Sundararajan



Anya  
Ganeshan

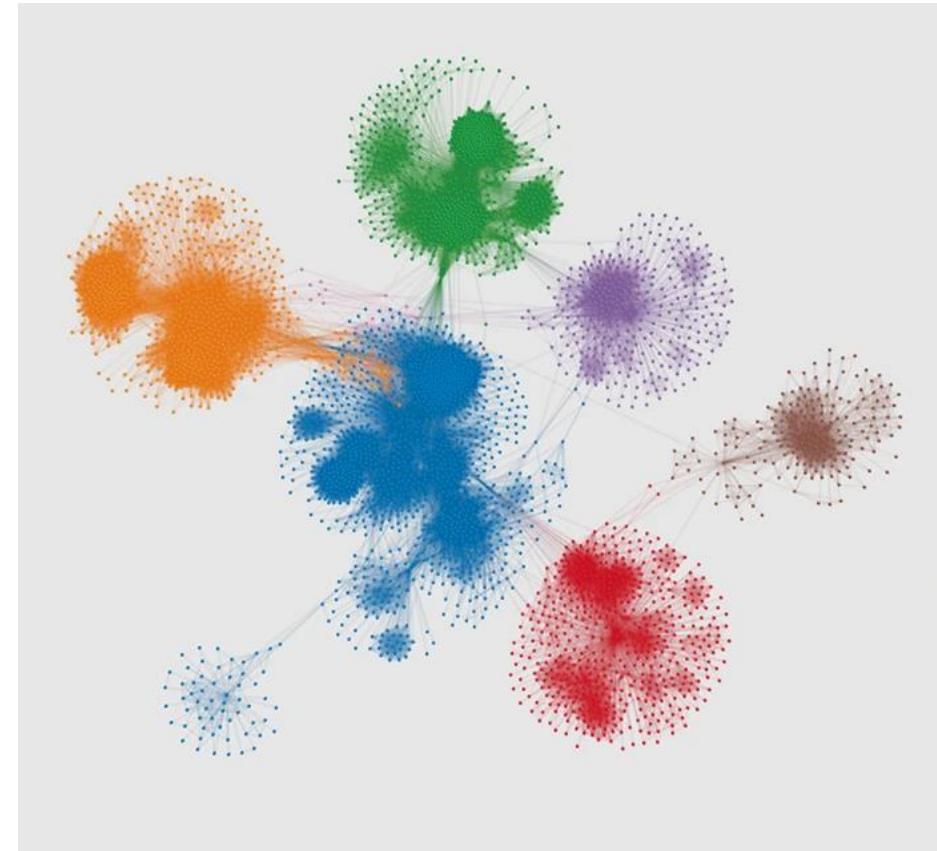


Sounak Bagchi



# Community Detection

- In complex networks, nodes cluster and form relatively dense groups – often called communities.
- Community detection is a fundamental graph algorithm with practical applications like ***fraud detection*** in Fintech and ***identity and access management*** in social networks



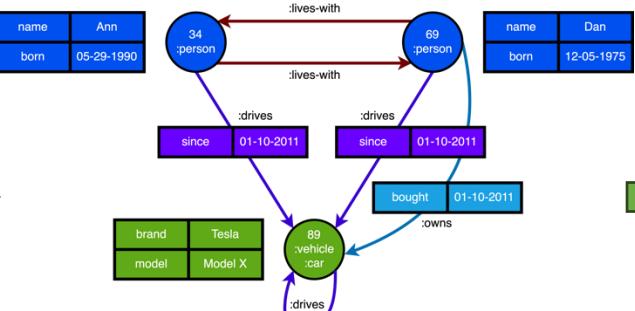
Vertices are Facebook users and edges represent Facebook friendships. Communities, represented by different colors.

Image credit: Fortunato, S., Newman, M.E.J. 20 years of network community detection. *Nat. Phys.* **18**, 848–850 (2022).

# Open-Source Massive-Scale (Property) Graph Analytics in Python with Arachne+Arkouda powered by Chapel

<b>id</b>	<b>label</b>	<b>name</b>	<b>born</b>	<b>brand</b>	<b>model</b>
34	person	Ann	1990	NULL	NULL
69	src id	dst id	relationship	since	bought
89	34	69	lives-with	NULL	NULL
89	69	34	lives-with	NULL	NULL
89	34	89	drives	2011	NULL
69	89	drives	2011	NULL	NULL
69	89	89	owns	NULL	2011
89	89	drives	NULL	NULL	NULL

Load in terabytes-sized CSVs, HDF5s, Parquets, etc.



Convert tabular data to graph format with all data closely maintained with vertex and edges.

<code>query_node_labels()</code>
<code>query_edge_relationships()</code>
<code>query_node_properties()</code>
<code>query_edge_properties()</code>

Extract (multiple) subgraphs by querying on attributes.

<code>bfs_layers()</code>
<code>subgraph_isomorphism()</code>
<code>square_counting()</code>
and more....!

Perform further analysis! Convert to NetworkX (if small enough).

```

1. import arkouda as ak
2. import arachne as ar
3.
4. ## Get src and dst from input file.
5.
6. graph = ar.PropGraph()
7. graph.add_edges_from(src,dst)
8.
9. ## Generate label_df and relationships_df from input file.
10.
11. graph.add_node_labels(label_df)
12. graph.add_edge_relationships(relationships_df)
13.
14. ## User generates labels_to_find and relationships_to_find.
15. returned_nodes = graph.query_labels(labels_to_find)
16. returned_edges = graph.query_relationships(relationships_to_find)
17.
18. subgraph_src = ak.in1d(returned_edges[0], returned_nodes)
19. subgraph_dst = ak.in1d(returned_edges[1], returned_nodes)
20.
21. kept_edges = subgraph_src & subgraph_dst
22.
23. subgraph_src = subgraph_src[kept_edges]
24. subgraph_dst = subgraph_dst[kept_edges]
25.
26. subgraph = ar.Graph()
27. subgraph.add_edges_from(subgraph_src, subgraph_dst)
28. ## Run some new operations on subgraph! Reference our HPEC22 paper ☺

```

Easily usable through NetworkX-like API. Data exchangeable between NetworkX, NumPy, SciPy, etc.



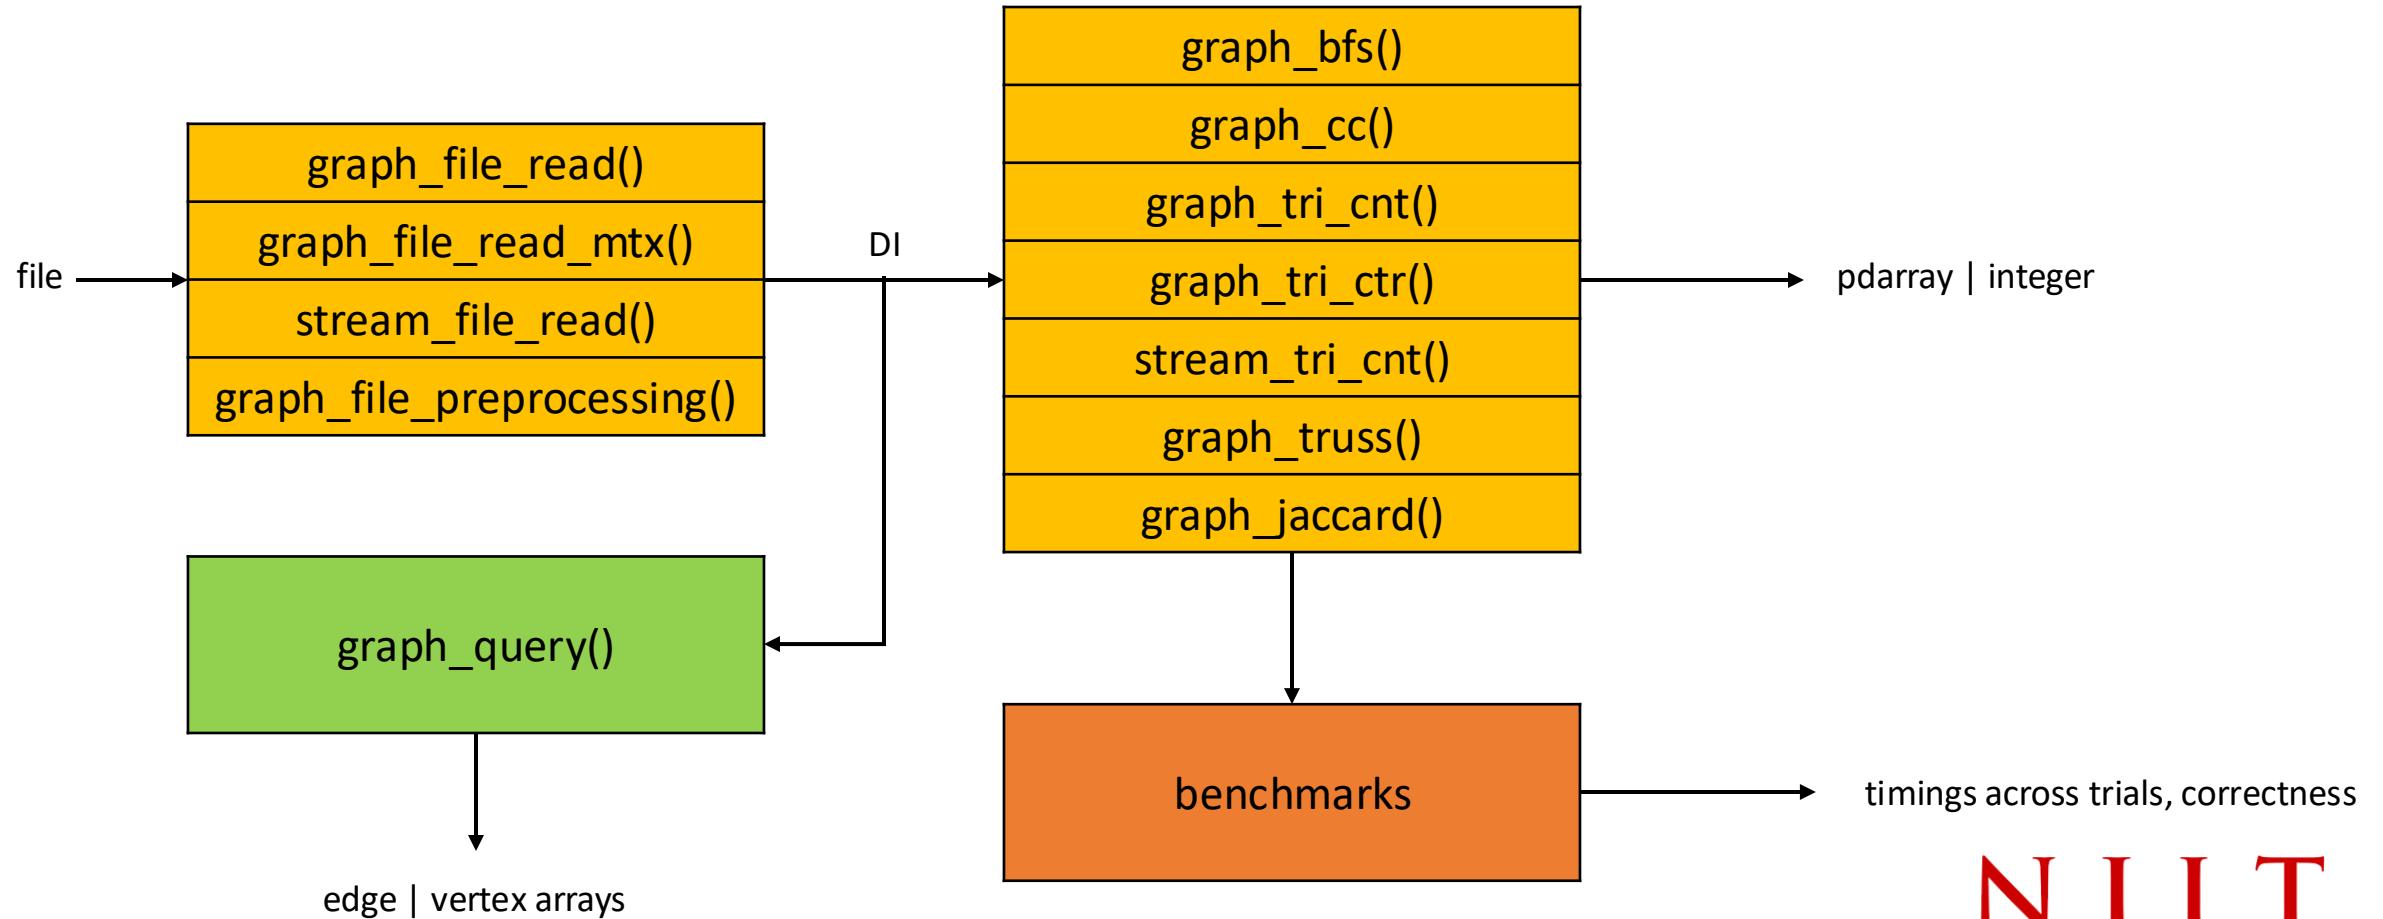
Runs on any hardware, data stays in the back-end, user calls API through Pythpm: powerful and productive. Server can run on supercomputers, Python API usable locally.

Kernel	Time(s)
<b>Loading graph into memory</b>	5.4
<b>Node label ETL and storage</b>	1.87
<b>Edge label ETL and storage</b>	80.29
<b>Node querying</b>	1.01
<b>Edge querying</b>	1.39
<b>Subgraph building</b>	1.22
<b>Breadth-first search</b>	2.45

Performance analyzing a 10-billion edge graph on 16 compute nodes with 128 cores each and 1TB memory. Subgraph generated contained 1-million edges. BFS executed from the 5 highest degree vertices and the average was taken.



# Modules of Arachne



# Graphs for Testing

Real-world

<b>Graphs</b>	<b>Edges</b>	<b>Vertices</b>	<b>CCs</b>	<b>Triangles</b>	<b>Max K</b>
<b>ca-GrQc</b>	14484	5242	354	48260	44
<b>ca-HepTh</b>	25973	9877	427	28339	32
<b>as-caida20071105</b>	53381	26475	1	36365	16
<b>facebook_combined</b>	88234	4039	1	1612010	97
<b>ca-CondMat</b>	93439	23133	567	173361	26
<b>ca-HepPh</b>	118489	12008	276	3358499	239
<b>email-Enron</b>	183831	36692	1065	727044	22
<b>ca-AstroPh</b>	198050	18772	289	1351441	57
<b>loc-brightkite_edges</b>	214078	58228	547	494728	43
<b>soc-Epinions1</b>	405740	75879	2	1624481	33
<b>amazon0601</b>	2443408	403394	7	3986507	11
<b>com-Youtube</b>	2987624	1134890	1	3056386	19
<b>friendster</b>	1806067135	65608366	1	4173724142	129

values found by our algorithms

Experiments were conducted on a high-performance server with 2 x Intel Xeon E5-2650 v3 @ 2.30GHz CPUs with 10 cores per CPU and a RAM capacity of 512GB.

few vertices, outperforms some algorithms less edges but more vertices.

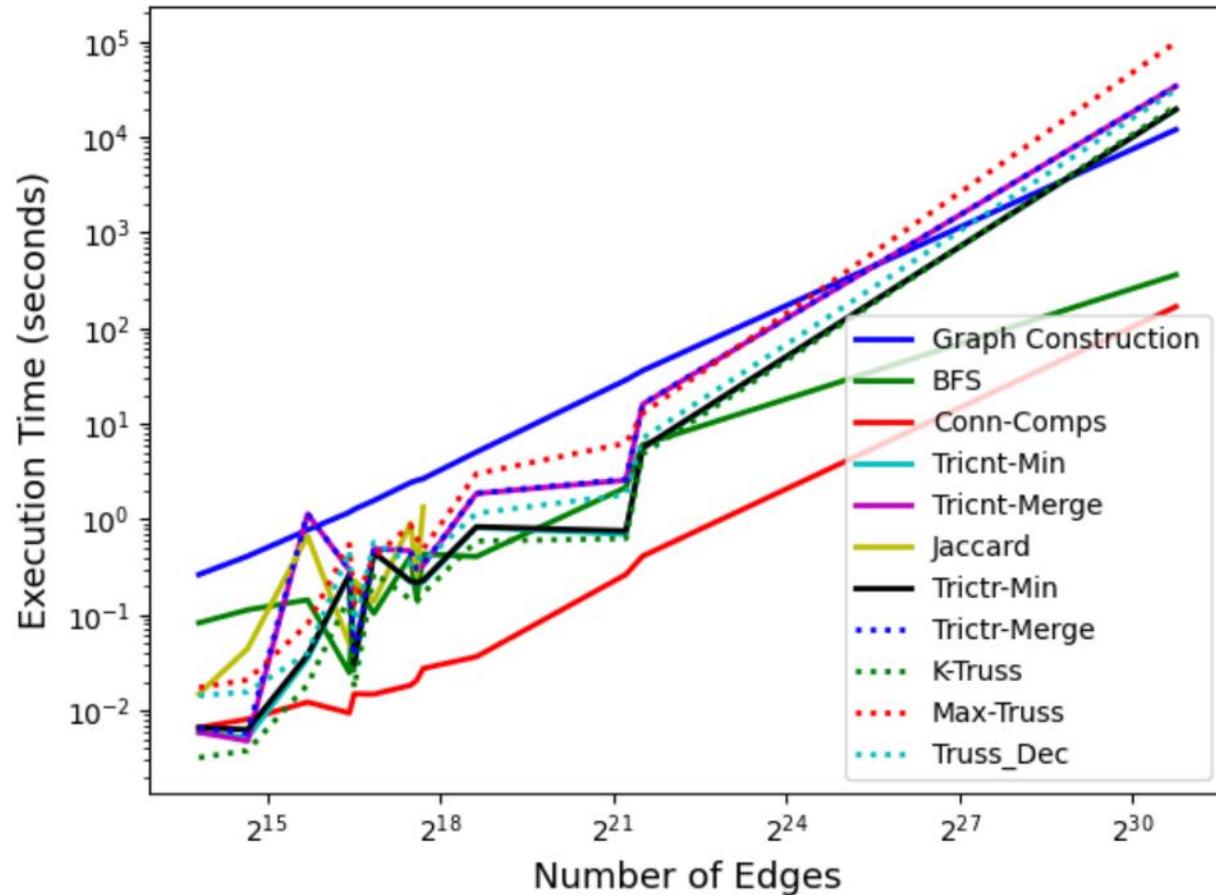
Synthetic

<b>delaunayn20</b>	<b>Edges</b>	<b>Vertices</b>	<b>CCs</b>	<b>Triangles</b>	<b>Max K</b>
<b>delaunayn21</b>	3145686	1048576	1	2109090	4
<b>delaunayn22</b>	6291408	2097152	1	4218386	4
<b>delaunayn23</b>	12582869	4194304	1	8436672	4
<b>delaunayn24</b>	25165784	8388608	1	16873359	4
	50331601	16777216	1	33746670	4

delaunayn10 - delaunayn19

# Arachne Results – Real-World Graphs

[Alvarado Rodriguez, Du, Patchett, Li, Bader 2022]

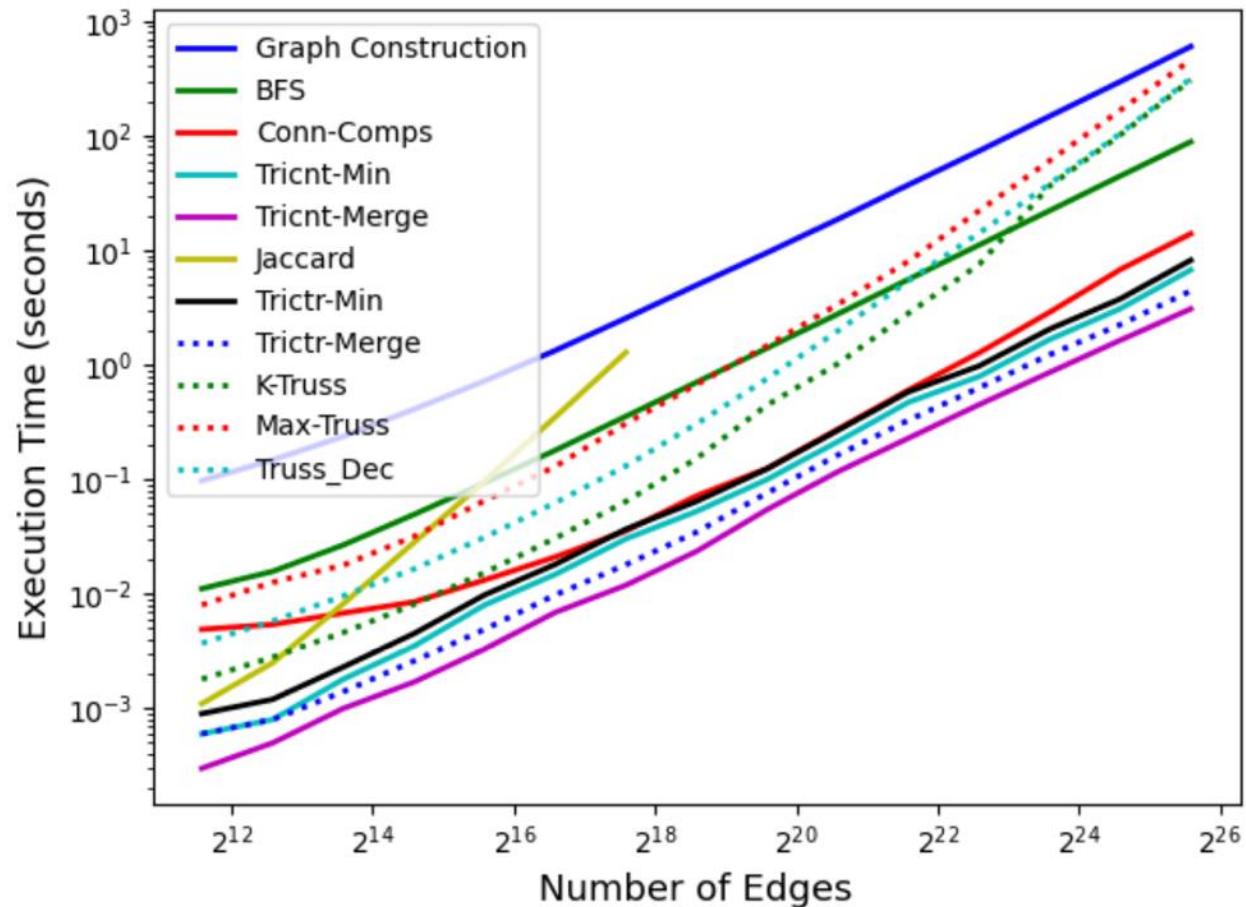


## Key Points:

1. Graph construction is time consuming but once the graph is built into memory all the algorithms can use it in a highly efficient way.
2. The structural properties of graphs can significantly affect execution times even for the same algorithm.

# Arachne Results – Synthetic Graphs

[Alvarado Rodriguez, Du, Patchett, Li, Bader 2022]



## Key Points:

1. Synthetic graphs demonstrate the scalability of our algorithms as the number of edges in a graph increase.
2. The memory requirements for each algorithm differ, hence the Jaccard coefficient algorithm encounters out of memory errors when the graph gets too big. Jaccard requires  $\frac{N \times N}{2}$  memory and  $\left(\frac{N}{P}\right)^2 \times \frac{M}{P}$  calculations.

# Breadth-First Search Improvements

Graph	num_vertices	num_edges	di (original)	di-norev	speedup
as-caida	26,475	53,381	2.22	1.22	1.82
delaunayn10	2,048	3,056	0.11	0.05	2.11
delaunayn20	1,058,576	3,145,686	90.49	47.61	1.90

Execution time in seconds on eight locales with 512GB memory and twenty processing units each.

```

while !SetCurF.isEmpty() do
    coforall loc in Locales do
        // parallel search on each locale
        forall i in SetCurF do
            if i is on current locale then
                SetNeighbor =  $\cup_k$ , k is a neighbor of i
                forall j in SetNeighbor do
                    if depth[j] == -1 then
                        SetNextF.add(j)
                    depth[j] = current_level + 1
                end
            end
        end
    end
end
SetCurF <=> SetNextF // exchange values
SetNextF.clear()
current_level+ = 1
end

```

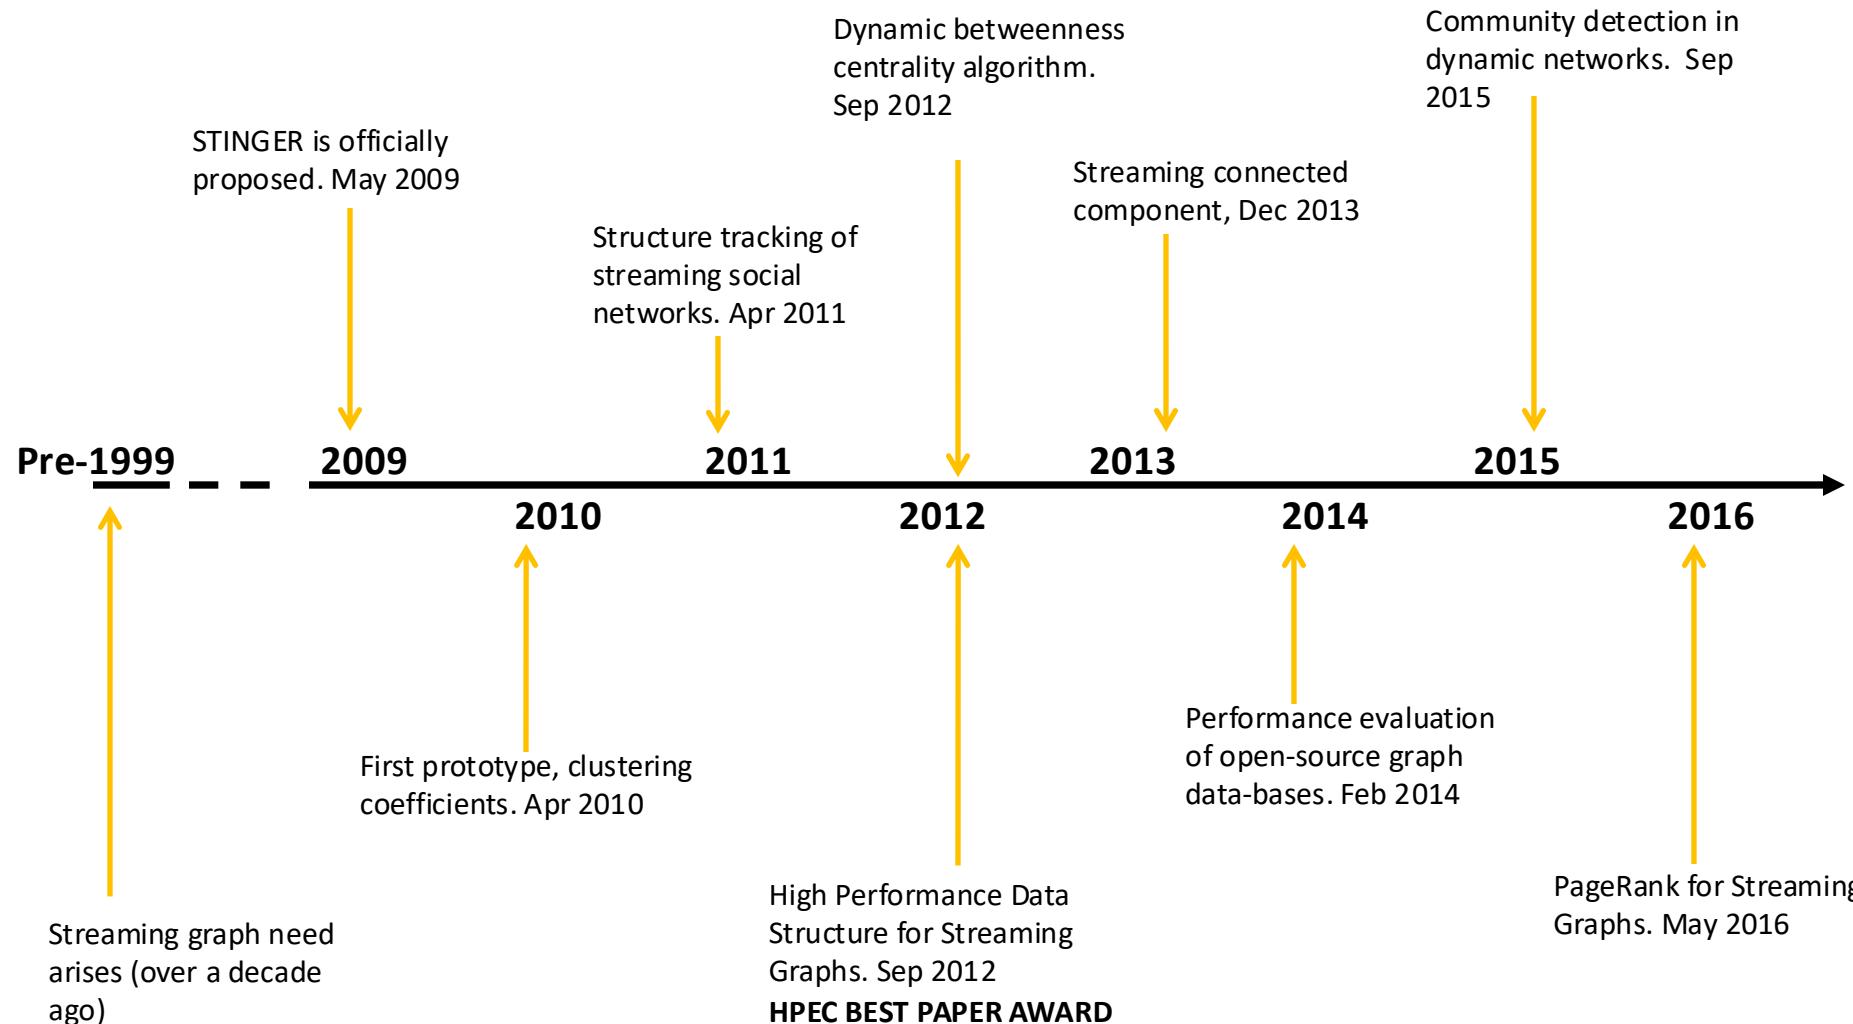
Performed twice in di(original) while iterating over the edges in (SRC,DST) and in (SRCr, DSTr) to get the full neighborhoods which can lead to twice the number of remote reads and writes!

- About 50% improvement in number of PUTs and GETs with di-norev by including full neighborhoods of each vertex contiguously in one array instead of maintaining reversed edges.
- No change in storage volume, 2m edges still stored.
- Similar changes could optimize the rest of our graph kernels.

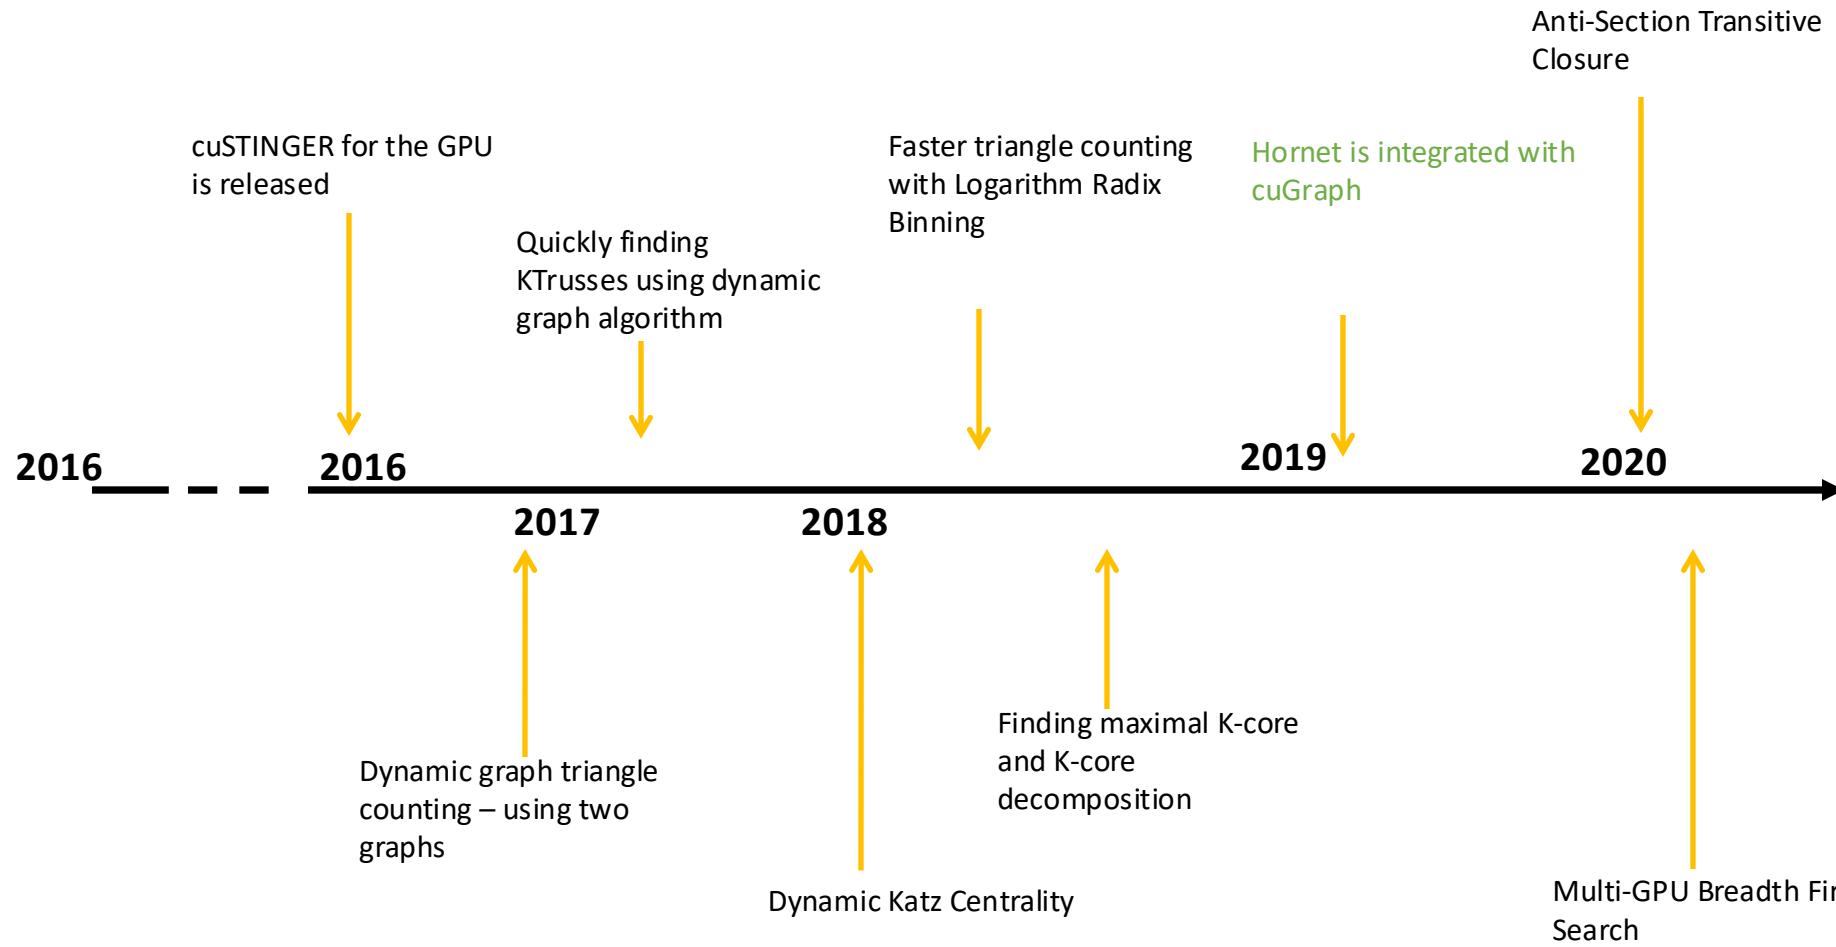
# STING Initiative: Focusing on Globally Significant Grand Challenges

- Many globally-significant grand challenges can be modeled by **Spatio-Temporal Interaction Networks and Graphs** (or “STING”).
- Emerging real-world graph problems include:
  - Detecting community structure in large social networks
  - Defending the nation against cyber-based attacks
  - Discovering insider threats (e.g. Ft. Hood shooter, WikiLeaks)
  - Improving the resilience of the electric power grid
  - Detecting and preventing disease in human populations.
- Unlike traditional applications in computational science and engineering, solving these problems at scale often raises new research challenges due to:
  - Sparsity and the lack of locality in the massive data
  - Design of parallel algorithms for massive, streaming data analytics
  - The need for new exascale supercomputers that are energy-efficient, resilient, and easy-to-program

# STINGER – Time Frame



# Hornet (GPU only) – Time Frame



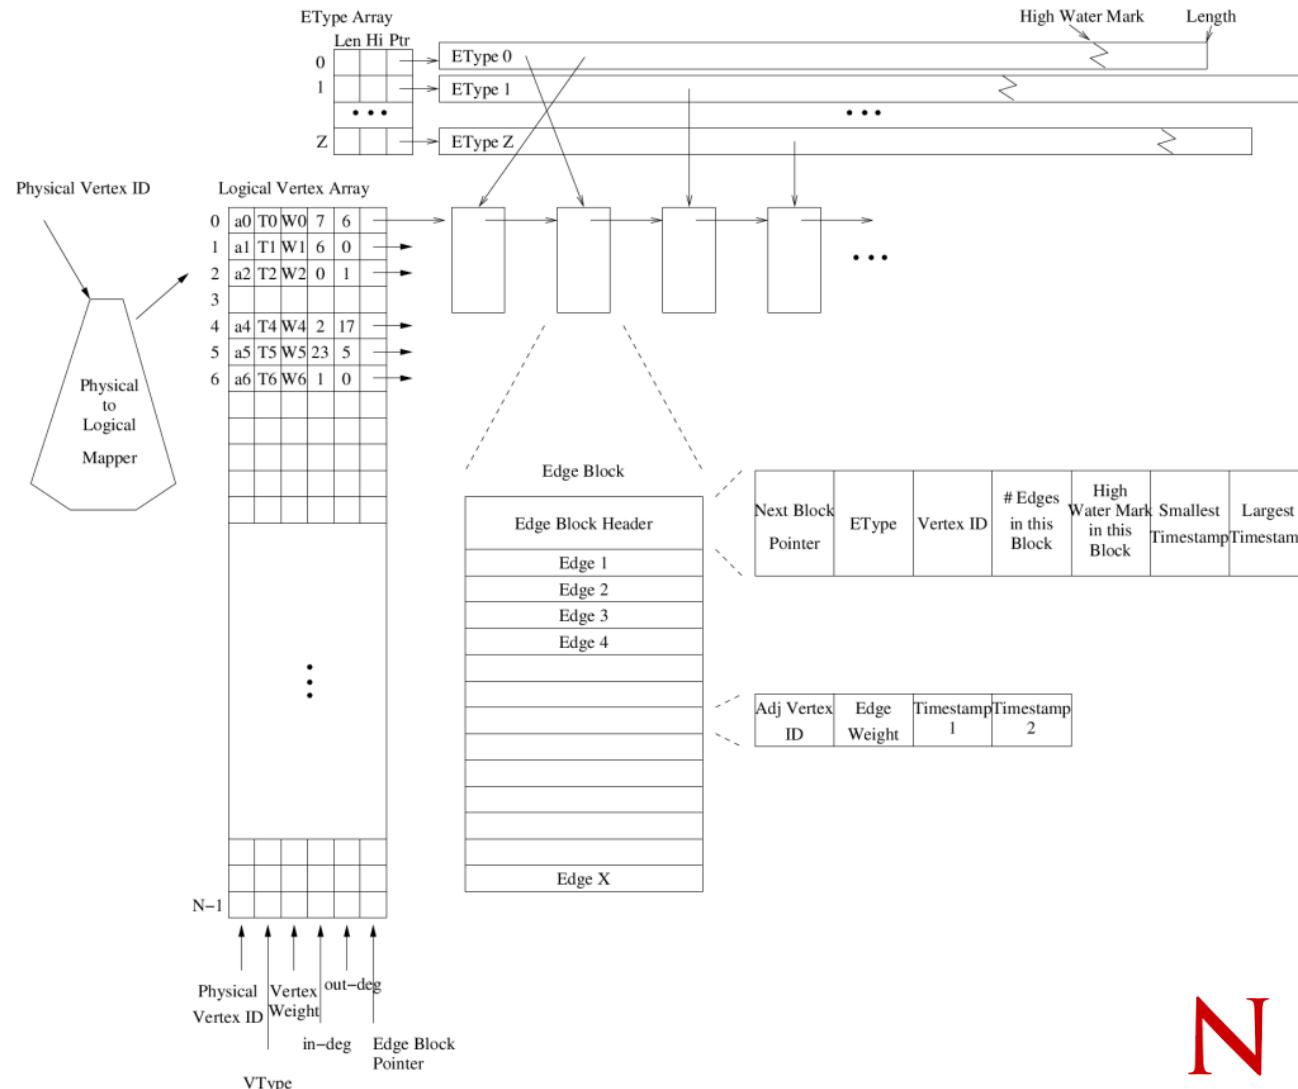
# STING Extensible Representation (STINGER)

## *Design goals*

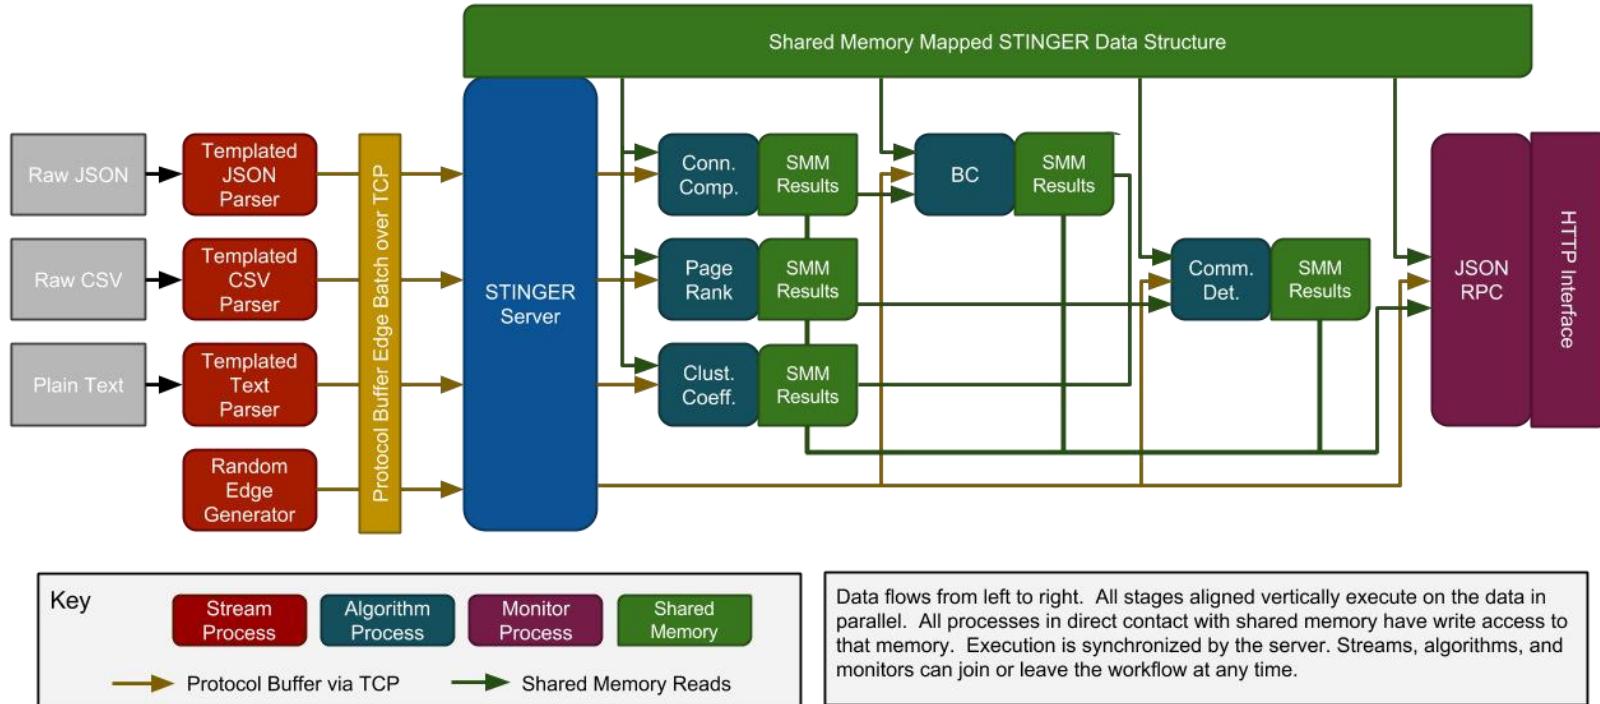
- Enable algorithm designers to implement dynamic graph algorithms with ease.
- Portable semantics for various platforms
- Good performance for all types of graph problems and algorithms
  - static and dynamic.
- Assumes globally addressable memory access
- Support multiple, parallel readers and a single writer
  - One server manages the graph data structures
  - Multiple analytics run in background with read-only permissions.

# STING Extensible Representation (STINGER)

- Semi-dense edge list blocks with free space
- Compactly stores timestamps, types, weights
- Maps from application IDs to storage IDs
- Deletion by negating IDs, separate compaction



# STING: High-level architecture



- △ Server: Graph storage, kernel orchestration
- △ OpenMP + sufficiently POSIX-ish
- △ Multiple processes for resilience

# STINGER as an analysis package

<http://www.stingergraph.com/>

**Anything that a static graph package can do (and a whole lot more):**

**Parallel agglomerative clustering:**

Find clusters that are optimized for a user-defined edge scoring function.

**K-core Extraction:**

Extract additional communities and filter noisy high-degree vertices.

**Classic breadth-first search:**

Performs a parallel breadth-first search of the graph starting at a given source vertex to find shortest paths.

**Parallel connected components:**

Finds the connected components in a static network.

**AND...**

**Streaming edge insertions and deletions:**

New edge insertions, updates, and deletions in batches or individually. Optimized to update at rates of over 3 million edges per second on graphs of one billion edges.

**Streaming clustering coefficients:**

Tracks the local and global clustering coefficients of a graph.

**Streaming connected components:**

Real time tracking of the connected components.

**Streaming Betweenness Centrality:**

Find the key points within information flows and structural vulnerabilities.

**Streaming community detection:**

Track and update the community structures within the graph as they change.

# Why not existing technologies?

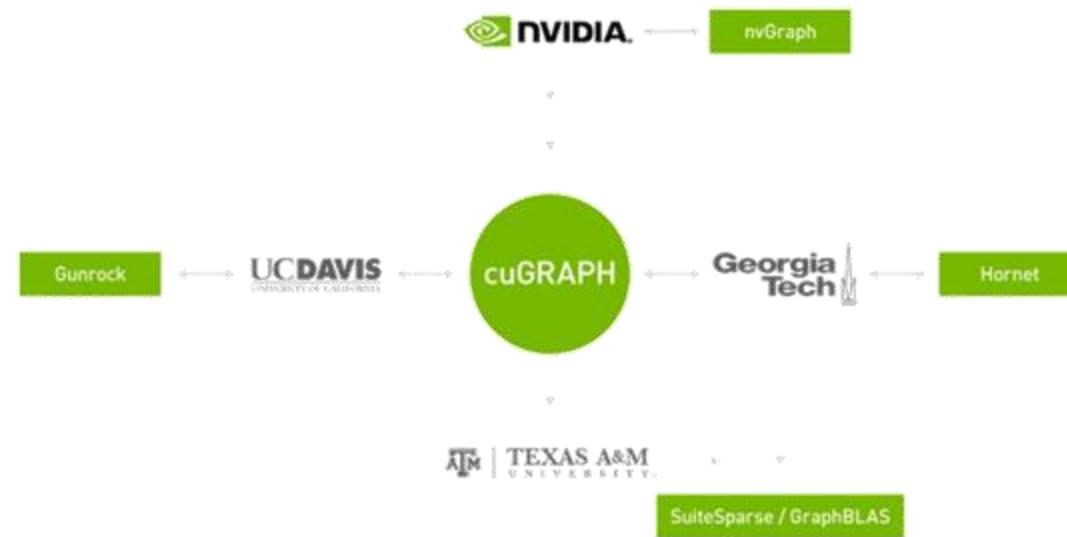
- Traditional SQL databases
  - Not structured to do any meaningful graph queries with any level of efficiency or timeliness
- Graph databases - mostly on-disk
  - Distributed disk can keep up with storing / indexing, but is simply too slow at random graph access to process on as the graph updates
- Hadoop and HDFS-based projects
  - Not really the right programming model for many structural queries over the entire graph, random access performance is poor
- Smaller graph libraries, processing tools
  - Can't scale, can't process dynamic graphs, frequently leads to impossible visualization attempts



# AI Lab (NVAIL) 2019, PI: Bader

## Building the Future of Graph Analytics with RAPIDS

“Prof. David Bader and his lab ... are leaders in high performance computing algorithms, with a focus on both static and dynamic graph algorithms. With Prof. Bader and his lab, we will work on the design and implementation of scalable graph algorithms and graph primitives for integrating into cuGRAPH, leveraging their Hornet framework.” – Sandra Skaff, NVIDIA, April 2019





# 2019 Facebook AI Systems Award: Scalable Graph Learning Algorithms

**Project Aim:** Develop scalable graph learning algorithms and implementations that open the door for learned graph models on massive graphs

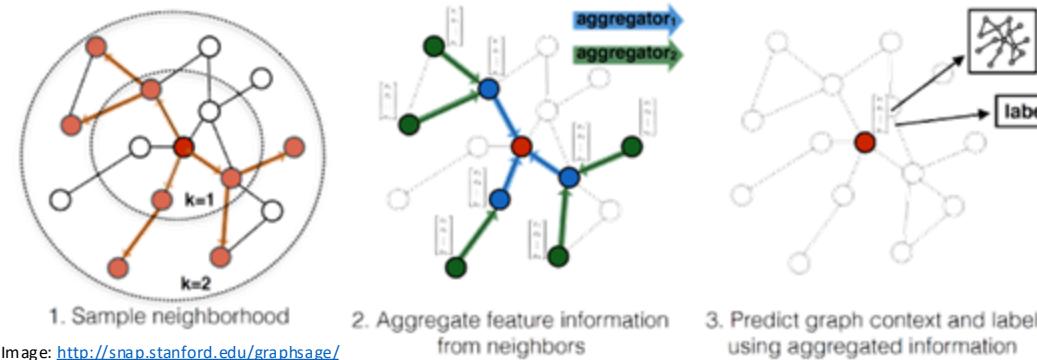


Image: <http://snap.stanford.edu/graphsage/>

Inductive Representation Learning on Large Graphs. W.L. Hamilton, R. Ying, and J. Leskovec arXiv:1706.02216 [cs.SI], 2017.

**Deep Learning (DL)** has significantly impacted the tasks of speech recognition, image classification, object detection and recommendation

**Complex tasks:** self-driving, super-human image recognition, recommendation engines, machine natural language translation, content selection, learning patterns of life

**Techniques used in DL:** convolutional neural networks (CNNs) → applicable for Euclidean data types and does not apply for Graphs

**Solution:** embedding graphs into a lower dimensional Euclidean space, generating a regular structure

1. developing a **scalable high performance graph learning system** based on GCNs algorithms, like GraphSage, by improving the workflow on shared-memory NUMA machines balancing computation between threads, optimizing data movement, and improving memory locality
2. **investigate graph learning algorithm:** specific decompositions and develop new strategies for graph learning that can inherently scale well while maintaining high accuracy
  - Explore decomposition results from graph theory, for example forbidden graphs and the Embedding Lemma and determine how to apply such results into the field of graph learning
  - Investigate whether these decompositions could assist in a dynamic graph setting

# Graphs are pervasive in large-scale data analysis

- **Sources** of massive data: peta- and exa-scale simulations, experimental devices, the Internet, scientific applications.
- **New challenges for analysis**: data sizes, heterogeneity, uncertainty, data quality.

## Astrophysics

Problem: Outlier detection.  
Challenges: massive datasets, temporal variations.  
Graph problems: clustering, matching.

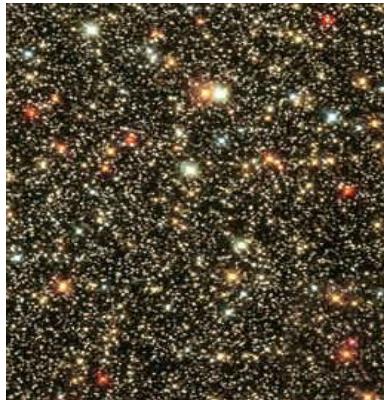
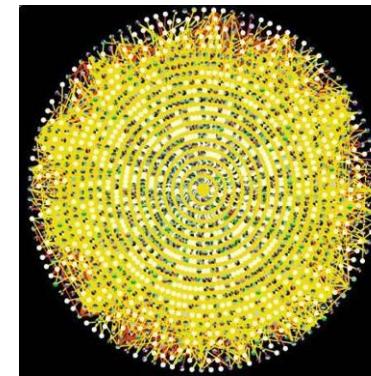


Image sources: (1) [http://physics.nmt.edu/images/astro/hst\\_starfield.jpg](http://physics.nmt.edu/images/astro/hst_starfield.jpg)  
(2,3) [www.visualComplexity.com](http://www.visualComplexity.com)

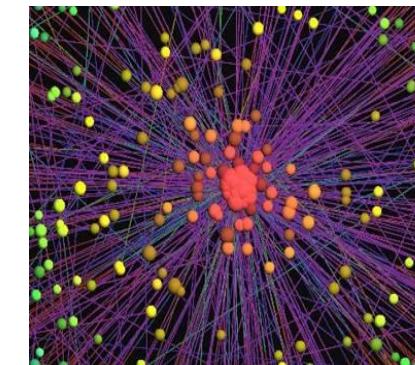
## Bioinformatics

Problem: Identifying drug target proteins.  
Challenges: Data heterogeneity, quality.  
Graph problems: centrality, clustering.



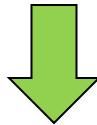
## Social Informatics

Problem: Discover emergent communities, model spread of information.  
Challenges: new analytics routines, uncertainty in data.  
Graph problems: clustering, shortest paths, flows.



# Characterizing Graph-theoretic computations

Input: Graph abstraction



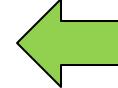
Problem: Find \*\*\*

- paths
- clusters
- partitions
- matchings
- patterns
- orderings



Graph algorithms

- traversal
- shortest path algorithms
- flow algorithms
- spanning tree algorithms
- topological sort
- .....



Factors that influence choice of algorithm

- graph sparsity ( $m/n$  ratio)
- static/dynamic nature
- weighted/unweighted, weight distribution
- vertex degree distribution
- directed/undirected
- simple/multi/hyper graph
- problem size
- granularity of computation at nodes/edges
- domain-specific characteristics



Graph problems are often recast as sparse linear algebra (e.g., partitioning) or linear programming (e.g., matching) computations

Streaming Analytics move us from reporting the news to predictive analytics

## Traditional HPC

- Great for “static” data sets.
- Massive scalability at the cost of programmability.
- Great for dense problems.
  - Sparse problems typically underutilize the system.



## Streaming Analytics

- Requires specialized analytics and data structures.
- Rapidly changing data.
- Low data re-usage.
  - Focused on memory operations and not FLOPS.



# Graph Data Science

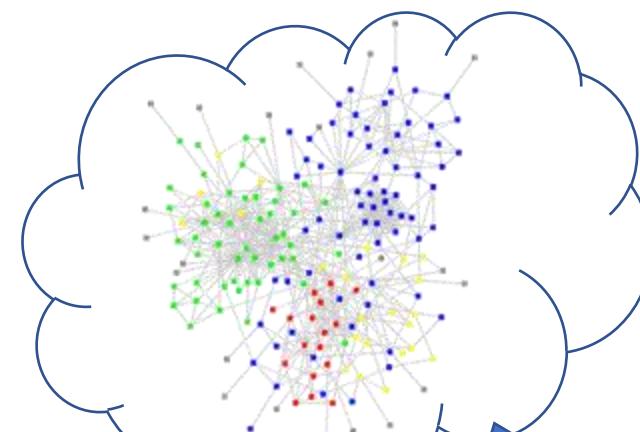
- Are there new graph techniques? Do they scale? Can the computational systems (algorithms, machines) handle massive networks with billions to trillions of items? Can the techniques tolerate noisy data, massive data, streaming data, etc. ...
- **Communities may overlap, exhibit different properties and sizes, and be driven by different models**
  - Detect communities (static or emerging)
  - Identify important individuals
  - Detect anomalous behavior
  - Given a community, find a representative member of the community
  - Given a set of individuals, find the best community that includes them
  - Find congestion, weak points, anomalies, surprises, ...

# Massive Streaming Graph Analytics



(A, B, t1, **poke**)  
(A, C, t2, **msg**)  
(A, D, t3, **view wall**)  
(A, D, t4, **post**)

(B, A, t2, **poke**)  
(B, A, t3, **view wall**)  
(B, A, t4, **msg**)



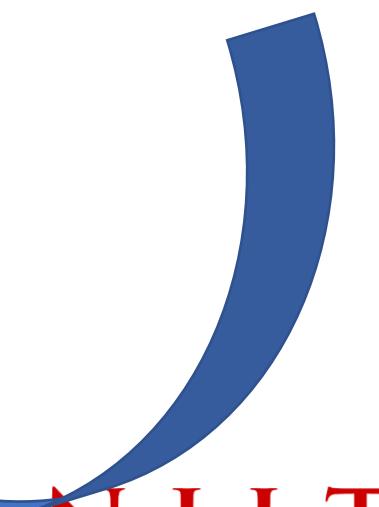
... e9 e8 e7 Billions of edges  
e6 e5 e4 e3 e2 e1 ...



David A. Bader



Q3?  
Q2?  
Q1?



N J I T  
New Jersey Institute  
of Technology

# Hierarchy of Interesting Analytics

## ► Extend single-shot graph queries to include time.

- Are there  $s-t$  paths between time  $T_1$  and  $T_2$ ?
- What are the important vertices at time  $T$ ?

## ► Use persistent queries to monitor properties.

- Does the path between  $s$  and  $t$  shorten drastically?
- Is some vertex suddenly very central?

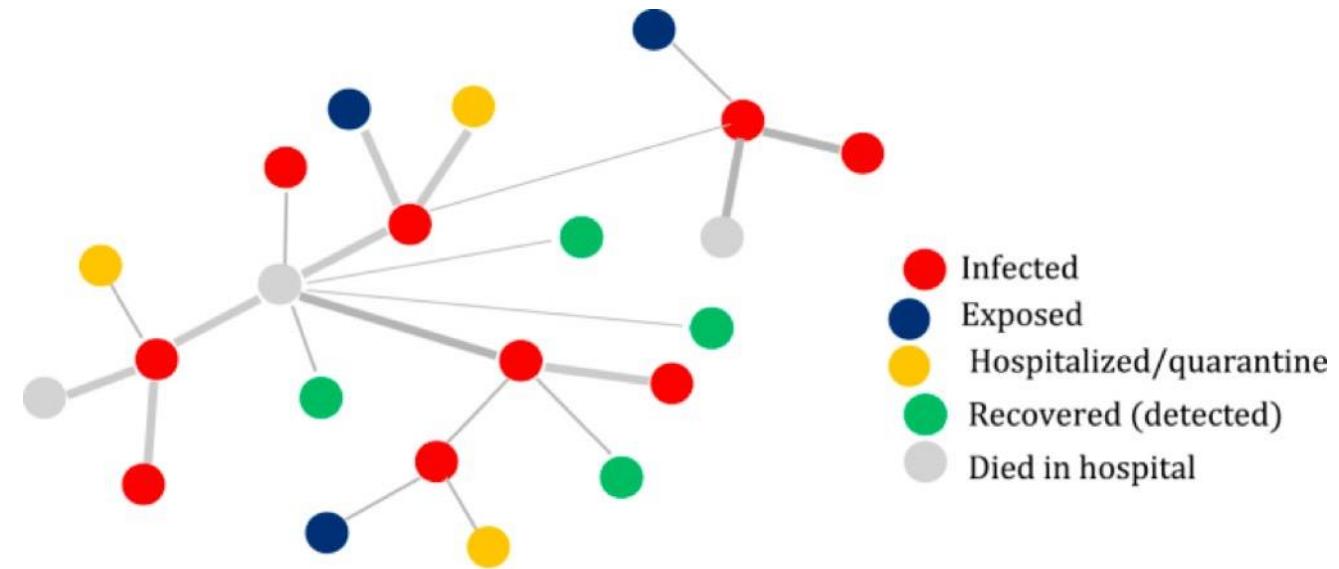
## ► Extend persistent queries to fully dynamic properties.

- Does a small community stay independent rather than merge with larger groups?
- When does a vertex jump between communities?

## ► New types of queries, new challenges...

# Modeling Pandemic Spread

- The graph represents the contact patterns between individuals in a population.
- Various graph algorithms can be used to simulate the spread of a pandemic.
  - Centrality measures such as eigenvector centrality can identify the most important vertices in the network
  - Visualization of the spread of the pandemic can be created to check the effects of intervention and control strategies.
- The dataset can be a million or even a trillion vertices.



[Alguliyev, Aliguliyev, Yusifov, 2020]

R. Alguliyev, R. Alguliyev, and F. Yusifov, "Graph modelling for tracking the COVID-19 pandemic spread." Infectious disease modelling, 6, 2021: 112-122

# Using Arachne with Arkouda (1/3)

```
In [2]: ak.connect("d-6-15-4", 5555)

connected to arkouda server tcp://*:5555

In [3]: # Read in the graph.
filename = "/home/gridsan/oarodriguez/biggraph_shared/Adata/simple.txt"
ne = 13
nv = 10
G = ar.graph_file_read(ne, nv, 2, 0, filename, 1, 0, 0, 0, 1)

13 10 2 0 /home/gridsan/oarodriguez/biggraph_shared/Adata/simple.txt 1 0 0 0 1

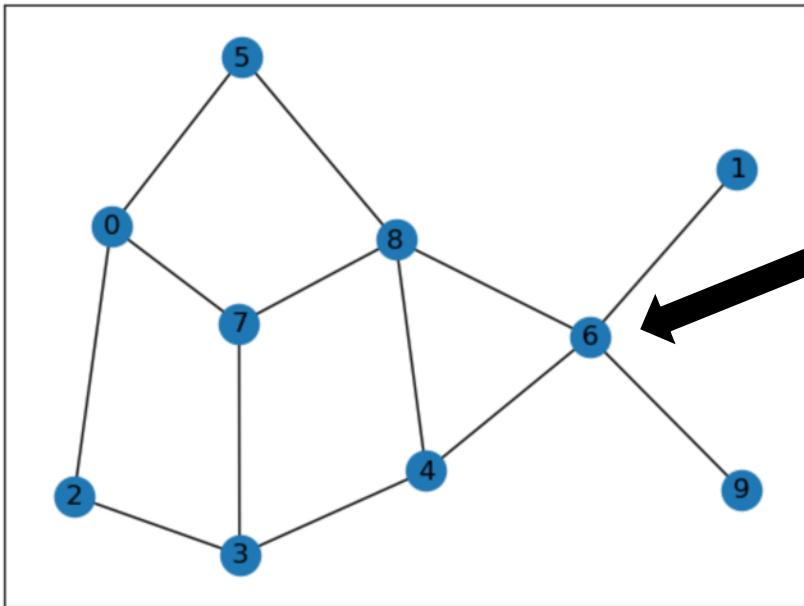
In [4]: # Add the edges of the graph to a list of tuples.
src = ar.graph_query(G, "src")
dst = ar.graph_query(G, "dst")

edges = []
for (u, v) in zip(src.to_ndarray(), dst.to_ndarray()):
    edges.append((u,v))

In [5]: # Display the graph with NetworkX.
nxG = nx.Graph()
nxG.add_edges_from(edges)

pos = nx.spring_layout(nxG, seed=225)
nx.draw_networkx(nxG, pos, with_labels=True)
plt.show()
```

# Using Arachne with Arkouda (2/3)



max degree! (also 8, but 6 is the first occurrence)

```
In [6]: # Get value of the maximum degree.
neighbour = ar.graph_query(G, "neighbour")
neighbourR = ar.graph_query(G, "neighbourR")
degrees = neighbour + neighbourR
print("The value of the maximum degree is: {}".format(ak.max(degrees)))
```

The value of the maximum degree is: 4

# Using Arachne with Arkouda (3/3)

## Breadth-First Search

```
In [7]: d = ar.graph_bfs(G, int(ak.argmax(degrees)), 0)
print(d)
```

```
[3 1 3 2 1 2 0 2 1 1]
```

```
In [8]: # Get the size of each level of BFS.
d_histogram = ak.histogram(d, bins=ak.max(d)+1)
print(d_histogram)
```

```
(array([0. , 0.75, 1.5 , 2.25]), array([1 4 3 2]))
```

