



# Chapel: Productive Parallel Programming

(under the influence of jetlag)

Brad Chamberlain, Chapel Team, Cray Inc.  
ETH Zürich  
April 25<sup>th</sup>, 2014\*



\* = Happy birthday, Mom!

# Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.



# Sustained Performance Milestones

## 1 GF – 1988: Cray Y-MP; 8 Processors

- Static finite element analysis



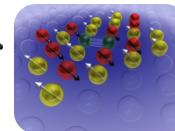
## 1 TF – 1998: Cray T3E; 1,024 Processors

- Modeling of metallic magnet atoms



## 1 PF – 2008: Cray XT5; 150,000 Processors

- Superconductive materials



## 1 EF – ~2018: Cray \_\_\_\_; ~10,000,000 Processors

- TBD

# Sustained Performance Milestones

## 1 GF – 1988: Cray Y-MP; 8 Processors

- Static finite element analysis
- Fortran77 + Cray autotasking + vectorization



## 1 TF – 1998: Cray T3E; 1,024 Processors

- Modeling of metallic magnet atoms
- Fortran + MPI (Message Passing Interface)



## 1 PF – 2008: Cray XT5; 150,000 Processors

- Superconductive materials
- C++/Fortran + MPI + vectorization



## 1 EF – ~2018: Cray \_\_\_\_; ~10,000,000 Processors

- TBD
- TBD: C/C++/Fortran + MPI + CUDA/OpenCL/OpenMP/OpenACC?

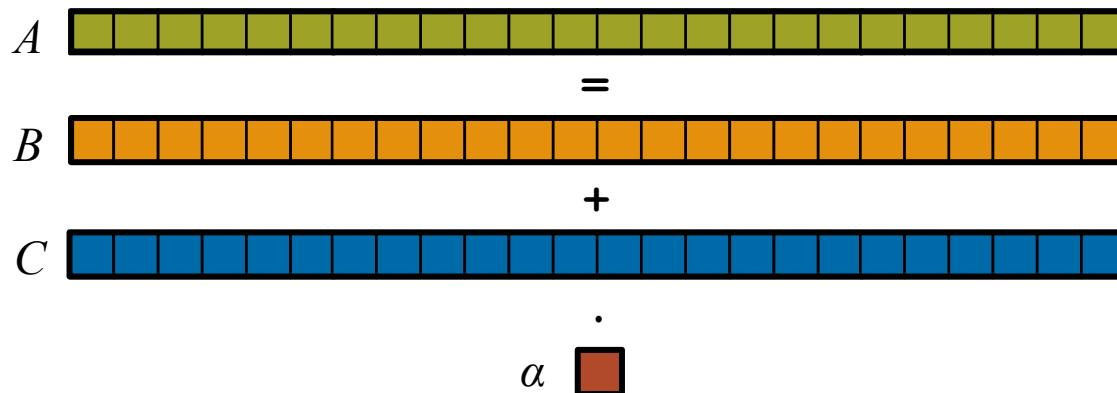
Or, perhaps something completely different?

# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures:**

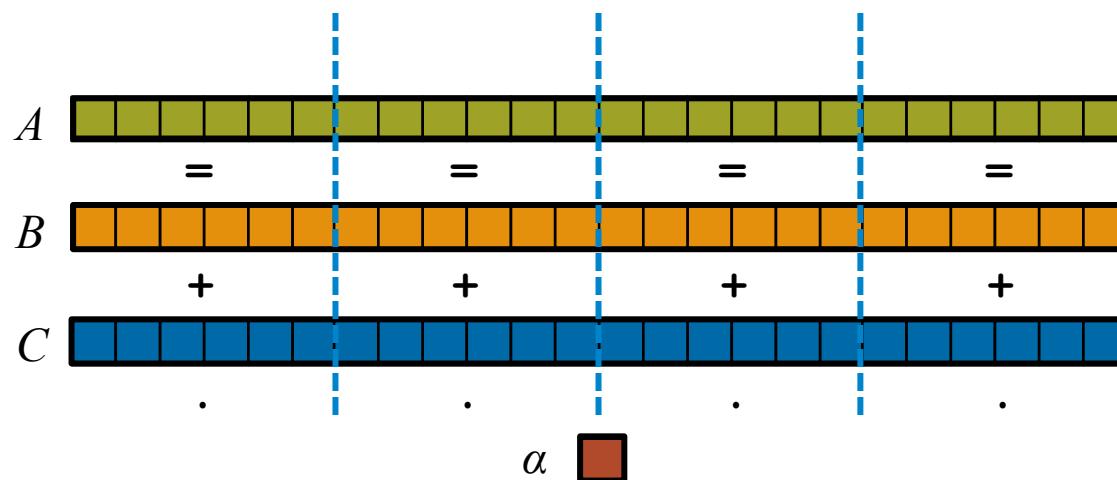


# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures, in parallel:**

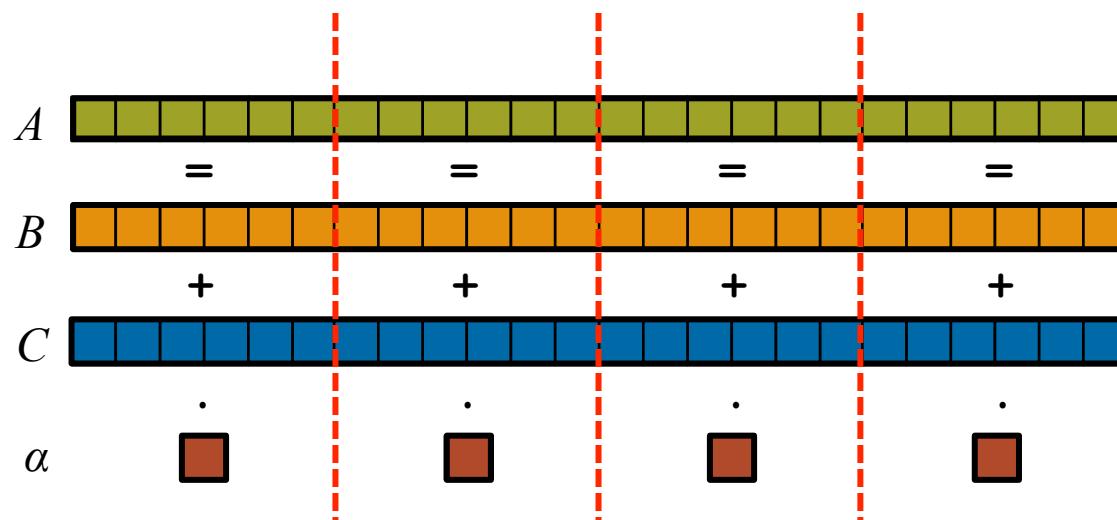


# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures, in parallel (distributed memory):**

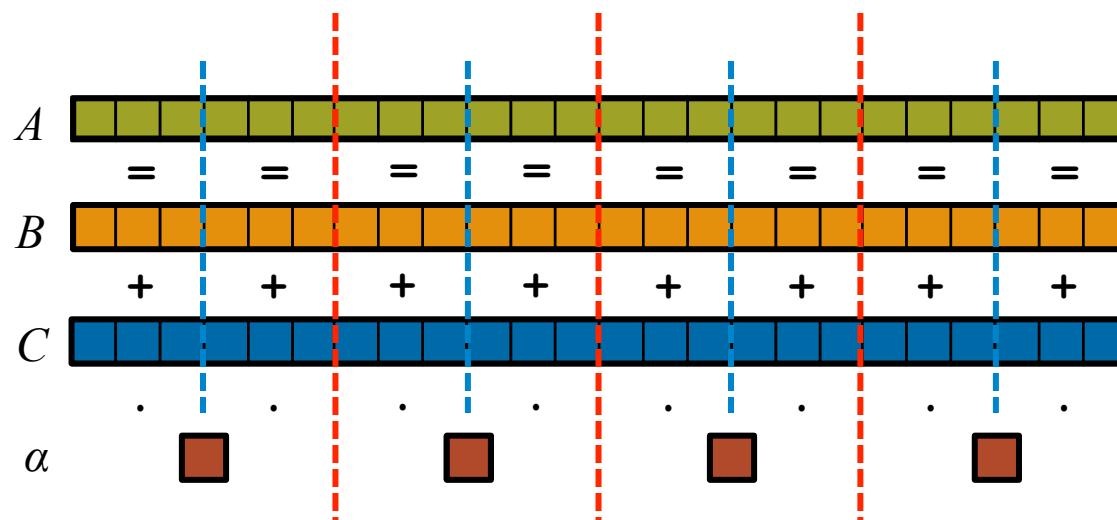


# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures, in parallel (distributed memory multicore):**



# STREAM Triad: MPI

MPI

```
#include <hpcc.h>

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

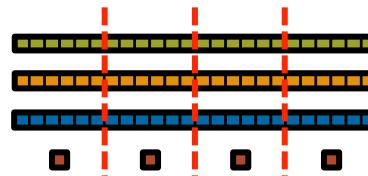
    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
                0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
                                       sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
```



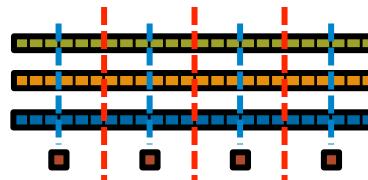
```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d).
\n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 0.0;
}
scalar = 3.0;

for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
```

# STREAM Triad: MPI+OpenMP



MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
                0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
                                       sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
```

```
        if (!a || !b || !c) {
            if (c) HPCC_free(c);
            if (b) HPCC_free(b);
            if (a) HPCC_free(a);
            if (doIO) {
                fprintf( outFile, "Failed to allocate memory (%d).
                \n", VectorSize );
                fclose( outFile );
            }
            return 1;
        }

#ifndef _OPENMP
#pragma omp parallel for
#endif
        for (j=0; j<VectorSize; j++) {
            b[j] = 2.0;
            c[j] = 0.0;
        }

        scalar = 3.0;

#ifndef _OPENMP
#pragma omp parallel for
#endif
        for (j=0; j<VectorSize; j++)
            a[j] = b[j]+scalar*c[j];

        HPCC_free(c);
        HPCC_free(b);
        HPCC_free(a);
```

# STREAM Triad: MPI+OpenMP vs. CUDA

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );
}
```

*HPC suffers from too many distinct notations for expressing parallelism and locality*

```
int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }

    ifdef _OPENMP
    #pragma omp parallel for
    #endiff
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }

    scalar = 3.0;
```

```
#ifdef _OPENMP
# pragma omp parallel for
#endiff
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);

    return 0;
```

## CUDA

```
#define N 2000000
```

```
int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid
```

```
set_array<<<dimGrid,dimBlock>>>(d_b, .5f, N);
set_array<<<dimGrid,dimBlock>>>(d_c, .5f, N);

scalar=3.0f;
STREAM_Triad<<<dimGrid,dimBlock>>>(d_b, d_c, d_a, scalar, N);
cudaThreadSynchronize();
```

```
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
```

```
_global_ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}
```

```
_global_ void STREAM_Triad( float *a, float *b, float *c,
                            float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
```

# Why so many programming models?

HPC has traditionally given users...

...low-level, *control-centric* programming models

...ones that are closely tied to the underlying hardware

...ones that support only a single type of parallelism

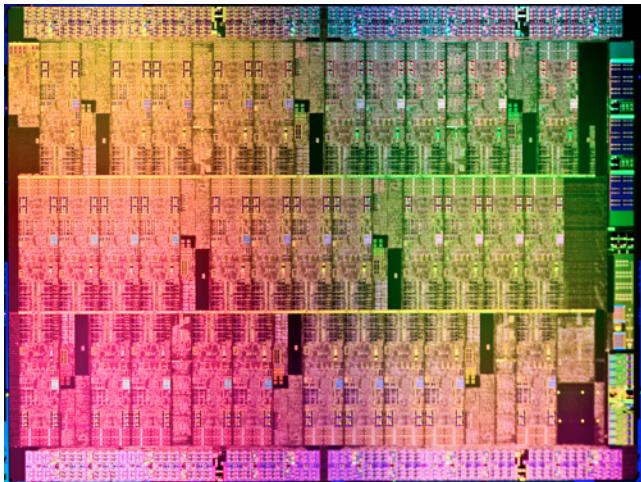
## Examples:

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	MPI	executable
Intra-node/multicore	OpenMP/pthreads	iteration/task
Instruction-level vectors/threads	pragmas	iteration
GPU/accelerator	CUDA/OpenCL/OpenACC	SIMD function/task

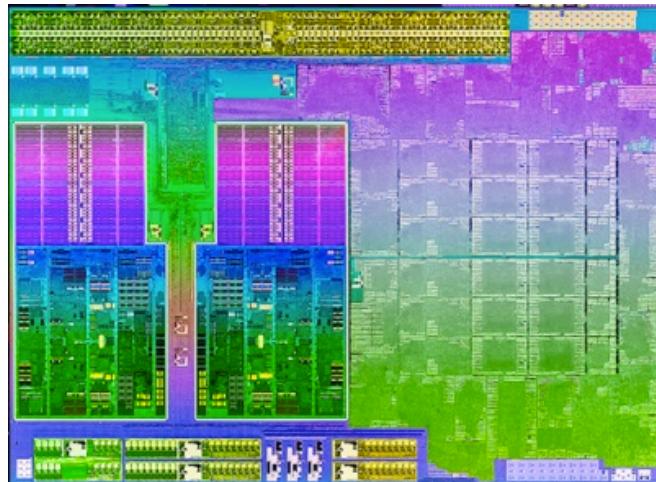
**benefits:** lots of control; decent generality; easy to implement

**downsides:** lots of user-managed detail; brittle to changes

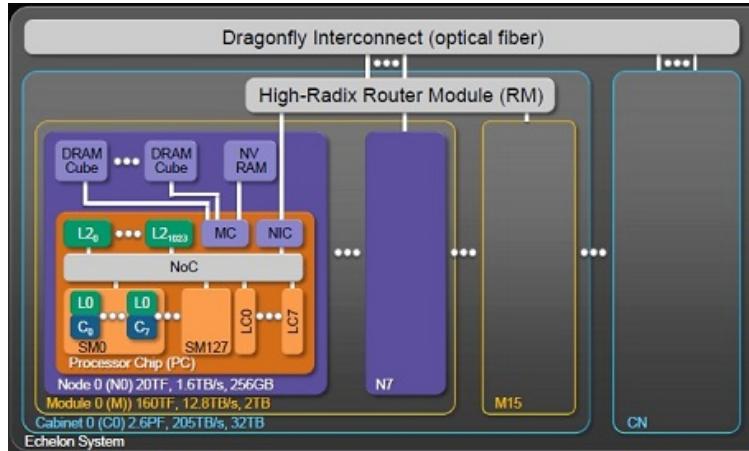
# Prototypical Next-Gen Processor Technologies



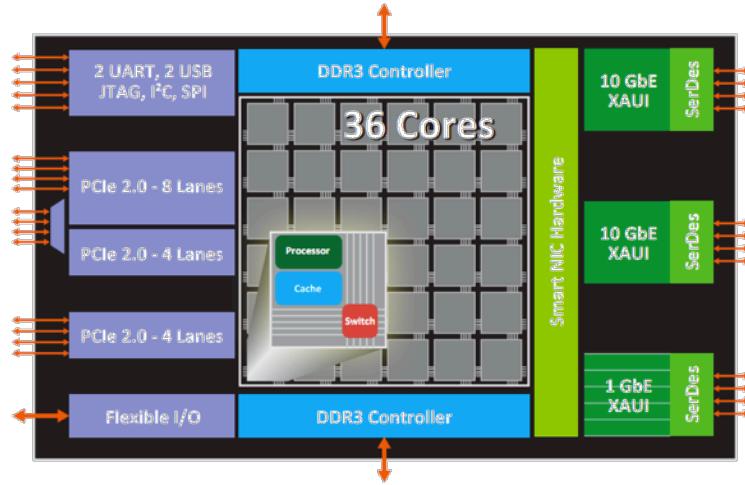
Intel MIC



AMD APU



Nvidia Echelon

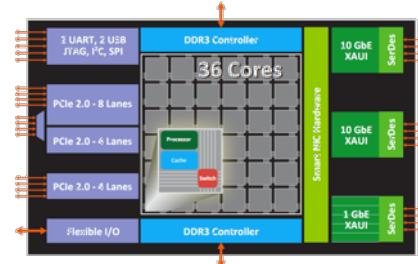
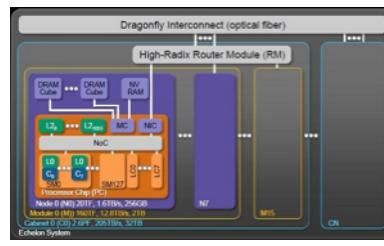
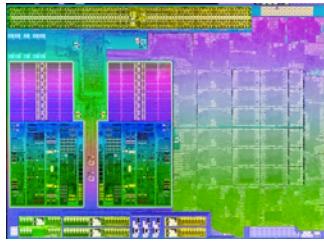
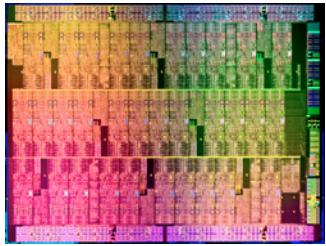


Tilera Tile-Gx

[http://download.intel.com/pressroom/images/Aubrey\\_Isle\\_die.jpg](http://download.intel.com/pressroom/images/Aubrey_Isle_die.jpg) <http://www.zdnet.com/amds-trinity-processors-take-on-intels-ivy-bridge-3040155225/>

<http://insidehpc.com/2010/11/26/nvidia-reveals-details-of-echelon-gpu-designs-for-exascale/> <http://tilera.com/sites/default/files/productbriefs/Tile-Gx%203036%20SB012-01.pdf>

# General Characteristics of These Architectures



- Increased hierarchy and/or sensitivity to locality
- Potentially heterogeneous processor/memory types

⇒ Next-gen programmers will have a lot more to think about at the node level than in the past

# (“Glad I’m not an HPC Programmer!”)

## A Possible Reaction:

“This is all well and good for HPC users, but I’m a mainstream desktop programmer, so this is all academic for me.”

## The Unfortunate Reality:

- Performance-minded mainstream programmers will increasingly need to deal with parallelism and locality too

# Rewinding a few slides...

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }

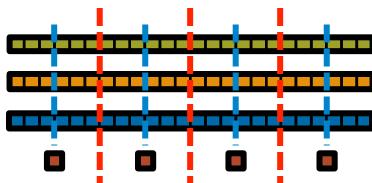
    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }

    scalar = 3.0;

    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);

    return 0;
}
```



## CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid

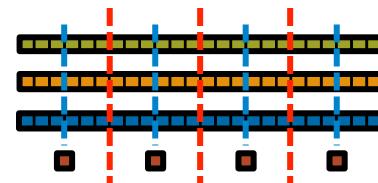
    set_array<<<dimGrid,dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid,dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid,dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad( float *a, float *b, float *c,
                             float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```



# STREAM Triad: Chapel

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *par
int myRank, commSize;
int rv, errCount;
MPI_Comm comm = MPI_COMM_WORLD;

MPI_Comm_size( comm, &commSize );
MPI_Comm_rank( comm, &myRank );

rv = HPCC_Stream( params, 0 == myR
MPI_Reduce( &rv, &errCount, 1, MPI
return errCount;
}

int HPCC_Stream(HPCC_Params *params,
register int j;
double scalar;

VectorSize = HPCC_LocalVectorSize(
a = HPCC_XMALLOC( double, VectorSi
b = HPCC_XMALLOC( double, VectorSi
c = HPCC_XMALLOC( double, VectorSi

if (!a || !b || !c) {
if (c) HPCC_free(c);
if (b) HPCC_free(b);
if (a) HPCC_free(a);
if (doIO) {

```

## Chapel

```
config const m = 1000,
alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

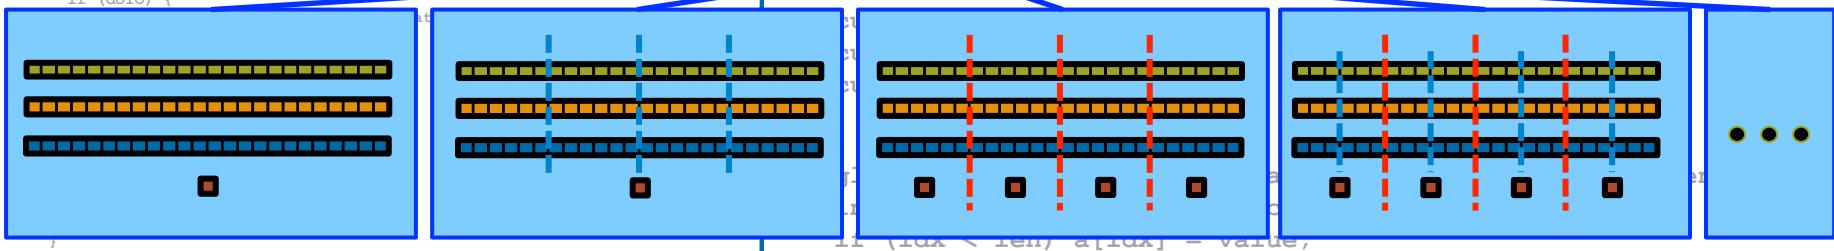
A = B + alpha * C;
```

dmapped ...;

the special sauce

```
N);
N);

_c, d_a, scalar, N);
```



Philosophy: Good language design can tease details of locality and parallelism away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.



# Outline

- ✓ Motivation
- Chapel Background and Themes
  - Survey of Chapel Concepts
  - Project Status and Next Steps



# What is Chapel?

- **An emerging parallel programming language**
  - Design and development led by Cray Inc.
    - in collaboration with academia, labs, industry
  - Initiated under the DARPA HPCS program
- **Overall goal: Improve programmer productivity**
  - Improve the **programmability** of parallel computers
  - Match or beat the **performance** of current programming models
  - Support better **portability** than current programming models
  - Improve the **robustness** of parallel codes
- **A work-in-progress**

# Chapel's Implementation

- Being developed as open source at SourceForge
- Licensed as BSD software
- **Target Architectures:**
  - Cray architectures
  - multicore desktops and laptops
  - commodity clusters and the cloud
  - systems from other vendors
  - *in-progress:* CPU+accelerator hybrids, manycore, ...

# Motivating Chapel Themes

- 1) General Parallel Programming**
- 2) Global-View Abstractions**
- 3) Multiresolution Design**
- 4) Control over Locality/Affinity**
- 5) Reduce HPC ↔ Mainstream Language Gap**

# Motivating Chapel Themes

- 1) General Parallel Programming**
- 2) Global-View Abstractions**
- 3) Multiresolution Design**
- 4) Control over Locality/Affinity**
- 5) Reduce HPC ↔ Mainstream Language Gap**

# 1) General Parallel Programming

With a unified set of concepts...

...express any parallelism desired in a user's program

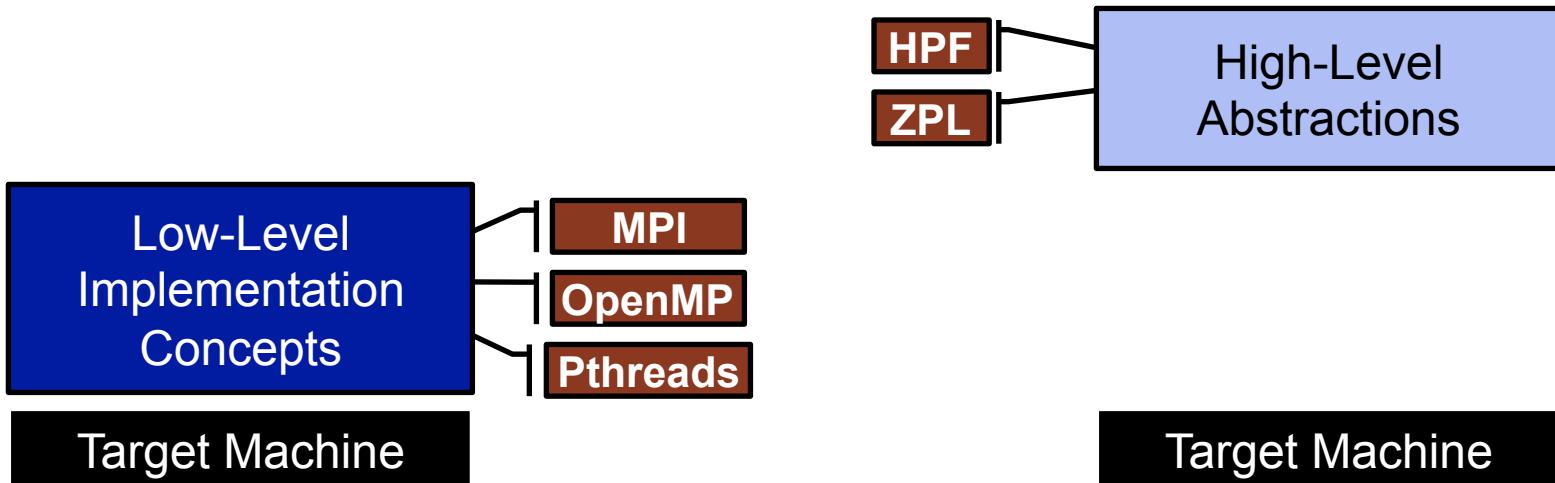
- **Styles:** data-parallel, task-parallel, concurrency, nested, ...
- **Levels:** model, function, loop, statement, expression

...target any parallelism available in the hardware

- **Types:** machines, nodes, cores, instructions

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	Chapel	executable/task
Intra-node/multicore	Chapel	iteration/task
Instruction-level vectors/threads	Chapel	iteration
GPU/accelerator	Chapel	SIMD function/task

### 3) Multiresolution Design: Motivation



*“Why is everything so tedious/difficult?”*

*“Why don’t my programs port trivially?”*

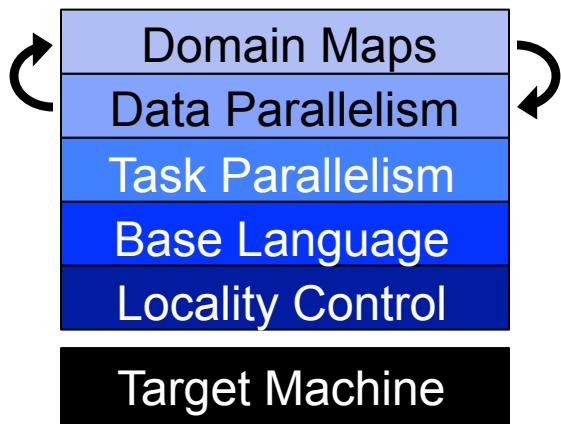
*“Why don’t I have more control?”*

### 3) Multiresolution Design

#### **Multiresolution Design:** Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for greater degrees of control

*Chapel language concepts*



- build the higher-level concepts in terms of the lower
- permit the user to intermix layers arbitrarily

## 5) Reduce HPC ↔ Mainstream Language Gap



### Consider:

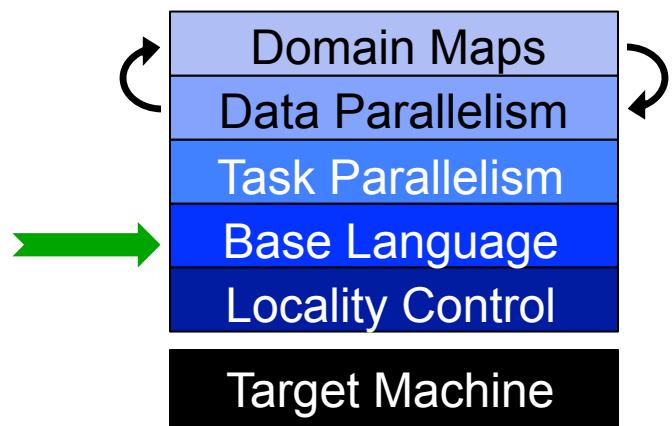
- Students graduate with training in Java, Matlab, Python, etc.
- Yet HPC programming is dominated by Fortran, C/C++, MPI

### We'd like to narrow this gulf with Chapel:

- to leverage advances in modern language design
- to better utilize the skills of the entry-level workforce...
- ...while not alienating the traditional HPC programmer
  - e.g., support object-oriented programming, but make it optional

# Outline

- ✓ Motivation
- ✓ Chapel Background and Themes
- Survey of Chapel Concepts



- Project Status and Next Steps

# Static Type Inference

```
const pi = 3.14,           // pi is a real
      coord = 1.2 + 3.4i, // coord is a complex...
      coord2 = pi*coord,  // ...as is coord2
      name = "brad",      // name is a string
      verbose = false;    // verbose is boolean

proc addem(x, y) {         // addem() has generic arguments
    return x + y;          // and an inferred return type
}

var sum = addem(1, pi),     // sum is a real
    fullname = addem(name, "ford"); // fullname is a string

writeln((sum, fullname));
```

(4.14, bradford)

# Range Types and Algebra

```
const r = 1..10;

printVals(r);
printVals(r # 3);
printVals(r by 2);
printVals(r by -2);
printVals(r by 2 # 3);
printVals(r # 3 by 2);
printVals(0.. #n);

proc printVals(r) {
    for i in r do
        write(r, " ");
    writeln();
}
```

```
1 2 3 4 5 6 7 8 9 10
1 2 3
1 3 5 7 9
10 8 6 4 2
1 3 5
1 3
0 1 2 3 4 ... n-1
```

# Iterators

```
iter fibonacci(n) {
  var current = 0,
       next = 1;
  for 1..n {
    yield current;
    current += next;
    current <=> next;
  }
}
```

```
for f in fibonacci(7) do
  writeln(f);
```

```
0
1
1
2
2
3
5
8
```

```
iter tiledRMO(D, tilesize) {
  const tile = {0..#tilesize,
                0..#tilesize};
  for base in D by tilesize do
    for ij in D[tile + base] do
      yield ij;
}
```

```
for ij in tiledRMO({1..m, 1..n}, 2) do
  write(ij);
```

```
(1,1) (1,2) (2,1) (2,2)
(1,3) (1,4) (2,3) (2,4)
(1,5) (1,6) (2,5) (2,6)
...
(3,1) (3,2) (4,1) (4,2)
```

# Zippered Iteration

```
for (i,f) in zip(0..#n, fibonacci(n)) do  
    writeln("fib #", i, " is ", f);
```

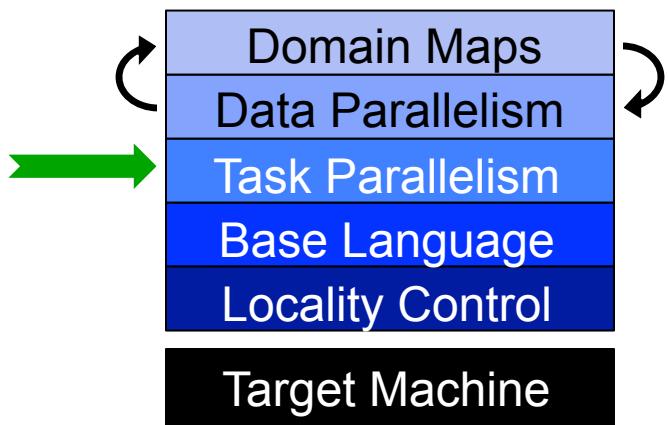
```
fib #0 is 0  
fib #1 is 1  
fib #2 is 1  
fib #3 is 2  
fib #4 is 3  
fib #5 is 5  
fib #6 is 8  
...
```

# Other Base Language Features

- **tuple types and values**
- **rank-independent programming features**
- **interoperability features**
- **compile-time features for meta-programming**
  - e.g., compile-time functions to compute types, parameters
- **OOP (value- and reference-based)**
- **argument intents, default values, match-by-name**
- **overloading, where clauses**
- **modules (for namespace management)**
- ...

# Outline

- ✓ Motivation
- ✓ Chapel Background and Themes
- Survey of Chapel Concepts



- Project Status and Next Steps

# Task Parallelism: Begin Statements

```
// create a fire-and-forget task for a statement
begin writeln("hello world");
writeln("goodbye");
```

## Possible outputs:

hello world  
goodbye

goodbye  
hello world

# Task Parallelism: Coforall Loops

```
// create a task per iteration
coforall t in 0..#numTasks {
    writeln("Hello from task ", t, " of ", numTasks);
} // implicit join of the numTasks tasks here

writeln("All tasks done");
```

## Sample output:

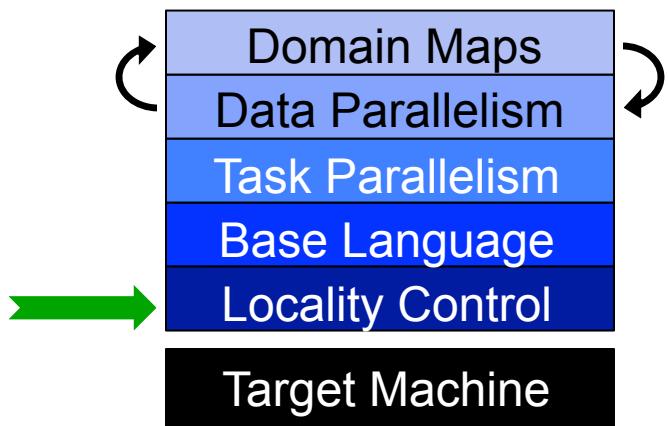
```
Hello from task 2 of 4
Hello from task 0 of 4
Hello from task 3 of 4
Hello from task 1 of 4
All tasks done
```

# Other Task Parallel Concepts

- **cobegin**: create tasks using compound statements
- **sync statements**: join unstructured tasks
- **serial statements**: conditionally squash parallelism
- **atomic variables**: support atomics ops, similar to modern C++
- **sync/single variables**: support producer/consumer patterns

# Outline

- ✓ Motivation
- ✓ Chapel Background and Themes
- Survey of Chapel Concepts



Theme 4: Control over  
Locality/Affinity

- Project Status and Next Steps

# The Locale Type

## Definition:

- Abstract unit of target architecture
- Supports reasoning about locality
- Capable of running tasks and storing variables
  - i.e., has processors and memory

**Typically:** A compute node (multicore processor or SMP)

# Defining Locales

- Specify # of locales when running Chapel programs

```
% a.out --numLocales=8
```

```
% a.out -nl 8
```

- Chapel provides built-in locale variables

```
config const numLocales: int = ...;  
const Locales: [0..#numLocales] locale = ...;
```

*Locales*



- User's main() begins executing on locale #0

# Locale Operations

- Locale methods support queries about the target system:

```
proc locale.physicalMemory(...) { ... }  
proc locale.numCores { ... }  
proc locale.id { ... }  
proc locale.name { ... }
```

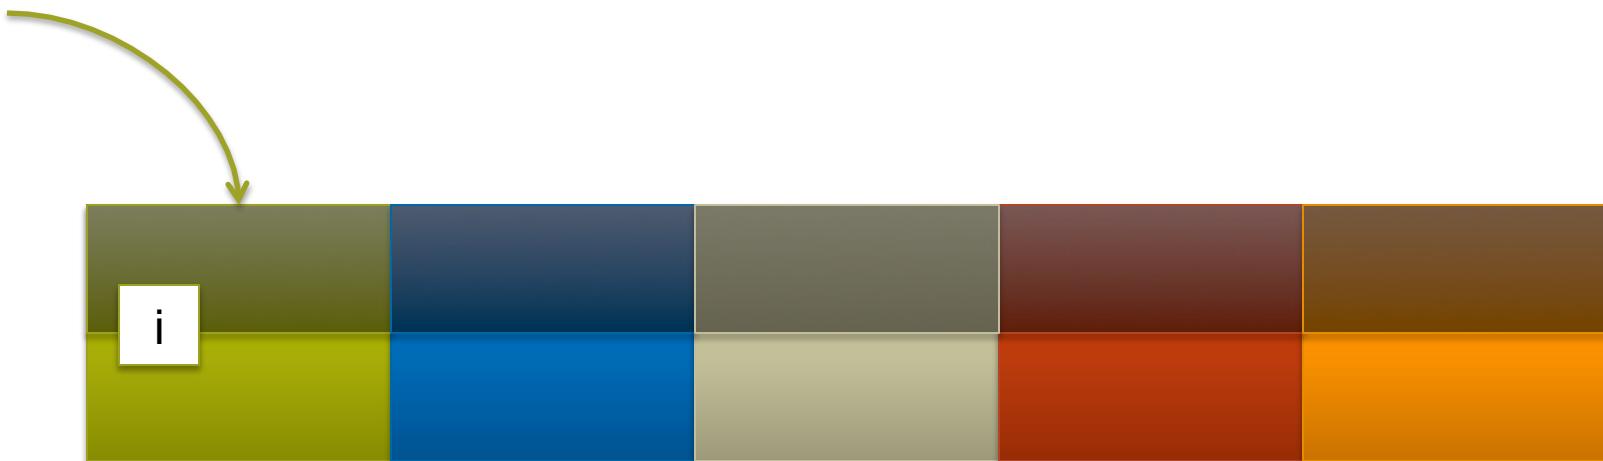
- On-clauses support placement of computations:

```
writeln("on locale 0");  
on Locales[1] do  
  writeln("now on locale 1");  
writeln("on locale 0 again");
```

```
begin on A[i,j] do  
  bigComputation(A);  
  
begin on node.left do  
  search(node.left);
```

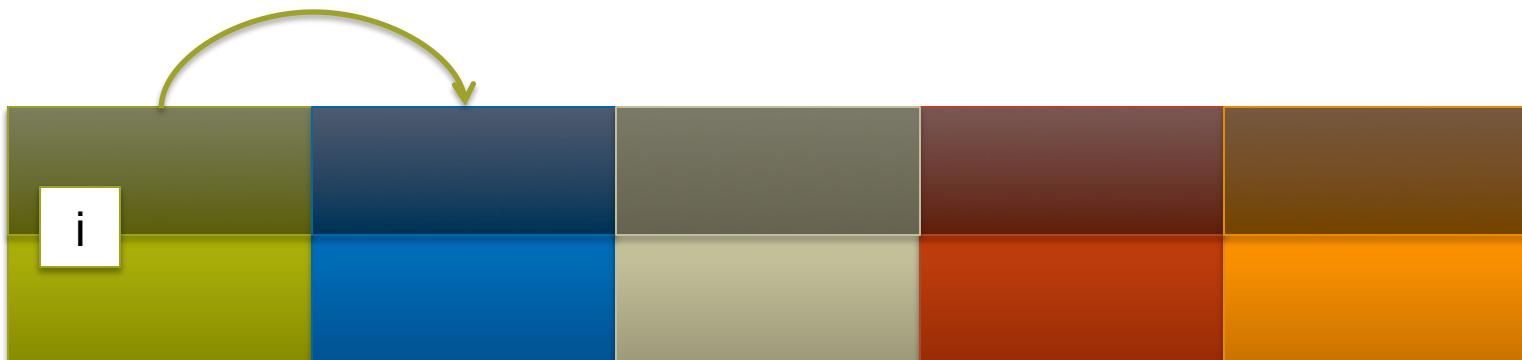
# Chapel: Scoping and Locality

```
var i: int;
```



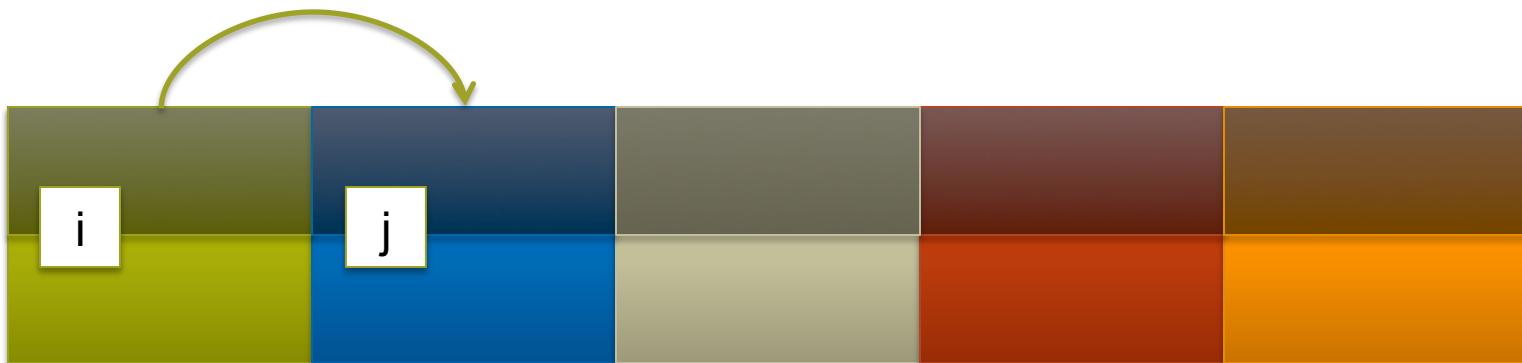
# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {
```



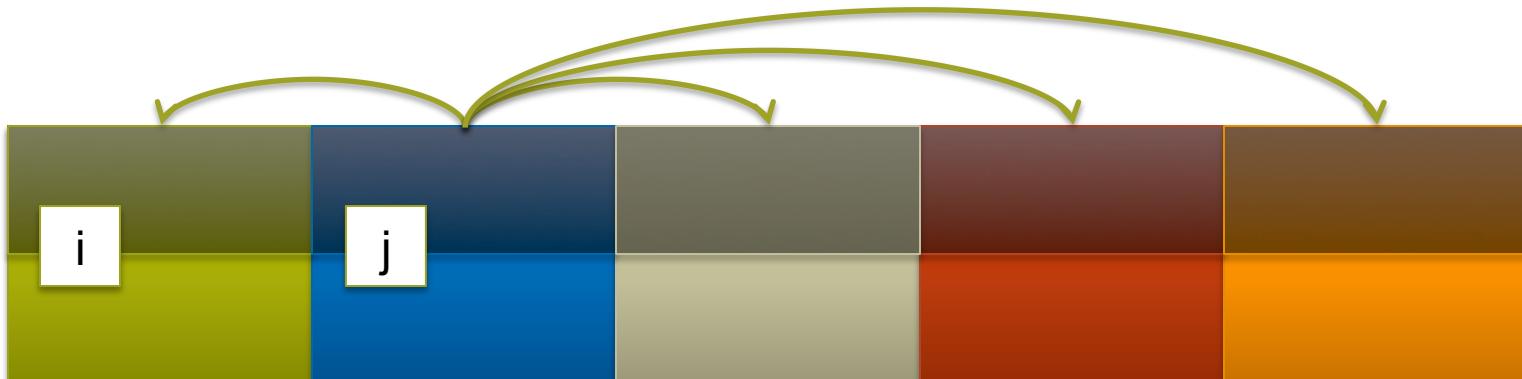
# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;
```



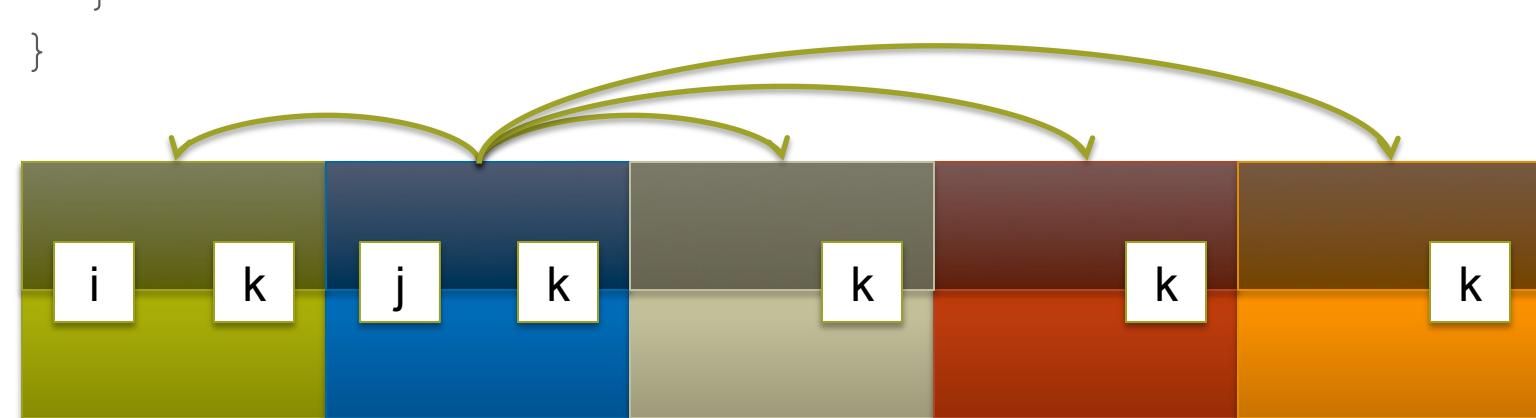
# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {
```



# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
  
            // within this scope, i, j, and k can be referenced;  
            // the implementation manages the communication for i and j  
        }  
    }  
}
```

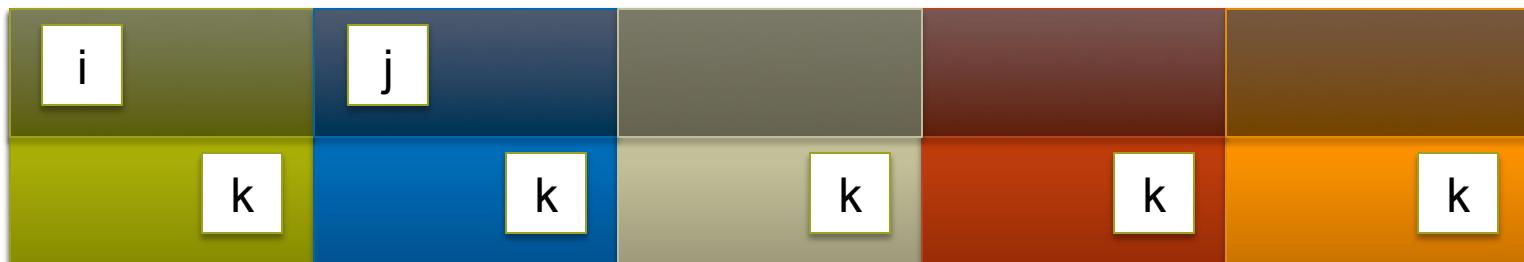


# Chapel and PGAS: Public vs. Private

## How public a variable is depends only on scoping

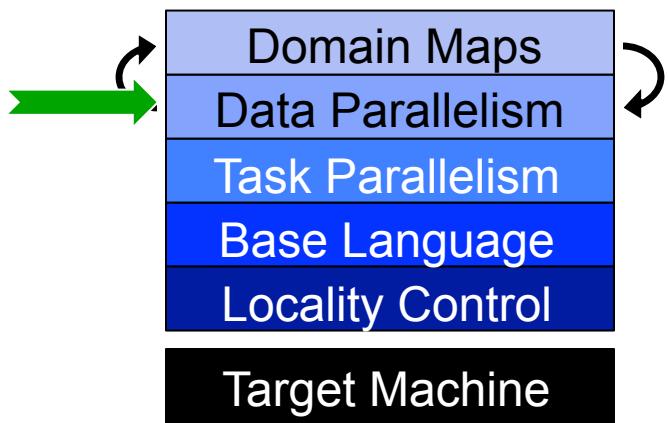
- who can see it?
- who actually bothers to refer to it non-locally?

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k = i + j;  
        }  
    }  
}
```



# Outline

- ✓ Motivation
- ✓ Chapel Background and Themes
- Survey of Chapel Concepts



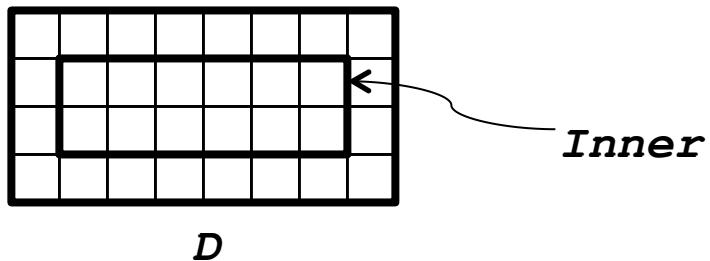
- Project Status and Next Steps

# Domains

## Domain:

- A first-class index set
- The fundamental Chapel concept for data parallelism

```
config const m = 4, n = 8;  
  
const D = {1..m, 1..n};  
const Inner = {2..m-1, 2..n-1};
```

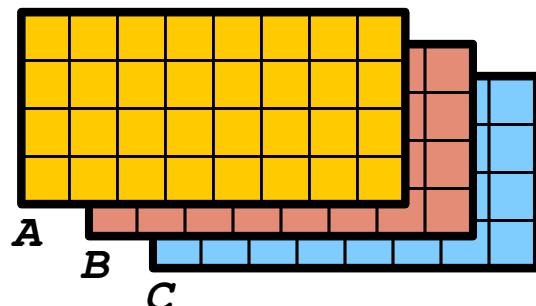


# Domains

## Domain:

- A first-class index set
- The fundamental Chapel concept for data parallelism
- Useful for declaring arrays and computing with them

```
config const m = 4, n = 8;  
  
const D = {1..m, 1..n};  
const Inner = {2..m-1, 2..n-1};  
  
var A, B, C: [D] real;
```



# Chapel Domain/Array Operations

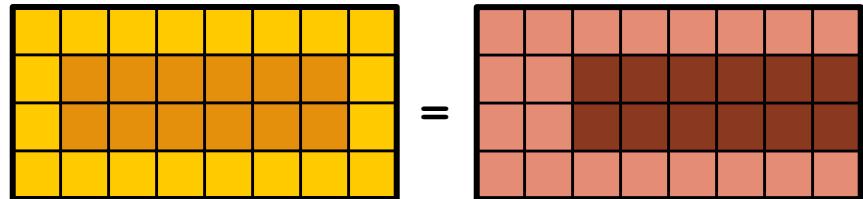
- Data Parallel Iteration

```
forall (i,j) in D do
    A[i,j] = i + j/10.0;
```

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

- Array Slicing; Domain Algebra

```
A[InnerD] = B[InnerD+(0,1)];
```



- Promotion of Scalar Functions and Operators

```
A = exp(B, C);
```

```
A = foo("hi", B, C);
```

```
A = B + alpha * C;
```

- And many others: reductions, scans, reallocation, reshaping, remapping, set operations, aliasing, ...

# Notes on Forall Loops

```
forall a in A do  
    writeln("Here is an element of A: ", a);
```

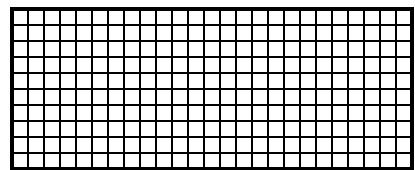
Typically:

- $1 \leq \# \text{Tasks} \ll \# \text{Iterations}$
- $\# \text{Tasks} \approx \text{amount of HW parallelism}$

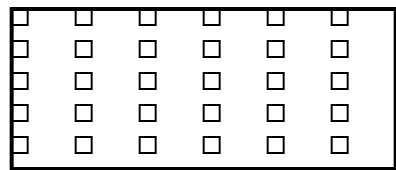
```
forall (a, i) in zip(A, 1..n) do  
    a = i / 10.0;
```

Like for loops, forall-loops may be zippered, and corresponding iterations will match up

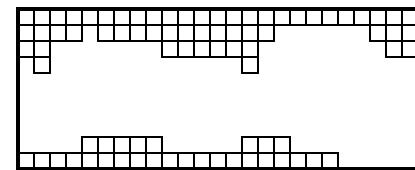
# Chapel Domain Types



*dense*



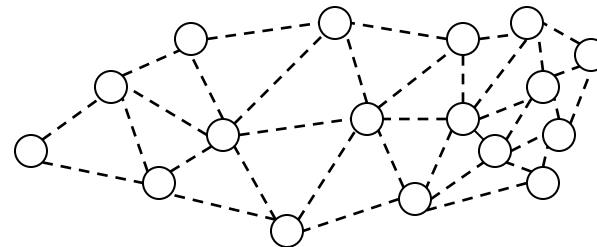
*strided*



*sparse*

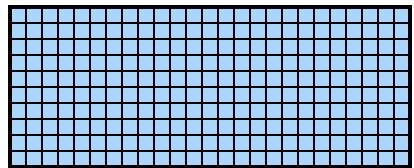


*associative*

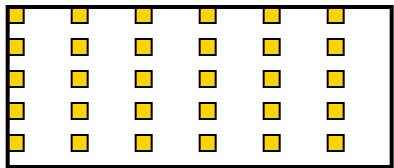


*unstructured*

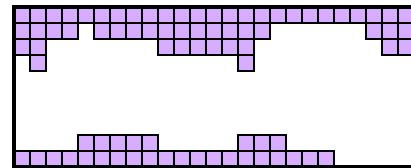
# Chapel Array Types



*dense*



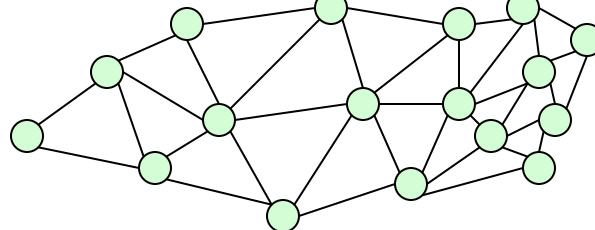
*strided*



*sparse*



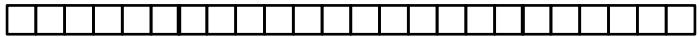
*associative*



*unstructured*

# STREAM Triad in Chapel

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```

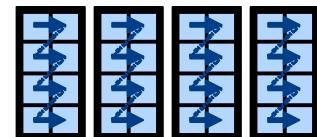
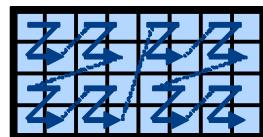
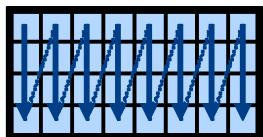
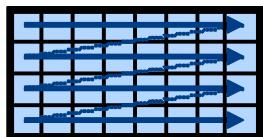


```
A = B + alpha * C;
```

# Data Parallelism Implementation Qs

## Q1: How are arrays laid out in memory?

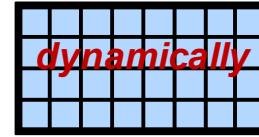
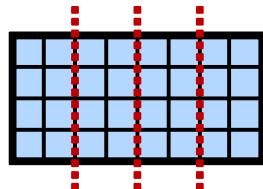
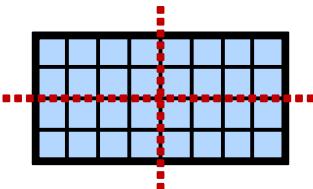
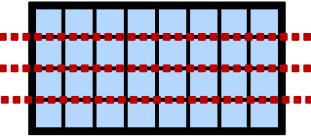
- Are regular arrays laid out in row- or column-major order? Or...?



- How are sparse arrays stored? (COO, CSR, CSC, block-structured, ...?)

## Q2: How are arrays stored by the locales?

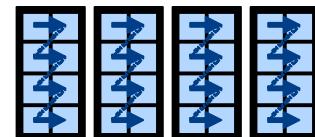
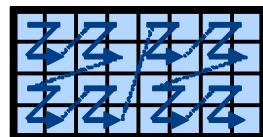
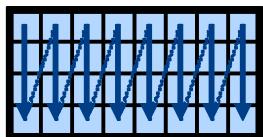
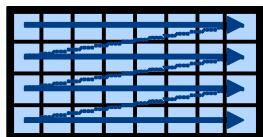
- Completely local to one locale? Or distributed?
- If distributed... In a blocked manner? cyclically? block-cyclically?  
recursively bisected? dynamically rebalanced? ...?



# Data Parallelism Implementation Qs

## Q1: How are arrays laid out in memory?

- Are regular arrays laid out in row- or column-major order? Or...?



- How are sparse arrays stored? (COO, CSR, CSC, block-structured, ...?)

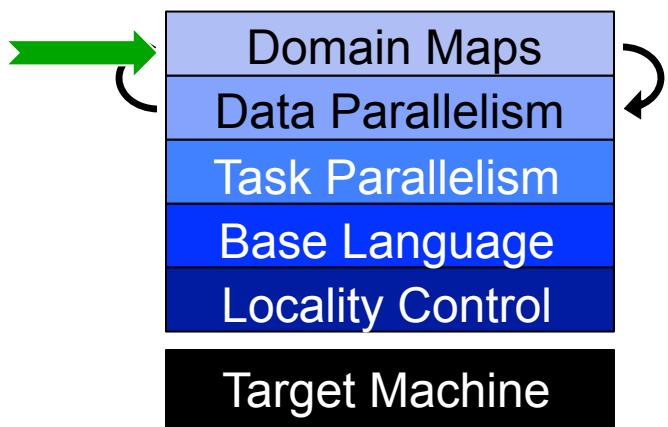
## Q2: How are arrays stored by the locales?

- Completely local to one locale? Or distributed?
- If distributed... In a blocked manner? cyclically? block-cyclically? recursively bisected? dynamically rebalanced? ...?

A: Chapel's *domain maps* are designed to give the user full control over such decisions

# Outline

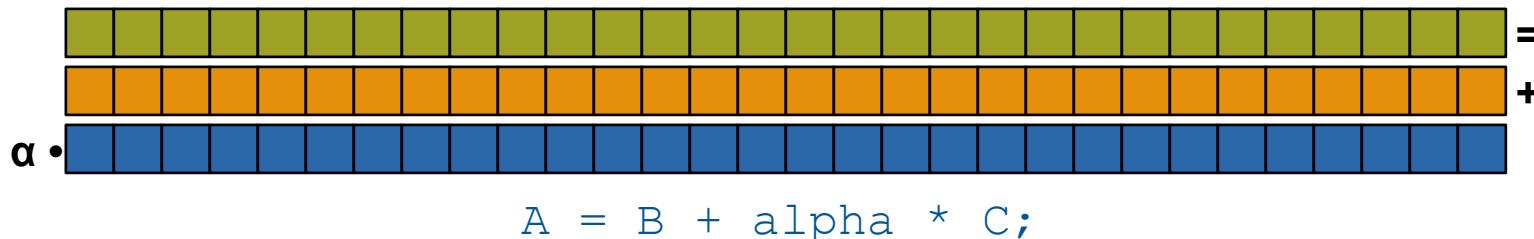
- ✓ Motivation
- ✓ Chapel Background and Themes
- Survey of Chapel Concepts



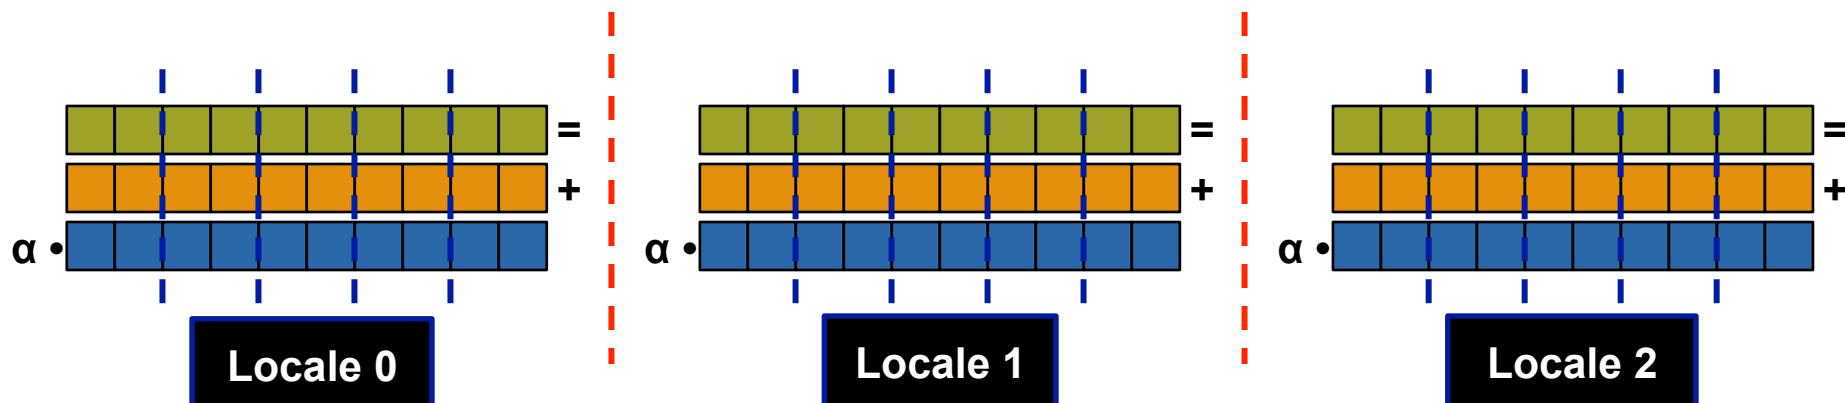
- Project Status and Next Steps

# Domain Maps

Domain maps are “recipes” that instruct the compiler how to map the global view of a computation...

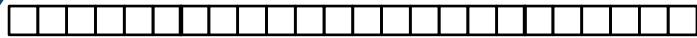


...to the target locales' memory and processors:

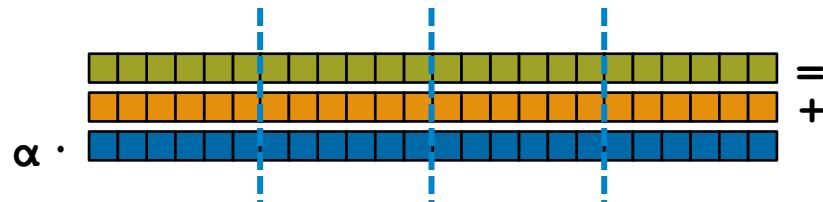


# STREAM Triad: Chapel (multicore)

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```

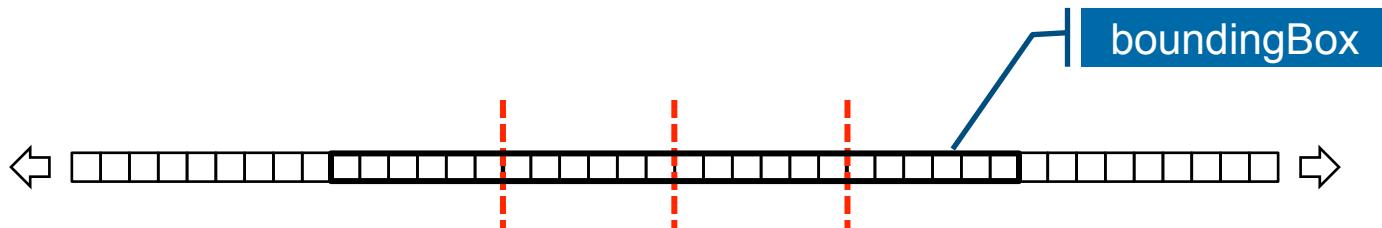


```
A = B + alpha * C;
```

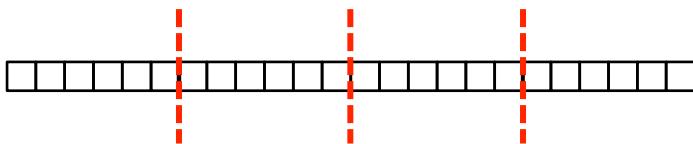
No domain map specified => use default layout

- current locale owns all domain indices and array values
- computation will execute using local processors only

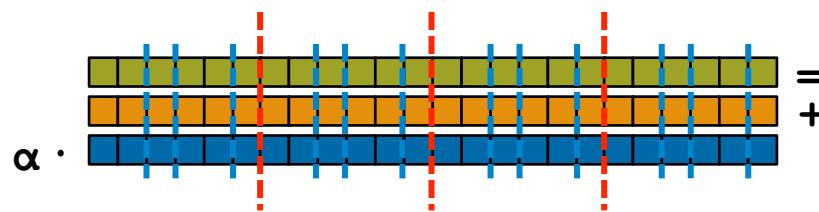
# STREAM Triad: Chapel (multilocale, blocked)



```
const ProblemSpace = {1..m}
dmapped Block(boundingBox={1..m});
```

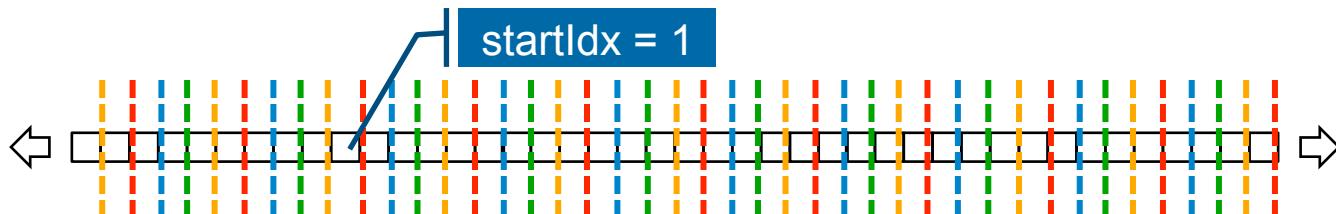


```
var A, B, C: [ProblemSpace] real;
```

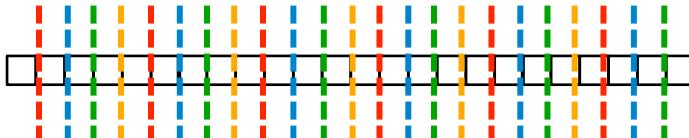


```
A = B + alpha * C;
```

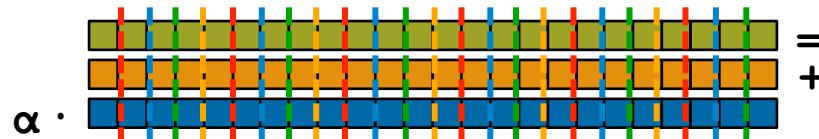
# STREAM Triad: Chapel (multilocale, cyclic)



```
const ProblemSpace = {1..m}
    dmapped Cyclic(startIdx=1);
```



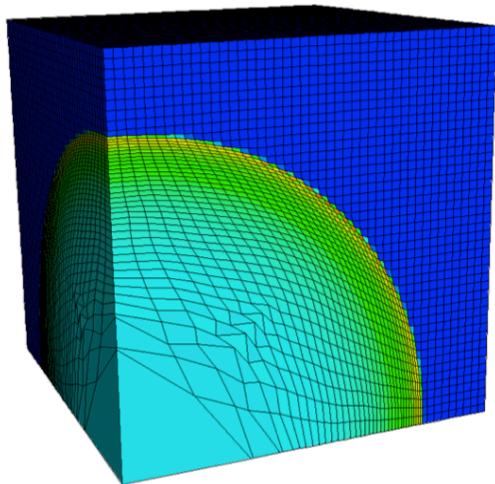
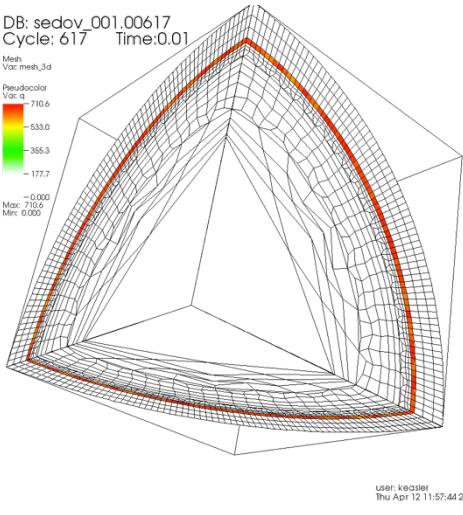
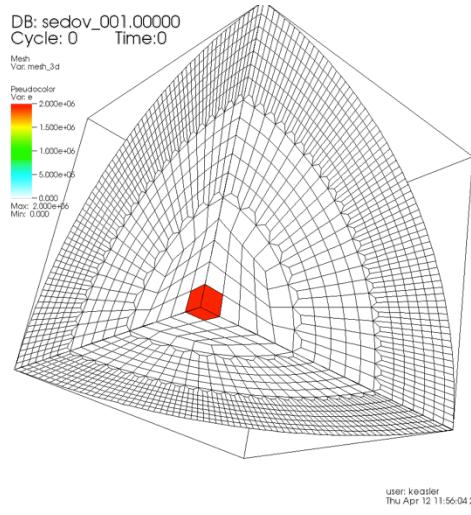
```
var A, B, C: [ProblemSpace] real;
```



```
A = B + alpha * C;
```

# LULESCH: a DOE Proxy Application

**Goal:** Solve one octant of the spherical Sedov problem (blast wave) using Lagrangian hydrodynamics for a single material



pictures courtesy of Rob Neely, Bert Still, Jeff Keasler, LLNL

# LULESH in Chapel



# LULESCH in Chapel

**1288 lines of source code**

plus    266 lines of comments  
        487 blank lines

(the corresponding C+MPI+OpenMP version is nearly 4x bigger)

This can be found in Chapel v1.9 in examples/benchmarks/lulesh/\*.chpl

# LULESCH in Chapel

This is all of the representation dependent code.  
It specifies:

- data structure choices
  - structured vs. unstructured mesh
  - local vs. distributed data
  - sparse vs. dense materials arrays
- a few supporting iterators

# LULESH in Chapel



Here is some sample representation-independent code  
`IntegrateStressForElems ()`  
LULESH spec, section 1.5.1.1 (2.)

# Representation-Independent Physics

```

proc IntegrateStressForElems(sigxx, sigyy, sigzz, determ) {
    forall k in Elems { ← parallel loop over elements
        var b_x, b_y, b_z: 8*real;
        var x_local, y_local, z_local: 8*real;
        localizeNeighborNodes(k, x, x_local, y, y_local, z, z_local); ← collect nodes neighboring this
        element; localize node fields
        var fx_local, fy_local, fz_local: 8*real;

        local {
            /* Volume calculation involves extra work for numerical consistency. */
            CalcElemShapeFunctionDerivatives(x_local, y_local, z_local,
                b_x, b_y, b_z, determ[k]);

            CalcElemNodeNormals(b_x, b_y, b_z, x_local, y_local, z_local);

            SumElemStressesToNodeForces(b_x, b_y, b_z, sigxx[k], sigyy[k], sigzz[k],
                fx_local, fy_local, fz_local);
        }
        for (noi, t) in elemToNodesTuple(k) { ← update node forces from
            fx[noi].add(fx_local[t]); element stresses
            fy[noi].add(fy_local[t]);
            fz[noi].add(fz_local[t]);
        }
    }
}

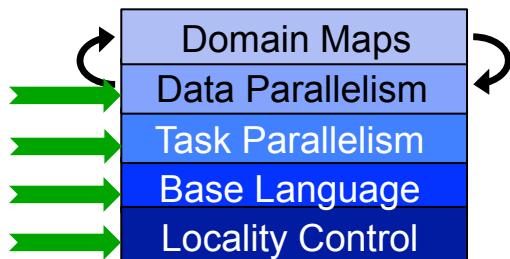
```

Because of domain maps, this code is independent of:

- structured vs. unstructured mesh
- shared vs. distributed data
- sparse vs. dense representation

# Chapel's Domain Map Philosophy

- 1. Chapel provides a library of standard domain maps**
  - to support common array implementations effortlessly
- 2. Expert users can write their own domain maps in Chapel**
  - to cope with any shortcomings in our standard library



- 3. Chapel's standard domain maps are written using the same end-user framework**
  - to avoid a performance cliff between “built-in” and user-defined cases

# Domain Map Descriptors

## Domain Map

**Represents:** a domain map value

**Generic w.r.t.:** index type

**State:** the domain map's representation

**Typical Size:**  $\Theta(1)$

**Required Interface:**

- create new domains

## Domain

**Represents:** a domain

**Generic w.r.t.:** index type

**State:** representation of index set

**Typical Size:**  $\Theta(1) \rightarrow \Theta(numIndices)$

**Required Interface:**

- create new arrays
- queries: size, members
- iterators: serial, parallel
- domain assignment
- index set operations

## Array

**Represents:** an array

**Generic w.r.t.:** index type, element type

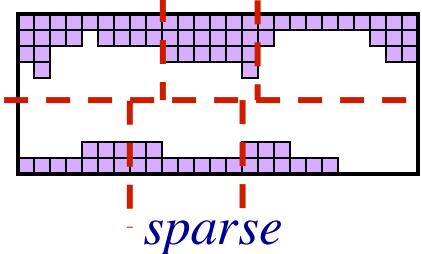
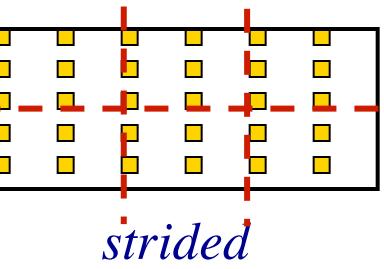
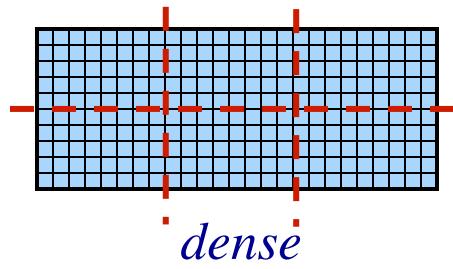
**State:** array elements

**Typical Size:**  $\Theta(numIndices)$

**Required Interface:**

- (re-)allocation of elements
- random access
- iterators: serial, parallel
- slicing, reindexing, aliases
- get/set of sparse “zero” values

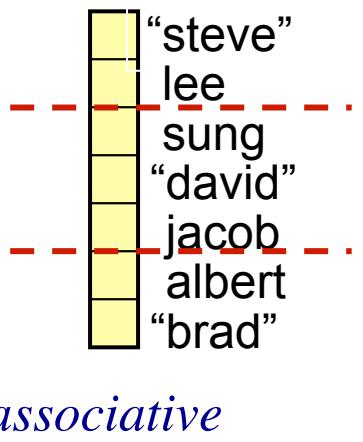
# All Domain Types Support Domain Maps



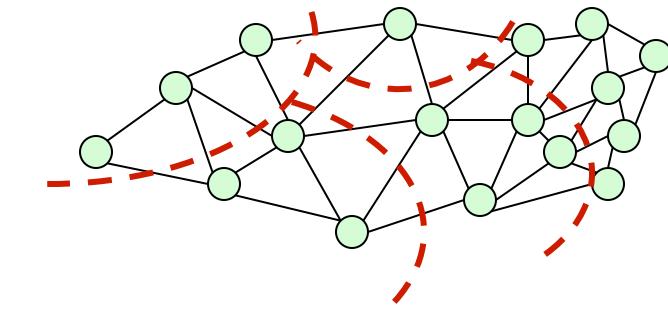
*dense*

*strided*

*sparse*



*associative*



*unstructured*

# For More Information on Domain Maps

**HotPAR'10:** *User-Defined Distributions and Layouts in Chapel: Philosophy and Framework*

Chamberlain, Deitz, Iten, Choi; June 2010

**CUG 2011:** *Authoring User-Defined Domain Maps in Chapel*

Chamberlain, Choi, Deitz, Iten, Litvinov; May 2011

## Chapel release:

- Current domain maps:  
  \$CHPL\_HOME/modules/dists/\*.chpl  
    layouts/\*.chpl  
    internal/Default\*.chpl
- Technical notes detailing the domain map interface for implementers:  
  \$CHPL\_HOME/doc/technotes/README.dsi

## Two Other Thematically Similar Features

- 1) **leader-follower iterators:** Permit users to specify the parallelism and work decomposition used by forall loops
  - including zippered forall loops
- 2) **locale models:** Permit users to model the target architecture and how Chapel should be implemented on it
  - e.g., how to manage memory, create tasks, communicate, ...

Like domain maps, these are...

- ...written in Chapel by expert users using lower-level features
  - e.g., task parallelism, on-clauses, base language features, ...
- ...available to the end-user via higher-level abstractions
  - e.g., forall loops, on-clauses, lexically scoped PGAS memory, ...

# Summary of this Section

- **Chapel avoids locking crucial implementation decisions into the language specification**
  - local and distributed parallel array implementations
  - parallel loop scheduling policies
  - target architecture models
- **Instead, these can be...**
  - ...specified in the language by an advanced user
  - ...swapped between with minimal code changes
- **The result cleanly separates the roles of domain scientist, parallel programmer, and compiler/runtime**



# Outline

- ✓ Motivation
- ✓ Chapel Background and Themes
- ✓ Survey of Chapel Concepts
- Project Status and Next Steps

# Implementation Status -- Version 1.9.0 (Apr 2014)

## Overall Status:

- **User-facing Features:** generally in good shape
  - some require additional attention (e.g., strings, OOP)
- **Multiresolution Features:** in use today
  - their interfaces are likely to continue evolving over time
- **Error Messages:** not always as helpful as one would like
  - correct code works well, incorrect code can be puzzling
- **Performance:** hit-or-miss depending on the idioms used
  - Chapel designed to ultimately support competitive performance
  - effort to-date has focused primarily on correctness

## This is a good time to:

- Try out the language and compiler
- Use Chapel for non-performance-critical projects
- Give us feedback to improve Chapel
- Use Chapel for parallel programming education

# Chapel and Education

- When teaching parallel programming, I like to cover:
  - data parallelism
  - task parallelism
  - concurrency
  - synchronization
  - locality/affinity
  - deadlock, livelock, and other pitfalls
  - performance tuning
  - ...
- I don't think there's been a good language out there...
  - for teaching *all* of these things
  - for teaching *some* of these things well at all
  - *until now:* We believe Chapel can play a crucial role here

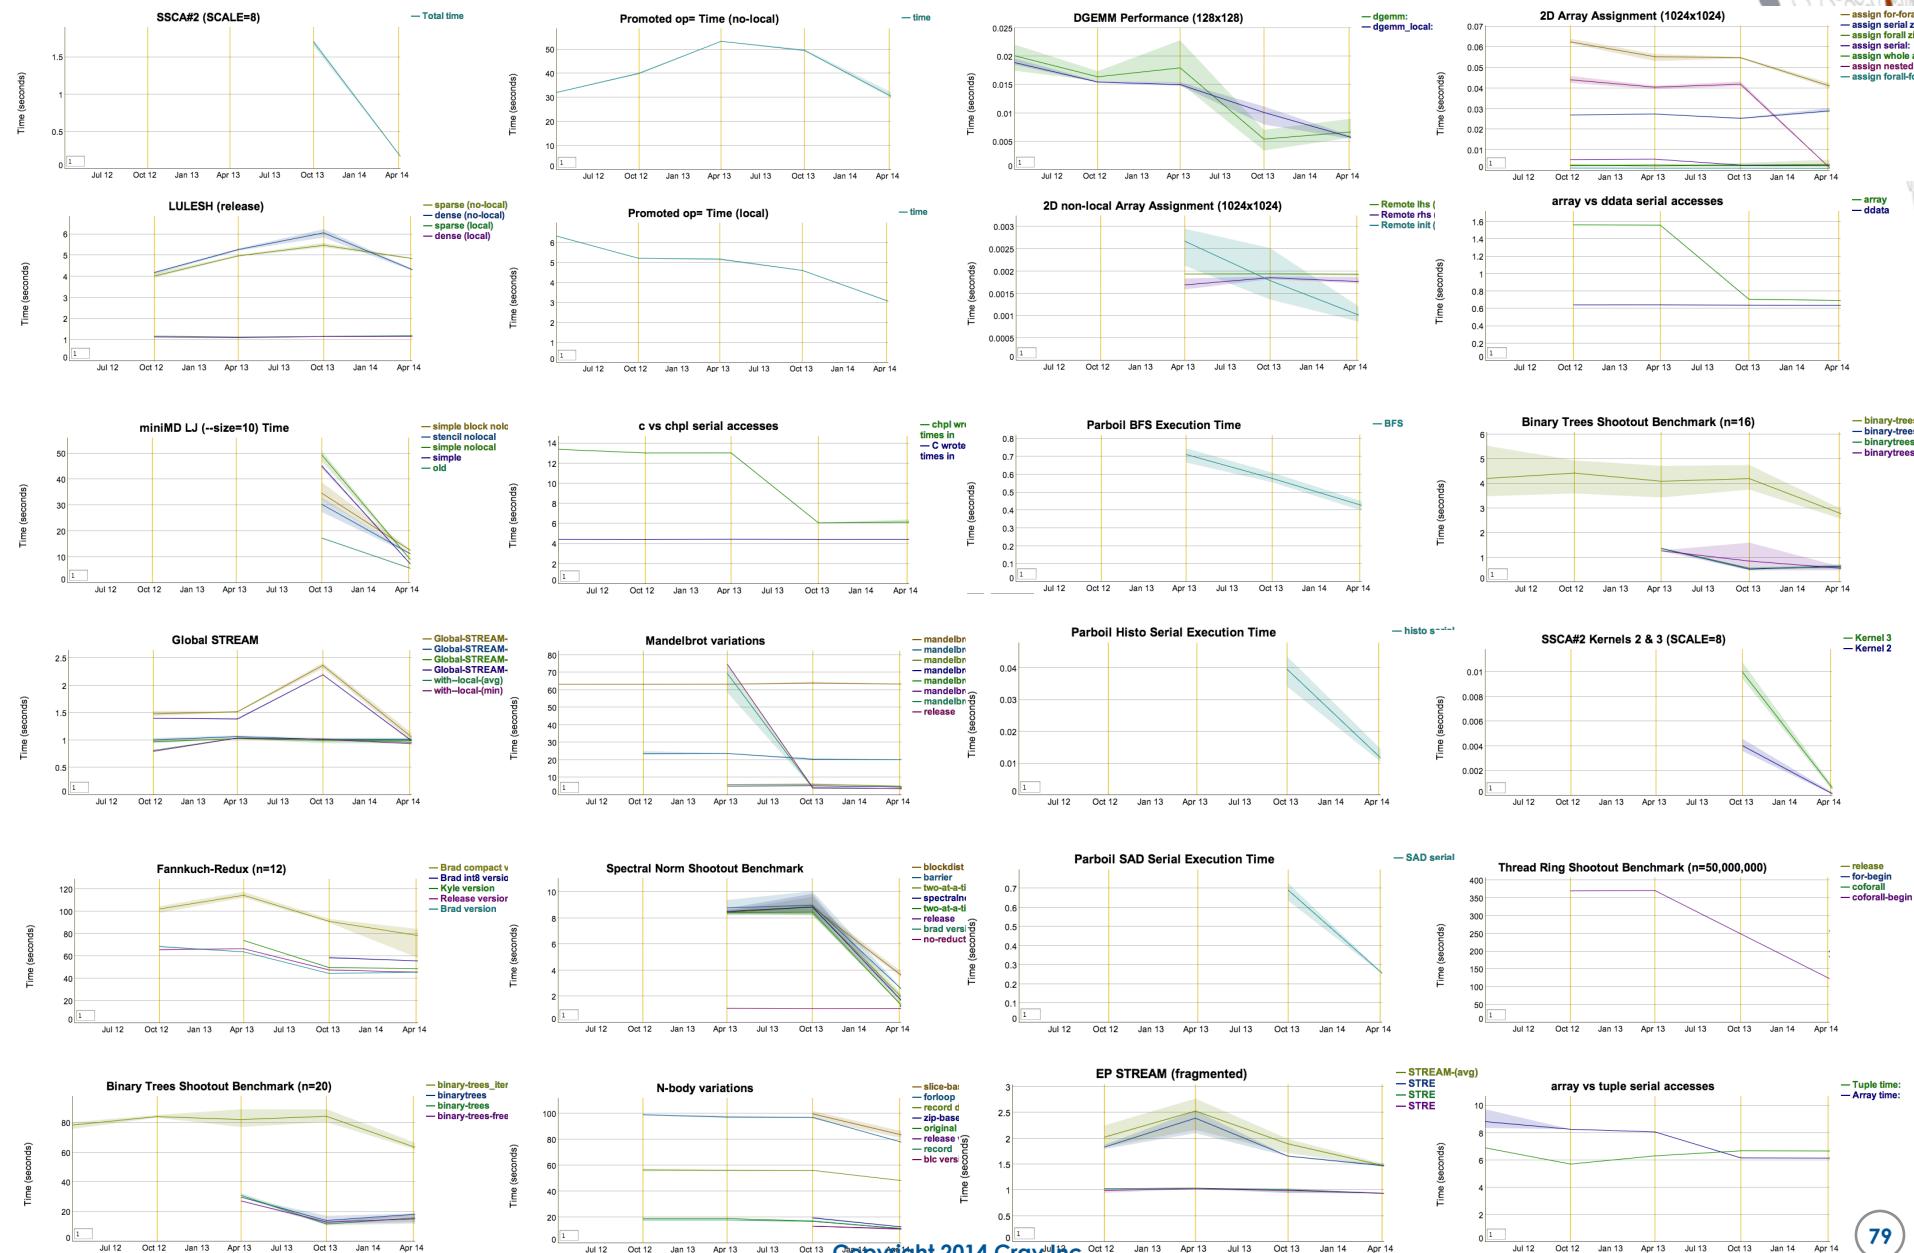
(see <http://chapel.cray.com/education.html> for more information and  
<http://cs.washington.edu/education/courses/csep524/13wi/> for my use of Chapel in class)

# Chapel: the next five years

- **Harden prototype to production-grade**
  - add/improve lacking features
  - optimize performance
- **Target more complex/modern compute node types**
  - e.g., Intel MIC, CPU+GPU, AMD APU, ...
- **Continue to grow the user and developer communities**
  - including nontraditional circles: desktop parallelism, “big data”
  - transition Chapel from Cray-managed to community-governed

# Chapel Execution Time Trends (v1.5–1.9)

(available for browsing at: <http://chapel.sourceforge.net/perf/>)



# The Cray Chapel Team (Summer 2013)





# Chapel...

...is a collaborative effort — join us!



Sandia National Laboratories



Lawrence Livermore  
National Laboratory



Lawrence Berkeley  
National Laboratory



# Summary

***Higher-level programming models can help insulate algorithms from parallel implementation details***

- yet, without necessarily abdicating control
- Chapel does this via its multiresolution design
  - here, we saw it principally in domain maps
    - leader-follower iterators and locale models are other examples
  - these avoid locking crucial policy decisions into the language

***We believe Chapel can greatly improve productivity***

...for current and emerging HPC architectures

...for emerging mainstream needs for parallelism and locality

# For More Information: Online Resources

## Chapel project page: <http://chapel.cray.com>

- overview, papers, presentations, language spec, ...

## Chapel SourceForge page: <https://sourceforge.net/projects/chapel/>

- release downloads, community mailing lists, repository, ...

## Mailing Aliases:

- [chapel\\_info@cray.com](mailto:chapel_info@cray.com): contact the team at Cray
- [chapel-users@lists.sourceforge.net](mailto:chapel-users@lists.sourceforge.net): user-oriented discussion list
- [chapel-developers@lists.sourceforge.net](mailto:chapel-developers@lists.sourceforge.net): developer discussion
- [chapel-education@lists.sourceforge.net](mailto:chapel-education@lists.sourceforge.net): educator discussion
- [chapel-bugs@lists.sourceforge.net](mailto:chapel-bugs@lists.sourceforge.net): public bug forum

# For More Information: Suggested Reading

## Overview Papers:

- A Brief Overview of Chapel, Chamberlain (pre-print of a chapter for *A Brief Overview of Parallel Programming Models*, edited by Pavan Balaji, to be published by MIT Press in 2014).
  - *a more detailed overview of Chapel's history, motivating themes, features*
- The State of the Chapel Union [slides], Chamberlain, Choi, Dumler, Hildebrandt, Iten, Litvinov, Titus. CUG 2013, May 2013.
  - *a high-level overview of the project summarizing the HPCS period*

## Blog Series:

- [Ten] Myths About Scalable Programming Languages, Chamberlain. IEEE Technical Committee on Scalable Computing (TCSC) Blog, (<https://www.ieeetcsc.org/activities/blog/>), April-November 2012.
  - *a series of technical opinion pieces designed to combat standard arguments against the development of high-level parallel languages*

# Legal Disclaimer

*Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.*

*Cray Inc. may make changes to specifications and product descriptions at any time, without notice.*

*All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.*

*Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.*

*Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.*

*Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.*

*The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, URIKA, and YARCDATA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.*

Copyright 2014 Cray Inc.



**CRAY**  
THE SUPERCOMPUTER COMPANY

<http://chapel.cray.com>

[chapel\\_info@cray.com](mailto:chapel_info@cray.com)

<http://sourceforge.net/projects/chapel/>