



Chapel: Overview and Features for Heterogeneity

Brad Chamberlain, Chapel Team, Cray Inc.

Heterogeneous C++ Working Group

February 5, 2018



Safe Harbor Statement



This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.



Plan for this morning



- Chapel context
- Overview of example core Chapel features
- Additional detail on features related to heterogeneity
- Feel free to interrupt for Questions / Discussion



COMPUTE

|

STORE

|

ANALYZE

Chapel and Heterogeneity in a Nutshell



- Chapel's design supports hardware heterogeneity
 - and in a way that is user-extensible without compiler changes
- That said, in practice...
 - ...virtually all of our work has been on homogeneous cases
 - ...we haven't spent much time on many of the hardest cases
 - e.g., no FPGA work, less GPU work than we'd like



What is Chapel?



COMPUTE



STORE



ANALYZE

What is Chapel?



Chapel: A productive parallel programming language

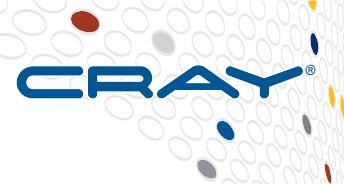
- portable
- open-source
- a collaborative effort

Goals:

- Support general parallel programming
 - “any parallel algorithm on any parallel hardware”
- Make parallel programming at scale far more productive



Scalable Parallel Programming Concerns



Q: What do HPC programmers need from a language?

A: *Serial Code*: Software engineering and performance

Parallelism: What should execute simultaneously?

Locality: Where should those tasks execute?

Mapping: How to map the program to the system?

Separation of Concerns: Decouple these issues

Chapel is a language designed to address these needs from first principles



Chapel and Other Languages



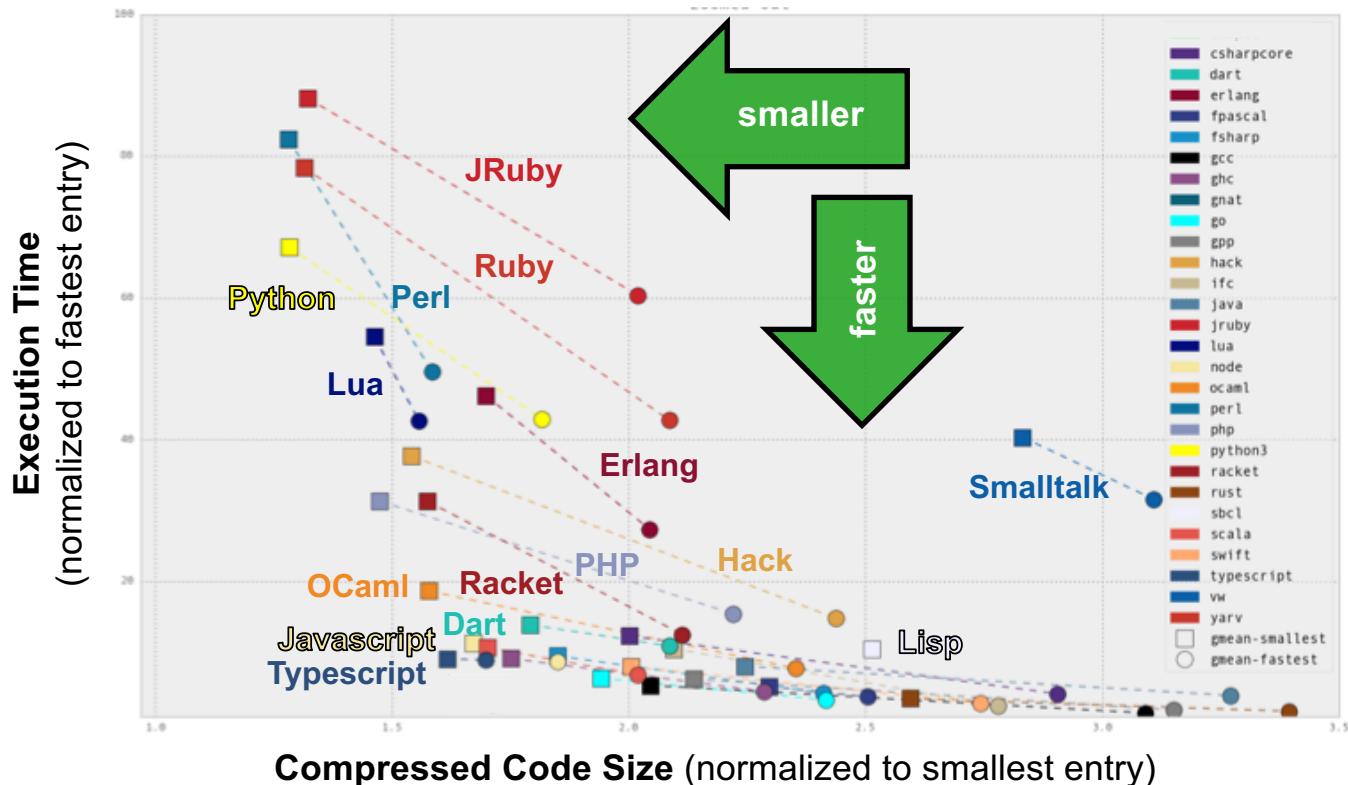
Chapel strives to be as...

- ...programmable as Python**
- ...fast as Fortran**
- ...scalable as MPI, SHMEM, or UPC**
- ...portable as C**
- ...flexible as C++**
- ...fun as [your favorite programming language]**



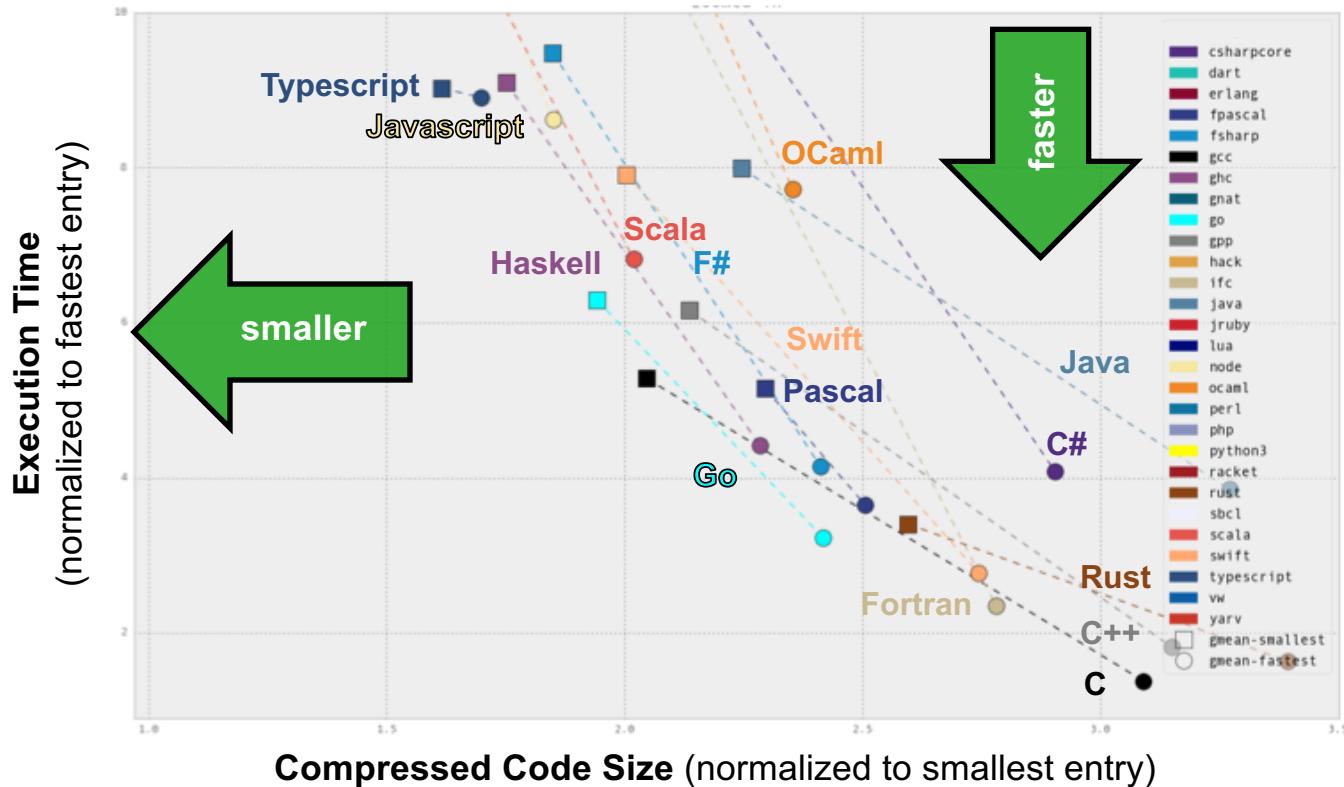
CLBG Cross-Language Summary

(Oct 2017 standings)



CLBG Cross-Language Summary

(Oct 2017 standings, zoomed in)



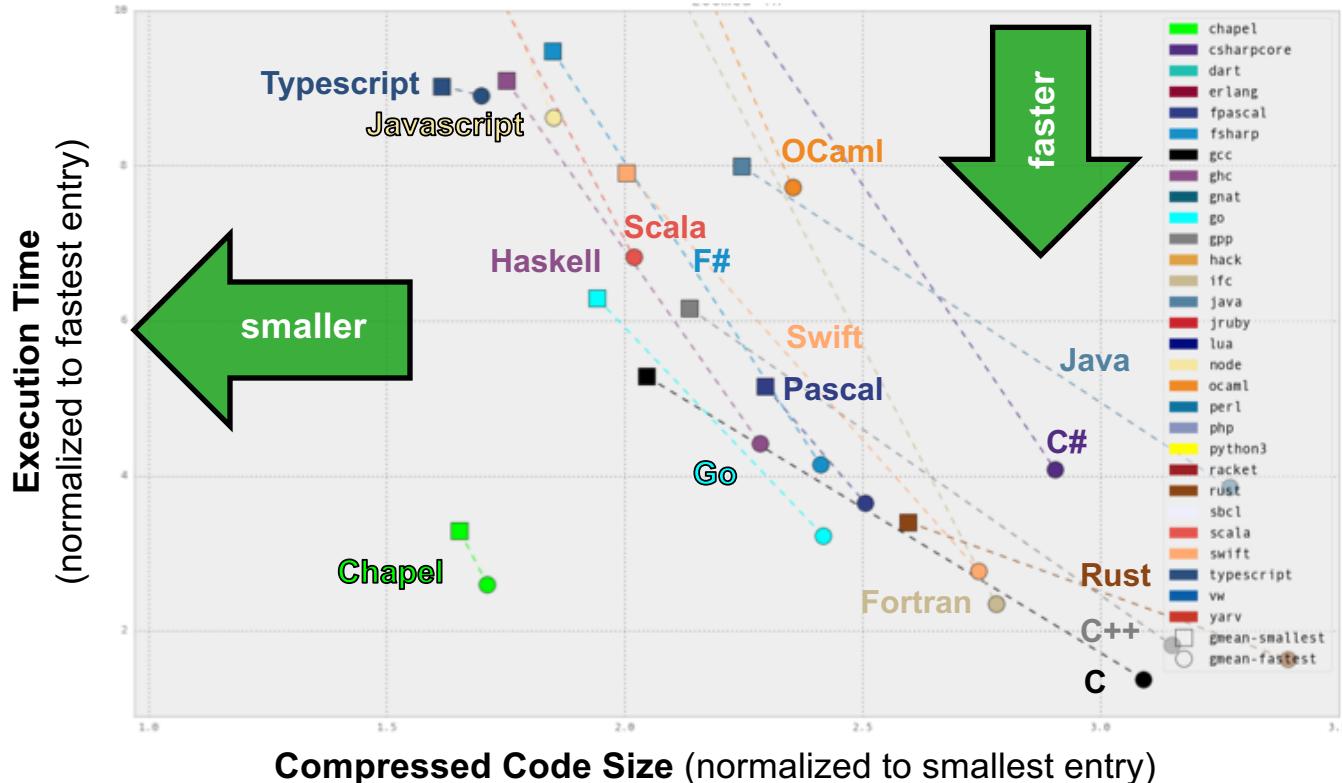
COMPUTE

STORE

ANALYZE

CLBG Cross-Language Summary

(Oct 2017 standings, zoomed in)

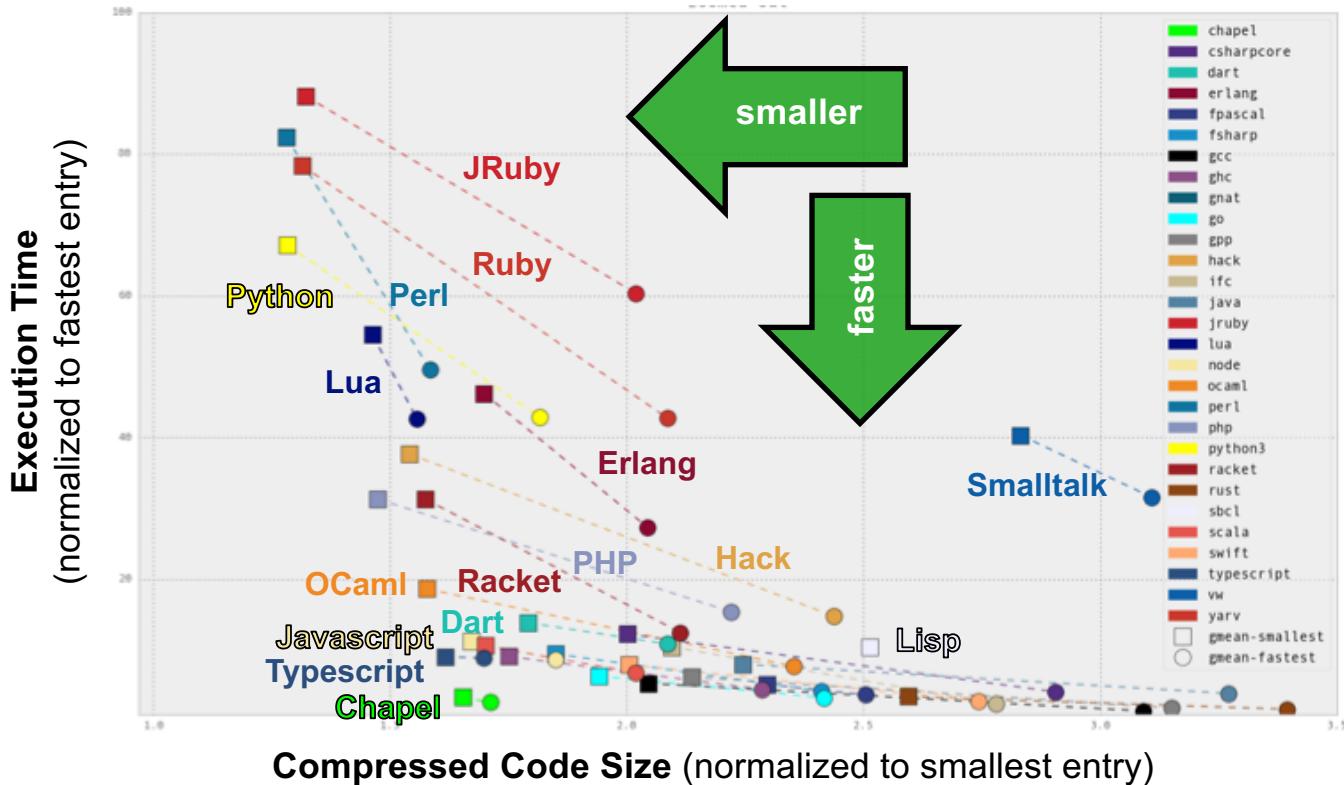


COMPUTE

STORE

ANALYZE

CLBG Cross-Language Summary (Oct 2017 standings)

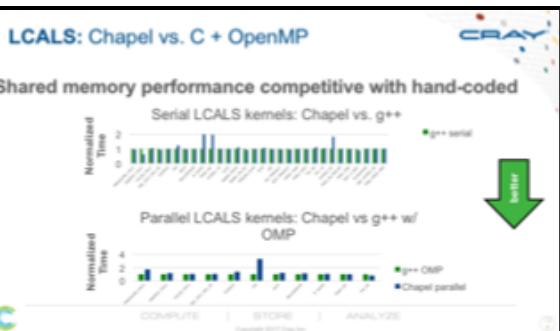
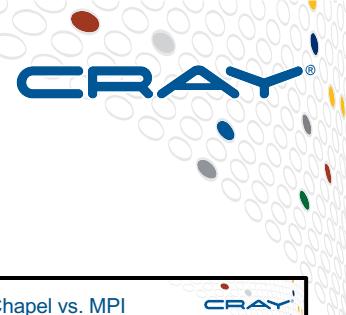


COMPUTE

STORE

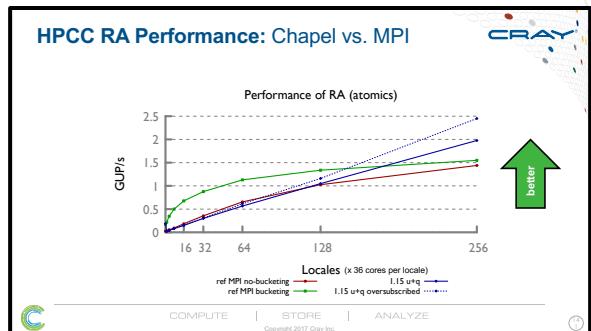
ANALYZE

Chapel Performance: HPC Benchmarks



LCALS

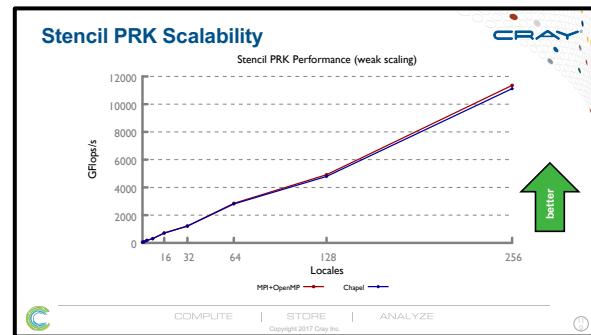
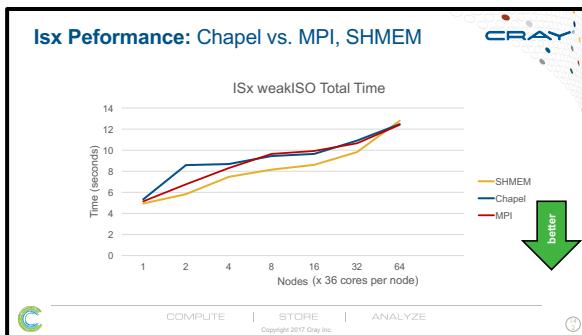
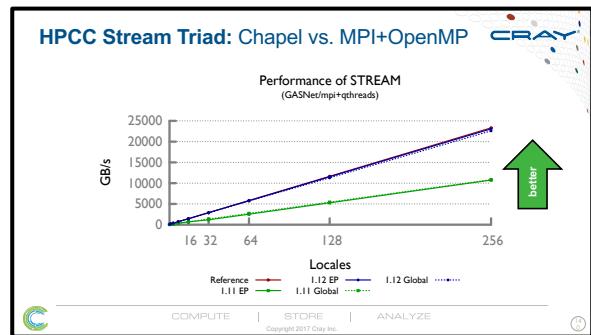
HPCC RA



STREAM
Triad

ISx

PRK
Stencil



COMPUTE

STORE

ANALYZE

Copyright 2018 Cray Inc.

Nightly performance graphs online
at: <https://chapel-lang.org/perf>

The Chapel Team at Cray (May 2017)



14 full-time employees + 2 summer interns + 2–4 GSoC students



COMPUTE

STORE

ANALYZE

Chapel Community Partners



Lawrence Berkeley
National Laboratory



Yale

(and several others...)

<https://chapel-lang.org/collaborations.html>



COMPUTE

STORE

ANALYZE



Introduction to Chapel, by Example



COMPUTE

|

STORE

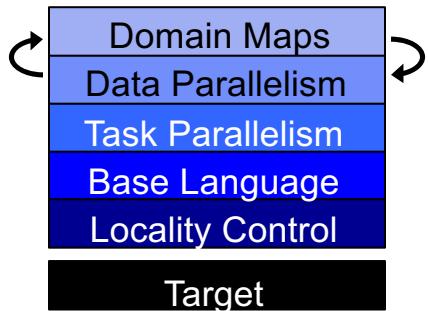
|

ANALYZE

Chapel language feature areas



Chapel language concepts



COMPUTE

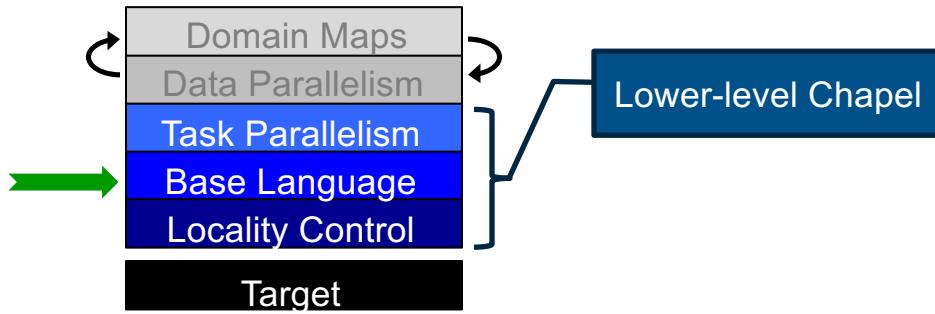
|

STORE

|

ANALYZE

Base Language



COMPUTE

STORE

ANALYZE

Base Language Features, by example



```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0  
1  
1  
2  
3  
5  
8  
...
```



Base Language Features, by example



Configuration declarations
(to avoid command-line argument parsing)
./a.out --n=1000000

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0
1
1
2
3
5
8
...
...
```



Base Language Features, by example



Modern iterators

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=gt; next;
    }
}
```

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0  
1  
1  
2  
3  
5  
8  
...
```



Base Language Features, by example

Static type inference for:

- arguments
- return types
- variables

```
iter fib(n)
  var current = 0,
      next = 1;

  for i in 1..n {
    yield current;
    current += next;
    current <=> next;
  }
}
```

```
config const n = 10;

for f in fib(n) do
  writeln(f);
```

```
0
1
1
2
3
5
8
...
...
```



Base Language Features, by example



```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

Zippered iteration

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```



Base Language Features, by example



Range types and operators

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```



COMPUTE

STORE

ANALYZE

Base Language Features, by example



```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

tuples

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```



COMPUTE

STORE

ANALYZE

Base Language Features, by example



```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

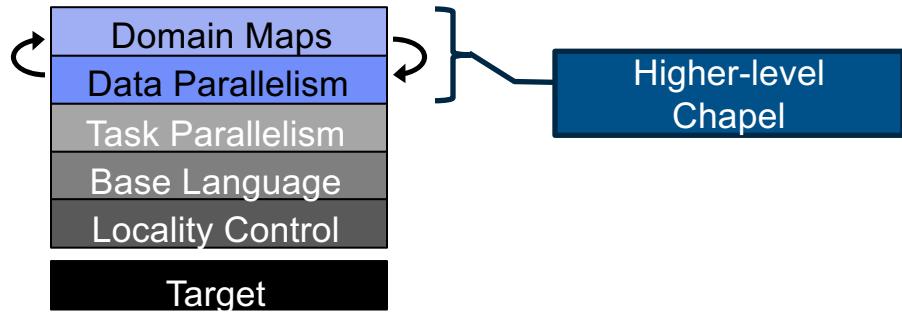
```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```



Data Parallelism in Chapel



Chapel language concepts



COMPUTE

|

STORE

|

ANALYZE

Data Parallelism, by example



dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

Data Parallelism, by example



Domains (Index Sets)

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

Data Parallelism, by example



Arrays

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

Data Parallelism, by example



Data-Parallel Forall Loops

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

Data Parallelism, by example



This is a shared memory program
Nothing has referred to remote
locales, explicitly or implicitly

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

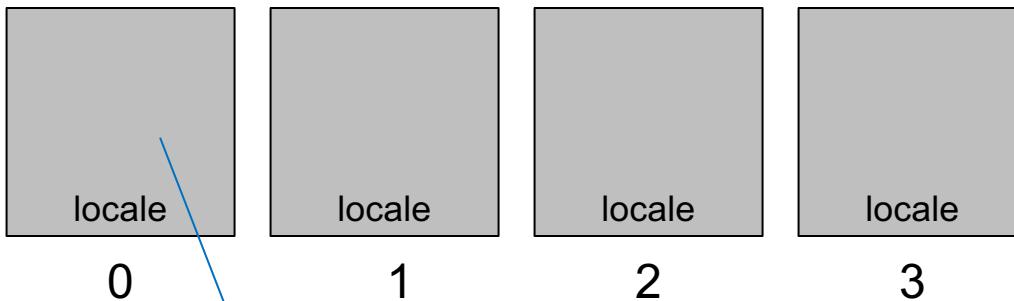
ANALYZE

Locales



- Unit of the target system useful for reasoning about locality
 - Each locale can run tasks and store variables
 - Has processors and memory (or can defer to something that does)
 - For most HPC systems, locale == compute node

Locales :



User's main() executes on locale #0



COMPUTE

STORE

ANALYZE

Data Parallelism, by example



This is a shared memory program
Nothing has referred to remote
locales, explicitly or implicitly

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

Distributed Data Parallelism, by example



Domain Maps
(Map Data Parallelism to the System)

dataParallel.chpl

```
use CyclicDist;
config const n = 1000;
var D = {1..n, 1..n}
        dmapped Cyclic(startIdx = (1,1));
var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5 --numLocales=4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

Distributed Data Parallelism, by example



magic?
HPF-like?
descriptive?

Not in the slightest...

- Lowering of code is well-defined
- User can control details
- Part of Chapel's *multiresolution philosophy*...

dataParallel.chpl

```
use CyclicDist;
config const n = 1000;
var D = {1..n, 1..n}
    dmapped Cyclic(startIdx = (1,1));
var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5 --numLocales=4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

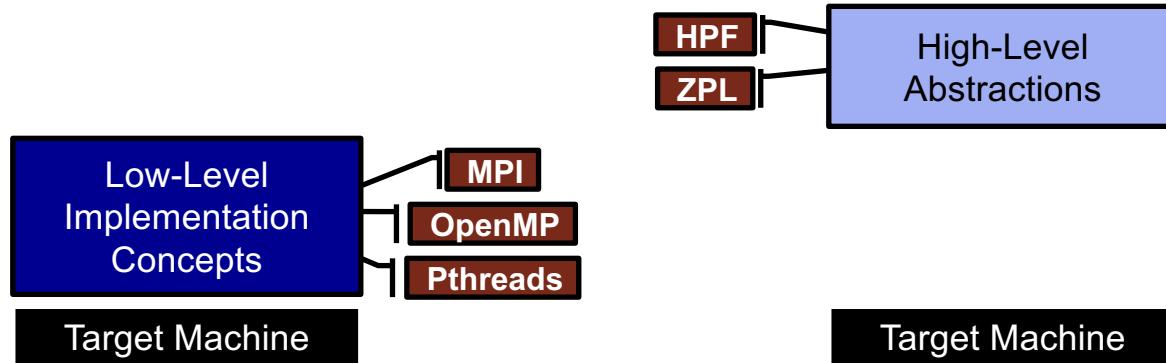


COMPUTE

STORE

ANALYZE

Chapel's Multiresolution Design: Motivation



"Why is everything so tedious/difficult?"

*"Why don't my programs trivially port
to new systems?"*

"Why don't I have more control?"

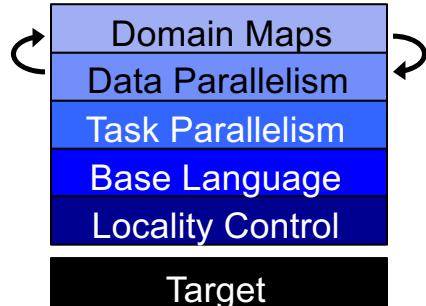


Chapel's Multiresolution Philosophy



Multiresolution Design: Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for greater degrees of control



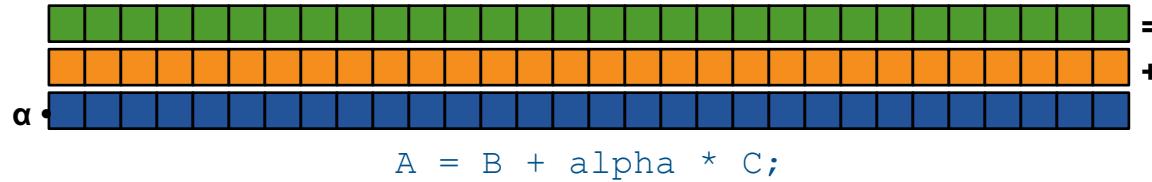
- build the higher-level concepts in terms of the lower
- permit users to intermix layers arbitrarily



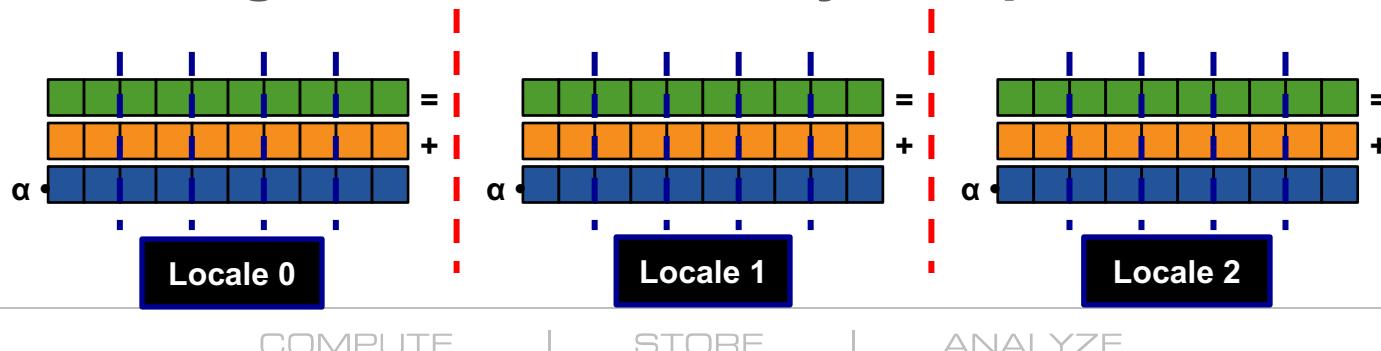
Domain Maps: A Multiresolution Feature



Domain maps are “recipes” that instruct the compiler how to map the global view of a computation...



...to the target locales' memory and processors:



Authoring Domain Maps



- **Users can write their own domain maps**
 - Implemented within Chapel itself
 - Create an object type per concept:
 - The domain map itself
 - A domain
 - An array
 - Make them satisfy a standard interface
 - e.g., arrays must support iteration, random access, etc.
 - Compiler targets this interface in implementing the language
 - Goal: make the language flexible, future-proof

Note: *all* Chapel arrays
are implemented this way



Distributed Data Parallelism, by example



dataParallel.chpl

```
use CyclicDist;
config const n = 1000;
var D = {1..n, 1..n}
        dmapped Cyclic(startIdx = (1,1));
var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5 --numLocales=4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

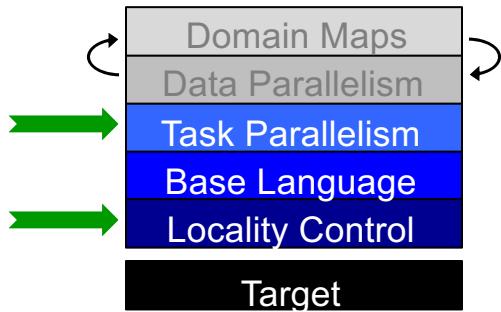


COMPUTE

STORE

ANALYZE

Task Parallelism and Locality Control



Task Parallelism and Locality, by example



taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```

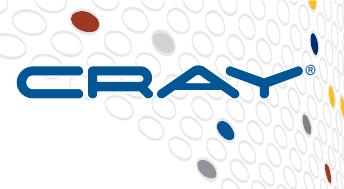


COMPUTE

STORE

ANALYZE

Task Parallelism and Locality, by example



Abstraction of
System Resources

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE

STORE

ANALYZE

Task Parallelism and Locality, by example



High-Level
Task Parallelism

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE

STORE

ANALYZE

Task Parallelism and Locality, by example



Control of Locality/Affinity

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE

STORE

ANALYZE

Task Parallelism and Locality, by example



Abstraction of
System Resources

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE

STORE

ANALYZE

Task Parallelism and Locality, by example



High-Level
Task Parallelism

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE

STORE

ANALYZE

Task Parallelism and Locality, by example



taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

Not seen here:
Data-centric task coordination
via atomic and full/empty vars

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE

STORE

ANALYZE

Task Parallelism and Locality, by example



taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.numPUs();
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE

STORE

ANALYZE

Parallelism and Locality: Distinct in Chapel



- This is a **parallel**, but local program:

```
coforall i in 1..msgs do  
    writeln("Hello from task ", i);
```

- This is a **distributed**, but serial program:

```
writeln("Hello from locale 0!");  
on Locales[1] do writeln("Hello from locale 1!");  
on Locales[2] do writeln("Hello from locale 2!");
```

- This is a **distributed parallel** program:

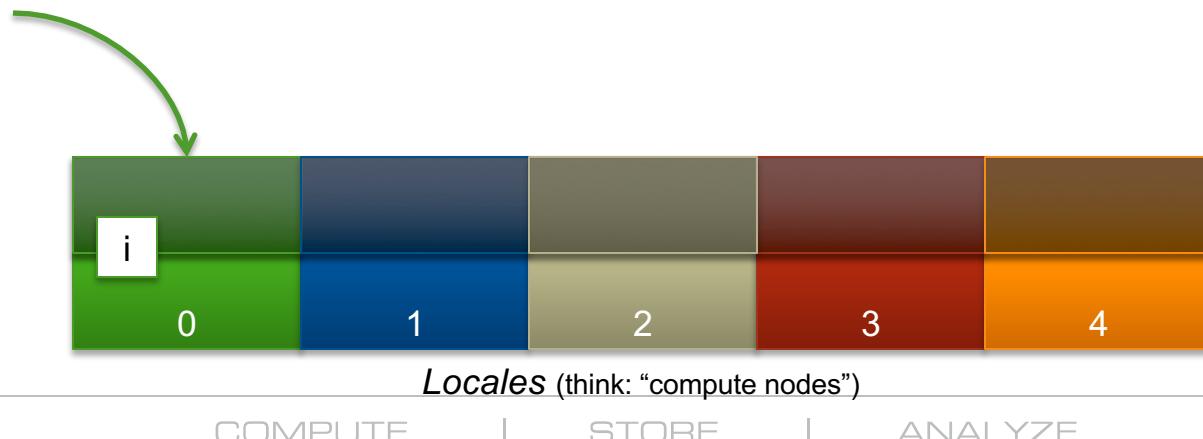
```
coforall i in 1..msgs do  
    on Locales[i%numLocales] do  
        writeln("Hello from task ", i,  
               " running on locale ", here.id);
```



Chapel: Scoping and Locality



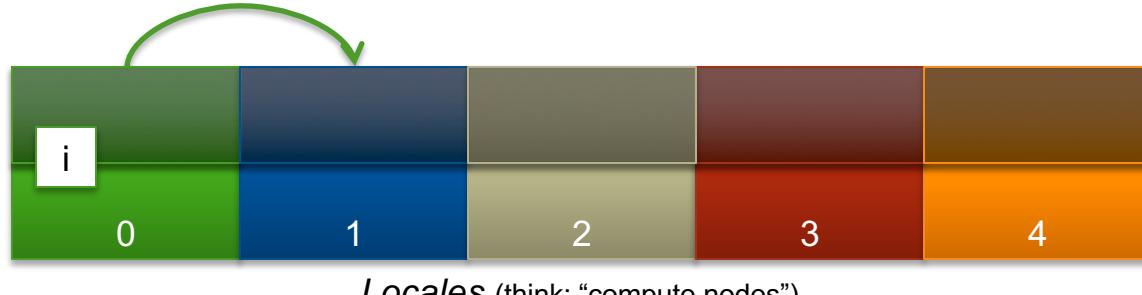
```
var i: int;
```



Chapel: Scoping and Locality



```
var i: int;  
on Locales[1] {
```



COMPUTE

STORE

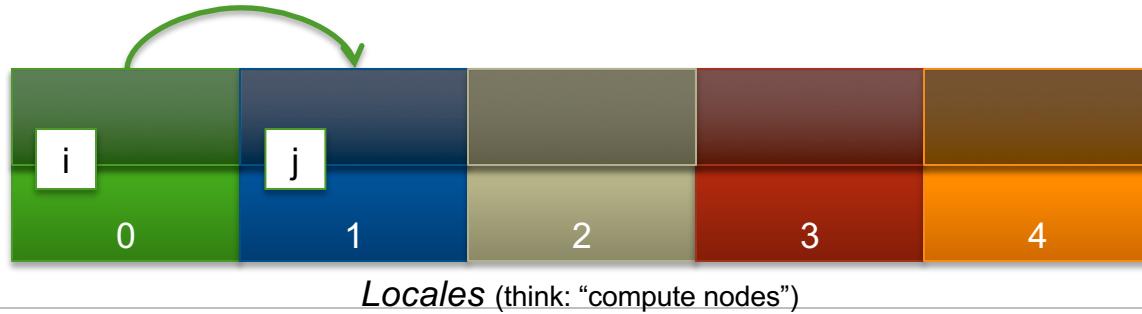
ANALYZE



Chapel: Scoping and Locality



```
var i: int;  
on Locales[1] {  
    var j: int;
```



COMPUTE

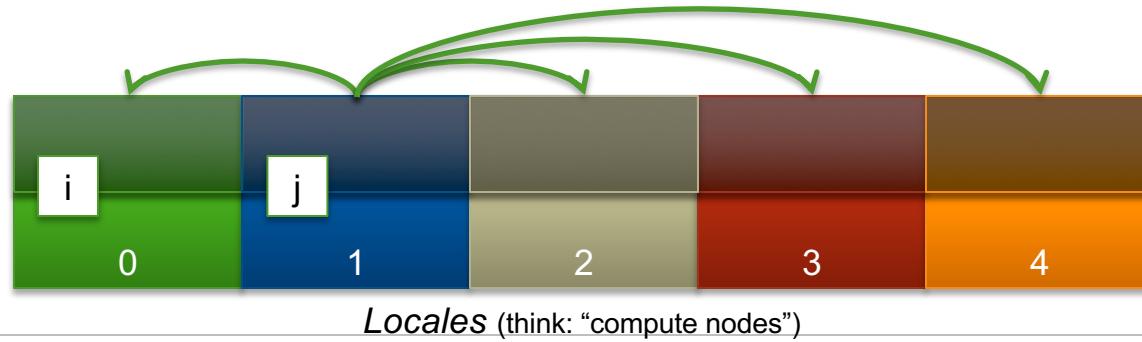
STORE

ANALYZE

Chapel: Scoping and Locality



```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {
```



COMPUTE

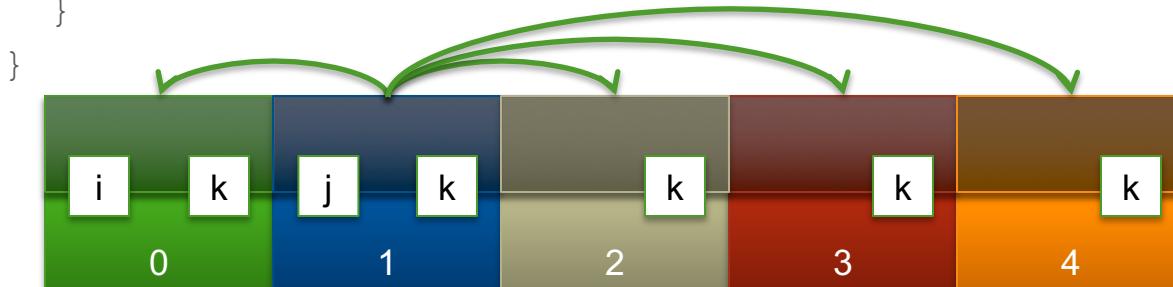
STORE

ANALYZE

Chapel: Scoping and Locality



```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
            ...  
        }  
    }  
}
```



Locales (think: “compute nodes”)

COMPUTE

STORE

ANALYZE



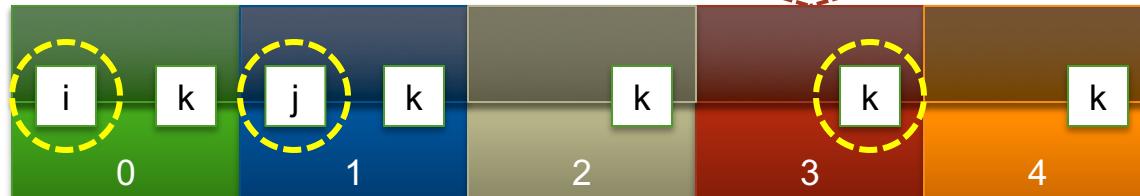
Chapel: Scoping and Locality



```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
            k = 2*i + j;  
        }  
    }  
}
```

OK to access *i*, *j*, and *k*
wherever they live

$k = 2*i + j;$



Locales (think: “compute nodes”)

COMPUTE

STORE

ANALYZE

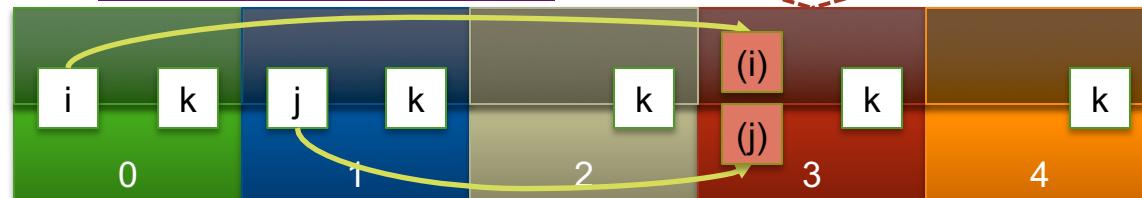


Chapel: Scoping and Locality



```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
            k = 2*i + j;  
        }  
    }  
}
```

here, *i* and *j* are remote, so
the compiler + runtime will
transfer their values



Locales (think: “compute nodes”)

COMPUTE

STORE

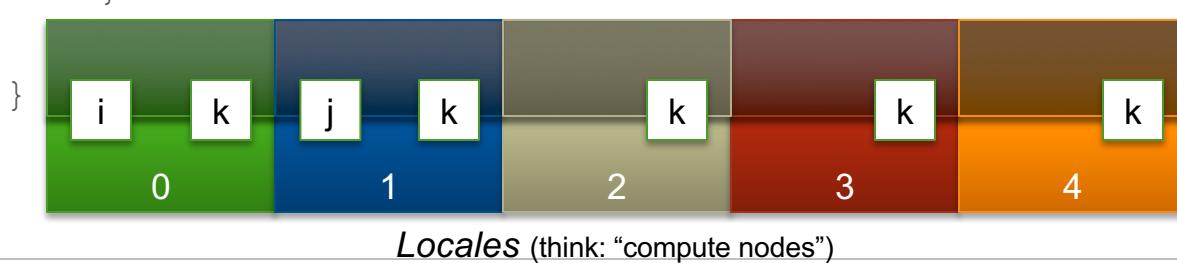
ANALYZE



Chapel: Locality queries



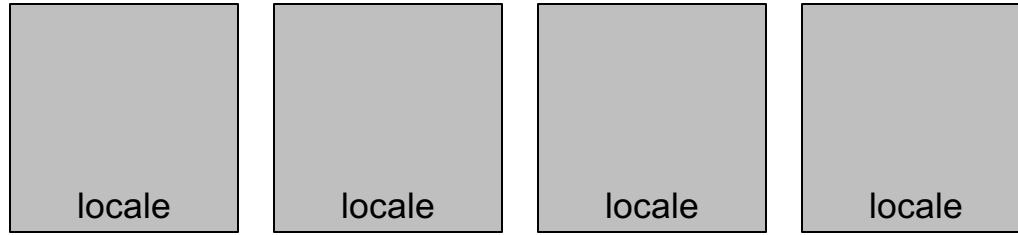
```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            ...here...           // query the locale on which this task is running  
            ...j.locale...      // query the locale on which j is stored  
            ...here.physicalMemory (...) ... // query system characteristics  
            ...here.runningTasks () ...     // query runtime characteristics  
        }  
    }  
}
```



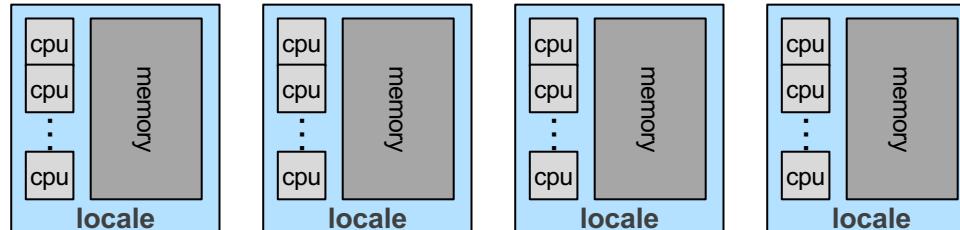
Classic Locales



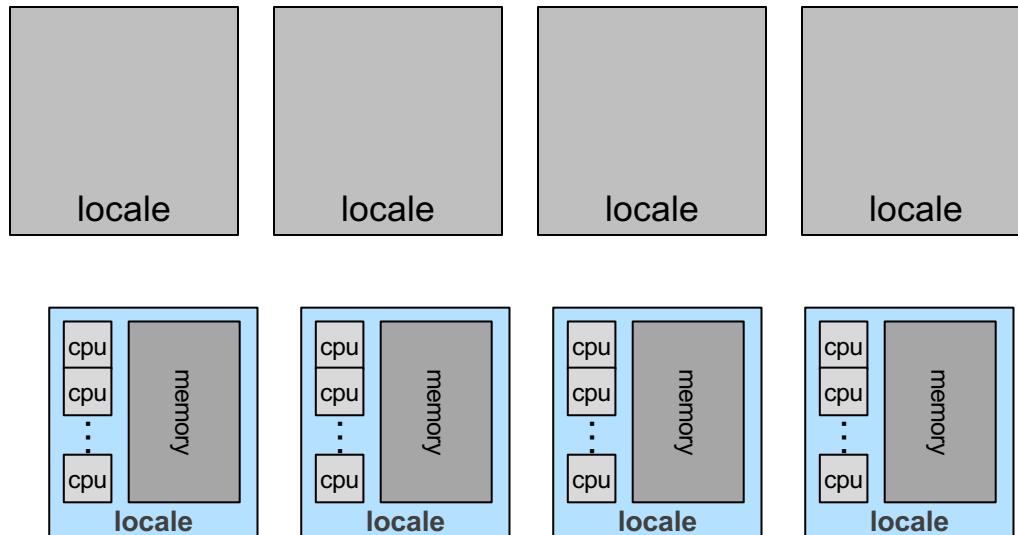
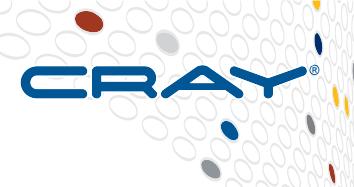
- Historically, Chapel's locales were black boxes
 - Intra-node concerns handled by compiler, runtime, OS



- This was sufficient when compute nodes were simple



Classic Locales



COMPUTE

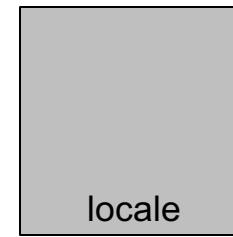
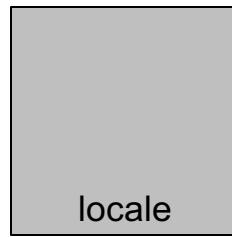
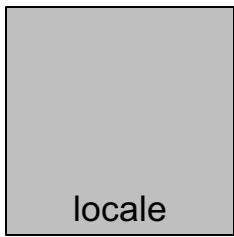
STORE

ANALYZE



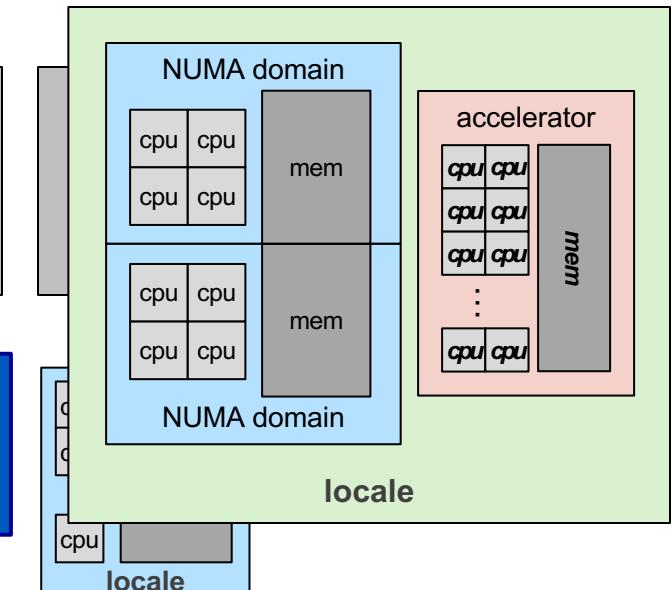
Classic Locales

- Classic model breaks down for more complex cases
 - E.g. multiple flavors of memory or processors



Could hope compilers will “simply get smart enough”

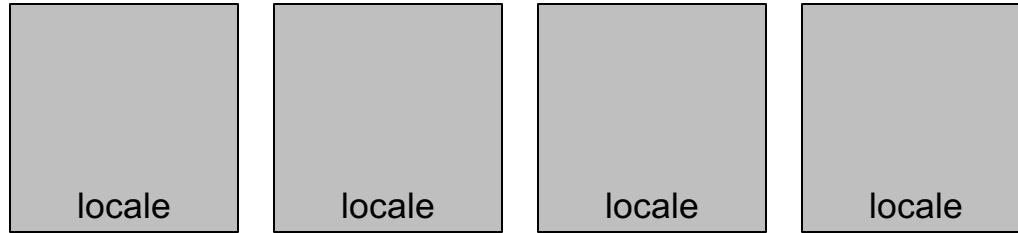
...but seems naïve and doesn't match Chapel's philosophy



Hierarchical Locales

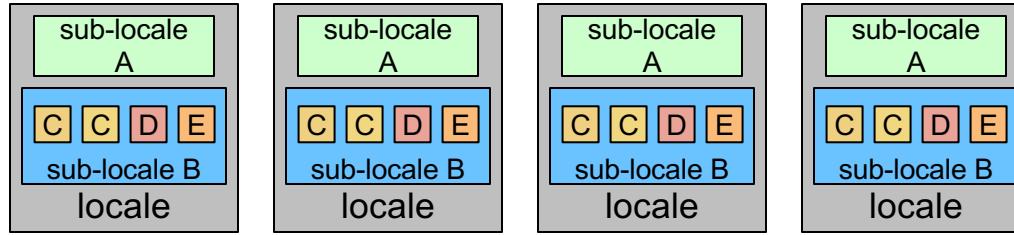


- So, we made locales hierarchical



Hierarchical Locales

- So, we made locales hierarchical
 - Locales can now themselves contain locales
 - E.g., an accelerator sub-locale, a scratchpad memory sub-locale



- Target sub-locales with on-clauses, as before
 - `on Locales[0].GPU do computationThatLikesGPUs();`
 - Ideally, hide such logic in abstractions: domain maps, parallel iterators
- Introduced a new multiresolution type: *locale models*

Chapel's Locale Models



- User-specified type representing locales
- Similar goals to domain maps:
 - Support user implementation of key high-level abstractions
 - Make language future-proof (w.r.t. emerging architectures)



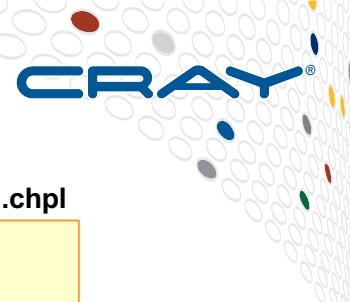
Authoring a Locale Model



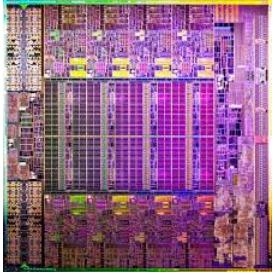
- **Creating a locale model:**
 - Create a top-level locale object type
 - In turn, it can contain fields representing sub-locales
 - Each locale / sub-locale type must meet a required interface:
 - **Memory:** How is it managed? (malloc, realloc, free)
 - **Tasking:** How do I launch and synchronize tasks?
 - **Communication:** How are data & control transferred between locales?
 - gets, puts, active messages
 - widening of pointers



An Example: The numa Locale Model

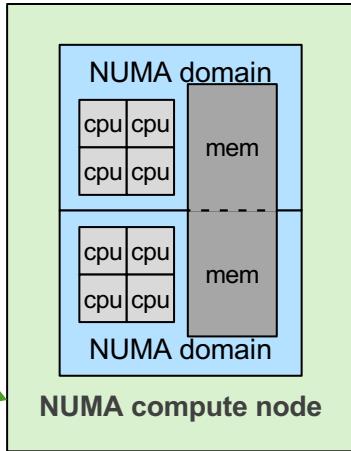


physical



<http://www1.picmag.com/media/images/337192-intel-xeon-e5-chip.jpg?thumb=1>

conceptual



\$CHPL_HOME/modules/.../numa/LocaleModel.chpl

```
class NumaDomain : AbstractLocaleModel {
    const sid: chpl_sublocID_t;
}

// The node model
class LocaleModel : AbstractLocaleModel {
    const numSublocales: int;
    var childSpace: domain(1);
    var childLocales: [childSpace] NumaDomain;
}

// support for memory management
proc chpl_here_alloc(size:int, md:int(16)) { ... }

// support for "on" statements
proc chpl_executeOn
    (loc: chpl_localeID_t, // target locale
     fn: int,             // on-body func idx
     args: c_void_ptr,   // func args
     args_size: int(32)   // args size
    ) { ... }

// support for tasking stmts: begin, cobegin, coforall
proc chpl_taskListAddCoStmt
    (subloc_id: int,           // target subloc
     fn: int,                 // body func idx
     args: c_void_ptr,        // func args
     ref tlist: _task_list,   // task list
     tlist_node_id: int       // task list owner
    ) { ... }
```



COMPUTE

STORE

ANALYZE

Copyright 2018 Cray Inc.

Locale Models: Status



- All Chapel compilations use a locale model
 - Set via environment variable or compiler flag
- Current locale models:
 - flat: the default, has no sublocales (as in the classic model)
 - numa: supports a sub-locale per NUMA domain within the node
 - knl: for Intel® Xeon Phi™: numa w/ sublocale for HBM/MCDRAM
- In practice...
 - we use the 'flat' locale model almost exclusively



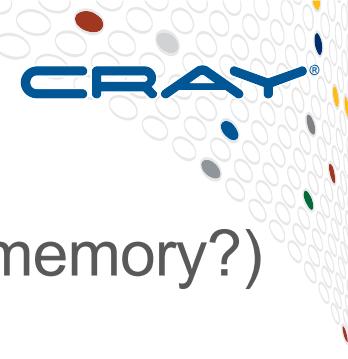
Performance: ‘numa’ vs. ‘flat’



- **Using ‘numa’ leads to performance overheads**
 - Local arrays must be “chunked” between the numa sublocales
 - Indexing must do extra work to pick the right chunk
- **Though ‘flat’ has no sub-locales, it’s also NUMA-aware**
 - First-touch heuristics used to map sub-arrays to NUMA domains
 - Yet array remains contiguous in memory ⇒ simple indexing
- **As a result, ‘numa’ rarely outperforms ‘flat’**



Challenges: Static analysis & locality



- Local vs. Remote Locales (distributed or shared memory?)
 - In general, given:

```
on x.locale ...
```

can't statically tell whether locale shares memory with `here` or not
 - May result in overheads due to conservatism
 - “Assuming it's remote, I'll introduce wide pointers & communication, ...”



Challenges: Memory-only locales



- What if a locale only represents memory? (say)
 - Have interface run tasks on other locales with processors
 - Choice of policy is up to locale model author
 - Round-robin, dynamic load-balance, nearest, ...



Wrapping Up



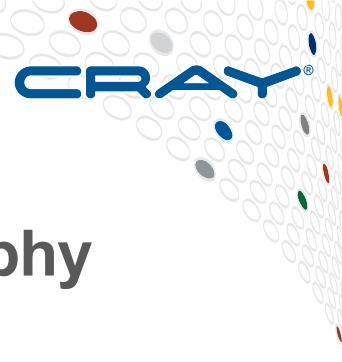
COMPUTE

|

STORE

|

ANALYZE



Summary

- **Chapel's design uses a multiresolution philosophy**
 - High-level for productivity
 - Low-level for control
 - User-extensible for flexibility, future-proof design
- **Locale models support mapping to new architectures**
 - Provide bridge from compiler to system resources
 - Enables targeting of heterogeneous resources
- **Chapel performance can match C+MPI+OpenMP**
 - With improvements in readability, writability, code size



Possible Discussion Topics



- Error-handling / Exceptions



COMPUTE

|

STORE

|

ANALYZE



Chapel Resources



COMPUTE



STORE



ANALYZE

Chapel Central: <https://chapel-lang.org/>





The Chapel Parallel Programming Language

What is Chapel?

Chapel is a modern programming language that is...

- **parallel:** contains first-class concepts for concurrent and parallel computation
- **productive:** designed with programmability and performance in mind
- **portable:** runs on laptops, clusters, the cloud, and HPC systems
- **scalable:** supports locality-oriented features for distributed memory systems
- **open-source:** hosted on [GitHub](#), permissively [licensed](#)

New to Chapel?

As an introduction to Chapel, you may want to...

- read a [blog article](#) or [book chapter](#)
- watch an [overview talk](#) or browse its [slides](#)
- [download](#) the release
- browse [sample programs](#)
- view [other resources](#) to learn how to trivially write distributed programs like this:

```
use CyclicDist;           // use the Cyclic distribution library
config const n = 100;      // use ./a.out --n=<n> to override this default
forall i in {1..n} mapped Cyclic(startIdx=1) do
    writeln("Hello from iteration ", i, " of ", n, " running on node ", here.id);
```

What's Hot?

- **Chapel 1.16** is now available—[download](#) a copy today!
- The [CHIUW 2018 call for participation](#) is now available!
- A recent [Cray blog post](#) reports on highlights from CHIUW 2017.
- Chapel is now one of the supported languages on [Try It Online!](#)
- Watch talks from [ACCU 2017](#), [CHIUW 2017](#), and [ATPESC 2016](#) on [YouTube](#).
- [Browse slides](#) from [PADAL](#), [EAGE](#), [EMBRACE](#), [ACCU](#), and other recent talks.
- See also: [What's New?](#)



COMPUTE

STORE

ANALYZE

How to Track Chapel

<http://facebook.com/ChapelLanguage>

<http://twitter.com/ChapelLanguage>

<https://www.youtube.com/channel/UCHmm27bYjhknK5mU7ZzPGsQ/chapel-announce@lists.sourceforge.net>

Chapel Programming Language (@ChapelLanguage)

We're pleased to note that Chapel is currently ranked 5th in the Computer Language Benchmarks Game's "fast-faster-fastest" graphs. That said, we're even prouder of how clear and concise the Chapel programs are relative to other entries that perform well.

http://benchmarksgame.alioth.debian.org/_which-programs-are...

How many times slower?

270 people reached

Boost Post

1 Like 18 Comments 1 Share

Russel Winder, Mykola Ratchevskyi and 2 others

Write a comment...

Vladimir Foltak measures how many programmers of given language care about that issue seriously. French times are a joke and can always be

Tweets 222 Following 12 Followers 129 Likes 32

Tweets Tweets & replies Media

Chapel Language @ChapelLanguage · 5h Doing interesting applications work in Chapel or another PGAS language? Submit it to the PAW 2017 workshop at #SC17. sourceyinstitute.github.io/PAW/

chapel.cray.com Joined March 2016

115 Photos and videos

The 2nd Annual PGAS Applications Workshop

Chapel Parallel Programming Language

Home Videos Playlists Channels About

Chapel videos

SC16 Chapel Tutorial Demo

Chapel Parallel Programming Language · 9 months ago • 397 views

This is a quick demo of the Chapel tutorial video for SC16. It shows a good way to get a quick taste of Chapel. All codes shown represent complete Chapel programs, not...

Chapel Productive, Multiresolution Parallel Programming | Brad Chamberlain, Cray, Inc.

SC16 Chapel Tutorial Demo · 9 months ago • 631 views

Presented at the Argonne Training Program on Extreme-Scale Computing, Summer 2016.

CH3W 2016 keynote: "Chapel in the (Cosmological) Wild", Nikhil Padmanabhan

Chapel Parallel Programming Language · 9 months ago • 277 views

This is Nikhil Padmanabhan's keynote talk from CH3W 2016: the 3rd Annual Chapel Implementers and Users workshop. The slides are available at...



COMPUTE

STORE

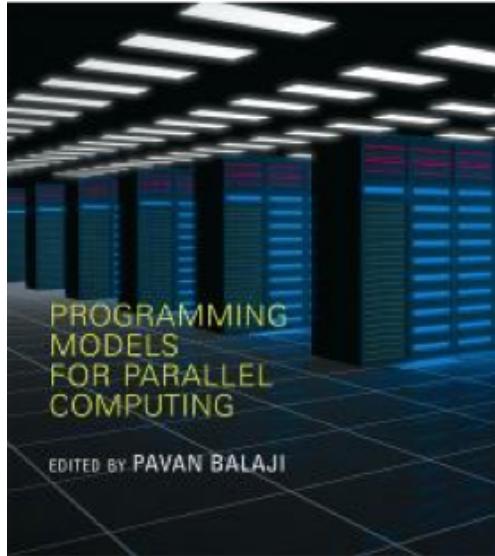
ANALYZE

Suggested Reading (healthy attention spans)



Chapel chapter from [Programming Models for Parallel Computing](#)

- a detailed overview of Chapel's history, motivating themes, features
- published by MIT Press, November 2015
- edited by Pavan Balaji (Argonne)
- chapter is now also available [online](#)



Other Chapel papers/publications available at <https://chapel-lang.org/papers.html>



COMPUTE

STORE

ANALYZE

Suggested Reading (short attention spans)



[CHIUW 2017: Surveying the Chapel Landscape](#), Cray Blog, July 2017.

- *a run-down of recent events*

[Chapel: Productive Parallel Programming](#), Cray Blog, May 2013.

- *a short-and-sweet introduction to Chapel*

[Six Ways to Say “Hello” in Chapel](#) (parts [1](#), [2](#), [3](#)), Cray Blog, Sep-Oct 2015.

- *a series of articles illustrating the basics of parallelism and locality in Chapel*

[Why Chapel?](#) (parts [1](#), [2](#), [3](#)), Cray Blog, Jun-Oct 2014.

- *a series of articles answering common questions about why we are pursuing Chapel in spite of the inherent challenges*

[\[Ten\] Myths About Scalable Programming Languages](#), [IEEE TCSC Blog](#)

(index available on chapel-lang.org “blog posts” page), Apr-Nov 2012.

- *a series of technical opinion pieces designed to argue against standard reasons given for not developing high-level parallel languages*



Chapel StackOverflow and GitHub Issues



A screenshot showing two side-by-side interfaces for the Chapel programming language. On the left is the StackOverflow Q&A site, displaying four questions related to Chapel. On the right is the GitHub repository page for 'chapel-lang/chapel', showing a list of 292 open issues. Both pages include search, filter, and sorting functionality.



COMPUTE

STORE

ANALYZE

Copyright 2018 Cray Inc.

Where to..



Submit bug reports:

[GitHub issues for chapel-lang/chapel](#): public bug forum
chapel_bugs@cray.com: for reporting non-public bugs

Ask User-Oriented Questions:

[StackOverflow](#): when appropriate / other users might care
[#chapel-users \(irc.freenode.net\)](#): user-oriented IRC channel
chapel-users@lists.sourceforge.net: user discussions

Discuss Chapel development

chapel-developers@lists.sourceforge.net: developer discussions
[#chapel-developers \(irc.freenode.net\)](#): developer-oriented IRC channel

Discuss Chapel's use in education

chapel-education@lists.sourceforge.net: educator discussions

Directly contact Chapel team at Cray: chapel_info@cray.com



Questions?



COMPUTE

|

STORE

|

ANALYZE

Legal Disclaimer



Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.

Cray Inc. may make changes to specifications and product descriptions at any time, without notice.

All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.

Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.

Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.

The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, and URIKA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.





CRAY
THE SUPERCOMPUTER COMPANY