

Chapel: A Domain-Specific Language for Productive Parallel Programming at Scale

Brad Chamberlain, Chapel Team, Cray Inc.
University of Washington, CSEP 590c
May 9th, 2016



Chapel: A Domain-Specific Language for Productive Parallel Programming at Scale

Brad Chamberlain, Chapel Team, Cray Inc.
University of Washington, CSEP 590c
May 9th, 2016





Chapel: An Extremely General Language for Productive Programming on Modern Systems

Brad Chamberlain, Chapel Team, Cray Inc.
University of Washington, CSEP 590c
May 9th, 2016



Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.

My Background

Affiliations:

- Graduated from UW in 2001
(and remain associated with UW as an affiliate professor)
- Spent a lost but educational year at a start-up
- Have worked at Cray Inc. since Fall 2002

R&D topics:

- Focus primarily on parallel language design & implementation
 - **at UW:** on the ZPL data-parallel language
 - **at start-up:** on a language for reconfigurable embedded computing
 - **at Cray:** on Chapel
- Most of my work has been in High-Performance Computing (HPC)
(but, mainstream computing is becoming more and more like HPC)
 - parallel processors, GPU computing, cloud computing, data analytics, ...
 - HPC challenges are becoming everyone's challenges: **parallelism & locality**

The Chapel Team at Cray (Spring 2015)



(Our team currently has an open language/compiler developer position)

Motivation for Chapel

Q: Why doesn't HPC programming have an equivalent to Python / Matlab / Java / C++ / (your favorite programming language here) ?

- one that makes it easy to get programs up and running quickly
- one that is portable across system architectures and scales
- one that bridges the HPC, data analysis, and mainstream communities

A: We believe this is due not to any particular technical challenge, but rather a lack of sufficient...

- ...long-term efforts
- ...resources
- ...community will
- ...co-design between developers and users
- ...patience

Chapel is our attempt to reverse this trend!

HPC's Status Quo: SPMD Programming

SPMD = Single Program, Multiple Data

- concept: write one program, run multiple copies of it in parallel
- arguably a “bottom-up” programming model design
 - “HPC systems can run lots of programs, so let’s get parallelism that way”
- often clumsy in practice


SPMD by Example

In pictures: “Apply a 3-Point Stencil to a vector”

Conceptual View

$$\begin{aligned} & (\text{brown bar}) \\ & + \text{ (red bar) } / 2 \\ & = \text{ (yellow bar) } \end{aligned}$$

SPMD View




SPMD by Example

In pictures: “Apply a 3-Point Stencil to a vector”

Conceptual View

$$\begin{aligned} & (\text{brown bar}) \\ & + (\text{pink bar}) / 2 \\ & = \text{yellow bar} \end{aligned}$$

SPMD View



4) Global-View Abstractions

In code: “Apply a 3-Point Stencil to a vector”




SPMD pseudo-code

```

proc main() {
    var n = 1000;
    var p = numProcs(),
        me = myProc(),
        myN = n/p,
    var A, B: [0..myN+1] real;

    if (me < p-1) {
        send(me+1, A[myN]);
        recv(me+1, A[myN+1]);
    }
    if (me > 0) {
        send(me-1, A[1]);
        recv(me-1, A[0]);
    }
    forall i in 1..myN do
        B[i] = (A[i-1] + A[i+1])/2;
}

```



Bug: Refers to uninitialized values at ends of A


COMPUTE

STORE

ANALYZE

4) Global-View Abstractions

In code: “Apply a 3-Point Stencil to a vector”



What is Chapel?

Chapel: An emerging parallel programming language

- extensible
- portable
- open-source
- a collaborative effort
- a work-in-progress
- a top-down language design

Goals:

- Support general parallel programming
 - “any parallel algorithm on any parallel hardware”
- Make parallel programming far more productive

What does “Productivity” mean to you?

Recent Graduates:

“something similar to what I used in school: Python, Matlab, Java, ...”

Seasoned HPC Programmers:

“that sugary stuff that I don’t need because I ~~was born to suffer~~
want full control
to ensure performance”

Computational Scientists:

“something that lets me express my parallel computations
without having to wrestle with architecture-specific details”

Chapel Team:


“something that lets computational scientists express what they want,
without taking away the control that HPC programmers want,
implemented in a language as attractive as recent graduates want.”

STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures:




STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel:




STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory):




STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory multicore):



STREAM Triad: MPI

MPI

```
#include <hpcc.h>

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );


    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
        0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
        sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
}
```



```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d).
        \n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 0.0;
}
scalar = 3.0;

for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
```

STREAM Triad: MPI+OpenMP

MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
                0, comm );


    return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
                                       sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
}

```



```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d).
        \n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

#ifndef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 0.0;
}

scalar = 3.0;

#ifndef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
}

```

STREAM Triad: MPI+OpenMP vs. CUDA

MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );

    return errCount;
}

int HPCC_Triad(HPCC_Params *params, FILE *outFile)
{
    int i, j, k;
    double scalar;
    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );
    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }

    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }

    scalar = 3.0;

    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
    return 0;
}
```

CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid

    set_array<<<dimGrid, dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid, dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid, dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

    __global__ void set_array(float *a, float value, int len) {
        int idx = threadIdx.x + blockIdx.x * blockDim.x;
        if (idx < len) a[idx] = value;
    }

    __global__ void STREAM_Triad( float *a, float *b, float *c,
                                float scalar, int len) {
        int idx = threadIdx.x + blockIdx.x * blockDim.x;
        if (idx < len) c[idx] = a[idx]+scalar*b[idx];
    }
}
```

COMPUTE

STORE

ANALYZE

Why so many programming models?

HPC tends to approach programming models bottom-up:

Given a system and its core capabilities...

...provide features that can access the available performance.

- portability, generality, programmability: not strictly necessary.

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	MPI	executable
Intra-node/multicore	OpenMP / pthreads	iteration/task
Instruction-level vectors/threads	pragmas	iteration
GPU/accelerator	CUDA / Open[CL MP ACC]	SIMD function/task

benefits: lots of control; decent generality; easy to implement
downsides: lots of user-managed detail; brittle to changes

Rewinding a few slides...

MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );

    return errCount;
}

int HPCC_LocalVectorSize(HPCC_Params *params, int len, double scalar);

VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );

a = HPCC_XMALLOC( double, VectorSize );
b = HPCC_XMALLOC( double, VectorSize );
c = HPCC_XMALLOC( double, VectorSize );

if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

#ifndef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 0.0;
}
scalar = 3.0;

#ifdef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
return 0;
}
```

CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid

    set_array<<<dimGrid, dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid, dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid, dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad( float *a, float *b, float *c,
                             float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```

COMPUTE

STORE

ANALYZE

STREAM Triad: Chapel

MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include<omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params,
int myRank, commSize;
int rv, errCount;
MPI_Comm comm = MPI_COMM_WORLD;

MPI_Comm_size( comm, &commSize );
MPI_Comm_rank( comm, &myRank );

rv = HPCC_Stream( params, 0 == myR
MPI_Reduce( &rv, &errCount, 1, MPI_
return errCount;

int HPCC_Stream(HPCC_Params *params,
register int j;
double scalar;
VectorSize = HPCC_LocalVectorSize();
a = HPCC_XMALLOC( double, VectorSi
b = HPCC_XMALLOC( double, VectorSi
c = HPCC_XMALLOC( double, VectorSi
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d)\n", VectorSize );
        fclose( outFile );
    }
}
```

Chapel

```
config const m = 1000,
alpha = 3.0;


const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

A = B + alpha * C;
```

the special sauce



Philosophy: Good, *top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

COMPUTE

STORE

ANALYZE

Outline

✓ Motivation

➤ Chapel's Design Themes

- Survey of Chapel Concepts

- Project Status and Next Steps

Design Themes for Chapel

- 1) General Parallel Programming**
- 2) Reduce HPC ↔ Mainstream Language Gap**
- 3) Global-View Abstractions**
- 4) Multiresolution Design**
- 5) Control over Locality/Affinity**

Design Themes for Chapel

- 1) General Parallel Programming
- 2) Reduce HPC ↔ Mainstream Language Gap
- 3) Global-View Abstractions
- 4) Multiresolution Design
- 5) Control over Locality/Affinity } We'll cover this one as we go

1) General Parallel Programming

With a unified set of concepts...

...target any hardware parallelism available in the system

- **Types:** machines, nodes, accelerators, cores, instructions

...express any software parallelism desired by the user

- **Styles:** data-parallel, task-parallel, concurrency, nested, ...
- **Levels:** model, function, loop, statement, expression

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	Chapel	task (or executable)
Intra-node/multicore	Chapel	iteration/task
Instruction-level vectors/threads	Chapel	iteration
GPU/accelerator	Chapel	SIMD function/task

2) Reduce HPC ↔ Mainstream Language Gap



Consider:

- Students graduate with training in Java, Matlab, Python, etc.
- Yet HPC programming is dominated by Fortran, C/C++, MPI, ...

We'd like to narrow this gulf with Chapel:

- to leverage advances in modern language design
- to better utilize the skills of the entry-level workforce...
...while not alienating the traditional HPC programmer
 - e.g., support object-oriented programming, but make it optional


3) Global-View Abstractions

In pictures: “Apply a 3-Point Stencil to a vector”

Conceptual View


$$\begin{aligned}
 & (\text{brown bar}) \\
 & + (\text{pink bar}) / 2 \\
 & = \text{yellow bar}
 \end{aligned}$$

SPMD View



3) Global-View Abstractions

In pictures: “Apply a 3-Point Stencil to a vector”



3) Global-View Abstractions

In code: “Apply a 3-Point Stencil to a vector”

Global-View

```
proc main() {
    var n = 1000;
    var A, B: [1..n] real;

    forall i in 2..n-1 do
        B[i] = (A[i-1] + A[i+1])/2;
}
```



Local-View (SPMD)

```
proc main() {
    var n = 1000;
    var p = numProcs(),
        me = myProc(),
        myN = n/p,
        myLo = 1,
        myHi = myN;
    var A, B: [0..myN+1] real;

    if (me < p-1) {
        send(me+1, A[myN]);
        recv(me+1, A[myN+1]);
    } else
        myHi = myN-1;
    if (me > 0) {
        send(me-1, A[1]);
        recv(me-1, A[0]);
    } else
        myLo = 2;
    forall i in myLo..myHi do
        B[i] = (A[i-1] + A[i+1])/2;
}
```



COMPUTE

STENCIL


ANALYZE

3) Global-View Abstractions: A Final Note

- A language may support both global- and local-view programming — in particular, Chapel does

```
proc main() {  
    coforall loc in Locales do  
        on loc do  
            MySPMDProgram(loc.id, Locales.numElements);  
  
}  
  
proc MySPMDProgram(myImageID, numImages) {  
    ...  
}
```

4) Multiresolution Design: Motivation



“Why is everything so tedious/difficult?”

*“Why don’t my programs trivially port
to new systems?”*


“Why don’t I have more control?”

4) Multiresolution Design

Multiresolution Design: Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for greater degrees of control


Chapel language concepts



- build the higher-level concepts in terms of the lower
- permit the user to intermix layers arbitrarily

Outline


- ✓ Motivation
- ✓ Chapel's Design Themes
- Survey of Chapel Concepts



● Project Status and Next Steps

Outline

- ✓ Motivation
- ✓ Chapel's Design Themes
- Survey of Chapel Concepts



● Project Status and Next Steps

Static Type Inference

```
const pi = 3.14,                      // pi is a real
      coord = 1.2 + 3.4i,              // coord is a complex...
      coord2 = pi*coord,              // ...as is coord2
      name = "brad",                  // name is a string
      verbose = false;                // verbose is boolean

proc addem(x, y) {                    // addem() has generic arguments
    return x + y;                     // and an inferred return type
}

var sum = addem(1, pi),               // sum is a real
    fullname = addem(name, "ford");   // fullname is a string

writeln((sum, fullname));
```

(4.14, bradford)

Range Types, Values, and Operators

```

const r = 1..10;

printVals(r);
printVals(r # 3);
printVals(r by 2);
printVals(r by -2);
printVals(r by 2 # 3);
printVals(r # 3 by 2);
printVals(0.. #n);

proc printVals(r) {
    for i in r do
        write(i, " ");
    writeln();
}

```

```

1 2 3 4 5 6 7 8 9 10
1 2 3
1 3 5 7 9
10 8 6 4 2
1 3 5
1 3
0 1 2 3 4 ... n-1

```

Iterators

```
iter fibonacci(n) {
  var current = 0,
       next = 1;
  for 1..n {
    yield current;
    current += next;
    current <=> next;
  }
}
```

```
for f in fibonacci(7) do
  writeln(f);
```

```
0
1
1
2
3
5
8
```

```
iter tiledRMO(D, tilesize) {
  const tile = {0..#tilesize,
                0..#tilesize};
  for base in D by tilesize do
    for ij in D[tile + base] do
      yield ij;
}
```

```
for ij in tiledRMO({1..m, 1..n}, 2) do
  write(ij);
```

```
(1,1) (1,2) (2,1) (2,2)
(1,3) (1,4) (2,3) (2,4)
(1,5) (1,6) (2,5) (2,6)
...
(3,1) (3,2) (4,1) (4,2)
```

Zippered Iteration

```
for (i,f) in zip(0..#n, fibonacci(n)) do  
    writeln("fib #", i, " is ", f);
```


```
fib #0 is 0  
fib #1 is 1  
fib #2 is 1  
fib #3 is 2  
fib #4 is 3  
fib #5 is 5  
fib #6 is 8  
...
```

Other Base Language Features

- tuple types and values
- interoperability features
- OOP (value- and reference-based)
- modules (for namespace management)
- rank-independent programming features
- compile-time features for meta-programming
 - e.g., compile-time functions to compute types, parameters
- argument intents, default values, match-by-name
- overloading, where clauses
- ...

Outline

- ✓ Motivation
- ✓ Chapel's Design Themes
- Survey of Chapel Concepts



- Project Status and Next Steps

Task Parallelism: Begin Statements

```
// create a fire-and-forget task for a statement
begin writeln("hello world");
writeln("goodbye");
```

Possible outputs:

hello world
goodbye

goodbye
hello world

Task Parallelism: Coforall Loops

```
// create a task per iteration
coforall t in 0..#numTasks {
    writeln("Hello from task ", t, " of ", numTasks);
} // implicit join of the numTasks tasks here

writeln("All tasks done");
```

Sample output:


```
Hello from task 2 of 4
Hello from task 0 of 4
Hello from task 3 of 4
Hello from task 1 of 4
All tasks done
```

Other Task Parallel Concepts

- **cobegin**s: create tasks using compound statements
- **atomic variables**: support atomics ops, similar to modern C++
- **sync/single variables**: support producer/consumer patterns
- **sync statements**: join unstructured tasks
- **serial statements**: conditionally squash parallelism

Outline

- ✓ Motivation
- ✓ Chapel's Design Themes
- Survey of Chapel Concepts



- Project Status and Next Steps

The Locale Type

Definition:

- Abstract unit of target architecture
- Supports reasoning about locality
 - defines “here vs. there” / “local vs. remote”
- Capable of running tasks and storing variables
 - i.e., has processors and memory

Typically: A compute node (multicore processor or SMP)

Getting started with locales

- Specify # of locales when running Chapel programs

```
% a.out --numLocales=8
```

```
% a.out -nl 8
```

- Chapel provides built-in locale variables

```
config const numLocales: int = ...;  
const Locales: [0..#numLocales] locale = ...;
```

Locales



- User's main() begins executing on locale #0

Locale Operations

- Locale methods support queries about the target system:

```
proc locale.physicalMemory(...) { ... }
proc locale.numCores { ... }
proc locale.id { ... }
proc locale.name { ... }
```

- On-clauses support placement of computations:

```
writeln("on locale 0");
on Locales[1] do
    writeln("now on locale 1");
writeln("on locale 0 again");
```

```
on A[i,j] do
    bigComputation(A);
on node.left do
    search(node.left);
```

Parallelism and Locality: Orthogonal in Chapel

- This is a **parallel**, but local program:

```
begin writeln("Hello world!");  
writeln("Goodbye!");
```

- This is a **distributed**, but serial program:

```
writeln("Hello from locale 0!");  
on Locales[1] do writeln("Hello from locale 1!");  
writeln("Goodbye from locale 0!");
```

- This is a **distributed and parallel** program:

```
begin on Locales[1] do writeln("Hello from locale 1!");  
on Locales[2] do begin writeln("Hello from locale 2!");  
writeln("Goodbye from locale 0!");
```

Time Check!


Do we have time for a sidebar on PGAS?

Yes? No!

PGAS Programming in a Nutshell

Global Address Space:

- permit parallel tasks to access variables by naming them
 - regardless of whether they are local or remote
 - compiler / library / runtime will take care of communication



Images / Threads / Locales / Places / etc. (think: “compute nodes”)

COMPUTE | STORE | ANALYZE


PGAS Programming in a Nutshell

Global Address Space:

- permit parallel tasks to access variables by naming them
 - regardless of whether they are local or remote
 - compiler / library / runtime will take care of communication

Partitioned:

- establish a strong model for reasoning about locality
 - every variable has a well-defined location in the system
 - local variables are typically cheaper to access than remote ones



Images / Threads / Locales / Places / etc. (think: “compute nodes”)


PGAS Programming in a Nutshell

Global Address Space:

- permit parallel tasks to access variables by naming them
 - regardless of whether they are local or remote
 - compiler / library / runtime will take care of communication

Partitioned:

- establish a strong model for reasoning about locality
 - every variable has a well-defined location in the system
 - local variables are typically cheaper to access than remote ones




Images / Threads / Locales / Places / etc. (think: “compute nodes”)

C O M P U T E | S T O R E | A N A L Y Z E


SPMD PGAS Languages (using a pseudo-language, not Chapel)

```
proc main() {  
    var i (*): int;           // declare a shared variable i
```



SPMD PGAS Languages (using a pseudo-language, not Chapel)

```
proc main() {  
    var i(*) : int;          // declare a shared variable i  
    i = 2*this_image();     // each image initializes its copy
```



SPMD PGAS Languages

(using a pseudo-language, not Chapel)

```
proc main() {  
    var i(*) : int;           // declare a shared variable i  
    i = 2*this_image();       // each image initializes its copy  
    var j: int;               // declare a private variable j
```



COMPUTE

STORE

ANALYZE


SPMD PGAS Languages (using a pseudo-language, not Chapel)

```

proc main() {
    var i(*) : int;           // declare a shared variable i
    i = 2*this_image();       // each image initializes its copy
    var j: int;               // declare a private variable j
    j = i( (this_image() + 1) % num_images() );
    // ^ access our neighbor's copy of i
    // communication is implemented by the compiler + runtime


    // How did we know our neighbor had an i?
    // Because it's SPMD – we're all running the same program.
}

```



Chapel and PGAS


- Chapel is PGAS, but unlike most, it's not inherently SPMD
 - ⇒ never think about “the other copies of the program”
 - ⇒ “global name/address space” comes from lexical scoping
 - as in traditional languages, each declaration yields one variable
 - variables are stored on the locale where the task declaring it is executing



COMPUTE | STORE | ANALYZE


Chapel: Scoping and Locality

```
var i: int;
```



Chapel: Scoping and Locality


```
var i: int;  
on Locales[1] {
```



COMPUTE | STORE | ANALYZE


Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;
```




Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {
```



Chapel: Scoping and Locality


```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
            // within this scope, i, j, and k can be referenced. For example:  
            k = 2*i + j;  
            // The implementation manages any communication.  
        }  
    }  
}
```



Chapel: Locality queries


```
var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales {
        on loc {
            var k: int;
```

...here... // query the locale on which this task is running
...j.locale... // query the locale on which j is stored



Outline


- ✓ Motivation
- ✓ Chapel's Design Themes
- Survey of Chapel Concepts



● Project Status and Next Steps

Outline


- ✓ Motivation
- ✓ Chapel's Design Themes
- Survey of Chapel Concepts



- Project Status and Next Steps

Data Parallelism By Example: STREAM Triad

```
const ProblemSpace = {1..m};
```




```
var A, B, C: [ProblemSpace] real;
```




```
A = B + alpha * C;
```

STREAM Triad: Chapel (multicore)

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```




```
A = B + alpha * C;
```

No domain map specified \Rightarrow use default layout


- current locale owns all domain indices and array values
- computation will execute using local processors only

COMPUTE | STORE | ANALYZE


STREAM Triad: Chapel (multilocale, blocked)



```
const ProblemSpace = {1..m}
dmapped Block(boundingBox={1..m});
```




```
var A, B, C: [ProblemSpace] real;
```




```
A = B + alpha * C;
```


STREAM Triad: Chapel (multilocale, cyclic)



```
const ProblemSpace = {1..m}
dmapped Cyclic(startIdx=1);
```



```
var A, B, C: [ProblemSpace] real;
```



```
A = B + alpha * C;
```

STREAM Triad: Chapel

MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include<omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params,
int myRank, commSize;
int rv, errCount;
MPI_Comm comm = MPI_COMM_WORLD;

MPI_Comm_size( comm, &commSize );
MPI_Comm_rank( comm, &myRank );

rv = HPCC_Stream( params, 0 == myR
MPI_Reduce( &rv, &errCount, 1, MPI_
return errCount;

int HPCC_Stream(HPCC_Params *params,
register int j;
double scalar;
VectorSize = HPCC_LocalVectorSize();
a = HPCC_XMALLOC( double, VectorSi
b = HPCC_XMALLOC( double, VectorSi
c = HPCC_XMALLOC( double, VectorSi
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d)\n", VectorSize );
        fclose( outFile );
    }
}
```

Chapel

```
config const m = 1000,
alpha = 3.0;


const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

A = B + alpha * C;
```

the special sauce








Philosophy: Good, *top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

COMPUTE

STORE




ANALYZE

Chapel has Many Types of Domains/Arrays



LULESCH: a DOE Proxy Application

Goal: Solve one octant of the spherical Sedov problem (blast wave) using Lagrangian hydrodynamics for a single material




pictures courtesy of Rob Neely, Bert Still, Jeff Keasler, LLNL


COMPUTE

STORE

ANALYZE



LULESCH in Chapel



COMPUTE

STORE

ANALYZE

LULESH in Chapel

1288 lines of source code

plus 266 lines of comments
 487 blank lines

(the corresponding C+MPI+OpenMP version is nearly 4x bigger)

This can be found in Chapel v1.9 in examples/benchmarks/lulesh/*.chpl

LULESCH in Chapel


This is the only representation-dependent code.
It specifies:

- data structure choices:
 - structured vs. unstructured mesh
 - local vs. distributed data
 - sparse vs. dense materials arrays
- a few supporting iterators


Domain maps insulate the rest of the application
from these choices

Domain Maps

Domain maps are “recipes” that instruct the compiler how to map the global view of a computation...




...to the target locales' memory and processors:



Chapel's Domain Map Philosophy

- 1. Chapel provides a library of standard domain maps**
 - to support common array implementations effortlessly

- 2. Expert users can write their own domain maps in Chapel**
 - to cope with any shortcomings in our standard library



- 3. Chapel's standard domain maps are written using the same end-user framework**
 - to avoid a performance cliff between “built-in” and user-defined cases

For More Information on Domain Maps

HotPAR'10: *User-Defined Distributions and Layouts in Chapel: Philosophy and Framework*
Chamberlain, Deitz, Iten, Choi; June 2010

CUG 2011: *Authoring User-Defined Domain Maps in Chapel*
Chamberlain, Choi, Deitz, Iten, Litvinov; May 2011

Chapel release:

- Technical notes detailing the domain map interface for implementers:
<http://chapel.cray.com/docs/latest/technotes/dsi.html>
- Current domain maps:
\$CHPL_HOME/modules/dists/*.chpl
layouts/*.chpl
internal/Default*.chpl

Two Other Thematically Similar Features

- 1) **parallel iterators:** Permit users to specify the parallelism and work decomposition used by forall loops
 - including zippered forall loops
- 2) **locale models:** Permit users to model the target architecture and how Chapel should be implemented on it
 - e.g., how to manage memory, create tasks, communicate, ...

Like domain maps, these are...

- ...written in Chapel by expert users using lower-level features
 - e.g., task parallelism, on-clauses, base language features, ...
- ...available to the end-user via higher-level abstractions
 - e.g., forall loops, on-clauses, lexically scoped PGAS memory, ...

Language Summary

HPC programmers deserve better programming models

Higher-level programming models can help insulate algorithms from parallel implementation details

- yet, without necessarily abdicating control
- Chapel does this via its multiresolution design
 - domain maps, parallel iterators, and locale models are all examples
 - avoids locking crucial policy decisions into the language definition

We believe Chapel can greatly improve productivity

...for current and emerging HPC architectures

...for HPC users and mainstream uses of parallelism at scale

Outline

- ✓ Motivation
- ✓ Chapel's Design Themes
- ✓ Survey of Chapel Concepts
- Project Status and Next Steps

Chapel is Portable

- Chapel's design is intended to be hardware-independent
- The current release requires:
 - a C/C++ compiler
 - a *NIX environment (Linux, OS X, BSD, Cygwin, ...)
 - POSIX threads
 - (for distributed execution): support for RDMA, MPI, or UDP
- Chapel can run on...
 - ...laptops and workstations
 - ...commodity clusters
 - ...the cloud
 - ...HPC systems from Cray and other vendors
 - ...modern processors like Intel Xeon Phi, GPUs*, etc.

* = academic work only; not yet supported in the official release

Chapel is Open-Source

- Chapel's development is hosted at GitHub
 - <https://github.com/chapel-lang>
- Chapel is licensed as Apache v2.0 software
- Instructions for download + install are online
 - see <http://chapel.cray.com/download.html> to get started

A Year in the Life of Chapel

- **Two major releases per year** (April / October)
 - ~a month later: detailed release notes
- **CHIUW:** Chapel Implementers and Users Workshop (May/June)
 - (3rd annual) CHIUW 2016 will be held at IPDPS (Chicago, IL)
- **SC** (Nov)
 - tutorials, panels, BoFs, posters, educator sessions, exhibits, ...
 - annual **CHUG (Chapel Users Group)** happy hour
- **Talks, tutorials, research visits, blog posts, ...** (year-round)



Chapel is a Collaborative Effort



Colorado
State
University



RICE®



LABORATORY FOR
TELECOMMUNICATIONS
SCIENCES

ETH Zürich

 Lawrence Livermore
National Laboratory

 Sandia National Laboratories



Lawrence Berkeley
National Laboratory

Argonne
NATIONAL LABORATORY

 OAK
RIDGE
National Laboratory


Pacific Northwest
NATIONAL LABORATORY
Proudly Operated by Battelle Since 1965

 東京大学
THE UNIVERSITY OF TOKYO

 UNIVERTAD
DE MÁLAGA




UNIVERSITY OF
MARYLAND

(and many others as well...)

<http://chapel.cray.com/collaborations.html>

COMPUTE

STORE


ANALYZE

Copyright 2016 Cray Inc.

Chapel is a Work-in-Progress

- **Currently being picked up by early adopters**
 - Users who try it generally like what they see
 - Last release got 1400+ downloads over six months
- **Most features are functional and working well**
 - some areas need improvements: strings, object-oriented features, ...
- **Performance can be hit-or-miss**
 - shared memory performance is often competitive with C+OpenMP
 - distributed memory performance needs more work
- **We are actively working to address these lacks**


Stream-EP Performance Over Time



Scalability: RA (atomics) Performance

● RA (atomics) summary

- 5x better performance for ugni-qthreads
- 3x better performance for ugni-muxed



Chapel Resources: For More Information



COMPUTE



STORE



ANALYZE

Chapel Websites

Project page: <http://chapel.cray.com>

- overview, papers, presentations, language spec, ...

GitHub: <https://github.com/chapel-lang>

- download Chapel; browse source repository; contribute code

Facebook: <https://www.facebook.com/ChapelLanguage>

Twitter: <https://twitter.com/ChapelLanguage>



The screenshot shows the Chapel Programming Language Facebook page. It features a large green and blue 'C' logo. The cover photo displays a snippet of Chapel code for task parallelism. The timeline shows a post from the page itself announcing its presence on Facebook, along with a link to the project page. Other posts include snippets of Chapel code and general information about the language.



The screenshot shows the Chapel Language Twitter account (@ChapelLanguage). The profile picture is the same green and blue 'C' logo. The bio describes Chapel as a productive parallel programming language for large-scale computing. A recent tweet discusses switching the memory allocator to jemalloc and includes a link to a performance benchmark chart. The account has 4 tweets, 1 follower, and 19 following.

COMPUTE


STORE

ANALYZE

Suggested Reading

Chapel chapter from ***Programming Models for Parallel Computing***

- a detailed overview of Chapel's history, motivating themes, features
- edited by Pavan Balaji, published by MIT Press
- an early draft is available online,
entitled [*A Brief Overview of Chapel*](#)



Other Chapel papers/publications available at <http://chapel.cray.com/papers.html>

Chapel Blog Articles

[Chapel: Productive Parallel Programming](#), Cray Blog, May 2013.

- *a short-and-sweet introduction to Chapel*

[Chapel Springs into a Summer of Code](#), Cray Blog, April 2016.

- *a run-down of some current events*

[Six Ways to Say “Hello” in Chapel](#) (parts [1](#), [2](#), [3](#)), Cray Blog, Sep-Oct 2015.

- *a series of articles illustrating the basics of parallelism and locality in Chapel*

[Why Chapel?](#) (parts [1](#), [2](#), [3](#)), Cray Blog, Jun-Oct 2014.

- *a series of articles answering common questions about why we are pursuing Chapel in spite of the inherent challenges*

[\[Ten\] Myths About Scalable Programming Languages](#), IEEE TCSC Blog (index available on [chapel.cray.com](#) “blog articles” page), Apr-Nov 2012.

- *a series of technical opinion pieces designed to argue against standard reasons given for not developing high-level parallel languages*

Chapel Mailing Lists

low-traffic (read-only):

chapel-announce@lists.sourceforge.net: announcements about Chapel

community lists:

chapel-users@lists.sourceforge.net: user-oriented discussion list

chapel-developers@lists.sourceforge.net: developer discussions

chapel-education@lists.sourceforge.net: educator discussions

chapel-bugs@lists.sourceforge.net: public bug forum

(subscribe at SourceForge: <http://sourceforge.net/p/chapel/mailman/>)

To contact the Cray team:

chapel_info@cray.com: contact the team at Cray

chapel_bugs@cray.com: for reporting non-public bugs

Questions?



COMPUTE

| STORE

| ANALYZE

Legal Disclaimer

Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.

Cray Inc. may make changes to specifications and product descriptions at any time, without notice.

All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.

Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.

Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.

The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, URIKA, and YARCDATA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.

Copyright 2016 Cray Inc.