

LOCALITY-BASED OPTIMIZATIONS IN THE CHAPEL COMPILER

Engin Kayraklıoglu, Elliot Ronaghan, Michael P. Ferguson, Bradford L. Chamberlain

Hewlett Packard Enterprise

engin@hpe.com

The 34th International Workshop on Languages and Compilers for Parallel Computing

October 13, 2021

WHAT IS CHAPEL?

Chapel: A modern parallel programming language

- portable & scalable
- open-source & collaborative

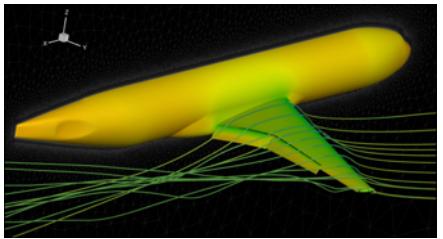


Goals:

- Support general parallel programming
- Make parallel programming at scale far more productive

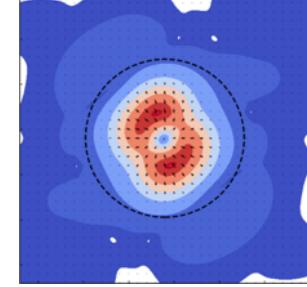


NOTABLE CURRENT APPLICATIONS OF CHAPEL



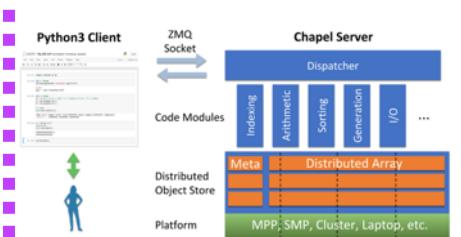
CHAMPS: 3D Unstructured CFD

Éric Laurendeau, Simon Bourgault-Côté,
Matthieu Parenteau, et al.
École Polytechnique Montréal
~48k lines of Chapel



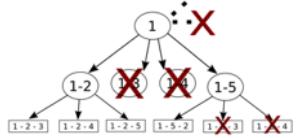
ChplUltra: Simulating Ultralight Dark Matter

Nikhil Padmanabhan, J. Luna Zagorac,
Richard Easter, et al.
Yale University / University of Auckland



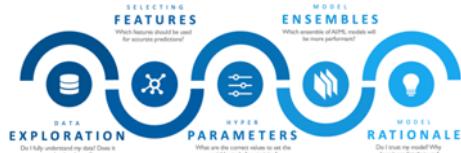
Arkouda: NumPy at Massive Scale

Mike Merrill, Bill Reus, et al.
US DOD
~16k lines of Chapel



ChOp: Chapel-based Optimization

Tiago Carneiro, Nouredine Melab, et al.
INRIA Lille, France



CrayAI: Distributed Machine Learning

Hewlett Packard Enterprise



Your Project Here?

CHAPEL

Distributed Memory Programming in Chapel

```
var myDomain = {1 .. n};
```

```
var myArray: [myDomain] int;
```

A “domain” is an index set in Chapel

This is from 1 to n, inclusive

```
for element in myArray do  
    element = 1;
```

Chapel arrays are declared over domains

This is an integer array over ‘myDomain’

A ‘for’ loop executes sequentially on the locale that started it

This iterates over array elements

CHAPEL

Distributed Memory Programming in Chapel

```
use BlockDist;
```

The ‘dmapped’ clause can be used to make a domain distributed

Now, ‘myDomain’ is block-distributed

```
var myDomain = {1..n} dmapped Block({1..n});
```

```
var myArray: [myDomain] int;
```

No change is necessary for array declaration

Now, ‘myArray’ is block-distributed

```
for element in myArray do
```

```
    element = 1;
```

‘for’ loops are always sequential and executes on the initiating locale

This loop behaves the same; accesses to ‘element’ can be remote

```
use BlockDist;  
  
var myDomain = {1..n} dmapped Block({1..n});  
var myArray: [myDomain] int;
```

A ‘forall’ loop can be distributed and parallel depending on the iterator it executes over

Now, this loop is parallel and distributed identically to ‘myArray’

```
forall element in myArray do  
    element = 1;
```

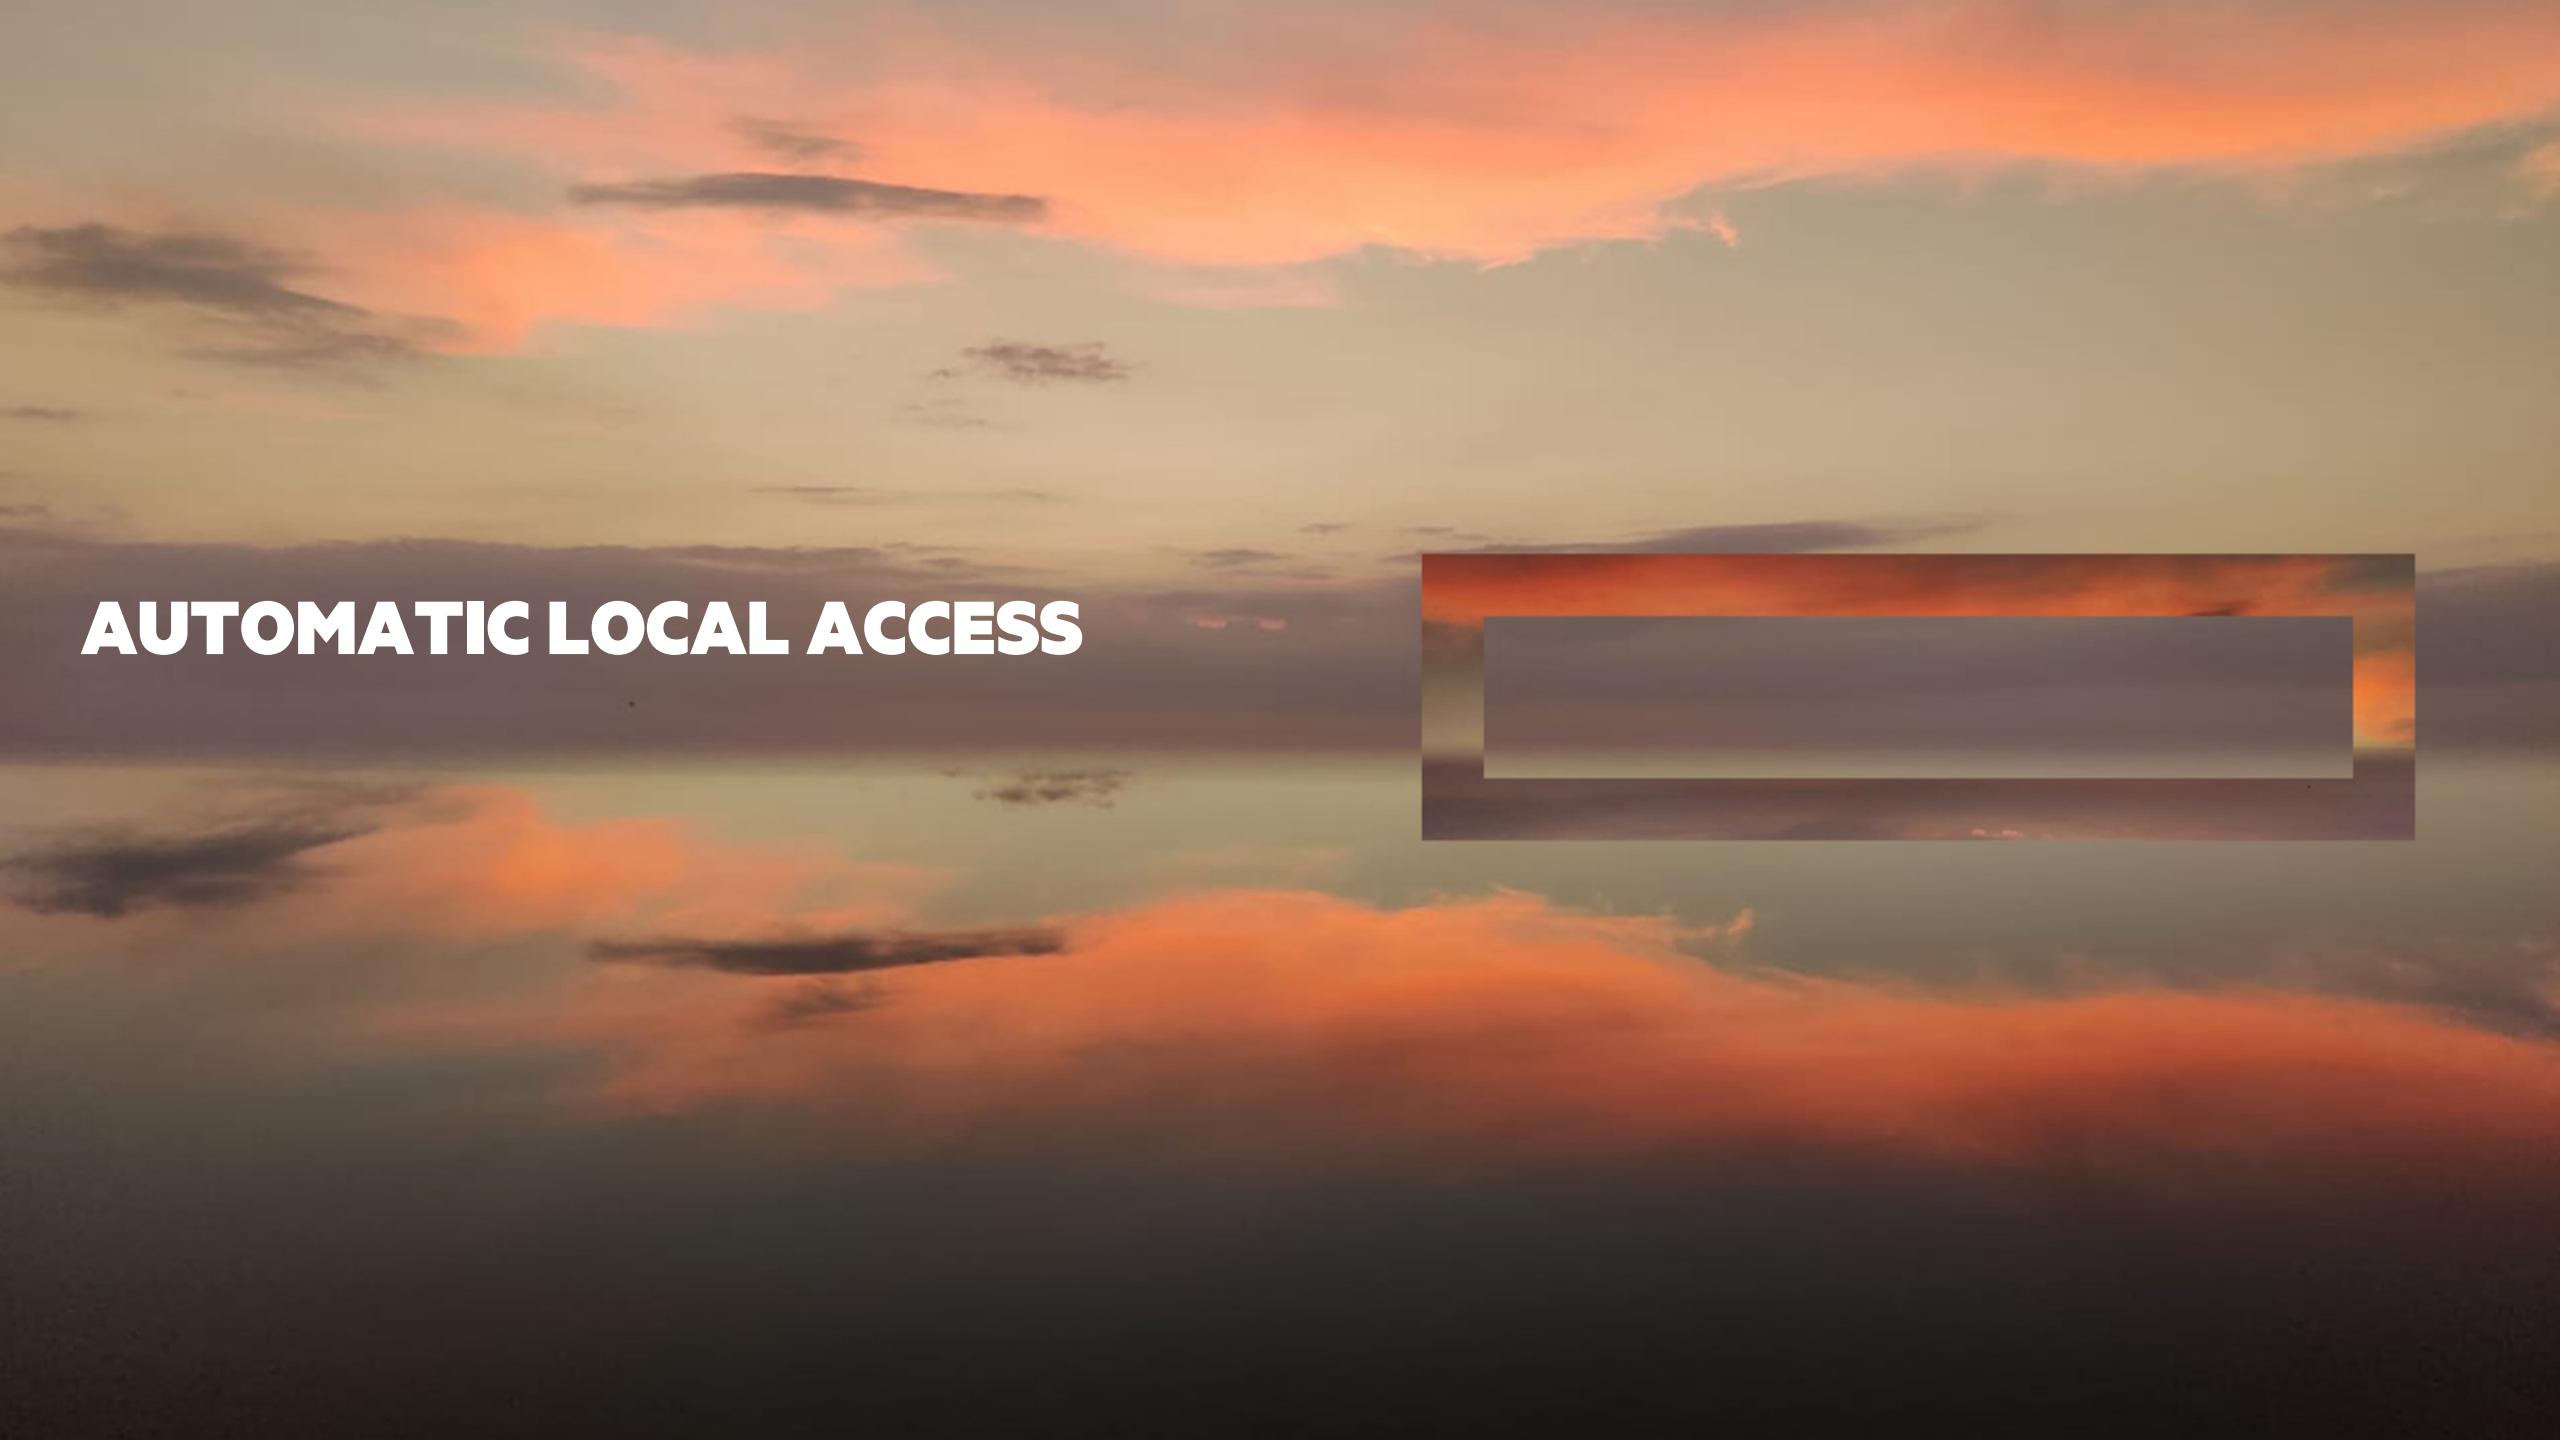
This implies that accesses to ‘element’ are always local

CHAPEL

For More Information

- **Michelle Strout's invited talk:** “Separating Parallel Performance Concerns Using Chapel”
 - 10:40 EDT, today
- **Web page:** chapel-lang.org
- **Development** on GitHub: github.com/chapel-lang/chapel
- **Mailing list** on Discourse: chapel.discourse.group
- **Public chat** on Gitter: gitter.im/chapel-lang/chapel
- **Questions** on StackOverflow: stackoverflow.com/questions/tagged/chapel
- **Talks** on YouTube: youtube.com/c/ChapelParallelProgrammingLanguage
- **News** on Twitter: [@ChapelLanguage](https://twitter.com/@ChapelLanguage)





A wide-angle photograph of a sunset or sunrise over a calm body of water. The sky is filled with soft, wispy clouds, with the light from the sun reflecting off the water and染色 the clouds in shades of orange, yellow, and pink. The overall atmosphere is peaceful and scenic.

AUTOMATIC LOCAL ACCESS



AUTOMATIC LOCAL ACCESS

Before This Optimization

- Three common idioms for implementing STREAM Triad in Chapel

```
var D = newBlockDom(1..n);  
var A, B, C: [D] int;
```

Idiom 1

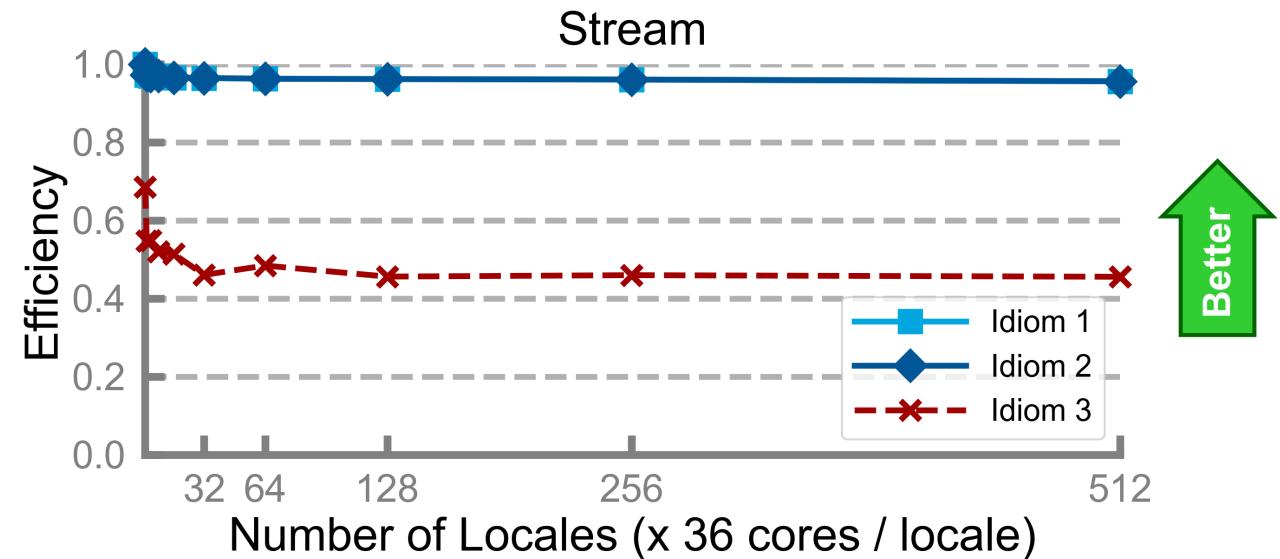
```
A = B + alpha * C;
```

Idiom 2

```
forall (a,b,c) in zip(A,B,C) do  
    a = b + alpha * c;
```

Idiom 3

```
forall i in D do  
    A[i] = B[i] + alpha * C[i];
```



Numbers are collected from a Cray XC with Aries interconnect and ugni communication layer

AUTOMATIC LOCAL ACCESS

Locality Check Overhead

```
var D = newBlockDom(1..n);  
var A, B, C: [D] int;
```

Idiom 3

```
forall i in D do  
  A[i] = B[i] + alpha * C[i];
```

```
proc array.access(idx: int) {  
  if isLocalIndex(idx) then  
    localAccess(idx);  
  else  
    nonLocalAccess(idx);  
}
```

A per-access check is the source of overhead

- This check can be avoided for all 3 accesses
- Because:
 - The ‘forall’ distribution is aligned with A’s ... because the loop is over A.domain
 - The loop index is the same as the access index
 - B and C’s distribution is aligned with A’s ... because they all share the same domain

AUTOMATIC LOCAL ACCESS

Examples

```
var D = newBlockDom({1..N});  
var A: [D] int, B: [D] int;
```

```
forall i in D do  
  A[i] = i;
```

optimized!

the array is indexed using the loop index

the array has the same domain as the loop

```
forall i in D do  
  A[i] = B[i];
```

optimized!

the arrays are indexed using the loop index

the arrays have the same domain as the loop

```
forall i in A.domain do  
  A[i] = B[i];
```

optimized!

the arrays are indexed using the loop index

loop is run over a domain query

the arrays have the same domain as the loop

AUTOMATIC LOCAL ACCESS

Dynamic Checks and Loop Versioning

- If the compiler cannot determine the domain of an array:
 - Equality of domains will be checked at execution time
 - Depending on that, an optimized or unoptimized version of the loop will be run

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int); // currently we can't infer 'B' has the same domain as 'A'
forall i in A.domain do
    A[i] = calculate(B[i]); // B[i] is local if A.domain == B.domain
                            // that can only be confirmed at execution time
```

- Terminology
 - ‘A[i]’ is a static candidate
 - ‘B[i]’ is a dynamic candidate



AUTOMATIC LOCAL ACCESS

Dynamic Checks and Loop Versioning

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
forall i in A.domain do
    A[i] = calculate(B[i]);
```

Static checks are created for both arrays

Dynamic check is created only for B

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
param staticCheckA = canUseLocalAccess(A, A.domain);
param staticCheckB = canUseLocalAccess(B, A.domain);
if staticCheckA || staticCheckB {
    const dynamicCheckB = canUseLocalAccessDyn(B, A.domain);
    if dynamicCheckB then
        forall i in A.domain do
            A.localAccess[i] = calculate(B.localAccess[i]);
    else
        forall i in A.domain do
            A.localAccess[i] = calculate(B[i]);
    } else {
        forall i in A.domain do
            A[i] = calculate(B[i]);
    }
```

AUTOMATIC LOCAL ACCESS

Dynamic Checks and Loop Versioning

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
forall i in A.domain do
    A[i] = calculate(B[i]);
```

Will be executed if

- A passes static checks
- B passes static and dynamic checks

Will be executed if

- A passes static checks
- B fails static or dynamic checks

Will be executed if

- Neither array passes static checks

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
param staticCheckA = canUseLocalAccess(A, A.domain);
param staticCheckB = canUseLocalAccess(B, B.domain);
if staticCheckA || staticCheckB {
    const dynamicCheckB = canUseLocalAccessDyn(B, B.domain);
    if dynamicCheckB then
        forall i in A.domain do
            A.localAccess[i] = calculate(B.localAccess[i]);
    else
        forall i in A.domain do
            A.localAccess[i] = calculate(B[i]);
    } else {
        forall i in A.domain do
            A[i] = calculate(B[i]);
    }
}
```

AUTOMATIC LOCAL ACCESS

After This Optimization: STREAM Triad

```
var D = newBlockDom(1..n);  
var A, B, C: [D] int;
```

Idiom 1

```
A = B + alpha * C;
```

Idiom 2

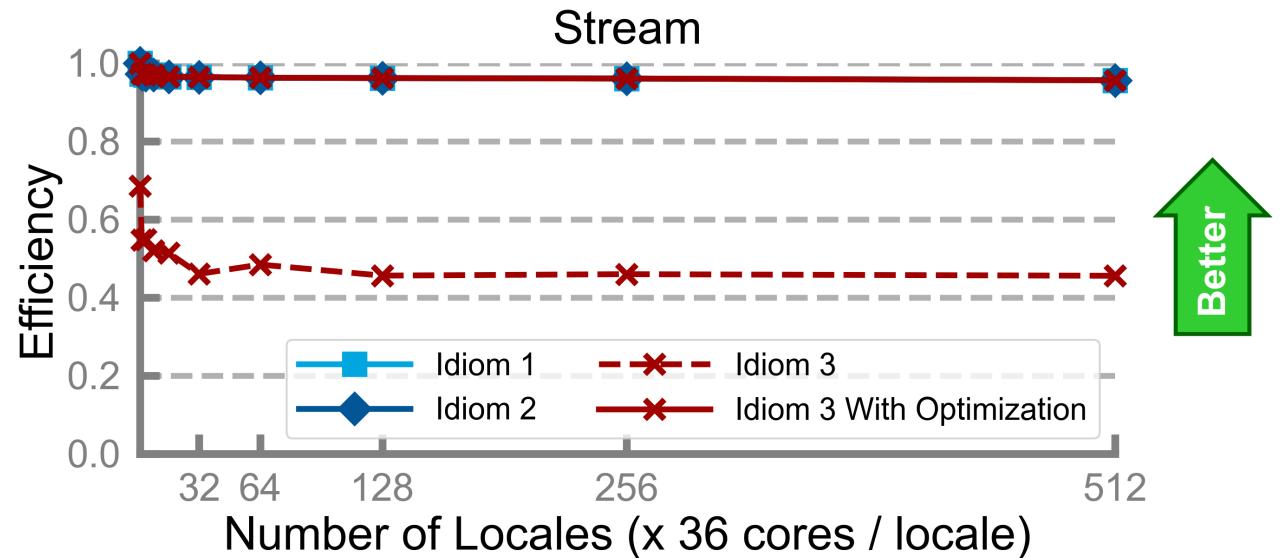
```
forall (a,b,c) in zip(A,B,C) do  
    a = b + alpha * c;
```

Idiom 3

```
forall i in A.domain do  
    A[i] = B[i] + alpha * C[i];
```

Reaches 96% efficiency at-scale

All idioms perform similarly



Numbers are collected from a Cray XC with Aries interconnect and ugni communication layer

AUTOMATIC LOCAL ACCESS

After This Optimization: NAS Parallel Benchmarks - FT

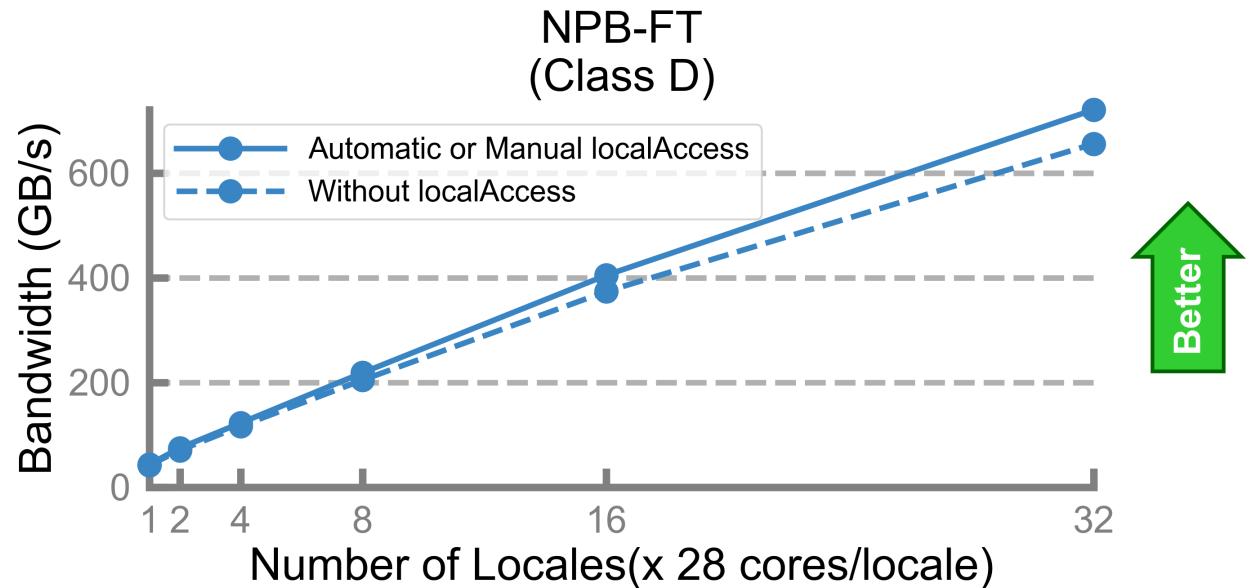
- Explicit 'localAccess' calls are no longer needed in NPB-FT

Kernel with 'localAccess' calls

```
forall ijk in DomT {  
    const elt = V.localAccess[ijk] *  
                T.localAccess[ijk];  
  
    V.localAccess[ijk] = elt;  
    Wt.localAccess[ijk] = elt;  
}
```

Kernel without 'localAccess' calls

```
forall ijk in DomT {  
    const elt = V[ijk] *  
                T[ijk];  
  
    V[ijk] = elt;  
    Wt[ijk] = elt;  
}
```

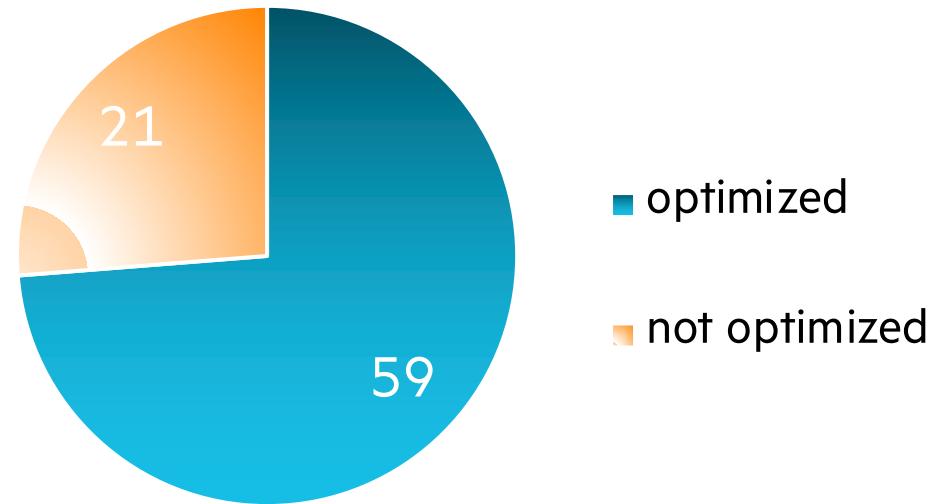


Numbers are collected from a Cray XC with Aries interconnect and ugni communication layer

AUTOMATIC LOCAL ACCESS

After This Optimization: chplUltra

- [chplUltra](#)^[1] is an Ultralight Dark Matter simulator written in Chapel
- We removed all explicit calls to *localAccess*
 - 80 places in total
 - 59 are optimized automatically
 - 21 were not optimized
 - The patterns where the optimization does not fire
 - 10 locality hard to detect due to complex alignments
 - 7 array access indices are not loop indices
 - 4 is not inside forall loops



[1] Nikhil Padmanabhan et al. "Simulating Ultralight Dark Matter in Chapel". In: 2020 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW). May 2020,

AUTOMATIC AGGREGATION



AUTOMATIC AGGREGATION

Before This Optimization

- The *indexgather* benchmark from the [bale](#)^[2] study tests random access performance

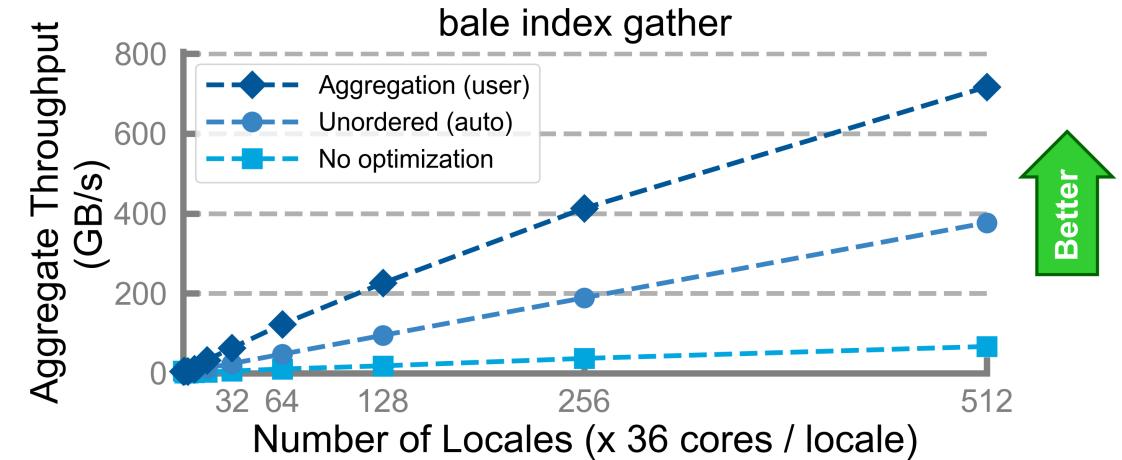
```
var D = newBlockDom(1..n);
var Src, Dst, inds: [D] int;
```

Straightforward *indexgather*

```
forall (d, i) in zip(Dst, inds) do
    d = Src[i];
```

Communication can be done in any order
The Chapel compiler already had
“unordered forall” optimization

Numbers are collected from a Cray XC with Aries interconnect and ugni communication layer



A Manual Approach for Data Aggregation

```
forall (d, i) in zip(Dst, inds) with (var agg = new SrcAggregator(int)) do
    agg.copy(d, Src[i]);
```

[2] <https://github.com/jdevinney/bale>

AUTOMATIC AGGREGATION

Connecting the Dots for Automatic Aggregation

Straightforward *indexgather*

```
forall (d, i) in zip(Dst, Inds) do  
    d = Src[i];
```

What does it take to make this transition automatically?

```
forall (d, i) in zip(Dst, Inds) with (var agg = new SrcAggregator(int)) do  
    agg.copy(d, Src[i]);
```

Q: Is it safe to complete communication in any order?

A: unordered forall optimization does that analysis

Q: Is exactly one side of the operation local?

A: automatic local access optimization does that analysis

Q: Do you have means to aggregate the data?

A: aggregator objects can do that

AUTOMATIC AGGREGATION

Examples

```
var D = newBlockDom({1..N});  
var A: [D] int, B: [D] int;
```

```
forall i in D do  
  A[i] = B[computeIndex(i)];
```

aggregated!

destination of copy is local

source of copy is likely not local

```
forall (a, i) in zip(A, 0..) do  
  B[computeIndex(i)] = a;
```

source is yielded by the first iterand, must be local

aggregated!

destination is likely not local

```
forall (i,a) in zip(A.domain, A) do  
  A[computeIndex(i)] = a;
```

source is yielded by the second iterand

but it is aligned with the first one

aggregated!

destination is likely not local

AUTOMATIC AGGREGATION

After This Optimization: *indexgather*

- The *indexgather* benchmark from the bale study tests random access performance

```
var D = newBlockDom(1..n);
var Src, Dst, inds: [D] int;
```

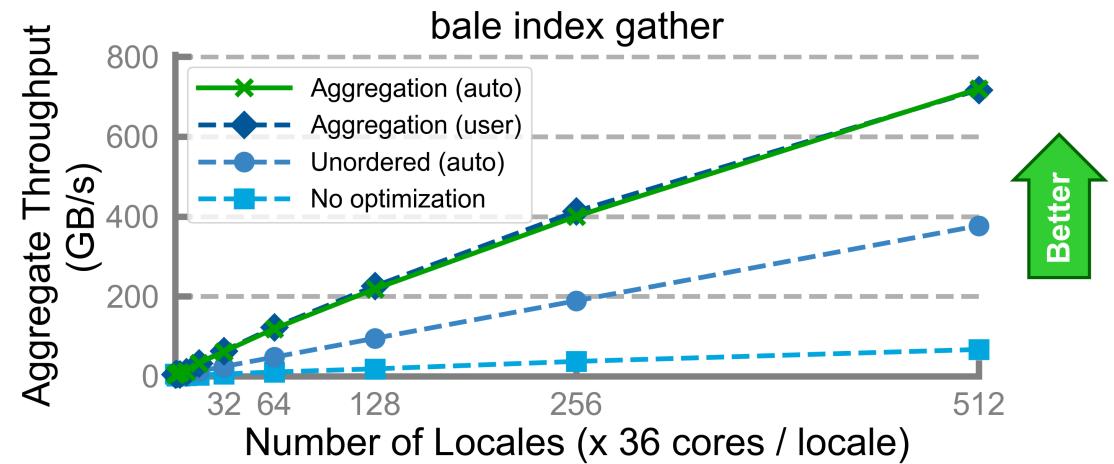
Straightforward *indexgather*

```
forall (d, i) in zip(Dst, inds) do
    d = Src[i];
```

Straightforward version performs identical
to the manually-aggregated version

```
forall (d, i) in zip(Dst, inds) with (var agg = new SrcAggregator(int)) do
    agg.copy(d, Src[i]);
```

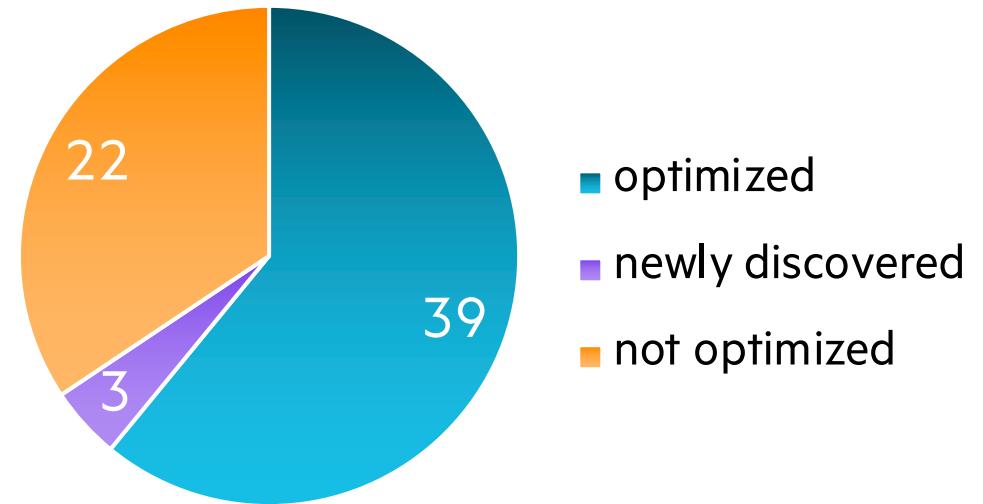
Numbers are collected from a Cray XC with Aries interconnect and ugni communication layer



AUTOMATIC AGGREGATION

After This Optimization: Arkouda

- Arkouda^[3] is a data analytics tool that has a Python client and a server implemented in Chapel
- We removed all the manual aggregation from the source
 - 61 places in total
 - 39 are optimized automatically
 - 22 are not optimized
 - 3 cases that were not using aggregators are now optimized
 - The patterns where the aggregation does not fire:
 - 9: aggregation is not based on ‘forall’ loops
 - 6: compiler cannot prove that unordered operation is safe
 - 3: locality is hard to detect
 - 2: aggregated copy is not in the last statement of the body
 - 1: one side of the assignment is defined within the loop body
 - 1: needs further investigation



[3] <https://github.com/Bears-R-Us/arkouda>

LIMITATIONS & NEXT STEPS

Automatic Local Access

- Can we do the same optimizations when the index is a complex expression?
 - Today: Access must be at same index as the loop index

Automatic Aggregation

- Support arbitrary operations
 - Today: Limited to copy operations (i.e., '=' operator)
- Improve worst-case performance
 - Today: If everything is local, aggregation adds overhead and can reduce performance by half
- Investigate multi-hop aggregation
 - Today: Per-locale buffers can have a large memory footprint significantly at-scale
- Expose aggregation as a user-facing language feature
 - Today: The aggregator objects are not in the documented part of the standard library



SUMMARY

- We have discussed two locality-based optimizations in the Chapel compiler

Automatic Local Access

- Avoids locality checks while accessing distributed arrays using indices
- Added to Chapel in version 1.23 (1.25 was released few weeks ago)
- On-by-default

Automatic Aggregation

- Aggregates some remote copy operations
- Added to Chapel in version 1.24 (1.25 was released few weeks ago)
- Off-by-default, enable with --auto-aggregation



The background image shows a wide-angle view of a sunset or sunrise over a body of water. The sky is filled with various shades of orange, yellow, and blue, with wispy clouds scattered across it. The horizon line is visible in the distance.

THANK YOU!



engin@hpe.com