



**Hewlett Packard  
Enterprise**


# **MAKING PARALLEL PROGRAMMING AND GPUS MORE ACCESSIBLE WITH CHAPEL**

---


Engin Kayraklıoglu  
May 31st, 2024

[engin@hpe.com](mailto:engin@hpe.com)  
[linkedin.com/in/engink](https://www.linkedin.com/in/engink)

# IT IS HARD TO AVOID GPUS IN HPC



# GPUS ARE EASY TO FIND... BUT DIFFICULT TO PROGRAM



- ... but programming for multiple nodes with GPUs appears to require at least 2 programming models
  - all of the models rely on C/C++/Fortran, which are different than the languages being taught these days
  - as a result, *using GPUs in HPC has a high barrier of entry*

**Chapel is an alternative for productive  
distributed/shared memory GPU programming in a vendor-neutral way.**

# WHAT IS CHAPEL?

---

**Chapel:** A modern parallel programming language

- portable & scalable
- open-source & collaborative



## Goals:

- Support general parallel programming
- Make parallel programming at scale far more productive

[chapel-lang.org](http://chapel-lang.org)



# **WHAT IS CHAPEL?**

---

## **Chapel works everywhere**

- you can develop on your laptop and have the code scale on a supercomputer
- GPUs can be targeted in a vendor-neutral way
- runs on Linux laptops/clusters, Cray systems, MacOS, WSL, AWS, Raspberry Pi
- shown to scale on Cray networks (Slingshot, Aries), InfiniBand, RDMA-Ethernet

## **Chapel makes distributed/shared memory parallel programming easy**

- data-parallel, locality-aware loops,
- ability to move execution and allocation to remote nodes,
- distributed arrays and bulk array operations
- different types of parallelism can be expressed with the same language features



# **WHAT IS CHAPEL?**

---

## **Chapel works everywhere**

- you can develop on your laptop and have the code scale on a supercomputer
- GPUs can be targeted in a vendor-neutral way
- runs on Linux laptops/clusters, Cray systems, MacOS, WSL, AWS, Raspberry Pi
- shown to scale on Cray networks (Slingshot, Aries), InfiniBand, RDMA-Ethernet


## **Chapel makes distributed/shared memory parallel programming easy**

- data-parallel, locality-aware loops,
- ability to move execution and allocation to remote nodes,
- distributed arrays and bulk array operations
- different types of parallelism can be expressed with the same language features




# APPLICATIONS OF CHAPEL

Active GPU efforts




**CHAMPS: 3D Unstructured CFD**

Laurendeau, Bourgault-Côté, Parenteau, Plante, et al.  
École Polytechnique Montréal




**Arkouda: Interactive Data Science at Massive Scale**

Mike Merrill, Bill Reus, et al.  
U.S. DoD




**ChOp: Chapel-based Optimization**

T. Carneiro, G. Helbecque, N. Melab, et al.  
INRIA, IMEC, et al.




**ChplUltra: Simulating Ultralight Dark Matter**

Nikhil Padmanabhan, J. Luna Zagorac, et al.  
Yale University et al.




**Lattice-Symmetries: a Quantum Many-Body Toolbox**

Tom Westerhout  
Radboud University




**Chapel-based Hydrological Model Calibration**

Marjan Asgari et al.  
University of Guelph



**Desk dot chpl: Utilities for Environmental Eng.**

Nelson Luis Dias  
The Federal University of Paraná, Brazil




**CrayAI HyperParameter Optimization (HPO)**

Ben Albrecht et al.  
Cray Inc. / HPE




**RapidQ: Mapping Coral Biodiversity**

Rebecca Green, Helen Fox, Scott Bachman, et al.  
The Coral Reef Alliance



**CHGL: Chapel Hypergraph Library**

Louis Jenkins, Cliff Joslyn, Jesun Firoz, et al.  
PNNL



**ChapQG: Layered Quasigeostrophic CFD**

Ian Grooms and Scott Bachman  
University of Colorado, Boulder et al.




**Your Application Here?**

(Images provided by their respective teams and used with permission)

# CORAL REEF SPECTRAL BIODIVERSITY


1. Read in a  $(M \times N)$  raster image of habitat data
2. Create a  $(P \times P)$  mask to find all points within a given radius.
3. Convolve this mask over the entire domain and perform a weighted reduce at each location.



Algorithmic complexity:  $O(MNP^3)$

Typically:

- $M, N > 10,000$
- $P \sim 400$



# CORAL REEF SPECTRAL BIODIVERSITY

---

```
proc convolve(InputArr, OutputArr) { // 3D Input, 2D Output
    for ... {
        tonOfMath();
    }
}

proc main() {
    var InputArr: ...;
    var OutputArr: ...;

    convolve(InputArr, OutputArr);
}
```



# CORAL REEF SPECTRAL BIODIVERSITY

```
proc convolve(InputArr, OutputArr) { // 3D Input, 2D Output
    foreach ... {
        tonOfMath();
    }
}

proc main() {
    var InputArr: ...;
    var OutputArr: ...;

    coforall loc in Locales do on loc {
        coforall gpu in here.gpus do on gpu {
            coforall task in 0..#numWorkers {
                var MyInputArr = InputArr[...];
                var MyOutputArr: ...;
                convolve(MyInputArr, MyOutputArr);
                OutputArr[...] = MyOutputArr;
            }
        }
    }
}
```

**Using a different loop flavor to enable GPU execution.**

**Multi-node, multi-GPU, multi-thread parallelism are expressed using the same language constructs.**


**High-level, intuitive array operations work across nodes and/or devices**

# CORAL REEF SPECTRAL BIODIVERSITY

```
proc convolve(InputArr, OutputArr) { //3D Input  
    foreach ... {  
        tonOfMath();  
    }  
}  
  
proc main() {  
    var InputArr: ...;  
    var OutputArr: ...;  
  
    coforall loc in Locales do on loc { //using parallelism  
        coforall gpu in here.gpus do on gpu { //using parallelism  
            coforall task in 0..#numWorkers { //using parallelism  
                var MyInputArr = InputArr[...];  
                var MyOutputArr: ...;  
                convolve(MyInputArr, OutputArr);  
                OutputArr[...] = MyOutputArr;  
            } } } }
```

## Runs on Frontier!

- 5x improvement going from 2 to 64 nodes
  - (from 16 to 512 GPUs)
- Straightforward code changes:
  - from sequential Chapel code
  - to GPU-enabled one
  - to multi-node, multi-GPU, multi-thread



- Scalability improvements coming soon!

# **WHAT WE WILL DISCUSS TODAY**

---

- Native GPU programming in Chapel using simple snippets
- Two stories from the community analyzing performance

## **What we will not discuss today:**


- Comprehensive list of Chapel features
  - (important ones will be covered)
- How GPU support is implemented
  - (happy to go over some backup slides, if there's interest)
- Everything you can do with GPUs using Chapel
  - (there's only so much time 😊 )



# **GPU PROGRAMMING IN CHAPEL**

# LOCALES IN CHAPEL

- In Chapel, a *locale* refers to a compute resource with...
  - processors, so it can run tasks
  - memory, so it can store variables
- For now, think of each compute node as being a locale




# KEY BUILT-IN TYPES AND VARIABLES RELATED TO LOCALES

`locale`: A type that represents system resources on which the program can run


`Locales`: An array of `locale` values

`here` : The `locale` on which the current task is executing




# KEY CONCERNS FOR SCALABLE PARALLEL COMPUTING

- parallelism:** Which tasks should run simultaneously?
- locality:** Where should tasks run? Where should data be allocated?



# PARALLELISM AND LOCALITY



```
⚙️ var A: [1..2, 1..2] real;  
on Locales[1] {  
    var B: [1..2, 1..2] real;  
    B = 2;  
    A = B;  
}  
  
writeln(A);
```

# PARALLELISM AND LOCALITY

```
forall b in B do  
  b = 2;
```

```
var A: [1..2, 1..2] real;
```


```
forall i in B.domain do  
  B[i] = 2;
```

```
on Locales[1] {  
  var B: [1..2, 1..2] real;  
  B = 2;  
  A = B;  
}
```

Can be expressed  
in different ways

```
writeln(A);
```

# PARALLELISM AND LOCALITY




```
⚙️ var A: [1..2, 1..2] real;  
      ⚙️  
for l in Locales do on l {  
  var B: [1..2, 1..2] real;  
  B = 2;  
  A = B;  
}  
      ⚙️
```

```
⚙️ writeln(A);
```

# PARALLELISM AND LOCALITY

CPU Core      Memory




The coforall loop creates  
a parallel task per iteration

```
⚙️ var A: [1..2, 1..2] real;  
⚙️  
coforall l in Locales do on l {  
    var B: [1..2, 1..2] real;  
    B = 2;  
    A = B;  
}  
⚙️ writeln(A);
```


# KEY CONCERNS FOR SCALABLE PARALLEL COMPUTING

- 1. parallelism:** Which tasks should run simultaneously?
- 2. locality:** Where should tasks run? Where should data be allocated?
  - complicating matters, compute nodes now often have GPUs with their own processors and memory




# KEY CONCERNS FOR SCALABLE PARALLEL COMPUTING

- 1. parallelism:** Which tasks should run simultaneously?
- 2. locality:** Where should tasks run? Where should data be allocated?
  - complicating matters, compute nodes now often have GPUs with their own processors and memory
  - we represent these as *sub-locales* in Chapel



# PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

CPU Core    GPU Core    Memory



```
const nProcs = 1,  
      nRows = nProcs,  
      nCols = 5;           ← We'll work with a single GPU in this step  
  
var A: [1..nRows, 1..nCols] real;  
     ← Number of rows and columns in the array  
  
on here.gpus[0] {  
    var B: [1..nRows, 1..nCols] real;  
    B = 2;  
    A = B;  
}  
  
writeln(A);
```


2D array of real values

# PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

CPU Core    GPU Core    Memory

```
const nProcs = here.gpus.size,  
nRows = nProcs,  
nCols = 5;
```

Now, we'll use all the local GPUs



Each iteration of 'coforall' will run in parallel

'on' now targets the yielded GPU

```
var A: [1..nRows, 1..nCols] real;
```

Iterating GPUs and 1.. in lockstep manner

```
coforall (gpu, gRow) in zip(here.gpus, 1..) {  
    on gpu {  
        var B: [1..nCols] real;  
        B = 2;  
        A[gRow, ...] = B;  
    }  
}
```


Per-GPU array is now 1D

Assigning B to only a single  
row of A

```
writeln(A);
```

# PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

CPU Core   GPU Core   Memory



```
const nProcs = here.gpus.size,  
nRows = nProcs,  
nCols = 5;  
  
var A: [1..nRows, 1..nCols] real;
```

```
coforall (gpu, gRow) in zip(here.gpus, 1..) {  
    on gpu do setRowToTwo(gRow);
```


```
}
```

```
writeln(A);  
  
proc setRowToTwo(row) {  
    var B: [1..nCols] real;  
    B = 2;  
    A[row, ..] = B;  
}
```

Small refactor to put the application logic in a function

# PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

CPU Core    GPU Core    Memory



```
const nProcs = here.gpus.size,  
      nRows = nProcs,  
      nCols = 5;  
  
var A: [1..nRows, 1..nCols] real;
```

```
coforall (gpu, gRow) in zip(here.gpus, 1..) {  
    on gpu do setRowToTwo(gRow);  
}
```


```
writeln(A);
```

```
proc setRowToTwo (row) { /* body omitted */ }
```

Body of this function will  
always be the same

# PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

CPU Core GPU Core Memory



```
const nProcs = here.gpus.size,  
nRows = Locales.size*nProcs,  
nCols = 5;
```

Now, we also use all locales

```
var A: [1..nRows, 1..nCols] real;
```

```
coforall (loc, cRow) in zip(Locales, 1.. by nProcs) {  
    on loc {
```

Iterate locales and starting row of each locale in a lockstep manner

```
        coforall (gpu, gRow) in zip(here.gpus, cRow..) {  
            on gpu do setRowToTwo(gRow);  
        }
```

```
    }  
}
```


For each locales' GPUs, we now start the offset at cRow instead of 1

```
writeln(A);
```

```
proc setRowToTwo (row) { /* body omitted */ }
```

# PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

CPU Core    GPU Core    Memory



```
const nProcs = here.gpus.size+1,  
      nRows = Locales.size*nProcs,  
      nCols = 5;  
  
var A: [1..nRows, 1..nCols] real;  
  
coforall (loc, cRow) in zip(Locales, 1.. by nProcs) {  
    on loc {  
        cobegin {  
            1 setRowToTwo (cRow);  
            2 coforall (gpu, gRow) in zip(here.gpus, cRow+1..) {  
                on gpu do setRowToTwo (gRow);  
            }  
        }  
    }  
}  
writeln(A);  
  
proc setRowToTwo (row) { /* body omitted */ }
```

+1 for the CPU per locale

The 2 statements in 'cobegin' will run in parallel

# DIFFERENT TYPES OF PARALLELISM EXPRESSED CONCISELY

```
const nProcs = here.gpus.size+1,  
      nRows = Locales.size*nProcs,  
      nCols = 5;  
  
var A: [1..nRows, 1..nCols] real;  
  
coforall (loc, cRow) in zip(Locales, 1.. by nProcs) do on loc {  
    cobegin {  
        setRowToTwo(cRow);  
        coforall (gpu, gRow) in zip(here.gpus, cRow+1..) do on gpu {  
            setRowToTwo(gRow);  
        }  
    }  
    writeln(A);  
  
proc setRowToTwo(row) {  
    var B: [1..nCols] real;  
    B = 2;  
    A[row, ..] = B;  
}
```

**GPU programming in Chapel doesn't require learning new concepts**

The only GPU-specific concept in the language  
is 'gpus' array on 'locale' type

**Full code in a single slide that will use**

- ✓ all nodes
- ✓ all CPU cores
- ✓ all GPU cores

**Made possible by Chapel's  
parallelism and locality constructs**

# **STORIES FROM THE CHAPEL COMMUNITY**

# CHAPEL PERFORMANCE ON DIFFERENT GPU AND CPUS

- Comparing Chapel's performance
  - ...against OpenMP, Kokkos, CUDA and HIP
  - ...on different GPU and CPUs
  - ...using BabelStream, miniBUDE and TeaLeaf
- Recently presented at
  - Heterogeneity in Computing Workshop (HCW)
  - In conjunction with IPDPS

## Performance Portability of the Chapel Language on Heterogeneous Architectures

Josh Milthorpe

*Oak Ridge National Laboratory*

Oak Ridge, Tennessee, USA

*Australian National University*

Canberra, Australia

ORCID: 0000-0002-3588-9896

Xianghao Wang

*Australian National University*

Canberra, Australia

Ahmad Azizi

*Australian National University*

Canberra, Australia

*Abstract*—A performance-portable application can run on a variety of different hardware platforms, achieving an acceptable level of performance without requiring significant rewriting for each platform. Several performance-portable programming models are now suitable for high-performance scientific application development, including OpenMP and Kokkos. Chapel is

other heterogeneous programming models that allow single-source programming for diverse hardware platforms.

We seek to answer the question: how well does Chapel support the development of *performance-portable* application codes compared to more widely-used programming models

Paper is available at [milthorpe.org/wp-content/uploads/2024/03/Milthorpe\\_HCW2024.pdf](http://milthorpe.org/wp-content/uploads/2024/03/Milthorpe_HCW2024.pdf)

# MINIBUDE

- Proxy for BUDE (a protein docking simulation)
  - The computation is very arithmetically intensive and makes significant use of trigonometric functions




Figure from: "Performance Portability of the Chapel Language on Heterogeneous Architectures". Josh Milthorpe (Oak Ridge National Laboratory, Australian National University), Xianghao Wang (Australian National University), Ahmad Azizi (Australian National University) Heterogeneity in Computing Workshop (**HCW**)

# BABELSTREAM

- Performs stream triad computation computing  $A = B + \alpha * C$  for arrays A, B, C and scalar  $\alpha$




Figure from: "Performance Portability of the Chapel Language on Heterogeneous Architectures". Josh Milthorpe (Oak Ridge National Laboratory, Australian National University), Xianghao Wang (Australian National University), Ahmad Azizi (Australian National University) Heterogeneity in Computing Workshop (**HCW**)

# TEALEAF

- Tealeaf simulates heat conduction over time
- On this application Chapel performed well on CPUs but not GPUs
  - We have some leads for performance issues and still investigating




Figure from: "Performance Portability of the Chapel Language on Heterogeneous Architectures". Josh Milthorpe (Oak Ridge National Laboratory, Australian National University), Xianghao Wang (Australian National University), Ahmad Azizi (Australian National University). Heterogeneity in Computing Workshop 2024 (**HCW**)

# NATIVE GPU PROGRAMMING IN CHAPEL AT SCALE

- Comparing Chapel's native GPU programming
  - ...against interoperability with HIP and CUDA
  - ...on Frontier and Perlmutter
  - ...using N-Queens as proxy for combinatorial optimization
- To be presented at Euro-Par 2024
  - 26-30 August
  - Madrid, Spain

## Investigating Portability in Chapel for Tree-based Optimization on GPU-powered Clusters

Tiago Carneiro<sup>1</sup>[0000-0002-6145-8352], Engin Kayraklioglu<sup>2</sup>[0000-0002-4966-3812], Guillaume Helbecque<sup>3,4</sup>[0000-0002-8697-3721], and Nouredine Melab<sup>4</sup>


<sup>1</sup> Interuniversity Microelectronics Centre (IMEC), Belgium  
[tiago.carneiro@imec.be](mailto:tiago.carneiro@imec.be)

<sup>2</sup> Hewlett Packard Enterprise, USA  
[engin@hpe.com](mailto:engin@hpe.com)


<sup>3</sup> University of Luxembourg, Luxembourg  
[guillaume.helbecque@uni.lu](mailto:guillaume.helbecque@uni.lu)

<sup>4</sup> Université de Lille, CNRS, Centrale Lille, Inria, UMR 9189 - CRISTAL - Centre de Recherche en Informatique Signal et Automatique de Lille, France  
[nouredine.melab@univ-lille.fr](mailto:nouredine.melab@univ-lille.fr)

# NATIVE GPU PROGRAMMING IN CHAPEL AT SCALE



(a) NVIDIA-based System



(b) AMD-based system

Figure from: "Investigating Portability in Chapel for Tree-Based Optimizations on GPU-powered Clusters". Tiago Carneiro, Engin Kayraklıoglu, Guillaume Helbecque, Nouredine Melab  
Europar 2024

# UPCOMING: KEYNOTE AT CHAPELCON '24

## A Case for Parallel-First Languages in Post-Serial, Accelerated World

Paul Sathre, Virginia Tech

June 7th, 2024

Parallel processors have finally dominated all scales of computing hardware, from the personal and portable to the ivory tower datacenters of yore. However, dominant programming models and pedagogy haven't kept pace, and languish in a post-serial mix of libraries and language extensions. Further, heterogeneity in the form of GPUs has dominated the performance landscape of the last decade, penetrating casual user markets thanks to data science, crypto and AI booms. Unfortunately, GPUs' performance remains largely constrained to expert users by the lack of more productive and portable programming abstractions. This talk will probe questions about how to rethink and democratize parallel programming for the masses. By reflecting on lessons learned from a decade and a half of accelerated computing, I'll show where Chapel 2.0 fits into the lineage of GPU computing, can capitalize on GPU momentum, and lead a path forward.



ChapelCon '24 is free and fully virtual  
[chapel-lang.org/ChapelCon24.html](http://chapel-lang.org/ChapelCon24.html)



# **SUMMARY**


# WHERE WE ARE TODAY

## Over ~3 years we have been steadily improving

- NVIDIA, AMD GPUs are supported
- Multiple nodes with multiple GPUs can be used
- Parallel tasks can use GPUs concurrently
- GPU features can be emulated on CPUs

## Mature enough to get started, big efforts are still underway

- Distributed arrays
- Intel support
- Improving language features to support GPU programming
- Performance improvements
- Bug fixes



# IF YOU WANT TO LEARN MORE ABOUT GPU PROGRAMMING IN CHAPEL

**GPU Programming Blog Series:** [chapel-lang.org/blog/series/gpu-programming-in-chapel/](https://chapel-lang.org/blog/series/gpu-programming-in-chapel/)

## Introduction to GPU Programming in Chapel

Posted on January 10, 2024.

Tags: [GPU Programming](#) [How-To](#)

By: [Daniel Fedorin](#)

## Chapel's High-Level Support for CPU-GPU Data Transfers and Multi-GPU Programming

Posted on April 25, 2024.

Tags: [GPU Programming](#) [How-To](#)

By: [Engin Kayraklıoglu](#)

**Technote:** <https://chapel-lang.org/docs/main/technotes/gpu.html>

- Anything and everything about our GPU support
  - configuration, advanced features, links to some tests, caveats/limitations
- More of a reference manual than a tutorial

## Previous talks

- **LinuxCon / Open Source Summit North America 2024 Talk:** GPU Programming in Chapel and a Live Demo
  - <https://youtu.be/5-jLdKduaJE?si=ezaz5mDORvmTjgRL>
- **CHIUW '23 Talk:** updates from May '22-May '23 period
  - <https://chapel-lang.org/CHIUW/2023/KayraklıogluSlides.pdf>
- **LCPC '22 Talk:** a lot of details on how the Chapel compiler works to create GPU kernels
  - <https://chapel-lang.org/presentations/Engin-SIAM-PP22-GPU-static.pdf>



# HPE DEVELOPER MEETUP

## Meetup for "Vendor-Neutral GPU Programming in Chapel"

Jul 31, 2024 08:00 AM PDT (-7 UTC)

Jade Abraham, Engin Kayraklioglu



speakers will discuss Chapel's GPU support in detail and collaborate with you to determine how it may help in your particular situation.



**Registration:**

[https://hpe.zoom.us/webinar/register/3117139444656/WN\\_ojVy9LR\\_QHSCGxeg21rj7A](https://hpe.zoom.us/webinar/register/3117139444656/WN_ojVy9LR_QHSCGxeg21rj7A)

**HPE developer meetups home page:**

<https://developer.hpe.com/campaign/meetups/>



# CHAPELCON '24 (FORMERLY CHI UW, THE CHAPEL IMPLEMENTERS AND USERS WORKSHOP)

Fully virtual, free registration

## Schedule Overview

- June 5<sup>th</sup>: **Tutorial Day**: Chapel and Arkouda tutorials
- June 6<sup>th</sup>: **Coding Day**: Opportunities for coding
- June 7<sup>th</sup>: **Conference Day**

## Keynote

### A Case for Parallel-First Languages in a Post-Serial, Accelerated World

Paul Sathre (*Virginia Tech*)



**Registration:**





## The Chapel Parallel Programming Language

### ChapelCon '24

The Chapel Event of the Year

June 5–7, 2024  
free and online in a virtual format

[Register Here](#)

ChapelCon '24 welcomes anyone with computing challenges that demand performance, particularly through parallelism and scalability. Our open-source

<https://chapel-lang.org/ChapelCon24.html>



## Chapel Language Blog

About Chapel Website Featured Series Tags Authors All Posts

### Introducing ChapelCon '24: The Chapel Event of the Year

Posted on April 1, 2024.

Tags: [ChapelCon](#) [CHI UW](#) [Community](#)

By: [Engin Kayraklıoglu](#)

If you are following Chapel's Discourse or social media, you might have seen that this year we are

<https://chapel-lang.org/blog/posts/chapelcon24/>

# CHAPEL RESOURCES

**Chapel homepage:** <https://chapel-lang.org>

- (points to all other resources)


**Blog:** <https://chapel-lang.org/blog/>

## Social Media:

- Facebook: [@ChapelLanguage](#)
- LinkedIn: <https://www.linkedin.com/company/chapel-programming-language/>
- Mastodon: [@ChapelProgrammingLanguage](#)
- X / Twitter: [@ChapelLanguage](#)
- YouTube: [@ChapelLanguage](#)

## Community Discussion / Support:

- Discourse: <https://chapel.discourse.group/>
- Gitter: <https://gitter.im/chapel-lang/chapel>
- Stack Overflow: <https://stackoverflow.com/questions/tagged/chapel>
- GitHub Issues: <https://github.com/chapel-lang/chapel/issues>



The Chapel Parallel Programming Language

**What is Chapel?**

Chapel is a programming language designed for productive parallel computing at scale.

**Why Chapel?** Because it simplifies parallel programming through elegant support for:

- data parallelism to trivially use the cores of a laptop, cluster, or supercomputer
- task parallelism to create concurrency within a node or across the system
- a global namespace supporting direct access to local or remote variables
- GPU programming in a vendor-neutral manner using the same features as above
- distributed arrays that can leverage thousands of nodes' memories and cores



**Chapel Characteristics**

- productive: code tends to be similarly readable/writable as Python
- scalable: runs on laptops, clusters, the cloud, and HPC systems
- fast: performance [competes with or beats](#) conventional HPC programming models
- portable: compiles and runs in virtually any \*nix environment
- open-source: hosted on [GitHub](#), permissively licensed
- production-ready: used in [real-world applications](#) spanning diverse fields

**New to Chapel?**

As an introduction to Chapel, you may want to...

- watch an [overview talk](#) or browse its [slides](#)
- read a [chapter-length](#) introduction to Chapel
- learn about [projects powered by Chapel](#)
- check out [performance highlights](#) like these:



- read about [GPU programming](#) in Chapel, or [watch a recent talk about it](#)
- [browse sample programs](#) or [learn how](#) to write distributed programs like this one:

```
use CyclicDist;           // use the Cyclic distribution library
config const n = 100;      // use --n=eval> when executing to override this def
forall i in Cyclic.createDomain(1..n) do
    writeln("Hello from iteration ", i, " of ", n, " running on node ", here.i)
```

**What's Hot?**

- [ChapelCon '24](#) is coming in June (online)—[Read](#) about it and [register](#) today
- Doing science in Python and needing more speed/scale? [Maybe we can help?](#)

# SUMMARY

---

- GPUs are becoming more and more common in HPC
  - However, programming GPUs is more challenging than programming CPUs
  - On multiple nodes, users are typically required to use multiple paradigms
- HPC and GPUs should be more accessible
  - from wider range of disciplines,
  - with varying levels of expertise, and
  - limited time to invest in programming
- Chapel wants to make HPC more accessible
  - Existing applications prove that Chapel delivers on the promise
  - Its growing support for GPU programming can:
    - enable programming GPUs in a productive and vendor-neutral way
    - provide an all-inclusive solution for programming in HPC
- The Chapel team at HPE would be excited to collaborate with AMD!
  - Please feel free to reach out: [engin@hpe.com](mailto:engin@hpe.com)



[chapel-lang.org](http://chapel-lang.org)

