



# Caching Puts and Gets in a PGAS Language Runtime

*Michael Ferguson*  
Cray Inc.

Daniel Buettner  
Laboratory for  
Telecommunication Sciences

September 17, 2015



# Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.





CC Flickr/Daniel Jolivet




CC Flickr/Ben Salter




CC Flickr/ajmexico


# TRAIN LATENCY

(8 HOUR TRIP, 60 TON CARS, 60 SEC/CAR)




# TRAIN BANDWIDTH






# INFINIBAND (IB) LATENCY

\* with small 10-node cluster, QDR IB



# INFINIBAND (IB) BANDWIDTH

\* with small 10-node cluster, QDR IB



# AGGREGATION



CC Flickr/ajmexico

# OVERLAP



CC Flickr/Barry Lewis

CACHE HELPS WITH BOTH!

# BACKGROUND: MEMORY MODEL ALLOWS PREFETCH AND WRITE-BEHIND



# Memory model for C11, C++11, Chapel: *data race free programs are sequentially consistent*

- See Adve, S.V., Boehm, H.-J. 2010. Memory models: a case for rethinking parallel languages and hardware. Communications of the ACM 53(8): 90–101. <http://cacm.acm.org/magazines/2010/8/96610-memory-models-a-case-for-rethinking-parallel-languages-and-hardware/fulltext>

# A RACY PROGRAM

Thread 1

```
x = 42;  
notify = 1;
```

Thread 2

```
while 0 == notify { /* wait */  
compute_with(x);
```

# A RACY PROGRAM

Thread 1

```
x = 42;
```

```
notify = 1;
```

Thread 2

```
while 0 == notify { /* wait */ }
```

```
compute_with(x);
```

Thread 1


```
r1 = 42;
```

```
notify = 1; x = r1;
```


Thread 2

```
r2 = notify; while 0 == r2 { /* wait */ }
```

```
compute_with(x);
```




compiler or processor



Compiler and processor would like to start loads earlier in order to hide memory latency. We'll call that *prefetch*.


Compiler *and* processor would like to complete stores later in order to hide memory latency. We'll call that *write behind*.

$B[i] = \dots$



---

COMPUTE | STORE | ANALYZE



# REMEMBER THE RACY PROGRAM?



Thread 1

```
x = 42;
```

```
notify = true;
```

Thread 2

```
while 0 == notify { /* wait */ }
```

```
compute_with(x);
```

Thread 1


```
r1 = 42;
```

```
notify = 1; x = r1;
```


Thread 2


```
r2 = notify; while 0 == r2 { /* wait */ }
```

```
compute_with(x);
```




compiler or processor





# COMMUNICATION OPTIMIZATION





# FIXING IT WITH A CACHE



# CACHE FOR REMOTE DATA

- **Goal: communication aggregation and overlap**
- **Bonus points: avoiding repeated communication**
- **Software cache in Chapel's runtime**
- **One cache per pthread**
- **Write-back cache with dirty bits**

# CACHE COHERENCY

- Simple, local coherency
- Discard all cached data on *acquire*
- Wait for pending operations on a *release*
- Strategy used in related work with UPC



# CACHE FEATURES

	<i>Overlap</i>		<i>Aggregation</i>	
	<i>GET</i>	<i>PUT</i>	<i>GET</i>	<i>PUT</i>
<i>Do PUTs in background</i>		X		
<i>Start one PUT per contiguous written region</i>				X
<i>Round GETs up to 64-byte cache lines</i>			X	
<i>Sequential read-ahead</i>	X		X	
<i>Programmer-provided prefetch hints*</i>	X			



# OTHER APPROACHES

A historical color photograph showing three men in military uniforms, likely World War I era, gathered around a mechanical device, possibly an engine or a piece of heavy machinery. One man in the foreground wears a dark flight helmet and goggles, looking down at the work. Another man to his right wears a dark helmet and a light-colored cap, also focused on the task. A third man is partially visible behind them. They are all wearing dark, worn uniforms. The background is dark and out of focus.

# WEAK MEMORY CONSISTENCY?

```
1 x starts at 0;  
...  
if someOption then  
2   x = 2;  
  if someOtherOption then  
3    x = 3;  
4 return x;
```

# WEAK MEMORY CONSISTENCY?

```
1 x starts at 0;  
...  
...  
2 PUT 2 into x;  
...  
3 PUT 3 into x;  
4 GET x;
```

Chapel

**result must be 3**

OpenSHMEM

**result could be 0, 2, or 3**

# COMPILER OPTIMIZATION?

```
for i in 1..100
{
    // PUT into B
    B[f(i)] = i;
}
```

Can the compiler  
prove these PUTs  
do not overlap?

**PUT 1** into  $B[f(1)]$

[

**PUT 2** into  $B[f(2)]$

[

**PUT 3** into  $B[f(3)]$

[

# COMPILER OPTIMIZATION?

```
for i in 1..100
{
    // PUT into B
    B[f(i)] = i;
}
```

PUT 1 into B[f(1)]  
I    I  
PUT 2 into B[f(2)]  
I    I  
PUT 2 into B[f(2)]

With a cache,  
conflicting access is  
handled at runtime.

# OVERLAPPING GETS

```
var A:[1..n] int;  
on Locales[1] {  
    var sum:int;  
    for i in 1..n do  
        sum += A[f(i)]  
}
```

We would like to overlap the GETs for  $A[f(i)]$  with each other

```
var A:[1..n] int;
on Locales[1] {
    var sum:int;
    var h: [0..k] handles;
    var bufs: [0..k] int;
    // Warm up loop
    for i in 1..k {
        // nonblocking GET A[f(i)] into bufs[i%k]
        h[i%k] = get_nb(bufs[i%k], A[f(i)])
    }
    for i in 1..n {
        wait (h[i%k]);
        sum += bufs[i%k];
        if i+k<=n {
            // nonblocking GET A[f(i+k)] into bufs[(i+k)%k]
            h[(i+k)%k] = get_nb(bufs[(i+k)%k], A[f(i+k)])
        }
    }
}
```

Explicit overlap is messy!

```
var A:[1..n] int;
on Locales[1] {
    var sum:int;
    // Optional warm up
    for i in 1..k do prefetch(A[f(i)]);
    for i in 1..n {
        if i+k <= n then prefetch(A[f(i+k)]);
        sum += A[f(i)]
    }
}
```

Much better!



# COMMUNICATION AGGREGATION

```
for i in 1..n do  
    B[i] = compute(i);
```

```
var localB:[1..n] int;  
for i in 1..n do  
    localB[i] = compute(i);  
B = localB;
```

```
for i in 1..n do  
    consume(A[i]);
```

```
var localA:[1..n] int = A;  
for i in 1..n do  
    consume(localA[i]);
```

**Simple, cache aggregates**

**Manual optimization reduces portability**

# PERFORMANCE




San Diego Air and Space Museum


# TEST CONFIGURATIONS

- Cray XC30™ system with 50 nodes, Aries network
  - GASNet Aries: GASNet with the aries conduit
  - Cray uGNI: native uGNI support for Chapel
- Cray CS400™ system with 200 nodes, FDR InfiniBand
  - GASNet ibv: GASNet with the InfiniBand Verbs conduit
- v1.9+ is revision 5ba6639
- v1.11+ is revision 6c635a1

# SYNTHETIC BENCHMARKS




# APPLICATION BENCHMARKS




# PREFETCH EXAMPLE

```
var A:[1..n] int;
on Locales[1] {
    var sum:int;
    // Optional warm up
    for i in 1..k do prefetch(A[f(i)]);
    for i in 1..n {
        if i+k <= n then prefetch(A[f(i+k)]);
        sum += A[f(i)]
    }
}
```

# PREFETCH EXAMPLE




# VS OPTIMIZATION



\*for GASNet/Aries, Iulesh improved 3.2x between v1.9+ and v1.11+

# VS OPTIMIZATION

\* with GASNet Aries v1.11+ configuration



**Cache for Remote Data:  
providing communication  
overlap and aggregation  
since Chapel v 1.10!**



# Legal Disclaimer

*Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.*

*Cray Inc. may make changes to specifications and product descriptions at any time, without notice.*

*All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.*

*Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.*

*Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.*

*Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.*

*The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, URIKA, and YARCDATA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.*

*Copyright 2014 Cray Inc.*





**CRAY**  
THE SUPERCOMPUTER COMPANY


<http://chapel.cray.com>

[chapel\\_info@cray.com](mailto:chapel_info@cray.com)

<https://github.com/chapel-lang/chapel/>

# Backup Slides

# APPLICATION BENCHMARKS



# ADDING IMPLIED FENCES

- ***acquire* and *release* triggered by task or on statement spawn, join, start, and finish**

```
release  
on {  
    acquire  
    ...  
    release  
}  
acquire  
}  
} acquire
```

```
sync {  
    release  
begin {  
    acquire  
    ....  
    release  
}  
    }  
} acquire
```

# LOOKING INSIDE



# CACHE ENTRY

- node
- address
- readahead trigger
- min sequence number
- max put sequence number
- max prefetch sequence number

64 byte cache lines

COMPUTE

Valid Line Bits

STORE



ANALYZE

Optional Dirty Bits

1024 byte cache page





# Pointer Tree



COMPUTE | STORE | ANALYZE

Inspired by “Two Level Tree Structure for Fast Pointer Lookup” by Hans J Boehm

# CACHE DATA STRUCTURES




Pointer Tree



Dirty LRU




Free Lists



Operations Queue COMPUTE



per task:  
last acquire sequence number



# WRITE BEHIND



Write Recorded in Dirty Bits, Page added to Dirty Queue








Flushed on release or  
when there are too many dirty pages

# READAHEAD




ra skip,len = 0



GET with 2  
earlier valid  
lines triggers  
synchronous  
readahead

ra skip=1 pg len = 1 pg





COMPUTE

I


ANALYZE




ra skip=1 pg len = 1 pg




ra skip,len=0




ra skip=1 pg len =2 pg



ra skip,len=0




ra skip=1 pg len =2 pg




GET here triggers  
more readahead



ra skip,len=0




ra skip=2 pg len =4 pg




STORE

# INFINIBAND (IB) LATENCY


\* with small 10-node cluster, QDR IB



# INFINIBAND (IB) BANDWIDTH



\* with small 10-node cluster, QDR IB



# Legal Disclaimer

*Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.*

*Cray Inc. may make changes to specifications and product descriptions at any time, without notice.*

*All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.*

*Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.*

*Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.*

*Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.*

*The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, URIKA, and YARCDATA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.*

*Copyright 2014 Cray Inc.*





**CRAY**  
THE SUPERCOMPUTER COMPANY

<http://chapel.cray.com>

[chapel\\_info@cray.com](mailto:chapel_info@cray.com)

<https://github.com/chapel-lang/chapel/>