



# Global-View Abstractions for User-Defined Reductions and Scans

PPoPP

March 29, 2006

**Steve Deitz**

**Brad Chamberlain**

**David Callahan**

**Larry Snyder**





# A Bit of Context



- ◆ DARPA **HPCS**: **H**igh **P**roductivity **C**omputing **S**ystems
  - HPLS: High Productivity Language Systems
- ◆ Increase productivity for HEC community by 2010

Productivity = Programmability +  
Performance +  
Portability +  
Robustness

- ◆ Revolutionary results (not evolutionary)
- ◆ Marketable to people other than program sponsors
- ◆ Phase II competition: Cray, IBM, and Sun

- ◆ *Chapel*: Cascade High-Productivity Language
- ◆ Overall goal:
  - Simplify the creation of parallel programs
  - Support evolution to production-grade codes
  - Emphasize generality
- ◆ Motivating Language Technologies:
  - 1) Global-view multithreaded parallel programming
  - 2) Locality-aware programming
  - 3) Object-oriented programming
  - 4) Generic programming and type inference

## ◆ Definition

The **reduce operator** takes a binary operation  $\oplus$  and a sequence of values  $(/a_1, a_2, \dots, a_n/)$  and returns the value

$$a_1 \oplus a_2 \oplus \dots \oplus a_n.$$

## ◆ Example

+ reduce  $(/1, 2, 3, 4, 5, 6, 7, 8, 9, 10/)$

→ 55

## ◆ Definition

The **scan operator** takes a binary operation  $\oplus$  and a sequence of values  $(/a_1, a_2, \dots, a_n/)$  and returns the sequence of values

$$(/a_1, a_1 \oplus a_2, \dots, a_1 \oplus a_2 \oplus \dots \oplus a_n/).$$

## ◆ Example

+ scan  $(/1, 2, 3, 4, 5, 6, 7, 8, 9, 10/)$

$\rightarrow (/1, 3, 6, 10, 15, 21, 28, 36, 45, 55/)$

- ◆ Performance of reduce and scan
  - Highly efficient implementation
    - ◆ Log-tree for reduce
    - ◆ Parallel-prefix for scan
  - Associative operator → parallel performance
  - Commutative operator → more efficient
  
- ◆ Application of reduce and scan
  - Common in scientific computing
    - ◆ Used to test for convergence in iterative algorithms
    - ◆ Used in matrix multiplication and sorting kernels
    - ◆ ~9% of MPI calls in NPB are reductions (static count)
  - Supported on today's systems

- ◆ Reduce and Scan Overview
- ◆ Local-View vs. Global-View Overview
- ◆ MPI and Chapel Abstraction for Reduce and Scan
- ◆ Application to MPI and Quantitative Results

## ◆ *Local-View Programming:*

- Programmer codes on a per-processor basis
  - ◆ Breaks distributed data structures into per-process chunks
  - ◆ Breaks work into per-process iterations/control flow

## ◆ *Global-View Programming:*


- Programmer codes independent of processors (mostly)
  - ◆ Relies on compiler to manage decomposition details
  - ◆ Provides guide to decomposition at a high level



◆ **Example:** “Apply a stencil to a vector”

**global-view**

$$\begin{array}{l}
 ( \text{[blue, blue, blue, blue, light blue, light blue]} \\
 + \text{[light blue, light blue, blue, blue, blue, blue]} ) / 2 \\
 = \text{[yellow, orange, orange, orange, orange, yellow]}
 \end{array}$$



$$\begin{array}{l}
 ( \text{[blue, light blue, light blue]} \\
 + \text{[light blue, light blue, blue]} ) / 2 \\
 = \text{[yellow, orange]}
 \end{array}
 \quad
 \begin{array}{l}
 ( \text{[blue, blue, light blue, light blue]} \\
 + \text{[light blue, light blue, blue, blue]} ) / 2 \\
 = \text{[orange, orange]}
 \end{array}
 \quad
 \begin{array}{l}
 ( \text{[blue, light blue, light blue]} \\
 + \text{[light blue, light blue, blue]} ) / 2 \\
 = \text{[orange, yellow]}
 \end{array}$$

## ◆ Example: “Apply a stencil to a vector”

### global-view

```
var n: int = 1000;
var a, b: [1..n] float;

forall i in 2..n-1 do
  b(i) = (a(i-1) + a(i+1))/2;
```

### local-view

```
var n: int = 1000;
var locN: int = n/numProcs;
var a, b: [0..locN+1] float;
var innerLo: int = 1;
var innerHi: int = locN;

if (iHaveRightNeighbor) {
  send(right, a(locN));
  recv(right, a(locN+1));
} else
  innerHi = locN-1;
if (iHaveLeftNeighbor) {
  send(left, a(1));
  recv(left, a(0));
} else
  innerLo = 2;
forall i in innerLo..innerHi do
  b(i) = (a(i-1) + a(i+1))/2;
```



# 3D Stencil in NAS MG



```
subroutine comm(u,n1,n2,n3,kk)
  use caf_intrinsic
  integer axis, dir, n1, n2, n3, k, ierr
  double precision u(n1, n2, n3)
  integer i3, i2, i1, buff_len, buff_id
  buff_id = 2 + dir
  buff_len = 0
  if( axis .eq. 1 )then
    if( dir .eq. -1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          buff_len = buff_len + 1
          buff(buff_len, buff_id) = u( 2,
            i2,i3)
        enddo
      enddo
      buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
        buff(1:buff_len, buff_id)
    else if( dir .eq. +1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          buff_len = buff_len + 1
          buff(buff_len, buff_id) = u( n1-1,
            i2,i3)
        enddo
      enddo
      buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
        buff(1:buff_len, buff_id)
    endif
  else if( axis .eq. 2 )then
    if( dir .eq. -1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          buff_len = buff_len + 1
          buff(buff_len, buff_id) = u( i1,
            i2,i3)
        enddo
      enddo
      buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
        buff(1:buff_len, buff_id)
    else if( dir .eq. +1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          buff_len = buff_len + 1
          buff(buff_len, buff_id) = u( i1-1,
            i2,i3)
        enddo
      enddo
      buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
        buff(1:buff_len, buff_id)
    endif
  endif
  if( axis .eq. 3 )then
    if( dir .eq. -1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          buff_len = buff_len + 1
          buff(buff_len, buff_id) = u( i1, i2,
            1)
        enddo
      enddo
      buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
        buff(1:buff_len, buff_id)
    else if( dir .eq. +1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          buff_len = buff_len + 1
          buff(buff_len, buff_id) = u( i1, i2,
            1)
        enddo
      enddo
      buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
        buff(1:buff_len, buff_id)
    endif
  endif
  return
end
```

```
do i3=2,n3-1
  do i1=1,n1
    buff_len = buff_len + 1
    buff(buff_len, buff_id) = u( i1,
      2,i3)
  enddo
enddo
buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
  buff(1:buff_len, buff_id)
else if( dir .eq. +1 )then
  do i3=2,n3-1
    do i1=1,n1
      buff_len = buff_len + 1
      buff(buff_len, buff_id) = u( i1, n2-
        1,i3)
    enddo
  enddo
  buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
    buff(1:buff_len, buff_id)
endif
endif
if( axis .eq. 3 )then
  if( dir .eq. -1 )then
    do i2=1,n2
      do i1=1,n1
        buff_len = buff_len + 1
        buff(buff_len, buff_id) = u(
          i1,i2,2)
      enddo
    enddo
    buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
      buff(1:buff_len, buff_id)
  else if( dir .eq. +1 )then
    do i2=1,n2
      do i1=1,n1
        buff_len = buff_len + 1
        buff(buff_len, buff_id) = u(
          i1,i2,n3-1)
      enddo
    enddo
    buff(1:buff_len, buff_id+1)[nbr(axis, dir, k)] =
      buff(1:buff_len, buff_id)
  endif
endif
return
end
subroutine take3( axis, dir, u, n1, n2, n3 )
  use caf_intrinsic
  implicit none
  include 'cafnpb.h'
  include 'globals.h'
  integer axis, dir, n1, n2, n3
  double precision u( n1, n2, n3 )
  integer buff_id, indx
  integer i3, i2, i1
  buff_id = 3 + dir
  indx = 0
  if( axis .eq. 1 )then
    if( dir .eq. -1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          indx = indx + 1

```

```
          u(n1,i2,i3) = buff(indx, buff_id )
        enddo
      enddo
    else if( dir .eq. +1 )then
      do i3=2,n3-1
        do i2=2,n2-1
          indx = indx + 1
          u(i1,i2,i3) = buff(indx, buff_id )
        enddo
      enddo
    endif
  endif
  if( axis .eq. 2 )then
    if( dir .eq. -1 )then
      do i3=2,n3-1
        do i1=1,n1
          indx = indx + 1
          u(i1,n2,i3) = buff(indx, buff_id )
        enddo
      enddo
    else if( dir .eq. +1 )then
      do i3=2,n3-1
        do i1=1,n1
          indx = indx + 1
          u(i1,n2,i3) = buff(indx, buff_id )
        enddo
      enddo
    endif
  endif
  if( axis .eq. 3 )then
    if( dir .eq. -1 )then
      do i2=1,n2
        do i1=1,n1
          indx = indx + 1
          u(i1,i2,n3) = buff(indx, buff_id )
        enddo
      enddo
    else if( dir .eq. +1 )then
      do i2=1,n2
        do i1=1,n1
          indx = indx + 1
          u(i1,i2,1) = buff(indx, buff_id )
        enddo
      enddo
    endif
  endif
  return
end
subroutine commp( axis, u, n1, n2, n3, kk )
  use caf_intrinsic
  implicit none
  include 'cafnpb.h'
  include 'globals.h'
  integer axis, dir, n1, n2, n3
  double precision u( n1, n2, n3 )
  integer i3, i2, i1, buff_len, buff_id
  integer i, kk, indx
  dir = -1
  buff_id = 3 + dir
  buff_len = nm2

```

```
do i=1,nm2
  buff(1, buff_id) = 0.000
enddo
dir = +1
buff_id = 3 + dir
buff_len = nm2
do i=1,nm2
  buff(1, buff_id) = 0.000
enddo
dir = +1
buff_id = 2 + dir
buff_len = 0
if( axis .eq. 1 )then
  do i3=2,n3-1
    do i2=2,n2-1
      buff_len = buff_len + 1
      buff(buff_len, buff_id) = u( n1-1,
        i2,i3)
    enddo
  enddo
endif
if( axis .eq. 2 )then
  do i3=2,n3-1
    do i1=1,n1
      buff_len = buff_len + 1
      buff(buff_len, buff_id) = u( i1, n2-
        1,i3)
    enddo
  enddo
endif
if( axis .eq. 3 )then
  do i3=2,n3-1
    do i2=2,n2-1
      buff_len = buff_len + 1
      buff(buff_len, buff_id) = u( i1, i2, n3-
        1)
    enddo
  enddo
endif
dir = -1
buff_id = 2 + dir
buff_len = 0
if( axis .eq. 1 )then
  do i3=2,n3-1
    do i2=2,n2-1
      buff_len = buff_len + 1
      buff(buff_len, buff_id) = u( 2, i2,i3)
    enddo
  enddo
endif
if( axis .eq. 2 )then
  do i3=2,n3-1
    do i1=1,n1
      buff_len = buff_len + 1
      buff(buff_len, buff_id) = u( i1,
        2,i3)
    enddo
  enddo
endif
if( axis .eq. 3 )then
  do i3=2,n3-1
    do i2=2,n2-1
      buff_len = buff_len + 1
      buff(buff_len, buff_id) = u( i1, i2, 2)
    enddo
  enddo
endif
do i=1,nm2
  buff(1,4) = buff(i,3)
  buff(1,2) = buff(i,1)
enddo
dir = -1

```

```
buff_id = 3 + dir
indx = 0
if( axis .eq. 1 )then
  do i3=2,n3-1
    do i2=2,n2-1
      indx = indx + 1
      u(n1,i2,i3) = buff(indx, buff_id )
    enddo
  enddo
endif
if( axis .eq. 2 )then
  do i3=2,n3-1
    do i1=1,n1
      indx = indx + 1
      u(i1,n2,i3) = buff(indx, buff_id )
    enddo
  enddo
endif
if( axis .eq. 3 )then
  do i3=2,n3-1
    do i1=1,n1
      indx = indx + 1
      u(i1,i2,n3) = buff(indx, buff_id )
    enddo
  enddo
endif
if( axis .eq. 1 )then
  do i3=2,n3-1
    do i2=2,n2-1
      indx = indx + 1
      u(i1,i2,i3) = buff(indx, buff_id )
    enddo
  enddo
endif
if( axis .eq. 2 )then
  do i3=2,n3-1
    do i1=1,n1
      indx = indx + 1
      u(i1,i2,i3) = buff(indx, buff_id )
    enddo
  enddo
endif
if( axis .eq. 3 )then
  do i3=2,n3-1
    do i2=2,n2-1
      indx = indx + 1
      u(i1,i2,1) = buff(indx, buff_id )
    enddo
  enddo
endif
return
end

```

- ◆ Reduce and Scan Overview
- ◆ Local-View vs. Global-View Overview
- ◆ MPI and Chapel Abstraction for Reduce and Scan
- ◆ Application to MPI and Quantitative Results

Compute the k minimum values in a sequence of values

E.g.    `mink(2) reduce (/8, 3, 5, 2, 7, 9, 10, 1, 4, 6/)`

→ (1, 2)


E.g.    `mink(2) reduce (/Cray, Gonzaga, IBM, Sun, Duke/)`

→ (Duke, Gonzaga)

◆ **Example:** “Reduce a sequence with Min2”


global-view

( / 8 , 3 , 5 , 2 , 7 / )



local-view

( / 8 , 3 / ) | ( / 5 , 2 , 7 / )



Local-view abstraction  
applies only to this part  
of the reduction

- ◆ Has two varieties
  - Reduce returns reduction on root process
  - Allreduce returns result back to all processes


```
int MPI_Reduce(  
    void* sendbuf,  
    void* recvbuf,  
    int count,  
    MPI_Datatype datatype,  
    MPI_Op op,  
    int root,  
    MPI_Comm comm)
```

```
int MPI_Allreduce(  
    void* sendbuf,  
    void* recvbuf,  
    int count,  
    MPI_Datatype datatype,  
    MPI_Op op,  
    MPI_Comm comm)
```

- ◆ Supports nine built-in reduction operators
- ◆ Accepts arbitrary combine function
- ◆ MPI\_Op assumes associativity, can be commutative


◆ **Example:** “Reduce a sequence with Min2”

global-view



Global-view abstraction  
needs to hide the details  
of communication

local-view



Local-view abstraction  
applies only to this part  
of the reduction



- ◆ Similar to classes in other OO languages
  - Fields are accessed via “.” syntax
  - Methods are invoked “.” syntax
  - Support multiple inheritance, dynamic dispatch
- ◆ Automatic default constructor

```
class Point {  
    var x : int, y : int;  
}  
class ColorfulPoint : Point {  
    var c : color;  
}  
var cp = ColorfulPoint(x=2,y=3,c=red);
```



Default constructor provides  
named argument passing

## ◆ Classes can be generic

```
class Triple {  
    type elt_type;  
    var x, y, z : elt_type;  
}  
def Triple.add(e : elt_type) {  
    x += e; y += e; z += e;  
}  
var t = Triple(float);  
t.add(1.0);
```

Class can be instantiated on different types

```
class reduction {
  type elt_type;
  param commutative = true;
  param associative = true;
  def accumulate(x : elt_type);
  def combine(s : reduction);
  def gen();
  def scan_gen(x : elt_type) do
    return gen();
  def red_gen() do
    return gen();
}
```

All reductions are generic to the sequence element type and two parameters that can be overridden

Default scan and reduce generate functions call a common generate function

```
def reduce(r, s) {
  for e in s do
    r.accumulate(e);
  return r.red_gen();
}
```

Serial reduce is implemented as is any overloaded operator

```

class mink : reduction {
  var k : int;
  var v : [1..k] elt_type = max(elt_type);
  def accumulate(x : elt_type) do
    if x < v(1) {
      v(1) = x;
      for i in 2..k do
        if v(i-1) < v(i) then
          (v(i-1),v(i)) = (v(i),(v(i-1)));
      }
  def combine(s: mink(elt_type)) do
    for i in 1..k do
      accumulate(s.v(i));
  def gen() do
    return v;
}

```

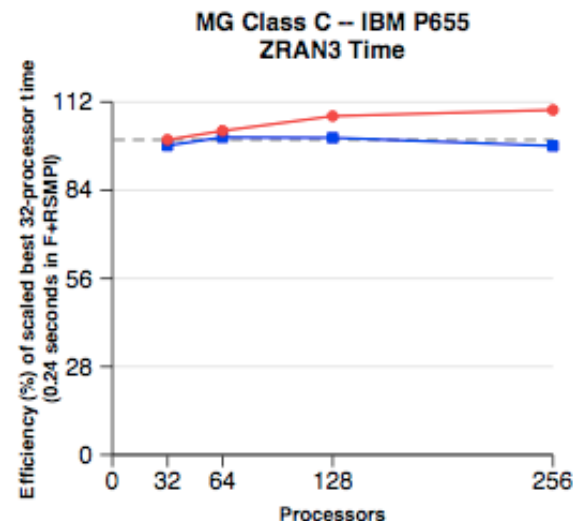
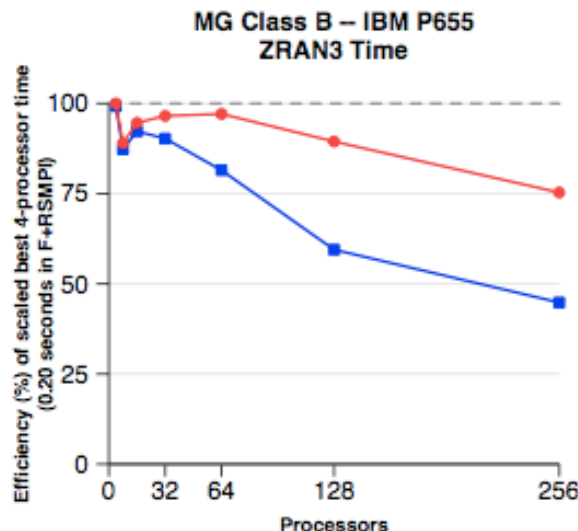

Default constructor stores  
maxes in k-element array.  
Operator class stores state.

Combine function takes  
other thread's locally  
accumulated state.

```
... mink(2) reduce (/8, 3, 5, 2, 7/) ...
```

- ◆ Reduce and Scan Overview
- ◆ Local-View vs. Global-View Overview
- ◆ MPI and Chapel Abstraction for Reduce and Scan
- ◆ Application to MPI and Quantitative Results


- ◆ RSMPI (Reduce and Scan in MPI)
  - Proof of concept of global view
  - Potential use for MPI programmers
- ◆ Minimal changes to MPI to support global view
  - Adds a construct to iterate over the local part of an ordered set
  - Uses a preprocessor to generate the accumulate loop based on the iterator construct.
  - Adds an operator as a collection of functions
    - ◆ Identity function
    - ◆ Accumulate function
    - ◆ Combine function
    - ◆ Generate function
  - Requires a single RSMPI\_Reduce call



Reduction: Find 10 minimum and 10 maximum values and locations

-- linear speedup  
 ■ F+MPI  
 ● F+RSMPI

F+MPI: Uses 10 min and 10 max reductions  
 F+RSMPI: Uses one big reduction



Reduction: Determine if array is sorted

C+MPI: Compare boundary elements,  
determine if local portion is sorted, and  
'sum' reduce local determinations

C+RSMPI: Use one big reduction



- ◆ High-level Chapel abstraction
  - Takes advantage of generic OOP capabilities
  - Contains both accumulate and combine function
  - Uses default constructor to produce identity
  - Supports generate function to return result that is different from the combining state
  - Supports a scan generate function that can return different result from a reduction
- ◆ RSMPI
  - Shows performance does not suffer at global view
  - Suggests path for incremental improvements
- ◆ Global-View
  - Orthogonal to changes in the distribution/layout

For more information:

<http://chapel.cs.washington.edu>

email: [chapel\\_info@cray.com](mailto:chapel_info@cray.com)