

Performance Optimizations

Chapel Team, Cray Inc.
Chapel version 1.10
October 2nd, 2014



COMPUTE | STORE | ANALYZE

Safe Harbor Statement



This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

2

Executive Summary



- **Performance optimizations did not receive a large effort in this release cycle**
 - Some changes that we expected to benefit performance did not
 - e.g., improvements to generated code (see later deck)
 - Yet, a few focused changes had a significant impact in key cases



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

3

Outline

- Loop-Invariant Code Motion (LICM) Improvements
- CHPL TARGET ARCH and --specialize
- Other Performance Optimizations



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

Loop-Invariant Code Motion (LICM) Improvements



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

LICM Improvements

CRAY

Background:

- LICM hoists invariant code out of the body of loops
 - improves execution performance
- ```
for i in 1..10 { var t = 10 * someConst; } // can hoist t
```
- LICM was only effective for `--local` compilations
  - did not hoist wide variables, ineffective for
    - `--no-local`
    - `multi-locale`
    - `numa`

### This Effort:

- Make LICM equally effective for all configurations
  - same code hoisted for `--local`, `--no-local`, `multi-locale`, `numa`



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

6

Checks in loop invariant code motion accidentally prevented wide variables from being hoisted, which meant LICM was largely ineffective for everything except for `--local`.

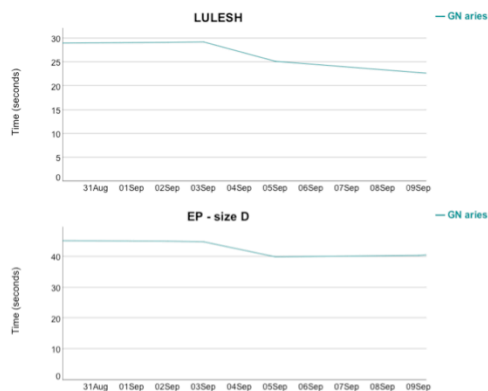
To fix this, the check that prevented hoisting of wide variables was removed, and two latent bugs were discovered and fixed.

## LICM Improvements: Impact

CRAY

- **Multi-locale performance impact**

- substantial communication count improvements
- performance improvements



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

7

There were some substantial communication (comm) count improvements. Most notably tests in test/performance/ferguson. These test the comm counts for some basic chapel communication idioms:

```
remote-array-read : #gets 200011 => 100012
remote-array-write : #gets 100011 => 12
remote-class-read : #gets 700011 => 600012
remote-class-write : #gets 400011 => 300012
remote-record-read : #gets 400011 => 300012
remote-record-write : #gets 100011 => 12
remote-tuple-read : #gets 400011 => 300012
remote-tuple-write : #gets 100011 => 12
```

There were a few other improvements in real tests such as SSCA

```
ssca2 kernel #3 : #gets 239 => 158
```


There would have been more comm count improvements but LICM is not very effective with bounds-checks on, and tests in test/performance/ferguson, SSCA, and the other improvements happen to throw -fast

## LICM Improvements: Impact (continued)



- **Single locale performance impact**

- some variations of Fannkuch took a performance hit
  - result of a latent bug that was discovered and fixed as part of this work
  - did not affect the fastest versions



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

8



## LICM Improvements: Next Steps

CRAY

- **Fix Fannkuch performance regression**
  - source of regression has been identified
  - fix would have come too late in the release cycle
    - was not a high priority since fastest versions were unaffected



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

9

## CHPL\_TARGET\_ARCH and --specialize



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

## CHPL\_TARGET\_ARCH : Background



- **C compilers can optimize for the target architecture**
  - We needed some way to expose this in Chapel
- **Some operations based on intrinsics are slow**
  - By default, C compilers won't emit instructions added by recent ISAs
- **A target architecture is required to get good vectorization**
  - No AVX or SSE newer than SSE2 without specifying the architecture



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

11

## CHPL\_TARGET\_ARCH : This Effort



- **New environment variable**

- CHPL\_TARGET\_ARCH
- Indicates what architecture to use for the runtime and executable
- Uses standardized names across the backend C compilers
  - Not all compilers support all architectures
- With Cray PrgEnv-\* the value is from the loaded craype-\* module

- **New compiler flag**

- --specialize
- Indicates that specialization should be performed
  - -march (GCC/Clang)
  - ...
- Implied by --fast by default



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

12


See README.chplenv for more details on the supported options for CHPL\_TARGET\_ARCH

## CHPL\_TARGET\_ARCH: Impact



- **popcount**

- ~7x speedup for 32 bit
- ~10x speedup for 64 bit



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

13


GCC's `__builtin_popcount` is slower than the straight C version without allowing specialization.

## CHPL\_TARGET\_ARCH: Impact

CRAY

- **array vs ddata serial accesses**

- ~1.4x slowdown for ddata
- Not clear why this happens



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

14

Little to no impact on other tests.

## CHPL\_TARGET\_ARCH: Next Steps



- Investigate why the timings went up for ddata



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

15

## Other Performance Optimizations



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.



## Other Performance Optimizations



- **Switch to Qthreads tasking layer (see 'Runtime' slides)**
  - and switch in thread counts from logical to physical cores by default
- **Improved the performance of the 1D array serial iterator**
  - previously, we were using a more expensive (strided) idiom
- **Made task counters use network atomics when available**
  - previously, they used processor atomics and active messages
- **Improved the performance of readstring()**
- **Reduced communication counts due to other refactorings**



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

17

## Overall Performance Optimization Priorities/Next Steps



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

## Overall Perf. Opt. Priorities/Next Steps



- **Focus on SIMD-ization of generated code**
- **Reduce amount of unnecessary communication code**
  - Relates to earlier item about --local vs. --no-local
- **Add support for standalone parallel iterators**
- **Optimize reductions**
- **Improve LICM for cases like Fannkuch that backslid**



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

19

## Legal Disclaimer



*Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.*

*Cray Inc. may make changes to specifications and product descriptions at any time, without notice.*

*All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.*

*Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.*

*Cray uses codenames internally to identify products that are in development and not yet publicly announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.*

*Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.*

*The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, URIKA, and YARCDATA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.*

*Copyright 2014 Cray Inc.*



COMPUTE | STORE | ANALYZE

Copyright 2014 Cray Inc.

20



**CRAY**  
THE SUPERCOMPUTER COMPANY

<http://chapel.cray.com>

[chapel\\_info@cray.com](mailto:chapel_info@cray.com)

<http://sourceforge.net/projects/chapel/>