

Library Improvements

Chapel version 1.20
September 19, 2019

- ✉ chapel_info@cray.com
- 🌐 chapel-lang.org
- 🐦 @ChapelLanguage



Outline

- New Modules
 - Collection Types
 - UnitTest
 - EpochManager
- Module Improvements
 - Sort
 - Reflection
 - LinearAlgebra



New Modules

- Collection Types
- UnitTest
- EpochManager



Collection Types

List, Map and Set



Collections: Background

- List, set, and map operations had been supported by arrays and domains

// Append to a "list"

```
var D1 = {1..0};  
  
var A: [D1] real;  
  
A.push_back(1.1);
```

// Add to a "map"

```
var D2: domain(string);  
  
var B: [D2] int;  
  
B["hello"] = 42;
```

Collections: Background

- Using arrays as collection types had some problems
 - Required (via assertion) that the array and its domain had a 1:1 relationship
 - resulted in surprising instabilities for users once domains were shared

// Append to a "list"

```
var D1 = {1..0};  
  
var A: [D1] real;  
  
A.push_back(1.1); // OK  
  
var C = A;  
  
A.push_back(2.2); // error
```

// Add to a "map"

```
var D2: domain(string);  
  
var B: [D2] int;  
  
B["hello"] = 42; // OK  
  
var D = B;  
  
B["bye"] = 33; // error
```

Collections: Background

- Using arrays as collection types had some problems
 - Permitted modifying ‘const’ domains

```
// Illegally modifies D1
const D1 = {1..0};

var A: [D1] real;
A.push_back(1.1);
```

```
// Illegally modifies D2
const D2: domain(string);

var B: [D2] int;
B["hello"] = 42;
```

Collections: This Effort

- Implemented new collection types
 - list – Replaces "vector-like" operations on arrays
 - map – Replaces "map-like" operations on arrays
 - set – Alternative to set operations on associative domains
- Deprecated list- and map-style operations on arrays
 - pop_back, push_back, pop_front, push_front
 - clear, insert, remove
 - |=, +=, ^=, &=

Collections: List

- The list type enables users to build up and iterate over a collection of elements

```
use List;  
  
var lst: list(int);      // Declare a list  
lst.extend(1..8);        // Extend with items from a range.  
writeln(lst);
```

- Parallel-safe operations can be used by setting "parSafe" to true at initialization

```
var lst: list(int, parSafe=true);  
  
coforall tid in 1..8 with (ref lst) do  
  
    lst.append(tid);
```

Collections: List

- Intended to replace "vector-like" array methods, which are now deprecated

```
var foo: [1..0] int;  
  
foo.push_back(4); // warning: push_back is deprecated – please use list.append  
foo.clear(); // warning: clear is deprecated – please use list.clear
```

- Can now be written as:

```
use List;  
  
var foo: list(int);  
  
foo.append(100);  
foo.clear();
```

Collections: Map

- The map type enables users to associate keys with values. It supports...
 - Assignment to any index of the key type
 - Reading from indices in the map
 - Iterators over the keys, values, or key-value pairs
 - Boolean set operators over the keys

// Boolean set operators

```
m1 |= m2;  
m1 &= m2;  
m1 ^= m2;
```

// Declarations and basic usage

```
use Map;  
var m1 = new Map(string, int);  
var m2 = new Map(string, int);  
m1["one"] = 1;  
m2["two"] = 2;
```

// Iterators

```
for k in m1 {...} // keys  
for v in m1.values() {...} // values  
for kv in m1.items() {...} // both
```

Collections: Map

- Intended to replace "map-like" methods on arrays, which are now deprecated

```
var D: domain(string);  
  
var A: [D] int;  
  
A["hello"] = 42; // warning: growing associative domains by assigning  
// to an array is deprecated
```

- Can now be written as:

```
var m = new map(string, int);  
  
m["hello"] = 42;
```

Collections: Set

- Offers a lightweight alternative to domains
- Similar to both list and map, set operations can be made parallel-safe

```
use Set;

var s1: set(int, parSafe=true);

coforall i in 1..8 with (ref s1) do
    s1.add(i);
writeln(s1 ^ s1);
```

Collections: Impact, Status

Impact:

- New list, set, and map data structures added to standard modules
- Vector-like and map-like methods on arrays are deprecated

Status:

- All three types have testing coverage
- List and map are used in the Mason, UnitTest, and TOML packages



Collections: Next Steps

Next Steps:

- Collect user feedback
- Investigate the performance of list and map
- Explore potential optimizations
- Make sure parallel safety strategies are appropriate
- Consider alternative underlying implementations

UnitTest



UnitTest: Background

- Existing options for writing Chapel user code tests are problematic
 - `start_test`
 - Not ideal for writing unit tests
 - Not a user-facing test feature, not supported by mason
 - `mason test`
 - Very limited capabilities by design
 - Only uses exit code of program to determine pass/fail status



UnitTest: This Effort

- Introduced a ‘UnitTest’ package module for writing unit tests in Chapel
- Implemented as a Google Summer of Code project

Student	Mentors
Krishna Kumar Dey	Ben Albrecht, Lydia Duncan, Sam Partee

- Functions are designated as tests through their signature
 - Test functions must take a ‘borrowed Test’, throw, and return nothing:

```
proc someTest (t: borrowed Test) throws
```

UnitTest: This Effort

- ‘mason test’ acts as the test runner both inside and outside of mason packages
 - Inside mason packages, runs tests in ‘test/’ directory
 - Outside mason packages, finds tests in current directory, unless path given
- Test runner compiles and runs tests
 - This allows resuming test suite after a halt

UnitTest: Test Assertions

- The ‘Test’ type provides JUnit-style assertion methods

```
proc testFoo(t: borrowed Test) throws {  
    t.assertTrue(FileSystem.isFile('foo'));  
    t.assertFalse(FileSystem.isDir('foo'));  
  
    t.assertEqual(here.id, 0);  
    t.assertNotEqual(here.id, 1);  
  
    t.assertGreaterThan(10, 1);  
    t.assertLessThan(1, 10);  
}
```



UnitTest: Example Usage

```
use UnitTest;

proc testUpper(test: borrowed Test) throws {
    test.assertEqual('foo'.toUpperCase(), 'FOO');
}

UnitTest.main();
```

```
> mason test example.chpl
```

```
-----
Ran 1 test in 11.3092 seconds
```

```
OK (passed = 1 )
```



UnitTest: Skipping Tests

- Metadata is attached to test functions via ‘Test’ methods
 - Tests can be skipped with or without a condition:

```
proc skip(reason: string) throws
```

```
proc skipIf(condition: bool, reason: string) throws
```

UnitTest: Skipped Test Example

```
proc testFoo(t: borrowed Test) throws {
    t.skipIf(CHPL_LLVM == 'none', "This test requires LLVM");
    ...
}

proc testFoo(t: borrowed Test) throws {
    t.skip( "This is not yet expected to work");
    ...
}
```



UnitTest: Test Dependencies

- Cross-test dependencies can be specified:
 - Functions are specified as first-class functions

```
proc dependsOn(tests: argType ...?n) throws
```



UnitTest: Test Dependency Example

```
use FileSystem, UnitTest;

proc testMoveDir(t: borrowed Test) throws {
    // testMkdir must be run before testMoveDir is run
    t.dependsOn(testMkdir);
    moveDir('foo', 'bar');
    t.assertTrue(isDir('bar'));
}

proc testMkdir(t: borrowed Test) throws {
    mkdir('foo');
    t.assertTrue(isDir('foo'));
}
```



UnitTest: Multilocale Tests

- Tests can specify how many locales they require

- Tests can list discrete numbers of locales supported:

```
proc addNumLocales(locale: int ...?n) throws
```

- Tests can provide a range of locales supported:

```
proc maxLocales(value: int) throws
```

```
proc minLocales(value: int) throws
```

- The test runner makes multiple passes with different numbers of locales ...
... such that each tests is run with a number of locales that it requires
- Running one test with N numLocales values requires wrapping it with N tests

UnitTest: Multilocale Test Example 1

```
proc foo(t: borrowed Test) throws {
    t.addNumLocales(4);
    // Will run with 4 locales only
}
```

```
proc bar(t: borrowed Test) throws {
    t.minLocales(8);
    t.maxLocales(16);
    // Will run with 16 locales only (defaults to max of locale bounds)
}
```

UnitTest: Multilocale Test Example 2

```
proc foo(t: borrowed Test) throws {
    t.addNumLocales(4, 8);
    // Will run with 8 locales only
}
```

```
proc bar(t: borrowed Test) throws {
    t.minLocales(8);
    t.maxLocales(16);
    // Will run with 8 locales only
}
```

UnitTest: Multilocale Test Example 3

```
proc foo() {  
    // Will run with 4 and 8 locales, because called from foo4 and foo8  
}  
  
proc foo4(t: borrowed Test) throws {  
    t.addNumLocales(4);  
  
    foo();  
}  
  
proc foo8(t: borrowed Test) throws {  
    t.addNumLocales(8);  
  
    foo();  
}
```

UnitTest: Status, Next Steps

Status:

- Chapel has an official user-facing test framework
- UnitTest is available as a package module
- ‘mason test’ is the UnitTest test runner

Next Steps:

- Support test suites as modules, records, and/or classes
- Support introspective assertions, similar to pytest
 - Allows a general assert function to give useful information when failing
- Respond to user feedback

EpochManager



EpochManager: Background

- Would like Chapel to include a library of lock-free data structures:
 - After all, Chapel has an emphasis on productivity and performance
- Two missing features preventing implementation of such structures
 1. Compare-and-swap is not implemented for Chapel classes
 - Challenge: 128-bit wide pointers in multilocale configurations
 2. A general solution to the ABA problem is needed
 - Challenge: Chapel is not garbage collected
 - Challenge: Generation counters would require 256-bit compare-and-swap
 - Challenge: Hazard Pointers are not sufficiently composable

EpochManager: This Effort

This Effort: Add prototype EpochManager and related package modules

- Implemented as a Google Summer of Code project

Student	Mentors
Garvit Dewan	Louis Jenkins Michael Ferguson

- EpochManager solves the ABA problem by deferring deletion in Epochs
- AtomicObjects provides compare-and-swap on local or wide pointers
- These enabled implementation of LockFreeQueue and LockFreeStack

Next Steps:

- Migrate 128-bit CAS support to runtime
- Handle nested usage of EpochManager

Module Improvements

- Sort
- Reflection
- LinearAlgebra



Sort Improvements



Sort: Background, This Effort

Background: A parallel radix sort was added in 1.19

- Only recursive sorts were run in parallel
 - count and bucketize were serial
- Not as fast as some C++ sort implementations
- No support for distributed sorting

This Effort:

- Improve the Sort module by fixing bugs and improving performance
- Prototype more performant algorithms and distributed sorting

Sort: Implementation Improvements

- For radix sort:
 - count step is now parallel
 - now supports floating point numbers as well as 'c_string'
- Arrays with domains using ranges with 'align' can now be sorted
- isSorted() is now efficient for Block-distributed arrays

Sort: Prototype Sorts

- Developed prototype two-array radix sort and distributed sort
 - These are included as undocumented features in the Sort module
- Two-array radix sort can offer significantly better performance
 - ~5x faster than the recursive implementation
 - but using 2x the space
- Distributed radix sort is functional but has scaling problems
- Described in a [CHIUW 2019 presentation](#)

Sort: Next Steps

Next Steps: Follow up on investigations to improve Sort module

- Decide if 2x space overhead is acceptable
- Investigate in-place parallel algorithms
- Resolve performance problems with distributed sort
- Investigate other distributed sort algorithms



Reflection Improvements



Reflection: Background, This Effort

Background: There was no way to query the location of the current code

- Users desired a way to know things like file name and line number
 - e.g. for logging

This Effort: Added source location query functions to the Reflection module

```
proc getLineNumber() param : int  
proc getFileName() param : string  
proc getRoutineName() param : string  
proc getModuleName() param : string
```

Reflection: Impact, Next Steps

Impact: Chapel programs can now query source locations

```
use Reflection;  
  
writeln(fileName(), ":", lineNumber(), " ",  
       moduleName(), ".", routineName());
```

Next Steps:

- Add functionality for nested routines and modules
 - 'get[Module/Routine]Name()' only return the nearest module/routine
 - Want to be able reason about the full stack in nested cases
- Enable getting location information where a routine is called

Linear Algebra Improvements



LinearAlgebra Module: Background, This Effort

CRAY
a Hewlett Packard Enterprise company

Background: Chapel's LinearAlgebra module provides linear algebra routines

This Effort: New routines and improved support for distributed and/or sparse data

- New routines for CSR and COO matrices:

```
proc isDiag(A) : bool;  
  
proc isHermitian(A) : bool;  
  
proc isSymmetric(A) : bool;
```

LinearAlgebra Module: This Effort

- The following are implemented as part of Google Summer of Code

Student	Mentors
Alvis Wong	Engin Kayraklioglu, Ben Albrecht

- New routines for linear system solution:

```
proc lu(A);  
proc solve(A, b);
```

- Miscellaneous routines:

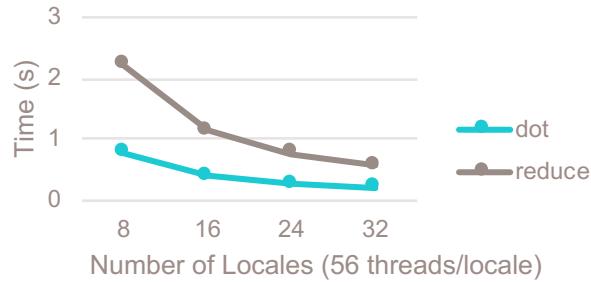
```
proc setDiag(X, offset=0, val=0);  
proc det(A);  
proc inv(A);  
proc jacobi(A, ref x, b, tol=0.0001, maxiter=1000);
```

- More support for distributed vectors/matrices: 'plus', 'minus', faster 'dot'

LinearAlgebra Module: Impact, Next Steps

Impact:

- Support for more routines, on different types of data (sparse or distributed)
- Improved dot performance



Next Steps: Continue improvements

- More support for distributed and/or sparse data
- GPU support

For More Information

For a more complete list of library-related changes in the 1.20 release, refer to the following sections of the [CHANGES.md](#) file:

- Deprecated / Removed Library Features
- Standard Library Modules
- Package Modules
- Bug Fixes



FORWARD LOOKING STATEMENTS

This presentation may contain forward-looking statements that involve risks, uncertainties and assumptions. If the risks or uncertainties ever materialize or the assumptions prove incorrect, the results of Hewlett Packard Enterprise Company and its consolidated subsidiaries ("Hewlett Packard Enterprise") may differ materially from those expressed or implied by such forward-looking statements and assumptions. All statements other than statements of historical fact are statements that could be deemed forward-looking statements, including but not limited to any statements regarding the expected benefits and costs of the transaction contemplated by this presentation; the expected timing of the completion of the transaction; the ability of HPE, its subsidiaries and Cray to complete the transaction considering the various conditions to the transaction, some of which are outside the parties' control, including those conditions related to regulatory approvals; projections of revenue, margins, expenses, net earnings, net earnings per share, cash flows, or other financial items; any statements concerning the expected development, performance, market share or competitive performance relating to products or services; any statements regarding current or future macroeconomic trends or events and the impact of those trends and events on Hewlett Packard Enterprise and its financial performance; any statements of expectation or belief; and any statements of assumptions underlying any of the foregoing. Risks, uncertainties and assumptions include the possibility that expected benefits of the transaction described in this presentation may not materialize as expected; that the transaction may not be timely completed, if at all; that, prior to the completion of the transaction, Cray's business may not perform as expected due to transaction-related uncertainty or other factors; that the parties are unable to successfully implement integration strategies; the need to address the many challenges facing Hewlett Packard Enterprise's businesses; the competitive pressures faced by Hewlett Packard Enterprise's businesses; risks associated with executing Hewlett Packard Enterprise's strategy; the impact of macroeconomic and geopolitical trends and events; the development and transition of new products and services and the enhancement of existing products and services to meet customer needs and respond to emerging technological trends; and other risks that are described in our Fiscal Year 2018 Annual Report on Form 10-K, and that are otherwise described or updated from time to time in Hewlett Packard Enterprise's other filings with the Securities and Exchange Commission, including but not limited to our subsequent Quarterly Reports on Form 10-Q. Hewlett Packard Enterprise assumes no obligation and does not intend to update these forward-looking statements.



THANK YOU

QUESTIONS?

-  chapel_info@cray.com
-  [@ChapelLanguage](https://twitter.com/ChapelLanguage)
-  chapel-lang.org



- cray.com 
- [@cray_inc](https://twitter.com/cray_inc) 
- linkedin.com/company/cray-inc-/ 