

User Application Optimizations

Chapel versions 1.21 / 1.22

April 9 / 16, 2020



chapel_info@cray.com



chapel-lang.org



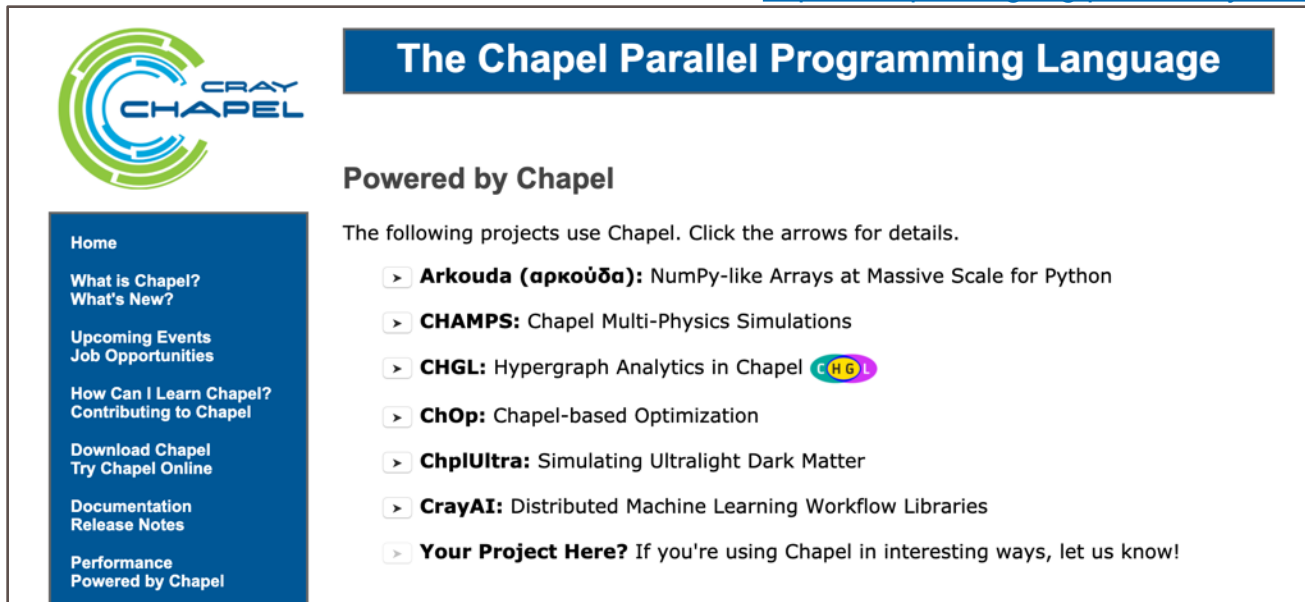
[@ChapelLanguage](https://twitter.com/ChapelLanguage)



User Applications: Background

- Chapel has seen an uptick in users with substantial applications

<https://chapel-lang.org/poweredby.html>




The screenshot shows the Cray Chapel website. On the left is a blue sidebar with navigation links: Home, What is Chapel? What's New?, Upcoming Events Job Opportunities, How Can I Learn Chapel? Contributing to Chapel, Download Chapel Try Chapel Online, Documentation Release Notes, and Performance Powered by Chapel. The main content area has a blue header 'The Chapel Parallel Programming Language'. Below it, the section 'Powered by Chapel' states 'The following projects use Chapel. Click the arrows for details.' and lists several projects with right-pointing arrows: Arkouda (αρκούδα), CHAMPS, CHGL (Hypergraph Analytics in Chapel, with a CHGL logo), ChOp, ChplUltra, CrayAI, and a link for 'Your Project Here?' with the text 'If you're using Chapel in interesting ways, let us know!'.

CRAY CHAPEL

The Chapel Parallel Programming Language

Powered by Chapel

The following projects use Chapel. Click the arrows for details.

- **Arkouda (αρκούδα):** NumPy-like Arrays at Massive Scale for Python
- **CHAMPS:** Chapel Multi-Physics Simulations
- **CHGL:** Hypergraph Analytics in Chapel 
- **ChOp:** Chapel-based Optimization
- **ChplUltra:** Simulating Ultralight Dark Matter
- **CrayAI:** Distributed Machine Learning Workflow Libraries
- **Your Project Here?** If you're using Chapel in interesting ways, let us know!

- Has resulted in optimization/tuning collaborations with the Chapel team

Outline

- [DistributedFFT \(from ChplUltra\)](#)
- [Arkouda](#)



DistributedFFT (from ChplUltra)

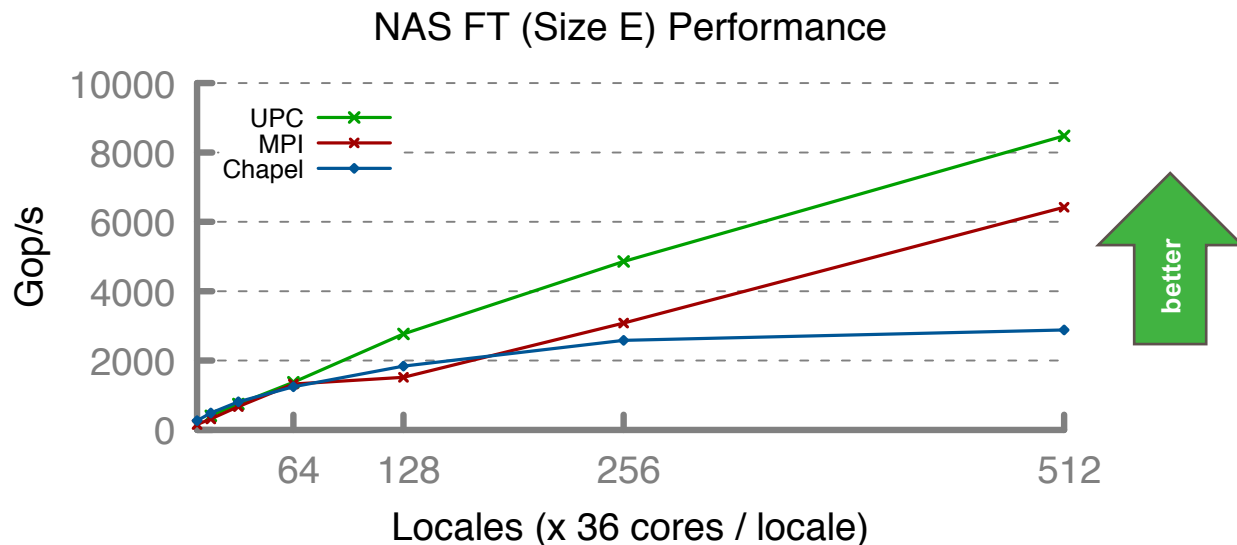


DistributedFFT: Background

- DistributedFFT provides a distributed 3D FFT
 - <https://github.com/npadmana/DistributedFFT>
- Used as the foundation for an Ultralight Dark Matter (ULDM) Simulation
 - Originally used Python/pyFFTW, but problem size soon exceeded single node
 - Chapel was chosen for distributed nature and high performance
- Chapel port uses FFTW for serial FFT computations
 - Chapel is responsible for data distribution, parallelism, and communication
 - Showcases Chapel's interoperability features

DistributedFFT: Background

- Initial Chapel port provided significant performance improvements over Python
 - Added NAS-FT benchmark to validate and compare to traditional HPC
 - Performance was decent, though lagged UPC/MPI at 64 nodes or more




DistributedFFT: This Effort

- Implemented several algorithmic optimizations
 - FFTW batching and plan caching, added comm/compute overlap
- Identified Chapel performance bottlenecks
 - Point-to-point array slice assignment has overhead
 - Only optimized for 'A = B', not 'A[locale0] = B[locale1]'
 - Local array assignment is not parallelized (uses serial 'memcpy')
- Opened Chapel issues for performance issues, worked around in the meantime
 - Using comm primitives (PUT/GET), manually parallelizing array assignment




DistributedFFT: Impact

- Significantly improved performance and scalability (3x faster at 512 nodes)
 - Performance ahead of MPI
 - On par with highly optimized non-blocking UPC



DistributedFFT: Impact

- Good performance across small, medium, and large problem sizes
 - Due to data decomposition, one-sided overlapping comm, FFTW batching



DistributedFFT: Next Steps

- Implement fixes for upstream Chapel performance issues
- Explore using cuFFT to offload FFT computations to GPUs

Arkouda



Arkouda: Background

- Arkouda provides NumPy-like arrays at HPC scale
 - A NumPy/Pandas Python interface, backed by Chapel
 - <https://github.com/mhmerrill/arkouda>
 - One of the largest open-source projects using Chapel
- Widespread adoption requires performing well on a wide variety of platforms
 - Has motivated several improvements to the Chapel compiler and runtime
- Some key operations use fine-grained communication (many 8-byte messages)
 - Cray Aries can achieve high rates for small messages
 - Other networks have much lower rates for small messages

Arkouda Aggregation: Background

- Previously, Arkouda used 'unorderedCopy()' for fine-grained messages
 - On Cray Aries, bulk copies achieve 8000 MB/s
 - unorderedCopy achieves 1000 MB/s (1/8 of bulk copy rate)
 - On 56 Gb FDR InfiniBand, bulk copies achieve 6000 MB/s
 - unorderedCopy only achieves 3 MB/s per node (1/2000 of bulk copy rate)
- In order to get performance on IB, small message aggregation is needed
 - Room for improvement on Aries as well

Arkouda Aggregation: This Effort

- Added copy aggregators to Arkouda
 - Aggregators must be created for each task
 - Have to specify type and whether source or destination is remote

```
forall i in D do
```

```
    unorderedCopy(revA[n-i], A[i]);
```

=>

```
forall i in D with (var agg = new DstAggregator(int)) do
```

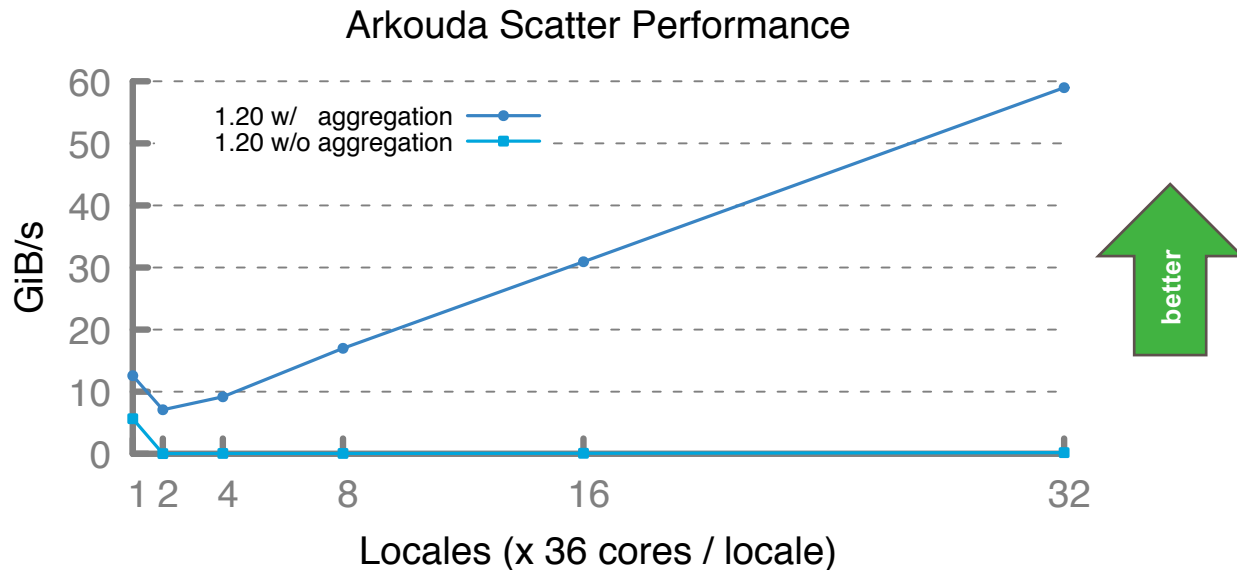
```
    agg.copy(revA[n-i], A[i]);
```

Arkouda Aggregation: Impact

- Significant performance improvements for operations with fine-grained comm
 - Following results are from:
 - 32-node Cray CS with 56 Gb InfiniBand network
 - 512-node Cray XC with Aries network
 - Per-node hardware is similar for both systems
 - 36-core Broadwell CPU
 - 128 GB RAM


Arkouda Aggregation: Cray CS Impact

- Significant performance improvements for operations with fine-grained comm
 - 300x speedup for scatters on 32 CS nodes




Arkouda Aggregation: Cray CS Impact

- Significant performance improvements for operations with fine-grained comm
 - 1200x speedup for sorting on 32 CS nodes



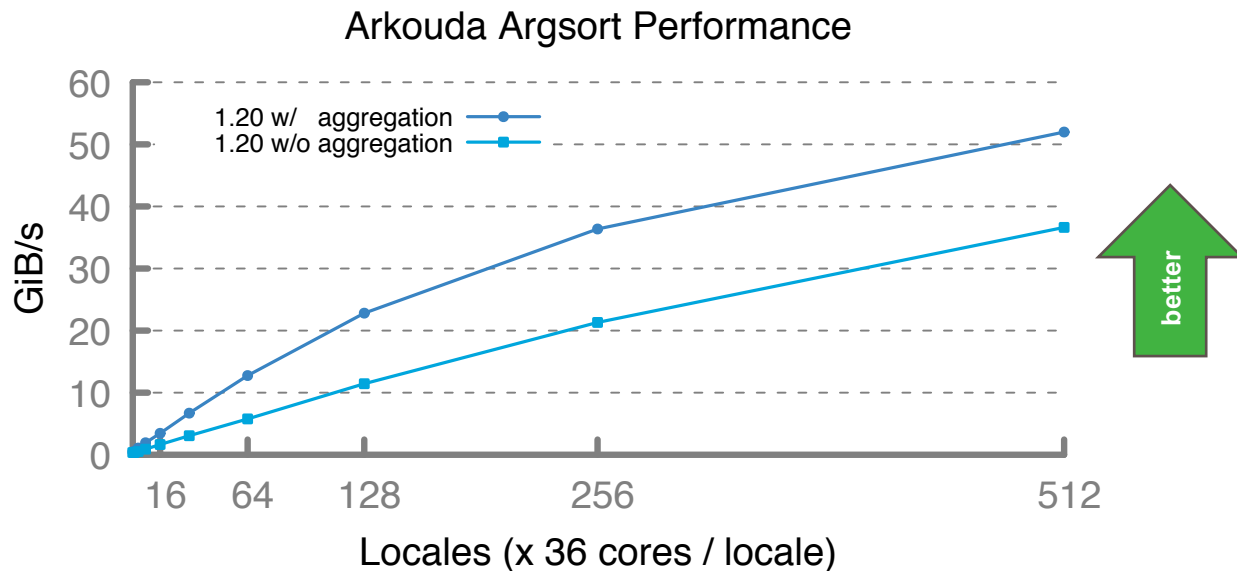
Arkouda Aggregation: Cray XC Impact

- Significant performance improvements for operations with fine-grained comm
 - 10% speedup for scatters on 512 XC nodes



Arkouda Aggregation: Cray XC Impact

- Significant performance improvements for operations with fine-grained comm
 - 40% speedup for sorting on 512 XC nodes



Arkouda Aggregation: Next Steps



- Optimize aggregation performance and reduce memory footprint
 - Current implementation is simple, lots of optimization opportunity
- Improve aggregation ease-of-use
 - Add utility functions for common idioms (gather/scatter)
 - Possibly enable with automatic unordered compiler optimization
- Add aggregation to Chapel's standard library

Arkouda Chapel Performance: This Effort

- Many 1.21 optimizations were motivated by, and benefitted, Arkouda
 - Optimizing on-statements for InfiniBand networks benefitted aggregation
 - Optimizing distributed array/domain creation benefitted most operations
 - Extending fast-followers improved array binary operations


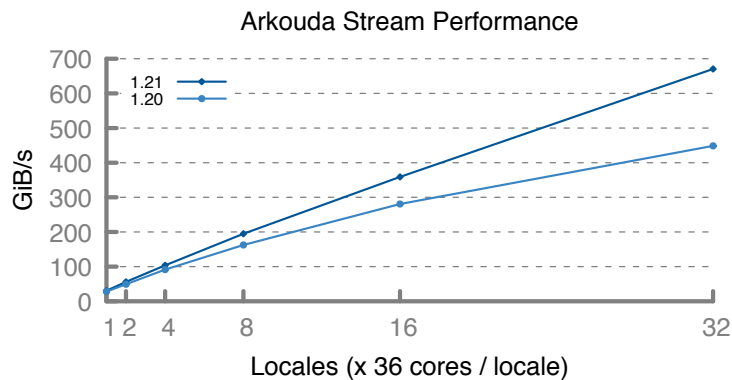
Arkouda Chapel Performance: Cray CS Impact

- Significant performance improvements on Cray CS (FDR InfiniBand)
 - At 32 locales: 40% faster Gather, 45% faster Scatter





Arkouda Chapel Performance: Cray CS Impact

- Significant performance improvements on Cray CS (FDR InfiniBand)
 - At 32 locales: 50% faster Stream, 15% faster ArgSORT





Arkouda Chapel Performance: Cray XC Impact

- Significant performance improvements on Cray XC (Aries)
 - At 512 locales: 20% faster Gather, 5% faster Scatter



Arkouda Chapel Performance: Cray XC Impact

- Significant performance improvements on Cray XC (Aries)
 - At 512 locales: 525% faster Stream, 10% faster Argsort



Arkouda: Additional Performance Improvements

- A few other Arkouda-specific optimizations were contributed upstream
 - Eliminated overhead for sorting negative integers
 - Optimized hashing used for string sorting
 - Requires extremely fast serial performance and scalable operations
 - Showcases Chapel's ability to program in the small and large

Arkouda: CI Testing

- Added continuous integration (CI) testing
 - Build and test Arkouda for each Pull Request

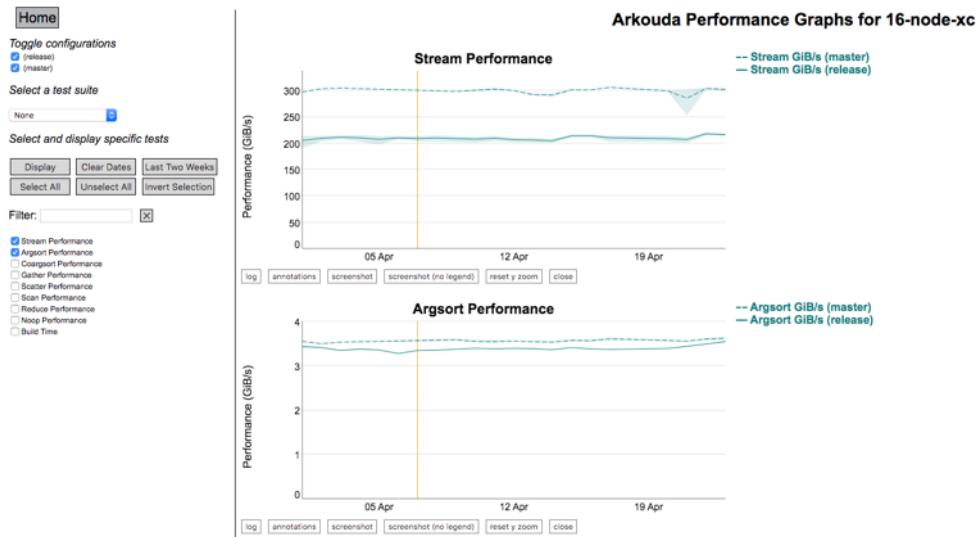
The screenshot displays a GitHub Pull Request interface for the title "Optimize coargsort for numeric arrays #330". It shows the pull request was merged 10 days ago. Below the header, there are tabs for Conversation (6), Commits (2), Checks (5), and Files changed (2). A green checkmark indicates the pull request is merged. The main section shows a list of CI jobs under the "CI" category, with "arkouda_tests_linux (chapel)" selected. The job details show a successful status with a search bar and a list of steps including "Set up job", "Initialize containers", "Run actions/checkout@v2", "Install dependencies", "Build/Install Arkouda", "Arkouda make check", "Arkouda unit tests", "Arkouda benchmark --correctness-only", "Post actions/checkout@v2", "Stop containers", and "Complete job".

Job Name	Status	Duration
lint	✓	
arkouda_tests_linux (chapel)	✓	
arkouda_tests_linux (chapel)	✓	
unit_tests_linux (chapel)	✓	
unit_tests_linux (chapel)	✓	

Step	Duration
Set up job	2s
Initialize containers	26s
Run actions/checkout@v2	2s
Install dependencies	17s
Build/Install Arkouda	12m 17s
Arkouda make check	6s
Arkouda unit tests	3s
Arkouda benchmark --correctness-only	8s
Post actions/checkout@v2	1s
Stop containers	2s
Complete job	0s


Arkouda: Performance testing

- Added benchmarking infrastructure and nightly performance testing
 - Runs single-node, 16-node CS, and 16-node XC
 - <https://chapel-lang.org/perf/arkouda/>



Arkouda: Summary

- Arkouda performance has significantly improved
 - Particularly on InfiniBand, which is now comparable to Aries



User Applications: Next Steps

- Continue to support user efforts
 - Please feel free to contact us if your Chapel code needs tuning
- Develop better tooling for profiling and performance investigation
- Improve tooling for building/testing/releasing Chapel packages and applications
 - i.e. improve Mason

FORWARD LOOKING STATEMENTS

This presentation may contain forward-looking statements that involve risks, uncertainties and assumptions. If the risks or uncertainties ever materialize or the assumptions prove incorrect, the results of Hewlett Packard Enterprise Company and its consolidated subsidiaries ("Hewlett Packard Enterprise") may differ materially from those expressed or implied by such forward-looking statements and assumptions. All statements other than statements of historical fact are statements that could be deemed forward-looking statements, including but not limited to any statements regarding the expected benefits and costs of the transaction contemplated by this presentation; the expected timing of the completion of the transaction; the ability of HPE, its subsidiaries and Cray to complete the transaction considering the various conditions to the transaction, some of which are outside the parties' control, including those conditions related to regulatory approvals; projections of revenue, margins, expenses, net earnings, net earnings per share, cash flows, or other financial items; any statements concerning the expected development, performance, market share or competitive performance relating to products or services; any statements regarding current or future macroeconomic trends or events and the impact of those trends and events on Hewlett Packard Enterprise and its financial performance; any statements of expectation or belief; and any statements of assumptions underlying any of the foregoing. Risks, uncertainties and assumptions include the possibility that expected benefits of the transaction described in this presentation may not materialize as expected; that the transaction may not be timely completed, if at all; that, prior to the completion of the transaction, Cray's business may not perform as expected due to transaction-related uncertainty or other factors; that the parties are unable to successfully implement integration strategies; the need to address the many challenges facing Hewlett Packard Enterprise's businesses; the competitive pressures faced by Hewlett Packard Enterprise's businesses; risks associated with executing Hewlett Packard Enterprise's strategy; the impact of macroeconomic and geopolitical trends and events; the development and transition of new products and services and the enhancement of existing products and services to meet customer needs and respond to emerging technological trends; and other risks that are described in our Fiscal Year 2018 Annual Report on Form 10-K, and that are otherwise described or updated from time to time in Hewlett Packard Enterprise's other filings with the Securities and Exchange Commission, including but not limited to our subsequent Quarterly Reports on Form 10-Q. Hewlett Packard Enterprise assumes no obligation and does not intend to update these forward-looking statements.



THANK YOU

QUESTIONS?



chapel_info@cray.com



[@ChapelLanguage](https://twitter.com/ChapelLanguage)



chapel-lang.org



cray.com

[@cray_inc](https://twitter.com/cray_inc)

[linkedin.com/company/cray-inc-/](https://linkedin.com/company/cray-inc/)

