

CHAPEL 1.23 RELEASE NOTES: PERFORMANCE OPTIMIZATIONS

Chapel Team
October 15, 2020

OUTLINE

- Array Optimizations
- Compilation Time Improvements
- Memory Improvements

ARRAY OPTIMIZATIONS

- [Automatic Local Access Optimization](#)
- [Improvements to Associative Types](#)
- [Array Tracking Optimization](#)
- [Constant Domain Optimization](#)
- [Parallel Array Initialization](#)
- [Parallel Array Assignment](#)
- [Array Swap Optimization](#)



AUTOMATIC LOCAL ACCESS OPTIMIZATION

AUTOMATIC LOCAL ACCESS OPTIMIZATION

Background

- Iterating over arrays/domains using 'forall' is a very common pattern in Chapel:

```
var D = newBlockDom({1..N});  
var A: [D] int;  
forall i in D do  
    A[i] = calculate(i);
```

loop is run over the domain of an array

the array is indexed using the loop index

- For distributed arrays, every 'A[i]' checks whether it is a local access
 - This check is overhead for this pattern: they are all guaranteed to be local
- Potential workarounds:

```
forall (a, i) in zip(A, A.domain) do  
    a = calculate(i);  
  
forall i in A.domain do  
    A.localAccess(i) = calculate(i);
```

clunky

AUTOMATIC LOCAL ACCESS OPTIMIZATION

This Effort

- Implemented a compiler analysis that replaces 'A[i]' with 'A.localAccess[i]'
 - The optimization is done statically if the compiler can prove that:
 - the loop domain supports the optimization
 - the array is indexed with the loop index symbol
 - the loop domain matches the array's domain
 - The optimization is subject to a dynamic check at execution time if:
 - the first two conditions above are met, but the compiler cannot prove that the loop and array domains match
- An example where the optimization can be done statically:

```
var D = newBlockDom( {1..10} );
var A: [D] int;
forall i in D do
  A[i] = calculate(i); // ==> A.localAccess[i] = calculate(i);
```



AUTOMATIC LOCAL ACCESS OPTIMIZATION

Arrays With Common Domains

- The optimization also applies to multiple arrays

```
var D = newBlockDom({1..N});  
var A: [D] int;  
var B: [D] int;  
forall i in D do  
    A[i] = calculate(B[i]);
```

array(s) indexed using the loop index

- Even when the loop domain is not explicit

```
var D = newBlockDom({1..N});  
var A: [D] int;  
var B: [D] int;  
forall i in A.domain do  
    A[i] = calculate(B[i]);
```

array(s) indexed using the loop index

loop is run over a domain query

array(s) have the same domain as the loop

AUTOMATIC LOCAL ACCESS OPTIMIZATION

Dynamic Checks

- If the compiler cannot determine the domain of an array:
 - Equality of domains will be checked at execution time
 - Depending on that, an optimized or unoptimized version of the loop will be run

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int); // currently we can't infer 'B' has the same domain as 'A'
forall i in A.domain do
    A[i] = calculate(B[i]); // B[i] is local if A.domain == B.domain
                            // that can only be confirmed at execution time
```

- Terminology
 - 'A' is a static candidate
 - 'B' is a dynamic candidate
- The compiler will clone loops if there are one or more dynamic candidates
 - This can increase compilation time



AUTOMATIC LOCAL ACCESS OPTIMIZATION

Dynamic Checks

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
param staticCheckA = canUseLocalAccess(A, A.domain);
param staticCheckB = canUseLocalAccess(B, B.domain);
if staticCheckA || staticCheckB {
    const dynamicCheckB = canUseLocalAccessDyn(B, A.domain);
    if dynamicCheckB then
        forall i in A.domain do
            A.localAccess[i] = calculate(B.localAccess[i]);
    else
        forall i in A.domain do
            A.localAccess[i] = calculate(B[i]);
} else {
    forall i in A.domain do
        A[i] = calculate(B[i]);
}
```

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
forall i in A.domain do
    A[i] = calculate(B[i]);
```

Static checks are created for both arrays

Dynamic check is created only for B

AUTOMATIC LOCAL ACCESS OPTIMIZATION

Dynamic Checks

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
param staticCheckA = canUseLocalAccess(A, A.domain);
param staticCheckB = canUseLocalAccess(B, B.domain);
if staticCheckA || staticCheckB {
    const dynamicCheckB = canUseLocalAccessDyn(B, B.domain);
    if dynamicCheckB then
        forall i in A.domain do
            A.localAccess[i] = calculate(B.localAccess[i]);
    else
        forall i in A.domain do
            A.localAccess[i] = calculate(B[i]);
} else {
    forall i in A.domain do
        A[i] = calculate(B[i]);
}
```

```
var A = newBlockArr({1..N}, int);
var B = newBlockArr({1..N}, int);
forall i in A.domain do
    A[i] = calculate(B[i]);
```

Will be executed if

- A passes static checks
- B passes static and dynamic checks

Will be executed if

- A passes static checks
- B fails static or dynamic checks

Will be executed if

- Neither array passes static checks

AUTOMATIC LOCAL ACCESS OPTIMIZATION

Dynamic Support for Subset Domains

- The optimization covers cases where the loop domain is a subset of the array domain

```
var D = newBlockDom({1..10});  
var A, B: [D] int;  
forall i in D.expand(-1) do  
  A[i] = calculate(B[i]);
```

Optimized upon a dynamic check



- It also detects iteration over (a subset of) the local subdomain of a distributed array's domain

```
var D = newBlockDom({1..10});  
var A, B: [D] int;  
coforall l in Locales do on l {  
  forall i in D.localSubdomain() do  
    A[i] = calculate(B[i]);  
  // ... or ...  
  forall i in D.localSubdomain().expand(-1) do  
    A[i] = calculate(B[i]);  
}
```

Optimized upon a dynamic check



AUTOMATIC LOCAL ACCESS OPTIMIZATION

Queried Domains in Array Formals

- Static optimization opportunities for array formals without domain queries are limited

```
proc foo(A, B) {  
    forall i in A.domain do  
        A[i] = calculate(B[i]);  
}
```

'A[i]' can be optimized statically

Currently, we can't determine whether B is an array early enough during compilation, so we use dynamic checks for it

- To avoid dynamic checks and loop cloning, be more explicit when multiple arguments share a domain

```
proc foo(A: [?D], B: [D]) {  
    forall i in A.domain do  
        A[i] = calculate(B[i]);  
}
```

We know that B is an array that has the same domain as the loop domain

AUTOMATIC LOCAL ACCESS OPTIMIZATION

Available Compiler Flags

- `--[no-]auto-local-access`
 - Enable/disable this optimization
 - Enabled by default
- `--[no-]dynamic-auto-local-access`
 - Enable/disable dynamic optimization
 - Enabled by default
 - Dynamic optimization results in loop cloning and can increase compilation time in some codes
- `--[no-]report-auto-local-access`
 - Enable/disable verbose output about the optimization steps
 - Disabled by default



AUTOMATIC LOCAL ACCESS OPTIMIZATION

Caveats

- The optimization is thwarted if
 - The locale changes between the 'forall' and the array access

```
forall i in A.domain do
    on Locales[X] do      // this statement can move the execution to another locale
        A[i] = calculate(i);
```

- The array index symbol is not identical to the loop index symbol

```
forall i in A.domain {
    const k = i;
    A[k] = calculate(i);
}
```



AUTOMATIC LOCAL ACCESS OPTIMIZATION

Caveats

- Zippered foralls are supported only if the loop index is expanded

```
forall (i,a) in zip(D, someIterator()) { } // the loop will be analyzed further
```

```
forall idx in zip(D, someIterator()) { } // the loop will not be analyzed further
```

- Indexing into shadow variables is not analyzed

```
forall i in D with (ref A) do  
    A[i] = calculate(i);
```

- Indexing into array views is not analyzed

```
var A = otherArr[2..10];  
forall i in A.domain do  
    A[i] = calculated(i);
```

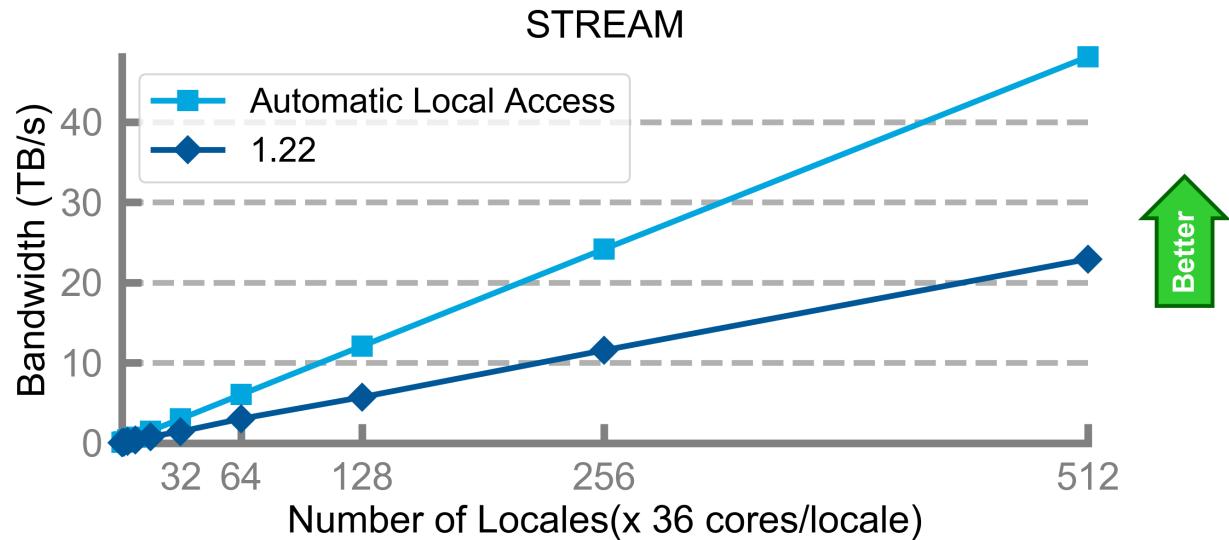


AUTOMATIC LOCAL ACCESS OPTIMIZATION

Impact

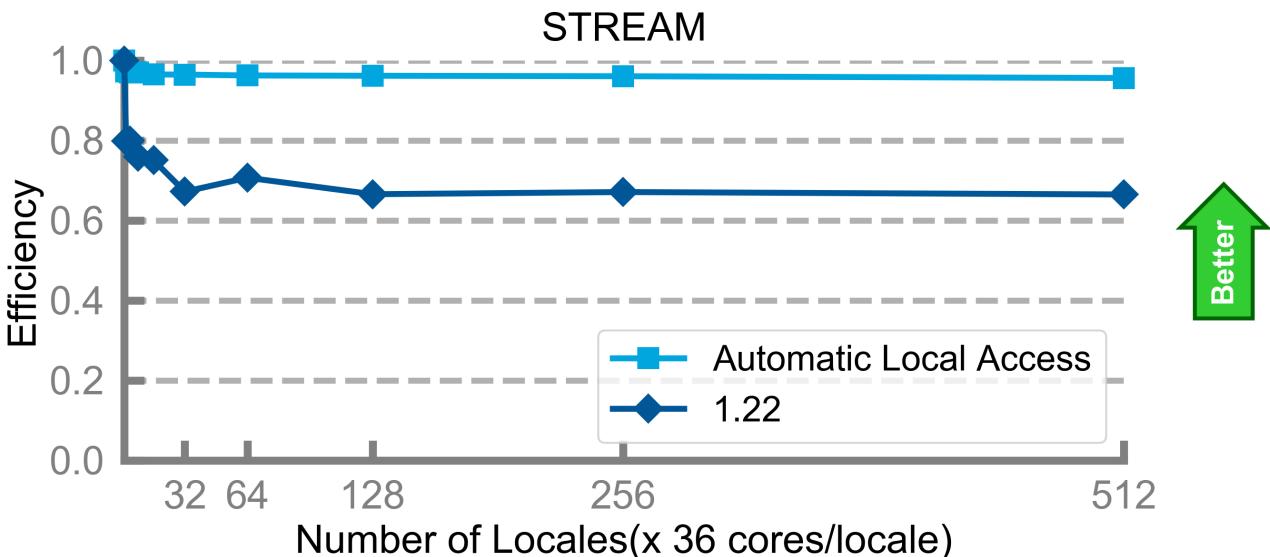
- Global STREAM with array indexing:

```
forall i in ProblemSpace do  
    A[i] = B[i] + alpha * C[i];
```



now essentially performs like other idioms:

```
forall (a, b, c) in zip(A, B, C) do  
    a = b + alpha * c;
```



or:

```
A = B + alpha * C;
```



AUTOMATIC LOCAL ACCESS OPTIMIZATION

Impact

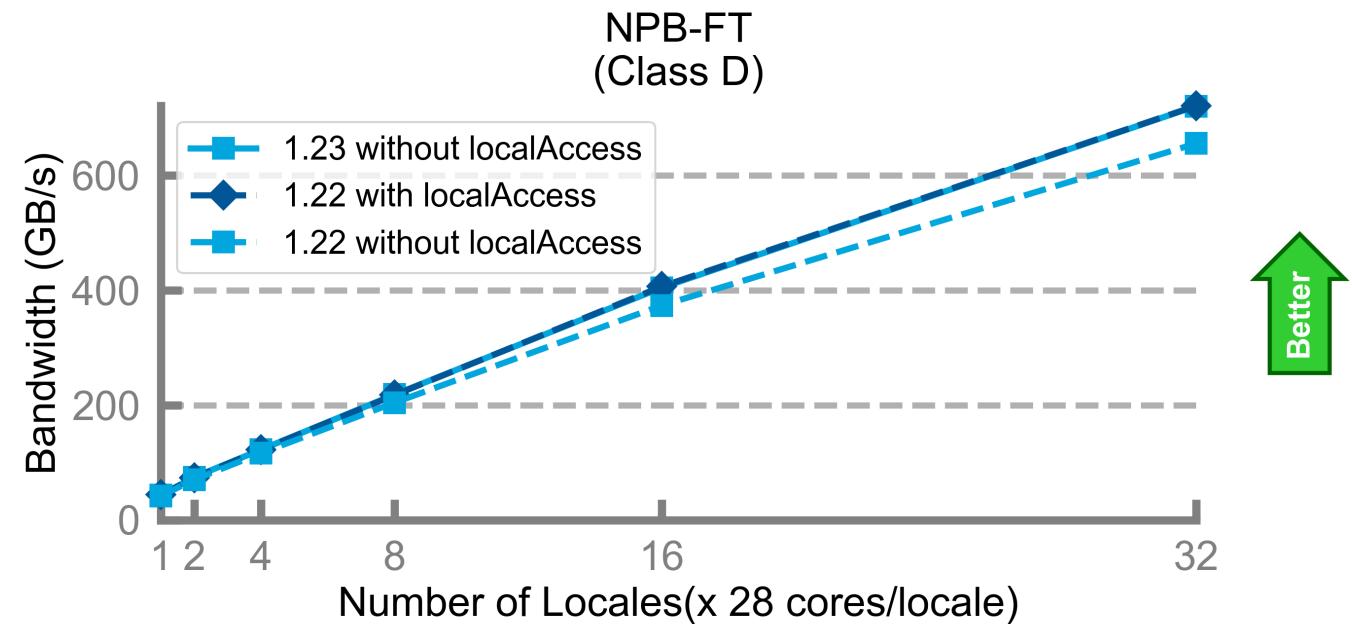
- Explicit 'localAccess' calls are no longer needed in NPB-FT

- Kernel with 'localAccess' calls

```
forall ijk in DomT {  
    const elt = V.localAccess[ijk] *  
                T.localAccess[ijk];  
  
    V.localAccess[ijk] = elt;  
    Wt.localAccess[ijk] = elt;  
}
```

- Kernel without 'localAccess' calls

```
forall ijk in DomT {  
    const elt = V[ijk] *  
                T[ijk];  
  
    V[ijk] = elt;  
    Wt[ijk] = elt;  
}
```



AUTOMATIC LOCAL ACCESS OPTIMIZATION

Next Steps

- Expand static check to certain array/domain operations, e.g.:

```
coforall l in Locales do on l {
    forall i in A.localSubdomain() do // localSubdomain always produces a subset
        A[i] = calculate(i);
    forall i in A.domain[someSlice] do // slicing always produces a subset
        A[i] = calculate(i)
}
```

- Accesses above will be optimized dynamically on Chapel 1.23, but we could optimize them statically
- Investigate how we can expand the analysis to affine accesses

```
forall i in A.domain do
    A[i] = calculate(A[i-1], A[i], A[i+1]);
```





IMPROVEMENTS TO ASSOCIATIVE TYPES

ASSOCIATIVE TYPES

Background and This Effort

Background: Historically, Chapel's lowest-level associative types were associative domains/arrays

- Hash table implementation was intertwined in domain/array implementation
 - Other types like set/map were built on top of associative domains/arrays
 - Wanted associative type for internal data structures, but associative domains created circular dependency

This Effort: Factored hash table implementation into an internal standalone type

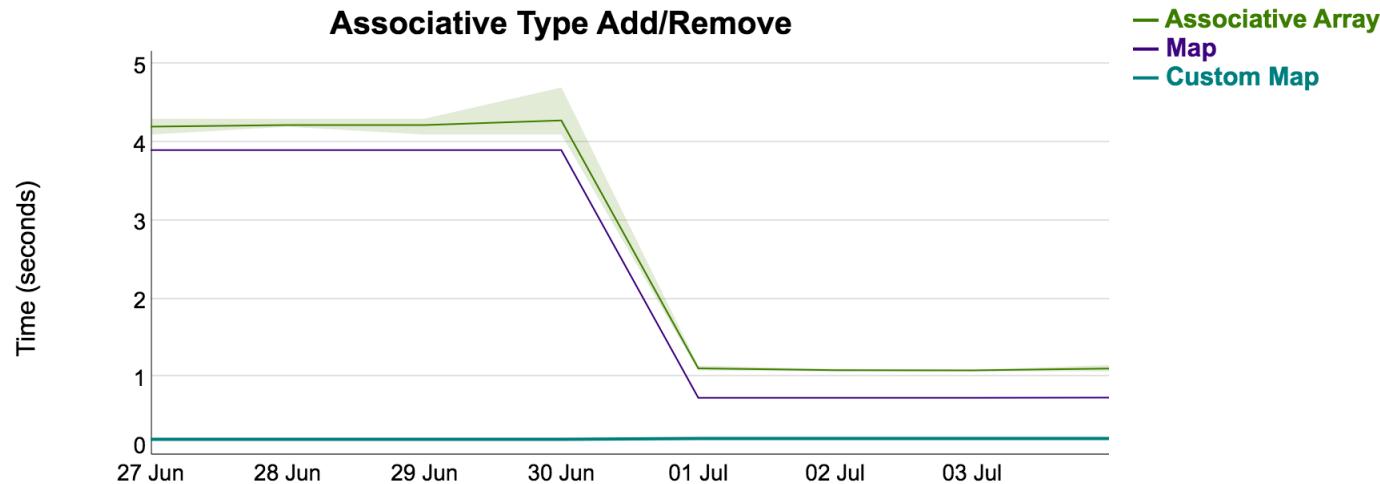
- Changed set/map types to use the standalone hash table, which enabled optimizations
- Further optimized hash table implementation, especially for repeated insertions/deletions



ASSOCIATIVE TYPES

Impact

- Significantly improved performance for associative types
 - Especially for repeated insertion/removal patterns identified by users



ARRAY TRACKING OPTIMIZATION



ARRAY TRACKING OPTIMIZATION

Background and This Effort

Background: Chapel domains track arrays declared over them

- Supports resizing arrays when their domain is modified:

```
var D = {1..10};  
var A: [D] int;  
var B: [D] int;  
D = {1..20};      // this resizes 'A' and 'B'
```

- Previously, domains tracked arrays with a linked list, which has $O(n)$ removal
- In many cases, arrays are removed in the opposite order that they are created, so $O(1)$ in practice
- However, for arrays-of-arrays that freed their array elements in parallel, $O(n)$ behavior occurred
 - Some user codes have suffered from this

This Effort: Switched from using a linked list to a hash table to track arrays

- Hash table insertion/removal is always $O(1)$



ARRAY TRACKING OPTIMIZATION

Impact

- Significantly reduced worst-case overheads for tracking arrays
 - ~700x speedup for task-intents with array-of-arrays

```
// Snippet from user n-body code
const nBodies = 10000;
const D = {0..#nBodies};
var forces: [D][0..#3] real;
forall d in D with (+ reduce forces) { ... }           // 486.5s -> 0.65s
```

- ~500x speedup for distributed array-of-arrays at 512 nodes

```
// Per-task timers from ISx, 9 timers in actual code
const D = newBlockDom(0..#numLocales*here.maxTaskPar);
var totalTimeSPMD, ...: [D][1..trials] real;           // 250.0s -> 0.5s
```



CONSTANT DOMAIN OPTIMIZATION



CONSTANT DOMAIN OPTIMIZATION

Background

- Tracking the arrays declared over a domain was optimized
 - However, tracking is only needed if the domain can be resized
 - Unnecessary if the domain is constant



CONSTANT DOMAIN OPTIMIZATION

This Effort

- Stop tracking arrays for domains declared 'const' or domain literals

```
const D = {1..10};  
var A: [D] int;           // no need to track A, 'D' is a constant
```

```
var B: [1..20] int;      // no need to track 'B', 1..20 is a constant
```

- An important case for this optimization is array-of-arrays

```
var A: [1..1_000_000] [1..5] int;    // no need to track 1 million arrays, 1..5 is a constant
```

- Add compiler analysis to detect domain creation/move/copy operations

- By *only* looking at variable/formal declarations
 - And *not* doing def/use analysis



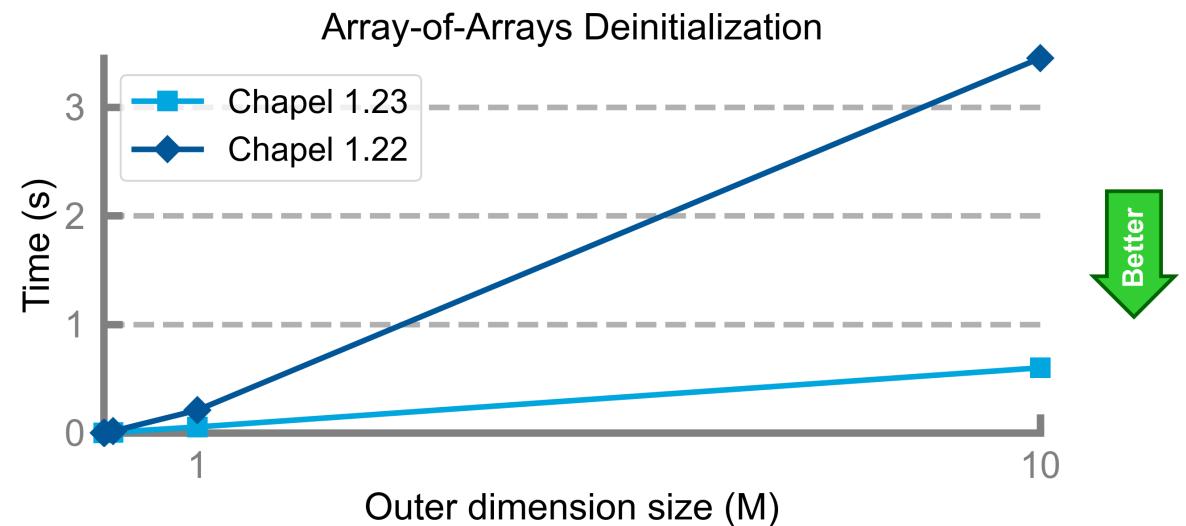
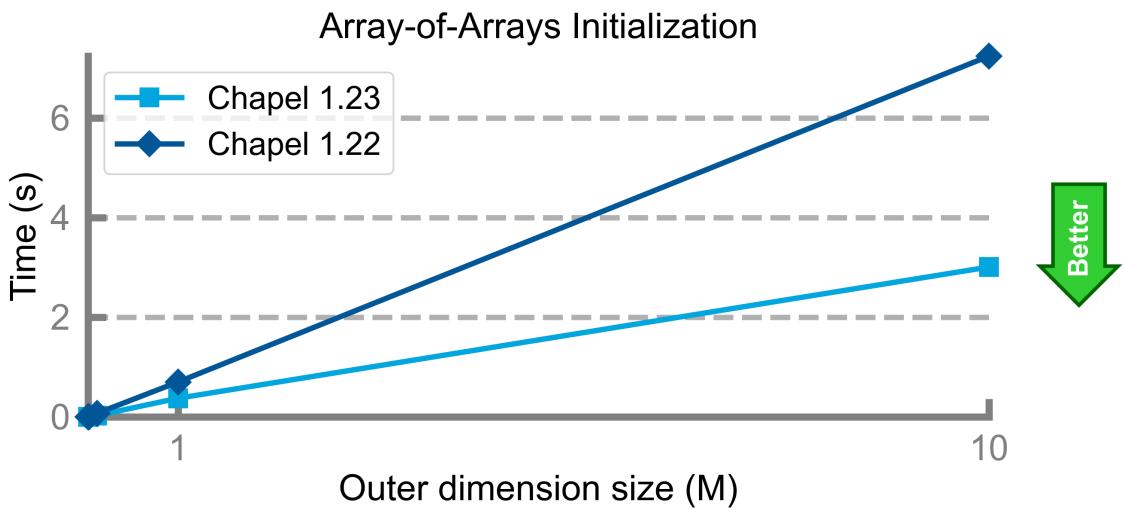
CONSTANT DOMAIN OPTIMIZATION

Impact

- More than 2x faster array initialization/deinitialization on constant domains

	Init (ns)	Deinit (ns)
Chapel 1.22	118	96
Chapel 1.23	51	47

- 2.5x faster initialization, 6x faster deinitialization for array-of-arrays



CONSTANT DOMAIN OPTIMIZATION

Next Steps

- Implement lighter-weight reference counting for domains
- More def/use analysis on domains and arrays can help cover some more cases
 - Passing a non-constant domain to a 'const ref' formal and defining an array on that formal
 - Domains that are declared 'var' but never modified
- Find answers for some semantic questions
 - Should we special-case domains w.r.t copy elision rules?
 - See <https://github.com/chapel-lang/chapel/issues/16431>



PARALLEL ARRAY INITIALIZATION



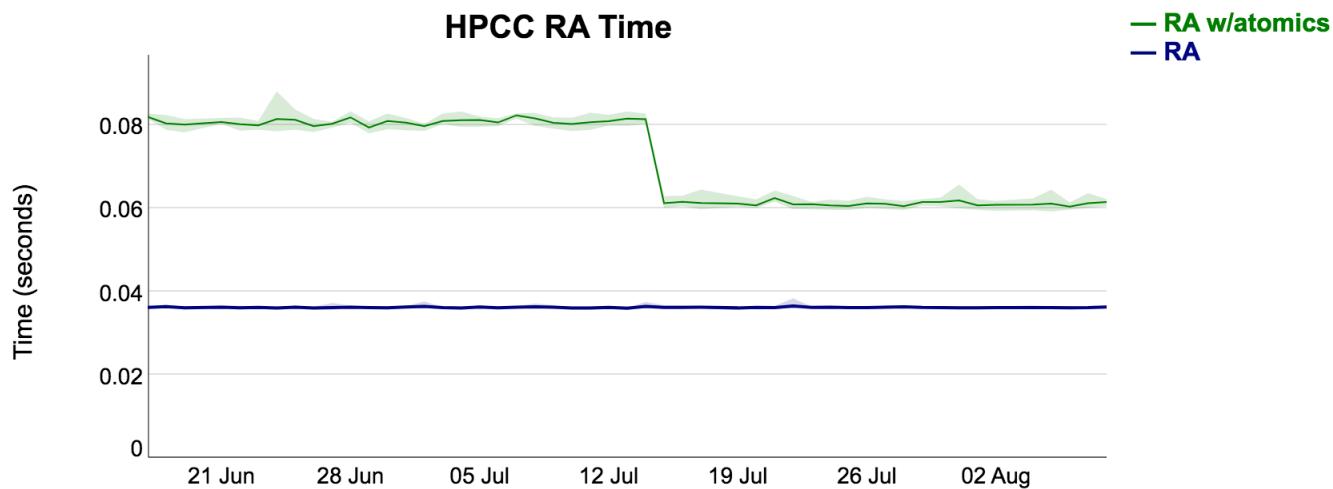
PARALLEL ARRAY INITIALIZATION

Background: Chapel initializes large numeric (integral/real/complex) arrays in parallel

- Performance issues with tracking a domain's arrays prevented parallelizing arrays-of-arrays
 - As a simplified proxy we only parallelized integral/real/complex arrays
 - Optimizing how arrays are tracked eliminated that performance issue

This Effort: Extend parallel initialization to all arrays

Impact: Better NUMA affinity for more arrays, which improves performance of parallel operations



PARALLEL ARRAY ASSIGNMENT

PARALLEL ARRAY ASSIGNMENT

Background and This Effort

Background:

- Large Chapel arrays are initialized in parallel
- However, array assignments were not parallel

```
var A: [1..n] int; // parallel default initialization  
var B: [1..n] int; // parallel default initialization  
  
A = B;           // this was done sequentially
```

- Especially in multi-socket systems, parallel 'memcpy's can improve the bandwidth significantly

This Effort:

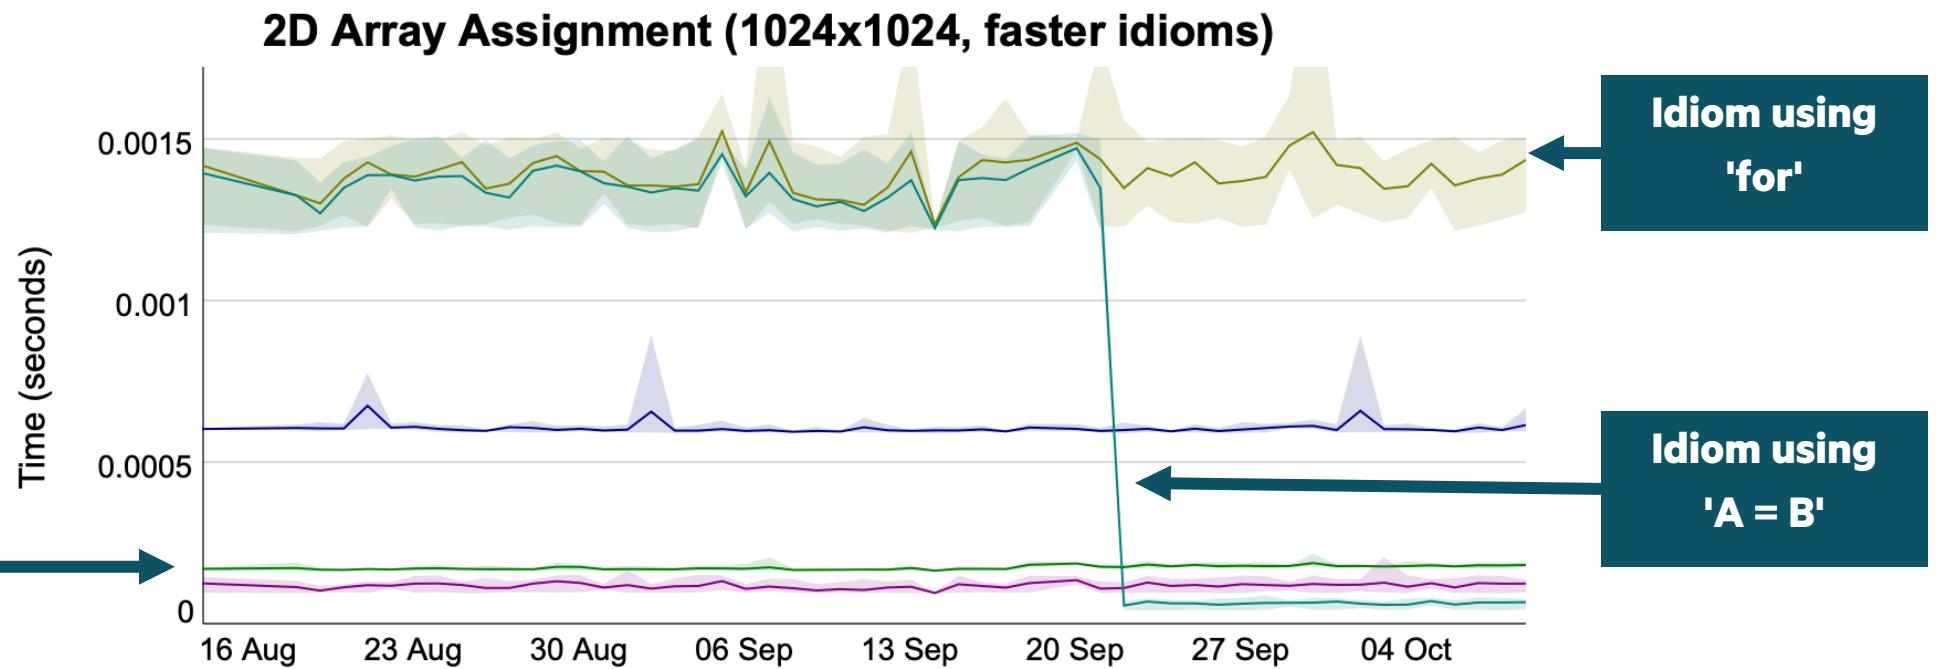
- Use parallel local copies for large array assignments if applicable



PARALLEL ARRAY ASSIGNMENT

Impact

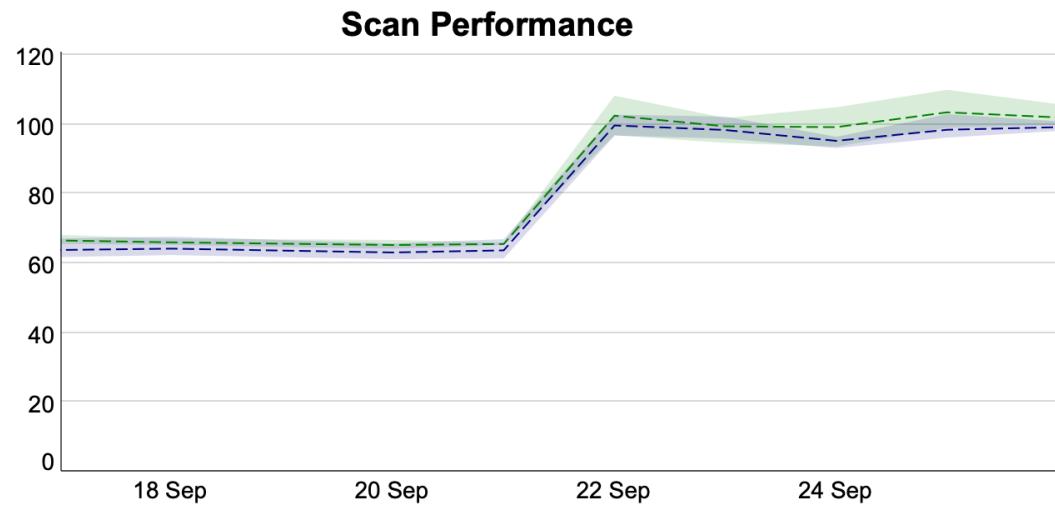
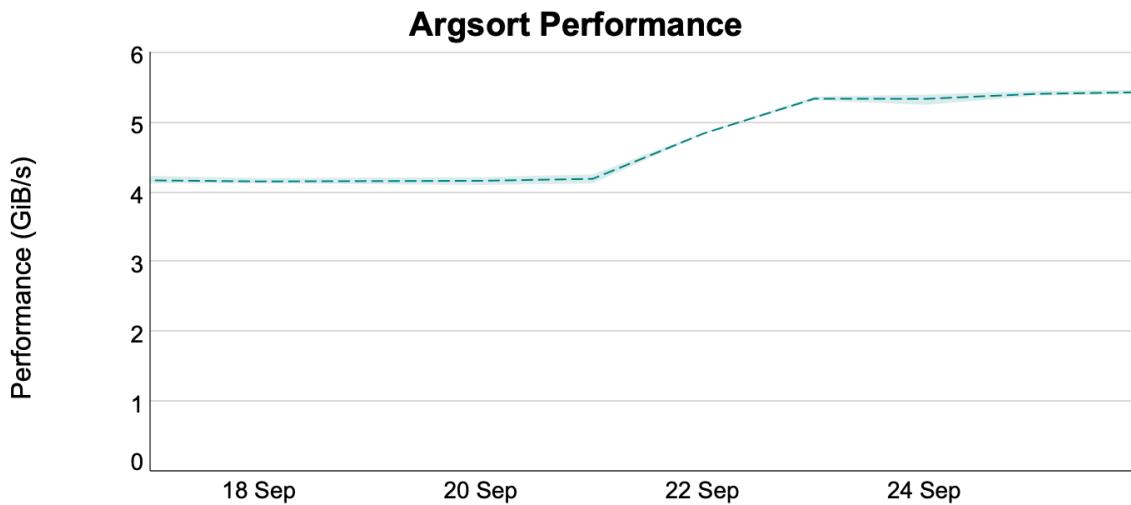
- Array copies are significantly faster



PARALLEL ARRAY ASSIGNMENT

Impact

- Arkouda performance improvements



PARALLEL ARRAY ASSIGNMENT

Next Steps

- Investigate making remote array copies parallel
 - Initial attempts resulted in some regressions



ARRAY SWAP OPTIMIZATION



ARRAY SWAP OPTIMIZATION

Background and This Effort

Background:

- Chapel supports a swap assignment operator ('<=>') for convenience and optimization opportunities
- Users have long requested that array swaps be performed using a pointer swap rather than per-element swaps
 - historically, this wasn't generally possible due to our implementation of array slices
 - once we switched to using array views, it enabled this optimization in many cases

This Effort: Implemented array swaps using pointer swaps for some common cases:

- default rectangular arrays that:
 - are the same size
 - are stored on the same locale
 - are not array views
- block-distributed arrays that:
 - have equivalent distributions



ARRAY SWAP OPTIMIZATION

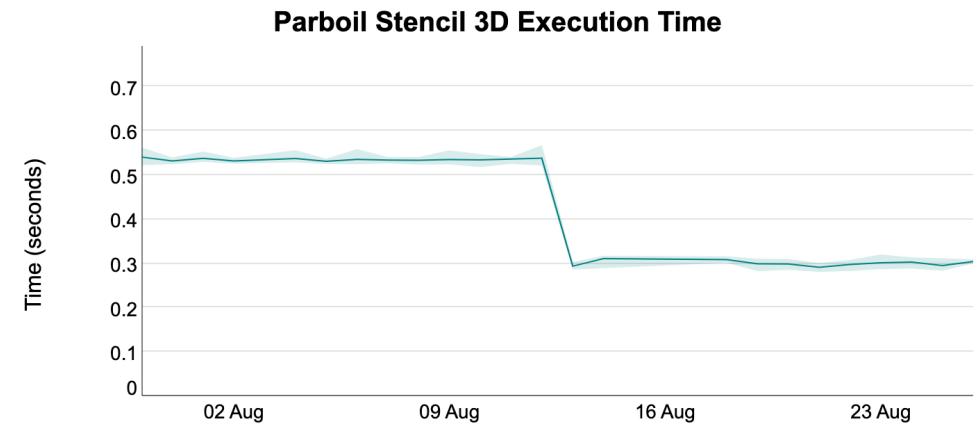
Impact

Impact: Turned array swaps for many cases from an $O(n)$ operation to $O(1)$ or $O(\#targetLocales)$

Array size	Local Array			Block Array (16 locales)		
	Before	After	Factor	Before	After	Factor
100M	32ms	~0.15ms	213x	67ms	2.7ms	24.8x
1B	310ms	~0.15ms	2070x	510ms	3.4ms	150x
10B	[OOM]	~0.15ms	N/A	5100ms	3.2ms	1590x

- Supports writing certain code patterns more productively, such as iterative stencil patterns:

```
var New, Old: [D] real;  
do {  
    New = computeStencil(Old);  
    const delta = max reduce abs(New - Old);  
    Old <=> New; //prepare for the next iteration  
}  
while delta > epsilon;
```



ARRAY SWAP OPTIMIZATION

Next Steps

Next Steps:

- Extend optimization to other array types and distributions
 - e.g., sparse arrays, Cyclic distributions, etc.
- Optimize other forms of array/sub-array swapping, for example:

```
A[i, ..] <=> A[j, ..]; // row swap — think about how to implement this efficiently on distributed arrays
A[.., i] <=> A[.., j]; // column swap — (these patterns appear in PNNL's work on CHGL)
```



COMPILATION TIME IMPROVEMENTS

- Single-Iteration Coforalls
- Other Compilation Time improvements

SINGLE-ITERATION COFORALLS



SINGLE-ITERATION COFORALLS

Background and This Effort

Background: ‘coforall’ loops create a distinct task per loop iteration

- Historically, many iterators would include special cases to avoid task creation for single-iteration coforalls

```
iter batch(r: range) {
    const numTasks = here.maxTaskPar - here.runningTasks() + 1;
    if numTasks == 1 then
        for i in r do
            yield i;
    else
        coforall tid in 0..<numTasks do
            for i in myChunk(tid, numTasks, r) do
                yield i;
}
```

This Effort: Optimize single-iteration coforalls

- Avoid task creation by having parent task run body directly
- Eliminate manipulation of atomic running tasks counter



SINGLE-ITERATION COFORALLS

Impact

- Significantly faster single-iteration coforalls

```
coforall 1..1 {} // ~13x faster with this optimization
```

```
coforall 1..here.maxTaskPar do  
    coforall 1..1 {} // ~90x faster with this optimization
```

- Single-iteration coforalls have little overhead now
 - Enabled removing special cases in iterators, reducing generated code size
 - ~3% faster compilation on average
 - ~15% faster Arkouda compilation





OTHER COMPILATION TIME IMPROVEMENTS

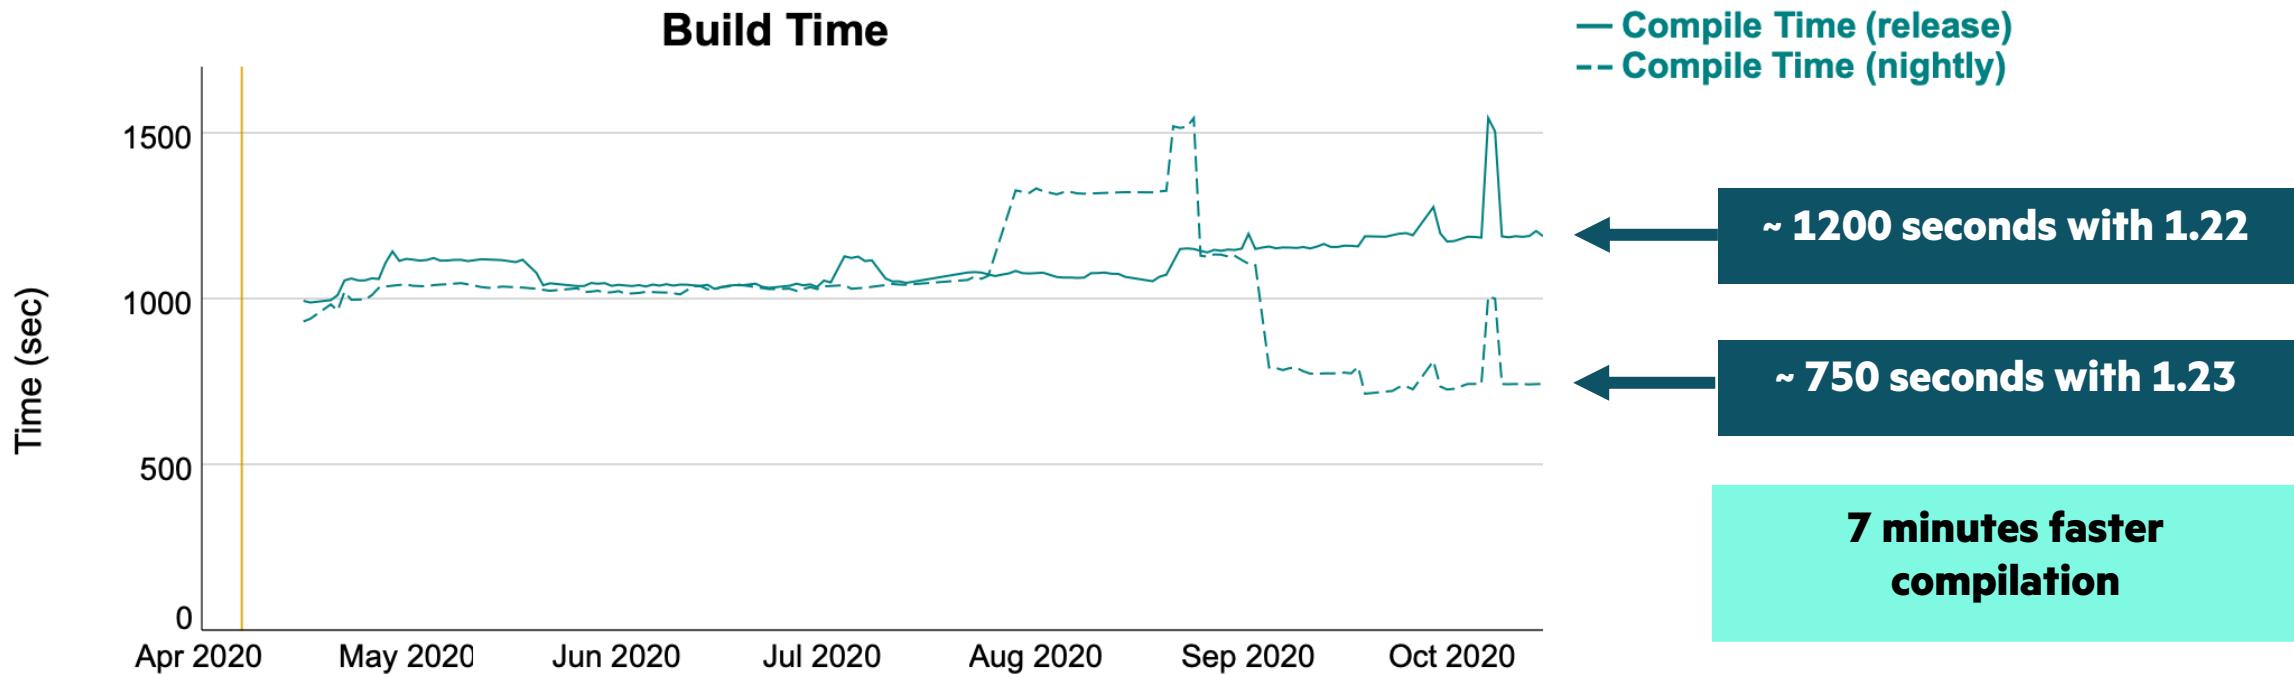
COMPILATION TIME IMPROVEMENTS

- Refactored formatted string implementation
 - Faster compilation for applications with lots of 'writef' and/or 'string.format' calls
 - ~30% faster Arkouda compilation
- Refactored several string/bytes operations
 - Reduced inlining with iterators and casts
 - ~9% faster compilation on average
 - ~3% faster Arkouda compilation
- Replaced some ‘where’-clauses with formal types
 - Fewer generic functions to resolve
 - ~7% faster compilation on average



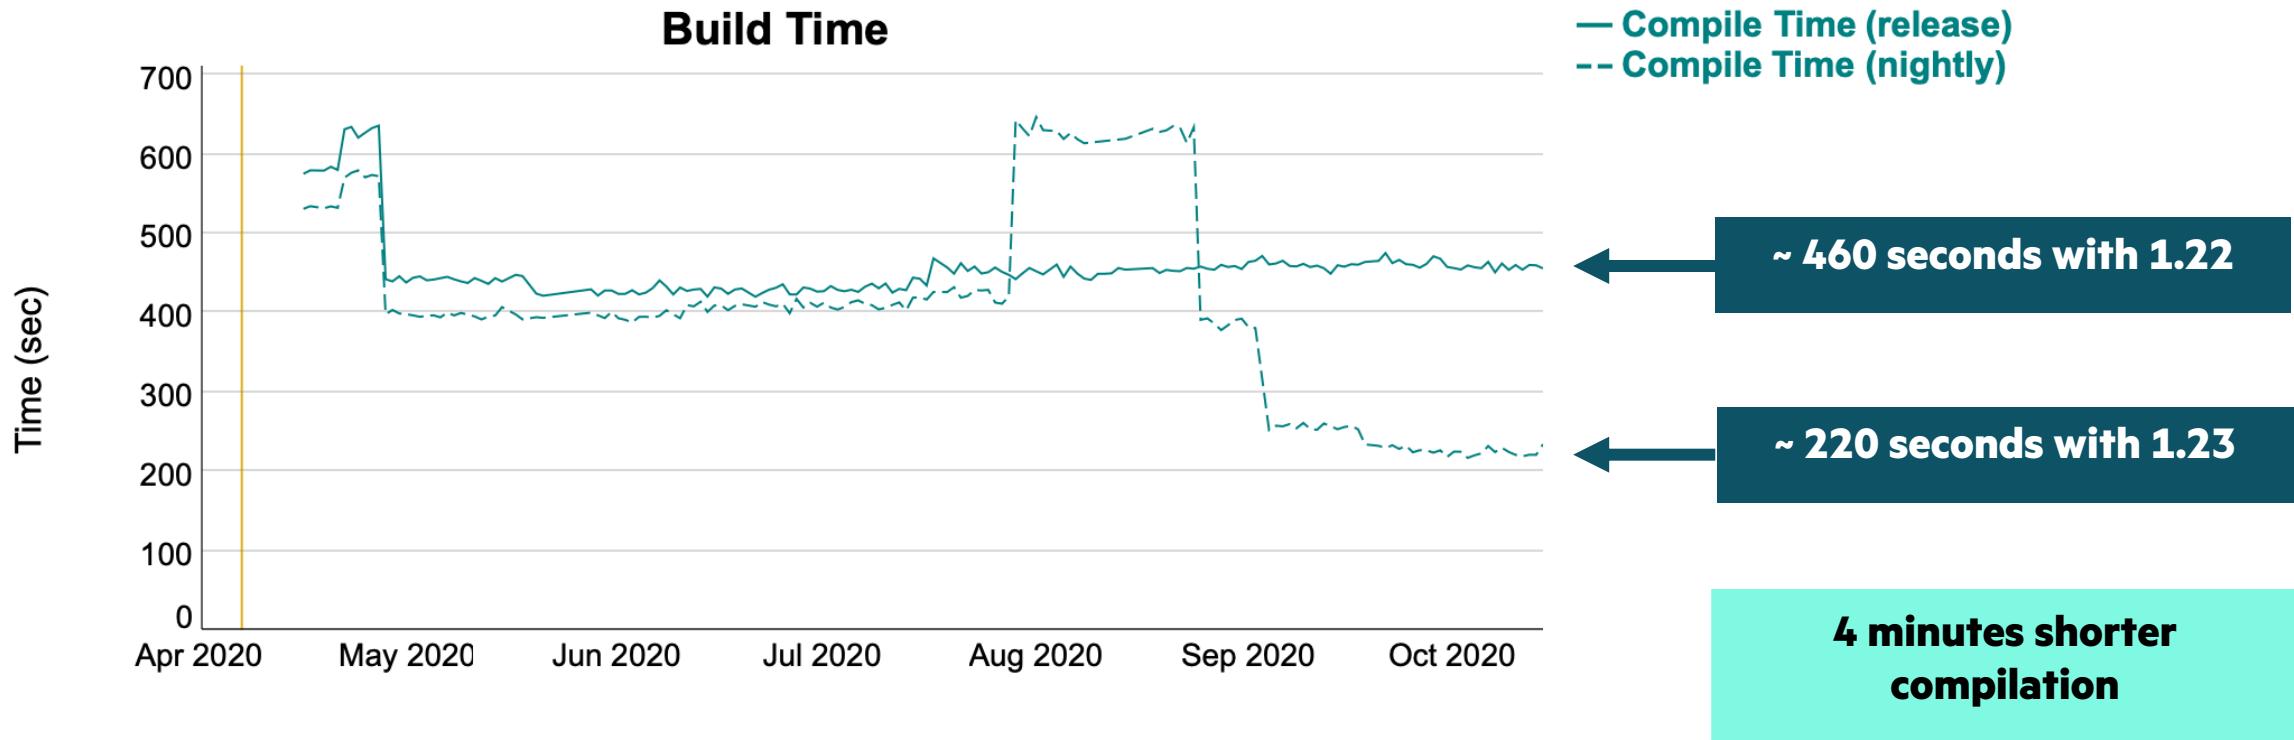
COMPILATION TIME IMPROVEMENTS

- Multi-locale Arkouda build time on Cray XC



COMPILATION TIME IMPROVEMENTS

- Single locale Arkouda build time



COMPILATION TIME IMPROVEMENTS

Next Steps

- More opportunities to reduce the generated code size and compilation time
 - We can stop inlining several array support functions
 - Need to investigate potential performance regressions
 - Iterator outlining
 - There are some large iterators that we inline even with '—no-fast'
 - Currently, non-inlined iterators generate even more code and are very slow
 - Investigate whether we can outline such iterators' bodies into helpers and inline smaller bodies



MEMORY IMPROVEMENTS

- Memory Fragmentation Improvements
- Memory Leak Improvements

MEMORY FRAGMENTATION IMPROVEMENTS



MEMORY FRAGMENTATION

Background and This Effort

Background: ‘jemalloc’ per-thread arenas can cause memory fragmentation

- Each thread allocates from a different arena to improve concurrent allocation performance
- Freed memory is not immediately returned to the system, but retained for later use to reduce system calls
- This leads to cross-thread fragmentation, which limits available memory for large allocations—for example:
 - thread/arena 0 allocates/frees a large array – had to grab memory from system, retains for future use
 - thread/arena 1 then does the same operation – cannot use arena 0 memory, must grab more from system
- This impacted configurations that allocate large arrays through ‘jemalloc’
 - Did not impact ugni, which uses a different allocation scheme for large arrays

This Effort: Use a single arena to satisfy large allocations

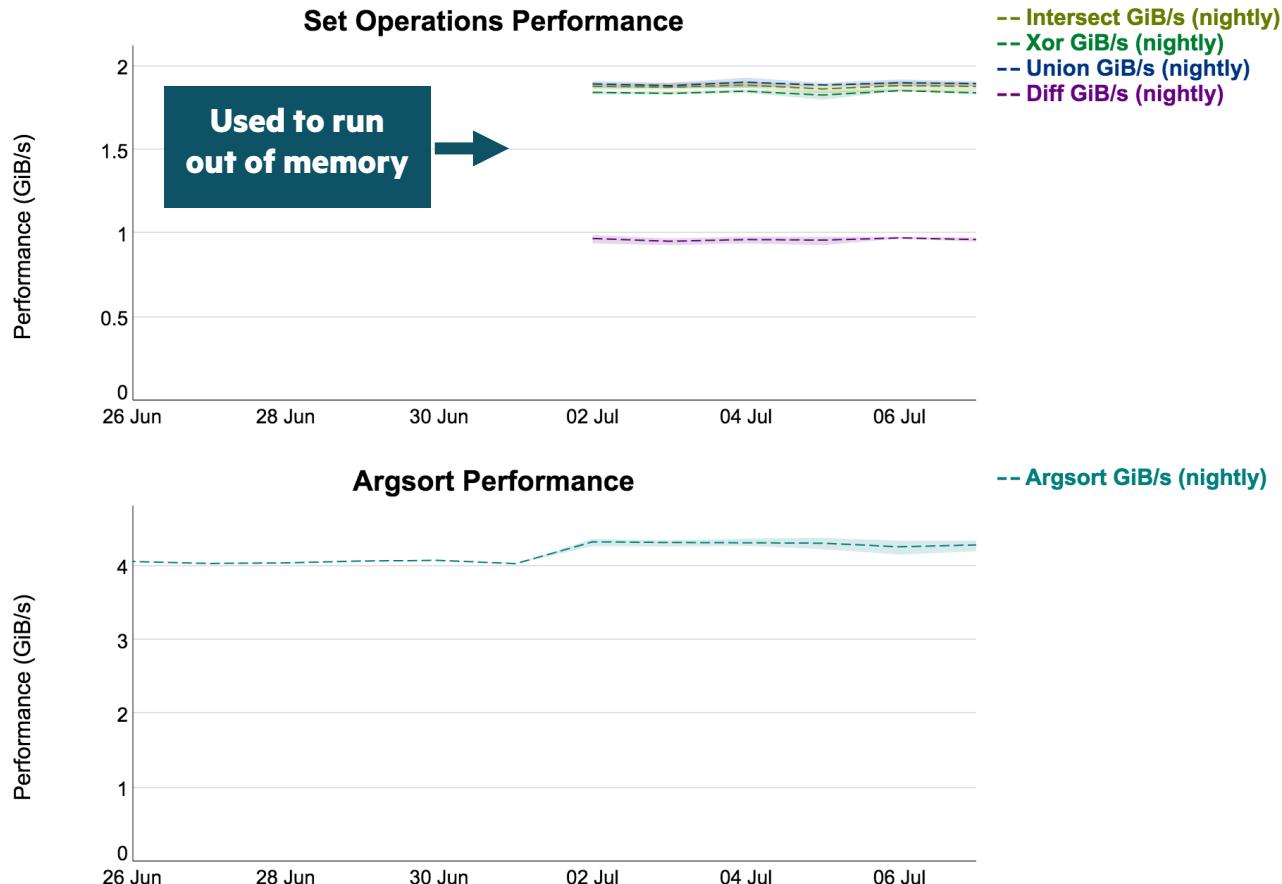
- Increases contention for large allocations, but concurrent large allocations are rare



MEMORY FRAGMENTATION

Impact

- Reduced memory fragmentation and improved performance for repeated array creation



A scenic landscape featuring a large, rugged mountain range in the background, with one prominent peak covered in snow. In the foreground, there is a calm body of water, likely a lake or river, reflecting the surrounding environment. A cluster of trees with vibrant autumn foliage (yellow and orange) stands on a small island or peninsula in the water. The sky is filled with dramatic, dark clouds.

MEMORY LEAK IMPROVEMENTS

MEMORY LEAKS

Background, This Effort and Next Steps

Background:

- Memory leaks have historically been tracked in graphs
 - Made sense when hundreds of tests leaked
 - Makes it cumbersome to triage leaks now that there are only a few leaking tests

This Effort:

- Converted multi-locale leak testing to a correctness test now that it has 0 leaks
- Classified remaining single-locale leaks into distinct bugs with smaller reproducers
 - We believe 24 leaking tests are coming from 8 different bugs
 - See <https://github.com/chapel-lang/chapel/issues/15623>

Next Steps:

- Investigate turning single-locale testing into correctness tests
 - Will require some adjustments for current known/expected leaks
- Close remaining single-locale leaks





OTHER PERFORMANCE IMPROVEMENTS

OTHER PERFORMANCE IMPROVEMENTS

For a more complete list of performance optimizations in the 1.23 release, refer to the following sections in the [CHANGES.md](#) file:

- ‘Performance Optimizations’
- ‘Memory Improvements’





THANK YOU

<https://chapel-lang.org>
@ChapelLanguage

