

Chapel Language Specification 0.702

Cray Inc
411 First Ave S, Suite 600
Seattle, WA 98104

Contents

1	Scope	1
2	Notation	3
3	Organization	5
4	Acknowledgments	7
5	Language Overview	9
5.1	Motivating Principles	9
5.2	Getting Started	10
5.3	Example Chapel Programs	11
5.3.1	Jacobi Method	11
5.3.2	Matrix and Vector Norms	16
5.3.3	Simple Producer-Consumer Program	20
5.3.4	Generic Stack Implementations	22
6	Lexical Structure	27
6.1	Comments	27
6.2	White Space	27
6.3	Case Sensitivity	27
6.4	Tokens	27
6.4.1	Identifiers	27
6.4.2	Keywords	28
6.4.3	Literals	28
6.4.4	Operators and Punctuation	28
6.4.5	Grouping Tokens	28
6.5	User-Defined Compiler Errors	29
7	Types	31
7.1	Primitive Types	31
7.1.1	The Bool Type	31
7.1.2	Signed and Unsigned Integral Types	31
7.1.3	Real Types	32
7.1.4	Complex Types	32
7.1.5	Imaginary Types	32
7.1.6	The String Type	33
7.1.7	Primitive Type Literals	33
7.2	Enumerated Types	34
7.3	Class Types	35
7.4	Record Types	35
7.5	Union Types	35
7.6	Tuple Types	35
7.7	Sequence Types	35
7.8	Domain and Array Types	35
7.9	Type Aliases	35

8	Variables	37
8.1	Variable Declarations	37
8.1.1	Default Initialization	38
8.1.2	Local Type Inference	38
8.2	Global Variables	38
8.3	Local Variables	38
8.4	Constants	39
8.4.1	Compile-Time Constants	39
8.4.2	Runtime Constants	39
8.5	Configuration Variables	40
9	Conversions	41
9.1	Implicit Conversions	41
9.1.1	Implicit Numeric Conversions	41
9.1.2	Implicit Enumeration Conversions	41
9.1.3	Implicit Class Conversions	42
9.1.4	Implicit Record Conversions	42
9.1.5	Implicit Compile-Time Constant Conversions	42
9.1.6	Implicit Statement Bool Conversions	42
9.2	Explicit Conversions	42
9.2.1	Explicit Numeric Conversions	42
9.2.2	Explicit Enumeration Conversions	42
9.2.3	Explicit Class Conversions	43
9.2.4	Explicit Record Conversions	43
10	Expressions	45
10.1	Literal Expressions	45
10.2	Variable Expressions	45
10.3	Call Expressions	45
10.3.1	Indexing Expressions	46
10.3.2	Member Access Expressions	46
10.4	The Query Expression	46
10.5	Casts	47
10.6	LValue Expressions	47
10.7	Operator Precedence and Associativity	48
10.8	Operator Expressions	48
10.9	Arithmetic Operators	49
10.9.1	Unary Plus Operators	49
10.9.2	Unary Minus Operators	49
10.9.3	Addition Operators	50
10.9.4	Subtraction Operators	50
10.9.5	Multiplication Operators	51
10.9.6	Division Operators	52
10.9.7	Modulus Operators	53
10.9.8	Exponentiation Operators	54
10.10	Bitwise Operators	54
10.10.1	Bitwise Complement Operators	54
10.10.2	Bitwise And Operators	54
10.10.3	Bitwise Or Operators	55
10.10.4	Bitwise Xor Operators	55
10.11	Shift Operators	55

10.12	Logical Operators	56
10.12.1	The Logical Negation Operator	56
10.12.2	The Logical And Operator	56
10.12.3	The Logical Or Operator	57
10.13	Relational Operators	57
10.13.1	Ordered Comparison Operators	57
10.13.2	Equality Comparison Operators	58
10.14	Miscellaneous Operators	59
10.14.1	The String Concatenation Operator	59
10.14.2	The Sequence Concatenation Operator	59
10.14.3	The Arithmetic Domain By Operator	59
10.14.4	The Arithmetic Sequence By Operator	59
10.15	Let Expressions	60
10.16	Conditional Expressions	60
11	Statements	61
11.1	Blocks	61
11.2	Block Level Statements	62
11.3	Expression Statements	62
11.4	Assignment Statements	62
11.5	The Conditional Statement	63
11.6	The Select Statement	63
11.7	The While and Do While Loops	64
11.8	The For Loop	65
11.8.1	Zipper Iteration	65
11.8.2	Tensor Product Iteration	65
11.8.3	Parameter For Loops	66
11.9	The Use Statement	66
11.10	The Type Select Statement	66
11.11	The Empty Statement	67
12	Modules	69
12.1	Module Definitions	69
12.2	Program Execution	69
12.2.1	The <i>main</i> Function	69
12.2.2	Command-Line Arguments	70
12.2.3	Module Execution	70
12.2.4	Programs with a Single Module	70
12.3	Using Modules	70
12.3.1	Explicit Naming	70
12.4	Nested Modules	71
12.5	Implicit Modules	71
13	Functions	73
13.1	Function Definitions	73
13.2	The Return Statement	74
13.3	Function Calls	74
13.4	Formal Arguments	74
13.4.1	Named Arguments	75
13.4.2	Default Values	75
13.5	Intents	75

13.5.1	The Blank Intent	75
13.5.2	The In Intent	76
13.5.3	The Out Intent	76
13.5.4	The Inout Intent	76
13.6	Variable Functions	76
13.6.1	Explicit Setter Functions	77
13.7	Function Overloading	77
13.8	Function Resolution	78
13.8.1	Identifying Visible Functions	78
13.8.2	Determining Candidate Functions	78
13.8.3	Determining More Specific Functions	79
13.9	Nested Functions	80
13.9.1	Accessing Outer Variables	80
13.10	Variable Length Argument Lists	80
14	Classes	83
14.1	Class Declarations	83
14.2	Class Assignment	83
14.3	Class Fields	83
14.3.1	Class Field Accesses	84
14.4	Class Methods	84
14.4.1	Class Method Declarations	84
14.4.2	Class Method Calls	84
14.4.3	The <i>this</i> Reference	85
14.4.4	Class Methods without Parentheses	85
14.4.5	The <i>this</i> Method	85
14.5	Class Constructors	86
14.5.1	The Default Constructor	86
14.6	Getters and Setters	86
14.7	Inheritance	87
14.7.1	Accessing Base Class Fields	87
14.7.2	Derived Class Constructors	87
14.7.3	Shadowing Base Class Fields	87
14.7.4	Overriding Base Class Methods	87
14.7.5	Inheriting from Multiple Classes	88
14.8	Class Promotion of Scalar Functions	88
14.9	Nested Classes	88
14.10	Automatic Memory Management	88
15	Records	89
15.1	Record Declarations	89
15.2	Class and Record Differences	89
15.2.1	Records as Value Classes	89
15.2.2	Record Inheritance	89
15.2.3	Record Assignment	90
15.3	Default Comparison Operators on Records	90

16 Unions	91
16.1 Union Declarations	91
16.1.1 Union Fields	91
16.2 Union Assignment	91
16.3 The Type Select Statement and Unions	91
17 Tuples	93
17.1 Tuple Expressions	93
17.2 Tuple Type Definitions	93
17.3 Tuple Assignment	93
17.4 Tuple Destructuring	94
17.4.1 Variable Declarations in a Tuple	94
17.4.2 Ignoring Values with Underscore	94
17.5 Homogeneous Tuples	94
17.5.1 Declaring Homogeneous Tuples	95
17.6 Tuple Indexing	95
17.7 Formal Arguments of Tuple Type	95
17.7.1 Formal Argument Declarations in a Tuple	95
18 Sequences	97
18.1 Sequence Literals	97
18.2 Sequence Type Definitions	97
18.3 Sequence Rank	97
18.4 Sequence Assignment	97
18.5 Iteration over Sequences	98
18.6 Sequence Concatenation	98
18.7 Sequence Indexing	98
18.7.1 Sequence Indexing by Integers	98
18.7.2 Sequence Indexing by Tuples	99
18.8 Sequence Promotion of Scalar Functions	99
18.8.1 Zipper Promotion	99
18.8.2 Tensor Product Promotion	100
18.9 Sequence Operators	100
18.10 Sequences in Logical Contexts	100
18.10.1 Sequences in Select Statements	101
18.11 Filtering Predicates	101
18.12 Methods and Functions on Sequences	101
18.12.1 The <i>length</i> Method	101
18.12.2 The <i>reverse</i> Method	101
18.12.3 The <i>spread</i> Function	102
18.12.4 The <i>transpose</i> Function	102
18.12.5 The <i>reshape</i> Function	102
18.13 Arithmetic Sequences	103
18.13.1 Strided Arithmetic Sequences	103
18.13.2 Querying the Bounds and Stride of an Arithmetic Sequence	103
18.13.3 Indefinite Sequences	103
18.14 Conversions Between Sequences and Tuples	104

19	Domains and Arrays	105
19.1	Domains	105
19.1.1	Domain Types	105
19.1.2	Index Types	106
19.1.3	Domain Assignment	106
19.1.4	Formal Arguments of Domain Type	106
19.1.5	Iteration over Domains	106
19.1.6	Domain Promotion of Scalar Functions	107
19.2	Arrays	107
19.2.1	Array Types	107
19.2.2	Array Indexing	108
19.2.3	Array Slicing	108
19.2.4	Array Assignment	108
19.2.5	Formal Arguments of Array Type	109
19.2.6	Iteration over Arrays	109
19.2.7	Array Promotion of Scalar Functions	109
19.2.8	Array Initialization	110
19.3	Arithmetic Domains and Arrays	110
19.3.1	Arithmetic Domain Literals	110
19.3.2	Arithmetic Domain Types	111
19.3.3	Strided Arithmetic Domains	111
19.3.4	Arithmetic Domain Indexing	111
19.3.5	Arithmetic Array Indexing	112
19.3.6	Arithmetic Array Slicing	112
19.3.7	Formal Arguments of Arithmetic Array Type	112
19.4	Sparse Domains and Arrays	112
19.4.1	Changing the Indices in Sparse Domains	113
19.5	Indefinite Domains and Arrays	113
19.5.1	Changing the Indices in Indefinite Domains	114
19.5.2	Testing Membership in Indefinite Domains	114
19.6	Opaque Domains and Arrays	114
19.7	Enumerated Domains and Arrays	115
19.8	Association of Arrays to Domains	115
19.9	Subdomains	115
19.10	Predefined Functions and Methods on Domains	116
19.11	Predefined Functions and Methods on Arrays	116
20	Iterators	117
20.1	Iterator Functions	117
20.2	The Yield Statement	117
20.3	Iterator Calls	117
20.3.1	Iterators in For and Forall Loops	117
20.3.2	Iterators as Sequences	117
20.4	The Structural Iterator Interface	118
21	Generics	119
21.1	Generic Functions	119
21.1.1	Formal Type Arguments	119
21.1.2	Formal Parameter Arguments	120
21.1.3	Formal Arguments without Types	120
21.1.4	Formal Arguments with Queried Types	120

21.1.5	Formal Arguments of Generic Type	121
21.1.6	Formal Arguments of Generic Array Types	121
21.2	Function Visibility in Generic Functions	121
21.3	Generic Types	121
21.3.1	Type Aliases in Generic Types	122
21.3.2	Parameters in Generic Types	122
21.3.3	Fields without Types	123
21.3.4	Fields of Generic Types	123
21.3.5	Generic Methods	123
21.3.6	The <i>elt_type</i> Type	123
21.4	Where Expressions	124
21.5	Example: A Generic Stack	124
22	Parallelism and Synchronization	125
22.1	The Forall Loop	125
22.1.1	Alternative Forall Loop Syntax	125
22.1.2	The Ordered Forall Loop	126
22.2	The Forall Expression	126
22.3	The Cobegin Statement	127
22.4	The Begin Statement	127
22.5	The Ordered Expression	127
22.6	The Serial Statement	128
22.7	Synchronization Variables	128
22.7.1	Single Variables	129
22.7.2	Sync Variables	130
22.7.3	Additional Synchronization Variable Functions	130
22.7.4	Synchronization Variables of Record and Class Types	131
22.8	Memory Consistency Model	131
22.9	Atomic Statement	131
23	Locality and Distribution	133
23.1	Locales	133
23.1.1	The Locale Type	133
23.1.2	Predefined Locales Array	133
23.1.3	Querying the Locale of a Variable	134
23.2	Specifying Locales for Computation	134
23.2.1	On	134
23.2.2	On and Iterators	135
23.3	Distributions	135
23.3.1	Distributed Domains	135
23.3.2	Distributed Arrays	136
23.3.3	Undistributed Domains and Arrays	136
23.4	Standard Distributions	136
23.5	User-Defined Distributions	136
24	Reductions and Scans	137
24.1	Reduction Expressions	137
24.2	Scan Expressions	137
24.3	User-Defined Reductions and Scans	138

25	Input and Output	139
25.1	The <i>file</i> type	139
25.2	Standard files <i>stdout</i> , <i>stdin</i> , and <i>stderr</i>	140
25.3	The <i>write</i> , <i>writeln</i> , and <i>read</i> functions	140
25.4	The <i>write</i> and <i>writeln</i> method on files	140
25.5	The <i>read</i> method on files	140
25.6	User-Defined <i>read</i> and <i>write</i> methods	140
25.7	Default <i>read</i> and <i>write</i> methods	141
26	Standard Modules	143
26.1	BitOps	143
26.2	Math	143
26.3	Random	147
26.4	Standard	148
26.5	Time	149
26.6	Types	149

1 Scope

Chapel is a new parallel programming language that is under development at Cray Inc. in the context of the DARPA High Productivity Language Systems initiative and the DARPA High Productivity Computing Systems initiative. This document specifies the Chapel language.

This document is a work in progress and is not definitive. In particular, it is not a standard.

2 Notation

Special notations are used in this specification to denote Chapel code and to denote Chapel syntax.

Chapel code is represented with a fixed-width font where keywords are bold and comments are italicised.

Example.

```
for i in D do    // iterate over domain D
    writeln(i);    // output indices in D
```

Chapel syntax is represented with standard syntax notation in which productions define the syntax of the language. A production is defined in terms of non-terminal (*italicized*) and terminal (non-italicized) symbols. The complete syntax defines all of the non-terminal symbols in terms of one another and terminal symbols.

A definition of a non-terminal symbol is a multi-line construct. The first line shows the name of the non-terminal that is being defined followed by a colon. The next lines before an empty line define the alternative productions to define the non-terminal.

Example. The production

```
bool-literal :
    true
    false
```

defines *bool-literal* to be either the symbol `true` or `false`.

In the event that a single line of a definition needs to break across multiple lines of text, more indentation is used to indicate that it is a continuation of the same alternative production.

As a short-hand for cases where there are many alternatives that define one symbol, the first line of the definition of the non-terminal may be followed by “one of” to indicate that the single line in the production defines alternatives for each symbol.

Example. The production

```
unary-operator : one of
    + - ~ !
```

is equivalent to

```
unary-operator :
    +
    -
    ~
    !
```

As a short-hand to indicate an optional symbol in the definition of a production, the subscript “opt” is suffixed to the symbol.

Example. The production

formal :
formal-tag identifier formal-type_{opt} default-expression_{opt}

is equivalent to

formal :
formal-tag identifier formal-type default-expression
formal-tag identifier formal-type
formal-tag identifier default-expression
formal-tag identifier

3 Organization

This specification is organized as follows:

- Section 1, Scope, describes the scope of this specification.
- Section 2, Notation, introduces the notation that is used throughout this specification.
- Section 3, Organization, describes the contents of each of the sections within this specification.
- Section 4, Acknowledgments, offers a note of thanks to people and projects.
- Section 5, Language Overview, describes Chapel at a high level.
- Section 6, Lexical Structure, describes the lexical components of Chapel.
- Section 7, Types, describes the types in Chapel and defines the primitive and enumerated types.
- Section 8, Variables, describes variables and constants in Chapel.
- Section 9, Conversions, describes the legal implicit and explicit conversions allowed between values of different types. Chapel does not allow for user-defined conversions.
- Section 10, Expressions, describes the serial expressions in Chapel.
- Section 11, Statements, describes the serial statements in Chapel.
- Section 12, Modules, describes modules, Chapel's abstraction to allow for name space management.
- Section 13, Functions, describes functions and function resolution in Chapel.
- Section 14, Classes, describes reference classes in Chapel.
- Section 15, Records, describes records or value classes in Chapel.
- Section 16, Unions, describes unions in Chapel.
- Section 17, Tuples, describes tuples in Chapel.
- Section 18, Sequences, describes sequences in Chapel.
- Section 19, Domains and Arrays, describes domains and arrays in Chapel. Chapel arrays are more general than arrays in many other languages. Domains are index sets, an abstraction that is typically not distinguished from arrays.
- Section 20, Iterators, describes iterator functions and a class iterator interface in Chapel.
- Section 21, Generics, describes Chapel's support for generic functions and types.
- Section 22, Parallelism and Synchronization, describes parallel expressions and statements in Chapel as well as synchronization constructs and atomic sections.
- Section 23, Locality and Distribution, describes constructs for managing locality and distributing data in Chapel.
- Section 24, Reductions and Scans, describes the built-in reductions and scans as well as structural interfaces to support user-defined reductions and scans.

- Section 25, Input and Output, describes support for input and output in Chapel, including file input and output..
- Section 26, Standard Modules, describes the standard modules that are provided with the Chapel language.

4 Acknowledgments

We would like to recognize the following people for their efforts and impact on the Chapel language and its implementation—David Callahan, Hans Zima, John Plevyak, Shannon Hoffswell, Roxana Diaconescu, Mark James, Mackale Joyner, and Robert Bocchino.

Chapel is a derivative of a number of parallel and distributed languages and takes ideas directly from them. These include the MTA extensions of C, HPF, and ZPL.

Chapel also takes many serial programming ideas from many other programming languages, especially C#, C++, Java, Fortran, and Ada.

The preparation of this specification was made easier and the final result greatly improved because of the good work that went in to the creation of other language standards and specifications, in particular the specifications of C# and C.

5 Language Overview

5.1 Motivating Principles

Chapel is a new programming language being developed by Cray Inc. as part of DARPA's High Productivity Computing Systems (HPCS) program to improve the productivity of programming parallel systems. There are four main motivating principles for the design of the Chapel language.

General Parallel Programming Chapel's first motivating principle is to support general parallel programming through the use of high-level language abstractions for expressing parallelism. Chapel supports a *global-view programming model* that raises the level of abstraction for the expression of both data and control flow as compared to parallel programming models currently used in production.

Global-view data structures are arrays and other data aggregates whose sizes and indices are expressed globally in spite of the fact that their implementations may distribute them across the memories of multiple nodes or *locales*.¹ This contrasts with most parallel languages used in practice, which tend to require users to partition distributed data aggregates into per-processor chunks, either manually or using language abstractions. As a simple example, to create a 0-based vector with n elements distributed between p locales, a language like Chapel that supports global-view data structures allows the user to declare the array to contain n elements and to refer to the array using the indices $0 \dots n - 1$. In contrast, most traditional approaches require the user to declare the array as p chunks of n/p elements each and to specify and manage inter-processor communication and synchronization explicitly (and the details can be messy if p does not divide n evenly). Moreover, the chunks are typically accessed using local indices on each processor (*e.g.*, $0..n/p$), requiring the user to explicitly translate between logical indices and those used by the implementation.

A *global view of control* means that a user's program commences execution with a single logical thread of control and then introduces additional parallelism through the use of certain language concepts. All parallelism in Chapel is implemented via multithreading, though these threads are created via high-level language concepts and managed by the compiler and runtime, rather than through explicit fork/join-style programming. An impact of this approach is that Chapel can express parallelism that is more general than the Single Program, Multiple Data (SPMD) model that today's most common parallel programming approaches use as the basis for their programming and execution models. Chapel's general support for parallelism does not preclude users from coding in an SPMD style if they wish.

Supporting general parallel programming also means targeting a broad range of parallel architectures. Chapel is designed to target a wide spectrum of HPC hardware including clusters of commodity processors and SMPs; vector, multithreading, and multicore processors; custom vendor architectures; distributed-memory, shared-memory, and shared address space architectures; and networks of any topology. Our portability goal is to have any legal Chapel program run correctly on all of these architectures, and for Chapel programs that express parallelism in an architecturally-neutral way to perform reasonably on all of them. Naturally, Chapel programmers can tune their codes to more closely match a particular machine's characteristics, though doing so may cause the program to be a poorer match for other architectures.

¹A *locale* in Chapel is a unit of the target architecture that supports computation and data storage. Locales are defined for an architecture such that a locale's threads will all have similar access times to any specific memory address. For commodity clusters, each of their (single-core) processors, multicore processors, or SMP nodes would be considered a locale.

Control of Locality A second principle in Chapel is to allow the user to optionally and incrementally specify where data and computation should be placed on the physical machine. We consider this control over program locality to be essential for achieving scalable performance on large machine sizes. Such control contrasts with shared-memory programming models which present the user with a flat memory model. It also contrasts with SPMD-based programming models in which such details are explicitly specified by the programmer on a process-by-process basis via the multiple cooperating program instances.

Object-Oriented Programming (OOP) A third principle in Chapel is support for object-oriented programming. OOP has been instrumental in raising productivity in the mainstream programming community due to its encapsulation of related data and functions into a single software component, its support for specialization and reuse, and its use as a clean mechanism for defining and implementing interfaces. Chapel supports objects in order to make these benefits available in a parallel language setting, and to provide a familiar paradigm for members of the mainstream programming community. Chapel supports traditional reference-based classes as well as value classes that are assigned and passed by value.

Chapel does not require the programmer to use an object-oriented style in their code, so that traditional Fortran and C programmers in the HPC community need not adopt a new programming paradigm in order to use Chapel effectively. Many of Chapel's standard library capabilities are implemented using objects, so such programmers may need to utilize a method-invocation style of syntax to use these capabilities. However, using such libraries does not necessitate broader adoption of OOP methodologies.

Generic Programming Chapel's fourth principle is support for generic programming and polymorphism. These features allow code to be written in a style that is generic across types, making it applicable to variables of multiple types, sizes, and precisions. The goal of these features is to support exploratory programming as in popular interpreted and scripting languages, and to support code reuse by allowing algorithms to be expressed without explicitly replicating them for each possible type. This flexibility at the source level is implemented by having the compiler create versions of the code for each required type signature rather than by relying on dynamic typing which would result in unacceptable runtime overheads for the HPC community.

Chapel's first two principles are designed to provide support for general, performance-oriented parallel programming through high-level abstractions. The second two principles are supported to help narrow the gulf that exists between parallel programming languages and mainstream programming and scripting languages.

5.2 Getting Started

A Chapel version of the standard "hello world" computation is given here:

```
writeln("Hello, world!");
```

This program contains a single line of code that makes a call to the standard `writeln` subroutine, passing it a string literal argument, `"Hello, world!"`. This call causes the string to be printed to the console when the program is executed.

In general, Chapel programs define code using one or more named *modules*, each of which supports top-level initialization code that is invoked the first time the module is used. Programs also define a single entry point via a subroutine named `main`. To facilitate exploratory programming, Chapel allows programmers to define modules using files rather than an explicit module declaration, and to omit the program entry point when the program only has a single user module. This example takes advantage of both these features.

Chapel code is stored in files with the extension `.chpl`. Assuming this program is stored in a file called `hello.chpl`, it would define a single user module, `hello`, whose name is taken from the filename. Since the file defines a module, the top-level code in the file defines the module's initialization code. And since the program is composed of the single `hello` module, it need not define an entry point. Thus, when the program is executed, the single `hello` module will be initialized by executing its top-level code, invoking the call to `writeln()`, and printing out the message.

To compile and run this program, execute the following commands at the system prompt:

```
> chpl -o hello hello.chpl
> ./hello
```

The following output will be printed to the console:

```
> Hello, world!
```

5.3 Example Chapel Programs

To introduce the Chapel language, four short example codes are presented and discussed. Each example highlights certain Chapel features, showing how they are used in the context of a program. These examples do not cover all of Chapel's features. They are intended to introduce the user to many of the basic serial features that are currently supported in the compiler and to demonstrate how to program with these features. As more features are supported in the Chapel compiler, more examples will be added to this section.

All examples in this section are included with the release of the Chapel compiler.

5.3.1 Jacobi Method

Description The following example Chapel program solves a system of finite difference equations for the Laplace equation using the Jacobi method. The program uses two-dimensional arrays, `x` and `xNew`, to store and calculate the approximate solution. At each iteration, the next approximation for the solution at each grid point, `xNew(i, j)` is calculated by computing the average of the four neighboring grid points, `x(i-1, j)`, `x(i, j-1)`, `x(i+1, j)`, and `x(i, j+1)`. After all entries in `xNew` are computed, `x` is assigned `xNew` and convergence is tested. If convergence has not been reached, the next approximation is calculated, and so on until the convergence test is met.

Chapel Features This program demonstrates how *arrays* are declared and used. In Chapel, arrays are declared using *domains*. Domains are sets of indices which may be distributed across multiple processors indicating how data and parallel work should be divided among the processors. An array is a mapping from the domain to a collection of variables. An array is thus defined using a domain, distributing the entries of the array according to the domain's distribution. Domains may also be used to define array *slices* and as *iterators*. There are five kinds of domains in the Chapel language: *arithmetic*, *sparse*, *indefinite*, *opaque* and *enumerated*. This example program uses arithmetic domains and arrays.

Code Listing The Chapel code for this example follows.

```

1  config var n = 5,                // size of nxn grid
2      epsilon = 0.00001,          // convergence tolerance
3      verbose = false;            // control for amount of output

5  def main() {
6      const ProblemSpace = [1..n, 1..n], // domain for interior points
7      BigDomain = [0..n+1, 0..n+1]; // domain with boundary points

9      var X, XNew: [BigDomain] real = 0.0; // X holds approximate solution
10                                         // XNew is work array

12     X[n+1, 1..n] = 1.0;

14     if (verbose) {
15         writeln("Initial configuration:");
16         writeln(X, "\n");
17     }

19     var iteration = 0,                // iteration counter
20         delta: real;                // convergence measure

22     do {
23         forall (i,j) in ProblemSpace do
24             XNew(i,j) = (X(i-1,j) + X(i+1,j) + X(i,j-1) + X(i,j+1)) / 4.0;

26         delta = max reduce abs(XNew[ProblemSpace] - X[ProblemSpace]);
27         X[ProblemSpace] = XNew[ProblemSpace];

29         iteration += 1;

31         if (verbose) {
32             writeln("iteration: ", iteration);
33             writeln(X);
34             writeln("delta: ", delta, "\n");
35         }
36     } while (delta > epsilon);

38     writeln("Jacobi computation complete.");
39     writeln("Delta is ", delta, " (< epsilon = ", epsilon, ")");
40     writeln("# of iterations: ", iteration);
41 }

```

Execution and Output Compiling and running this program gives the following output:

```

> ./a.out
Jacobi computation complete.
Delta is 9.92124e-06 (< epsilon = 1e-05)
# of iterations: 60

```

It is possible to run a different sized problem, to use a different convergence tolerance, or enable more output without recompiling this program. There are three variables defined in lines 1 - 3, *n*, *epsilon* and *verbose*, which are *configuration variables* and can be set at the time of program execution through command line switches. Executing the following command line sequence,

```
> ./a.out --verbose=true --n=2 --epsilon=0.01
```

results in overriding the default values for `verbose`, `n` and `epsilon`, producing the following output.

```
Initial configuration:
0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0
0.0 1.0 1.0 0.0

iteration: 1
0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0
0.0 0.25 0.25 0.0
0.0 1.0 1.0 0.0
delta: 0.25

iteration: 2
0.0 0.0 0.0 0.0
0.0 0.0625 0.0625 0.0
0.0 0.3125 0.3125 0.0
0.0 1.0 1.0 0.0
delta: 0.0625

iteration: 3
0.0 0.0 0.0 0.0
0.0 0.09375 0.09375 0.0
0.0 0.34375 0.34375 0.0
0.0 1.0 1.0 0.0
delta: 0.03125

iteration: 4
0.0 0.0 0.0 0.0
0.0 0.109375 0.109375 0.0
0.0 0.359375 0.359375 0.0
0.0 1.0 1.0 0.0
delta: 0.015625

iteration: 5
0.0 0.0 0.0 0.0
0.0 0.117188 0.117188 0.0
0.0 0.367188 0.367188 0.0
0.0 1.0 1.0 0.0
delta: 0.0078125

Jacobi computation complete.
Delta is 0.0078125 (< epsilon = 0.01)
# of iterations: 5
```

Implementation Details This example program begins with the declaration of the three configuration variables. Note that these variables do not contain a type in their declaration:

```
config var n = 5,
           epsilon = 0.00001,
           verbose = false;
```

Instead the types of these variables are inferred from their initial values: `n` is an integer, `epsilon` is a real type, and `verbose` is a boolean type. In Chapel, if a variable declaration contains an initialization expression, it is optional to include a type specification. More details about variable declarations are given in §8.1.

After the configuration variables are declared, the `main` function is defined. The first two lines of `main` define two arithmetic domains:

```
def main() {
  const ProblemSpace = [1..n, 1..n],
        BigDomain = [0..n+1, 0..n+1];
```

Both domains are declared to be `const`, indicating that the values for the domains remain constant during the execution of the program. They are defined using arithmetic sequences, and are, thus, arithmetic domains. Because the domains are initialized in their declaration, it is not necessary to specify that they are of domain type. In this case, `: domain(2)` was omitted from both domain declarations. `BigDomain` is essentially the `ProblemSpace` domain with additional boundary rows and columns. By defining `ProblemSpace` to be the interior points of the larger grid, `BigDomain`, the Jacobi computation can be cleanly specified in one line, line 24, for just the interior points, eliminating the need to write special case computations for the boundary points.

After the domains are declared, the arrays are declared using `BigDomain`, and the last row of `X` is set to one.

```
var X, XNew: [BigDomain] real = 0.0;

X[n+1, 1..n] = 1.0;
```

Because they are defined using `BigDomain`, the arrays `X` and `XNew` are of size $n+2 \times n+2$. They are declared to be of type `real` and all elements of the arrays are initialized to zero. When an array is assigned a scalar value, as in the initialization to zero, each element of the array is assigned the scalar value. In the following line, the notation `n+1, 1..n` indicates a *slice* of the array `X` is to be assigned the value one. Each entry in this slice, which is the last row of `X`, is set to one.

Lines 14 - 17 print the initial configuration of the problem, if `verbose` is set to true.

```
if (verbose) {
  writeln("Initial configuration:");
  writeln(X, "\n");
}
```

The `writeln` function outputs strings and values of variables that are passed as arguments, followed by a line return. When an array is passed as an argument, it is output in row-major order with linefeeds after each row. Including the character string, `"\n"` inserts an additional line return.

Lines 19 - 20 contain the remaining variable declarations:

```
var iteration = 0,
    delta: real;
```

The variable `delta` must be specified to be of type `real` since it does not have an initialization expression.

The computational loop in lines 22 - 36 performs the Jacobi method. This loop is executed until `delta` is less than or equal to `epsilon`. Each iteration of the `do while` loop computes the next approximate solution using a `forall` loop.


```

do{
  forall (i,j) in ProblemSpace do
    XNew(i,j) = (X(i-1,j) + X(i+1,j) + X(i,j-1) + X(i,j+1)) / 4.0;

```

This loop uses the `ProblemSpace` domain as an *iterator*. Iterators for loops provide a value or a set of values to be used in each iteration of the loop. In this case, each iteration is indexed by a *tuple* of indices (i, j) , which avoids the use of a nest of two loops, one loop for the index i and one loop for the index j . The compiler determines how this loop is made parallel, according to the default distribution for the `ProblemSpace` domain. Note that the `ProblemSpace` domain is the set of interior points of x 's domain. All of the references to x in this loop are defined, and do not go out of bounds.

After the next approximate solution is computed and stored in `XNew`, the change between `XNew` and `X`, for just the interior points, is calculated and stored in `delta`:

```

delta = max reduce abs(XNew[ProblemSpace] - X[ProblemSpace]);

```

The *built-in reduction*, `max reduce` computes the maximum value of the expression that follows it. In this case the expression is the computed array of absolute values of the difference between the `XNew` and `X` arrays. If `XNew` and `X` are distributed, then this reduction is computed in parallel accordingly. The use of `ProblemSpace` with these two arrays indicates that the slice of the arrays corresponding to the interior points are to be used in this calculation.

Next, `x` is updated with the new approximate solution and the iteration number is advanced.

```

X[ProblemSpace] = XNew[ProblemSpace];

iteration += 1;

```

The domain `ProblemSpace` is used in the assignment of `XNew` to `X`, updating only the slice of `X` corresponding to the interior points of the array.

The remaining part of the loop optionally provides output about the iteration.

```

if (verbose) {
  writeln("iteration: ", iteration);
  writeln(X);
  writeln("delta: ", delta, "\n");
}
while (delta > epsilon);

```

If `verbose` is set to true, then information about the current approximate solution and `delta` is output. Then, `delta` is compared to the convergence tolerance, `epsilon`. If `delta` is not small enough, then the `do while` loop continues. Otherwise, the loop exits.

The program ends with three lines of output.

```

writeln("Jacobi computation complete.");
writeln("Delta is ", delta, " (< epsilon = ", epsilon, ")");
writeln("# of iterations: ", iteration);
}

```

Information about the convergence is given, printing the values for `epsilon` and the iteration counter.

5.3.2 Matrix and Vector Norms

Description The following example contains a module `Norm` which provides `norm` functions that compute either a vector or matrix norm, depending on the rank of the array that is passed as the argument. There are four norm type options: 1-norm, 2-norm, infinity norm and Frobenius norm. For vectors, all four options are implemented in this module. For matrices, all options but the 2-norm are provided. The module uses a variable of enumerated type `normType` to indicate the choice of norm.

Chapel Features This example demonstrates the definition and use of modules, generic functions and function overloading. When the `Norm` module is used, the user may call `norm(x)` or `norm(x, normType)` where `x` is any array and `normType` is the *enumerated type* as defined in the `Norm` module. The `norm` function is *overloaded* with four separate function definitions for `norm`, based on the the rank of the input array, or the number of formal arguments used in the `norm` function call. Each of these four function definitions is a *generic function*, not specifying the type of the array argument `x`. Generic functions allow for code reuse and readability.

Program Listing The following program gives the definition of the `Norm` module followed by a module which tests `Norm`, demonstrating different calls to the `norm` function.

```

1  module Norm{
2    enum normType {norm1, norm2, normInf, normFrob};

4    def norm(x: [], p: normType) where x.rank == 1 {
5      // vector norm routine
6      select (p) {
7        when norm1 do return + reduce abs(x);
8        when norm2 do return sqrt(+ reduce (abs(x)*abs(x)));
9        when normInf do return max reduce abs(x);
10       when normFrob do return sqrt(+ reduce (abs(x)*abs(x)));
11       otherwise halt("Unexpected norm type");
12     }
13   }

15  def norm(x: [?D], p: normType) where x.rank == 2 {
16    // matrix norm routine
17    select (p) {
18      when norm1 do
19        return max reduce [j in D(2)] (+ reduce abs(x[D(1), j]));

21      when norm2 do
22        halt("Haven't implemented 2-norm for 2D arrays yet");

24      when normInf do
25        return max reduce [i in D(1)] (+ reduce abs(x[i, D(2)]));

27      when normFrob do return sqrt(+ reduce (abs(x)*abs(x)));

29      otherwise halt("Unexpected norm type");
30    }
31  }

33  def norm(x: [], p: normType) where x.rank > 2 {
34    compilerError("Norms not implemented for array ranks > 2D");
35  }

37  def norm(x: []) {
38    // default norm routine

```

```

39     select (x.rank) {
40         when 1 do return norm(x, norm2);
41         when 2 do return norm(x, normFrob);
42         otherwise compilerError("Norms not implemented for array ranks > 2D");
43     }
44 }
45 }

47 module TestNorm {
48     use Norm;

50     def testNorm(arr: []) {
51         // test all possible norms of arr
52         var testType = if (arr.rank == 1) then "vector" else "matrix";
53         writeln("Test of ", testType, " norms. Array = ");
54         writeln(arr);
55         writeln("1-norm = ", norm(arr, norm1));
56         if (arr.rank == 1) then
57             writeln("2-norm = ", norm(arr, norm2));
58         writeln("infinity norm = ", norm(arr, normInf));
59         writeln("frobenius norm = ", norm(arr, normFrob));
60         writeln("default norm = ", norm(arr));
61         writeln();
62     }

64     def main() {
65         // test vector norms:
66         const D1 = [1..4];
67         var a:[D1] real;
68         a = 2.0;
69         testNorm(a);

71         // test matrix norms:
72         const D2 = [1..2,1..2];
73         var b:[D2] real;
74         b = 2.0;
75         testNorm(b);
76     }
77 }

```

Execution and Output After the definition of the Norm module, a TestNorm module is defined in lines 47 - 77, giving an example of how the norm functions can be used in a program.

On line 48, use Norm, indicates that the Norm module is to be used when resolving functions in the TestNorm module. A testNorm function is defined, taking arr as an argument. Based on arr.rank, all of the valid norm options are tested, along with the generic norm(arr) function call. This testNorm function is called to test two arrays, a and b. The array a is a one-dimensional array, as defined by the domain D1. The array b is a two-dimensional array, as defined by the domain D2. Each array has a total of four elements, and each is initialized to 2.0.

The output of this program is given below. The first set of norms is computed for a which is a vector and the second set is computed for b which is a matrix. Even though the vector and matrix in this example contain the same number of elements with the same values, some of the computed norms are different between the vector and the matrix. Thus, different implementations of the norm functions are used to compute the norms, depending on whether the input is a vector or matrix.

Test of vector norms. Array =

```

2.0 2.0 2.0 2.0
1-norm = 8.0
2-norm = 4.0
infinity norm = 2.0
frobenius norm = 4.0
default norm = 4.0

```

```

Test of matrix norms.  Array =
2.0 2.0
2.0 2.0
1-norm = 4.0
infinity norm = 4.0
frobenius norm = 4.0
default norm = 4.0

```

Implementation Details The `Norm` module begins with defining `normType` to be an enumerated type with constant values `norm1`, `norm2`, `normInf`, `normFrob`. For three of the four `norm` function definitions (see lines 4, 15, 33), one of the formal arguments is of `normType`, indicating which type of norm is to be computed. If `norm` is called without the `normType` argument, then the default norm for vectors is the 2-norm (see line 40) and the default norm for matrices is the Frobenius norm (see line 41).

The first `norm` function is defined for vectors. This definition begins with line 4:

```
def norm(x: [], p: normType) where x.rank == 1 {
```

The function definition contains a `where x.rank == 1` clause, indicating that this version of `norm` is to be used when the input array `x` has just one dimension, and is thus, a vector. The arguments for this version of `norm` are `x` of generic array type and `p` of `normType`. Since all norm types for the vector case can be computed with whole array operations, there is no need to specify a domain type for `x`. It is good practice to specify `x` as an array type by using `[]` since the `norm` function is only defined for arrays. The compiler can detect errors if other, non-array types are passed as arguments to `norm`.

The body of the vector version of the `norm` function is a `select` statement on `p` in lines 6 - 12.

```

select (p) {
  when norm1 do return + reduce abs(x);
  when norm2 do return sqrt(+ reduce (abs(x)*abs(x)));
  when normInf do return max reduce abs(x);
  when normFrob do return sqrt(+ reduce (abs(x)*abs(x)));
  otherwise halt("Unexpected norm type");
}

```

For the first four cases of the `select` statement, when `p` is `norm1`, `norm2`, `normInf` or `normFrob` a reduction operator is used along with `abs(x)` to compute the value that is returned. The `abs` function is promoted over the `x` array, computing the absolute values of each entry of `x`. The 1-norm of a vector is the sum of the absolute values of the entries of `x`, computed with the sum reduction expression, `+ reduce abs(x)`. The 2-norm of a vector is the sum of the squares of the absolute values of the entries of `x`, computed with the sum reduction expression, `+ reduce (abs(x)*abs(x))`. The infinity norm of a vector is the maximum entry of `x` in absolute value, which can be computed with the maximum reduction expression `max reduce abs(x)`. The Frobenius norm of a vector is the same computation as the 2-norm of a vector.

The last case of the `select` statement, the `otherwise` clause, results in the program halting, outputting the string that is passed to the `halt` function indicating that `p` has an unexpected value.

The second `norm` function is defined for matrices. This definition begins with line 15:

```
def norm(x: [?D], p: normType) where x.rank == 2 {
```

This version of the function will be used when `x` is a matrix, that is the rank of the array `x` is 2. The arguments for this version of `norm` are `x` of generic array type and `p` of `normType`. For this matrix version, the domain of `x` is needed to express the norm computations. So, the domain type is *queried* and set to `D` with the expression `[?D]`.

The body of the matrix version of `norm` is a select statement on `p` in lines 17 - 30. The `norm1` case, the 1-norm of a matrix, is the maximum absolute column sum of the matrix.

```
  when norm1 do
    return max reduce [j in D(2)] (+ reduce abs(x[D(1), j]));
```

The return expression is the maximum reduction of a `forall` loop, which computes a sum reduction for each iteration. It uses `D(1)` and `D(2)`, which are the first and second dimension, respectively, of the domain `D`. The shorthand version of a `forall` loop is used, `[j in D(2)]`, indicating that for all `j in D(2)`, the sum of the absolute values across the rows, `+ reduce abs(x[D(1), j])`, is computed. Then, the maximum value of these `j` absolute column sums is computed with `max reduce`.

The `norm2` case, the 2-norm of a matrix, is not implemented. Since the 2-norm of a matrix A is the square root of the maximum eigenvalue of the matrix $A^T A$ (or $A^H A$ in the case where A is complex), this norm has not been included in this simple example code. This case will halt with a message indicating this unimplemented status.

The `normInf` case, the infinity norm of a matrix, is the maximum absolute row sum.

```
  when normInf do
    return max reduce [i in D(1)] (+ reduce abs(x[i, D(2)]));
```

For all `i in D(1)`, the sum of the absolute values across columns, `+ reduce abs(x[i, D(2)])` is computed. Then the maximum value of these `i` absolute row sums is computed with `max reduce`.

The `normFrob` case which is the Frobenius norm of a matrix, is the same computation as the 2-norm of a vector.

```
  when normFrob do return sqrt(+ reduce (abs(x)*abs(x)));
```

The Frobenius norm is the square root of the sum of the squares of absolute values of all entries in the matrix.

The final case of the select statement, the `otherwise` clause, halts with a message that the norm type is unexpected.

The third version of the `norm` function with arguments `x` and `p`, given in lines 33 - 35 is for the case where `x.rank > 2`. The module was designed to only compute norms of vectors and matrices. Calling the `norm` function with three-dimensional or higher arrays should give an error at compile time. In this case, an appropriate compiler error is given using the `compilerError` function.

```
  compilerError("Norms not implemented for array ranks > 2D");
```

For more information about user-defined compiler errors, see §6.5.

The final version of the `norm` function is given in lines 37 - 44.

```

def norm(x: []) {
  select (x.rank) {
    when 1 do return norm(x, norm2);
    when 2 do return norm(x, normFrob);
    otherwise compilerError("Norms not implemented for array ranks > 2D");
  }
}

```

This version has one formal argument, `x`, allowing the user to omit the norm type. It calls the other `norm` functions with default values for `normType`. For vectors, the default norm is defined to be the 2-norm. For matrices, the default norm is defined to be the Frobenius norm. This version of `norm` is a select statement on `x.rank`. Depending on the rank of `x`, the `norm` function is called with the appropriate default norm type or a compiler error is given.

5.3.3 Simple Producer-Consumer Program

Description The following example demonstrates a simple producer and consumer program. The *producer computation* sets the value of a sync variable and the *consumer computation* prints the value of the same variable.

Chapel Features The program contains a `begin` statement to implement two concurrent computations, and a `sync` variable to coordinate between the two. Sync variables have extra state associated with them to indicate whether they are logically *full* or *empty*. A sync variable is intended to be accessed by multiple concurrent computations, each of which may change the contents and state of the variable. A sync variable can be read when its state is full. Attempts to read an empty sync variable from one computation will suspend execution until another computation changes its state to full. Once a sync variable is able to be read, its state is atomically set to empty. Conversely, a sync variable can be written only when its state is empty. Attempts to write to a full sync variable will suspend until its state is empty. Once the variable is written to, the state is atomically set to full. In addition, there are functions that read and write sync variables which override this default behavior. More information about sync variables and their functions is in §22.7.2 and §22.7.3.

Program Listing The Chapel code for this example is given below.

```

1 use Time;

3 config var numIterations: int = 5,
4           sleepTime: uint = 2;

6 var s: sync int;

8 begin { // create consumer computation
9   for c in 1..numIterations do
10     writeln("consumer got ", s);
11 }

13 // producer computation
14 for p in 1..numIterations {
15   sleep(sleepTime);
16   s = p;
17 }

```

Execution and Output When this program executes, a consumer computation is created with the `begin` statement. However, the consumer computation cannot read `s` before it has been written to. So, execution begins with the producer computation which assigns the value 1 to `s` after calling the `sleep` function². Once `s` has the value 1, the consumer computation can read it and print it. Execution switches back and forth between the producer and consumer computation for `numIterations`.

Running the program with the default values for `numIterations` results in:

```
consumer got 1
consumer got 2
consumer got 3
consumer got 4
consumer got 5
```

Each line is printed after a delay of `sleepTime` seconds.

Since the variables `numIterations` and `sleepTime` are *configuration variables*, they can be reset at execution time. For example, to change the number of iterations to 10 and to change the number of seconds to sleep to 5, the following execution command may be used:

```
> a.out --numIterations=10 --sleepTime=5
```

Implementation Details The program uses the `sleep` function in line 15, which is provided in the standard Chapel module, `Time`. So, the first line, `use Time`, is needed to include the `Time` module when resolving function calls in this program.

Lines 3 - 4 give the declarations and initial default values for the two variables, `numIterations` and `sleepTime`. The `config` keyword in front of these two variable declarations indicates that these are configuration variables. Configuration variables can be set to override their default values at program execution time, through the use of command line switches.

In line 4, the variable `s` is declared to be a `sync` variable of type `int`. Since `s` is not initialized, its state is empty at the beginning of the program's execution. The variable `s` will be used to synchronize the exchange of data between the producer and consumer computations.

The remainder of the code defines the consumer and producer computations. The new computation that is created with the `begin` statement will be referred to as the consumer computation. The continuing computation will be referred to as the producer computation. Execution control switches between the producer and consumer computations as the state of the `sync` variable `s` changes when it is read and written to.

The `begin` statement in line 8 creates a new computation to execute the `for` loop in lines 9 - 10.

```
begin {
  for c in 1..numIterations do
    writeln("consumer got ", s);
}
```

²The call to the `sleep` function is used to mimic some amount of computational time. It is not necessary to use a `sleep` function when synchronizing between concurrent computations.

This loop is indexed by `c`, which iterates over the arithmetic sequence `1..numIterations`. The variable `c` is a new variable defined for the scope of the loop. Its type is inferred to be integer from the integer arithmetic sequence that follows. During each iteration of the `for` loop, the `writeln` function is called which outputs the string literal `"consumer got "` and the value of the sync variable `s`, followed by a line break. In order to print the value stored in `s`, the sync variable is read only when its state is full. Once `s` is successfully read, its state will be set to empty. The next iteration follows, and the consumer computation will suspend execution until the state of `s` is made full again. Since the consumer computation only reads `s`, it will be the producer computation that changes the states of `s` to full.

The producer computation executes the `for` loop in lines 14 - 17.

```
for p in 1..numIterations {
    sleep(sleepTime);
    s = p;
}
```

This loop is indexed by `p` which iterates over the same arithmetic sequence as the consumer computation. Like the index `c`, `p` is inferred to be an integer. During each iteration of this `for` loop, the `sleep` function is called with the argument `sleepTime`. This `sleep` function is provided in the `Time` standard module, and it causes the producer computation to `sleepTime` seconds. After returning from the `sleep` function, the producer computation assigns `s` to be the iteration number `p`. Because `s` is a sync variable, this assignment is executed only when the state of `s` is empty. Once `s` is written, its state is changed to full and the next iteration of the producer loop follows. Since the producer computation only writes `s`, it will be the consumer computation that changes the state of `s` to empty allowing the next iteration to write its iteration number to `s`.

The program terminates after `numIterations` of both the consumer and producer loops.

5.3.4 Generic Stack Implementations

Description Two implementations of a generic stack type, `Stack`, are given below. The first defines `Stack` to be a linked list of `MyNode` objects. In the second implementation, `Stack` is an array. Both implementations of `Stack` define the methods, `push`, `pop` and `isEmpty?`.

Chapel Features These examples demonstrate *classes* and *records*. Chapel classes and records are both structured data types containing fields and methods. Classes are reference types while records are value types. The examples also use the *unspecified type alias* `itemType` as the generic type for the items in the stack and define generic stack methods. To use either version, a type for `itemType` must be specified when `Stack` is instantiated. More information about type aliases can be found in §21.3.1. The array implementation of the generic stack demonstrates the association of arrays to domains. In this example, the array's size is doubled by reassigning its domain to one that is twice in size.

Sample Stack A simple example of how this generic stack can be used:

```
var stack1: Stack(string);
stack1.push("one");
stack1.push("two");
stack1.push("three");
writeln(stack1.pop());
writeln(stack1.pop());
writeln(stack1.pop());
```


In this simple example, the variable `stack1` is declared to be a stack of string type. By specifying `string` as an input, the default constructor for `Stack` will initialize the type alias for the generic stack to be a string. The strings, "one", "two" and "three" are pushed on `stack1`, and then three items are popped off of `stack1`, in the reverse order from how they were put on the stack. The output of this example is:

```
three
two
one
```

Linked List Implementation This implementation of a generic stack defines a *class*, `MyNode` and a *record*, `Stack`. In the following code, the reference pointers of `MyNode` objects are used to point to the top of the stack and to point between items in the stack, thus implementing the stack as a linked list.

```
1 class MyNode {
2   type itemType;
3   var item: itemType;
4   var next: MyNode(itemType);
5 }

7 record Stack {
8   type itemType;
9   var top: MyNode(itemType);

11  def push(item: itemType) {
12    top = MyNode(itemType, item, top);
13  }

15  def pop() {
16    if isEmpty? then
17      halt("attempt to pop an item off an empty stack");
18    var oldTop = top;
19    top = top.next;
20    return oldTop.item;
21  }

23  def isEmpty? return top == nil;
24 }
```

Linked List Implementation Details The code begins with a definition of the `MyNode` class in lines 1 - 5.

```
class MyNode {
  type itemType;
  var item: itemType;
  var next: MyNode(itemType);
```

Objects of type `MyNode` are used to store the generic items in the stack as a linked list. There are three fields in `MyNode`: `itemType`, `item` and `next`. `MyNode` objects are instantiated within `Stack` to have the same `itemType` as the `Stack`. The `item` field holds the data and `next` is a pointer to the next `MyNode` object in the linked list.

The `Stack` record contains two fields:

```
record Stack {
  type itemType;
  var top: MyNode(itemType);
```

When `Stack` is instantiated, a type is specified for the type alias, `itemType`. The `top` field is a pointer to the top of the stack, which is a `MyNode` object of `itemType`. When the stack is first instantiated, `top` is set to `nil`. To add and remove items from the stack, the `push` and `pop` methods are used.

The `push` method is given in lines 11 - 13. This method adds an item to the top of the stack by resetting `top` to point to a new `MyNode` object. This new `top` object stores the added item and sets its `next` field to the previous `top` object of the stack.

```
def push(item: itemType) {
    top = MyNode(itemType, item, top);
}
```

The default constructor for `MyNode` is called to create a new object with `itemType` and `top` fields of the `Stack` instance on which `pop` is called, and the formal argument `item` of the `pop` function call. To reference fields of an instance of a structured type, `this` is used. In this case, `this.top` and `this.itemType` are implicit in the uses of `top` and `itemType`.

The `pop` method is given in lines 15 - 21.

```
def pop() {
    if isEmpty? then
        halt("attempt to pop an item off an empty stack");
    var oldTop = top;
    top = top.next;
    return oldTop.item;
}
```

This method returns the item at the top of the stack and resets `top` to point to the next object in the stack. First, a call to the `Stack` method `isEmpty?` is made to determine if the stack is empty. If it is, the program halts with a message indicating that an attempt was made to pop an item off of an empty stack. Otherwise, `top` is reset, and the appropriate item is returned. In this method, `this.isEmpty?` and `this.top` are implicit when `isEmpty?` and `top` are used. The memory of the `oldTop` object is freed, through automatic garbage collection.

The `isEmpty?` method, which is used to check if the stack is empty in the `pop` method, is defined to be:

```
def isEmpty? return top == nil;
```

Array Implementation The following implementation of a generic stack uses the array, `data`, to store the generic items in the stack, and the counter, `numItems`, to track the number of items in the stack and to indicate the index of the top item of the stack. The number of items in the stack and the size of `data` are checked when a new item is pushed on the stack. If necessary, the size of `data` is doubled. This example demonstrates how the size of an array is increased by increasing its domain.

```
1 record Stack {
2     type itemType;
3     var numItems: int = 0;
4     var data: [1..2] itemType;

6     def push(item: itemType) {
7         var height = data.numElements;
8         if numItems == height then
9             data.domain = [1..height*2];
10            data(numItems+1) = item;
11            numItems += 1;
12    }
```

```

14  def pop() {
15      if isEmpty? then
16          halt("attempt to pop an item off an empty stack");
17          numItems -= 1;
18          return data(numItems+1);
19      }

21  def isEmpty? return numItems == 0;
22  }

```

Array Implementation Details There are three fields defined for this Stack record, `itemType`, an integer variable, `numItems` and an array `data` of `itemType`. When first instantiated, the stack is empty and `numItems` is initialized to zero. The data array is declared with an anonymous one-dimensional domain, `[1..2]`.

The same three methods, `push`, `pop`, and `isEmpty?` are defined for this array implementation. Like the linked list implementation, there are no explicit references to `this` when accessing fields and methods.

In `push`, lines 6 - 12, `item` is stored in the `data(numItems+1)` and the `numItems` counter is incremented.

```

def push(item: itemType) {
    var height = data.numElements;
    if numItems == height then
        data.domain = [1..height*2];
        data(numItems+1) = item;
        numItems += 1;
}

```

Before `item` can be added to the stack, the size of `data` must be checked to determine if the array's size needs to be increased to accomodate another item being added to the stack. The method `numElements` is predefined for arrays, returning the total number of elements in an array. The variable `height` is set to the total number of elements in `data`, which is the size of the current array allocated to store items in the stack. If the number of items in the stack, `numItems` equals `height`, then more storage in `data` is needed. To increase the size of an array, the size of its domain is increased. By resetting `data`'s domain to a domain of twice the size, the array `data` itself is now doubled in size, and more items can be pushed onto the stack.

In `pop`, the top item in the stack, as indicated by the `numItem` counter, is returned, if the stack is not empty.

```

def pop() {
    if isEmpty? then
        halt("attempt to pop an item off an empty stack");
        numItems -= 1;
        return data(numItems+1);
}

```

The `pop` method first checks to see if the stack is empty. If it is, the program halts indicating that there was an attempt to pop an item off of an empty stack. Otherwise, the `numItems` counter is decremented and `data(numItems+1)` is returned as the popped item.

The `isEmpty?` method checks to see if the stack is empty, that is if `numItems` equals zero.

```

def isEmpty? return numItems == 0;

```


6 Lexical Structure

This section describes the lexical components of Chapel programs.

6.1 Comments

Two forms of comments are supported. All text following the consecutive characters `//` and before the end of the line is in a comment. All text following the consecutive characters `/*` and before the consecutive characters `*/` is in a comment.

Comments, including the characters that delimit them, are ignored by the compiler. If the delimiters that start the comments appear within a string literal, they do not start a comment but rather are part of the string literal.

6.2 White Space

White-space characters are spaces, tabs, and new-lines. Aside from delimiting comments and tokens, they are ignored by the compiler.

6.3 Case Sensitivity

Chapel is a case sensitive language so identifiers that are identical except of the case of the characters are still different.

6.4 Tokens

Tokens include identifiers, keywords, literals, operators, and punctuation.

6.4.1 Identifiers

An identifier in Chapel is a sequence of characters that must start with a letter, lower-case or upper-case, or an underscore, and can include lower-case letters, upper-case letters, digits, the underscore, and the question mark.

Example. The following are legal identifiers:

```
x, x1e, xt3, isLegalChapelIdentifier?, legal_chapel_identifier
```

6.4.2 Keywords

The following keywords are reserved:

atomic	begin	bool	break	by
class	cobegin	complex	config	const
continue	def	distributed	do	domain
else	enum	false	for	forall
goto	if	imag	in	int
inout	iterator	let	locale	module
nil	of	on	ordered	otherwise
out	param	pragma	real	record
reduce	return	scan	select	seq
serial	single	sync	then	true
type	uint	union	use	var
when	where	while	yield	

6.4.3 Literals

Literal values for primitive types are described in 7.1.7.

6.4.4 Operators and Punctuation

The following special characters are interpreted by the syntax of the language specially:

symbols	use
= += -= *= /= **= %= &= = ^= &&= = #= <= >=	assignment
..	arithmetic sequences
...	variable argument lists
&& !	logical operators
& ^ ~ << >>	bitwise operators
== != <= >= < >	relational operators
+ - * / % **	arithmetic operators
#	sequence concatenation operator
:	types
;	statement separator
,	expression separator
.	member access
?	query types
" '	string delimiters

6.4.5 Grouping Tokens

The following braces are part of the Chapel language:

braces	use
()	parenthesization, function calls, and tuples
[]	domains, square tuples, forall expressions, and function calls
(/ /)	sequence literals
{ }	type scopes and blocks

6.5 User-Defined Compiler Errors

The special compiler error statement given by

```
compiler-error-statement :
  compilerError ( string-literal ) ;
```

invokes a compiler error if the function that the statement is located within may be called when the program is executed and the statement is not eliminated by parameter folding.

The compiler error is defined to be the string literal and it points to the spot in the Chapel program where the function containing the *compiler-error-statement* is called from.

7 Types

Chapel is a statically typed language with a rich set of types. These include a set of predefined primitive types, enumerated types, classes, records, unions, tuples, sequences, domains, and arrays. This section defines the primitive types, enumerated types, and type aliases.

Programmers can define their own enumerated types, classes, records, unions, and type aliases in type declaration statements summarized by the following syntax:

```
type-declaration-statement :  
  enum-declaration-statement  
  class-declaration-statement  
  record-declaration-statement  
  union-declaration-statement  
  type-alias-declaration-statement
```

Classes are discussed in §14. Records are discussed in §15. Unions are discussed in §16. Tuples are discussed in §17. Sequences are discussed in §18. Domains and arrays are discussed in §19.

7.1 Primitive Types

The primitive types include the following types: `bool`, `int`, `uint`, `real`, `complex`, `imag`, `string`, and `locale`. These primitive types are defined in this section except for the `locale` type which is defined in §23.1.1.

7.1.1 The Bool Type

Chapel defines a logical data type designated by the symbol `bool` with the two predefined values `true` and `false`.

The relational operators return values of `bool` type and the logical operators operate on values of `bool` type.

Some statements require expressions of `bool` type and Chapel supports a special conversion of values to `bool` type when used in this context (§9.1.6). For example, an integer can be used as the condition in a conditional statement. It is converted to `false` if it is zero, and otherwise, it is converted to `true`.

7.1.2 Signed and Unsigned Integral Types

The integral types can be parameterized by the number of bits used to represent them. The default signed integral type, `int`, and the default unsigned integral type, `uint`, are 32 bits.

The integral types and their ranges are given in the following table:

Type	Minimum Value	Maximum Value
<code>int(8)</code>	-128	127
<code>uint(8)</code>	0	255
<code>int(16)</code>	-32768	32767
<code>uint(16)</code>	0	65535
<code>int(32), int</code>	-2147483648	2147483647
<code>uint(32), uint</code>	0	4294967295
<code>int(64)</code>	-9223372036854775808	9223372036854775807
<code>uint(64)</code>	0	18446744073709551615

The unary and binary operators that are pre-defined over the integral types operate with 32- and 64-bit precision. Using these operators on integral types represented with fewer bits results in a coercion according to the rules defined in §9.1.

7.1.3 Real Types

Like the integral types, the real types can be parameterized by the number of bits used to represent them. The default real type, `real`, is 64 bits. The real types that are supported are machine-dependent, but usually include `real(32)` and `real(64)`, and sometimes include `real(128)`.

Arithmetic over real values follows the IEEE 754 standard.

7.1.4 Complex Types

Like the integral and real types, the complex types can be parameterized by the number of bits used to represent them. A complex number is composed of two real numbers so the number of bits used to represent a complex is twice the number of bits used to represent the real numbers. The default complex type, `complex`, is 128 bits; it consists of two 64-bit real numbers. The complex types that are supported are machine-dependent, but usually include `complex(64)` and `complex(128)`, and sometimes include `complex(256)`.

The real and imaginary components can be accessed via the methods `re` and `im`. The type of these components is `real`.

Example. Given a complex number `3.14+2.72i`, the expressions `c.re` and `c.im` refer to `3.14` and `2.72` respectively.

7.1.5 Imaginary Types

The imaginary types can be parameterized by the number of bits used to represent them. The default imaginary type, `imag`, is 64 bits. The imaginary types that are supported are machine-dependent, but usually include `imag(32)` and `imag(64)`, and sometimes include `imag(128)`.

Rationale. The imaginary type is included to avoid numeric instabilities and under-optimized code stemming from always coercing real values to complex values with a zero imaginary part.

7.1.6 The String Type

Strings are a primitive type designated by the symbol `string`. Their length is unbounded.

Characters in a string can be accessed via the `substring` method on strings. This method takes an integer i and returns the i th character in the string.

Example. The first character of a string `s` can be selected by the method call `s.substring(1)`.

7.1.7 Primitive Type Literals

Bool literals are designated by the following syntax:

```
bool-literal: one of
true false
```

Signed and unsigned integer literals are designated by the following syntax:

```
integer-literal:
  digits
  0 'x' hexadecimal-digits
  0 'b' binary-digits

digits:
  digit
  digit digits

digit: one of
  0 1 2 3 4 5 6 7 8 9

hexadecimal-digits:
  hexadecimal-digit
  hexadecimal-digit hexadecimal-digits

hexadecimal-digits: one of
  0 1 2 3 4 5 6 7 8 9 A B C D E F a b c d e f

binary-digits:
  binary-digit
  binary-digit binary-digits

binary-digits: one of
  0 1
```

Suffixes, like those in C, are not necessary. The type of an integer literal is the first type of the following that can hold the value of the digits: `int`, `int(64)`, `uint(64)`. Explicit conversions are necessary to change the type of the literal to another integer size.

Real literals are designated by the following syntax:

```
real-literal:
  digitsopt . digit digitsopt exponent-partopt

exponent-part:
  'e' signopt digits

sign: one of
  + -
```

The type of a real literal is `real`. Explicit conversions are necessary to change the type of the literal to another real size.

Note that real literals require that a digit follow the decimal point. This is necessary to avoid an ambiguity in interpreting `2.e+2` that arises if a method called `e` is defined on integers.

Imaginary literals are designated by the following syntax:

```
imaginary-literal :
  real-literal i
  integer-literal i
```

A complex number is specified by adding or subtracting an imaginary literal with a real literal. Alternatively, a 2-tuple literal of expressions of integer or real type can be cast to a complex. These expressions can be literals, but do not need to be. To create a complex literal or parameter, they must be literals or parameters.

Example. The following codes represent the same complex literal:

```
2.0i,    0.0+2.0i,    (0.0,2.0):complex.
```

String literals are designated by the following syntax:

```
string-literal :
  " characters "
  ' characters '

characters :
  character
  character characters

character :
  any-character
```

Implementation note. Strings are currently restricted to ASCII characters. In a future version of Chapel, strings will be defined over alphabets to allow for more exotic characters.

7.2 Enumerated Types

Enumerated types are declared with the following syntax:

```
enum-declaration-statement :
  enum identifier { enum-constant-list } ;

enum-constant-list
  enum-constant
  enum-constant , enum-constant-list

enum-constant :
  identifier init-partopt

init-part :
  = expression
```

An enumerated type defines a set of named constants. These are associated with parameters of integral type. Each enumerated type is a distinct type.

7.3 Class Types

The class type defines a type that contains variables and constants, called fields, and functions, called methods. Classes are defined in §14. The class type can also contain type aliases and parameters. Such a class is generic and is defined in §21.

7.4 Record Types

The record type is similar to a class type; the primary difference is that a record is a value rather than a reference. The difference between classes and records is elaborated on in §15.

7.5 Union Types

The union type defines a type that contains one of a set of variables. Like classes and records, unions may also define methods. Unions are defined in §16.

7.6 Tuple Types

A tuple is a light-weight record that consists of one or more anonymous fields. If all the fields are of the same type, the tuple is homogeneous. Tuples are defined in §17.

7.7 Sequence Types

A sequence defines an ordered set of values of some type. Sequences are defined in §18.

7.8 Domain and Array Types

Domains are index sets. Arrays are types that contain a set of zero or more elements all of the same type. The elements are referenced via indices that are in the domain that the array is declared over. Domains and arrays are defined in §19.

7.9 Type Aliases

Type aliases are declared with the following syntax:

```
type-alias-declaration-statement :
  type type-alias-declaration ;

type-alias-declaration :
  identifier type-partopt
  identifier type-partopt , type-alias-declaration

type-part :
  = type
```

A type alias is a symbol that aliases any type as specified in the *type-part*. A use of a type alias has the same meaning as using the type specified by *type-part* directly.

The *type-part* is optional in the definition of a class or record. Such a type alias is called an unspecified type alias. Classes and records that contain type aliases, specified or unspecified, are generic (§21.3.1).

8 Variables

A variable is a symbol that represents memory. Chapel is a statically-typed, type-safe language so every variable has a type that is known at compile-time and the compiler enforces that values assigned to the variable can be stored in that variable as specified by its type.

8.1 Variable Declarations

Variables are declared with the following syntax:

```

variable-declaration-statement :
  configopt variable-kind variable-declaration ;

variable-kind : one of
  param const var

variable-declaration-list :
  variable-declaration
  variable-declaration , variable-declaration-list

variable-declaration :
  identifier-list type-partopt initialization-part
  identifier-list type-part

identifier-list :
  identifier
  identifier , identifier-list

type-part :
  : type
  : synchronization-type type

initialization-part :
  = expression

```

A *variable-declaration-statement* is used to define one or more variables. If the statement is a top-level module statement, the variables are global; otherwise they are local. Global variables are discussed in §8.2. Local variables are discussed in §8.3.

The optional keyword `config` specifies that the variables are configuration variables, described in Section §8.5.

The *variable-kind* specifies whether the variables are parameters (`param`), constants (`const`), or regular variables (`var`). Parameters are compile-time constants whereas constants are runtime constants. Both levels of constants are discussed in §8.4.

Multiple variables can be defined in the same variable declaration list. All variables defined in the same *identifier-list* are defined to have the same type and initialization expression.

The *type-part* of a variable declaration specifies the type of the variable. It is optional if the *initialization-part* is specified. If the *type-part* is omitted, the type of the variable is inferred using local type inference described in §8.1.2.

The *initialization-part* of a variable declaration specifies an initial expression to assign to the variable. If the *initialization-part* is omitted, the variable is initialized to a default value described in §8.1.1.

8.1.1 Default Initialization

If a variable declaration has no initialization expression, a variable is initialized to the default value of its type. The default values are as follows:

Type	Default Value
bool	false
int(*)	0
uint(*)	0
real(*)	0.0
imag(*)	0.0i
complex(*)	0.0 + 0.0i
string	" "
enums	first enum constant
classes	nil
records	default constructed record
sequences	empty sequence
arrays	elements are default values
tuples	components are default values

8.1.2 Local Type Inference

If the type is omitted from a variable declaration, the type of the variable becomes the type of the initialization expression.

8.2 Global Variables

Variables declared in statements that are in a module but not in a function or block within that module are global variables. Global variables can be accessed anywhere within that module after the declaration of that variable. They can also be accessed in other modules that use that module.

8.3 Local Variables

Local variables are variables that are not global. Local variables are declared within block statements. They can only be accessed within the scope of that block statement (including all inner nested block statements and functions).

A local variable only exists during the execution of code that lies within that block statement. This time is called the lifetime of the variable. When execution has finished within that block statement, the local variable and the storage it represents is removed. Variables of class type are the sole exception. Constructors of class types create storage that is not associated with any scope. Such storage is managed automatically as discussed in §14.10.

8.4 Constants

Constants are divided into two categories: parameters, specified with the keyword `param`, are compile-time constants and constants, specified with the keyword `const`, are runtime constants.

8.4.1 Compile-Time Constants

A compile-time constant or parameter must have a single value that is known statically by the compiler. Parameters are restricted to primitive and enumerated types.

Parameters can be assigned expressions that are parameter expressions. Parameter expressions are restricted to the following constructs:

- Literals of primitive type.
- Parenthesized parameter expressions.
- Casts of parameter expressions to primitive or enumerated types.
- Applications of the unary operators `+`, `-`, `!`, and `~` on operands that are `bool` or integral parameter expressions.
- Applications of the binary operators `+`, `-`, `*`, `/`, `%`, `**`, `&&`, `||`, `!`, `&`, `|`, `^`, `~`, `<<`, `>>`, `==`, `!=`, `<=`, `>=`, `<`, and `>` on operands that are `bool` or integral parameter expressions.
- The conditional expression where the condition is a parameter and the then- and else-expressions are parameters.

There is an expectation that parameters will be expanded to more types and more operations, and that functions that return parameters will be introduced, in the future.

8.4.2 Runtime Constants

Constants, as opposed to parameters, do not have the restrictions that are associated with parameters. Constants can be any type. They require an initialization expression and contain the value of that expression throughout their lifetime.

Variables of class type that are constants are constant references. The fields of the class can be modified, but the variable always points to the object that it was initialized to reference.

8.5 Configuration Variables

If the keyword `config` precedes the keyword `var`, `const`, or `param`, the variable, constant, or parameter is called a configuration variable, configuration constant, or configuration parameter respectively. Such variables, constants, and parameters must be global.

The initialization of these variables can be set via implementation dependent means, such as command-line switches or environment variables. The initialization expression in the program is ignored if the initialization is alternatively set.

Configuration parameters are set during compilation time via compilation flags or other implementation dependent means.

Example. A configuration parameter is set via a compiler flag. It may be used to control the target that is being compiled. For example, the code

```
config param target: string = "XT3";
```

sets a string parameter `target` to "XT3". This can be checked to compile different code for this target.

9 Conversions

A conversion allows an expression of one type to be converted into another type. Conversions can be either implicit or explicit.

Implicit conversions can occur during an assignment (from the expression on the right-hand side to the variable on the left-hand side) or during a function call (from the actual expression to the formal argument). An implicit conversion does not require a cast.

Explicit conversions require a cast in the code. Casts are defined in §10.5. Explicit conversions are supported between more types than implicit conversions, but explicit conversions are not supported between all types.

9.1 Implicit Conversions

Implicit conversions are allowed between numeric types (§9.1.1), from enumerated types to numeric types (§9.1.2), between class types (§9.1.3), and between record types (§9.1.4). A special set of implicit conversions are allowed from compile-time constants of type `int` and `int(64)` to other smaller numeric types if the value is in the range of the smaller numeric type (§9.1.5). Lastly, implicit conversions are supported from integral and class types to `bool` in the context of a statement (§9.1.6).

9.1.1 Implicit Numeric Conversions

Let i , j , and k range over the constants 8, 16, 32, and 64 when parameterizing `int` and `uint` and over the constants 32, 64, and 128 when parameterizing `real`, `imag`, and `complex`. The implicit numeric conversions are as follows:

- From `bool` to `int(j)`, `uint(j)`, or `string`
- From `int(i)` to `int(j)`, `real(k)`, `complex(k)`, or `string` where $j > i$
- From `uint(i)` to `int(j)`, `uint(j)`, `real(k)`, `complex(k)`, or `string` where $j > i$
- From `real(i)` to `real(j)`, `complex(k)`, or `string` where $j > i$ and $k \geq 2i$
- From `imag(i)` to `imag(j)`, `complex(k)`, or `string` where $j > i$ and $k \geq 2i$
- From `complex(i)` to `complex(j)`, or `string` where $j > i$

The implicit numeric conversions do not result in any loss of information except for the conversions from any of the `int` and `uint` types to any of the `real` and `complex` types and from any of the `real`, `imag`, and `complex` types to `string` where there is a loss of precision.

9.1.2 Implicit Enumeration Conversions

An expression that is an enumerated type can be implicitly converted to any integral type as long as all of the constants defined by the enumerated type are within range of the integral type. It can also be implicitly converted to `string` where the string is the name of the enumerated constant.

9.1.3 Implicit Class Conversions

An expression of class type C can be implicitly converted to another class type D provided that C is derived from D .

9.1.4 Implicit Record Conversions

An expression of record type C can be implicitly converted to another record type D provided that C is derived from D .

9.1.5 Implicit Compile-Time Constant Conversions

The following two implicit conversions of parameters are supported:

- A parameter of type `int(32)` can be implicitly converted to `int(8)`, `int(16)`, or any unsigned integral type if the value of the parameter is within the range of the target type.
- A parameter of type `int(64)` can be implicitly converted to `uint(64)` if the value of the parameter is nonnegative.

9.1.6 Implicit Statement Bool Conversions

In the condition of an if-statement, while-loop, and do-while-loop, the following implicit conversions are supported:

- An expression of integral type is taken to be true if it is non-zero and is otherwise false.
- An expression of a class type is taken to be true if it is not nil and is otherwise false.

9.2 Explicit Conversions

The explicit conversions are a superset of the implicit conversions.

9.2.1 Explicit Numeric Conversions

Explicit conversions are allowed from any numeric type to any other numeric type, bool, or string, and vice versa.

9.2.2 Explicit Enumeration Conversions

Explicit conversions are allowed from any enumerated types to any numeric type, bool, or string, and vice versa.

9.2.3 Explicit Class Conversions

An expression of class type C can be explicitly converted to another class type D provided that C is derived from D or D is derived from C . In the event that D is derived from C , the runtime type of D must be a C .

9.2.4 Explicit Record Conversions

An expression of record type C can be explicitly converted to another record type D provided that C is derived from D . There are no explicit record conversions that are not also implicit record conversions.

10 Expressions

This section defines expressions in Chapel. For all expressions are described in §22.2.

The syntax for an expression is given by:

```
expression :
  literal-expression
  variable-expression
  member-access-expression
  call-expression
  query-expression
  cast-expression
  lvalue-expression
  unary-expression
  binary-expression
  let-expression
  if-expression
  forall-expression
```

10.1 Literal Expressions

A literal value for any of the built-in types is a literal expression. These are defined where the type is defined. The list of literal values is given by the following syntax:

```
literal-expression :
  bool-literal
  integer-literal
  real-literal
  imaginary-literal
  string-literal
  sequence-literal
  domain-literal
```

10.2 Variable Expressions

A use of a variable is itself an expression. The syntax of a variable expression is given by:

```
variable-expression :
  identifier
```

10.3 Call Expressions

The syntax to call a function is given by:

```
call-expression :
  expression ( named-expression-list )

named-expression-list :
  named-expression
  named-expression , named-expression-list

named-expression :
  expression
  identifier = expression
```

A *call-expression* is resolved to a particular function according to the algorithm for function resolution described in §13.8.

A *named-expression* is an expression that may be optionally named. The optional *identifier* represents a named actual argument described in §13.4.1.

10.3.1 Indexing Expressions

Indexing into arrays, sequences, tuples, and domains shares the same syntax of a call expression. Indexing, at its core, is nothing more than a call to the indexing function defined on these types.

10.3.2 Member Access Expressions

Member access expressions are call expressions to members of classes, records, or unions. The syntax for a member access is given by:

member-access-expression :
expression . *identifier*

The member access may be an access of a field or a function inside a class, record, or union.

10.4 The Query Expression

A query expression is used to query a type or value within a formal argument type expression. The syntax of a query expression is given by:

query-expression :
 ? *identifier*

Querying is restricted to querying the type of a formal argument, the element type of an formal argument that is an array, the domain of a formal argument that is an array, or the size of a primitive type.

Example. The following code defines a generic function where the type of the first parameter is queried and stored in the type alias τ and the domain of the second argument is queried and stored in the variable D :

```
def foo(x: ? $\tau$ , y: [?D]  $\tau$ ) {
  for i in D do
    y[i] = x;
}
```

The type alias τ is used to specify the element type of array y . Arrays passed to this function must have element type y . The body of the function iterates over the domain of y captured in variable D and assigns the value of argument x to each element in array y .

There is an expectation that query expressions will be allowed in more places in the future.

10.5 Casts

A cast is specified with the following syntax:

```
cast-expression :  
  expression : type
```

The expression is converted to the specified type. Except for the casts listed below, casts are restricted to valid explicit conversions (§9.2).

The following casts have special meanings and do not correspond to an explicit conversion:

- A cast to the keyword `seq` converts a tuple to a sequence as described in §18.14.
- A cast to a parameter expression of integral type converts a sequence to a tuple as described in §18.14.
- A cast from a 2-tuple to `complex` converts the 2-tuple into a complex where the first component becomes the real part and the second component becomes the imaginary part. The size of the complex is determined from the size of the components based on implicit conversions.
- A cast from any primitive type to a string literal that is a C-style format string creates a formatted string based on that format.

10.6 LValue Expressions

An *lvalue* is an expression that can be used on the left-hand side of an assignment statement, passed to a formal argument of a function that has `out` or `inout` intent, or returned by a variable function. Valid lvalue expressions include the following:

- Variable expressions.
- Member access expressions.
- Call expressions that are either setters or variable functions.
- Indexing expressions.
- Let expressions where the inner expression is an lvalue expression.
- Conditional expressions where the then- and else-expressions are lvalue expressions.

LValue expressions are given by the following syntax:

```
lvalue-expression :  
  variable-expression  
  member-access-expression  
  call-expression  
  let-expression  
  conditional-expression .
```

The syntax is more relaxed than the definition above. For example, not all *call-expressions* are lvalues.

10.7 Operator Precedence and Associativity

The following table summarizes the precedence of operators and their associativity. Operators listed earlier have higher precedence than those listed later.

operators	associativity	use
.	left	member access
() []	left	function call, index expression
:	left	cast
**	right	exponentiation
unary + - ~	right	sign and bitwise negation
* / %	left	multiply, divide, and modulus
+ -	left	plus and minus
&	left	bitwise and
^	left	bitwise xor
<< >>	left	shift left and shift right
	left	bitwise or
<= >= < >	left	ordered comparison
== !=	left	equality comparison
!	right	logical negation
&&	left	logical and
	left	logical or
#	left	sequence concatenation
..	left	arithmetic sequences
in	left	forall expressions
by	left	striding sequences
if	left	conditional expressions
reduce scan	left	reductions and scans
,	left	comma separated expressions

10.8 Operator Expressions

The application of operators to expressions is itself an expression. The syntax of a unary expression is given by:

```

unary-expression :
    unary-operator expression

unary-operator: one of
    + - ~ !

```

The syntax of a binary expression is given by:

```

binary-expression :
    expression binary-operator expression

binary-operator: one of
    + - * / % ** & | ^ << >> && || == != <= >= < > #

```

The operators are defined in subsequent sections.

10.9 Arithmetic Operators

This section describes the predefined arithmetic operators. These operators can be redefined over different types using operator overloading (§13.7).

All arithmetic operators are implemented over integral types of size 32 and 64 bits only. For example, adding two 8-bit integers is done by first converting them to 32-bit integers and then adding the 32-bit integers. The result is a 32-bit integer.

10.9.1 Unary Plus Operators

The unary plus operators are predefined as follows:

```
def +(a: int(32)): int(32)
def +(a: int(64)): int(64)
def +(a: uint(32)): uint(32)
def +(a: uint(64)): uint(64)
def +(a: real(32)): real(32)
def +(a: real(64)): real(64)
def +(a: real(128)): real(128)
def +(a: imag(32)): imag(32)
def +(a: imag(64)): imag(64)
def +(a: imag(128)): imag(128)
def +(a: complex(32)): complex(32)
def +(a: complex(64)): complex(64)
def +(a: complex(128)): complex(128)
```

For each of these definitions, the result is the value of the operand.

10.9.2 Unary Minus Operators

The unary minus operators are predefined as follows:

```
def -(a: int(32)): int(32)
def -(a: int(64)): int(64)
def -(a: uint(64))
def -(a: real(32)): real(32)
def -(a: real(64)): real(64)
def -(a: real(128)): real(128)
def -(a: imag(32)): imag(32)
def -(a: imag(64)): imag(64)
def -(a: imag(128)): imag(128)
def -(a: complex(32)): complex(32)
def -(a: complex(64)): complex(64)
def -(a: complex(128)): complex(128)
```

For each of these definitions that return a value, the result is the negation of the value of the operand. For integral types, this corresponds to subtracting the value from zero. For real and imaginary types, this corresponds to inverting the sign. For complex types, this corresponds to inverting the signs of both the real and imaginary parts.

It is an error to try to negate a value of type `uint(64)`. Note that negating a value of type `uint(32)` first converts the type to `int(64)` using an implicit conversion.

10.9.3 Addition Operators

The addition operators are predefined as follows:

```
def +(a: int(32), b: int(32)): int(32)
def +(a: int(64), b: int(64)): int(64)
def +(a: uint(32), b: uint(32)): uint(32)
def +(a: uint(64), b: uint(64)): uint(64)
def +(a: uint(64), b: int(64))
def +(a: int(64), b: uint(64))

def +(a: real(32), b: real(32)): real(32)
def +(a: real(64), b: real(64)): real(64)
def +(a: real(128), b: real(128)): real(128)

def +(a: imag(32), b: imag(32)): imag(32)
def +(a: imag(64), b: imag(64)): imag(64)
def +(a: imag(128), b: imag(128)): imag(128)

def +(a: complex(64), b: complex(64)): complex(64)
def +(a: complex(128), b: complex(128)): complex(128)
def +(a: complex(256), b: complex(256)): complex(256)

def +(a: real(32), b: imag(32)): complex(64)
def +(a: imag(32), b: real(32)): complex(64)
def +(a: real(64), b: imag(64)): complex(128)
def +(a: imag(64), b: real(64)): complex(128)
def +(a: real(128), b: imag(128)): complex(256)
def +(a: imag(128), b: real(128)): complex(256)

def +(a: real(32), b: complex(64)): complex(64)
def +(a: complex(64), b: real(32)): complex(64)
def +(a: real(64), b: complex(128)): complex(128)
def +(a: complex(128), b: real(64)): complex(128)
def +(a: real(128), b: complex(256)): complex(256)
def +(a: complex(256), b: real(128)): complex(256)

def +(a: imag(32), b: complex(64)): complex(64)
def +(a: complex(64), b: imag(32)): complex(64)
def +(a: imag(64), b: complex(128)): complex(128)
def +(a: complex(128), b: imag(64)): complex(128)
def +(a: imag(128), b: complex(256)): complex(256)
def +(a: complex(256), b: imag(128)): complex(256)
```

For each of these definitions that return a value, the result is the sum of the two operands.

It is a compile-time error to add a value of type `uint(64)` and a value of type `int(64)`.

Addition over a value of real type and a value of imaginary type produces a value of complex type. Addition of values of complex type and either real or imaginary types also produces a value of complex type.

10.9.4 Subtraction Operators

The subtraction operators are predefined as follows:

```
def -(a: int(32), b: int(32)): int(32)
def -(a: int(64), b: int(64)): int(64)
def -(a: uint(32), b: uint(32)): uint(32)
def -(a: uint(64), b: uint(64)): uint(64)
```

```

def -(a: uint(64), b: int(64))
def -(a: int(64), b: uint(64))

def -(a: real(32), b: real(32)): real(32)
def -(a: real(64), b: real(64)): real(64)
def -(a: real(128), b: real(128)): real(128)

def -(a: imag(32), b: imag(32)): imag(32)
def -(a: imag(64), b: imag(64)): imag(64)
def -(a: imag(128), b: imag(128)): imag(128)

def -(a: complex(64), b: complex(64)): complex(64)
def -(a: complex(128), b: complex(128)): complex(128)
def -(a: complex(256), b: complex(256)): complex(256)

def -(a: real(32), b: imag(32)): complex(64)
def -(a: imag(32), b: real(32)): complex(64)
def -(a: real(64), b: imag(64)): complex(128)
def -(a: imag(64), b: real(64)): complex(128)
def -(a: real(128), b: imag(128)): complex(256)
def -(a: imag(128), b: real(128)): complex(256)

def -(a: real(32), b: complex(64)): complex(64)
def -(a: complex(64), b: real(32)): complex(64)
def -(a: real(64), b: complex(128)): complex(128)
def -(a: complex(128), b: real(64)): complex(128)
def -(a: real(128), b: complex(256)): complex(256)
def -(a: complex(256), b: real(128)): complex(256)

def -(a: imag(32), b: complex(64)): complex(64)
def -(a: complex(64), b: imag(32)): complex(64)
def -(a: imag(64), b: complex(128)): complex(128)
def -(a: complex(128), b: imag(64)): complex(128)
def -(a: imag(128), b: complex(256)): complex(256)
def -(a: complex(256), b: imag(128)): complex(256)

```

For each of these definitions that return a value, the result is the value obtained by subtracting the second operand from the first operand.

It is a compile-time error to subtract a value of type `uint(64)` from a value of type `int(64)`, and vice versa.

Subtraction of a value of real type from a value of imaginary type, and vice versa, produces a value of complex type. Subtraction of values of complex type from either real or imaginary types, and vice versa, also produces a value of complex type.

10.9.5 Multiplication Operators

The multiplication operators are predefined as follows:

```

def *(a: int(32), b: int(32)): int(32)
def *(a: int(64), b: int(64)): int(64)
def *(a: uint(32), b: uint(32)): uint(32)
def *(a: uint(64), b: uint(64)): uint(64)
def *(a: uint(64), b: int(64))
def *(a: int(64), b: uint(64))

def *(a: real(32), b: real(32)): real(32)
def *(a: real(64), b: real(64)): real(64)
def *(a: real(128), b: real(128)): real(128)

```

```

def *(a: imag(32), b: imag(32)): real(32)
def *(a: imag(64), b: imag(64)): real(64)
def *(a: imag(128), b: imag(128)): real(128)

def *(a: complex(64), b: complex(64)): complex(64)
def *(a: complex(128), b: complex(128)): complex(128)
def *(a: complex(256), b: complex(256)): complex(256)

def *(a: real(32), b: imag(32)): imag(32)
def *(a: imag(32), b: real(32)): imag(32)
def *(a: real(64), b: imag(64)): imag(64)
def *(a: imag(64), b: real(64)): imag(64)
def *(a: real(128), b: imag(128)): imag(128)
def *(a: imag(128), b: real(128)): imag(128)

def *(a: real(32), b: complex(64)): complex(64)
def *(a: complex(64), b: real(32)): complex(64)
def *(a: real(64), b: complex(128)): complex(128)
def *(a: complex(128), b: real(64)): complex(128)
def *(a: real(128), b: complex(256)): complex(256)
def *(a: complex(256), b: real(128)): complex(256)

def *(a: imag(32), b: complex(64)): complex(64)
def *(a: complex(64), b: imag(32)): complex(64)
def *(a: imag(64), b: complex(128)): complex(128)
def *(a: complex(128), b: imag(64)): complex(128)
def *(a: imag(128), b: complex(256)): complex(256)
def *(a: complex(256), b: imag(128)): complex(256)

```

For each of these definitions that return a value, the result is the product of the two operands.

It is a compile-time error to multiply a value of type `uint(64)` and a value of type `int(64)`.

Multiplication of values of imaginary type produces a value of real type. Multiplication over a value of real type and a value of imaginary type produces a value of imaginary type. Multiplication of values of complex type and either real or imaginary types produces a value of complex type.

10.9.6 Division Operators

The division operators are predefined as follows:

```

def /(a: int(32), b: int(32)): int(32)
def /(a: int(64), b: int(64)): int(64)
def /(a: uint(32), b: uint(32)): uint(32)
def /(a: uint(64), b: uint(64)): uint(64)
def /(a: uint(64), b: int(64))
def /(a: int(64), b: uint(64))

def /(a: real(32), b: real(32)): real(32)
def /(a: real(64), b: real(64)): real(64)
def /(a: real(128), b: real(128)): real(128)

def /(a: imag(32), b: imag(32)): real(32)
def /(a: imag(64), b: imag(64)): real(64)
def /(a: imag(128), b: imag(128)): real(128)

def /(a: complex(64), b: complex(64)): complex(64)
def /(a: complex(128), b: complex(128)): complex(128)

```

```

def /(a: complex(256), b: complex(256)): complex(256)

def /(a: real(32), b: imag(32)): imag(32)
def /(a: imag(32), b: real(32)): imag(32)
def /(a: real(64), b: imag(64)): imag(64)
def /(a: imag(64), b: real(64)): imag(64)
def /(a: real(128), b: imag(128)): imag(128)
def /(a: imag(128), b: real(128)): imag(128)

def /(a: real(32), b: complex(64)): complex(64)
def /(a: complex(64), b: real(32)): complex(64)
def /(a: real(64), b: complex(128)): complex(128)
def /(a: complex(128), b: real(64)): complex(128)
def /(a: real(128), b: complex(256)): complex(256)
def /(a: complex(256), b: real(128)): complex(256)

def /(a: imag(32), b: complex(64)): complex(64)
def /(a: complex(64), b: imag(32)): complex(64)
def /(a: imag(64), b: complex(128)): complex(128)
def /(a: complex(128), b: imag(64)): complex(128)
def /(a: imag(128), b: complex(256)): complex(256)
def /(a: complex(256), b: imag(128)): complex(256)

```

For each of these definitions that return a value, the result is the quotient of the two operands.

It is a compile-time error to divide a value of type `uint(64)` by a value of type `int(64)`, and vice versa.

Division of values of imaginary type produces a value of real type. Division over a value of real type and a value of imaginary type produces a value of imaginary type. Division of values of complex type and either real or imaginary types produces a value of complex type.

10.9.7 Modulus Operators

The modulus operators are predefined as follows:

```

def %(a: int(32), b: int(32)): int(32)
def %(a: int(64), b: int(64)): int(64)
def %(a: uint(32), b: uint(32)): uint(32)
def %(a: uint(64), b: uint(64)): uint(64)
def %(a: uint(64), b: int(64))
def %(a: int(64), b: uint(64))

```

For each of these definitions that return a value, the result is the remainder when the first operand is divided by the second operand.

It is a compile-time error to take the remainder of a value of type `uint(64)` and a value of type `int(64)`, and vice versa.

There is an expectation that the predefined modulus operators will be extended to handle real, imaginary, and complex types in the future.

10.9.8 Exponentiation Operators

The exponentiation operators are predefined as follows:

```
def **(a: int(32), b: int(32)): int(32)
def **(a: int(64), b: int(64)): int(64)
def **(a: uint(32), b: uint(32)): uint(32)
def **(a: uint(64), b: uint(64)): uint(64)
def **(a: uint(64), b: int(64))
def **(a: int(64), b: uint(64))

def **(a: real(32), b: real(32)): real(32)
def **(a: real(64), b: real(64)): real(64)
def **(a: real(128), b: real(128)): real(128)
```

For each of these definitions that return a value, the result is the value of the first operand raised to the power of the second operand.

It is a compile-time error to take the exponent of a value of type `uint(64)` by a value of type `int(64)`, and vice versa.

There is an expectation that the predefined exponentiation operators will be extended to handle imaginary and complex types in the future.

10.10 Bitwise Operators

This section describes the predefined bitwise operators. These operators can be redefined over different types using operator overloading (§13.7).

10.10.1 Bitwise Complement Operators

The bitwise complement operators are predefined as follows:

```
def ~(a: bool): bool
def ~(a: int(32)): int(32)
def ~(a: int(64)): int(64)
def ~(a: uint(32)): uint(32)
def ~(a: uint(64)): uint(64)
```

For each of these definitions, the result is the bitwise complement of the operand.

10.10.2 Bitwise And Operators

The bitwise and operators are predefined as follows:

```
def &(a: bool, b: bool): bool
def &(a: int(32), b: int(32)): int(32)
def &(a: int(64), b: int(64)): int(64)
def &(a: uint(32), b: uint(32)): uint(32)
def &(a: uint(64), b: uint(64)): uint(64)
def &(a: uint(64), b: int(64))
def &(a: int(64), b: uint(64))
```


For each of these definitions that return a value, the result is computed by applying the logical and operation to the bits of the operands.

It is a compile-time error to apply the bitwise and operator to a value of type `uint(64)` and a value of type `int(64)`, and vice versa.

10.10.3 Bitwise Or Operators

The bitwise or operators are predefined as follows:

```
def |(a: bool, b: bool): bool
def |(a: int(32), b: int(32)): int(32)
def |(a: int(64), b: int(64)): int(64)
def |(a: uint(32), b: uint(32)): uint(32)
def |(a: uint(64), b: uint(64)): uint(64)
def |(a: uint(64), b: int(64))
def |(a: int(64), b: uint(64))
```

For each of these definitions that return a value, the result is computed by applying the logical or operation to the bits of the operands.

It is a compile-time error to apply the bitwise or operator to a value of type `uint(64)` and a value of type `int(64)`, and vice versa.

10.10.4 Bitwise Xor Operators

The bitwise xor operators are predefined as follows:

```
def ^(a: bool, b: bool): bool
def ^(a: int(32), b: int(32)): int(32)
def ^(a: int(64), b: int(64)): int(64)
def ^(a: uint(32), b: uint(32)): uint(32)
def ^(a: uint(64), b: uint(64)): uint(64)
def ^(a: uint(64), b: int(64))
def ^(a: int(64), b: uint(64))
```

For each of these definitions that return a value, the result is computed by applying the XOR operation to the bits of the operands.

It is a compile-time error to apply the bitwise xor operator to a value of type `uint(64)` and a value of type `int(64)`, and vice versa.

10.11 Shift Operators

This section describes the predefined shift operators. These operators can be redefined over different types using operator overloading (§13.7).

The shift operators are predefined as follows:

```

def <<(a: int(32), b): int(32)
def >>(a: int(32), b): int(32)
def <<(a: int(64), b): int(64)
def >>(a: int(64), b): int(64)
def <<(a: uint(32), b): uint(32)
def >>(a: uint(32), b): uint(32)
def <<(a: uint(64), b): uint(64)
def >>(a: uint(64), b): uint(64)

```

The type of the second actual argument must be any integral type.

The << operator shifts the bits of *a* left by the integer *b*. The new low-order bits are set to zero.

The >> operator shifts the bits of *a* right by the integer *b*. When *a* is negative, the new high-order bits are set to one; otherwise the new high-order bits are set to zero.

The value of *b* must be non-negative.

10.12 Logical Operators

This section describes the predefined logical operators. These operators can be redefined over different types using operator overloading (§13.7).

10.12.1 The Logical Negation Operator

The logical negation operator is predefined as follows:

```
def !(a: bool): bool
```

The result is the logical negation of the operand.

10.12.2 The Logical And Operator

The logical and operator is predefined over `bool` type. It returns `true` if both operands evaluate to `true`; otherwise it returns `false`. If the first operand evaluates to `false`, the second operand is not evaluated and the result is `false`.

The logical and operator over expressions *a* and *b* given by

```
a && b
```

is evaluated as the expression

```
if a.false? then false else b.true?
```

The methods `false?` and `true?` are predefined over `bool` type as follows:

```
def bool.true? return this;
def bool.false? return !this;
```

Overloading the logical and operator over other types is accomplished by overloading the `false?` and `true?` methods over other types.

10.12.3 The Logical Or Operator

The logical or operator is predefined over `bool` type. It returns true if either operand evaluate to true; otherwise it returns false. If the first operand evaluates to true, the second operand is not evaluated and the result is true.

The logical or operator over expressions `a` and `b` given by

```
a && b
```

is evaluated as the expression

```
if a.true? then true else b.true?
```

The methods `false?` and `true?` are predefined over `bool` type as described in §10.12.2. Overloading the logical or operator over other types is accomplished by overloading the `false?` and `true?` methods over other types.

10.13 Relational Operators

This section describes the predefined relational operators. These operators can be redefined over different types using operator overloading (§13.7).

10.13.1 Ordered Comparison Operators

The “less than” comparison operators are predefined as follows:

```
def <(a: int(32), b: int(32)): bool
def <(a: int(64), b: int(64)): bool
def <(a: uint(32), b: uint(32)): bool
def <(a: uint(64), b: uint(64)): bool
def <(a: real(32), b: real(32)): bool
def <(a: real(64), b: real(64)): bool
def <(a: real(128), b: real(128)): bool
def <(a: imag(32), b: imag(32)): bool
def <(a: imag(64), b: imag(64)): bool
def <(a: imag(128), b: imag(128)): bool
```

The result of `a < b` is true if `a` is less than `b`; otherwise the result is false.

The “greater than” comparison operators are predefined as follows:

```
def >(a: int(32), b: int(32)): bool
def >(a: int(64), b: int(64)): bool
def >(a: uint(32), b: uint(32)): bool
def >(a: uint(64), b: uint(64)): bool
def >(a: real(32), b: real(32)): bool
def >(a: real(64), b: real(64)): bool
def >(a: real(128), b: real(128)): bool
def >(a: imag(32), b: imag(32)): bool
def >(a: imag(64), b: imag(64)): bool
def >(a: imag(128), b: imag(128)): bool
```

The result of `a > b` is true if `a` is greater than `b`; otherwise the result is false.

The “less than or equal to” comparison operators are predefined as follows:

```
def <=(a: int(32), b: int(32)): bool
def <=(a: int(64), b: int(64)): bool
def <=(a: uint(32), b: uint(32)): bool
def <=(a: uint(64), b: uint(64)): bool
def <=(a: real(32), b: real(32)): bool
def <=(a: real(64), b: real(64)): bool
def <=(a: real(128), b: real(128)): bool
def <=(a: imag(32), b: imag(32)): bool
def <=(a: imag(64), b: imag(64)): bool
def <=(a: imag(128), b: imag(128)): bool
```

The result of `a <= b` is true if `a` is less than or equal to `b`; otherwise the result is false.

The “greater than or equal to” comparison operators are predefined as follows:

```
def >=(a: int(32), b: int(32)): bool
def >=(a: int(64), b: int(64)): bool
def >=(a: uint(32), b: uint(32)): bool
def >=(a: uint(64), b: uint(64)): bool
def >=(a: real(32), b: real(32)): bool
def >=(a: real(64), b: real(64)): bool
def >=(a: real(128), b: real(128)): bool
def >=(a: imag(32), b: imag(32)): bool
def >=(a: imag(64), b: imag(64)): bool
def >=(a: imag(128), b: imag(128)): bool
```

The result of `a >= b` is true if `a` is greater than or equal to `b`; otherwise the result is false.

10.13.2 Equality Comparison Operators

The equality comparison operators are predefined over `bool` and the numeric types as follows:

```
def ==(a: int(32), b: int(32)): bool
def ==(a: int(64), b: int(64)): bool
def ==(a: uint(32), b: uint(32)): bool
def ==(a: uint(64), b: uint(64)): bool
def ==(a: real(32), b: real(32)): bool
def ==(a: real(64), b: real(64)): bool
def ==(a: real(128), b: real(128)): bool
def ==(a: imag(32), b: imag(32)): bool
def ==(a: imag(64), b: imag(64)): bool
def ==(a: imag(128), b: imag(128)): bool
def ==(a: complex(64), b: complex(64)): bool
def ==(a: complex(128), b: complex(128)): bool
def ==(a: complex(256), b: complex(256)): bool
```

The result of `a == b` is true if `a` and `b` contain the same value; otherwise the result is false. The result of `a != b` is equivalent to `!(a == b)`.

The equality comparison operators are predefined over classes as follows:

```
def ==(a: object, b: object): bool
def !=(a: object, b: object): bool
```

The result of `a == b` is true if `a` and `b` reference the same storage location; otherwise the result is false. The result of `a != b` is equivalent to `!(a == b)`.

Default equality comparison operators are generated for records if the user does not define them. These operators are described in §15.3.

The equality comparison operators are predefined over strings as follows:

```
def ==(a: string, b: string): bool
def !=(a: string, b: string): bool
```

The result of `a == b` is true if the sequence of characters in `a` matches exactly the sequence of characters in `b`; otherwise the result is false. The result of `a != b` is equivalent to `!(a == b)`.

10.14 Miscellaneous Operators

This section describes several miscellaneous operators. These operators can be redefined over different types using operator overloading (§13.7).

10.14.1 The String Concatenation Operator

The string concatenation operator is predefined as follows:

```
def +(a: string, b: string): string
```

The result is the concatenation of `a` followed by `b`.

Example. Since integers can be implicitly converted to strings, an integer can be appended to a string using the string concatenation operator. The code

```
"result: "+i
```

where `i` is an integer appends the value of `i` to the string literal. If `i` is 3, then the resulting string would be `"result: 3"`.

10.14.2 The Sequence Concatenation Operator

The operator `#` is predefined on sequences and their element types. It is described in §18.6.

10.14.3 The Arithmetic Domain By Operator

The operator `by` is predefined on arithmetic domains. It is described in §19.3.3.

10.14.4 The Arithmetic Sequence By Operator

The operator `by` is predefined on arithmetic sequences. It is described in §18.13.1.

10.15 Let Expressions

A let expression allows variables to be declared at the expression level and used within that expression. The syntax of a let expression is given by:

```
let-expression :  
  let variable-declaration-list in expression
```

The scope of the variables is the let-expression.

Example. Let expressions are useful for defining variables in the context of expression. In the code

```
let x: real = a*b, y = x*x in 1/y
```

the value determined by `a*b` is computed and converted to type `real` if it is not already a `real`. The square of the `real` is then stored in `y` and the result of the expression is the reciprocal of that value.

10.16 Conditional Expressions

A conditional expression is given by the following syntax:

```
conditional-expression :  
  if expression then expression else expression  
  if expression then expression
```

The conditional expression is evaluated in two steps. First, the expression following the `if` keyword is evaluated. Then, if the expression evaluated to true, the expression following the `then` keyword is evaluated and taken to be the value of this expression. Otherwise, the expression following the `else` keyword is evaluated and taken to be the value of this expression. In both cases, the unselected expression is not evaluated.

The ‘else’ keyword can be omitted only when the conditional expression is immediately nested inside a `forall` expression. Such an expression is used to filter predicates as described in §18.11.

11 Statements

Chapel is an imperative language with statements that may have side effects. Statements allow for the sequencing of program execution. They are as follows:

```

statement :
  block-statement
  expression-statement
  conditional-statement
  select-statement
  while-do-statement
  do-while-statement
  for-statement
  param-for-statement
  return-statement
  yield-statement
  module-declaration-statement
  function-declaration-statement
  type-declaration-statement
  variable-declaration-statement
  use-statement
  type-select-statement
  empty-statement
  cobegin-statement
  begin-statement
  serial-statement
  atomic-statement
  on-statement

```

The declaration statements are discussed in the sections that define what they declare. Module declaration statements are defined in §12. Function declaration statements are defined in §13. Type declaration statements are defined in §7. Variable declaration statements are defined in §8. Return statements are defined in §13.2. Yield statements are defined in §20.2.

The *cobegin-statement* is defined in §22.3. The *begin-statement* is defined in §22.4. The *serial-statement* is defined in §22.6. The *atomic-statement* is defined in §22.9. The *on-statement* is defined in §23.2.1.

11.1 Blocks

A block is a statement or a possibly empty list of statements that form their own scope. A block is given by

```

block-statement :
  { statementsopt }
  { }

statements :
  statement
  statement statements

```

Variables defined within a block are local variables (§8.3).

The statements within a block are executed serially unless the block is in a *cobegin* statement (§22.3).

11.2 Block Level Statements

A block level statement is a category of statement that is sometimes called for by the language syntax. A block level statement is given by

```
block-level-statement :
    block-statement
    conditional-statement
    select-statement
    while-do-statement
    for-statement
    param-for-statement
    return-statement
    yield-statement
    type-select-statement
    empty-statement
    cobegin-statement
    begin-statement
    serial-statement
    atomic-statement
    on-statement
```

Block level statements are part of the language to avoid the excessive and unnecessary use of curly brackets. For example, function bodies are not required to be blocks, but must be block level statements.

11.3 Expression Statements

The expression statement evaluates an expression solely for side effects. The syntax for an expression statement is given by

```
expression-statement :
    expression ;
```

11.4 Assignment Statements

An assignment statement assigns the value of an expression to another expression that can appear on the left-hand side of the operator, for example, a variable. Assignment statements are given by

```
assignment-statement :
    lvalue-expression assignment-operator expression

assignment-operator: one of
    = += -= *= /= %= **= &= |= ^= &&= ||= #= <<= >>=
```

The expression on the right-hand side of the assignment operator is evaluated first; it can be any expression. The expression on the left hand side must be a valid lvalue (§10.6). It is evaluated second and then assigned the value.

The assignment operators that contain a binary operator as a prefix is a short-hand for applying the binary operator to the left and right-hand side expressions and then assigning the value of that application to the already evaluated left-hand side. Thus, for example, `x += y` is equivalent to `x = x + y` where the expression `x` is evaluated once.

Values of one primitive or enumerated type can be assigned to another primitive or enumerated type if an implicit coercion exists between those types (§9.1).

The validity and semantics of assigning between classes (§14.2), records (§15.2.3), unions (§16.2), tuples (§17.3), sequences (§18.4), domains (§19.1.3), and arrays (§19.2.4) is discussed in these later sections.

11.5 The Conditional Statement

The conditional statement allows execution to choose between two statements based on the evaluation of an expression of `bool` type. The syntax for a conditional statement is given by

```
conditional-statement :
  if expression then statement else-partopt
  if expression block-level-statement else-partopt

else-part :
  else statement
```

A conditional statement evaluates an expression of `bool` type. If the expression evaluates to true, the first statement in the conditional statement is executed. If the expression evaluates to false and the optional else-clause exists, the statement following the `else` keyword is executed.

If the expression is a parameter, the conditional statement is folded by the compiler. If the expression evaluates to true, the first statement replaces the conditional statement. If the expression evaluates to false, the second statement, if it exists, replaces the conditional statement; if the second statement does not exist, the conditional statement is removed.

If the statement that immediately follows the optional `then` keyword is a conditional statement and it is not in a block, the else-clause is bound to the nearest preceding conditional statement without an else-clause.

Each statement embedded in the *conditional-statement* has its own scope whether or not an explicit block surrounds it.

11.6 The Select Statement

The select statement is a multi-way variant of the conditional statement. The syntax is given by:

```
select-statement :
  select expression { when-statements }

when-statements :
  when-statement
  when-statement when-statements

when-statement :
  when expression-list do statement
  when expression-list block-level-statement
  otherwise statement

expression-list :
  expression
  expression , expression-list
```

The expression that follows the keyword `select`, the select expression, is compared with the list of expressions following the keyword `when`, the case expressions, using the equality operator `==`. If the expressions cannot be compared with the equality operator, a compile-time error is generated. The first case expression that contains an expression where that comparison is `true` will be selected and control transferred to the associated statement. If the comparison is always `false`, the statement associated with the keyword `otherwise`, if it exists, will be selected and control transferred to it. There may be at most one `otherwise` statement and its location within the select statement does not matter.

Each statement embedded in the *when-statement* has its own scope whether or not an explicit block surrounds it.

11.7 The While and Do While Loops

There are two variants of the while loop in Chapel. The syntax of the while-do loop is given by:

```
while-do-statement :
    while expression do statement
    while expression block-level-statement
```

The syntax of the do-while loop is given by:

```
do-while-statement :
    do statement while expression ;
```

In both variants, the expression evaluates to a value of type `bool` which determines when the loop terminates and control continues with the statement following the loop.

The while-do loop is executed as follows:

1. The expression is evaluated.
2. If the expression evaluates to `false`, the statement is not executed and control continues to the statement following the loop.
3. If the expression evaluates to `true`, the statement is executed and control continues to step 1, evaluating the expression again.

The do-while loop is executed as follows:

1. The statement is executed.
2. The expression is evaluated.
3. If the expression evaluates to `false`, control continues to the statement following the loop.
4. If the expression evaluates to `true`, control continues to step 1 and the the statement is executed again.

In this second form of the loop, note that the statement is executed unconditionally the first time.

11.8 The For Loop

The for loop iterates over sequences, domains, arrays, iterators, or any class that implements the structural iterator interface. The syntax of the for loop is given by:

```
for-statement :
  for index-expression in iterator-expression do statement
  for index-expression in iterator-expression block-level-statement

index-expression :
  expression

iterator-expression :
  expression
```

The index-expression can be an identifier or a tuple of identifiers. The identifiers are declared to be new variables for the scope of this statement.

If the iterator-expression is a tuple, the components of the tuple must support iteration, e.g., a tuple of arrays, and those components are iterated over using a zipper iteration defined in §11.8.1. If the iterator-expression is a tuple delimited by square brackets, the components of the tuple must support iteration and these components are iterated over using a tensor product iteration defined in §11.8.2.

11.8.1 Zipper Iteration

When multiple iterators are iterated over in a zipper context, on each iteration, each expression is iterated over, the values are returned by the iterators in a tuple and assigned to the index, and the statement is executed.

The shape of each iterator, the rank and the extents in each dimension, must be identical.

Example. The output of

```
for (i, j) in (1..3, 4..6) do
  write(i, " ", j, " ");
```

is “1 4 2 5 3 6”.

11.8.2 Tensor Product Iteration

When multiple iterators are iterated over in a tensor product context, they are iterated over as if they were nested in distinct for loops. There is no constraint on the iterators as there is in the zipper context.

Example. The output of

```
for (i, j) in [1..3, 4..6] do
  write(i, " ", j, " ");
```

is “1 4 1 5 1 6 2 4 2 5 2 6 3 4 3 5 3 6”. The statement is equivalent to

```
for i in 1..3 do
  for j in 4..6 do
    write(i, " ", j, " ");
```

11.8.3 Parameter For Loops

Parameter for loops are unrolled by the compiler so that the index variable is a parameter rather than a variable. The syntax for a parameter for loop statement is given by:

```
param-for-statement :
    for param identifier in arithmetic-sequence do statement
    for param identifier in arithmetic-sequence block-level-statement
```

Parameter for loops are restricted to iteration over arithmetic sequence literals the bounds of which must also be parameters. The loop is then unrolled for each iteration.

11.9 The Use Statement

The use statement makes symbols in a module available without accessing them via the module name. The syntax of the use statement is given by:

```
use-statement :
    use module-name ;

module-name :
    identifier
    module-name . module-name
```

The use statement makes symbols in the module's scope available in the scope where the use statement occurs.

It is an error for a variable, type or module to be defined both by a use statement and by a declaration in the same scope. Functions may be overloaded in this way.

11.10 The Type Select Statement

A type select statement has two uses. It can be used to determine the type of a union, as discussed in §16.3. In its more general form, it can be used to determine the types of one or more values using the same mechanisms used to disambiguate function definitions. Its syntax is given by:

```
type-select-statement :
    type select expression-list { type-when-statements }

type-when-statements :
    type-when-statement
    type-when-statement type-when-statements

type-when-statement :
    when type-list do statement
    when type-list block-level-statement
    otherwise statement

expression-list :
    expression
    expression , expression-list

type-list :
    type
    type , type-list
```

Call the expressions following the keyword `select`, the select expressions. The number of select expressions must be equal to the number of types following each of the `when` keywords. Like the `select` statement, one of the statements associated with a `when` will be executed. In this case, that statement is chosen by the function resolution mechanism. The select expressions are the actual arguments, the types following the `when` keywords are the types of the formal arguments for different anonymous functions. The function that would be selected by function resolution determines the statement that is executed. If none of the functions are chosen, the the statement associated with the keyword `otherwise`, if it exists, will be selected.

As with function resolution, this can result in an ambiguous situation. Unlike with function resolution, in the event of an ambiguity, the first statement in the list of `when` statements is chosen.

11.11 The Empty Statement

An empty statement has no effect. The syntax of an empty statement is given by

```
empty-statement :  
    ;
```


12 Modules

Chapel supports modules to manage name spaces. Every symbol, including variables, functions, and types, are associated with some module.

Module definitions are described in §12.1. A program consists of one or more modules. The execution of a program and command-line arguments are described in §12.2. Module uses and explicit naming of symbols is described in §12.3. Nested modules are described in §12.4. The relation between files and modules is described in §12.5.

12.1 Module Definitions

A module is declared with the following syntax:

```
module-declaration-statement :  
  module identifier block-statement
```

A module's name is specified after the module keyword. The *block-statement* opens the module's scope. Symbols defined in this block statement are defined in the module's scope.

Module declaration statements may only be top-level statements in files or top-level statements in other modules. A module that is declared in another module is called a nested module (§12.4).

12.2 Program Execution

Chapel programs start by executing the main function (§12.2.1). The main function takes no arguments but command-line arguments can be passed to a program via a global sequence of strings called `argv` (§12.2.2). Command-line flags can be passed to a program via configuration variables, as discussed in §8.5.

12.2.1 The *main* Function

The main function must be called `main` and must have zero arguments. There can be only one main function in all of the modules that make up a program. Every main function starts by using the module that it is defined in, and thus executing the top-level code in that module (§12.2.3).

The main function can be omitted if there is only a single module in the program other than the standard modules, as discussed in §12.2.4.

12.2.2 Command-Line Arguments

A predefined variable is used to capture arguments to the execution of a program. It has this declaration:

```
var argv: seq of string;
```

The number of arguments passed to the program execution can be queried with the sequence `length` function as in

```
argv.length
```

Implementation note. There is no support for the variable `argv`. Only configuration variables allow arguments to be passed to the execution of a program.

12.2.3 Module Execution

Top-level code in a module is executed the first time that module is used via a *use-statement*.

12.2.4 Programs with a Single Module

To aid in exploratory programming, if a program is defined in a single module that uses only standard modules, the module need not define a main function. A default main function is created in this case. This main function sole executable statement is to call the initialize function of that module.

Example. The code

```
writeln("Hello World!");
```

is a legal and complete Chapel program. The module declaration is taken to be the file as described in §12.5.

12.3 Using Modules

Modules can be used by code outside of that module. This allows access to the symbols in the modules without the need for explicit naming (§12.3.1). Using modules is accomplished via the use statement as defined in §11.9.

12.3.1 Explicit Naming

All symbols can be named explicitly with the following syntax:

```
module-named-identifier :  
module-identifier . identifier
```

```
module-identifier :  
identifier
```

This allows two variables that have the same name to be distinguished based on the name of their module. If code requires using symbols from two different modules that have the same name, explicit naming is needed to disambiguate between the two symbols. Explicit naming can also be used instead of using a module.

12.4 Nested Modules

A nested module is a module that is defined inside another module, the outer module. Nested modules automatically have access to all of the symbols in the outer module. However, the outer module needs to explicitly use a nested module to have access to its symbols.

Example. A nested module can be used without using the outer module by explicitly naming the module in the use statement. The code

```
use libmsl blas;
```

uses a module named `blas` that is nested inside a module named `libmsl`.

12.5 Implicit Modules

Multiple modules can be defined in the same file and do not need to bear any relation to the file in terms of their names. As a convenience, a module declaration statement can be omitted if it is the sole module defined in a file. In this case, the module takes its name from the file.

13 Functions

This section defines functions. Methods and iterators are functions and most of this section applies to them as well. They are defined separately in §20 and §14.4.

13.1 Function Definitions

Functions are declared with the following syntax:

```

function-declaration-statement :
    def function-name argument-list var-clauseopt
      return-typeopt where-clauseopt block-level-statement

function-name :
    identifier
    = identifier
    operator-name

operator-name : one of
    + - * / % ** && || ! == <= >= < > << >> & | ^ ~ #

argument-list :
    ( formalsopt )

formals :
    formal
    formal , formals

formal :
    formal-tag identifier formal-typeopt default-expressionopt
    formal-tag identifier formal-typeopt variable-argument-expression

formal-type :
    : type
    : TQUESTION identifier

default-expression :
    = expression

variable-argument-expression :
    ... expression
    ... TQUESTION identifier

formal-tag : one of
    in out inout param type

var-clause :
    var

where-clause :
    where expression

```

Operator overloading is supported in Chapel on the operators listed above under operator name. Operator and function overloading is discussed in §13.7.

The intents in, out, and inout are discussed in §13.5. The formal tags param and type make a function generic and are discussed in §21. If the formal argument's type is elided, generic, or prefixed with a question mark, the function is also generic and is discussed in §21.

Default expressions allow for the omission of actual arguments at the call site, resulting in the implicit passing of a default value. Default values are discussed in §13.4.2.

Functions can take a variable number of arguments. Such functions are discussed in §13.10.

The optional `var` clause defines a variable function discussed in §13.6. A variable setter function can be defined explicitly by prefixing the function's name with the assignment operator. This type of variable function is discussed in §13.6.1.

The optional `where` clause is only applicable if the function is generic. It is discussed in §21.4.

13.2 The Return Statement

The return statement must appear in a function. It exits that function, returning control to the point at which that function was called. It can optionally return a value. The syntax of the return statement is given by

```
return-statement :  
    return expressionopt ;
```

Example. The following code defines a function that returns the sum of three integers:

```
def sum(i1: int, i2: int, i3: int)  
    return i1 + i2 + i3;
```

13.3 Function Calls

Functions are called in call expressions described in §10.3. The function that is called is resolved according to the algorithm described in §13.8.

13.4 Formal Arguments

Chapel supports an intuitive formal argument passing mechanism. An argument is passed by value unless it is a class, array, or domain in which case it is passed by reference.

Intents (§13.5) result in potential assignments to temporary variables during a function call. For example, passing an array by intent `in`, a temporary array will be created.

13.4.1 Named Arguments

A formal argument can be named at the call site to explicitly map an actual argument to a formal argument.

Example. In the code

```
def foo(x: int, y: int) { ... }

foo(x=2, y=3);
foo(y=3, x=2);
```

named argument passing is used to map the actual arguments to the formal arguments. The two function calls are equivalent.

Named arguments are sometimes necessary to disambiguate calls or ignore arguments with default values. For a function that has many arguments, it is sometimes good practice to name the arguments at the callsite for compiler-checked documentation.

13.4.2 Default Values

Default values can be specified for a formal argument by appending the assignment operator and a default expression to the declaration of the formal argument. If the actual argument is omitted from the function call, the default expression is evaluated when the function call is made and the evaluated result is passed to the formal argument as if it were passed from the call site.

Example. In the code

```
def foo(x: int = 5, y: int = 7) { ... }

foo();
foo(7);
foo(y=5);
```

default values are specified for the formal arguments `x` and `y`. The three calls to `foo` are equivalent to the following three calls where the actual arguments are explicit: `foo(5, 7)`, `foo(7, 7)`, and `foo(5, 5)`. Note that named arguments are necessary to pass actual arguments to formal arguments but use default values for arguments that are specified earlier in the formal argument list.

13.5 Intents

Intents allow the actual arguments to be copied to a formal argument and also to be copied back.

13.5.1 The Blank Intent

If the intent is omitted, it is called a blank intent. In such a case, the value is copied in using the assignment operator. Thus classes are passed by reference and records are passed by value. Arrays and domains are an exception because assignment does not apply from the actual to the formal. Instead, arrays and domains are passed by reference.

With the exception of arrays, any argument that has blank intent cannot be assigned within the function.

13.5.2 The In Intent

If `in` is specified as the intent, the actual argument is copied to the formal argument as usual, but it may also be assigned to within the function. This assignment is not reflected back at the call site.

If an array is passed to a formal argument that has `in` intent, a copy of the array is made via assignment. Changes to the elements within the array are thus not reflected back at the call site. Domains cannot be passed to a function via the `in` intent.

13.5.3 The Out Intent

If `out` is specified as the intent, the actual argument is ignored when the call is made, but after the call, the formal argument is assigned to the actual argument at the call site. The actual argument must be a valid lvalue. The formal argument can be assigned to and read from within the function.

The formal argument cannot not be generic and is treated as a variable declaration. Domains cannot be passed to a function via the `out` intent.

13.5.4 The Inout Intent

If `inout` is specified as the intent, the actual argument is both passed to the formal argument as if the `in` intent applied and then copied back as if the `out` intent applied. The formal argument can be generic and takes its type from the actual argument. Domains cannot be passed to a function via the `inout` intent. The formal argument can be assigned to and read from within the function.

13.6 Variable Functions

A variable function is a function that can be assigned a value. Note that a variable function does not return a reference. That is, the reference cannot be captured.

A variable function is specified by following the argument list with the `var` keyword. A variable function must return an expression that can be assigned.

When a variable function is called on the left-hand side of an assignment statement, the expression that is normally returned is instead assigned the result of the expression on the right-hand side of the assignment statement. Note that the right-hand side expression is evaluated before the variable function is called. Otherwise a variable function evaluates normally and returns the result of the expression that it returns.

Example. The following code creates a function that can be interpreted as a simple two-element array where the elements are actually global variables:

```
var x, y = 0;

def A(i: int) var {
  if i < 0 || i > 1 then
    halt("array access out of bounds");
  return if i == 0 then x else y;
}
```

This function can be assigned to in order to write to the “elements” of the array as in

```
A(0) = 1;
A(1) = 2;
```

It can be called as an expression to access the “elements” as in

```
writeln(A(0) + A(1));
```

This code outputs the number 3.

13.6.1 Explicit Setter Functions

Variable functions can be created by overloading a normal function with an explicit setter function. Alternatively, an explicit setter function can be defined without a normal counterpart to create a function that can only be assigned values.

An explicit setter function is defined by prepending the assignment operator to the function name. Setter functions can only be called from the left-hand side of an assignment statement.

Setter functions require an extra formal argument (that must be the last argument). This argument is passed the value on the right-hand side of the assignment statement.

Example. The following code defines a function `A` and a setter function `A`. These two definitions of `A` are equivalent to the single definition in the previous section.

```
var x, y = 0;

def A(i: int) {
  if i < 0 || i > 1 then
    halt("array access out of bounds");
  return if i == 0 then x else y;
}

def =A(i: int, value: int) {
  if i < 0 || i > 1 then
    halt("array access out of bounds");
  (if i == 0 then x else y) = value;
}
```

This code has the extra constraint that the expression on the right-hand side of the assignment statement needs to be able to be passed to an argument of type `int`. For type `int`, the constraints between assignment and argument passing are the same, so these functions are equivalent to in the previous section.

In the variable function, there is no type constraint on the expression other than any constraint placed on it by the assignment statement. The equivalent setter function would be generic in respect to the last argument.

13.7 Function Overloading

Functions that have the same name but different argument lists are called overloaded functions. Function calls to overloaded functions are resolved according to the algorithm in §13.8.

Operator overloading is achieved by defining a function with a name specified by that operator. The operators that may be overloaded are listed in the following table:

arity	operators
unary	+ - ! ~
binary	+ - * / % ** && ! == <= >= < > << >> & ^ #

The arity and precedence of the operator must be maintained when it is overloaded. Operator resolution follows the same algorithm as function resolution.

13.8 Function Resolution

Given a function call, the function that the call resolves to is determined according to the following algorithm:

- Identify the set of visible functions. A visible function is any function with the same name that satisfies the criteria in §13.8.1.
- From the set of visible functions, determine the set of candidate functions. A function is a candidate if the function is a valid application of the function call's actual arguments as determined in §13.8.2. A compiler error occurs if there are no candidate functions.
- From the set of candidate functions, the most specific function is determined. The most specific function is a candidate function that is more specific than every other candidate function. If there is no function that is more specific than every other candidate function, the function call is ambiguous and a compiler error occurs. The term *more specific function* is defined in §13.8.3.

13.8.1 Identifying Visible Functions

A function is a visible function to a function call if the name of the function is the same as the name of the function call and the function is defined or used in a lexical outer scope.

Additionally, functions that have arguments of class type are considered globally visible and so are always visible regardless of the location of their definition.

13.8.2 Determining Candidate Functions

A function is a candidate function if there is a *valid mapping* from the function call to the function and each actual argument is mapped to a formal argument that is a *legal argument mapping*.

Valid Mapping A function call is mapped to a function according to the following steps:

- Each actual argument that is passed by name is matched to the formal argument with that name. If there is no formal argument with that name, there is no valid mapping.
- The remaining actual arguments are mapped in order to the remaining formal arguments in order. If there are more actual arguments than formal arguments, there is no valid mapping. If any formal argument that is not mapped to by an actual argument does not have a default value, there is no valid mapping.
- The valid mapping is the mapping of actual arguments to formal arguments plus default values to formal arguments that are not mapped to by actual arguments.

Legal Argument Mapping An actual argument of type T_A can be mapped to a formal argument of type T_F if any of the following conditions hold:

- T_A and T_F are the same type.
- There is an implicit coercion from T_A to T_F .
- T_A is derived from T_F .
- T_A is scalar promotable to T_F .

13.8.3 Determining More Specific Functions

Given two functions F_1 and F_2 , F_1 is determined to be more specific than F_2 by the following steps:

- If at least one of the legal argument mappings to F_1 is a *more specific argument mapping* than the corresponding legal argument mapping to F_2 and none of the legal argument mappings to F_2 is a more specific argument mapping than the corresponding legal argument mapping to F_1 , then F_1 is more specific.
- If F_1 shadows F_2 , then F_1 is more specific.
- Otherwise, F_1 is not more specific than F_2 .

Given an argument mapping, M_1 , from an actual argument, A , of type T_A to a formal argument, $F1$, of type T_{F1} and an argument mapping, M_2 , from the same actual argument to a formal argument, $F2$, of type T_{F2} , the more specific argument mapping is determined by the following steps:

- If T_{F1} and T_{F2} are the same type and $F1$ is an instantiated parameter, M_1 is more specific.
- If T_{F1} and T_{F2} are the same type and $F2$ is an instantiated parameter, M_2 is more specific.
- If M_1 requires scalar promotion and M_2 does not require scalar promotion, M_2 is more specific.
- If M_2 requires scalar promotion and M_1 does not require scalar promotion, M_1 is more specific.
- If $F1$ is generic over all types and $F2$ is not generic over all types, M_2 is more specific.

- If F_2 is generic over all types and F_1 is not generic over all types, M_1 is more specific.
- If T_{F_1} and T_{F_2} are the same type, neither mapping is more specific.
- If T_A and T_{F_1} are the same type, M_1 is more specific.
- If T_A and T_{F_2} are the same type, M_2 is more specific.
- If T_{F_1} is derived from T_{F_2} , then M_1 is more specific.
- If T_{F_2} is derived from T_{F_1} , then M_2 is more specific.
- If there is an implicit coercion from T_{F_1} to T_{F_2} , then M_1 is more specific.
- If there is an implicit coercion from T_{F_2} to T_{F_1} , then M_2 is more specific.
- If T_{F_1} is any `int` type and T_{F_2} is any `uint` type, M_1 is more specific.
- If T_{F_2} is any `int` type and T_{F_1} is any `uint` type, M_2 is more specific.
- Otherwise neither mapping is more specific.

13.9 Nested Functions

A function defined in another function is called a nested function. Nesting of functions may be done to arbitrary degrees, i.e., a function can be nested in a nested function.

Nested functions are only visible to function calls within the scope in which they are defined. An exception is to a function that has an argument that is a class type. Such functions are globally visible.

13.9.1 Accessing Outer Variables

Nested functions may refer to variables defined in the function in which they are nested. If the function has class arguments, and is thus globally visible, it is a compiler error to refer to a variable in the outer function.

Rationale. It may be too strict to make this a compiler error. Are there advantages to making this a runtime error?

13.10 Variable Length Argument Lists

Functions can be defined to take a variable number of arguments. This allows the call site to pass a different number of actual arguments.

If the variable argument expression is an identifier prepended by a question mark, the number of arguments is variable. Alternatively, the variable expression can evaluate to an integer parameter value requiring the call site to pass that number of arguments to the function.

In the function, the formal argument is a tuple of the actual arguments.

Example. The code

```
def mywriteln(x: int ...?k) {  
  for param i in 1..k do  
    writeln(x(i));  
}
```

defines a function called `mywriteln` that takes a variable number of arguments and then writes them out on separate lines. The parameter for-loop (§11.8.3) is unrolled by the compiler so that `i` is a parameter, rather than a variable. This function can be made generic (§21) to take arguments of different types by eliding the type.

A tuple of variables arguments can be passed to a function that takes variable arguments by destructuring the tuple in a tuple destructuring expression. The syntax of this expression is given by

```
tuple-destructuring-expression :  
( ... expression )
```

In this expression, the tuple defined by *expression* is expanded in place to represent its components. This allows for the forwarding of variable arguments as variable arguments.

14 Classes

Classes are an abstraction of a data structure where the storage location is allocated independent of the scope of the variable of class type. Each call to the constructor creates a new data object and returns a reference to the object. Storage is reclaimed automatically as described in §14.10.

14.1 Class Declarations

A class is defined with the following syntax:

```

class-declaration-statement :
    class identifier class-inherit-type-listopt {
        class-statement-list }

class-inherit-expression-list :
    class-type
    class-type , inherit-expression-list

class-statement-list :
    class-statement
    class-statement class-statement-list

class-statement :
    type-declaration-statement
    function-declaration-statement
    variable-declaration-statement

```

A *class-declaration-statement* defines a new type symbol specified by the identifier. Classes inherit data and functionality from other classes if the *inherit-type-list* is specified. Inheritance is described in §14.7.

The body of a class declaration consists of a sequence of statements where each of the statements either defines a variable (called a field), a function (called a method), or a type.

If a class contains a type alias or a parameter, the class is generic. Generic classes are described in §21.

14.2 Class Assignment

Classes are assigned by reference. After an assignment from one variable of class type to another, the variables reference the same storage location.

14.3 Class Fields

Variables and constants declared within class declarations define fields within that class. (Parameters make a class generic.) Fields define the storage associated with a class.

Example. The code

```

class Actor {
    var name: string;
    var age: uint;
}

```

defines a new class type called `Actor` that has two fields: the string field `name` and the unsigned integer field `age`.

14.3.1 Class Field Accesses

The field in a class is accessed via a member access expression as described in §10.3.2. Fields in a class can be modified via an assignment statement where the left-hand side of the assignment is a member access.

Example. Given a variable `anActor` of type `Actor`, defined above, the code

```

var s: string = anActor.name;
anActor.age = 27;

```

reads the field `name` and assigns the value to the variable `s`, and assigns the storage location in the object `anActor` associated with the field `age` the value 27.

14.4 Class Methods

A method is a function that is bound to a class. A method is called by passing an instance of the class to the method via a special syntax that is similar to a field access.

14.4.1 Class Method Declarations

Methods are declared with the following syntax:

```

method-declaration-statement :
    def type-binding function-name argument-listopt var-clauseopt
      return-typeopt where-clauseopt block-level-statement

type-binding :
    identifier .

```

If a method is declared within the lexical scope of a class, record, or union, the type binding can be omitted and is taken to be the innermost class, record, or union that the method is defined in.

14.4.2 Class Method Calls

A method is called by using the member access syntax as described in §10.3.2 where the accessed expression is the name of the method.

Example. A method to output information about an instance of the `Actor` class can be defined as follows:

```
def Actor.print() {
  writeln("Actor ", name, " is ", age, " years old");
}
```

This method can be called on an instance of the `Actor` class, `anActor`, by writing `anActor.print()`.

14.4.3 The *this* Reference

The instance of a class is passed to a method using special syntax. It does not appear in the argument list to the method. The reference `this` is an alias to the instance of the class on which the method is called.

Example. Let class `C`, method `foo`, and function `bar` be defined as

```
class C {
  def foo() {
    bar(this);
  }
}
def bar(c: C) { }
```

Then given an instance of `C` called `c`, the method call `c.foo()` results in a call to `bar` where the argument is `c`.

14.4.4 Class Methods without Parentheses

Methods do not require parentheses if they have empty argument lists. Methods declared without parentheses around empty argument lists must be called without parentheses.

Example. Given the definitions

```
class C {
  def foo { }
  def bar() { }
}
```

and an instance of `C` called `c`, then the method `foo` can be called by writing `c.foo` and the method `bar` can be called by writing `c.bar()`. It is an error to apply parentheses to `foo` or omit them from `bar`.

14.4.5 The *this* Method

A method declared with the name `this` allows a class to be “indexed” similarly to how a tuple, sequence, or array is indexed. Indexing into a class has the semantics of calling a method on the class named `this`. There is no other way to call a method called `this`. The `this` method must be declared with parentheses even if the argument list is empty.

Example. In the following code, the `this` method is used to create a class that acts like a simple array that contains three integers indexed by 1, 2, and 3.

```

class ThreeArray {
  var x1, x2, x3: int;
  def this(i: int) var {
    select i {
      when 1 do return x1;
      when 2 do return x2;
      when 3 do return x3;
    }
    halt("ThreeArray index out of bounds: ", i);
  }
}

```

14.5 Class Constructors

A class constructor is defined by declaring a method with the same name as the class. The constructor is used to create instances of the class. When the constructor is called, memory is allocated to store a class instance.

14.5.1 The Default Constructor

A default constructor is automatically created for every class in the Chapel program. This constructor is defined such that it has one argument for every field in the class. Each of the arguments has a default value.

The default constructor is very useful but its generality in terms of having one argument for each field all of which have default values makes it slightly difficult for the user to create their own constructor. It is expected that in many simple cases, the default constructor will be all that is necessary.

Example. Given the class

```

class C {
  def x: int;
  def y: real = 3.14;
  def z: string = "Hello, World!";
}

```

then instances of the class can be created by calling the default constructor as follows:

- The call `C()` is equivalent to `C(0, 3.14, "Hello, World")!`.
- The call `C(2)` is equivalent to `C(2, 3.14, "Hello, World")!`.
- The call `C(z=" ")` is equivalent to `C(0, 3.14, " ")`.
- The call `C(0, 0.0, " ")` is equivalent to `C(0, 0.0, " ")`.

14.6 Getters and Setters

All field accesses are resolved via getter and setter methods that are defined in the class with the same name as the field. A setter is defined as an explicit setter function (§13.6.1). Default getters and setters are defined that simply access or set the field if the user does not define their own.

Example. In the code


```
class C {  
  var x: int;  
  def =x(value: int) {  
    if value < 0 then  
      halt("x assigned negative value");  
    x = value;  
  }  
}
```

a setter is defined for field `x` that ensures that `x` is never assigned a negative value.

14.7 Inheritance

A “derived” class can inherit from one or more other classes by specifying those classes, the base classes, following the name of the derived class in the declaration of the derived class. When inheriting from multiple base classes, only one of the base classes may contain fields. The other classes can only define methods. Note that a class can still be derived from a class that contains fields which is itself derived from a class that contains fields.

14.7.1 Accessing Base Class Fields

A derived class contains data associated with the fields in its base classes. The fields can be accessed in the same way that they are accessed in their base class unless the getter or setter method is overridden in the derived class, as discussed in §14.7.4.

14.7.2 Derived Class Constructors

Derived class constructors automatically call the default constructor of the base class. There is an expectation that a more standard way of chaining constructor calls will be supported.

14.7.3 Shadowing Base Class Fields

A field in the derived class can be declared with the same name as a field in the base class. Such a field shadows the field in the base class in that it is always referenced when it is accessed in the context of the derived class. There is an expectation that there will be a way to reference the field in the base class but this is not defined at this time.

14.7.4 Overriding Base Class Methods

If a method in a derived class is declared with the identical signature as a method in a base class, then it is said to override the base class’s method. Such a method is a candidate for dynamic dispatch in the event that a variable that has the base class type references a variable that has the derived class type.

The identical signature requires that the names, types, and order of the formal arguments be identical.

14.7.5 Inheriting from Multiple Classes

Implementation note. Multiple inheritance is not yet supported.

A class can be derived from multiple base classes provided that only one of the base classes contains fields either directly or from base classes that it is derived from. The methods defined by the other base classes can be overridden.

14.8 Class Promotion of Scalar Functions

A class can be defined to promote scalar functions by defining an iterator in the class named `this` and specifying a return type. The return type indicates the type that the class promotes. The body of the `this` iterator is ignored. The class must also implement the iterator interface as described in §20.4.

There is an expectation that class promotion will be implemented in a different way in the future.

14.9 Nested Classes

Implementation note. Nested classes are not yet supported.

A class defined within another class is a nested class.

14.10 Automatic Memory Management

Implementation note. Memory allocated to store class objects is not yet reclaimed.

Memory associated with class instances is reclaimed automatically when there is no way for the current program to reference this memory. The programmer does not need to free the memory associated with class instances.

15 Records

A record is a data structure that is like a class but has value semantics. The key differences between records and classes are described in this section.

15.1 Record Declarations

A record is defined with the following syntax:

```

record-declaration-statement :
    record identifier inherit-type-listopt {
        record-statement-list }

record-inherit-expression-list :
    record-type
    record-type , inherit-expression-list

record-statement-list :
    record-statement
    record-statement record-statement-list

record-statement :
    type-declaration-statement
    function-declaration-statement
    variable-declaration-statement

```

The only difference between record and class declarations is that the `record` keyword replaces the `class` keyword.

15.2 Class and Record Differences

The main differences between records and classes are that records are value classes. They do not need to be reclaimed since the data is recovered when the variable goes out of scope, do not support dynamic dispatch, and are assigned by value.

15.2.1 Records as Value Classes

A record is not a reference to a storage location that contains the data in the record but is more like a variable of a primitive type. A record directly contains the data associated with the fields in the record.

15.2.2 Record Inheritance

When a record is derived from a base record, it contains the data in the base record. The difference between record inheritance and class inheritance is that there is no dynamic dispatch. The record type of a variable is the exact type of that variable.

15.2.3 Record Assignment

In record assignment, the fields of the record on the left-hand side of the assignment are assigned the values in the fields of the record on the right-hand side of the assignment. When a base record is assigned a derived record, just the fields that exist in the base record are assigned.

15.3 Default Comparison Operators on Records

Default functions to overload `==` and `!=` are defined for records if there is none defined for the record in the Chapel program. The default implementation of `==` applies `==` to each field of the two argument records and reduces the result with the `&&` operator. The default implementation of `!=` applies `!=` to each field of the two argument records and reduces the result with the `||` operator.

16 Unions

Unions have the semantics of records, however, only one field in the union can contain data at any particular point in the program's execution. Unions are safe so that an access to a field that does not contain data is a runtime error. When a union is constructed, it is in an unset state so that no field contains data.

16.1 Union Declarations

A union is defined with the following syntax:

```
union-declaration-statement :
    union identifier { union-statement-list }

union-statement-list :
    union-statement
    union-statement union-statement-list

union-statement :
    type-declaration-statement
    function-declaration-statement
    variable-declaration-statement
```

16.1.1 Union Fields

Union fields are accessed in the same way that record fields are accessed. It is a runtime error to access a field that is not currently set.

Union fields should not be specified with initialization expressions.

16.2 Union Assignment

Union assignment is by value. The field set by the union on the right-hand side of the assignment is assigned to the union on the left-hand side of the assignment and this same field is marked as set.

16.3 The Type Select Statement and Unions

The type-select statement can be applied to unions to access the fields in a safe way by determining the type of the union.

Implementation note. The type-select statement is not yet implemented on unions.

17 Tuples

A tuple is an ordered set of components that allows for the specification of a light-weight record with anonymous fields.

17.1 Tuple Expressions

A tuple expression is a comma-separated list of expressions that is enclosed in parentheses. The number of expressions is the size of the tuple and the types of the expressions determine the component types of the tuple.

The syntax of a tuple expression is given by:

```
tuple-expression :
  ( expression-list )

expression-list :
  expression
  expression , expression-list
```

Example. The statement

```
var x = (1, 2);
```

defines a variable `x` that is a 2-tuple containing the values 1 and 2.

17.2 Tuple Type Definitions

A tuple type is a comma-separated list of types. The number of types in the list defines the size of the tuple, which is part of the tuple's type. The syntax of a tuple type is given by:

```
tuple-type :
  ( type-list )

type-list :
  type
  type , type-list
```

Example. Given a tuple expression `(1, 2)`, the type of the tuple value is `(int, int)`, referred to as a 2-tuple of integers.

17.3 Tuple Assignment

In tuple assignment, the components of the tuple on the left-hand side of the assignment operator are each assigned the components of the tuple on the right-hand side of the assignment. The assignments are simultaneous so that each component expression on the right-hand side is fully evaluated before being assigned to the left-hand side.

17.4 Tuple Destructuring

When a tuple expression appears on the left-hand side of an assignment statement, the expression on the right-hand side is said to be *destructured*. The components of the tuple on the right-hand side are assigned to each of the component expressions on the left-hand side. This assignment is simultaneous in that the right-hand side is evaluated before the assignments are made.

Example. Given two variables of the same type, *x* and *y*, they can be swapped by the following single assignment statement:

```
(x, y) = (y, x);
```

17.4.1 Variable Declarations in a Tuple

Variables can be defined in a tuple to facilitate capturing the values from a function that returns a tuple. The extension to the syntax of variable declarations is as follows:

```
tuple-variable-declaration-statement :
    configopt variable-kind tuple-variable-declaration ;

tuple-variable-declaration :
    ( identifier-list ) type-partopt initialization-part
    ( identifier-list ) type-part
```

The identifiers defined within the *identifier-list* are declared to be new variables in the scope of the statement. The *type-part* and/or *initialization-part* defines a tuple that is destructured when assigned to the new variables.

17.4.2 Ignoring Values with Underscore

If an underscore appears as a component in a tuple expression in a destructuring context, the expression on the right-hand side is ignored, though it is still evaluated.

Implementation note. Underscores are currently not treated specially. The only way to ignore values when destructuring a tuple is to put them into variables that are never read.

17.5 Homogeneous Tuples

A homogeneous tuple is a special-case of a general tuple where the types of the components are identical. Homogeneous tuples have less restrictions for how they can be indexed (§17.6).

17.5.1 Declaring Homogeneous Tuples

A homogeneous tuple type may be specified with the following syntax if it appears as a top-level type in a variable declaration, formal argument declaration, return type specification, or type alias declaration:

```
homogeneous-tuple-type :
    integer-parameter-expression * type

integer-parameter-expression :
    expression
```

The homogeneous tuple type specification is syntactic sugar for the type explicitly replicated a number of times equal to the *integer-parameter-expression*.

Example. The following types are equivalent:

```
3*int      (int, int, int)
```

17.6 Tuple Indexing

A tuple may be indexed into by an integer. Indexing a tuple is given by the following syntax:

```
tuple-indexing-expression :
    expression ( integer-expression )
```

The result of indexing a tuple by integer k is the value of the k th component. If the tuple is not homogeneous, the tuple can only be indexed by an integer parameter. This ensures that the type of the indexing expression is known at compile-time.

17.7 Formal Arguments of Tuple Type

Implementation note. Formal arguments of tuple type are treated as if they were records. Conversions are not applied to the components.

17.7.1 Formal Argument Declarations in a Tuple

Implementation note. Formal arguments cannot be grouped together in a tuple.

18 Sequences

A sequence is an ordered set of elements of the same type.

18.1 Sequence Literals

Literal sequences are delimited by the braces `(/` and `/)`. The expressions enclosed in these braces are elements of the sequence. These expressions do not need to be literals themselves.

Example. The following code defines a sequence of integers:

```
(/1, 1, 2, 3, 5, 8, 13, 21/)
```

18.2 Sequence Type Definitions

A sequence type can be specified with an explicit element type using the following syntax:

```
sequence-type :
  seq of type
  seq ( type )
  seq ( elt_type = type )
```

The element type of a sequence can be referred to by its name `elt_type`. If `s` is a sequence, then the type `s.elt_type` refers to the type of its elements.

18.3 Sequence Rank

The rank of a sequence is determined as follows:

- If the element type of the sequence is not a sequence, then the rank of the sequence is 1.
- If the element type of the sequence is a sequence of rank k , then the rank of the sequence is $k + 1$.

Example. The rank of the sequence specified by the type

```
seq of seq of int
```

is two.

18.4 Sequence Assignment

Sequence assignment is by value.

18.5 Iteration over Sequences

Sequences can be iterated over within the context of a for or forall loop. The type of the index is the element type of the sequence.

Example. In the code

```
var ss: seq of string = ("/one", "two", "three/");
for s in ss do
  writeln(s);
```

the sequence of strings `ss` is iterated over with index `s` of type string. The output to this code is

```
one
two
three
```

18.6 Sequence Concatenation

The operator `#` is used to concatenate two sequences and to append or prepend an element to a sequence. In each case, the operands are unchanged and a new sequence is returned. When applied to two sequences, the new sequence is the concatenation of the two sequences. When applied to an element and a sequence, the new sequence is a copy of the sequence with the element prepended to it. When applied to a sequence and an element, the new sequence is a copy of the sequence with the element appended to it.

The concatenation, prepend, and append operations apply to expressions of types according to the rules of function resolution as specified by the following function prototypes:

```
def #(s1: seq, s2: seq) where s1.type == s2.type // concatenate
def #(e, s: seq) where e.type:s.elts_type // prepend
def #(s: seq, e) where e.type:s.elts_type // append
```

18.7 Sequence Indexing

Sequences can be indexed into by applying a function call to them. They can be indexed into by integers or tuples of integers.

18.7.1 Sequence Indexing by Integers

If `s` is a sequence and `i` is an integer, then the expression `s(i)` is evaluated to return an element in the sequence. A runtime error occurs if `i` is zero or the absolute value of `i` is greater than the length of the sequence. Let the elements of the sequence be denoted as e_1, e_2, \dots, e_n where n is the length of the sequence. If `i` is positive, then `s(i)` is the value e_i . If `i` is negative, then `s(i)` is the value e_{n-i+1} .

18.7.2 Sequence Indexing by Tuples

If s is a sequence and t is a tuple of integers of size k , then the expression $s(t)$ indexes into the sequence s k times using the integers in the tuple. In this case, s must be a sequence whose rank is at least as great as the size of t . If s has rank less than the size of the tuple, then the result is a sequence.

If the integers in tuple t are denoted as i_1, i_2, \dots, i_k , then the expression $s(t)$ is equivalent to the expression $s(i_1)(i_2)\dots(i_k)$.

A sequence can also be indexed by multiple integers. The integers are collected into a tuple and the sequence is then indexed by the tuple.

18.8 Sequence Promotion of Scalar Functions

Sequence promotion of a scalar function is defined over the sequence and its *leaf element type*. If the sequence has rank one, then the leaf element type is the element type of the sequence. If the sequence has rank greater than one, then the leaf element type is defined to be the first element type encountered by recursively indexing into the sequence such that the element type is not a sequence.

If a variable of type equal to the leaf element type of the sequence can be passed to a function's formal argument according to the rules of function resolution in §13.8, then the sequence can be passed to this function and the function is said to be sequence promoted. In general, it is said to be scalar promoted since types other than sequences may result in the promotion.

18.8.1 Zipper Promotion

Consider a function f with formal arguments s_1, s_2, \dots that are sequence promoted and formal arguments a_1, a_2, \dots that are not promoted. The call

```
f(s1, s2, ..., a1, a2, ...)
```

is equivalent to

```
[(e1, e2, ...) in (s1, s2, ...)] f(e1, e2, ..., a1, a2, ...)
```

The usual constraints of zipper iteration apply to zipper promotion so the sequences must have the same shape.

The result of the promotion is a sequence of the same rank but the leaf element type is the return type of the function that is promoted.

Example. Given a function defined as

```
def foo(i: int, j: int) {
  write(i, " ", j, " ");
}
```

and a call to this function written

```
foo((/1, 2, 3/), (/4, 5, 6/));
```

then the output is “1 4 2 5 3 6”.

18.8.2 Tensor Product Promotion

If the function f were called by using square brackets instead of parentheses, the equivalent rewrite would be

```
[(e1, e2, ...) in [s1, s2, ...]] f(e1, e2, ..., a1, a2, ...)
```

There are no constraints on tensor product promotion.

The result of the promotion is a sequence of rank equal to the sum of the ranks of the promoted arguments. The leaf element type of the sequence is the return type of the function that is promoted.

Example. Given a function defined as

```
def foo(i: int, j: int) {
    write(i, " ", j, " ");
}
```

and a call to this function written

```
foo[(/1, 2, 3/), (/4, 5, 6/)];
```

then the output is “1 4 1 5 1 6 2 4 2 5 2 6 3 4 3 5 3 6”.

18.9 Sequence Operators

Any operator that can be applied to the leaf element type of a sequence can be applied to the sequence according to the rules of zipper promotion.

Example. Given two sequences of integers

```
var s1 = (/1, 2, 3, 4, 5/);
var s2 = (/5, 4, 3, 2, 1/);
```

the sequence that is returned by applying the equality operator to the sequences is a sequence of bool values. The code

```
writeln(s1 == s2);
```

produces the output “false false true false false.”

18.10 Sequences in Logical Contexts

When a sequence expression is used as a top-level expression in the condition of a while statement, a do-while statement or an if statement, the sequence is promoted to a sequence of bool values following the implicit statement bool conversion rule (§9.1.6) applied to the leaf element type of the sequence.

The sequence of bool values is then implicitly reduced using the $\&\&$ operator to produce a single bool value. Reductions are defined in §24.

Implementation note. Neither the implicit conversion nor the implicit reduction is currently implemented.

18.10.1 Sequences in Select Statements

When a sequence expression is used as a top-level expression in the condition of a select statement, there are two interpretations. If the condition in the when expression is itself a sequence, the equality operator is used to compare the sequences and then an implicit `&&` reduction is applied to produce a single bool value. If the condition in the when expression is a scalar, the equality operator is used to compare the sequences and then an implicit `||` reduction is applied to produce a single bool value.

Implementation note. The implicit reduction is currently not implemented. Sequences cannot be used in the condition of a select statement or in the when expression.

18.11 Filtering Predicates

An if expression that is immediately enclosed by a forall expression does not require an else part. The result of the forall expression is a sequence of rank one.

Example. The following expression returns every other element in a sequence starting with the first:

```
[i in 1..s.length] if i % 2 == 1 then s(i)
```

18.12 Methods and Functions on Sequences

18.12.1 The *length* Method

```
def seq.length
```

The length method requires no parentheses and returns the number of elements in a sequence.

18.12.2 The *reverse* Method

```
def seq.reverse(dim: int = 1)
```

The reverse method returns the elements of a sequence in reverse order. An optional second argument is an int. If the value of this argument is d , then the rank of the first argument must be at least d . When d is one, then the sequence is reversed. When d is greater than one, the expression `s.reverse(d)` is equivalent to

```
[e in s] e.reverse(d-1)
```

Implementation note. The `dim` argument is not implemented.

18.12.3 The *spread* Function

```
def spread(s: seq, length: int, dim: int = 1)
```

The spread function takes a sequence of rank k and returns a new sequence of rank $k + 1$. When `dim` is equal to 1, the result is a sequence where every element is equal to s . The length of this sequence is specified by `length`. When `dim` is greater than one, we generate the sequence:

```
[e in s] spread(e, length, dim-1)
```

Implementation note. The spread function is not implemented.

18.12.4 The *transpose* Function

```
def transpose(s: seq, dims = (2,1))
```

The transpose function will reorder both the values and change the shape of the sequence. The optional `dims` argument is a tuple that corresponds to a permutation of the values $1..p$ where p is less than or equal to the rank of the input sequence. This list defines a permutation of the subscripts such that the following relationship holds between the input and output sequences:

$$s'(i'_1, \dots, i'_k) = s(i_1, \dots, i_k)$$

where

$$i_j = \begin{cases} i'_{dims(j)} & \text{if } j \leq p \\ i'_j & \text{otherwise} \end{cases}$$

There is a somewhat complex requirement in the shape of the input sequence so that this relation is well-defined. In the simple case of a rank-2 input, we require that all elements of the sequence have the same length. In the general case we require all sequences selected by a q -tuple to have the same length whenever an index position q is less than p .

Implementation note. The transpose function is not implemented.

18.12.5 The *reshape* Function

```
reshape(s: seq, shape, fill)
```

This reshape function returns a sequence whose leaves are the same as the first argument, in the same order but whose shape matches that of `shape`. The `shape` argument might be a sequence or it might be a tuple of ints. When it is a sequence, the output will conform to that sequence. When the shape is a tuple, then the shape of the output conforms to the shape of the arithmetic index set that would be determined by $1..shape$.

If present, the `fill` argument specifies a value to be used to pad the sequence if the number of leaf values in that sequence is too few to conform with `shape`. If the input sequence has too many values, it is truncated.

Implementation note. The reshape function is not implemented.

18.13 Arithmetic Sequences

Arithmetic sequences contain an ordered set of values of integral type that can be specified with a low bound l , a high bound h , and a stride s . If the stride is negative, the values contained by the arithmetic sequence are $h, h-s, h-2s, h-3s, \dots$ such that all of the values in the sequence are greater than l . If the stride is positive, the values contained by the arithmetic sequence are $l, l+s, l+2s, l+3s, \dots$ such that all of the values in the sequence are less than h .

An arithmetic sequence is specified by the syntax

```
arithmetic-sequence-literal:  
    expression .. expression
```

The first expression is taken to be the lower bound, the second expression is taken to be the upper bound. The stride of the arithmetic sequence is 1 and can be modified with the `by` operator.

The element type of the arithmetic sequence type is determined by the type of the low and high bound. It is either `int`, `uint`, `int(64)`, or `uint64`. The type is determined by conceptually adding the low and high bounds together.

All operations supported over sequences are supported over arithmetic sequences. So arithmetic sequences can, for example, sequence promote functions and be iterated over. The result of the sequence promotion is a sequence but not an arithmetic sequence.

18.13.1 Strided Arithmetic Sequences

The `by` operator can be applied to any arithmetic sequence to create a strided arithmetic sequence. It is predefined over an arithmetic sequence and an integer to yield a new arithmetic sequence that is strided by the integer. Striding of strided sequences is accomplished by multiplying the strides.

18.13.2 Querying the Bounds and Stride of an Arithmetic Sequence

```
def arithmetic sequence.low: elt_type  
def arithmetic sequence.high: elt_type  
def arithmetic sequence.stride: int
```

These routines respectively return the low bound, the high bound, and the stride of the arithmetic sequence. The type of the returned low and high bound is the element type of the arithmetic sequence.

18.13.3 Indefinite Sequences

An indefinite arithmetic sequence is specified by the syntax

```
indefinite-arithmetic-sequence-literal:  
    expression ..  
    .. expression
```

Indefinite arithmetic sequences can be iterated over with zipper iteration and their shape conforms to the shape of the sequences they are being iterated over.

Example. The code

```
for i in (1..5, 3..) do
  write(i);
```

produces the output “(1, 3)(2, 4)(3, 5)(4, 6)(5, 7)”.

It is an error to zip an indefinite arithmetic sequence with an arithmetic sequence that does not have the same sign stride.

Indefinite arithmetic sequences can be used to index into arithmetic sequences, sequences, arrays, and strings. In these cases, the elided bound conforms to the expression being indexed.

Implementation note. Indexing by indefinite arithmetic sequences is not yet supported.

18.14 Conversions Between Sequences and Tuples

A value of homogenous tuple type can be converted to a sequence by applying a cast to the keyword `seq`. The elements in the tuple become the elements in the sequence and the sequence’s length is the size of the tuple.

A value of sequence type can be converted to a homogeneous tuple type by applying a cast to a parameter of type `int`. If the parameter has value k then the tuple is of size k and the first k elements in the sequence are the elements in the tuple.

19 Domains and Arrays

A *domain* is a description of a collection of names for data. These names are referred to as the *indices* of the domain. All indices for a particular domain are values with some common type. Valid types for indices are primitive types and class references or unions, tuples or records whose fields are valid types for indices. This excludes sequences, domains, and arrays. Like sequences, domains have a rank and a total order on their elements. An *array* is generically a function that maps from a *domain* to a collection of variables. Chapel supports a variety of kinds of domains and arrays defined over those domains as well as a mechanism to allow application specific implementations of arrays.

Arrays abstract mappings from sets of values to variables. This key use of data structures coupled with the generic syntactic support for array usage increases software reusability. By separating the sets of values into their own abstraction, domains, distributions can be associated with sets rather than variables. This enables the orthogonality of data distributions. Distributions are discussed in §23.

19.1 Domains

Domains are first-class ordered sets of indices. There are five kinds of domains:

- Arithmetic domains are rectilinear sets of Cartesian indices that can have an arbitrary rank.
- Sparse domains are subsets of indices in arithmetic domains.
- Indefinite domains are sets of indices where the type of the index is some type that is not an array, domain, or sequence. Indefinite domains define dictionaries or associative arrays implemented via hash tables.
- Opaque domains are sets of anonymous indices. Opaque domains define graphs and unspecified sets.
- Enumerated domains are sets of constants defined by some enumerated type.

19.1.1 Domain Types

Domain types vary based on the kind of the domain. The type of an arithmetic domain is parameterized by the rank of the domain and the integral type of the indices. The type of a sparse domain is parameterized by the type of the arithmetic domain that defines the superset of its indices. The type of an indefinite domain is parameterized by the type of the index. The type of an opaque domain is unique. The type of an enumerated domain is parameterized by the enumerated type.

Example. In the code

```
var D: domain(2) = [1..n, 1..n];
```

D is defined as a two-dimensional arithmetic domain and is initialized to contain the set of indices (i, j) for all i and j such that $i \in 1, 2, \dots, n$ and $j \in 1, 2, \dots, n$.

19.1.2 Index Types

Each domain has a corresponding *index* type which is the type of the domain's indices qualified by its status as an index. Variables of this type can be declared using the following syntax:

```
index-type :  
    int ( domain-expression )
```

If the type of the indices of the domain is `int`, then the index type can be converted into this type.

A value with a type that is the same as the type of the indices in a domain but is not the index type can be converted into the index type using a special “method” called `index`.

Example. In the code

```
var j = D.index(i);
```

the type of the variable `j` is the index type of domain `D`. The variable `i`, which must have the same type as the underlying type of the indices of `D`, is verified to be in domain `D` before it is assigned to `j`.

Values of index type are known to be valid and may have specialized representations to facilitate accessing arrays defined for that domain. It may therefore be less expensive to access arrays using values of appropriate index type rather than values of the more general type the domain is defined over.

Implementation note. In the current implementation, the index type is not distinguished from the underlying type of the indices. The `index` method is not yet implemented.

19.1.3 Domain Assignment

Domain assignment is by value. If arrays are declared over a domain, domain assignment impacts these arrays as discussed in §19.8, but the arrays remain associated with the same domain regardless of the assignment.

19.1.4 Formal Arguments of Domain Type

Domains are passed to functions by reference. Formal arguments that receive domains are aliases of the actual arguments. It is a compile-time error to pass a domain to a formal argument that has a non-blank intent.

19.1.5 Iteration over Domains

All domains support iteration via `forall` and `for` loops over the indices in the set that the domain defines. The type of the indices returned by iterating over a domain is the index type of the domain.

19.1.6 Domain Promotion of Scalar Functions

Domain promotion of a scalar function is defined over the domain type and the type of the indices of the domain (not the index type). Domain promotion has the same semantics as sequence promotion where the scalar type is the indices of the domain and the promotion type is the domain type.

Example. Given an array `A` with element type `int` declared over a one-dimensional domain `D` with integral type `int`, then the array can be assigned the values given by the indices in the domain by writing

```
A = D;
```

19.2 Arrays

Arrays associate variables or elements with the sets of indices in a domain. Arrays must be declared over domains and have a specified element type.

19.2.1 Array Types

The type of an array is parameterized by the type of the domain that it is declared over and the element type of the array. Array types are given by the following syntax:

```
array-type :  
  [ domain-expression ] type  
  
domain-expression :  
  expression
```

The *domain-expression* must specify a domain that the array can be declared over. This can be a domain literal. If it is a domain literal, the square brackets around the domain literal can be omitted.

Example. In the code

```
var A: [D] real;
```

`A` is declared to be an array over domain `D` with elements of type `real`.

Implementation note. Arrays of arrays are not currently supported.

19.2.2 Array Indexing

Arrays can be indexed by indices in the domain they are declared over. The indexing results in an access of the element that is mapped by this index.

Example. If `A` is an array with element type `real` declared over a one-dimensional arithmetic domain `[1..n]`, then the first element in `A` can be accessed via the expression `A(1)` and set to zero via the assignment `A(1) = 0.0`.

Indexing into an array with a domain is call array slicing and is discussed in the next section.

Arithmetic arrays also support indexing over the components of their indices for multidimensional arithmetic domains (where the indices are tuples), as described in §19.3.5.

19.2.3 Array Slicing

An array can be indexed by a domain that has the same type as the domain which the array was declared over. Indexing in this manner has the effect of array slicing. The result is a new array declared over the indexing domain where the elements in the array alias the elements in the array being indexed.

Example. Given the definitions

```
var OuterD: domain(2) = [0..n+1, 0..n+1];
var InnerD: domain(2) = [1..n, 1..n];
var A, B: [OuterD] real;
```

the assignment given by

```
A(InnerD) = B(InnerD);
```

assigns the elements in the interior of `B` to the elements in the interior of `A`.

Arithmetic arrays also support slicing by indexing into them with arithmetic sequences or tuples of arithmetic sequences as described in §19.3.6.

19.2.4 Array Assignment

Array assignment is by value. Arrays can be assigned arrays, sequences, or domains. If `A` is an lvalue of array type and `B` is an expression of either array, sequence, or domain type, then the assignment

```
A = B;
```

is equivalent to

```
forall (i,e) in (A.domain,B) do
  A(i) = e;
```

If the zipper iteration is illegal, then the assignment is illegal. Notice that the assignment is implemented with the semantics of a `forall` loop.

Arrays can also be assigned single values of their element type. In this case, each element in the array is assigned this value. If e is an expression of the element type of the array or a type that can be implicitly converted to the element type of the array, then the assignment

```
A = e;
```

is equivalent to

```
forall i in A.domain do
  A(i) = e;
```

19.2.5 Formal Arguments of Array Type

Arrays are passed to functions by reference. Formal arguments that receive arrays are aliases of the actual arguments. The ordinary rule that disallows assignment to formal arguments of blank intent does not apply to arrays.

When a formal argument has array type, the element type of the array can be omitted and/or the domain of the array can be queried or omitted. In such cases, the argument is generic and is discussed in §21.1.6.

If a non-queried domain is specified in the array type of a formal argument, the domain must match the domain of the actual argument. This is verified at runtime. There is an exception if the domain is an arithmetic domain; it is described in §19.3.7.

19.2.6 Iteration over Arrays

All arrays support iteration via `forall` and `for` loops over the elements mapped to by the indices in the array's domain.

19.2.7 Array Promotion of Scalar Functions

Array promotion of a scalar function is defined over the array type and the element type of the array. Array promotion has the same semantics as sequence promotion where the scalar type is the element type of the array and the promotion type is the array type. The only difference between sequence promotion and array promotion is that if a function returns a value, the promoted function returns an array of those values rather than a sequence of those values. The array is defined over the same domain as the array that was passed to the function. In the event of zipper promotion over multiple arrays or both arrays and sequences, the promoted function returns a value based on the first argument to the function that enables promotion.

Implementation note. In the current implementation, promotion always returns sequences.

Example. Whole array operations is a special case of array promotion of scalar functions. In the code

```
A = B + C;
```

if A , B , and C are arrays, this code assigns each element in A the element-wise sum of the elements in B and C .

19.2.8 Array Initialization

By default, the elements in an array are initialized to the default values associated with the element type of the array. There is an expectation that this default initialization can be overridden for performance reasons by explicitly marking the array type or variable.

The initialization expression in the declaration of an array can be based on the indices in the domain using special array declaration syntax that replaces both the type and initialization specifications of the declaration:

```

special-array-declaration :
    identifier-list indexed-array-type-part initialization-part

indexed-array-type-part :
    : array-type-forall-expression type

array-type-forall-expression :
    [ identifier in domain-expression ]

initialization-part :
    = expression

```

In this code, the *array-type-forall-expression* is syntactic sugar for surrounding the *initialization-part* with this basic forall-expression.

Given a domain expression *D*, an element type *t*, an expression *e* that is of type *t* or that can be implicitly converted to type *t*, then the declaration of array *A* given by

```
var A: [i in D] t = e;
```

is equivalent to

```
var A: [D] t = [i in D] e;
```

The scope of the forall expression is as in the rewritten part so the expression *e* can include references to index *i*.

19.3 Arithmetic Domains and Arrays

An arithmetic domain is a rectilinear set of Cartesian indices. Arithmetic domains are specified as a tuple of arithmetic sequences enclosed in square brackets.

19.3.1 Arithmetic Domain Literals

An arithmetic domain literal is specified by the following syntax:

```

arithmetic-domain-literal :
    [ arithmetic-sequence-expression-list ]

arithmetic-sequence-expression-list :
    arithmetic-sequence-expression
    arithmetic-sequence-expression , arithmetic-sequence-expression-list

arithmetic-sequence-expression :
    expression

```


Example. The expression `[1..5, 1..5]` defines a 5×5 arithmetic domain with the indices $(1, 1), (1, 2), \dots, (5, 5)$.

19.3.2 Arithmetic Domain Types

The type of an arithmetic domain is determined from the rank of the arithmetic domain (the number of arithmetic sequences that define it) and by an underlying integral type called the *dimensional index type* which must be identical to each of the integral element types of the arithmetic sequences that define the arithmetic domain. By default, the dimensional index type of an arithmetic domain is `int`.

The arithmetic domain type is specified by the syntax of a function call to the keyword `domain` that takes at least an argument called `rank` that is a parameter of type `int` and optionally an integral type named `dim_type`.

Example. The expression `[1..5, 1..5]` defines an arithmetic domain with type `domain(2, int)`.

19.3.3 Strided Arithmetic Domains

If the arithmetic sequences that define an arithmetic domain are strided, then the arithmetic domain is said to be strided.

The `by` operator can be applied to any arithmetic domain to create a strided arithmetic domain. It is predefined over an arithmetic domain and an integer or a tuple of integers. In the integer case, the arithmetic sequences in each dimension are strided by the integer. In the tuple of integers case, the size of the tuple must match the rank of the domain; the integers stride each dimension of the domain. If the arithmetic sequences are already strided, the strides applied by the `by` operator are multiplied to the strides of the arithmetic sequences.

19.3.4 Arithmetic Domain Indexing

Arithmetic domains support indexing by a value of type `int` that is at least one and no more than the rank of the array. Indexing into an arithmetic domain returns the arithmetic sequence associated with that dimension.

Example. In the code

```
for i in D(1) do
  for j in D(2) do
    writeln(A(i, j));
```

domain `D` is iterated over by two nested loops. The first dimension of `D` is iterated over in the outer loop. The second dimension is iterated over in the inner loop.

19.3.5 Arithmetic Array Indexing

In addition to being indexed by indices defined by their arithmetic domains, arithmetic arrays can be indexed directly by values of the dimensional index type where the number of values is equal to the rank of the array. This has the semantics of first moving the values into a tuple and then indexing into the array.

Example. Given the definition

```
var ij = (i,j);
```

the indexing expressions `A(ij)` and `A(i,j)` are equivalent.

19.3.6 Arithmetic Array Slicing

In addition to slicing an arithmetic array by an arithmetic domain, arithmetic arrays also support slicing by arithmetic sequences directly. If each dimension is indexed by an arithmetic sequence, this is equivalent to slicing the domain by an arithmetic domain defined by those arithmetic sequences.

Implementation note. It is currently required that each dimension be indexed by an arithmetic sequence. There is an expectation that indexing some dimensions directly by values of integral type will result in an array slice of a different rank.

19.3.7 Formal Arguments of Arithmetic Array Type

Formal arguments of arithmetic array type allow an arithmetic domain to be specified that does not match the arithmetic domain of the actual arithmetic array that is passed to the formal argument. In this case, the shape (size in each dimension and rank) of the domain of the actual array must match the shape of the domain of the formal array. The indices are translated in the formal array, which is a reference to the actual array.

Example. In the code

```
def foo(X: [1..5] int) { ... }
var A: [1..10 by 2] int;
foo(A);
```

the array `A` is strided and its elements can be indexed by the odd integers between one and nine. In the function `foo`, the array `X` references array `A` and the same elements can be indexed by the integers between one and five.

19.4 Sparse Domains and Arrays

Implementation note. Sparse domains are not yet implemented.

A sparse domain type is given by the syntax

```

sparse-domain-type :
    sparse domain ( domain-expression )

domain-expression :
    expression

```

A sparse domain is a domain that contains a subset of the indices in the domain specified by the *domain-expression*, sometimes called the *base domain*.

Arrays declared over sparse domains can be indexed by all of the indices in the base domain. If the index is not part of the sparse domain, the element returned is called the *unrepresented element*. It is an error to assign a value to the unrepresented element by indexing into the array and assigning it a value. The unrepresented element can be set to any value but by default contains the default value associated with the element type of the array.

19.4.1 Changing the Indices in Sparse Domains

Indices can be added to or removed from sparse domains. Sparse domains support a method `add` that takes an index and adds this index to the sparse domain. All arrays declared over this sparse domain can now be assigned values corresponding to this index.

Sparse domains support a method `remove` that takes an index and removes this index from the sparse domain. The values in the arrays indexed by the removed index are lost.

The operators `+=` and `-=` have special semantics for sparse domains; they are interpreted as calls to the `add` and `remove` methods respectively. The statement

```
D += i;
```

is equivalent to

```
D.add(i);
```

Similarly, the statement

```
D -= i;
```

is equivalent to

```
D.remove(i);
```

19.5 Indefinite Domains and Arrays

An indefinite domain type can be defined over any scalar type and is given by the following syntax:

```

sparse-domain-type :
    domain ( scalar-type )

scalar-type :
    type

```

A scalar type is any primitive type, tuple of scalar types, or class, record, or union where all of the fields have scalar types. Enumerated types are scalar types but domains declared over enumerated types are described in §19.7. Arrays declared over indefinite domains are dictionaries mapping from values to variables.

19.5.1 Changing the Indices in Indefinite Domains

Like with sparse domains, indices can be added or removed to indefinite domains. Indefinite domains support a method `add` that takes an index and adds this index to the indefinite domain. All arrays declared over this indefinite domain can now access elements corresponding to this index.

Indefinite domains support a method `remove` that takes an index and removes this index from the indefinite domain. The values in the arrays indexed by the removed index are lost.

The operators `+=` and `-=` have special semantics for indefinite domains; they are interpreted as calls to the `add` and `remove` methods respectively. The statement

```
D += i;
```

is equivalent to

```
D.add(i);
```

Similarly, the statement

```
D -= i;
```

is equivalent to

```
D.remove(i);
```

19.5.2 Testing Membership in Indefinite Domains

An indefinite domain supports a `member?` method that can test whether a particular value is part of the index set. It returns `true` if the index is in the indefinite domain and otherwise returns `false`.

19.6 Opaque Domains and Arrays

Implementation note. Opaque domains are not yet implemented.

An opaque domain is a form of indefinite domain where there is no algebra on the types of the indices. The indices are, in essence, opaque. The opaque domain type is given by the following syntax:

```
opaque-domain:
  opaque domain
```

New index values for opaque domains are explicitly requested via a method called `new`. Indices can be removed by a method called `remove`.

Opaque domains permit more efficient implementations than indefinite domains under the assumption that creation of new domain index values is rare.

19.7 Enumerated Domains and Arrays

Implementation note. Enumerated domains are not yet implemented.

Enumerated domains are a special case of indefinite domains where the indices are the constants defined by an enumerated type. Enumerated domains do not support the `add` or `remove` methods. All of the constants defined by the enumerated type are indices into the enumerated domain.

An enumerated domain is specified as an indefinite domain would be except the type is an enumerated type rather than some other scalar type.

19.8 Association of Arrays to Domains

When an array is declared, it is linked during execution to the domain over which it was declared. This linkage is constant and cannot be changed.

When indices are added or removed from a domain, the change impacts the arrays declared over this particular domain. In the case of adding an index, an element is added to the array and initialized to the default value associated with the element type. In the case of removing an index, the element in the array is removed.

When a domain is reassigned a new value, the array is also impacted. Values that could be indexed by both the old domain and the new domain are preserved in the array. Values that could only be indexed by the old domain are lost. Values that can only be indexed by the new domain have elements added to the new array and initialized to the default value associated with their type.

For performance reasons, there is an expectation that a method will be added to domains to allow non-preserving assignment, *i.e.*, all values in the arrays associated with the assigned domain will be lost.

19.9 Subdomains

Implementation note. Subdomains are not yet implemented.

A subdomain is a domain whose indices are indices of a *base domain*. A subdomain is specified by the following syntax:

```
subdomain-type :
  domain ( domain-expression )
```

The ordering of the indices in the subdomain is consistent with the ordering of the indices in the base domain.

Subdomains are verified during execution even as domains are reassigned. The indices in a subdomain are known to be indices in a domain, allowing for fast bounds-checking.

In the case of arithmetic domains, the subdomain literal may be composed of indefinite arithmetic sequences. In such cases, the omitted bounds of the indefinite arithmetic sequences are taken from the bounds of the base domain.

19.10 Predefined Functions and Methods on Domains

There is an expectation that this list of predefined functions and methods will grow.

```
def Domain.numIndices: dim_type
```

Returns the number of indices in the domain.

19.11 Predefined Functions and Methods on Arrays

There is an expectation that this list of predefined functions and methods will grow.

```
def Array.numElements: this.domain.dim_type
```

Returns the number of elements in the array.

20 Iterators

An iterator is a function that conceptually returns a sequence of values rather than simply a single value. Classes can be viewed as iterators if they implement a structural iterator interface.

20.1 Iterator Functions

The syntax of a function declaration is identical to that of a function declaration except that the keyword `def` is replaced with the keyword `iterator`. The body of the iterator may include `yield` statements alongside `return` statements. When a `yield` is encountered, the value is returned, but the iterator is not finished evaluating. It will continue from the point after the `yield` and can yield or return more values. When a `return` is encountered, the value is returned and the iterator finishes. An iterator also completes after the last statement in the iterator function is executed.

20.2 The Yield Statement

Yield statement can only appear in iterators. The syntax of the yield statement is given by

```
yield-statement :  
    yield expression ;
```

20.3 Iterator Calls

Iterator functions can be called within `for` or `forall` loops, in which case they are executed in an interleaved manner with the body of the loop, or can be called in any expression context, in which case they evaluate to a sequence of values.

20.3.1 Iterators in For and Forall Loops

When an iterator is accessed via a `for` or `forall` loop, the iterator is evaluated alongside the loop body in an interleaved manner. For each iteration, the iterator yields a value and the body is executed.

20.3.2 Iterators as Sequences

If an iterator function is accessed outside of the context of a `for` or `forall` loop iterator expression, then the iterator is iterated over in total and the expression evaluates to a sequence that contains the values returned by the iterator on each iteration.

Example. Given an iterator

```
iterator squares(n: int): int {  
    for i in 1..n do  
        yield i * i;  
    }
```

the expression `squares(5)` evaluates to the sequence `(/1, 4, 9, 16, 25/)`.

20.4 The Structural Iterator Interface

There is a structural interface that allows a class or record to be treated as if it were an iterator. The iterator interface is important for user-defined distributions.

Implementation note. This section describes the current structural iterator interface. This does not yet support optimized iteration for rank greater than one. As such, this iterator interface is incomplete.

The iterator interface defines iteration over a class or record by a cursor of some type, called the cursor type. The values returned by the iterator are of a possibly different type, called the value type. A class or record `classType` supports the iterator interface if it defines the following functions for cursor type `cursorType` and value type `valueType`:

```
def classType.getHeadCursor(): cursorType
def classType.getNextCursor(cursor: cursorType): cursorType
def classType.getValue(cursor: cursorType): valueType;
def classType.isValidCursor?(cursor: cursorType): bool
```

Iteration over a class or record `C` of type `classType` defined by

```
for i in C do
    ; // body of loop
```

is equivalent to the following loop:

```
var cursor = C.getHeadCursor();
while C.isValidCursor?(cursor) {
    var i = C.getValue(cursor);
    ; // body of loop
    cursor = C.getNextCursor(cursor);
}
```


21 Generics

Chapel supports generic functions and types that are parameterizable over both types and parameters. The generic functions and types look similar to non-generic functions and types already discussed.

21.1 Generic Functions

A function is generic if any of the following conditions hold:

- Some formal argument is specified with an intent of `type` or `param`.
- Some formal argument has no specified type and no default value.
- Some formal argument is specified with a queried type.
- The type of some formal argument is a generic type, e.g., `seq`.
- The type of some formal argument is an array type where either the element type is queried or omitted or the domain is queried or omitted.

These conditions are discussed in the next sections.

21.1.1 Formal Type Arguments

If a formal argument is specified with intent `type`, then a type must be passed to the function at the call site. A copy of the function is instantiated for each unique type that is passed to this function at a call site. The formal argument has the semantics of a type alias.

Example. The following code defines a function that takes two types at the call site and returns a 2-tuple of sequences where the element types of the two sequences are defined by the two type arguments:

```
def buildTupleOfSeqs(type t, type tt)
    return (seq(t), seq(tt));
```

This function is instantiated with “normal” function call syntax where the arguments are types:

```
var tupleOfSeqs = buildTupleOfSeqs(int, string);
tupleOfSeqs(1) == 1;
tupleOfSeqs(2) == "hello";
```

21.1.2 Formal Parameter Arguments

If a formal argument is specified with intent `param`, then a parameter must be passed to the function at the call site. A copy of the function is instantiated for each unique parameter that is passed to this function at a call site. The formal argument is a parameter.

Example. The following code defines a function that takes an integer parameter `p` at the call site as well as a regular actual argument of integer type `x`. The function returns a homogeneous tuple of size `p` where each component in the tuple has the value of `x`.

```
def fillTuple(param p: int, x: int) {
  var result: p*int;
  for param i in 1..p do
    result(i) = x;
  return result;
}
```

The function call `fillTuple(3, 3)` returns a 3-tuple where each component contains the value 3.

21.1.3 Formal Arguments without Types

If the type of a formal argument is omitted, the type of the formal argument is taken to be the type of the actual argument passed to the function at the call site. A copy of the function is instantiated for each unique actual type.

Example. The example from the previous section can be extended to be generic on a parameter as well as the actual argument that is passed to it by omitting the type of the formal argument `x`. The following code defines a function that returns a homogeneous tuple of size `p` where each component in the tuple is initialized to `x`:

```
def fillTuple(param p: int, x) {
  var result: p*x.type;
  for param i in 1..p do
    result(i) = x;
  return result;
}
```

In this function, the type of the tuple is taken to be the type of the actual argument. The call `fillTuple(3, 3.14)` returns a 3-tuple of real values `(3.14, 3.14, 3.14)`. The return type is `(real, real, real)`.

21.1.4 Formal Arguments with Queried Types

If the type of a formal argument is specified as a queried type, the type of the formal argument is taken to be the type of the actual argument passed to the function at the call site. A copy of the function is instantiated for each unique actual type. The queried type has the semantics of a type alias.

Example. The example from the previous section can be rewritten to use a queried type for clarity:

```
def fillTuple(param p: int, x: ?t) {
  var result: p*t;
  for param i in 1..p do
    result(i) = x;
  return result;
}
```

21.1.5 Formal Arguments of Generic Type

If the type of a formal argument is a generic type, the type of the formal argument is taken to be the type of the actual argument passed to the function at the call site with the constraint that the type of the actual argument is an instantiation of the generic type. A copy of the function is instantiated for each unique actual type.

Example. The following code defines a function that takes an actual argument that is a sequence and outputs the elements in a sequence without any space between the elements. The function is generic on the element type of the sequence.

```
def output(s: seq) {
  for e in s do
    write(e);
}
```

The generic types `integral` and `numeric` are generic types that can only be instantiated with, respectively, the signed and unsigned integral types and all of the numeric types.

21.1.6 Formal Arguments of Generic Array Types

If the type of a formal argument is an array where either the domain or the element type is queried or omitted, the type of the formal argument is taken to be the type of the actual argument passed to the function at the call site. If the domain is omitted, the domain of the formal argument is taken to be the domain of the actual argument.

21.2 Function Visibility in Generic Functions

Function visibility in generic functions is altered depending on the instantiation. When resolving function calls made within visible functions, the visible functions are taken from any call site at which the function is instantiated for each particular instantiation.

21.3 Generic Types

A class or record is generic if any of the following conditions hold:

- The class contains a specified or unspecified type alias.

- The class contains a field that is a parameter.
- The class contains a field that has no type and no initialization expression.
- The class contains a field where the type of the field is generic.

21.3.1 Type Aliases in Generic Types

Type aliases defined in a class or a record can be unspecified type aliases; type aliases that are not bound to a type. If a class or record contains an unspecified type alias, the aliased type must be specified whenever the type is used.

A type alias defined in a class or record is accessed as if it were a field. Moreover, it becomes an argument with intent `type` to the default constructor for that class or record. This makes the default constructor generic. When the default constructor is instantiated, the type is instantiated where the type bound to the type alias is set to be the type passed to the default constructor.

Example. The following code defines a class called `Node` that implements a linked list data structure. It is generic over the type of the element contained in the linked list.

```
class Node {
  type elt_type;
  var data: elt_type;
  var next: Node(elt_type);
}
```

The call `Node(real, 3.14)` creates a node in the linked list that contains the value `3.14`. The `next` field is set to `nil`. The type specifier `Node` is a generic type and cannot be used to define a variable. The type specifier `Node(real)` denotes the type of the `Node` class instantiated over `real`. Note that the type of the `next` field is specified as `Node(elt_type)`; the type of `next` is the same type as the type of the object that it is a field of.

21.3.2 Parameters in Generic Types

Parameters defined in a class or record do not require an initialization expression. If they do not have an initialization expression, the parameter must be specified whenever the type is used.

A parameter defined in a class or record is accessed as if it were a field. This access returns a parameter. Parameters defined in classes or records become arguments with intent `param` to the default constructor for that class or record. This makes the default constructor generic. When the default constructor is instantiated, the type is instantiated where the parameter is bound to the parameter passed to the default constructor.

Example. The following code defines a class called `IntegerTuple` that is generic over an integer parameter which defines the number of components in the class.

```
class IntegerTuple {
  param size: int;
  var data: size*int;
}
```

The call `IntegerTuple(3)` creates an instance of the `IntegerTuple` class that is instantiated over parameter `3`. The field `data` becomes a 3-tuple of integers. The type of this class instance is `IntegerTuple(3)`. The type specified by `IntegerTuple` is a generic type.

21.3.3 Fields without Types

If a field in a class or record has no specified type or initialization expression, the class or record is generic over the type of that field. The field must be specified when the class or record is constructed or specified. The field becomes an argument to the default constructor that has no specified type and no default value. This makes the default constructor generic. When the default constructor is instantiated, the type is instantiated where the type of the field becomes the type of the actual argument passed to the default constructor.

Example. The following code defines another class called `Node` that implements a linked list data structure. It is generic over the type of the element contained in the linked list. This code does not specify the element type directly in the class as a type alias but rather leaves omits the type from the `data` field.

```
class Node {
  var data;
  var next: Node(data) = nil;
}
```

A node with integer element type can be defined in the call to the constructor. The call `Node(1)` defines a node with the value 1. The code

```
var list = Node(1);
list.next = Node(2);
```

defines a two-element list with nodes containing the values 1 and 2.

21.3.4 Fields of Generic Types

If a field in a class or record is specified to have a generic type, then the class or record is generic over the type of this field and the type of the field is constrained to be an instantiation of the field's specified generic type.

21.3.5 Generic Methods

All methods bound to generic classes or records are generic over the implicit `this` argument and any other argument that is generic.

21.3.6 The *elt_type* Type

The common idiom of parameterizing a collection-oriented data type by a single element type has special syntactic support given by

```
of-type :
  type of type
```

This syntax is a short-hand for passing the second type by name `elt_type` as the only argument to the first type. Given the definition of `Node` in the example in §21.3.1, one can specify the type `Node(real)` or `Node(elt_type=real)` by writing `Node of real`.

21.4 Where Expressions

The instantiation of a generic function can be constrained by *where clauses*. A where clause is specified in the definition of a function (§13.1). When a function is instantiated, the expression in the where clause must be a parameter expression and must evaluate to either `true` or `false`. If it evaluates to `false`, the instantiation is rejected and the function is not a possible candidate for function resolution. Otherwise, the function is instantiated.

Example. Given two overloaded function definitions

```
def foo(x) where x.type == int { ... }
def foo(x) where x.type == real { ... }
```

the call `foo(3)` resolves to the first definition because when the second function is instantiated the where clause evaluates to false.

21.5 Example: A Generic Stack

```
class MyNode {
  type itemType;           // type of item
  var item: itemType;      // item in node
  var next: MyNode(itemType); // reference to next node (same type)
}

record Stack {
  type itemType;           // type of items
  var top: MyNode(itemType); // top node on stack linked list

  def push(item: itemType) {
    top = MyNode(itemType, item, top);
  }

  def pop() {
    if isEmpty? then
      halt("attempt to pop an item off an empty stack");
    var oldTop = top;
    top = top.next;
    return oldTop.item;
  }

  def isEmpty? return top == nil;
}
```

22 Parallelism and Synchronization

Chapel is an explicitly parallel programming language. Parallelism is introduced into a program by the user with the following three constructs: `forall`, `cobegin`, and `begin`. In addition, some operations on arrays, domains, and sequences are executed in parallel. Synchronization is provided with *synchronization variables* and *atomic* statements. To avoid any unintended implications, the terms *computation* and *sub-computation* will be used to refer to distinct, concurrently executing portions of the program.

22.1 The Forall Loop

The forall loop is a variant of the for loop that allows for the concurrent execution of the loop body. The for loop is described in §11.8. The syntax for the forall loop is given by

```
forall-statement :
  forall index-expression in iterator-expression do statement
  forall index-expression in iterator-expression block-level-statement
```

The forall loop evaluates the loop body once for each element in the sequence returned by the *iterator-expression*. Each instance of the forall loop's statement may be executed concurrently with each other, but this is not guaranteed. The compiler and runtime determine the actual concurrency based on the specification of the iterator of the loop. The keyword `ordered`, described in §22.1.2, can be used to constrain the parallelism to give a partial order on the sequence returned by an iterator.

Control continues with the statement following the forall loop only after each iteration has been completely evaluated. Control transfers out of a loop body via `break`, `continue`, and `return` are not permitted. Control can be transferred out of the loop via a `yield` statement.

Example. In the code

```
forall i in 1..N do
  a(i) = b(i);
```

the user has stated that the element-wise assignments can execute concurrently. This loop may be performed serially, with maximum concurrency where each loop body iteration instance is executed in a separate computation, or somewhere in between.

Implementation note. The forall loop is currently executed serially.

22.1.1 Alternative Forall Loop Syntax

The forall loop may be alternatively specified with a more concise syntax given by:

```
alternative-forall-statement :
  [index-expression in iterator-expression] statement
```

The semantics are unchanged.

Example. The previous forall example can be alternatively written as:

```
[i in 1..N] a(i) = b(i);
```

22.1.2 The Ordered Forall Loop

By default a forall loop allows complete concurrent evaluation of the iterator expression and among the loop instances. The keyword `ordered` can be used to constrain the general parallelism among instances of the loop to that expressed by an iterator. This allows an iterator to both define a sequence of values and to impose a partial order on that sequence. If the iterator expression is a sequence value, there is no effect. This has the same semantics as with the ordered expression which is explained in §22.5. The syntax is:

```
ordered-forall-statement :
  ordered forall index-expression in iterator-expression do statement
  ordered forall index-expression in iterator-expression block-level-statement
```

Example. In the code

```
ordered forall i in walk(root) do
  work(i);

iterator walk(n: node) {
  yield n;
  forall c in 0..n.numOfChildren {
    yield n.child[c];
  }
}
```

there is a constraint on the parallel execution such that the function `work` is evaluated on a node before any of its immediate children nodes. The work on sibling nodes can be executed concurrently.

Implementation note. The ordered forall loop is currently executed serially.

22.2 The Forall Expression

With syntax similar to the alternative forall loop statement, a forall expression can be used to enable concurrent evaluation of sub-expressions. The sub-expressions are evaluated once for each element in the iterator expression. The syntax of a forall expression is given by

```
forall-expression :
  [index-expression in iterator-expression] expression
```

The semantics of the forall expression are that a sequence of the expression is evaluated. However, for efficiency, a sequence may not be generated if the semantics are the same.

Example.

```
[i in S] f(i);
```

The function `f` is evaluated for each index in `S` and the result of this expression is a sequence containing the evaluated expressions.

Implementation note. Forall expressions are evaluated serially.

22.3 The Cobegin Statement

The `cobegin` statement is used to create parallelism among statements within a block statement. The `cobegin` statement syntax is

```
cobegin-statement :  
cobegin block-statement
```

Each statement within the block statement is executed concurrently and is considered a separate computation. Control continues after all of the statements within the block statement have been evaluated.

As with the `forall` loop, control transfers are not permitted either into or out of the `cobegin`'s block statement. Similarly, `yield` statements are allowed.

Variables declared in the `cobegin` statement are *single variables*, described in §22.7.1.

22.4 The Begin Statement

The `begin` statement spawns a computation to execute a statement. Control continues simultaneously with the statement following the `begin` statement. The `begin` statement is an unstructured way to create a new computation that is executed only for its side-effects. The syntax for the `begin` statement is given by

```
begin-statement :  
begin statement
```

The following statements cannot be contained in `begin`-statements: `break`-statements, `continue`-statements, `yield`-statements, and `return`-statements.

22.5 The Ordered Expression

Implementation note. The ordered expression is not yet implemented.

The `ordered` keyword can be used as an unary operator to suppress parallel execution among instances of an expression that can involve side-effects to memory. The `ordered` keyword does not inhibit parallelism within the sub-expression. The syntax is:

```
ordered-expression :  
ordered expression
```

Example. In the code

```
ordered [i in S] f(i)
```

`f` is a function and `S` is a sequence. Each instance of `f(i)` is executed once for each value in `S` and in sequence order. The `ordered` constraint does not propagate to inhibit parallelism within `f`.

22.6 The Serial Statement

Implementation note. The serial statement is not yet implemented.

The `serial` statement can be used to dynamically control the degree of parallelism. The syntax is:

```
serial-statement :
    serial expression block-level-statement
```

where the expression evaluates to a bool type. Independent of that value, the block-level statement is evaluated. If the expression is true, any dynamically encountered forall loop or cobegin statement is executed serially within the current computation. Any dynamically encountered begin-statement is executed serially with the current computation; no new computation is spawned. Control continues to the statement following the begin-statement after the begin-statement finishes.

Example. In the code

```
ordered forall i in walk(root) do
    work(i);

iterator walk(n: node) {
    yield n;
    serial n.depth > 4 forall c in 0..n.numOfChildren {
        yield n.child[c];
    }
}
```

the serial statement inhibits concurrent execution on the tree for nodes that are deeper than four levels in the tree.

There is an expectation that functions that may be executed in a serial context are cloned to avoid the overhead of testing and suppressing parallelism.

22.7 Synchronization Variables

Synchronization variables are used to coordinate computations that share data. The use of and assignment to these variables implicitly controls the execution order of the computation. There are two kinds of synchronization variables, *single* and *sync* variables. A single variable can only be assigned once during its lifetime. A sync variable can be assigned multiple times during its lifetime.

The normal use of and assignment to a synchronization variable is well suited for producer-consumer data sharing. Additional functions on synchronization variable are provided such that other traditional synchronization primitives, such as semaphores and mutexes, can be constructed.

22.7.1 Single Variables

A single (assignment) variable can only be assigned once during its lifetime. A use of a single variable before it is assigned causes the computation's execution to be suspended until the variable is assigned. Otherwise, the use proceeds as with normal variables and the computation continues. After a single assignment variable is assigned, all computations with pending uses resume in an unspecified order. A single variable is specified with a single type given by the following syntax:

```
single-type :  
  single type
```

Example. In the code

```
class Tree {  
  var is_leaf : bool;  
  var left    : Tree;  
  var right   : Tree;  
  var value   : int;  
  
  def sum() {  
    if (is_leaf) then  
      return value;  
  
    var x : single int;  
    begin x = left.sum();  
    var y = right.sum();  
    return x+y;  
  }  
}
```

the single variable `x` is assigned by an asynchronous computation created with the `begin` statement. The computation returning the sum waits on the reading of `x` until it has been assigned.

While a `cobegin` might be a more suitable formulation, this fragment creates an asynchronous computation to compute the sum of the left sub-tree while the main computation continues with the right sub-tree. The final reference to variable `x` will be delayed until the assignment to `x` completes and that value will be used as a summand.

When a single variable has an initializer, the evaluation of that initializer is implicitly performed as an asynchronous computation.

Example. The code

```
var x: single int = left.sum;
```

is equivalent to

```
var x: single int;  
x = left.sum;
```

Any variable declaration within a `cobegin` statement is implicitly treated as a single variable for references in other statements of the `cobegin` statement.

Example. In the code

```

def sum() {
  if (is_leaf) then
    return value;
  var z;
  cobegin {
    var x = left.sum();
    var y = right.sum();
    z = x+y;
  }
  return z;
}

```

the computation with assignment to `z` waits for the other computations to assign to `x` and `y` before it references `x` and `y` in order to assign to `z`. The variables `x` and `y` are implicitly single.

22.7.2 Sync Variables

A sync variable generalizes the single assignment variable to permit multiple assignments to the variable. A sync variable is logically either *full* or *empty*. When it is empty, computations that attempt to read that variable are suspended until the variable becomes full by the next assignment to it, which atomically changes the state to full. When the variable is full, a read of that variable consumes the value and atomically transitions the state to empty. If there is more than one computation waiting on a sync variable, one is non-deterministically selected to use the variable and resume execution. The other computations continue to wait for the next assignment.

If a computation attempts to assign to a sync variable that is full, the computation is suspended and the assignment is delayed. When the sync variable becomes empty, the computation is resumed and the assignment proceeds, transitioning the state back to full. If there are multiple computations attempting such an assignment, one is non-deterministically selected to proceed and the other assignments continue to wait until the sync variable is emptied again.

A sync variable is specified with a sync type given by the following syntax:

```

sync-type :
  sync

```

22.7.3 Additional Synchronization Variable Functions

Synchronization variables support additional methods that can be used to bypass their semantics to provide new ones. For sync variable `s`, the following functions are defined:

```

writeFE(s, v) // wait for full, assign s=v, and leave empty
writeXF(s, v) // no wait, assign s=v, and leave full
writeXE(s, v) // no wait, assign s=v, and leave empty
readFF(s)     // wait for full, leave full, and return s's value
readXF(s)     // no wait, leave full, and return s's value
readXX(s)     // no wait, leave F/E unchanged, and return s's value

```

For single variables `s` only `readFF` is defined.

22.7.4 Synchronization Variables of Record and Class Types

A variable of record or class type can be a single or sync variable. The semantics of single and sync variables are applied only to the variable and not to accesses of individual fields. A record or class type may have synchronization variable fields to get synchronization semantics on individual field accesses.

22.8 Memory Consistency Model

This section is forthcoming.

22.9 Atomic Statement

Implementation note. Atomic statements are not yet implemented.

The atomic statement creates an atomic transaction of a statement. The statement is executed with transaction semantics in that the statement executes entirely, the statement appears to have completed in a single order and serially with respect to other atomic statements, and no variable assignment is visible until the statement has completely executed.

This definition of an atomic statement is sometimes called *strong atomicity* because the semantics are atomic to the entire program. *Weak atomicity* is defined so that an atomic statement is atomic only with respect to other atomic statements. If the performance implications of strong atomicity are not tolerable, the semantics of atomic transactions may be revisited, and could become weaker.

The syntax for the atomic statement is given by:

```
atomic-statement :
  atomic statement
```

Example. The following code illustrates one possible use of atomic statements:

```
var found = false;
atomic {
  if head == obj {
    found = true;
    head = obj.next;
  } else {
    var last = head;
    while last != null {
      if last.next == obj {
        found = true;
        last.next = object.next;
        break;
      }
      last = last.next;
    }
  }
}
```

Inside the atomic statement is a sequential implementation of removing a particular object denoted by `obj` from a singly linked list. This is an operation that is well-defined, assuming only one computation is attempting it at a time. The atomic statement ensures that, for example, the value of `head` does not change after it is first in the first comparison and subsequently read to initialize `last`. The variables eventually owned by this computation are `found`, `head`, `obj`, and the various `next` fields on examined objects.

The effect of an atomic statement is dynamic.

Example. If there is a method associated with a list that removes an object, that method may not be parallel safe, but could be invoked safely inside an atomic statement:

```
atomic found = head.remove(obj);
```

23 Locality and Distribution

Implementation note. Programs can currently only run on a single locale. The abstractions described here are not yet implemented.

Chapel provides high-level abstractions that allow programmers to exploit locality by defining the affinity of data and computation. This is accomplished by associating both data objects and computations with abstract *locales*. To provide a higher-level mechanism, Chapel allows a mapping from domains to locales to be specified. This mapping is called a *distribution* and it guides that placement of variables associated with arrays and the placement of subcomputations defined over the domain.

Throughout this section, the term *local* refers to data that is associated with the locale that a computation is running on and *remote* refers to data that is not. We assume that there is some execution overhead associated with accessing data that may be remote compared to data known to be local.

23.1 Locales

A locale abstracts a processor or node in a parallel computer system, or the basic component in the computer system where local memory can be accessed uniformly.

23.1.1 The Locale Type

The identifier `locale` is a primitive type that abstracts a locale as described above. Both data and computations can be associated with a value of locale type. The only operators defined over locales are the equality and inequality comparison operators.

23.1.2 Predefined Locales Array

A predefined configuration variable defines the *execution environment* for a program. This environment is defined by the following definitions:

```
config const numLocales: int;
const Locales: [1..numLocales] locale;
const Global: locale;
```

The environment consists of constants which are fixed when the program begins execution. The variable `Global` holds a special value of `locale` type that can be distinct from the values stored in `Locales`. This value is used to denote an object or computation that has no defined affinity.

When a program starts, a single computation is executing. It is running on the locale given by `Locales(1)`.

23.1.3 Querying the Locale of a Variable

Every variable v is associated with some locale which can be queried using the following syntax:

```
locale-access :  
    expression . locale
```

When the *expression* is a class type, the locale is where the object is located rather than where the *expression* may be located.

23.2 Specifying Locales for Computation

When execution is proceeding on some locale, a computation can be associated with a different locale in two ways: via distributions as discussed in §23.3 or with an *on-statement* as discussed below.

23.2.1 On

The *on* statement controls on which locale a computation or data should be placed. The syntax of the *on* statement is given by

```
on-statement :  
    on expression do statement  
    on expression block-level-statement
```

If the *expression* is a value of `locale` type, the *statement* or *block-level-statement* is executed on the locale specified directly by the expression. Otherwise, the expression must be a variable and the locale is taken to be the locale where the variable is located. Execution continues after the *on-statement* after execution of the *statement* or *block-level-statement* completes.

If the locale that the *expression* refers to is equal to `Global`, then the locale is unspecified and is determined by the runtime and/or compiler.

Example. A common idiom is to use *on* in conjunction with *forall* to access an array decomposed over multiple locales. The code

```
forall i in D do on A(i) {  
    // some computation  
}
```

executes each iteration of the *forall* loop on the locale where the element of $A(i)$ is located.

By default, when new variables and data objects are created, they are created in the locale where the computation is running. This locale can be changed by using the *on* keyword. Variables can be defined within an *on-statement* to define them on a particular locale.

23.2.2 On and Iterators

When a loop iterates over a sequence specified by an iterator, *on*-statements inside the iterator control where the corresponding loop body is executed.

Example. An iterator over a distributed tree might include an iterator over the nodes as defined in the following code:

```
class Tree {
  var left, right: Tree;
  iterator nodes {
    on this yield this;
    if left then
      forall t in left.nodes do
        yield t;
    if right then
      forall t in right.nodes do
        yield t;
  }
}
```

Given this code and a binary tree of type `Tree` stored in variable `tree`, then we can use the `nodes` iterator to iterate over the tree with the following code:

```
forall t in tree.nodes {
  // body executed on t as specified in nodes
}
```

Here, each instance of the body of the `forall` loop is executed on the locale where the corresponding object `t` is located. This is specified in the `nodes` iterator where the `on` keyword is used. In the case of zipper or tensor product iteration, the location of execution is taken from the first iterator. This can be overridden by explicitly using `on` in the body of the loop or by reordering the product of iteration.

23.3 Distributions

A mapping from domain index values to locales is called a *distribution*.

23.3.1 Distributed Domains

A domain for which a distribution is specified is referred to as a *distributed domain*. A domain supports a method, `locale`, that maps index values in the domain to locales that correspond to the domain's distribution.

Iteration over a distributed domain implicitly executes the control computation in the domain of the associated locale. Similarly, when iterating over the elements of an array defined over a distributed domain, the controlled computations are determined by the distribution of the domain. If there are conflicting distributions in product iterations, the locale of the computation is taken to be the first component in the product.

Example. If `D` is a distributed domain, then in the code

```
forall d in D {
  // body
}
```

the body of the loop is executed in the locale where the index `d` maps to by the distribution of `D`.

23.3.2 Distributed Arrays

Arrays defined over a distributed domain will have the element variables stored on the locale determined by the distribution. Thus, if d is an index of distributed domain D and A is an array defined over that domain, then $A(d).locale$ is the locale to which d maps to according to D .

23.3.3 Undistributed Domains and Arrays

If a domain or an array does not have a distributed part, the domain or array is not distributed and exists only on the locale on which it is defined.

23.4 Standard Distributions

Standard distributions include the following:

- The block distribution `Block`
- The cyclic distribution `Cyclic`
- The block-cyclic distribution `BlockCyclic`
- The cut distribution `Cut`

A design goal is that all standard distributions are defined with the same mechanisms that user-defined distributions (§23.5) are defined with.

23.5 User-Defined Distributions

This section is forthcoming.

24 Reductions and Scans

Chapel provides a set of built-in reductions and scans with parallel semantics, a mechanism for defining more reductions and scans with efficient implementations, and syntact support to make reductions and scans easy to use.

24.1 Reduction Expressions

The syntax for a reduction expression is given by:

```

reduce-expression :
  reduce-operator reduce expression
  type reduce expression

reduce-scan-operator: one of
  + * && || & | ^ min max

```

The expression on the right-hand side of the reduction can be of any type that can be iterated over, *e.g.*, a sequence or array.

The built-in reductions are defined in *reduce-scan-operator*. These include, in order, sum, product, logical and, logical or, bitwise and, bitwise or, bitwise exclusive or, minimum, and maximum.

User-defined reductions are specified by preceding the keyword `reduce` by the class type that implements the reduction interface as described in §24.3.

24.2 Scan Expressions

The syntax for a scan expression is given by:

```

scan-expression :
  reduce-scan-operator scan expression
  type scan expression

```

The expression on the right-hand side of the scan can be of any type that can be iterated over, *e.g.*, a sequence or array.

The built-in scans are defined in *reduce-scan-operator*. These are identical to the built-in reductions and are described in §24.1.

User-defined scans are specified by preceding the keyword `scan` by the class type that implements the scan interface as described in §24.3.

24.3 User-Defined Reductions and Scans

User-defined reductions and scans are supported via class definitions where the class implements a structural interface. The definition of this structural interface is forthcoming. The following paper sketched out such an interface:

S. J. Deitz, D. Callahan, B. L. Chamberlain, and L. Snyder. **Global-view abstractions for user-defined reductions and scans.** In *Proceedings of the Eleventh ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, 2006.

25 Input and Output

Chapel provides a built-in `file` type to handle input and output to files using functions and methods called `read`, `write`, and `writeln`.

25.1 The *file* type

The `file` type contains the following fields:

- The `filename` field is a `string` that contains the name of the file.
- The `path` field is a `string` that contains the path of the file.
- The `mode` field is a `string` that indicates whether the file is being read or written.
- The `style` field can be set to `text` or `binary` to specify that reading from or writing to the file should be done with text or binary formats.

These fields can be modified any time that the file is closed.

The `mode` field supports the following strings:

- `"r"` The file can be read.
- `"w"` The file can be written.

Implementation note. The `style` field is not yet implemented. All input and output is done in text mode. All files must be text files.

There is an expectation that there will be more styles to control the default reading and writing methods.

The `file` type supports the following methods:

- The `open` method opens the file for reading and/or writing.
- The `close` method closes the file for reading and/or writing.
- The `isOpen` method returns `true` if the file is open for reading and/or writing, and otherwise returns `false`.
- The `flush()` method flushes the file, finishing outstanding reading and writing.

Additionally, the `file` type supports the methods `read`, `write`, and `writeln` for input and output as discussed in §25.4 and §25.5.

25.2 Standard files *stdout*, *stdin*, and *stderr*

The files `stdout`, `stdin`, and `stderr` are predefined and map to standard output, standard input, and standard error as implemented in a platform dependent fashion.

25.3 The *write*, *writeln*, and *read* functions

The built-in function `write` can take an arbitrary number of arguments and writes each of the arguments out in turn to `stdout`. The built-in function `writeln` has the same semantics as `write` but outputs an *end-of-line* character after writing out the arguments. The built-in function `read` can take an arbitrary number of arguments and reads each of the arguments in turn from `stdin`.

These functions are wrappers for the methods on files described next.

Example. The `writeln` wrapper function allows for a simple implementation of the *Hello-World* program:

```
writeln("Hello, World!");
```

25.4 The *write* and *writeln* method on files

The `file` type supports methods `write` and `writeln` for output. These methods are defined to take an arbitrary number of arguments. Each argument is written in turn by calling the `write` method bound to that type. Default `write` methods are bound to any type that the user does not explicitly create one for.

A lock is used to ensure that output is serialized across multiple computations.

25.5 The *read* method on files

The `file` type supports a method `read` that takes an arbitrary number of arguments. Each argument is read in turn by calling a method also bound to the `file` type that takes a single argument of that type.

25.6 User-Defined *read* and *write* methods

To define the output for a given type, the user must define a method called `write` on that type that takes a single argument of file type. If such a method does not exist, a default method is created.

25.7 Default *read* and *write* methods

Default `write` methods are created for all types for which a user `write` method is not defined. They have the following semantics:

- **arrays** Outputs the elements of the array in row-major order where rows are separated by line-feeds and blank lines are used to separate other dimensions.
- **domains** Outputs the dimensions of the domain enclosed by [and].
- **sequences** Outputs the elements of the sequence in order delimited by the sequence delimiters (/ and /) and separated by commas.
- **arithmetic sequences** Outputs the lower bound of the sequence followed by . . followed by the upper bound of the sequence. If the stride of the sequence is not one, the output is additionally followed by the word `by` followed by the stride of the sequence.
- **tuples** Outputs the components of the tuple in order delimited by (and), and separated by commas.
- **classes** Outputs the values within the fields of the class prefixed by the name of the field and the character =. Each field is separated by a comma. The output is delimited by { and }.
- **records** Outputs the values within the fields of the class prefixed by the name of the field and the character =. Each field is separated by a comma. The output is delimited by (and).

Default `read` methods are created for all types for which a user `read` method is not defined. The default `read` methods are defined to read in the output of the default `write` method.

26 Standard Modules

This section describes predefined functions that are available to any Chapel program as well as a set of standard modules that, when used, define a set of functions and types available to Chapel programs. The standard modules include the following:

<code>BitOps</code>	Bit manipulation routines
<code>Math</code>	<i>(used by default)</i> Math routines
<code>Random</code>	Random number generation routines
<code>Standard</code>	<i>(used by default)</i> Basic routines
<code>Time</code>	Types and routines related to time
<code>Types</code>	<i>(used by default)</i> Routines related to primitive types

There is an expectation that each of these modules will be extended and that more standard modules will be defined.

26.1 BitOps

The module `BitOps` defines routines that manipulate the bits of values of integral types.

```
def bitPop(i: integral): int
```

Returns the number of bits set to one in the integral argument `i`.

```
def bitMatMultOr(i: uint(64), j: uint(64)): uint(64)
```

Returns the bitwise matrix multiplication of `i` and `j` where the values of `uint(64)` type are treated as 8×8 bit matrices and the combinator function is bitwise or.

```
def bitRotLeft(i: integral, shift: integral): i.type
```

Returns the value of the integral argument `i` after rotating the bits to the left `shift` number of times.

```
def bitRotRight(i: integral, shift: integral): i.type
```

Returns the value of the integral argument `i` after rotating the bits to the right `shift` number of times.

26.2 Math

The module `Math` defines routines for mathematical computations. This module is used by default; there is no need to explicitly use this module. The `Math` module defines routines that are derived from and implemented via the standard C routines defined in `math.h`.

```
def abs(i: int(?w)): int(w)
def abs(i: uint(?w)): uint(w)
def abs(x: real): real
def abs(x: complex): real
```

Returns the absolute value of the argument.

```
def acos(x: real): real
```

Returns the arc cosine of the argument. It is an error if x is less than -1 or greater than 1 .

```
def acosh(x: real): real
```

Returns the inverse hyperbolic cosine of the argument. It is an error if x is less than 1 .

```
def asin(x: real): real
```

Returns the arc sine of the argument. It is an error if x is less than -1 or greater than 1 .

```
def asinh(x: real): real
```

Returns the inverse hyperbolic sine of the argument.

```
def atan(x: real): real
```

Returns the arc tangent of the argument.

```
def atan2(y: real, x: real): real
```

Returns the arc tangent of the two arguments. This is equivalent to the arc tangent of y / x except that the signs of y and x are used to determine the quadrant of the result.

```
def atanh(x: real): real
```

Returns the inverse hyperbolic tangent of the argument. It is an error if x is less than -1 or greater than 1 .

```
def cbrt(x: real): real
```

Returns the cube root of the argument.

```
def ceil(x: real): real
```

Returns the value of the argument rounded up to the nearest integer.

```
def conjg(a: complex(?w)): complex(w)
```

Returns the conjugate of a .

```
def cos(x: real): real
```

Returns the cosine of the argument.

```
def cosh(x: real): real
```

Returns the hyperbolic cosine of the argument.

```
def erf(x: real): real
```

Returns the error function of the argument defined as

$$\frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$$

for the argument x .

```
def erfc(x: real): real
```

Returns the complementary error function of the argument. This is equivalent to $1.0 - \text{erf}(x)$.

```
def exp(x: real): real
```

Returns the value of e raised to the power of the argument.

```
def exp2(x: real): real
```

Returns the value of 2 raised to the power of the argument.

```
def expm1(x: real): real
```

Returns one less than the value of e raised to the power of the argument.

```
def floor(x: real): real
```

Returns the value of the argument rounded down to the nearest integer.

```
def lgamma(x: real): real
```

Returns the natural logarithm of the absolute value of the gamma function of the argument.

```
def log(x: real): real
```

Returns the natural logarithm of the argument. It is an error if the argument is less than or equal to zero.

```
def log10(x: real): real
```

Returns the base 10 logarithm of the argument. It is an error if the argument is less than or equal to zero.

```
def loglp(x: real): real
```

Returns the natural logarithm of $x+1$.

```
def log2(i: int(?w)): int(w)
def log2(i: uint(?w)): uint(w)
def log2(x: real): real
```

Returns the base 2 logarithm of the argument. It is an error if the argument is less than or equal to zero.

```
def nearbyint(x: real): real
```

Returns the rounded integral value of the argument determined by the current rounding direction.

```
def rint(x: real): real
```

Returns the rounded integral value of the argument determined by the current rounding direction.

```
def round(x: real): real
```

Returns the rounded integral value of the argument. Cases halfway between two integral values are rounded towards zero.

```
def sin(x: real): real
```

Returns the sine of the argument.

```
def sinh(x: real): real
```

Returns the hyperbolic sine of the argument.

```
def sqrt(x: real): real
```

Returns the square root of the argument. It is an error if the argument is less than zero.

```
def tan(x: real): real
```

Returns the tangent of the argument.

```
def tanh(x: real): real
```

Returns the hyperbolic tangent of the argument.

```
def tgamma(x: real): real
```

Returns the gamma function of the argument defined as

$$\int_0^{\infty} t^{x-1} e^{-t} dt$$

for the argument x .

```
def trunc(x: real): real
```

Returns the nearest integral value to the argument that is not larger than the argument in absolute value.

26.3 Random

The module `Random` supports the generation of pseudo-random values and streams of values. The current interface is minimal and should be expected to grow and evolve over time.

```
class RandomStream
```

Implements a pseudo-random stream of values. Our current implementation generates the values using a linear congruential generator. In future versions of this module, the `RandomStream` class will offer a wider variety of algorithms for generating pseudo-random values.

```
const RandomStream.seed: int(64)
```

The seed value for the random stream. If no seed is specified in the constructor, the millisecond value of the current time is used. The seed value must be an odd integer. If an even integer is supplied, the class constructor will increment it to obtain an odd integer.

```
def RandomStream.fillRandom(x:[?D] real)
```

Fill the argument array, `x`, with the next $|D|$ values of the pseudo-random stream. Arrays of arbitrary rank can be passed to this routine, causing the 1D stream of values to be mapped to the array elements according to the array's default iteration order. Once our implementation supports distributed arrays, this routine is intended to fill the array's values in parallel.

```
def RandomStream.fillRandom(x:[?D] complex)
```

Similar to the previous routine, but for use with arrays of complex values. The elements are filled in the same order as above except that pairs of values from the stream are assigned to each element, the first to the real component, the second to the imaginary. As this module matures, we will support `fillRandom` for arrays of other element types as well.

```
SeedGenerator
```

A symbol that can be used to generate seed values for the `RandomStream` class.

```
SeedGenerator.clockMS
```

Generates a seed value using the milliseconds value from the current time. As this module matures, `SeedGenerator` will support additional mechanisms for generating seed values.

```
def fillRandom(x:[], initseed: int(64))
```

A routine provided for convenience to support filling an array `x` with pseudo-random values without explicitly constructing an instance of the `RandomStream` class, useful for filling a single array or multiple arrays which require no coherence between them. The `initseed` parameter corresponds to the `seed` member of the `RandomStream` class and will default to the milliseconds value of the current time if no seed value is provided.

26.4 Standard

def `ascii(s: string): int`

Returns the ASCII code number of the first letter in the argument `s`.

def `assert(test: bool) {`

Exits the program if `test` is false and prints to standard error the location in the Chapel code of the call to `assert`. If `test` is true, no action is taken.

def `assert(test: bool, args ...?numArgs) {`

Exits the program if `test` is false and prints to standard error the location in the Chapel code of the call to `assert` as well as the rest of the arguments to the call. If `test` is true, no action is taken.

def `complex.re: real`

Returns the real component of the complex number.

def `complex.im: real`

Returns the imaginary component of the complex number.

def `complex.=re(f: real)`

Sets the real component of the complex number to `f`.

def `complex.=im(f: real)`

Sets the imaginary component of the complex number to `f`.

def `exit(status: int)`

Exits the program with code `status`.

def `halt() {`

Exits the program and prints to standard error the location in the Chapel code of the call to `halt` as well as the rest of the arguments to the call.

def `halt(args ...?numArgs) {`

Exits the program and prints to standard error the location in the Chapel code of the call to `halt` as well as the rest of the arguments to the call.

def `length(s: string): int`

Returns the number of characters in the argument `s`.

def `max(x, y...?k)`

Returns the maximum of the arguments when compared using the “greater-than” operator. The return type is inferred from the types of the arguments as allowed by implicit coercions.

def `min(x, y...?k)`

Returns the minimum of the arguments when compared using the “less-than” operator. The return type is inferred from the types of the arguments as allowed by implicit coercions.

def `string.substring(x): string`

Returns a value of string type that is a substring of the base expression. If `x` is `i`, a value of type `int`, then the result is the *i*th character. If `x` is an arithmetic sequence, the result is the substring where the characters in the substring are given by the values in the arithmetic sequence.

26.5 Time

The module `Time` defines routines that query the system time and a record `Timer` that is useful for timing portions of code.

record `Timer`

A timer is used to time portions of code. Its semantics are similar to a stopwatch.

enum `TimeUnits` { `microseconds`, `milliseconds`, `seconds`, `minutes`, `hours` };

The enumeration `TimeUnits` defines units of time. These units can be supplied to routines in this module to specify the desired time units.

def `getCurrentDate(): (int, int, int)`

Returns the year, month, and day of the month as integers. The year is the year since 0. The month is in the range 1 to 12. The day is in the range 1 to 31.

def `getCurrentTime(unit: TimeUnits = seconds): real`

Returns the elapsed time since midnight in the units specified.

def `Timer.clear()`

Clears the elapsed time stored in the `Timer`.

def `Timer.elapsed(unit: TimeUnits = seconds): real`

Returns the cumulative elapsed time, in the units specified, between calls to `start` and `stop`. If the timer is running, the elapsed time since the last call to `start` is added to the return value.

def `Timer.start()`

Start the timer. It is an error to start a timer that is already running.

def `Timer.stop()`

Stops the timer. It is an error to stop a timer that is not running.

def `sleep(t: uint)`

Delays the computation for `t` seconds.

26.6 Types

def `numBits(type t): int`

Returns the number of bits used to store the values of type `t`. This is implemented for all numeric types and `bool`.

def `max(type t): t`

Returns the maximum value that can be stored in type `t`. This is implemented for all numeric types.

def `min(type t): t`

Returns the minimum value that can be stored in type `t`. This is implemented for all numeric types.

Index

- &, 54
- &&, 56
- &&=, 62
- &=, 62
- (/, 97
- *, 51
- * tuples, 95
- **, 54
- **=, 62
- *=, 62
- +, 50, 59
- + (unary), 49
- +=, 62
- , 50
- (unary), 49
- =, 62
- /, 52
- /=, 62
- <, 57
- <<, 55
- <<=, 62
- <=, 58
- =, 62
- ==, 58
- >, 57
- >=, 58
- >>, 55
- >>=, 62
- ?, 27, 46
- #, 59, 98
- #=, 62
- %, 53
- %=, 62
- ~, 54
- ^, 55
- ^=, 62
- _, 94
- argv, 70
- arithmetic sequences, 103
 - indefinite, 103
 - integral element type, 103
 - literals, 103
 - strided, 103
- arrays, 11, 105, 107
 - association to domains, 22
 - slicing, 15
 - arithmetic, 110
 - arithmetic, strided, 111
 - as formal arguments, 18, 109, 112
 - assignment, 14, 108
 - association to domains, 115
 - distributed, 136
 - enumerated, 115
 - indefinite, 113
 - indexing, 108
 - initialization, 110
 - opaque, 114
 - predefined functions, 116
 - promotion, 109
 - slice, 112
 - slicing, 14, 108
 - sparse, 112
 - types, 107
- assignment, 62
 - tuples, 93
- atomic, 131
- atomic transactions, 131
- automatic memory management, 88
- begin, 20, 127
- block, 61
- block level statement, 62
- bool, 31
- by, 59
- case sensitivity, 27
- casts, 47
- class, 83
- classes, 22, 83
 - assignment, 83
 - constructors, 86
 - declarations, 83
 - fields, 83
 - generic, 121
 - getters, 86
 - indexing, 85
 - inheritance, 87
 - methods, 84
 - methods without parentheses, 85
 - nested, 88
 - promotion, 88
 - setters, 86
- cobegin, 127
- command-line arguments, 70
- comments, 27

- compiler errors
 - user-defined, 29
- complex
 - casts from tuples, 47
- complex, 32
- conditional
 - expression, 60
 - statement, 63
- conditional statement
 - dangling else, 63
- config, 12, 40
- configuration variables, 21
- const, 39
- constants
 - compile-time, 39
 - runtime, 39
- conversions
 - bool, 42
 - class, 42, 43
 - enumeration, 41, 42
 - explicit, 42
 - implicit, 41
 - numeric, 41, 42
 - parameter, 42
 - record, 42, 43
- def, 73
- default values, 75
- distributions, 135
- domains, 11, 105
 - arithmetic, 14, 110
 - arithmetic literals, 110
 - arithmetic, strided, 111
 - as formal arguments, 106
 - assignment, 106
 - association to arrays, 115
 - distributed, 135
 - enumerated, 115
 - indefinite, 113
 - index types, 106
 - opaque, 114
 - predefined functions, 116
 - promotion, 107
 - sparse, 112
 - subdomains, 115
 - types, 11, 105
- dynamic dispatch, 87
- else, 60, 63
- elt_type, 123
- enumerated types, 18, 34
- execution environment, 133
- exploratory programming, 70
- expression
 - as a statement, 62
- expression statement, 62
- fields
 - generic types, 123
 - without types, 123
- file type, 139
 - methods, 139
 - standard files, 140
- for, 65, 66
- for loops, 65
 - parameters, 66
- forall, 125
- forall expressions, 126
- forall loops, 125
 - alternative syntax, 125
 - ordered, 126
- formal arguments, 74
 - arithmetic arrays, 112
 - array types, 121
 - defaults, 75
 - domains, 106
 - generic types, 121
 - naming, 75
 - queried types, 120
 - tuples, 95
 - without types, 120
- function calls, 45, 74
- functions, 16, 73
 - as lvalues, 76, 77
 - candidates, 78
 - generic, 16, 119
 - most specific, 79
 - nested, 80
 - overloading, 16, 77
 - setters, 77
 - syntax, 73
 - variable number of arguments, 80
 - visible, 78
 - with class arguments, 78
- generics
 - function visibility, 121
 - functions, 16, 119
 - methods, 123
 - types, 121
- Global, 133, 134
- high, 103

- identifiers, 27
- if, 60, 63
- imaginary, 32
- in, 76
- indexing, 46
- inheritance, 87
- inout, 76
- int, 31
- integral, 121
- intents, 75
 - in, 76
 - inout, 76
 - out, 76
 - param, 120
 - type, 119
- iterator, 117
- iterators, 15, 117
 - and sequences, 117
 - on, 135
 - structural interface, 118
- Jacobi method
 - example, 11
- keywords, 28
- length
 - on sequences, 101
- let, 60
- literals
 - primitive type, 33
- local, 133
- locale, 133, 134
- Locales, 133
- locales, 133
- low, 103
- lvalue, 47
- main, 69
- member access, 46, 84
- memory consistency model, 131
- module, 69
- modules, 16, 69
 - and files, 71
 - nested, 71
 - using, 66, 70
- multiple inheritance, 88
- named arguments, 75
- numeric, 121
- numLocales, 133
- on, 134
- operators
 - arithmetic, 49
 - associativity, 48
 - bitwise, 54
 - logical, 56
 - overloading, 77
 - precedence, 48
 - relational, 57
- ordered, 126, 127
- out, 76
- param, 39, 66
- parameters, 39
 - configuration, 40
 - in classes or records, 122
- read, 140
 - default methods, 141
 - on files, 140
- read, 140
- readFF, 130
- readXF, 130
- readXX, 130
- real, 32
- record, 22
- record, 89
- records, 89
 - assignment, 90
 - differences with classes, 89
 - equality, 90
 - generic, 121
 - inequality, 90
 - inheritance, 89
- reductions, 15
- remote, 133
- reserved words, 28
- reshape, 102
- return, 74
- reverse, 101
- scalar promotion, 99
 - tensor product iteration, 100
 - zipper iteration, 99
- select, 18, 63
- seq, 97
- sequences
 - and conditional expressions, 101
 - and conditional statements, 100
 - and select statements, 101
 - and while statements, 100
 - arithmetic, 103

- assignment, 97
 - cast to tuples, 47
 - casts from tuples, 104
 - indexing, 98, 99
 - iteration, 98
 - literals, 97
 - promotion, 99
 - rank, 97
 - types, 97
- serial, 128
- single, 129
- spread, 102
- statement, 61
- stride, 103
- string, 33
- strings
 - format string casts, 47
- subdomains, 115
- sync, 130
- synchronization variables, 20, 128
 - built-in functions on, 130
 - implicit in `cobegin`, 129
 - of class type, 131
 - of record type, 131
 - single, 129
 - sync, 130
- tensor product iterator, 65
- then, 60, 63
- this, 85, 88
- transpose, 102
- tuples, 15, 93
 - assignment, 93
 - cast to sequences, 47
 - casts from sequences, 104
 - destructuring, 94
 - homogeneous, 94
 - indexing, 95
 - types, 93
 - variable declarations, 94
- type aliases, 22, 35
 - in classes or records, 122
- type inference, 14, 38
- type select statements, 66
- types
 - primitive, 31
 - fields, 91
 - type select, 91
 - use, 66
- variables
 - configuration, 12, 40
 - declarations, 37
 - default initialization, 38
 - global, 38
 - local, 38
- when, 63
- where, 124
- while, 64
- while loops, 64
- white space, 27
- write, 140
 - default methods, 14, 141
 - on files, 140
- write, 140
- writeFE, 130
- writeln, 14, 140
- writeXE, 130
- writeXF, 130
- yield, 117
- zipper iteration, 65
- uint, 31
- union, 91
- unions, 91
 - assignment, 91