

# Chapel Background & Overview

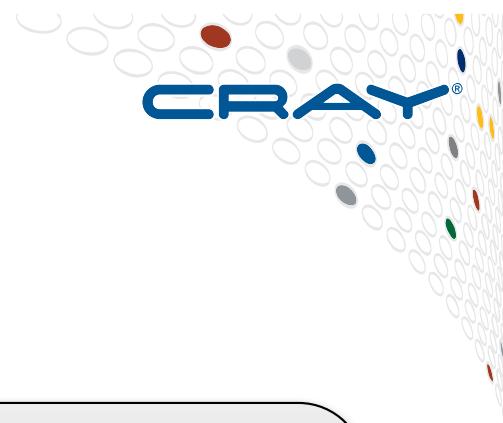


---

COMPUTE

| STORE


| ANALYZE



# Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.





# Motivation for Chapel

**Q: Can a single language be...**


- ...as productive as Python?
- ...as fast as Fortran?
- ...as portable as C?
- ...as scalable as MPI?
- ...as fun as <your favorite language here>?

**A: We believe so.**



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# The Challenge

**Q: So why don't we have such languages already?**

**A: ~~Technical challenges?~~**

- while they exist, we don't think this is the main issue...

**A: Due to a lack of...**

- ...long-term efforts
- ...resources
- ...community will
- ...co-design between developers and users
- ...patience

***Chapel is our attempt to reverse this trend***



COMPUTE | STORE | ANALYZE


Copyright 2017 Cray Inc.

# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures:




# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel:




# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory):




# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory multicore):



# STREAM Triad: MPI



MPI

```
#include <hpcc.h>

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );


    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
        0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
        sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
```



```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory
(%d).\n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 1.0;
}

scalar = 3.0;

for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
```



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# STREAM Triad: MPI+OpenMP



## MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
        0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
        sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory
(%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }

#ifndef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 1.0;
    }

    scalar = 3.0;

#ifndef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
```



COMPUTE

STORE

ANALYZE

# STREAM Triad: MPI+OpenMP vs. CUDA

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );

    return errCount;
}

int HPCC_LocalVectorSize(HPCC_Params *params, int len, double scalar);
double VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );

a = HPCC_XMALLOC( double, VectorSize );
b = HPCC_XMALLOC( double, VectorSize );
c = HPCC_XMALLOC( double, VectorSize );

if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

#ifdef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 1.0;
}
scalar = 3.0;

#ifdef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++) {
    a[j] = b[j]+scalar*c[j];
}

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);

return 0;
}
```

## CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid

    set_array<<<dimGrid, dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid, dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid, dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad( float *a, float *b, float *c,
                             float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```



COMPUTE

STORE

ANALYZE



# Why so many programming models?

HPC tends to approach programming models bottom-up:

Given a system and its core capabilities...

...provide features that can access the available performance.

- portability, generality, programmability: not strictly necessary.

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	MPI	executable
Intra-node/multicore	OpenMP / pthreads	iteration/task
Instruction-level vectors/threads	pragmas	iteration
GPU/accelerator	CUDA / Open[MP CL ACC]	SIMD function/task

**benefits:** lots of control; decent generality; easy to implement

**downsides:** lots of user-managed detail; brittle to changes



# Rewinding a few slides...

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );

    return errCount;
}

int HPCC_LocalVectorSize(HPCC_Params *params, int len) {
    double scalar;
    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );
    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }
    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 1.0;
    }
    scalar = 3.0;
    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];
    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
    return 0;
}
```

## CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid
        set_array<<<dimGrid, dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid, dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid, dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad( float *a, float *b, float *c,
                             float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```



COMPUTE

STORE

ANALYZE

# STREAM Triad: Chapel

## MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *pa,
int myRank, commSize;
int rv, errCount;
MPI_Comm comm = MPI_COMM_WORLD;
MPI_Comm_size(comm, &commSize);
MPI_Comm_rank(comm, &myRank);

rv = HPCC_Stream( params, 0 == myR
MPI_Reduce( &rv, &errCount, 1, MPI
return errCount;

int HPCC_Stream(HPCC_Params *params,
register int j;
double scalar;
VectorSize = HPCC_LocalVectorSize();
a = HPCC_XMALLOC( double, VectorSi
b = HPCC_XMALLOC( double, VectorSi
c = HPCC_XMALLOC( double, VectorSi

if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d). \n", VectorSi
        fclose( outFile );
    }
}
```

## Chapel

```
config const m = 1000,
alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 1.0;

A = B + alpha * C;
```


**dmapped** ...;

the special sauce

N);  
N);

c, d\_a, scalar, N);

CudaThreadSynchronize();



Philosophy: Good, *top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.



COMPUTE

STORE

ANALYZE



# What is Chapel?

**Chapel:** A productive parallel programming language

- portable
- open-source
- a collaborative effort

## Goals:

- Support general parallel programming
  - “any parallel algorithm on any parallel hardware”
- Make parallel programming at scale far more productive



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# What does “Productivity” mean to you?

## Recent Graduates:

“something similar to what I used in school: Python, Matlab, Java, ...”

## Seasoned HPC Programmers:

“that sugary stuff that I don’t need because I ~~was born to suffer~~  
want full control  
to ensure performance”

## Computational Scientists:

“something that lets me express my parallel computations  
without having to wrestle with architecture-specific details”

## Chapel Team:

“something that lets computational scientists express what they want,  
without taking away the control that HPC programmers want,  
implemented in a language as attractive as recent graduates want.”



COMPUTE | STORE | ANALYZE



# Chapel is Portable

- Chapel is designed to be hardware-independent
- The current release requires:
  - a C/C++ compiler
  - a \*NIX environment (Linux, OS X, BSD, Cygwin, ...)
  - POSIX threads
  - UDP, MPI, or RDMA (if distributed memory execution is desired)
- Chapel can run on...
  - ...laptops and workstations
  - ...commodity clusters
  - ...the cloud
  - ...HPC systems from Cray and other vendors
  - ...modern processors like Intel Xeon Phi, GPUs\*, etc.

\* = academic work only; not yet supported in the official release



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# Chapel is Open-Source

- Chapel's development is hosted at GitHub
  - <https://github.com/chapel-lang>
- Chapel is licensed as Apache v2.0 software
- Instructions for download + install are online
  - see <http://chapel.cray.com/download.html> to get started



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# The Chapel Team at Cray (May 2016)



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel Community R&D Efforts



THE GEORGE  
WASHINGTON  
UNIVERSITY  
WASHINGTON, DC



Lawrence Berkeley  
National Laboratory



Yale


(and several others...)

<http://chapel.cray.com/collaborations.html>



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# Outline

- ✓ Chapel Motivation and Background
- Chapel in a Nutshell
- Chapel Project: Past, Present, Future



COMPUTE | STORE | ANALYZE


Copyright 2017 Cray Inc.

# Chapel's Multiresolution Philosophy


## ***Multiresolution Design:*** Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for greater degrees of control


*Chapel language concepts*



- build the higher-level concepts in terms of the lower
- permit the user to intermix layers arbitrarily



# Base Language



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# Base Language Features, by example

```
iter fib(n) {  
    var current = 0,  
        next = 1;  
  
    for i in 1..n {  
        yield current;  
        current += next;  
        current <=> next;  
    }  
}
```

```
config const n = 10;  
  
for f in fib(n) do  
    writeln(f);
```

```
0  
1  
1  
2  
3  
5  
8  
...
```





# Base Language Features, by example

## Modern iterators

```
iter fib(n) {  
    var current = 0,  
        next = 1;  
  
    for i in 1..n {  
        yield current;  
        current += next;  
        current <=> next;  
    }  
}
```

```
config const n = 10;  
  
for f in fib(n) do  
    writeln(f);
```

```
0  
1  
1  
2  
3  
5  
8  
...
```



COMPUTE

|

STORE

|

ANALYZE

# Base Language Features, by example

Configuration declarations  
(to avoid command-line argument parsing)  
./a.out --n=1000000

```
iter fib(n) {  
    var current = 0,  
        next = 1;  
  
    for i in 1..n {  
        yield current;  
        current += next;  
        current <=> next;  
    }  
}
```

```
config const n = 10;  
  
for f in fib(n) do  
    writeln(f);
```

```
0  
1  
1  
2  
3  
5  
8  
...
```



# Base Language Features, by example

Static type inference for:

- arguments
- return types
- variables

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <= next;
    }
}
```

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0
1
1
2
3
5
8
...
...
```



# Base Language Features, by example

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```

Zippered iteration



# Base Language Features, by example

Range types and operators

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```





# Base Language Features, by example

```
iter fib(n) {  
    var current = 0,  
        next = 1;  
  
    for i in 1..n {  
        yield current;  
        current += next;  
        current <=> next;  
    }  
}
```

tuples

```
config const n = 10;  
  
for (i,f) in zip(0..#n, fib(n)) do  
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0  
fib #1 is 1  
fib #2 is 1  
fib #3 is 2  
fib #4 is 3  
fib #5 is 5  
fib #6 is 8  
...
```





# Base Language Features, by example


```
iter fib(n) {  
    var current = 0,  
        next = 1;  
  
    for i in 1..n {  
        yield current;  
        current += next;  
        current <=> next;  
    }  
}
```

```
config const n = 10;  
  
for (i,f) in zip(0..#n, fib(n)) do  
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0  
fib #1 is 1  
fib #2 is 1  
fib #3 is 2  
fib #4 is 3  
fib #5 is 5  
fib #6 is 8  
  
...
```



# Task Parallelism



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Task Parallelism, Locality Control, by example



taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Task Parallelism, Locality Control, by example

Abstraction of  
System Resources

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writelf("Hello from task %n of %n "+
                "running on %s\n",
            tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



# Task Parallelism, Locality Control, by example

High-Level  
Task Parallelism

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writelf("Hello from task %n of %n "+
                "running on %s\n",
            tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```





# Task Parallelism, Locality Control, by example

Control of Locality/Affinity

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Task Parallelism, Locality Control, by example

Abstraction of  
System Resources

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writelf("Hello from task %n of %n "+
                "running on %s\n",
            tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



# Task Parallelism, Locality Control, by example

High-Level  
Task Parallelism

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writelf("Hello from task %n of %n "+
                "running on %s\n",
            tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```





# Task Parallelism, Locality Control, by example

Not seen here:

Data-centric task coordination  
via atomic and full/empty vars

taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Task Parallelism, Locality Control, by example



taskParallel.chpl

```
coforall loc in Locales do
    on loc {
        const numTasks = here.maxTaskPar;
        coforall tid in 1..numTasks do
            writef("Hello from task %n of %n "+
                "running on %s\n",
                tid, numTasks, here.name);
    }
```

```
prompt> chpl taskParallel.chpl -o taskParallel
prompt> ./taskParallel --numLocales=2
Hello from task 1 of 2 running on n1033
Hello from task 2 of 2 running on n1032
Hello from task 2 of 2 running on n1033
Hello from task 1 of 2 running on n1032
```



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Parallelism and Locality: Orthogonal in Chapel

- This is a **parallel**, but local program:

```
coforall i in 1..msgs do  
    writeln("Hello from task ", i);
```

- This is a **distributed**, but serial program:


```
writeln("Hello from locale 0!");  
on Locales[1] do writeln("Hello from locale 1!");  
on Locales[2] do writeln("Hello from locale 2!");
```

- This is a **distributed parallel** program:

```
coforall i in 1..msgs do  
    on Locales[i%numLocales] do  
        writeln("Hello from task ", i,  
               " running on locale ", here.id);
```

# Higher-Level Features

*Chapel language concepts*





# Data Parallelism, by example

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE | STORE | ANALYZE

# Data Parallelism, by example

Domains (Index Sets)

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

# Data Parallelism, by example

Arrays

dataParallel.chpl

```
config const n = 1000;  
  
var D = {1..n, 1..n};  
  
var A: [D] real;  
  
forall (i,j) in D do  
    A[i,j] = i + (j - 0.5)/n;  
  
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel  
prompt> ./dataParallel --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE

# Data Parallelism, by example

## Data-Parallel Forall Loops

dataParallel.chpl

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE | STORE | ANALYZE

# Distributed Data Parallelism, by example

Domain Maps  
(Map Data Parallelism to the System)

dataParallel.chpl

```
use CyclicDist;
config const n = 1000;
var D = {1..n, 1..n}
dmapped Cyclic(startIdx = (1,1));
var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5 --numLocales=4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE

STORE

ANALYZE



# Distributed Data Parallelism, by example

dataParallel.chpl


```
use CyclicDist;
config const n = 1000;
var D = {1..n, 1..n}
        dmapped Cyclic(startIdx = (1,1));
var A: [D] real;
forall (i,j) in D do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl -o dataParallel
prompt> ./dataParallel --n=5 --numLocales=4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# Outline

- ✓ Chapel Motivation and Background
- ✓ Chapel in a Nutshell
- Chapel Project: Past, Present, Future



COMPUTE | STORE | ANALYZE


Copyright 2017 Cray Inc.



# Chapel's Origins: HPCS

## DARPA HPCS: High Productivity Computing Systems


- **Goal:** improve productivity by a factor of 10x
- **Timeframe:** Summer 2002 – Fall 2012
- Cray developed a new system architecture, network, software stack...
  - this became the very successful Cray XC30™ Supercomputer Series



...and a new programming language: Chapel  
(at that point, essentially a research prototype)



COMPUTE | STORE | ANALYZE



# Chapel's 5-year push

- Based on positive user response to Chapel under HPCS, Cray undertook a five-year effort to improve it
  - we're just completing our fourth year
- Focus Areas:
  1. Improving **performance** and scaling
  2. **Fixing** immature aspects of the language and implementation
    - e.g., strings, memory management and leaks, OOP, error handling, ...
  3. **Porting** to emerging architectures
    - Intel Xeon Phi, accelerators, heterogeneous processors and memories, ...
  4. Improving **interoperability**
  5. Growing the Chapel user and developer **community**
    - including non-scientific computing communities
  6. Exploring transition of Chapel **governance** to a neutral, external body



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# A Year in the Life of Chapel

- **Two major releases per year** (April / October)
  - ~a month later: detailed [release notes](#)
  - latest release: Chapel 1.15, released April 6<sup>th</sup> 2017
    - release notes due to be published this week or next
- **CHIUW: Chapel Implementers and Users Workshop** (~June)
  - (4<sup>th</sup> annual) [CHIUW 2017](#), June 1-2 at IPDPS (Orlando, FL)
  - talks from members of the broad community + a Chapel code camp
- **SC** (Nov)
  - tutorials, panels, BoFs, posters, educator sessions, exhibits, ...
  - annual **CHUG (Chapel Users Group)** happy hour
- **Talks, tutorials, research visits, blog posts, ...** (year-round)



COMPUTE

|


STORE

|

ANALYZE

# Chapel is a Work-in-Progress

- Currently being picked up by early adopters
  - ~3000+ downloads per year across two releases



- Users who try it generally like what they see



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# A notable early adopter

## Chapel in the (Cosmological) Wild

1:00 – 2:00

**Nikhil Padmanabhan**, Yale University Professor, Physics & Astronomy

**Abstract:** This talk aims to present my personal experiences using Chapel in my research. My research interests are in observational cosmology; more specifically, I use large surveys of galaxies to constrain the evolution of the

The image shows a YouTube video player interface. At the top left is the YouTube logo. To its right is a search bar with the placeholder "Search". Below the search bar is a navigation bar with icons for "Videos", "Playlists", and "Channels". The main content area features a thumbnail image of a man (Nikhil Padmanabhan) giving a presentation in front of a projection screen. The video title is "CHIUW 2016 keynote: 'Chapel in the (Cosmological) Wild', Nikhil Padmanabhan". Below the title, it says "Chapel Parallel Programming Language", "1 month ago • 86 views", and a partial description: "This is Nikhil Padmanabhan's keynote talk from CHIUW 2016: the 3rd Annual Chapel Implementers and Users workshop. The slides are availabl...". A timestamp "56:14" is visible in the bottom right corner of the thumbnail.



COMPUTE

STORE

ANALYZE

Copyright 2017 Cray Inc.



# Chapel: Top 3 Historical Barriers to Use

## 3. Core Language Feature Improvements

- Historical problems that are now much better:
  - strings, memory leaks, memory management
- Areas that have improved, but are still in-progress:
  - initializers (constructor replacement), error-handling




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Memory Leak Improvements

- Effort in recent years has dramatically reduced leaks
  - most remaining cases are due to user-level leaks in tests themselves





# Chapel: Top 3 Traditional Barriers to Use

## 3. Core Language Feature Improvements

- Historical problems that are now much better:
  - strings, memory leaks, memory management, interoperability, generics
- Areas that have improved, but are still in-progress:
  - initializers (constructor replacement), error-handling

## 2. Access to Standard Libraries

- Situation has improved significantly over past few years
  - Several core libraries added:
    - BigInteger, BitOps, DateTime, FileSystem, Random, Reflection, Spawn, ...
  - As well as access to many standard libraries / technologies:
    - BLAS, Curl, FFTW, Futures, HDFS, LAPACK, LinearAlgebra, MPI, ZMQ, ...

## 1. Performance



- Particularly for old-school HPC users, performance is crucial
- That said, as of this month's release, we're reaching parity more often



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Single-Locale Improvements in Execution Time



- Single-locale is increasingly on par with C / C++ / OpenMP



COMPUTE

STORE

ANALYZE

# Computer Language Benchmarks Game (CLBG)



## The Computer Language Benchmarks Game

64-bit quad core data set

Will your toy benchmark program be faster if you write it in a different programming language? It depends how you write it!

### Which programs are fast?

Which are succinct? Which are efficient?

<u>Ada</u>	<u>C</u>	<u>Chapel</u>	<u>Clojure</u>	<u>C#</u>	<u>C++</u>
<u>Dart</u>	<u>Erlang</u>	<u>F#</u>	<u>Fortran</u>	<u>Go</u>	<u>Hack</u>
<u>Haskell</u>	<u>Java</u>	<u>JavaScript</u>	<u>Lisp</u>	<u>Lua</u>	
<u>OCaml</u>	<u>Pascal</u>	<u>Perl</u>	<u>PHP</u>	<u>Python</u>	
<u>Racket</u>	<u>Ruby</u>	<u>JRuby</u>	<u>Rust</u>	<u>Scala</u>	
<u>Smalltalk</u>	<u>Swift</u>	<u>TypeScript</u>			

## Website that supports cross-language game / comparisons

- 13 toy benchmark programs
- exercises key features like:
  - memory management
  - tasking and synchronization
  - vectorization
  - big integers
  - strings and regular expressions
- specific approach prescribed

## Take results w/ grain of salt

- other programs may be different
  - not to mention other programmers
  - specific to this platform / OS / ...

That said, it's one of the only games in town...



COMPUTE

STORE

ANALYZE

# Computer Language Benchmarks Game (CLBG)



## The Computer Language Benchmarks Game

64-bit quad core data set

Will your toy benchmark program be faster if you write it in a different programming language? It depends how you write it!

**Which programs are fast?**

Which are succinct? Which are efficient?

<u>Ada</u>	<u>C</u>	<u>Chapel</u>	<u>Clojure</u>	<u>C#</u>	<u>C++</u>
<u>Dart</u>	<u>Erlang</u>	<u>F#</u>	<u>Fortran</u>	<u>Go</u>	<u>Hack</u>
<u>Haskell</u>	<u>Java</u>	<u>JavaScript</u>	<u>Lisp</u>	<u>Lua</u>	
<u>OCaml</u>	<u>Pascal</u>	<u>Perl</u>	<u>PHP</u>	<u>Python</u>	
<u>Racket</u>	<u>Ruby</u>	<u>JRuby</u>	<u>Rust</u>	<u>Scala</u>	
	<u>Smalltalk</u>	<u>Swift</u>	<u>TypeScript</u>		

## Chapel's approach to CLBG:


- want to know how we compare
- strive for entries that are elegant rather than heroic
  - e.g., “Want to learn how program x works? Check out the Chapel version.”



COMPUTE

STORE

ANALYZE



# CLBG: Website

Can sort results by execution time, code size, memory or CPU use:

The Computer Language Benchmarks Game						
chameneos-redux						
<u>description</u>						
program source code, command-line and measurements						
x	source	secs	mem	gz	cpu	cpu load
1.0	<a href="#">C gcc #5</a>	<b>0.60</b>	820	2863	2.37	100% 100% 98% 100%
1.2	<a href="#">C++ g++ #5</a>	<b>0.70</b>	3,356	1994	2.65	100% 100% 91% 92%
1.7	<a href="#">Lisp SBCL #3</a>	<b>1.01</b>	55,604	2907	3.93	97% 96% 99% 99%
2.3	<a href="#">Chapel #2</a>	<b>1.39</b>	76,564	1210	5.43	99% 99% 98% 99%
3.3	<a href="#">Rust #2</a>	<b>2.01</b>	56,936	2882	7.81	97% 98% 98% 98%
5.6	<a href="#">C++ g++ #2</a>	3.40	1,880	2016	11.88	100% 51% 100% 100%
6.8	<a href="#">Chapel</a>	4.09	66,584	1199	16.25	100% 100% 100% 100%
8.0	<a href="#">Java #4</a>	<b>4.82</b>	37,132	1607	16.73	98% 98% 54% 99%
8.5	<a href="#">Haskell GHC</a>	<b>5.15</b>	8,596	989	9.26	79% 100% 2% 2%
10	<a href="#">Java</a>	6.13	53,760	1770	8.78	42% 45% 41% 16%
10	<a href="#">Haskell GHC #4</a>	6.34	6,908	989	12.67	99% 100% 2% 1%
11	<a href="#">C# .NET Core</a>	<b>6.59</b>	86,076	1400	22.96	99% 82% 78% 91%
11	<a href="#">Go</a>	<b>6.90</b>	832	1167	24.19	100% 96% 56% 100%
13	<a href="#">Go #2</a>	7.59	1,384	1408	27.65	91% 99% 99% 78%
13	<a href="#">Java #3</a>	7.94	53,232	1267	26.86	54% 96% 98% 94%

The Computer Language Benchmarks Game						
chameneos-redux						
<u>description</u>						
program source code, command-line and measurements						
x	source	secs	mem	gz	cpu	cpu load
1.0	<a href="#">Erlang</a>	58.90	28,668	<b>734</b>	131.19	62% 60% 51% 53%
1.0	<a href="#">Erlang HiPE</a>	59.39	25,784	<b>734</b>	131.58	60% 56% 56% 54%
1.1	<a href="#">Perl #4</a>	5 min	14,084	<b>785</b>	7 min	40% 40% 29% 28%
1.1	<a href="#">Racket</a>	5 min	132,120	<b>791</b>	5 min	1% 0% 0% 100%
1.1	<a href="#">Racket #2</a>	175.88	116,488	842	175.78	100% 1% 1% 0%
1.2	<a href="#">Python 3 #2</a>	236.84	7,908	<b>866</b>	5 min	24% 48% 27% 45%
1.3	<a href="#">Ruby</a>	90.52	9,396	<b>920</b>	137.53	35% 35% 35% 34%
1.3	<a href="#">Ruby JRuby</a>	48.78	628,968	<b>928</b>	112.15	65% 60% 49% 58%
1.3	<a href="#">Go #5</a>	11.05	832	<b>957</b>	32.48	75% 74% 75% 73%
1.3	<a href="#">Haskell GHC #4</a>	6.34	6,908	<b>989</b>	12.67	99% 100% 2% 1%
1.3	<a href="#">Haskell GHC</a>	5.15	8,596	989	9.26	79% 100% 2% 2%
1.6	<a href="#">OCaml #3</a>					32% 38% 37% 39%
1.6	<a href="#">Go</a>					100% 96% 56% 100%
1.6	<a href="#">Chapel</a>					0% 100% 100% 100%
1.6	<a href="#">Chapel #2</a>					99% 99% 98% 99%

gz == code size metric  
strip comments and extra whitespace, then gzip



COMPUTE

STORE

ANALYZE

Copyright 2016 Cray Inc.



# CLBG: Chapel Standings as of Apr 20<sup>th</sup>

- 8 / 13 programs in top-20 smallest:
  - two #1 smallest:  
**n-body**  
**thread-ring**
  - 2 others in the top-5 smallest:  
**pidigits**  
**spectral-norm**
  - 1 other in the top-10 smallest:  
**regex-redux**
  - 3 others in the top-20 smallest:  
**chameneos-redux**  
**mandelbrot**  
**meteor-contest**
- 12 / 13 programs in top-20 fastest:
  - one #1 fastest:  
**pidigits**
  - 3 others in the top-5 fastest:  
**chameneos-redux**  
**meteor-contest**  
**thread-ring**
  - 3 others in the top-10 fastest:  
**fannkuch-redux**  
**fasta**  
**mandelbrot**
  - 5 others in the top-20 fastest:  
**binary-trees**  
**k-nucleotide**  
**n-body**  
**regex-redux**  
**spectral-norm**





# CLBG: Website

Can also compare languages pair-wise:

The Computer Language Benchmarks Game						
Chapel programs versus Go <u>all other Chapel programs &amp; measurements</u>						
by benchmark task performance						
<u>regex-redux</u>						
source	secs	mem	gz	cpu	cpu load	
<u>Chapel</u>	<b>10.02</b>	1,022,052	477	19.68	99%	72% 14% 12%
<u>Go</u>	29.51	352,804	798	61.51	77%	49% 43% 40%
<u>binary-trees</u>						
source	secs	mem	gz	cpu	cpu load	
<u>Chapel</u>	<b>14.32</b>	324,660	484	44.15	100%	58% 78% 75%
<u>Go</u>	34.77	269,068	654	132.04	95%	97% 95% 95%
<u>fannkuch-redux</u>						
source	secs	mem	gz	cpu	cpu load	
<u>Chapel</u>	<b>11.38</b>	46,056	728	45.18	100%	99% 99% 100%
<u>Go</u>	15.81	1,372	900	62.92	100%	100% 99% 99%



COMPUTE

STORE

ANALYZE

# CLBG: Website



Can also browse program source code (but this requires actual thought):

```
proc main() {
    printColorEquations();

    const group1 = [i in 1..popSize1] new Chameneos(i, ((i-1)%3):Color);
    const group2 = [i in 1..popSize2] new Chameneos(i, colors10[i]);

    cobegin {
        holdMeetings(group1, n);
        holdMeetings(group2, n);
    }

    print(group1);
    print(group2);

    for c in group1 do delete c;
    for c in group2 do delete c;
}

// Print the results of getNewColor() for all color pairs.
// proc printColorEquations() {
//     for c1 in Color do
//         for c2 in Color do
//             writeln(c1, " + ", c2, " -> ", getNewColor(c1, c2));
//     writeln();
// }

// Hold meetings among the population by creating a shared meeting
// place, and then creating per-chameneos tasks to have meetings.
// proc holdMeetings(population, numMeetings) {
//     const place = new MeetingPlace(numMeetings);

//     coforall c in population do          // create a task per chameneos
//         c.haveMeetings(place, population);

//     delete place;
// }
```

*excerpt from 1210.gz Chapel #2 entry*

```
void get_affinity(int* is_smp, cpu_set_t* affinity1, cpu_set_t* affinity2)
{
    cpu_set_t
    FILE*
    char
    char const*
    int
    int
    int
    int
    int
    size_t
    size_t

    active_cpus;
    f;
    buf [2048];
    pos;
    cpu_idx;
    physical_id;
    core_id;
    cpu_cores;
    apic_id;
    cpu_count;
    i;

    char const*
    size_t
    char const*
    size_t
    char const*
    size_t
    char const*
    size_t

    processor_str      = "processor";
    processor_str_len = strlen(processor_str);
    physical_id_str   = "physical id";
    physical_id_str_len = strlen(physical_id_str);
    core_id_str        = "core id";
    core_id_str_len   = strlen(core_id_str);
    cpu_cores_str      = "cpu cores";
    cpu_cores_str_len = strlen(cpu_cores_str);

    CPU_ZERO(&active_cpus);
    sched_getaffinity(0, sizeof(active_cpus), &active_cpus);
    cpu_count = 0;
    for (i = 0; i != CPU_SETSIZE; i += 1)
    {
        if (CPU_ISSET(i, &active_cpus))
        {
            cpu_count += 1;
        }
    }

    if (cpu_count == 1)
    {
        is_smp[0] = 0;
        return;
    }

    is_smp[0] = 1;
    CPU_ZERO(affinity1);
```

*excerpt from 2863.gz C gcc #5 entry*



COMPUTE

STORE

ANALYZE

# CLBG: Website



Can also browse program source code (but this requires actual thought):

```
proc main() {
    printColorEquations();

    const group1 = [i in 1..popSize1] new Chameneos(i, c);
    const group2 = [i in 1..popSize2] new Chameneos(i, c);

    cobegin {
        holdMeetings(group1, n);
        holdMeetings(group2, n);
    }

    print(group1);
    print(group2);

    for c in group1 do delete c;
    for c in group2 do delete c;
}

// Print the results of getNewColor() for all color pairs.
// proc printColorEquations() {
//     for c1 in Color do
//         for c2 in Color do
//             writeln(c1, " + ", c2, " -> ", getNewColor(c1, c2));
//             writeln();
// }

// Hold meetings among the population by creating a shared
// place, and then creating per-chameneos tasks to have
// them meet.
// proc holdMeetings(population, numMeetings) {
//     const place = new MeetingPlace(numMeetings);

//     coforall c in population do          // create a task
//         c.haveMeetings(place, population);

//     delete place;
// }
```

```
void get affinity(int* is_smp, cpu_set_t* affinity1, cpu_set_t* affinity2)

cobegin {
    holdMeetings(group1, n);
    holdMeetings(group2, n);
}

char const* processor_str      = "processor";
size_t processor_str_len     = strlen(processor_str);
char const* physical_id_str   = "physical id";
size_t physical_id_str_len   = strlen(physical_id_str);
char const* core_id_str       = "core id";
size_t core_id_str_len        = strlen(core_id_str);
char const* cpu_cores_str     = "cpu cores";
size_t cpu_cores_str_len      = strlen(cpu_cores_str);

proc holdMeetings(population, numMeetings) {
    const place = new MeetingPlace(numMeetings);

    coforall c in population do          // create a task
        c.haveMeetings(place, population);

    delete place;
}

is_smp[0] = 1;
CPU_ZERO(affinity1);
```

excerpt from 1210.gz Chapel #2 entry

excerpt from 2863.gz C gcc #5 entry




COMPUTE

STORE

ANALYZE

## Site summary: relative performance (sorted by geometric mean)



# CLBG: Website

- site has a sound philosophy about too-easy answers

**We want easy answers**, but easy answers are often incomplete or wrong. You and I know, there's more we should understand:

stories

details


fast?

conclusions

- yet, most readers probably still jump to conclusions
  - execution time dominates default (or only) views of results
  - it's simply human nature
- we're interested in elegance as well as performance
  - elegance is obviously in the eye of the beholder
    - we compare source codes manually
    - but then use CLBG's code size metric as a quantitative stand-in
  - want to be able to compare both axes simultaneously
  - to that end, we used scatter plots to compare implementations



# Chapel entries



COMPUTE


|

STORE

|

ANALYZE

# Chapel vs. 9 other languages



COMPUTE

STORE

ANALYZE

# Chapel vs. 9 other languages (zoomed out)







COMPUTE

STORE

ANALYZE

# Multi-locale Improvements in Execution Time

- Multi-locale performance is improving significantly as well



# 3 Key Multi-Locale Communication Benchmarks



## STREAM Triad:

- measures embarrassingly / pleasingly parallel computation

## RA:


- measures random updates to a large distributed array

## ISx:


- measures bucket-exchange idiom



COMPUTE | STORE | ANALYZE



# STREAM Triad: Chapel vs. MPI Scalability




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# RA: Chapel vs. MPI Scalability



Performance of RA (atomics)




COMPUTE

STORE

ANALYZE

# ISx: Performance Summary


- Gathered on Cray XC with default problem size
  - reference versions



# ISx: Performance Summary

- Gathered on Cray XC with default problem size
  - adding Chapel, six months ago:


ISx weakISO Total Time



# ISx: Performance Summary



- Gathered on Cray XC with default problem size
  - adding Chapel, today:



COMPUTE | STORE | ANALYZE


Copyright 2017 Cray Inc.

# ISx: Performance Summary



- Gathered on Cray XC with default problem size
  - dropping the old Chapel timings, and zooming in:

ISx weakISO Total Time



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.




# Overview Summary

- Chapel has nice features for parallelism and locality
- Traditional reasons for not using Chapel are falling away
  - performance specifically is becoming less of a concern with time
- Aiming for a “version 2.0 release” over the near year or so
  - intent: no further breaking changes after that point



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.



# High-level Questions about Chapel?




COMPUTE

|

STORE

|

ANALYZE



# Full-size CLBG Scatter Plots



COMPUTE

|


STORE

|

ANALYZE

Copyright 2017 Cray Inc.


# Chapel vs. C




COMPUTE | STORE | ANALYZE


Copyright 2017 Cray Inc.

# Chapel vs. C (zoomed out)






# Chapel vs. C++



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. C++ (zoomed out)




COMPUTE


STORE

ANALYZE

# Chapel vs. Fortran




# Chapel vs. Fortran (zoomed out)



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.


# Chapel vs. Go



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.


# Chapel vs. Go (zoomed out)



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.


# Chapel vs. Rust




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. Rust (zoomed out)



# Chapel vs. Swift




COMPUTE


STORE

ANALYZE

# Chapel vs. Swift (zoomed out)



# Chapel vs. Java



COMPUTE


|

STORE

|

ANALYZE


# Chapel vs. Java (zoomed out)




COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. Scala




# Chapel vs. Scala (zoomed out)






# Chapel vs. Python



COMPUTE | STORE | ANALYZE

Copyright 2017 Cray Inc.

# Chapel vs. Python (zoomed out)





# Legal Disclaimer

*Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.*

*Cray Inc. may make changes to specifications and product descriptions at any time, without notice.*

*All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.*

*Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.*

*Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.*

*Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.*

*The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, and URIKA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.*

