



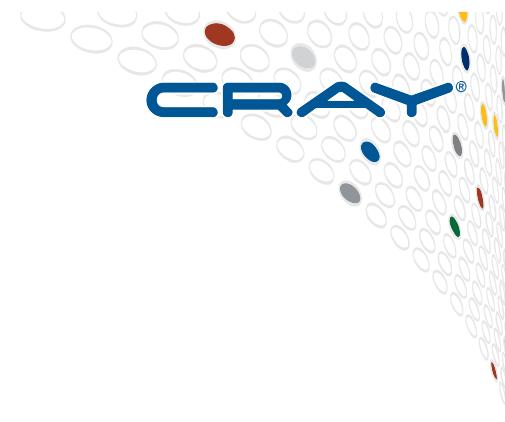
Locality / Affinity Features



COMPUTE

| STORE

| ANALYZE



What is a Locale?

Definition:

- Abstract unit of target architecture
- Supports reasoning about locality
 - defines “here vs. there” / “local vs. remote” / “cheap vs. \$\$\$”
- Capable of running tasks and storing variables
 - i.e., has processors and memory

Typically: A compute node (multicore processor or SMP)



COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.



Getting started with locales

- Users specify # of locales when running Chapel programs

```
% a.out --numLocales=8
```

```
% a.out -nl 8
```

- Chapel provides built-in locale variables

```
config const numLocales: int = ...;  
const Locales: [0..#numLocales] locale = ...;
```

Locales

L0	L1	L2	L3	L4	L5	L6	L7
----	----	----	----	----	----	----	----

- User's main () begins executing on locale #0



Locale Operations

- Locale methods support queries about the target system:

```
proc locale.physicalMemory(...) { ... }  
proc locale.numPUs() { ... }  
proc locale.id { ... }  
proc locale.name { ... }
```

- On-clauses support placement of computations:

```
writeln("on locale 0");  
  
on Locales[1] do  
    writeln("now on locale 1");  
  
writeln("on locale 0 again");
```

```
on A[i,j] do  
    bigComputation(A);  
  
on node.left do  
    search(node.left);
```

Parallelism and Locality: Orthogonal in Chapel



- This is a **parallel**, but local program:

```
begin writeln("Hello world!");  
writeln("Goodbye!");
```

- This is a **distributed**, but serial program:

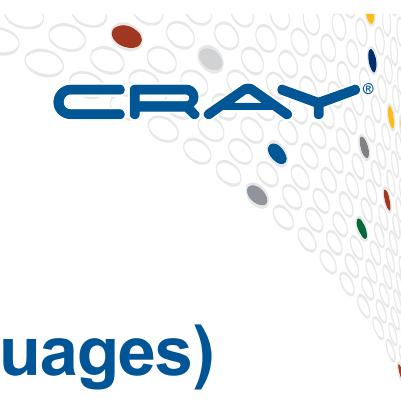
```
writeln("Hello from locale 0!");  
on Locales[1] do writeln("Hello from locale 1!");  
writeln("Goodbye from locale 0!");
```

- This is a **distributed and parallel** program:

```
begin on Locales[1] do writeln("Hello from locale 1!");  
on Locales[2] do begin writeln("Hello from locale 2!");  
writeln("Goodbye from locale 0!");
```

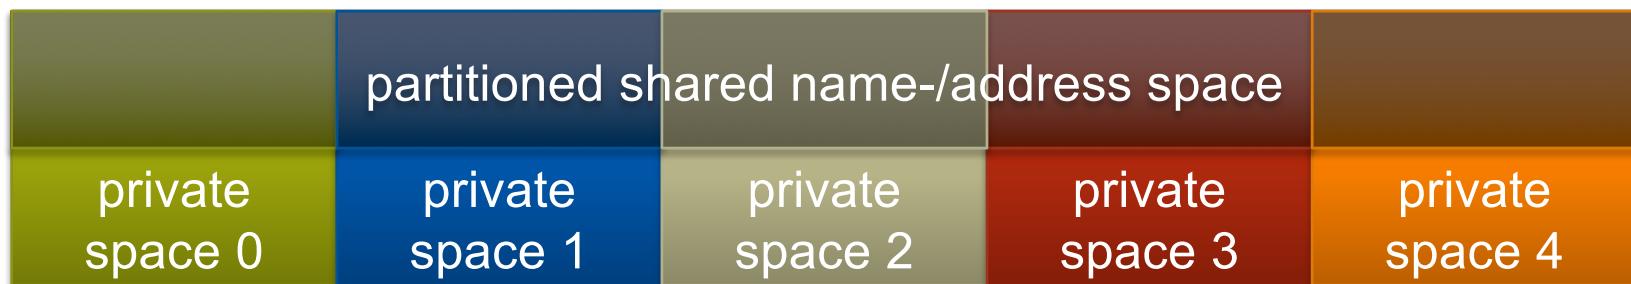


Partitioned Global Address Space (PGAS) Languages



(Or perhaps: partitioned global namespace languages)

- **abstract concept:**
 - support a shared namespace on distributed memory
 - permit parallel tasks to access remote variables by naming them
 - establish a strong sense of ownership
 - every variable has a well-defined location
 - local variables are cheaper to access than remote ones
- **traditional PGAS languages have been SPMD in nature**
 - best-known examples: Fortran Co-Arrays, UPC



COMPUTE

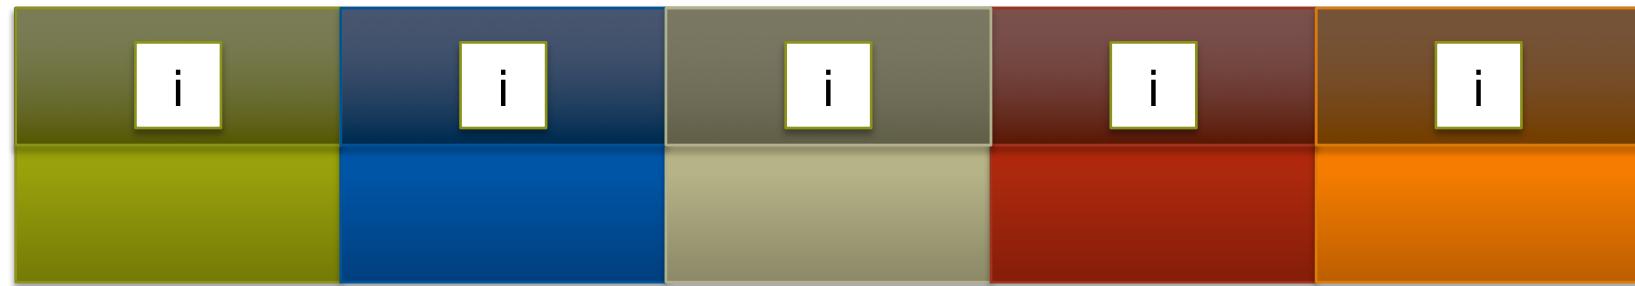
STORE

ANALYZE



SPMD PGAS Languages (using a pseudo-language, not Chapel)

```
shared int i(*) ;           // declare a shared variable i
```



COMPUTE | STORE | ANALYZE

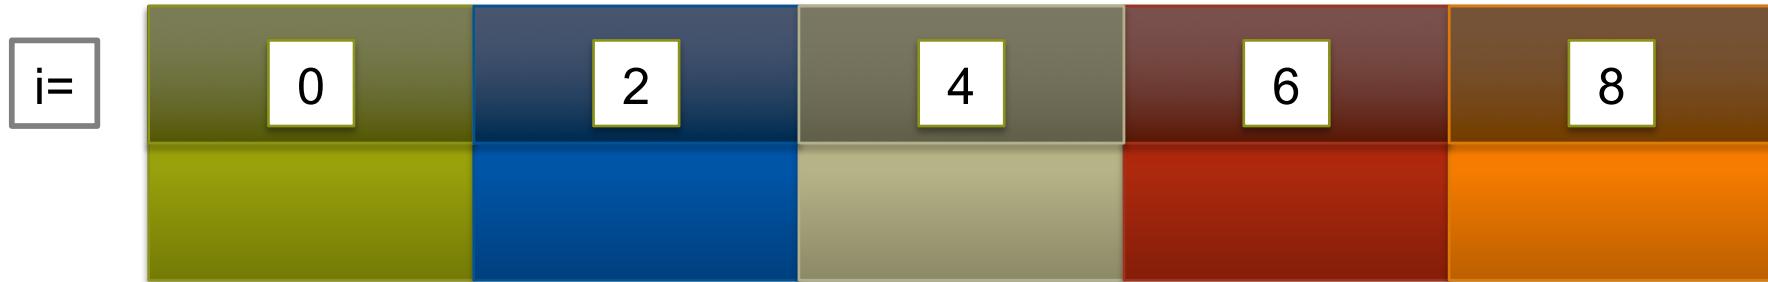
Copyright 2018 Cray Inc.

SPMD PGAS Languages

(using a pseudo-language, not Chapel)

```
shared int i(*) ;           // declare a shared variable i

function main() {
    i = 2*this_image();    // each image initializes its copy
```

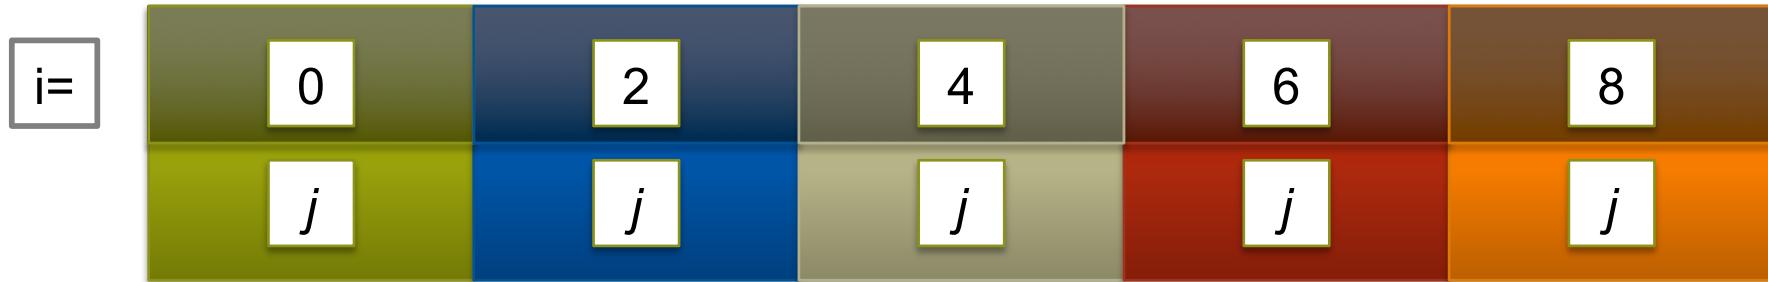




SPMD PGAS Languages

(using a pseudo-language, not Chapel)

```
shared int i(*) ;           // declare a shared variable i  
function main() {  
    i = 2*this_image();     // each image initializes its copy  
  
private int j;              // declare a private variable j
```



COMPUTE

STORE

ANALYZE

Copyright 2018 Cray Inc.

SPMD PGAS Languages

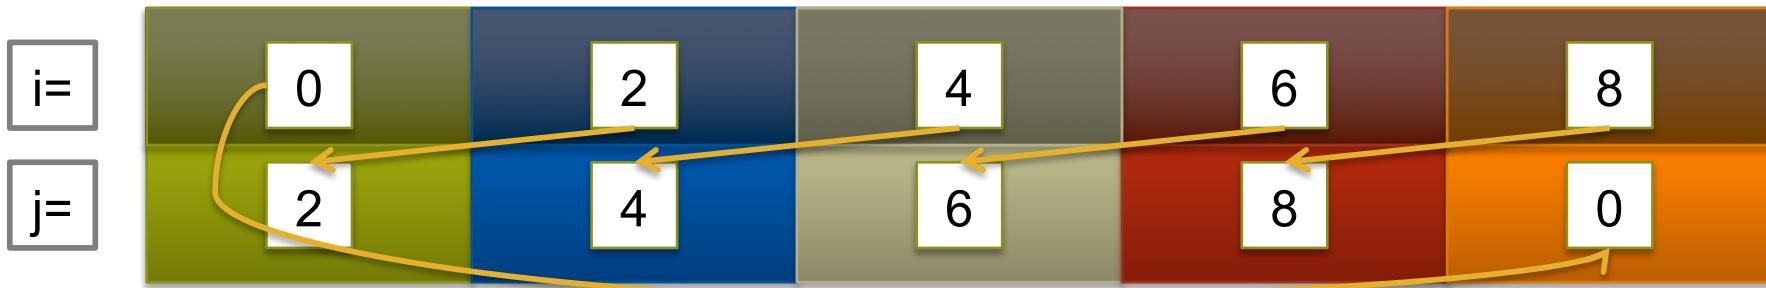
(using a pseudo-language, not Chapel)

```

shared int i(*) ;           // declare a shared variable i

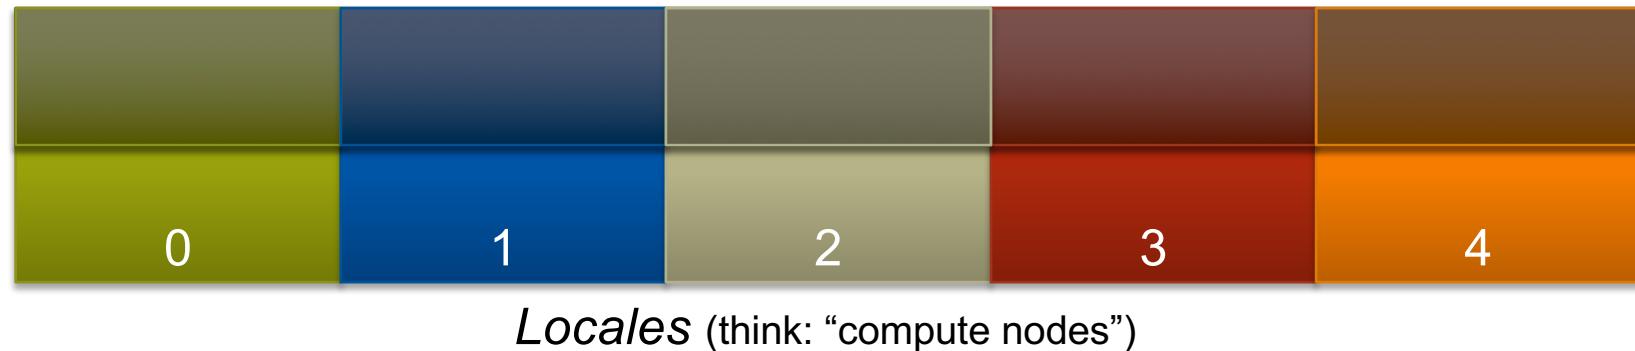
function main() {
    i = 2*this_image();    // each image initializes its copy
    barrier();
    private int j;          // declare a private variable j
    j = i( (this_image() + 1) % num_images() );
    // ^ access our neighbor's copy of i
    // communication is implemented by the compiler + runtime
    // Q: How did we know our neighbor had an i? A: Because it's SPMD – we're
    // all running the same program. (Simple, but restrictive)

```



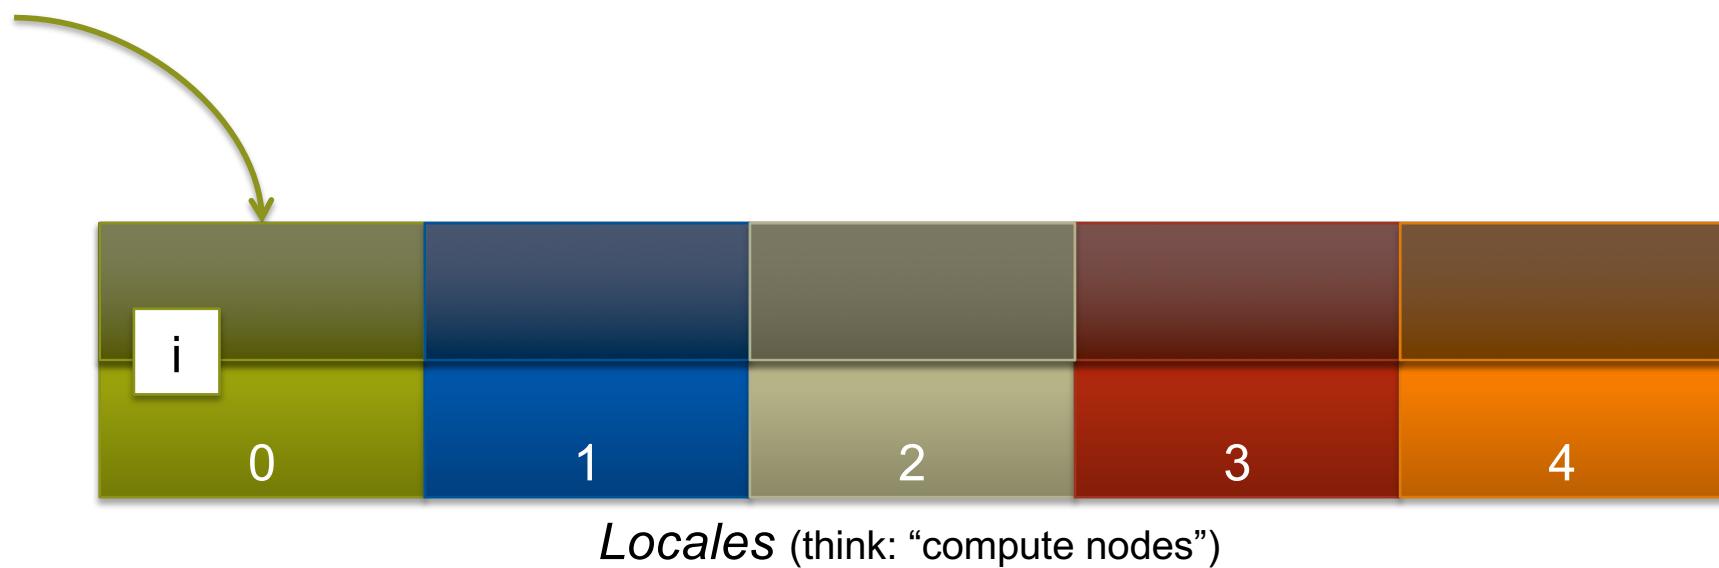
Chapel and PGAS

- Chapel is PGAS, but unlike most, it's not inherently SPMD
 - never think about “the other copies of the program”
 - “global name/address space” comes from lexical scoping
 - as in traditional languages, each declaration yields one variable
 - variables are stored on the locale where the task declaring it is executing



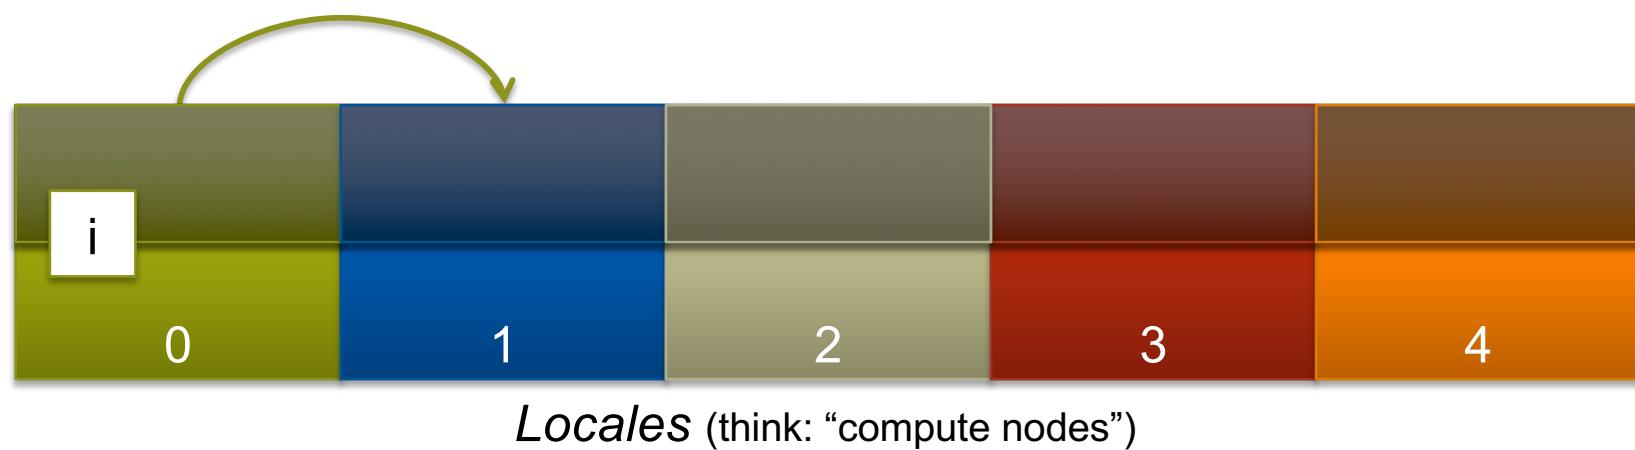
Chapel: Scoping and Locality

```
var i: int;
```



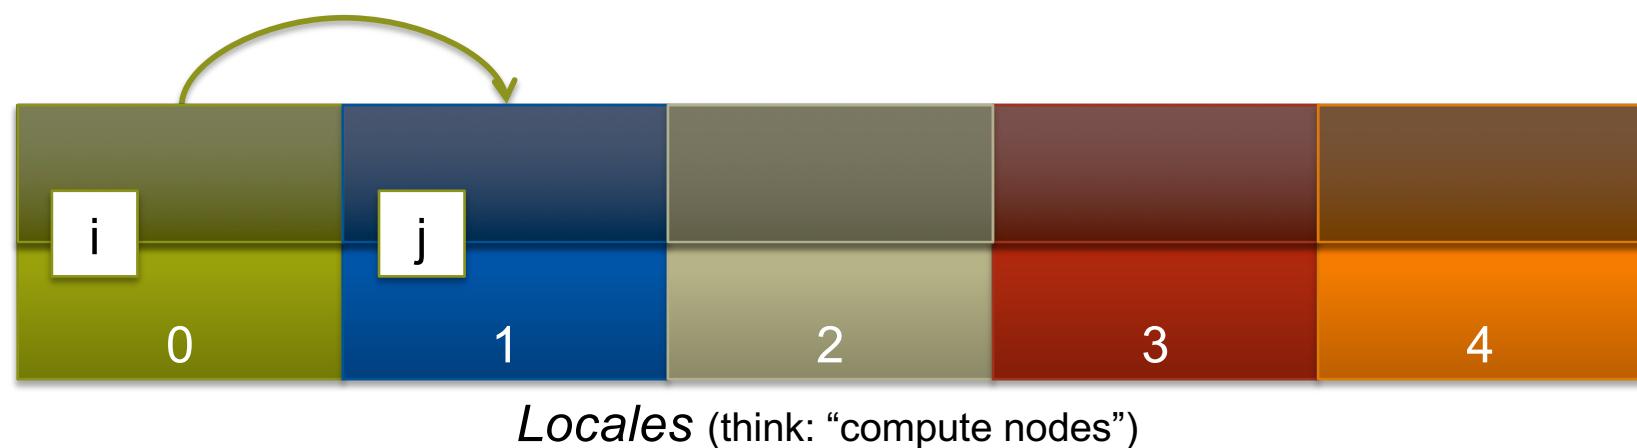
Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {
```



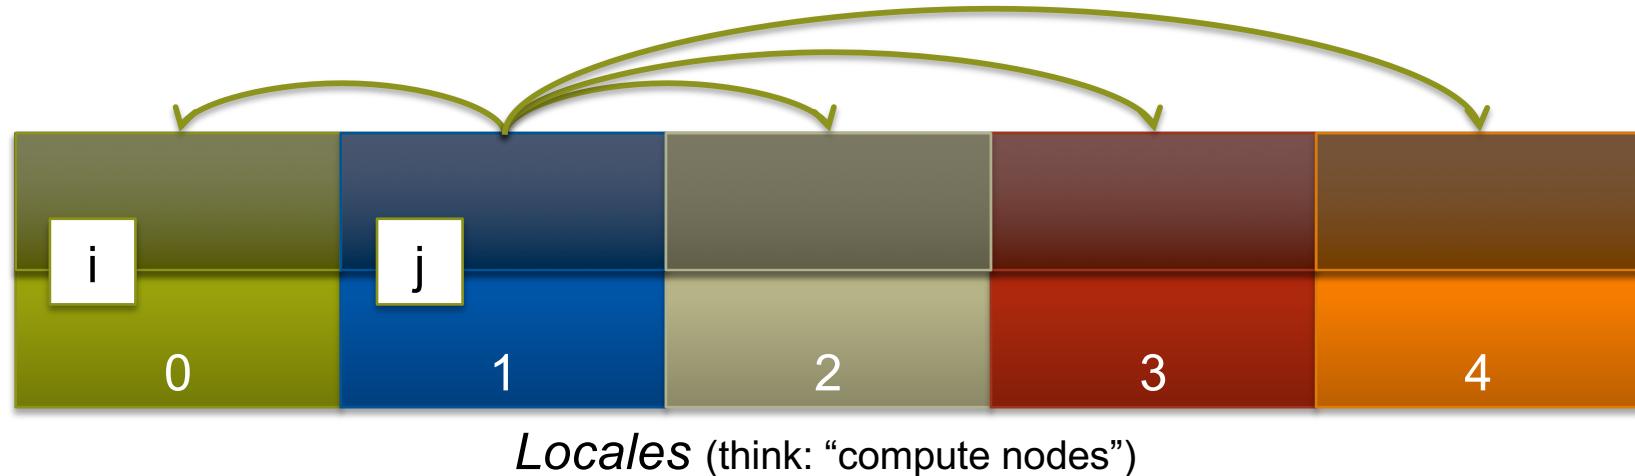
Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;
```



Chapel: Scoping and Locality

```
var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales {
        on loc {
```

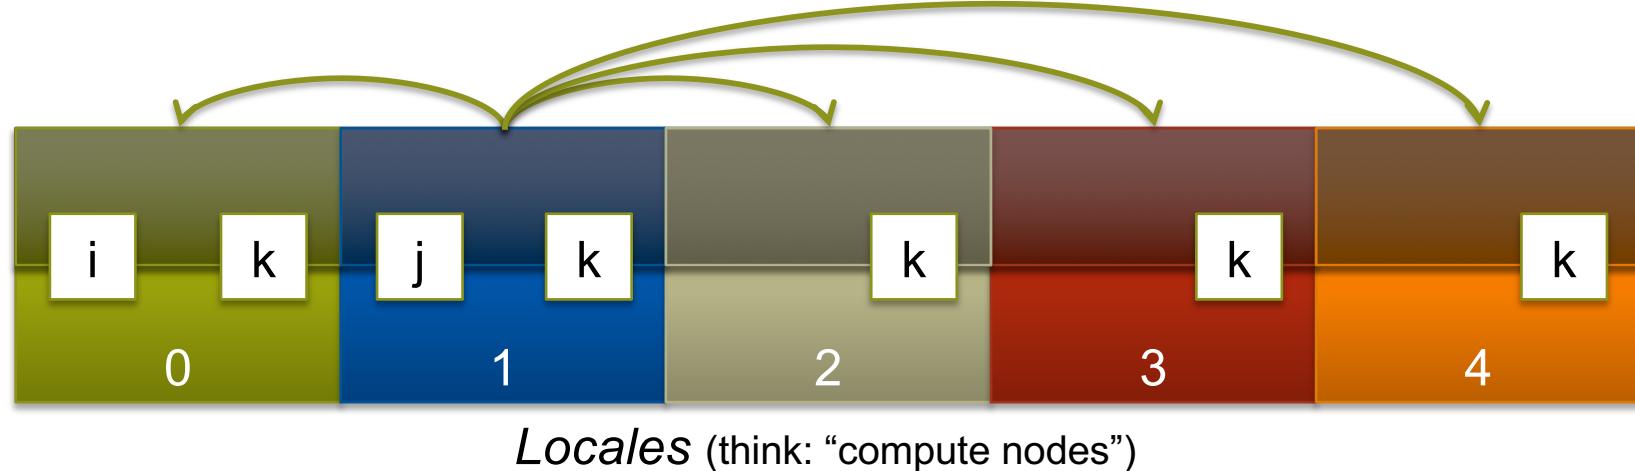


Chapel: Scoping and Locality

```

var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales {
        on loc {
            var k: int;
            ...
        }
    }
}

```



Chapel: Scoping and Locality

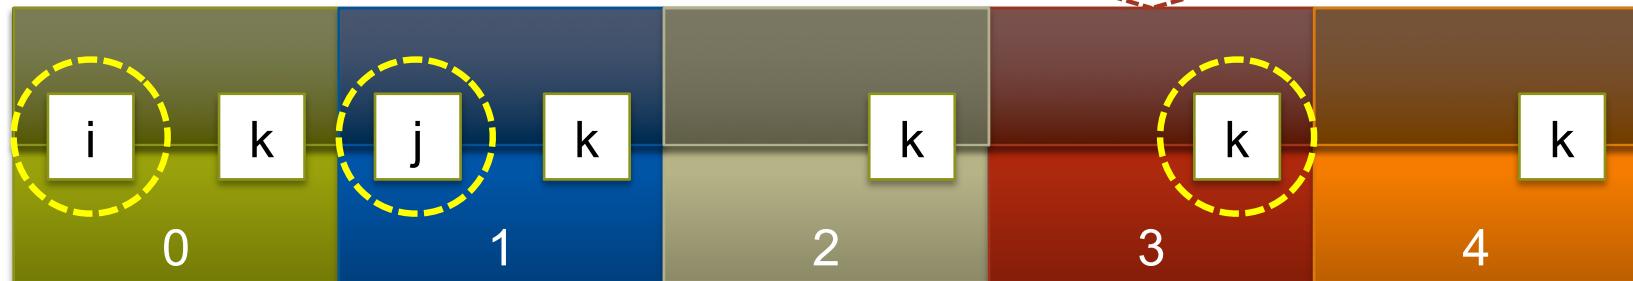
```

var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales {
        on loc {
            var k: int;
            k = 2*i + j;
        }
    }
}

```

OK to access i , j , and k
wherever they live

$k = 2*i + j;$



Locales (think: “compute nodes”)



Chapel: Scoping and Locality

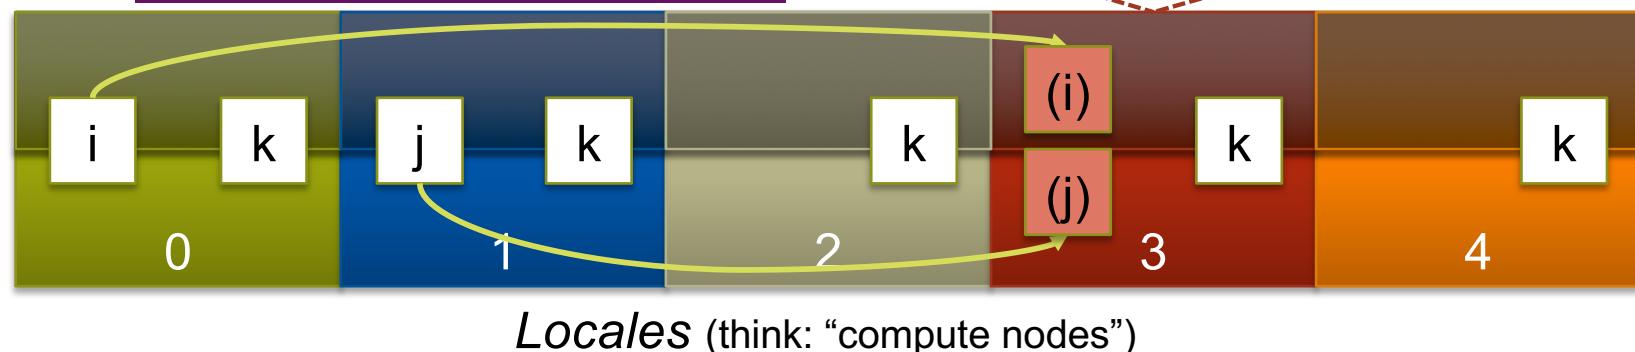
```

var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales {
        on loc {
            var k: int;
            k = 2*i + j;
        }
    }
}

```

here, *i* and *j* are remote, so
the compiler + runtime will
transfer their values

`k = 2*i + j;`



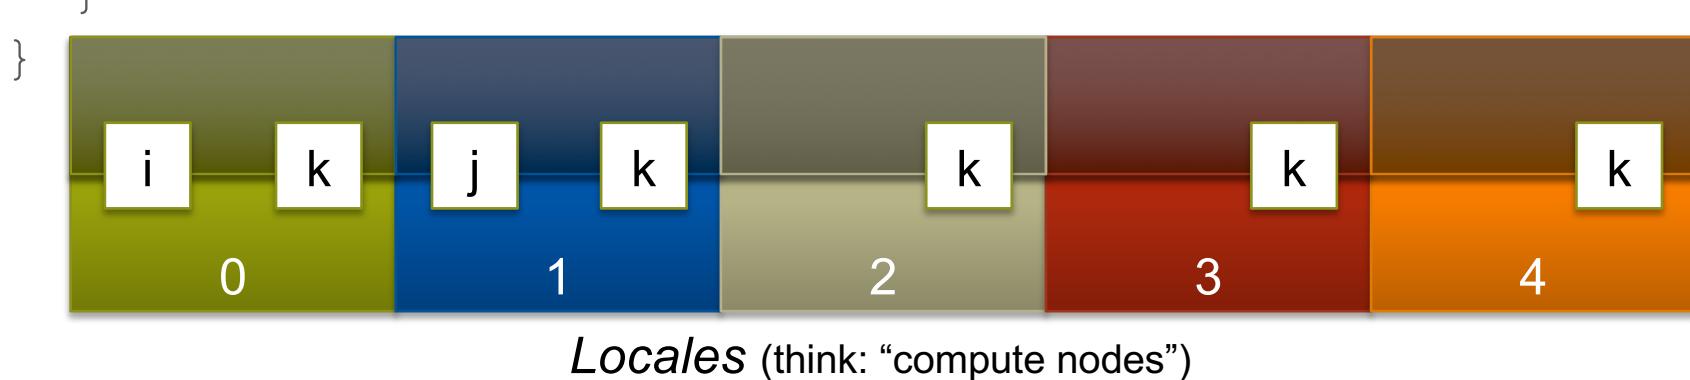
Chapel: Locality queries

```

var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales {
        on loc {
            var k: int;

                ...here...          // query the locale on which this task is running
                ...j.locale...      // query the locale on which j is stored
        }
    }
}

```



Chapel: Locality queries

```

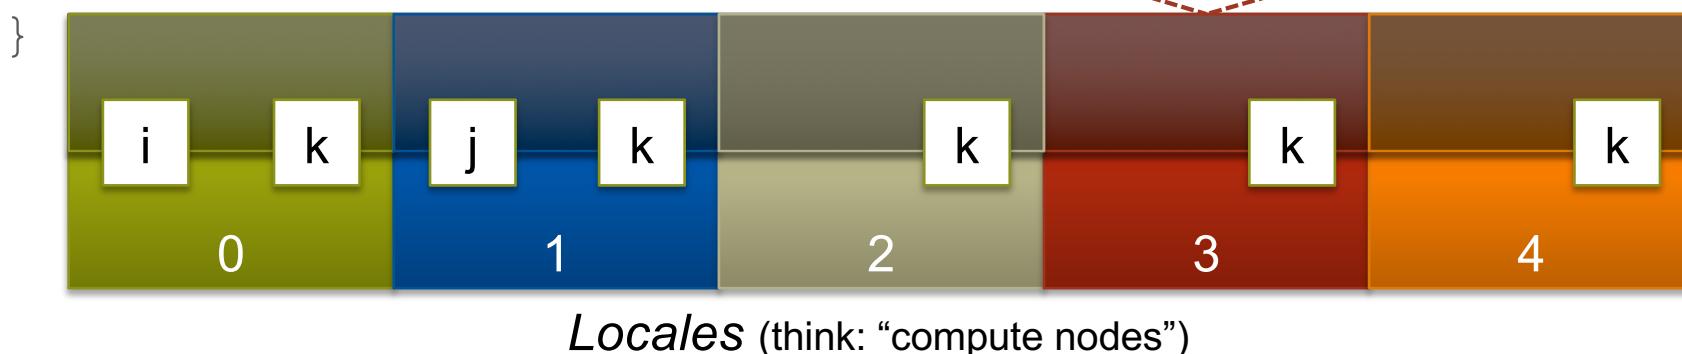
var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales {
        on loc {
            var k: int;
        }
    }
}

```

Recall: scalars have ‘const in’ task intents by default

...here... // query the locale on which this task is running
...j.locale... // query the locale on which i is stored

here.id = 3
j.locale.id = 3



Locales (think: “compute nodes”)



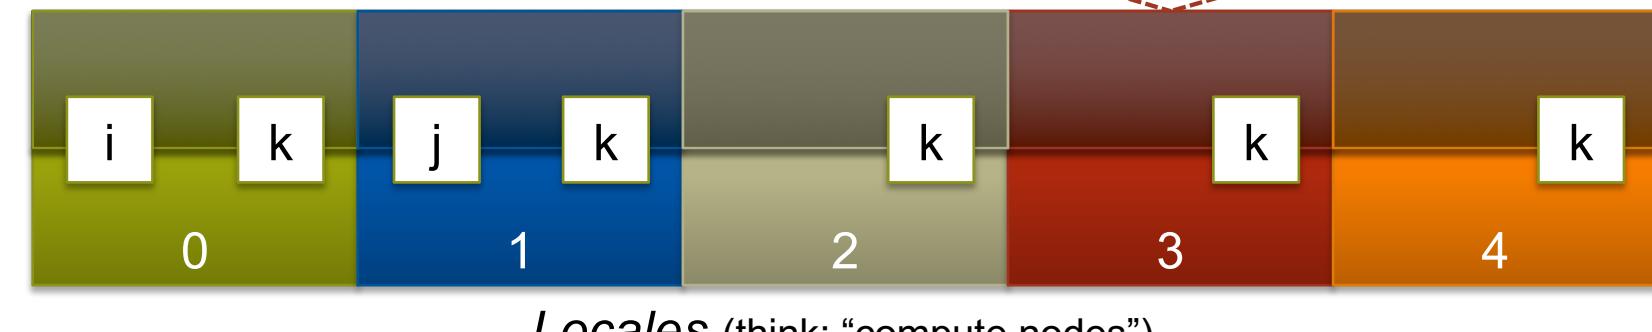
Chapel: Locality queries

```

var i: int;
on Locales[1] {
    var j: int;
    coforall loc in Locales with (ref i, ref j) {
        on loc {
            var k: int;

            ...here...          // query the locale on which this task is running
            ...j.locale...      // query the locale on which i is stored
        }
    }
}

```



Querying a Variable's Locale

- **Syntax**

```
locale-query-expr:  
    expr . locale
```

- **Semantics**

- Returns the locale on which *expr* is stored

- **Example**

```
var i: int;  
on Locales[1] {  
    var j: int;  
    writeln((i.locale.id, j.locale.id)); // outputs (0,1)  
}
```



Here

- **Built-in locale variable**

```
const here: locale;
```

- **Semantics**

- Refers to the locale on which the task is executing

- **Example**

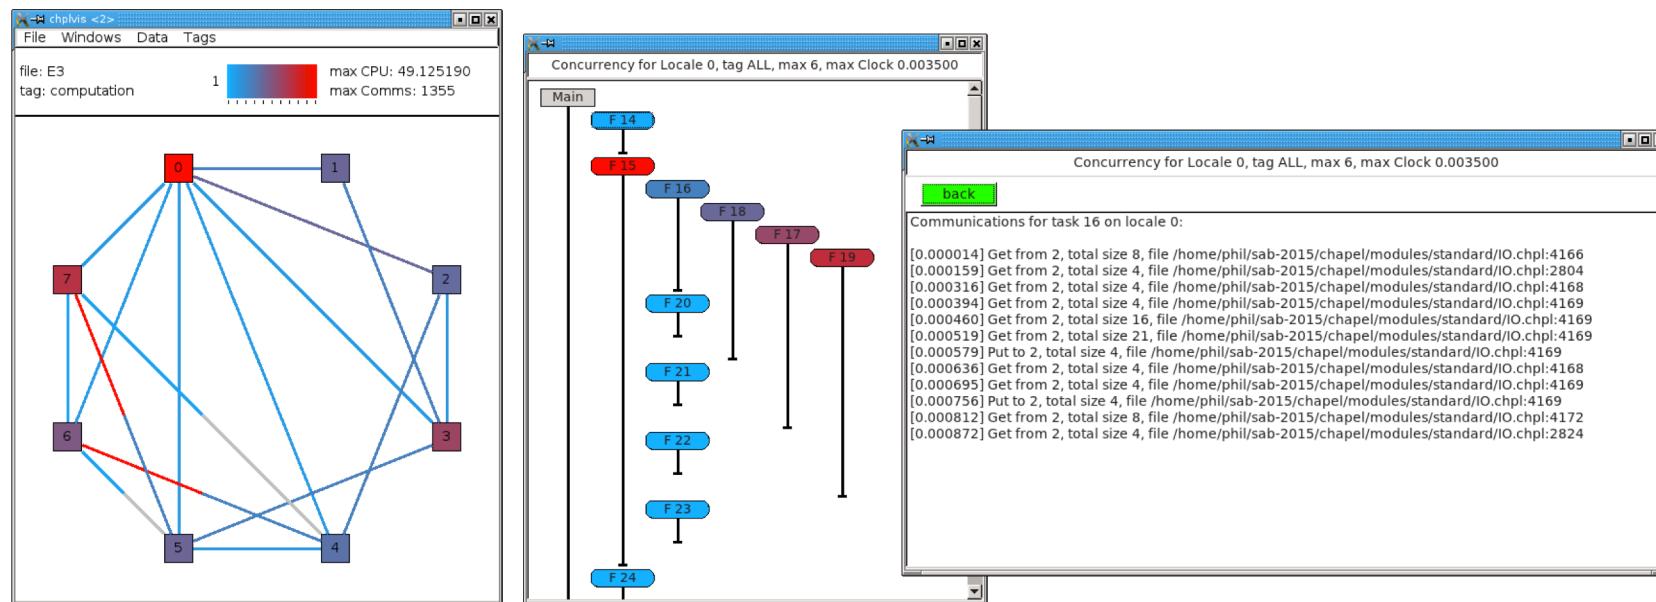
```
writeln(here.id);      // outputs 0
on Locales[1] do
  writeln(here.id);  // outputs 1

on myC do
  if (here == Locales[0]) then ...
```



Reasoning about Communication

- Though implicit, users can reason about communication
 - semantic model is explicit about where data is placed / tasks execute
 - execution-time queries support reasoning about locality
 - e.g., `here`, `x.locale`
 - tools should also play a role here
 - e.g., [CommDiagnostics](#) module which can count communications
 - e.g., `chplvis`, contained in the release (developed by Phil Nelson, WWU)



Rearranging Locales

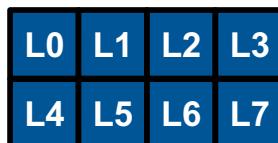
Create locale views with standard array operations:

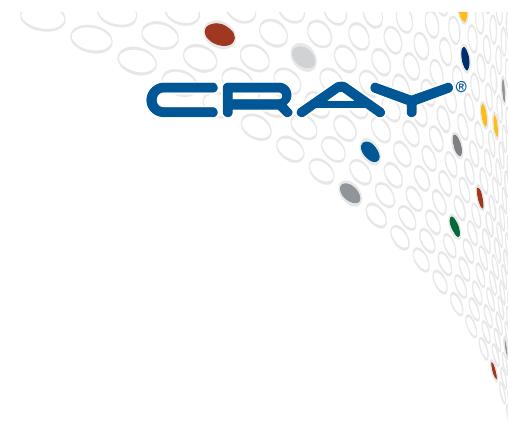
```
var TaskALocs = Locales[0..1];  
var TaskBLocs = Locales[2..];  
  
var Grid2D = reshape(Locales, {1..2, 1..4});
```

Locales: 

TaskALocs: 

TaskBLocs: 

Grid2D: 



Questions about (low-level) locality in Chapel?



COMPUTE

|

STORE

|

ANALYZE

Copyright 2018 Cray Inc.



Legal Disclaimer

Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.

Cray Inc. may make changes to specifications and product descriptions at any time, without notice.

All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.

Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.

Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.

The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, and URIKA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.

